

---

# Corner Detection

CS 543 / ECE 549 – Saurabh Gupta

# Why extract keypoints?

---

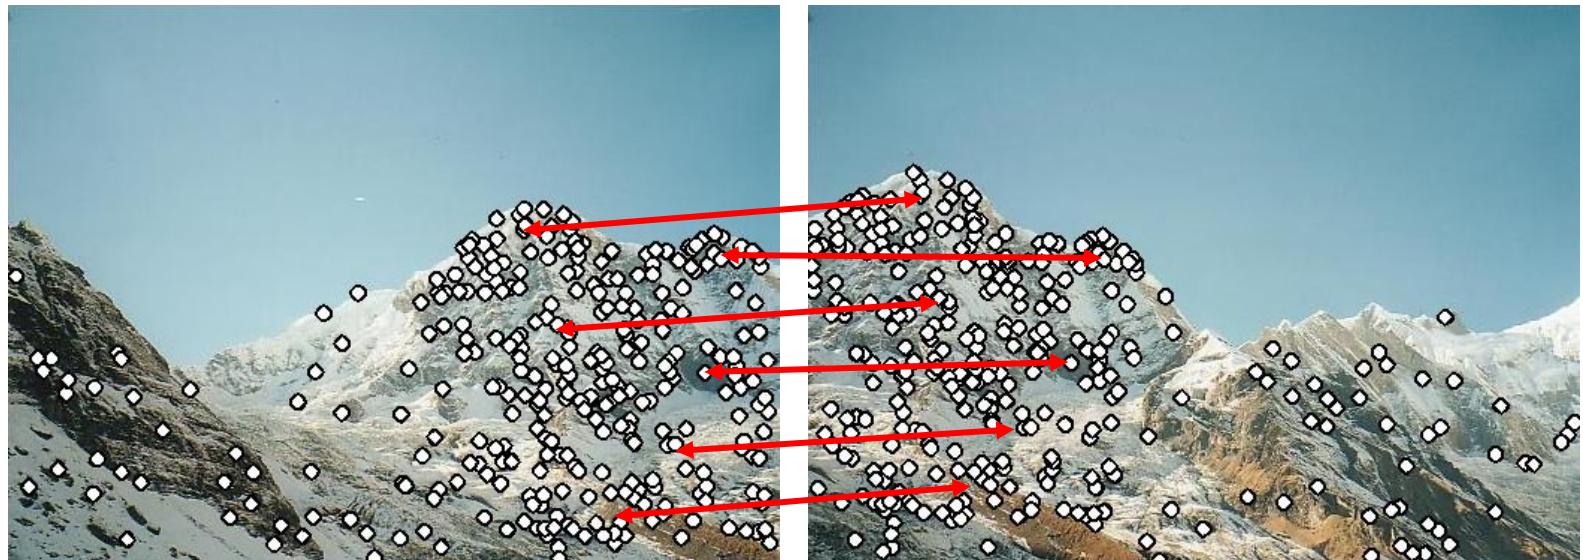
- Motivation: panorama stitching
  - We have two images – how do we combine them?



# Why extract keypoints?

---

- Motivation: panorama stitching
  - We have two images – how do we combine them?



Step 1: extract keypoints

Step 2: match keypoint features

# Why extract keypoints?

---

- Motivation: panorama stitching
  - We have two images – how do we combine them?



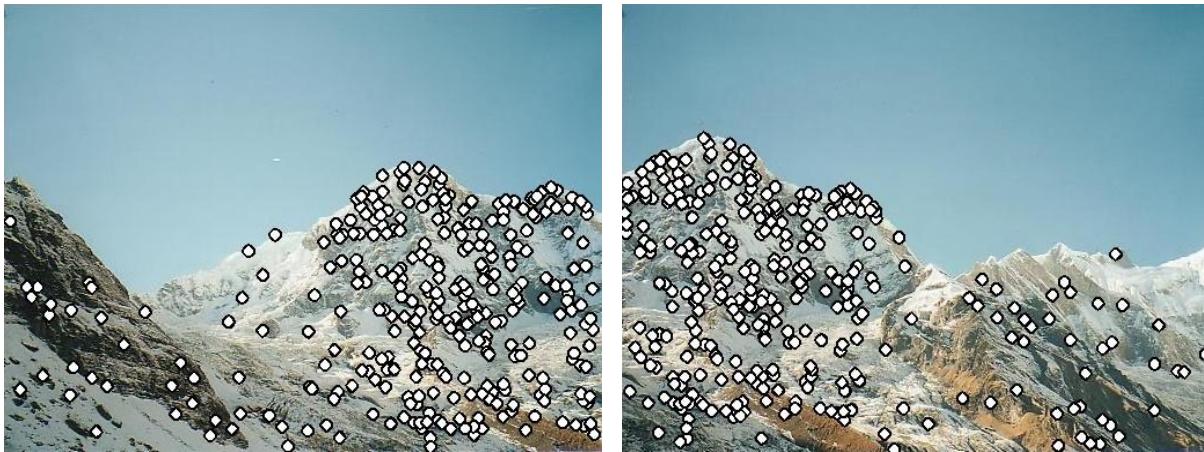
Step 1: extract keypoints

Step 2: match keypoint features

Step 3: align images

# Characteristics of good keypoints

---



- **Compactness and efficiency**
  - Many fewer keypoints than image pixels
- **Saliency**
  - Each keypoint is distinctive
- **Locality**
  - A keypoint occupies a relatively small area of the image; robust to clutter and occlusion
- **Repeatability**
  - The same keypoint can be found in several images despite geometric and photometric transformations

# Applications

---

Keypoints are used for:

- Image alignment
- 3D reconstruction
- Motion tracking
- Robot navigation
- Database indexing and retrieval
- Object recognition



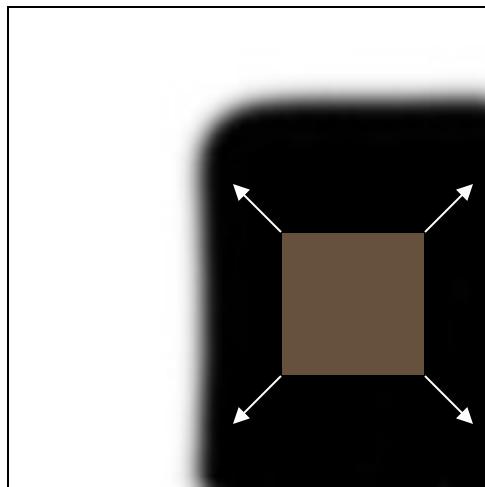
# Corner detection: Basic idea

---

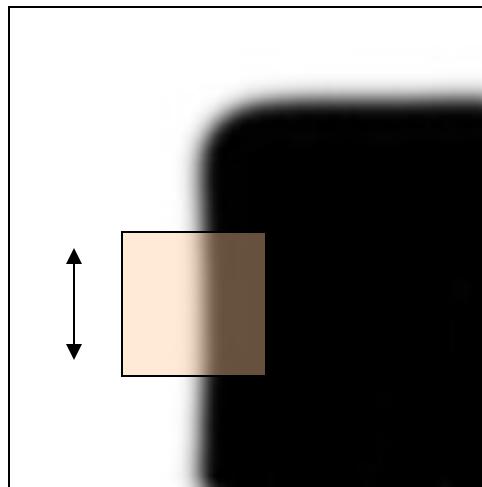
# Corner detection: Basic idea

---

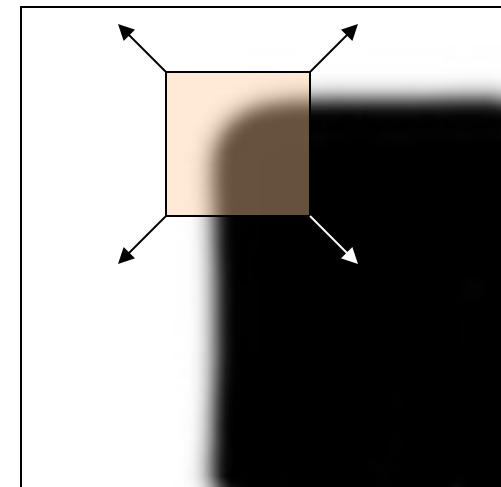
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give a *large change* in intensity



“flat” region:  
no change in  
all directions



“edge”:  
no change  
along the edge  
direction



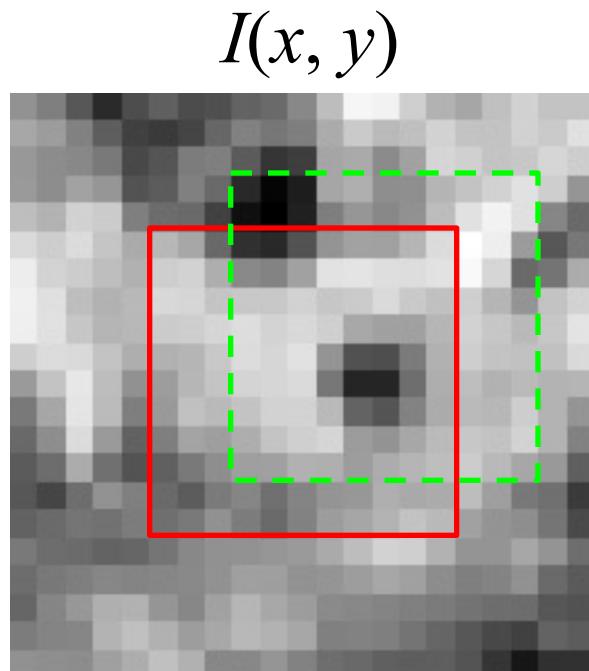
“corner”:  
significant  
change in all  
directions

# Corner Detection: Derivation

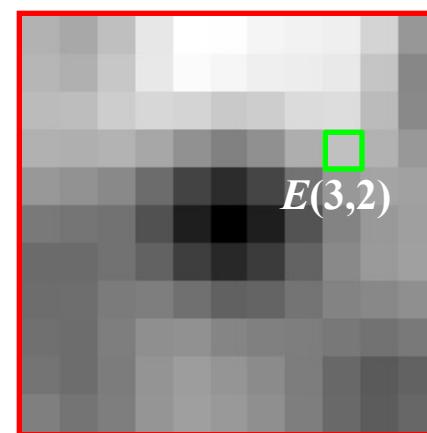
---

Change in appearance of window  $W$  for the shift  $[u, v]$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$



$$E(u, v)$$



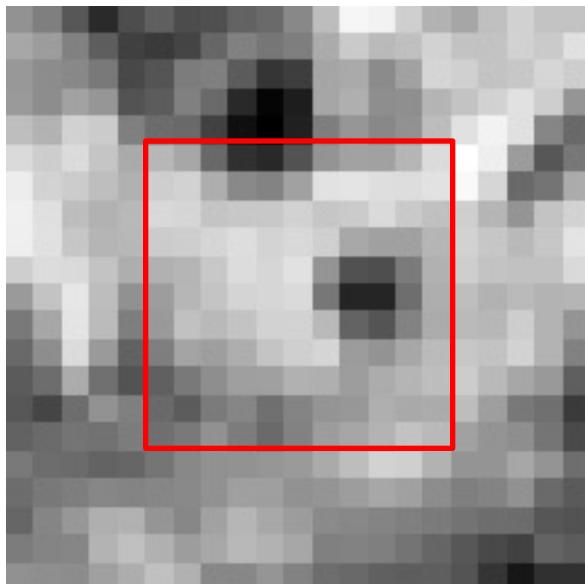
# Corner Detection: Derivation

---

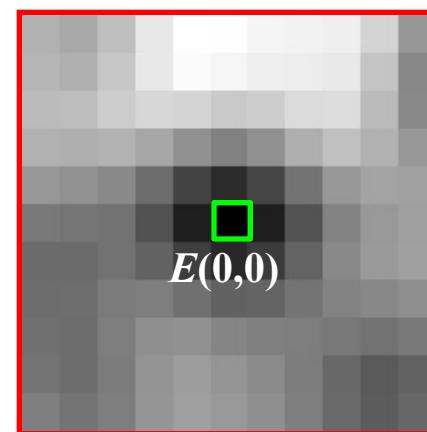
Change in appearance of window  $W$  for the shift  $[u, v]$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

$$I(x, y)$$



$$E(u, v)$$



# Corner Detection: Derivation

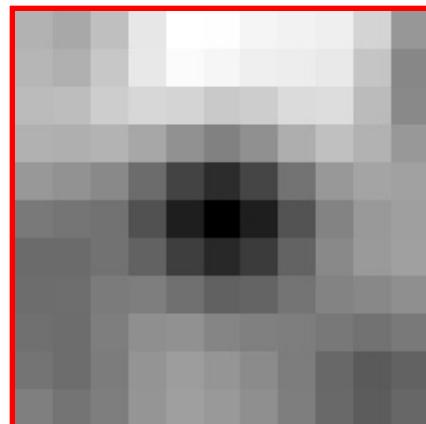
---

Change in appearance of window  $W$  for the shift  $[u, v]$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

$$E(u, v)$$



# Corner Detection: Derivation

---

First-order Taylor approximation for small motions  $[u, v]$ :

$$I(x + u, y + v) \approx I(x, y) + I_x u + I_y v$$

Let's plug this into  $E(u, v)$ :

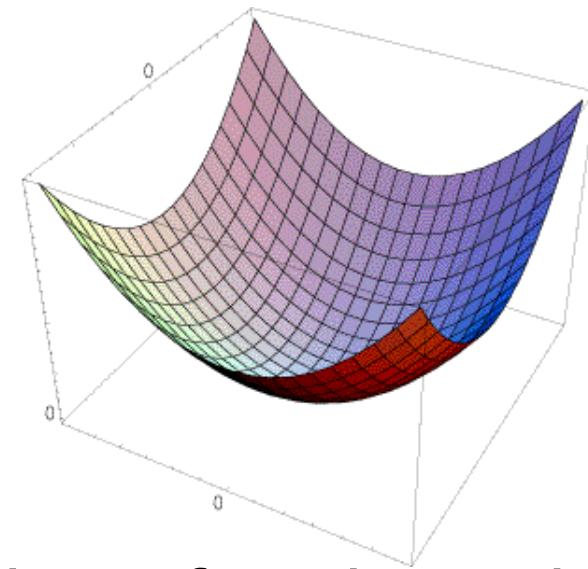
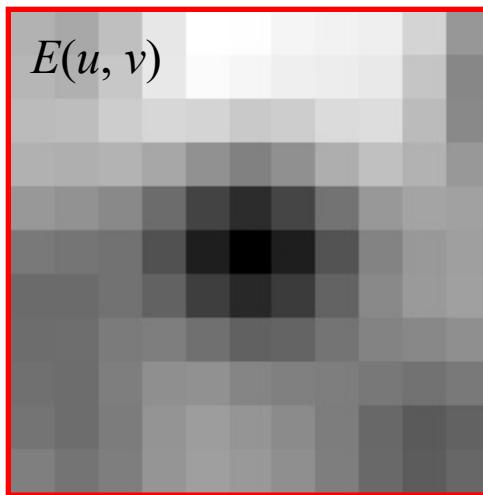
$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

# Corner Detection: Derivation

---

$E(u,v)$  can be locally approximated by a quadratic surface:

$$E(u,v) \approx u^2 \sum_{x,y} I_x^2 + 2uv \sum_{x,y} I_x I_y + v^2 \sum_{x,y} I_y^2$$



In which directions does this surface have the fastest/slowest change?

# Corner Detection: Derivation

---

$E(u,v)$  can be locally approximated by a quadratic surface:

$$E(u,v) \approx u^2 \sum_{x,y} I_x^2 + 2uv \sum_{x,y} I_x I_y + v^2 \sum_{x,y} I_y^2$$
$$= \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

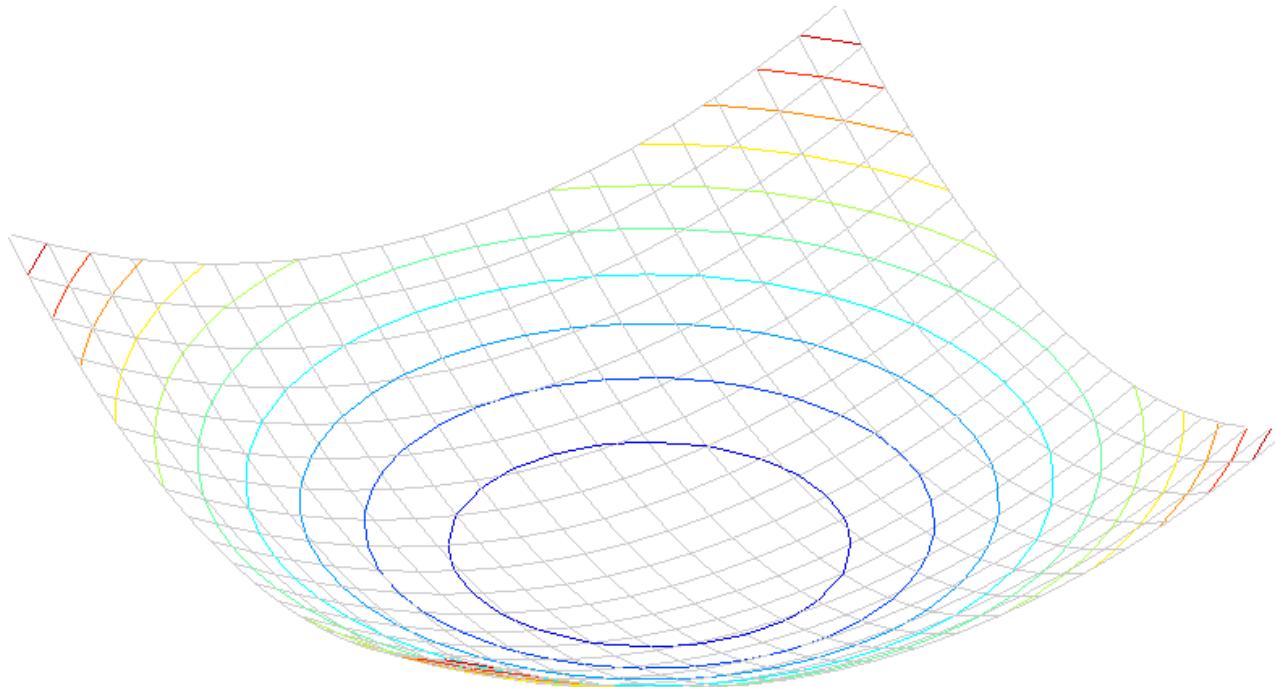
*Second moment matrix  $M$*

# Interpreting the second moment matrix

---

A horizontal “slice” of  $E(u, v)$  is given by the equation of an ellipse:

$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

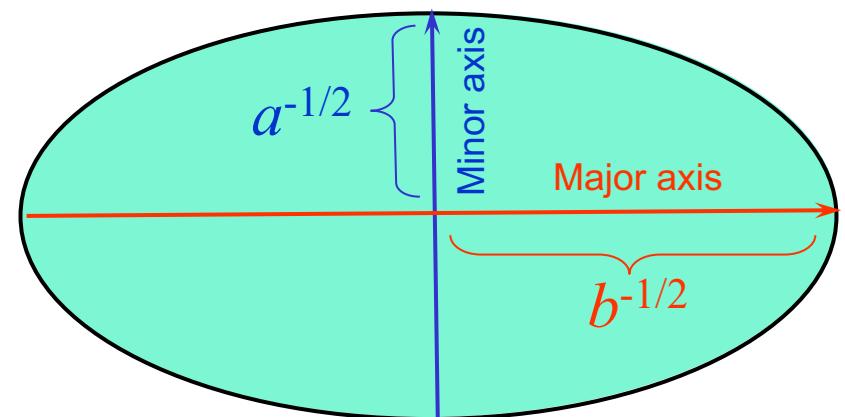


# Interpreting the second moment matrix

Consider the axis-aligned case (gradients are either horizontal or vertical):

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix}$$

$$\begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = 1$$



# Interpreting the second moment matrix

---

Consider the axis-aligned case (gradients are either horizontal or vertical):

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

If either  $a$  or  $b$  is close to 0, then this is **not** a corner, so we want locations where both are large

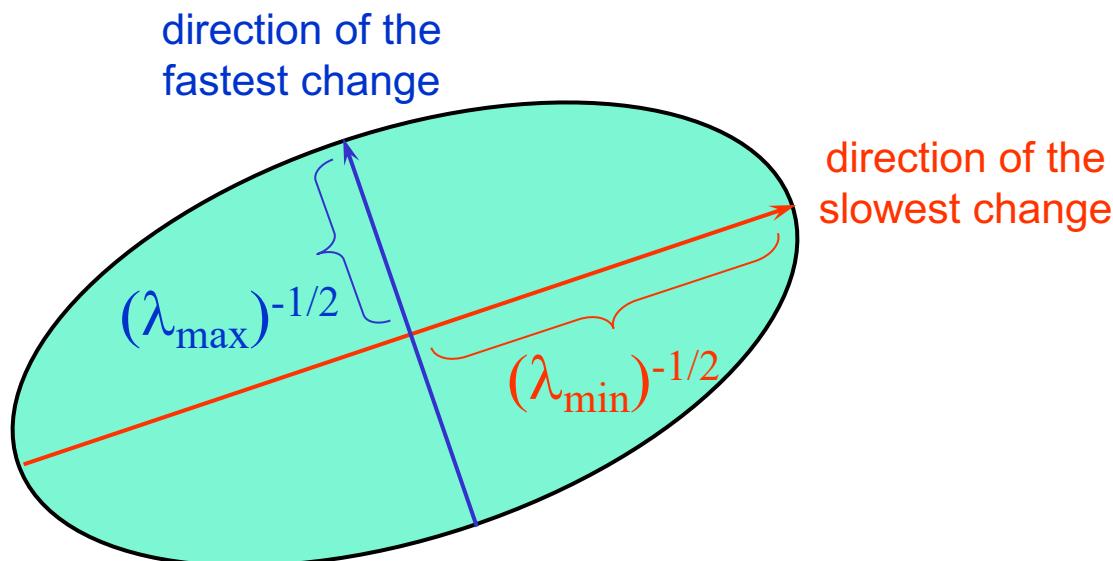
# Interpreting the second moment matrix

---

In the general case, need to *diagonalize M*:

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by  $R$ :



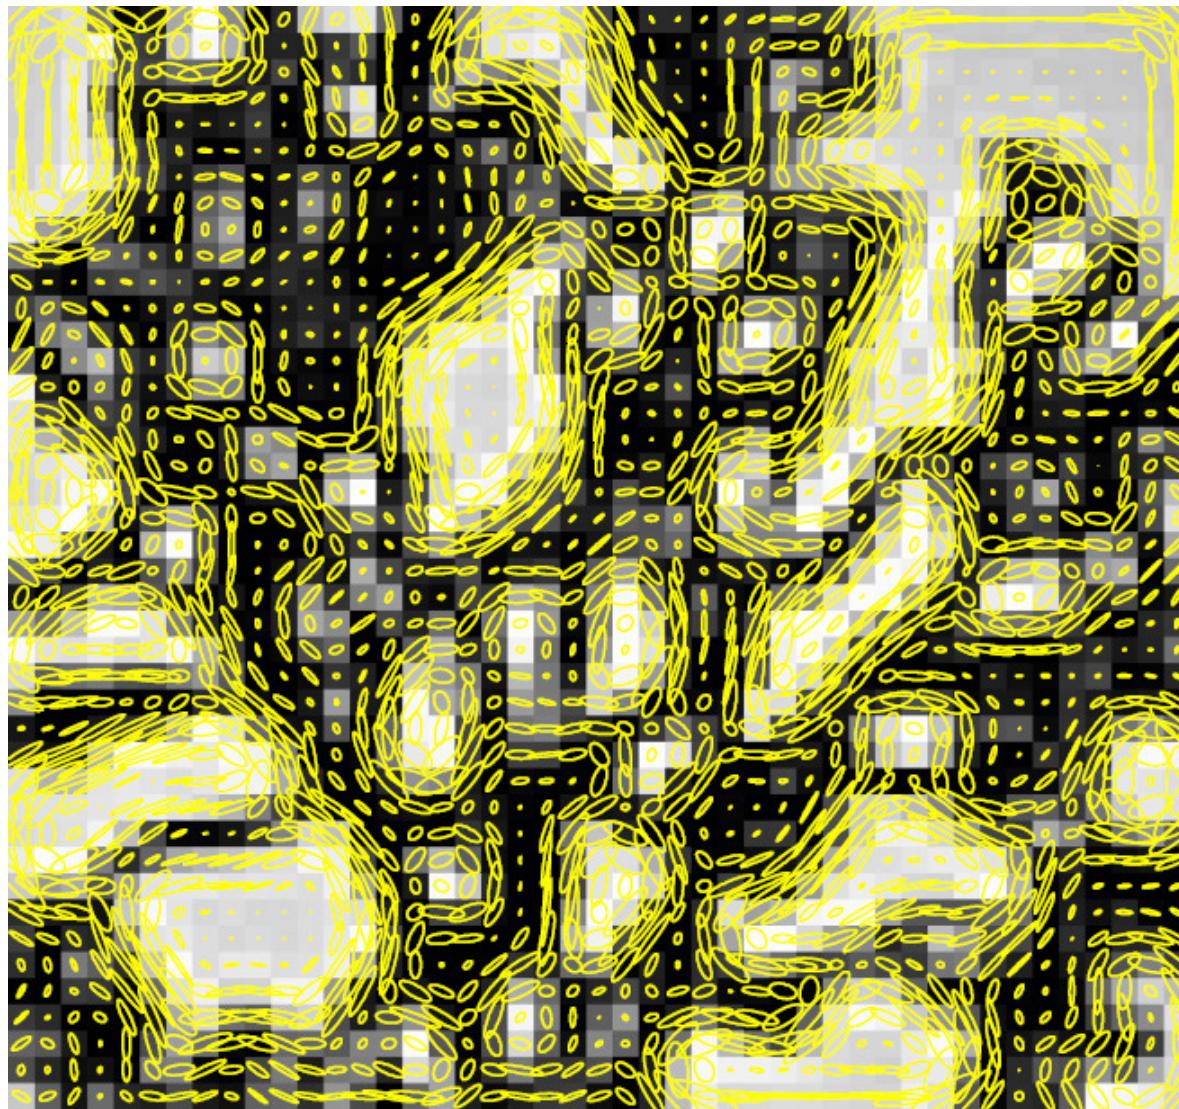
# Visualization of second moment matrices

---



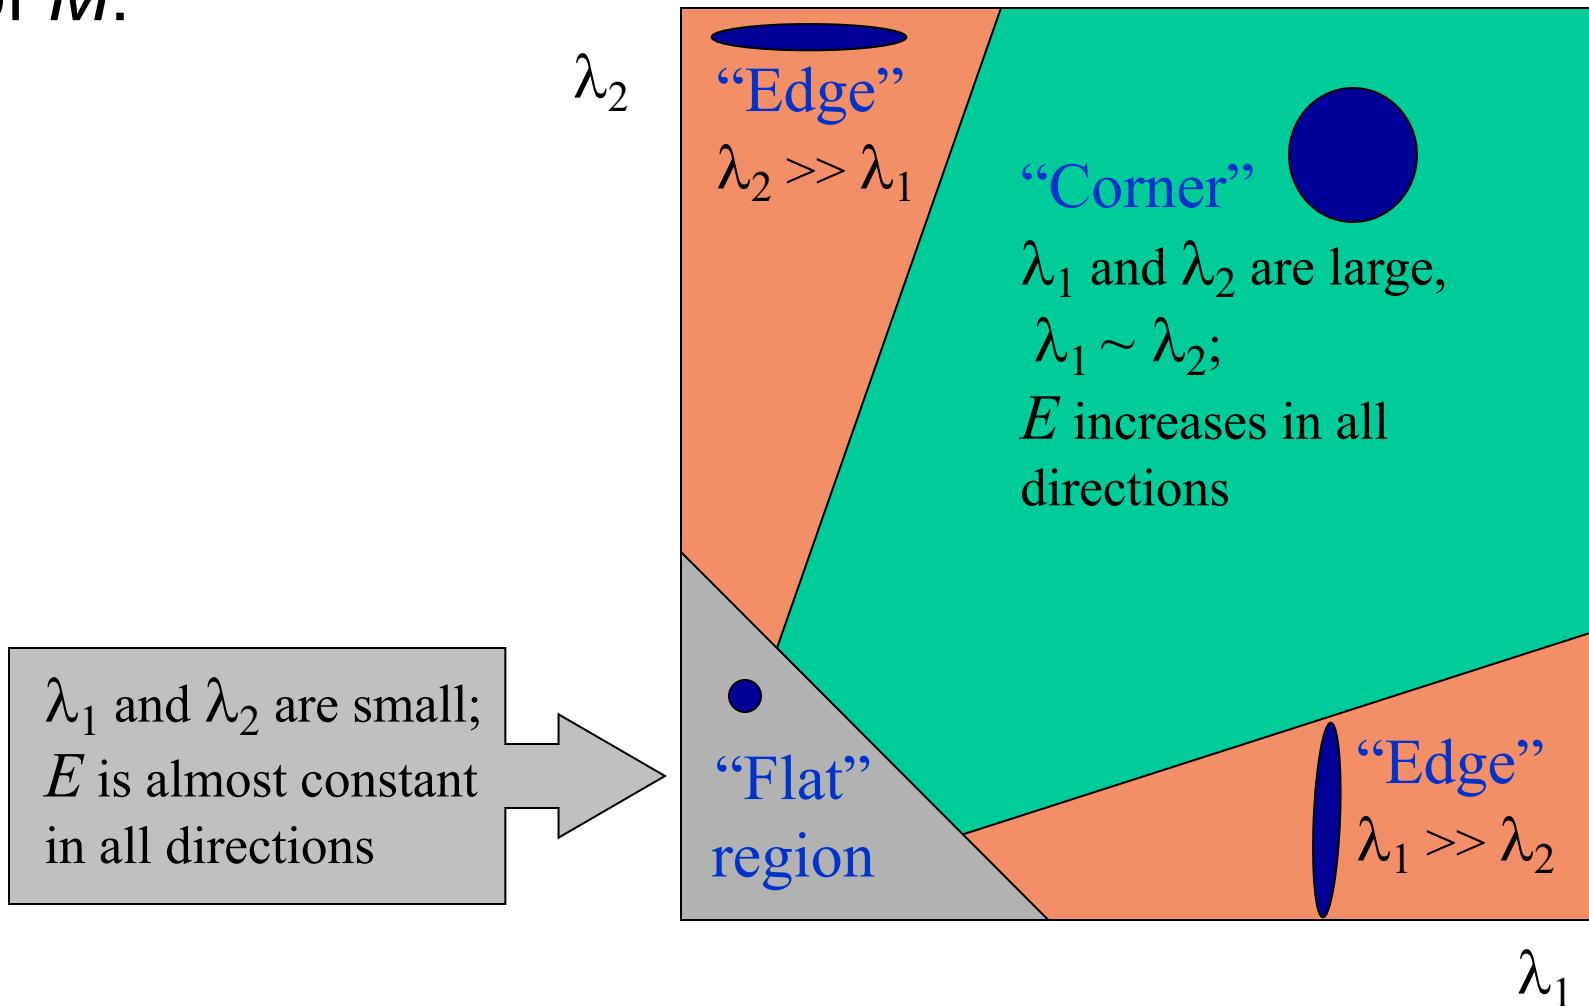
# Visualization of second moment matrices

---



# Interpreting the eigenvalues

Classification of image points using eigenvalues of  $M$ :

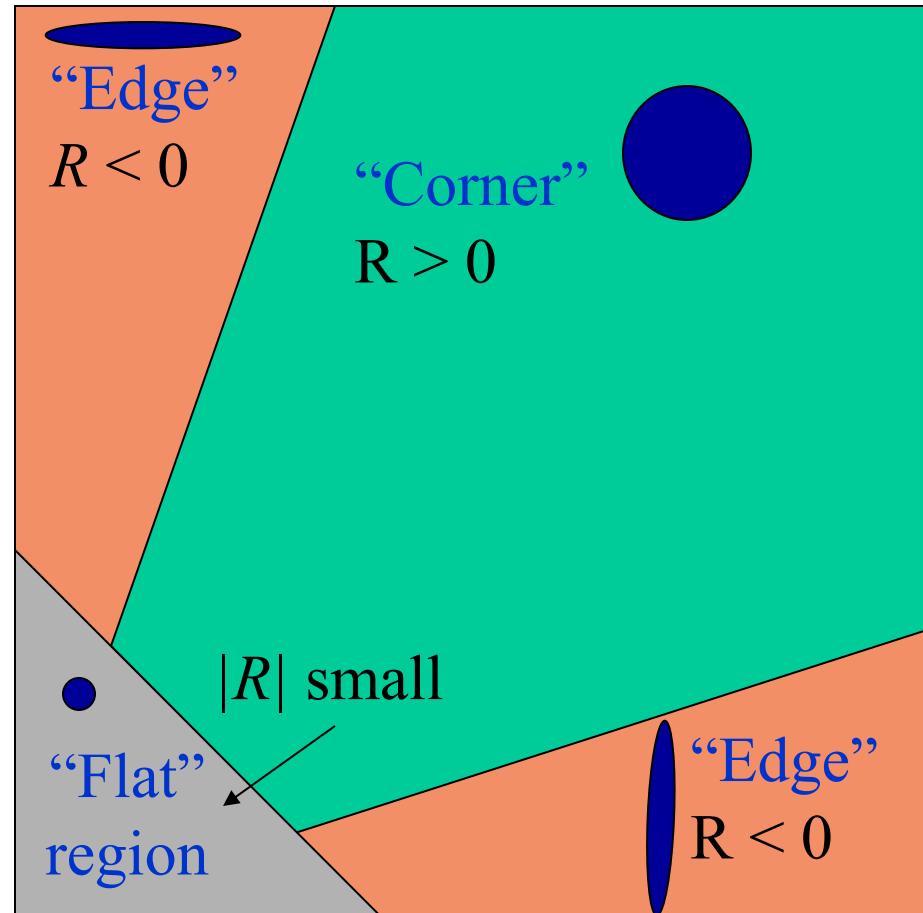


# Corner response function

---

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : constant (0.04 to 0.06)



# The Harris corner detector

---

1. Compute partial derivatives at each pixel
2. Compute second moment matrix  $M$  in a Gaussian window around each pixel:

$$M = \begin{bmatrix} \sum_{x,y} w(x,y) I_x^2 & \sum_{x,y} w(x,y) I_x I_y \\ \sum_{x,y} w(x,y) I_x I_y & \sum_{x,y} w(x,y) I_y^2 \end{bmatrix}$$

C.Harris and M.Stephens, [A Combined Corner and Edge Detector](#),  
*Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.  
Source: L. Lazebnik

# The Harris corner detector

---

1. Compute partial derivatives at each pixel
2. Compute second moment matrix  $M$  in a Gaussian window around each pixel
3. Compute corner response function  $R$

C.Harris and M.Stephens, [A Combined Corner and Edge Detector](#),

*Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.

# Harris Detector: Steps

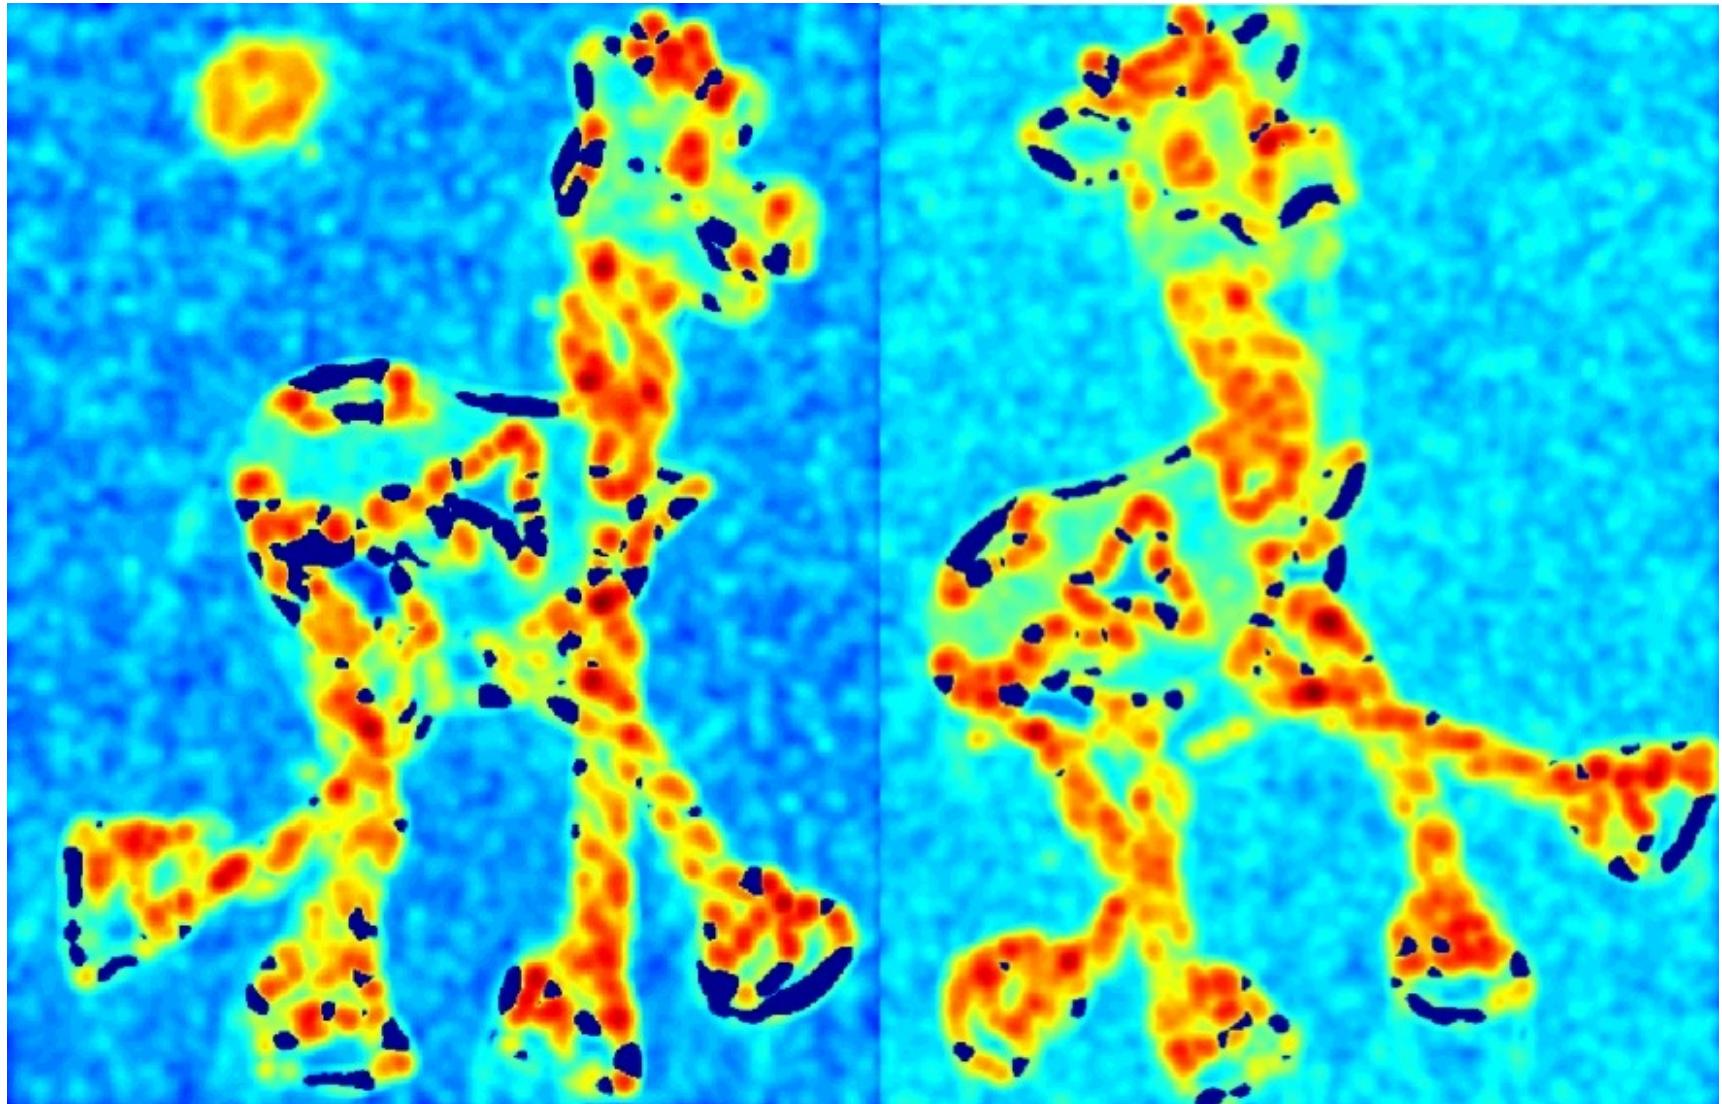
---



# Harris Detector: Steps

---

Compute corner response  $R$



# The Harris corner detector

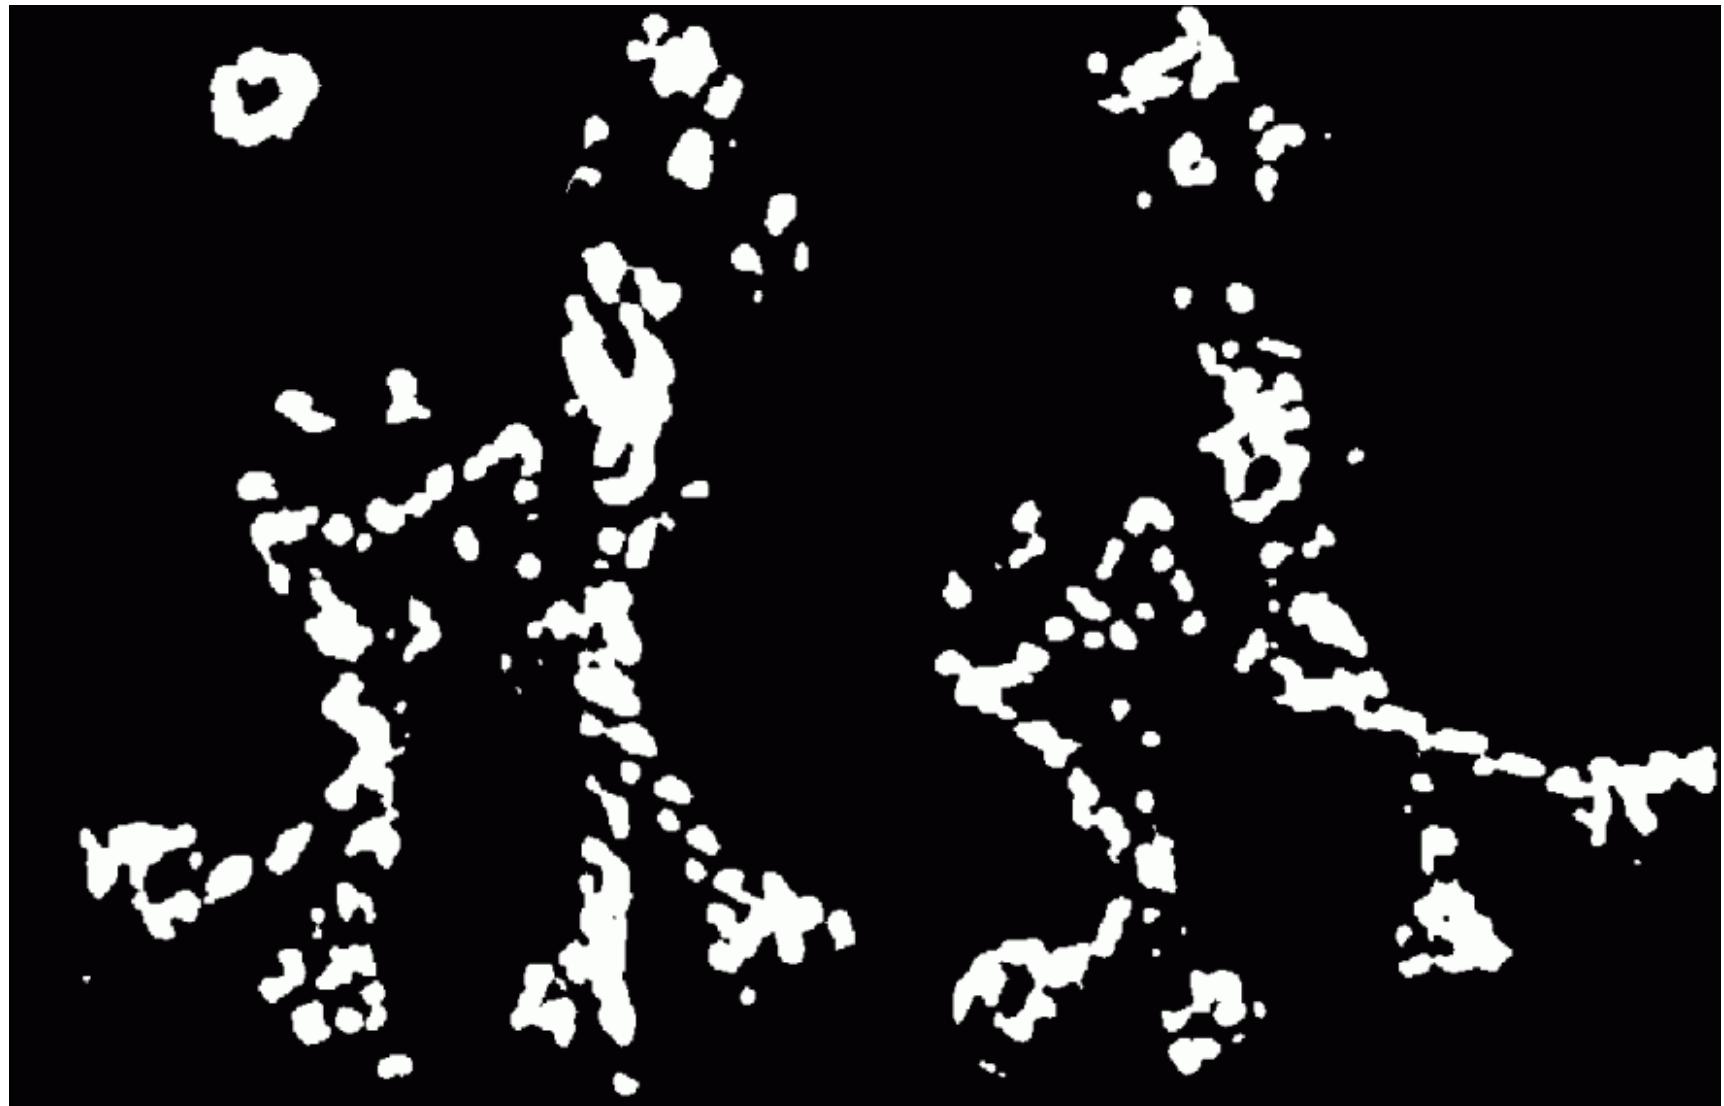
---

1. Compute partial derivatives at each pixel
2. Compute second moment matrix  $M$  in a Gaussian window around each pixel
3. Compute corner response function  $R$
4. Threshold  $R$
5. Find local maxima of response function  
(nonmaximum suppression)

# Harris Detector: Steps

---

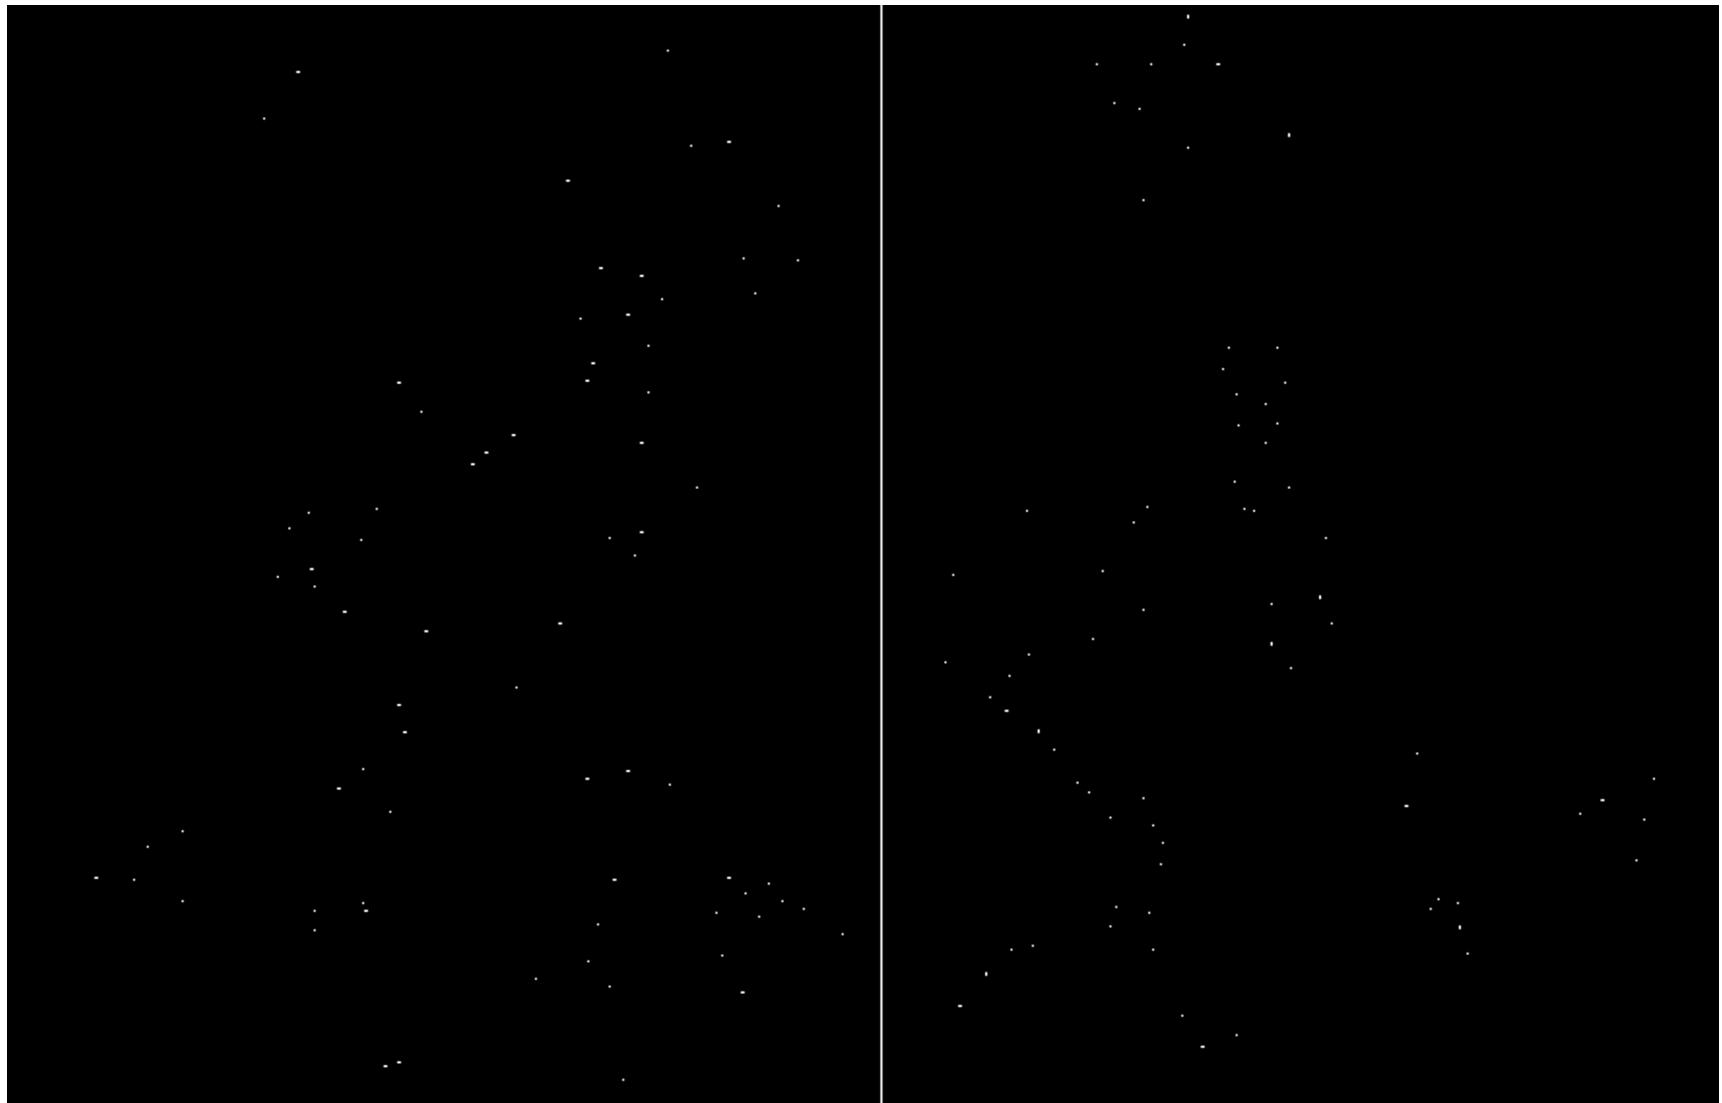
Find points with large corner response:  $R > \text{threshold}$



# Harris Detector: Steps

---

Take only the points of local maxima of  $R$



# Harris Detector: Steps

---



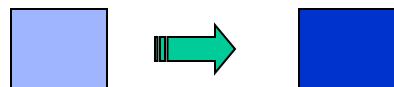
# Robustness of corner features

---

- What happens to corner features when the image undergoes geometric or photometric transformations?

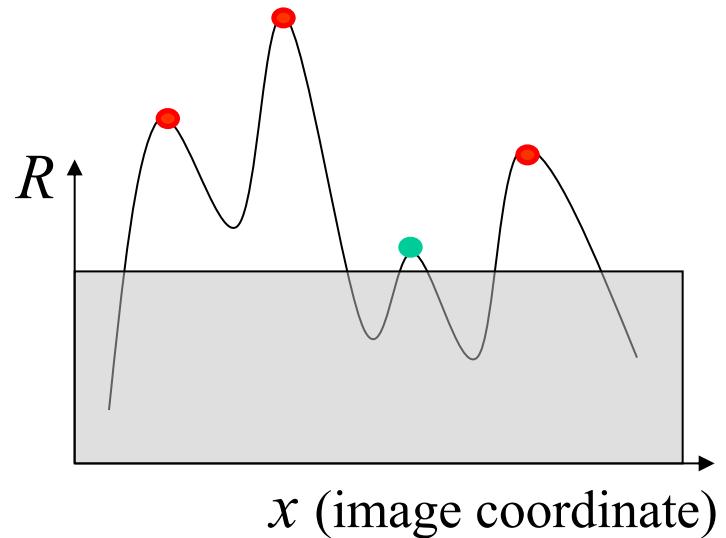
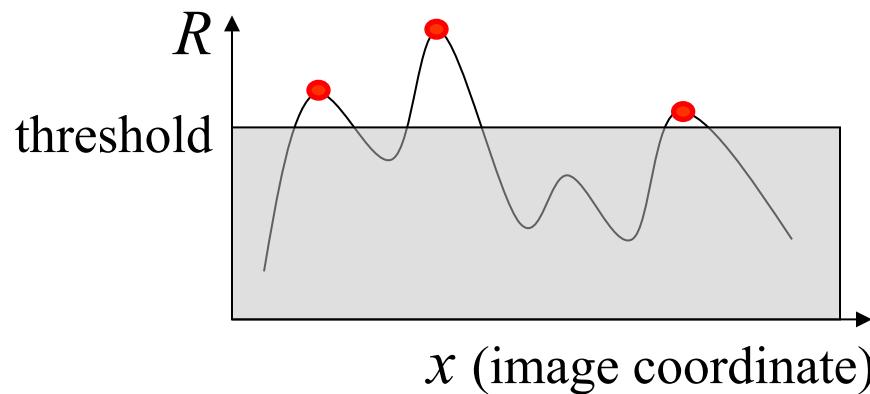


# Affine intensity change



$$I \rightarrow a I + b$$

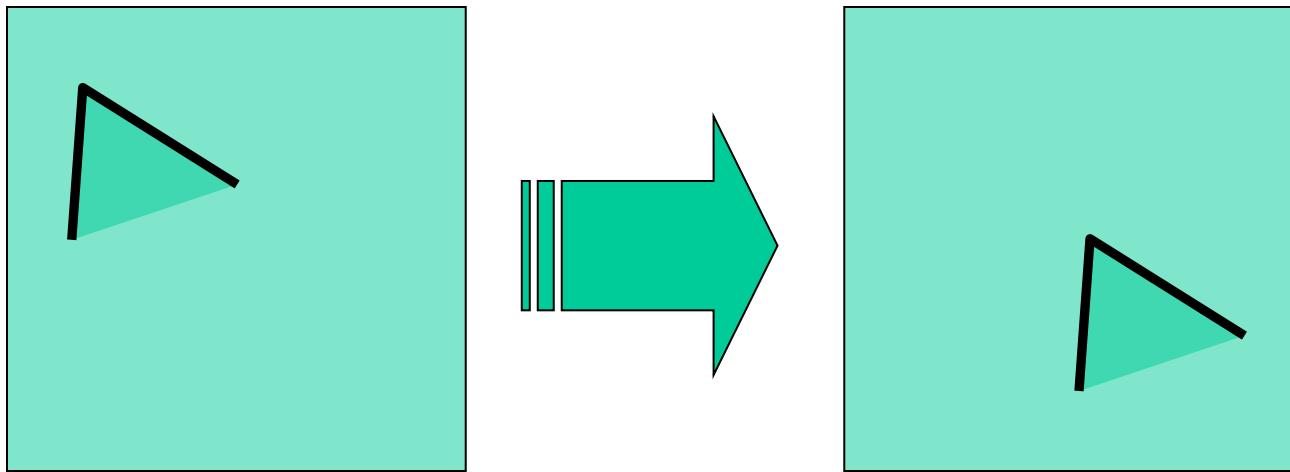
- Only derivatives are used, so invariant to intensity shift  $I \rightarrow I + b$
- Intensity scaling:  $I \rightarrow a I$



*Partially invariant to affine intensity change*

# Image translation

---

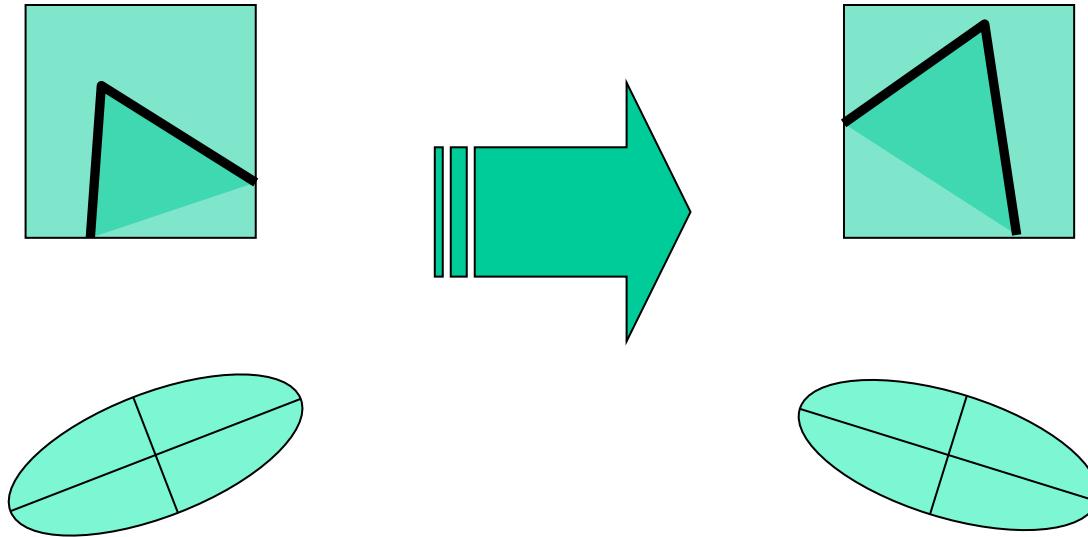


- Derivatives and window function are shift-invariant

Corner location is *covariant* w.r.t. translation

# Image rotation

---

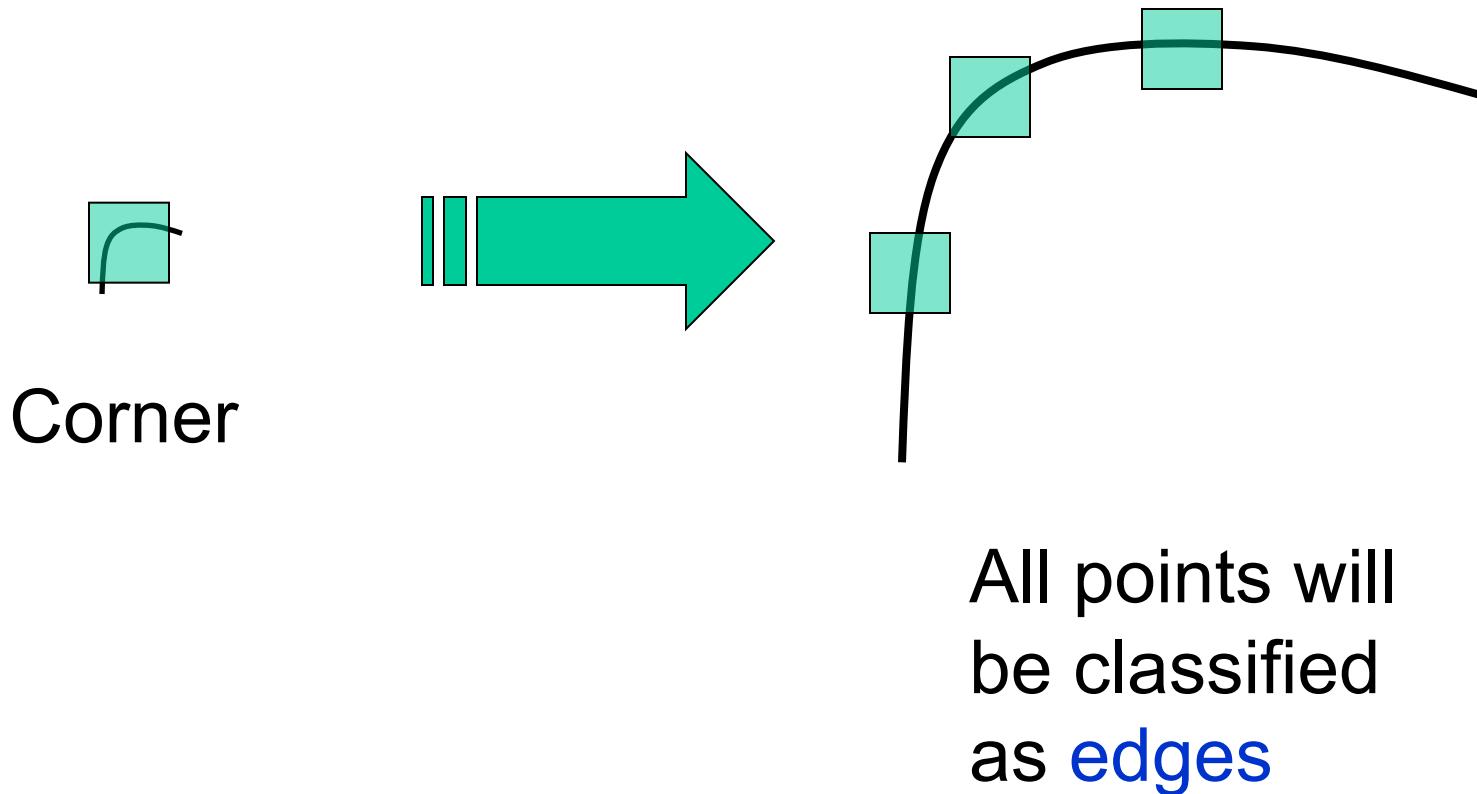


Second moment ellipse rotates but its shape  
(i.e. eigenvalues) remains the same

Corner location is covariant w.r.t. rotation

# Scaling

---



Corner location is not covariant w.r.t. scaling!

# Optical flow

---



Many slides adapted from S. Seitz, R. Szeliski, M. Pollefeys

**Slides from S. Lazebnik.**

# What direction is the object moving?

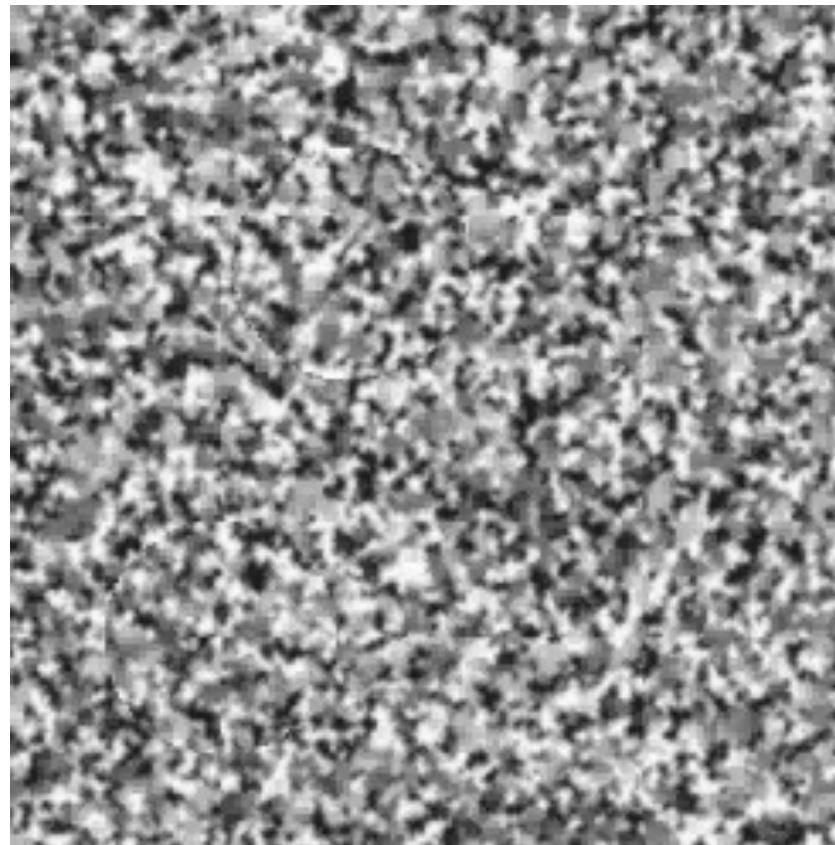
---

- A. Left right
- B. Up down
- C. Top-left to bottom-right
- D. Bottom left to top-right

# Motion is a powerful perceptual cue

---

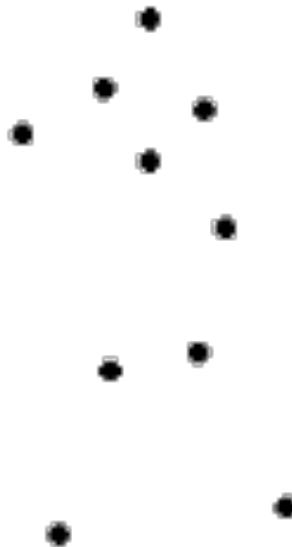
- Sometimes, it is the only cue



# Motion is a powerful perceptual cue

---

- Even “impoverished” motion data can evoke a strong percept

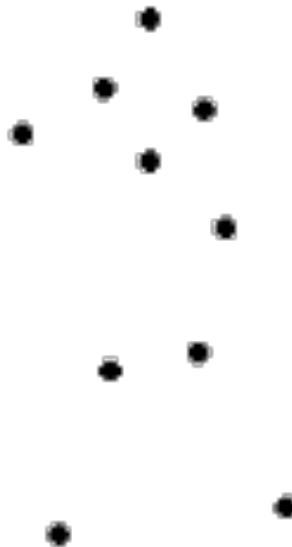


G. Johansson, "Visual Perception of Biological Motion and a Model For Its Analysis",  
*Perception and Psychophysics 14, 201-211, 1973.*

# Motion is a powerful perceptual cue

---

- Even “impoverished” motion data can evoke a strong percept



G. Johansson, "Visual Perception of Biological Motion and a Model For Its Analysis",  
*Perception and Psychophysics 14*, 201-211, 1973.

# Uses of motion in computer vision

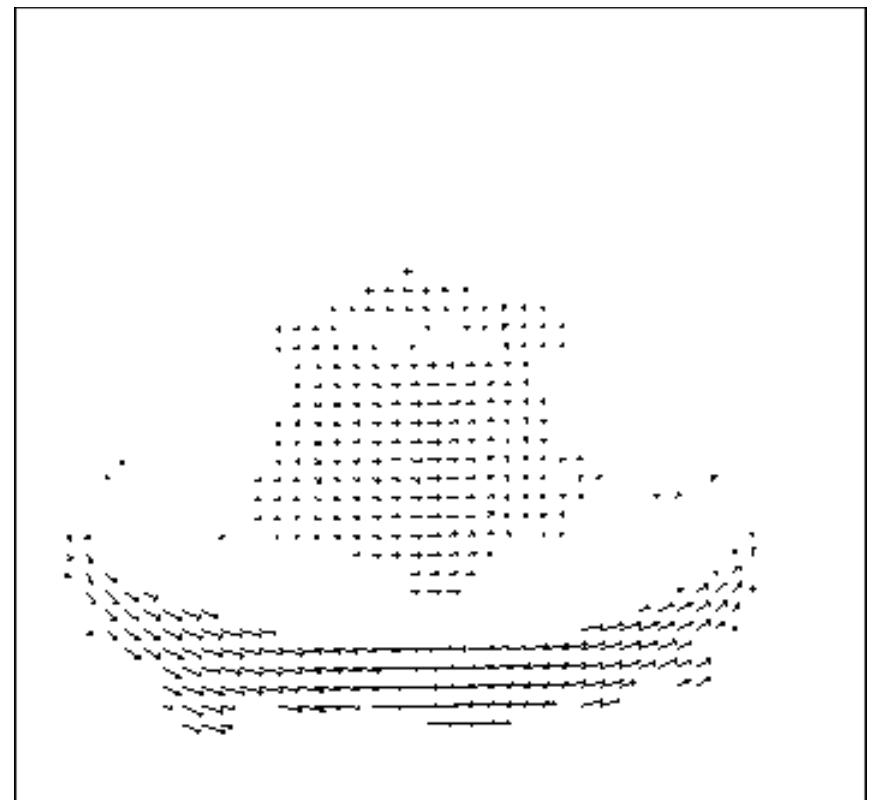
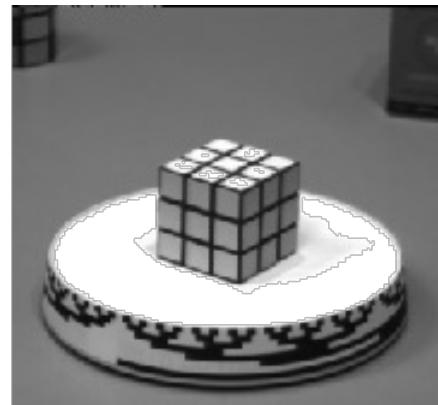
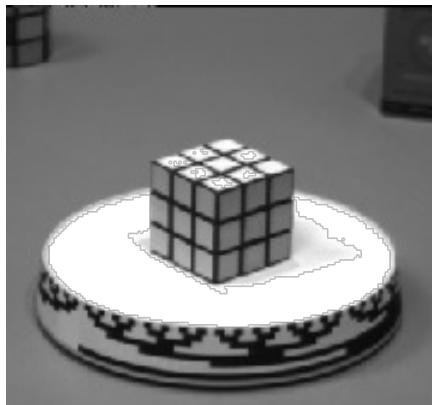
---

- 3D shape reconstruction
- Object segmentation
- Learning and tracking of dynamical models
- Event and activity recognition
- Self-supervised and predictive learning

# Motion field

---

- The motion field is the projection of the 3D scene motion into the image



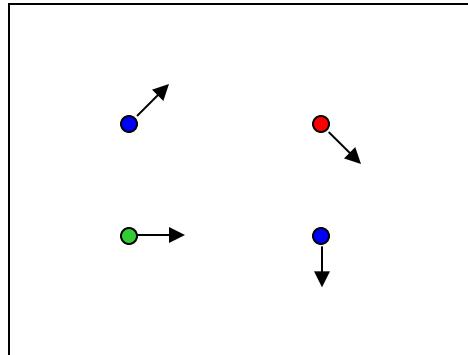
# Optical flow

---

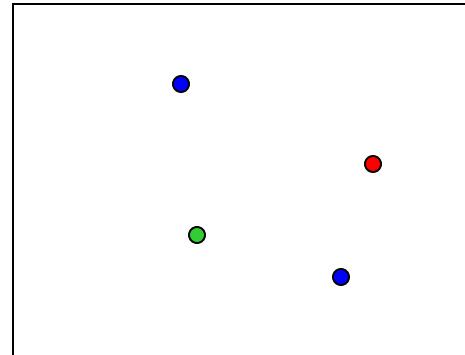
- **Definition:** optical flow is the *apparent* motion of brightness patterns in the image
- Ideally, optical flow would be the same as the motion field
- Have to be careful: apparent motion can be caused by lighting changes without any actual motion
  - Think of a uniform rotating sphere under fixed lighting vs. a stationary sphere under moving illumination

# Estimating optical flow

---



$I(x,y,t-1)$

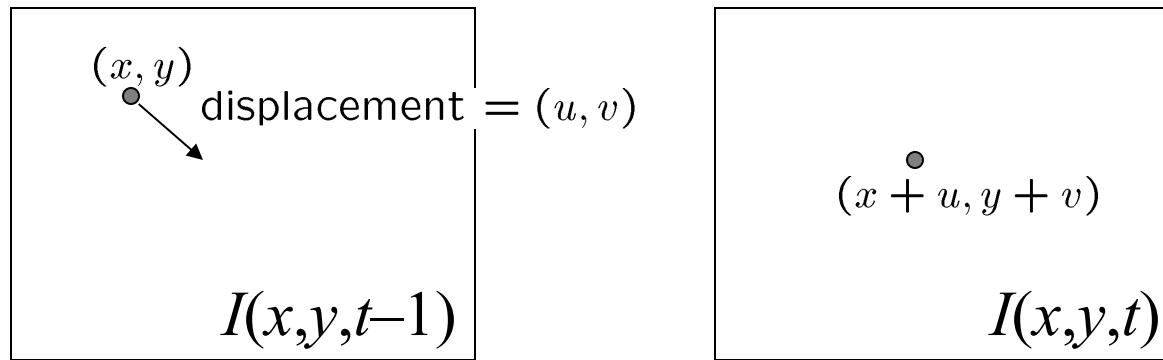


$I(x,y,t)$

- Given two subsequent frames, estimate the apparent motion field  $u(x,y)$  and  $v(x,y)$  between them
- Key assumptions
  - Brightness constancy:** projection of the same point looks the same in every frame
  - Small motion:** points do not move very far
  - Spatial coherence:** points move like their neighbors

# The brightness constancy constraint

---



Brightness Constancy Equation:

$$I(x, y, t - 1) = I(x + u(x, y), y + v(x, y), t)$$

Linearizing the right side using Taylor expansion:

$$I(x, y, t - 1) \approx I(x, y, t) + I_x u(x, y) + I_y v(x, y) + I_t$$

Hence,  $I_x u + I_y v + I_t \approx 0$

# The brightness constancy constraint

---

$$I_x u + I_y v + I_t = 0$$

- How many equations and unknowns per pixel?
  - One equation, two unknowns
- What does this constraint mean?

$$\nabla I \cdot (u, v) + I_t = 0$$

- The component of the flow perpendicular to the gradient (i.e., parallel to the edge) is unknown!

# The brightness constancy constraint

---

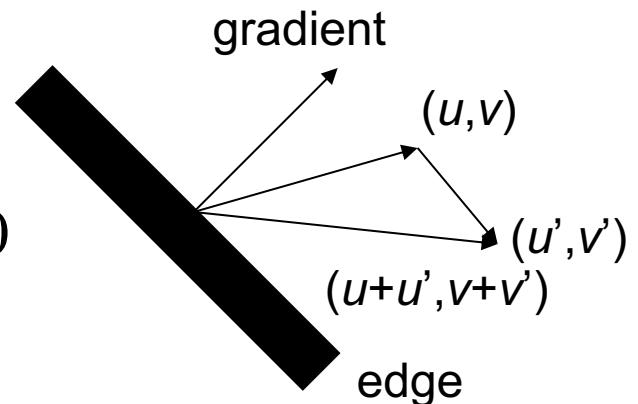
$$I_x u + I_y v + I_t = 0$$

- How many equations and unknowns per pixel?
  - One equation, two unknowns
- What does this constraint mean?

$$\nabla I \cdot (u, v) + I_t = 0$$

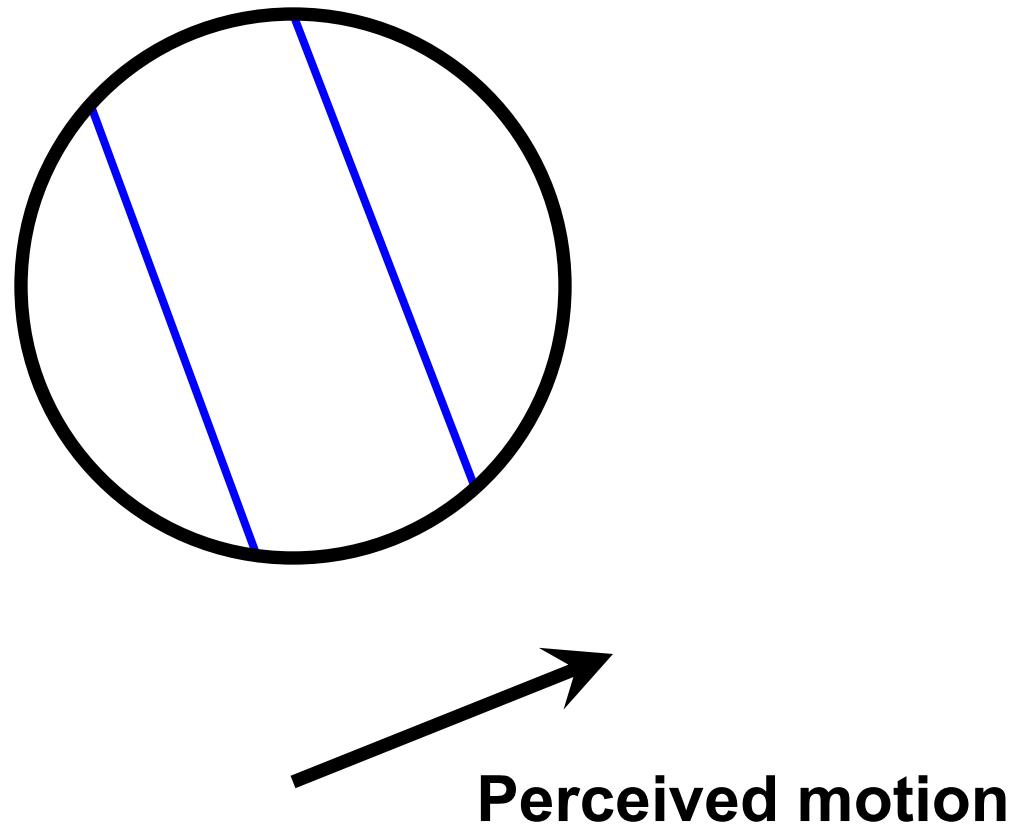
- The component of the flow perpendicular to the gradient (i.e., parallel to the edge) is unknown!

If  $(u, v)$  satisfies the equation,  
so does  $(u+u', v+v')$  if  $\nabla I \cdot (u', v') = 0$



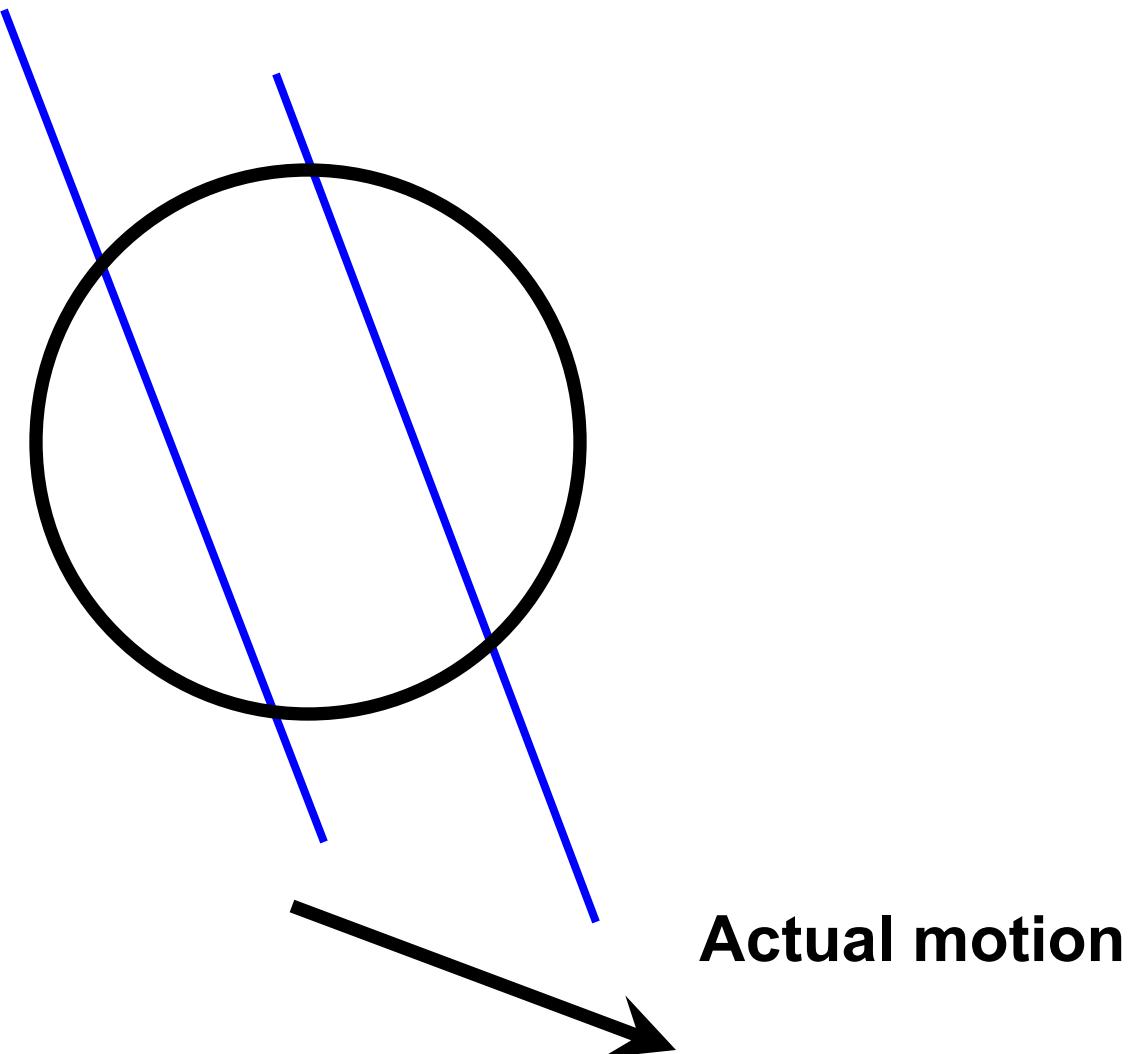
# The aperture problem

---



# The aperture problem

---



# The barber pole illusion

---



[http://en.wikipedia.org/wiki/Barberpole\\_illusion](http://en.wikipedia.org/wiki/Barberpole_illusion)

# The barber pole illusion

---



[http://en.wikipedia.org/wiki/Barberpole\\_illusion](http://en.wikipedia.org/wiki/Barberpole_illusion)

# Solving the aperture problem

---

- How to get more equations for a pixel?
- **Spatial coherence constraint:** assume the pixel's neighbors have the same  $(u, v)$ 
  - E.g., if we use a  $5 \times 5$  window, that gives us 25 equations per pixel

$$\nabla I(\mathbf{x}_i) \cdot [u, v] + I_t(\mathbf{x}_i) = 0$$

$$\begin{bmatrix} I_x(\mathbf{x}_1) & I_y(\mathbf{x}_1) \\ I_x(\mathbf{x}_2) & I_y(\mathbf{x}_2) \\ \vdots & \vdots \\ I_x(\mathbf{x}_n) & I_y(\mathbf{x}_n) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{x}_1) \\ I_t(\mathbf{x}_2) \\ \vdots \\ I_t(\mathbf{x}_n) \end{bmatrix}$$

B. Lucas and T. Kanade. [An iterative image registration technique with an application to stereo vision](#). In *Proceedings of the International Joint Conference on Artificial Intelligence*, pp. 674–679, 1981.

# Lucas-Kanade flow

---

- Linear least squares problem:

$$\begin{bmatrix} I_x(\mathbf{x}_1) & I_y(\mathbf{x}_1) \\ I_x(\mathbf{x}_2) & I_y(\mathbf{x}_2) \\ \vdots & \vdots \\ I_x(\mathbf{x}_n) & I_y(\mathbf{x}_n) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{x}_1) \\ I_t(\mathbf{x}_2) \\ \vdots \\ I_t(\mathbf{x}_n) \end{bmatrix}$$

- When is this system solvable?

B. Lucas and T. Kanade. [An iterative image registration technique with an application to stereo vision](#). In *Proceedings of the International Joint Conference on Artificial Intelligence*, pp. 674–679, 1981.

Source: L. Lazebnik

# Lucas-Kanade flow

---

- Linear least squares problem:

$$\begin{bmatrix} I_x(\mathbf{x}_1) & I_y(\mathbf{x}_1) \\ I_x(\mathbf{x}_2) & I_y(\mathbf{x}_2) \\ \vdots & \vdots \\ I_x(\mathbf{x}_n) & I_y(\mathbf{x}_n) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{x}_1) \\ I_t(\mathbf{x}_2) \\ \vdots \\ I_t(\mathbf{x}_n) \end{bmatrix}$$

$$\begin{array}{ccc} \mathbf{A} & \mathbf{d} & \mathbf{b} \\ n \times 2 & 2 \times 1 & n \times 1 \end{array}$$

- Solution given by  $(\mathbf{A}^T \mathbf{A})\mathbf{d} = \mathbf{A}^T \mathbf{b}$

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$\mathbf{M} = \mathbf{A}^T \mathbf{A}$  is the second moment matrix!

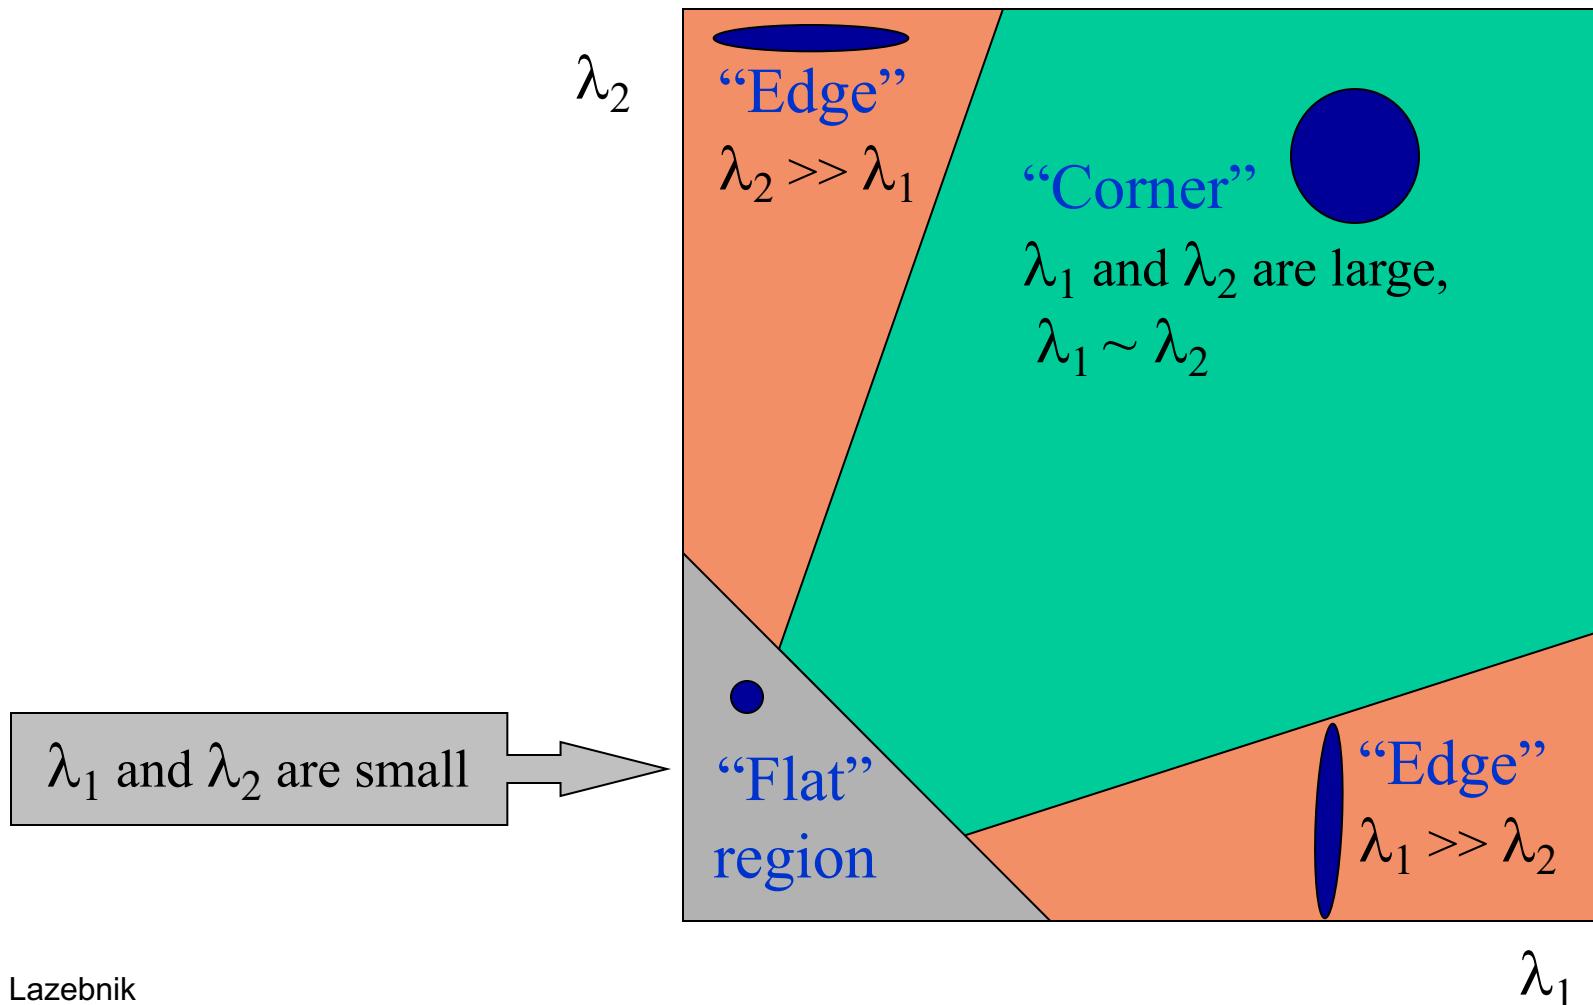
(summations are over all pixels in the window)

B. Lucas and T. Kanade. [An iterative image registration technique with an application to stereo vision](#). In *Proceedings of the International Joint Conference on Artificial Intelligence*, pp. 674–679, 1981.

# Recall: second moment matrix

---

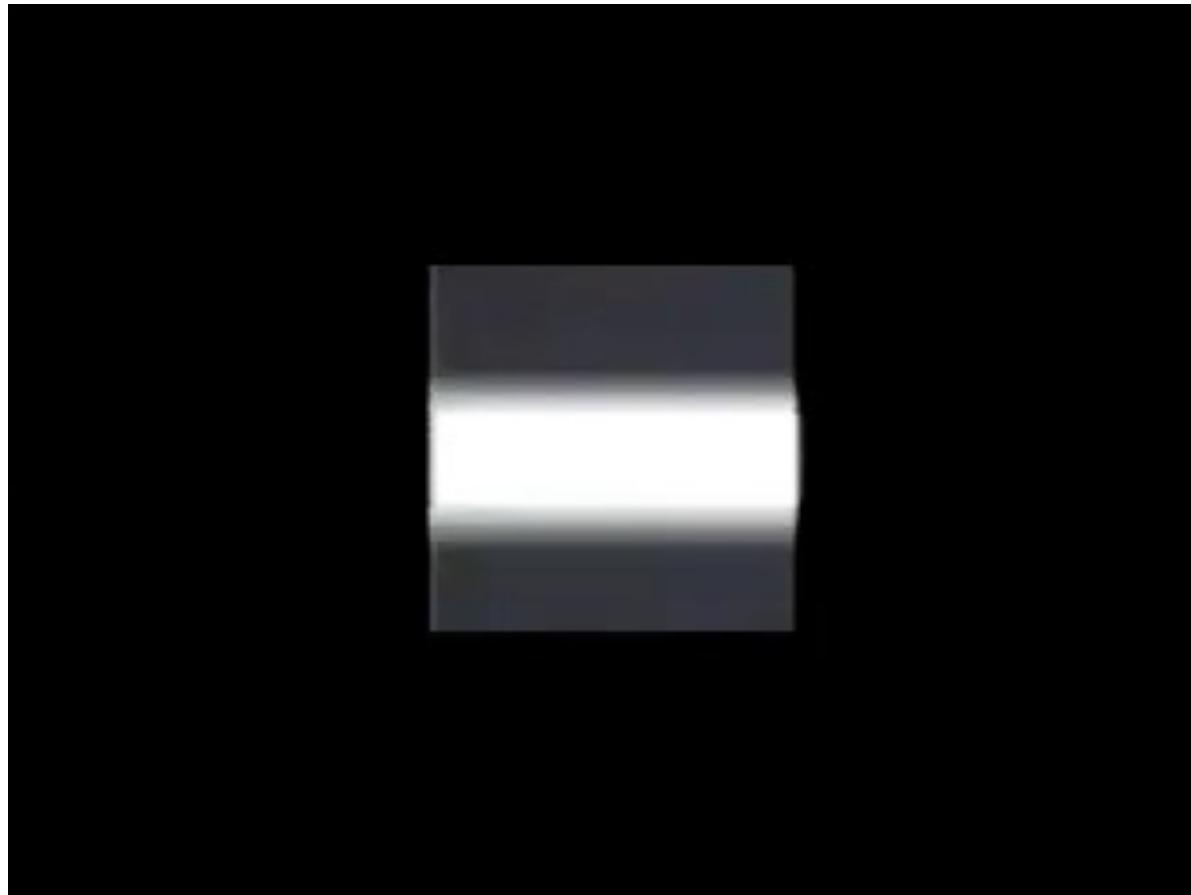
- Estimation of optical flow is well-conditioned precisely for regions with high “cornerness”:



# Conditions for solvability

---

- “Bad” case: single straight edge



# Conditions for solvability

---

- “Good” case



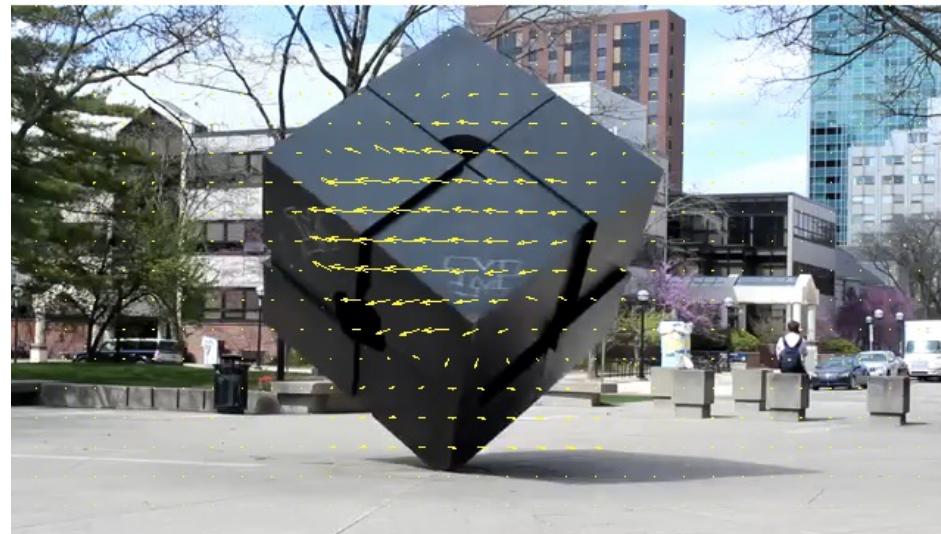
# Lucas-Kanade flow example

---

Input frames



Output



Source: [MATLAB Central File Exchange](#)

# Errors in Lucas-Kanade

---

- The motion is large (larger than a pixel)
- A point does not move like its neighbors
- Brightness constancy does not hold

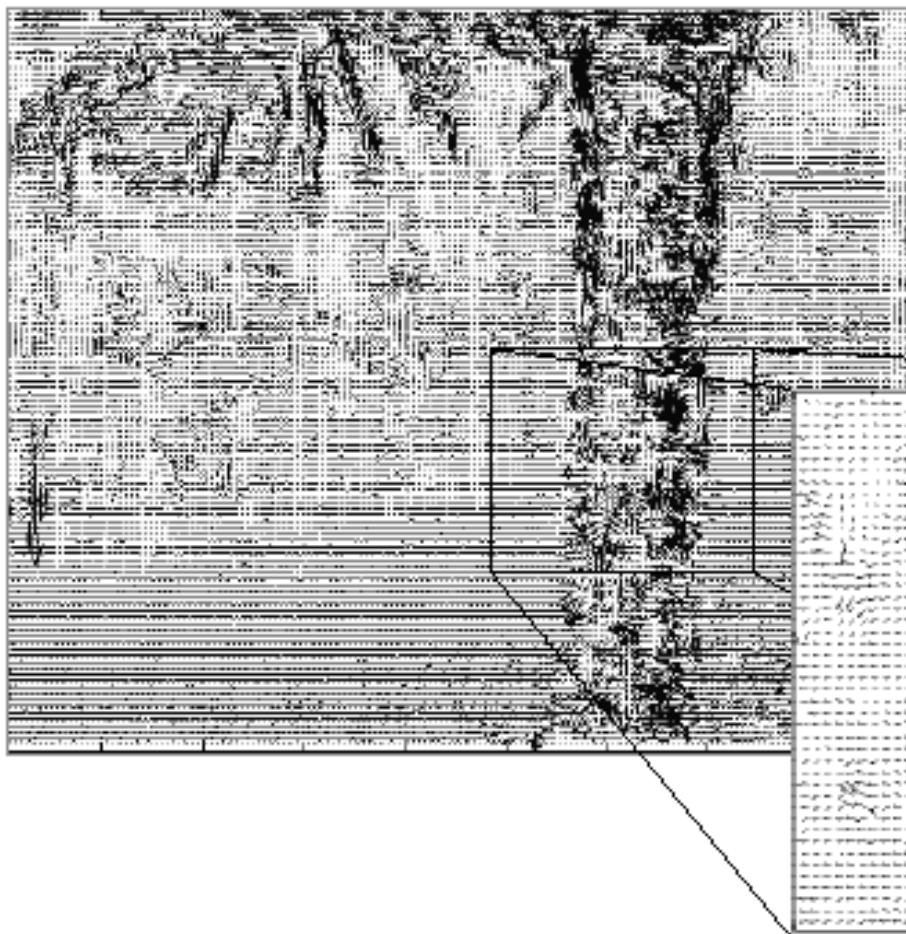
# “Flower garden” example

---

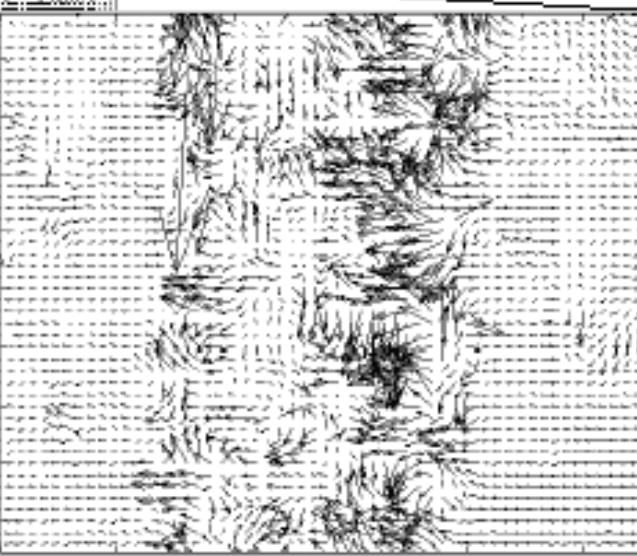


# “Flower garden” example

---

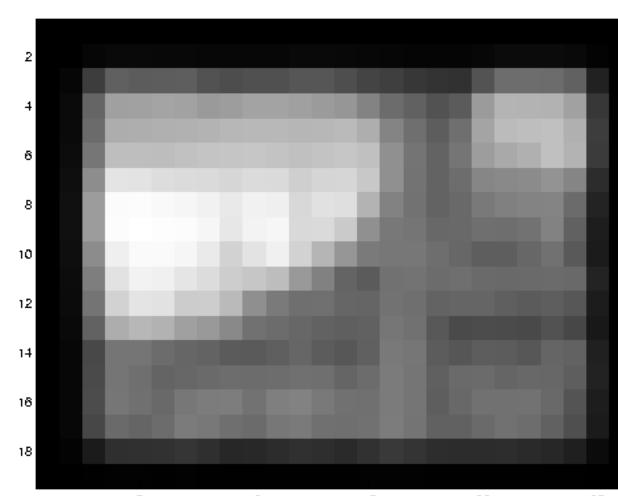
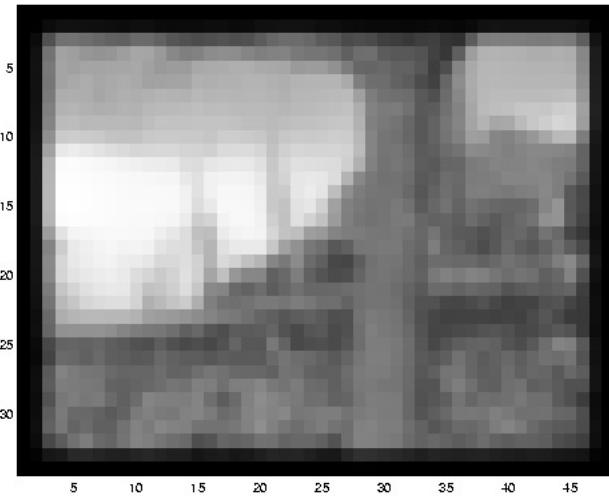
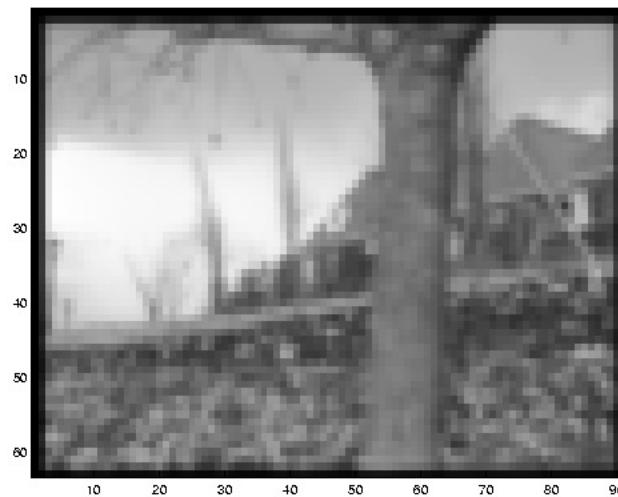


Lucas-Kanade  
fails in areas of  
large motion



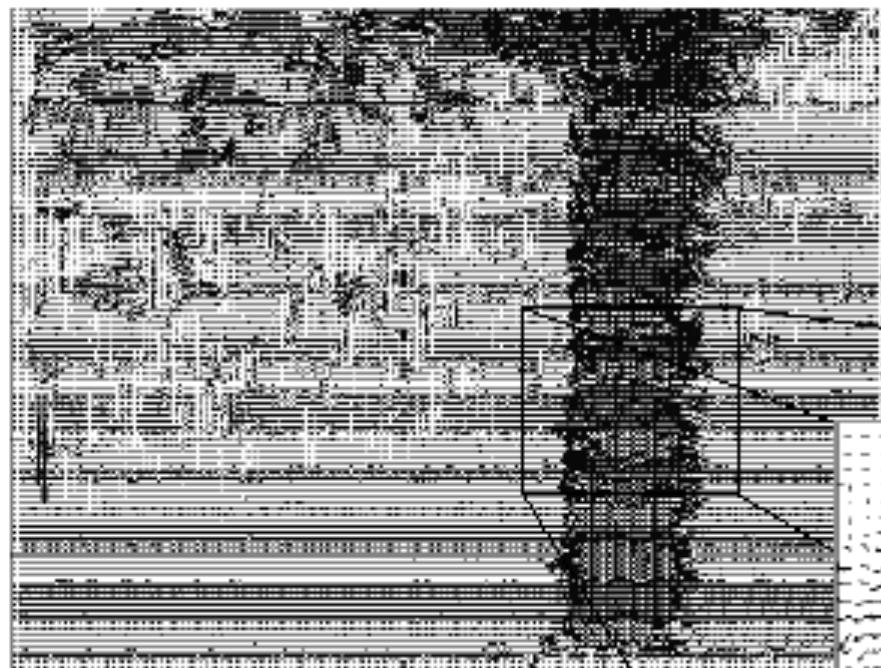
# Multi-resolution estimation

---

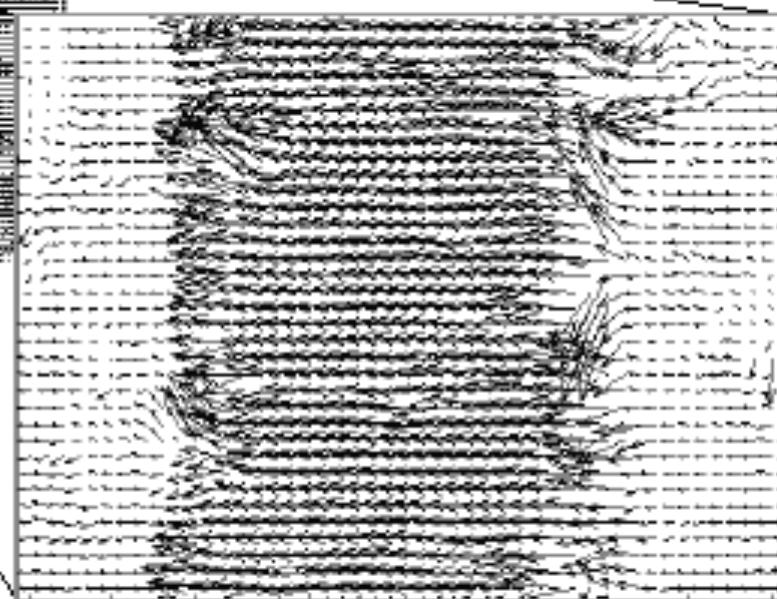


# Multi-resolution estimation

---

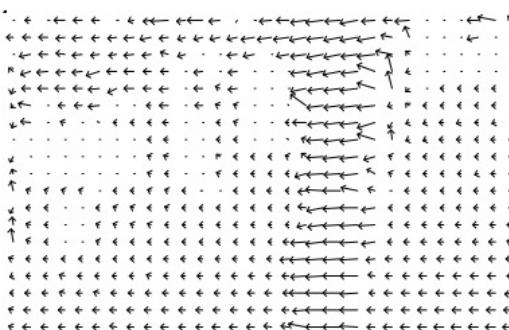


With multi-resolution  
estimation

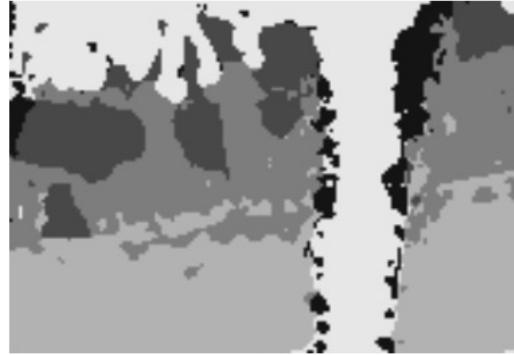


# Fixing the errors in Lucas-Kanade

- The motion is large (larger than a pixel)
  - Multi-resolution estimation, iterative refinement
  - Feature matching
- A point does not move like its neighbors
  - Motion segmentation



(a)



(b)



(c)

Figure 11: (a) The optic flow from multi-scale gradient method. (b) Segmentation obtained by clustering optic flow into affine motion regions. (c) Segmentation from consistency checking by image warping. Representing moving images with layers.

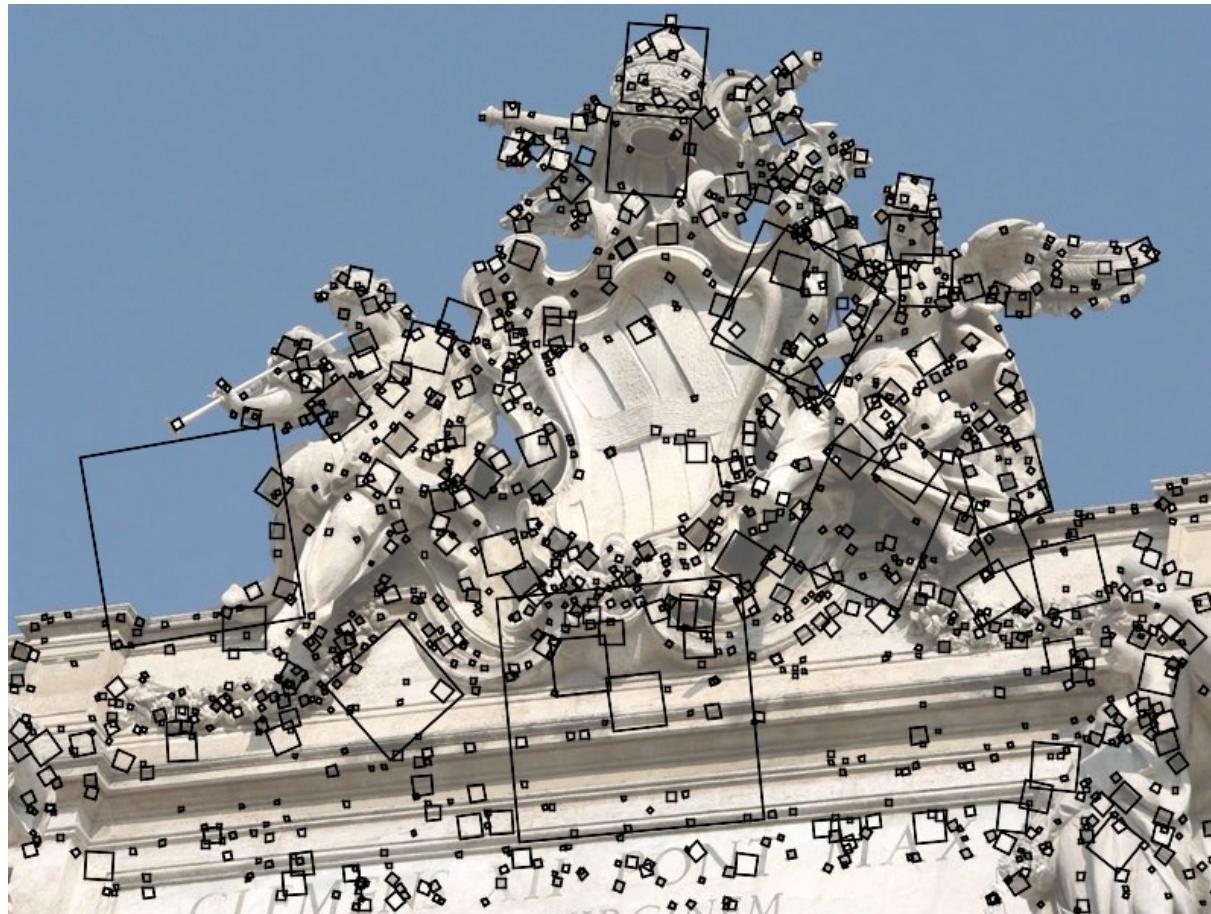
# Fixing the errors in Lucas-Kanade

---

- The motion is large (larger than a pixel)
  - Multi-resolution estimation, iterative refinement
  - Feature matching
- A point does not move like its neighbors
  - Motion segmentation
- Brightness constancy does not hold
  - Feature matching

# SIFT keypoint detection

---

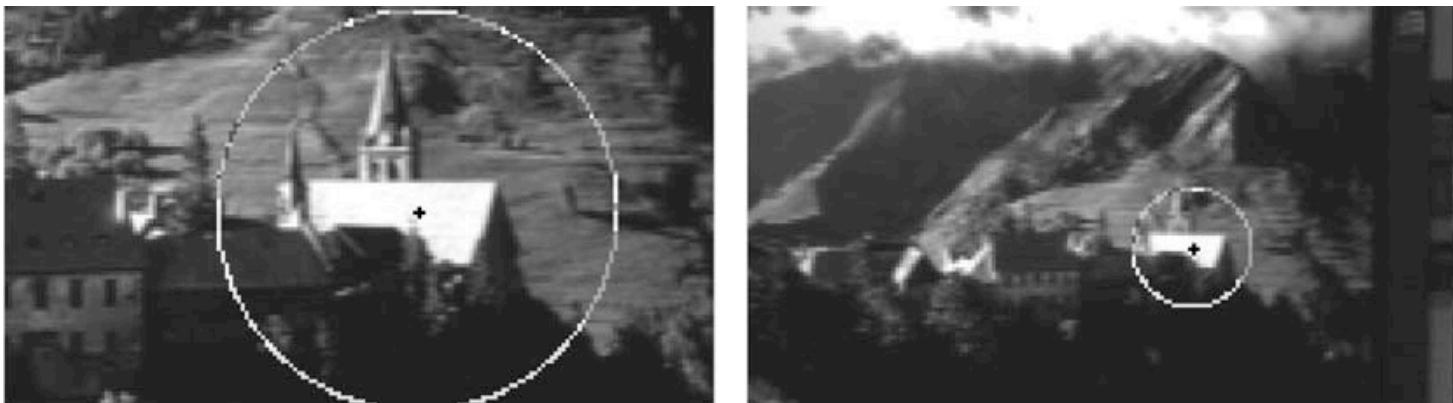


D. Lowe, [Distinctive image features from scale-invariant keypoints](#),  
*IJCV* 60 (2), pp. 91-110, 2004

# Keypoint detection with scale selection

---

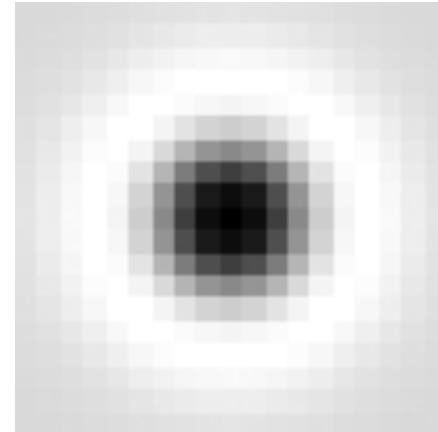
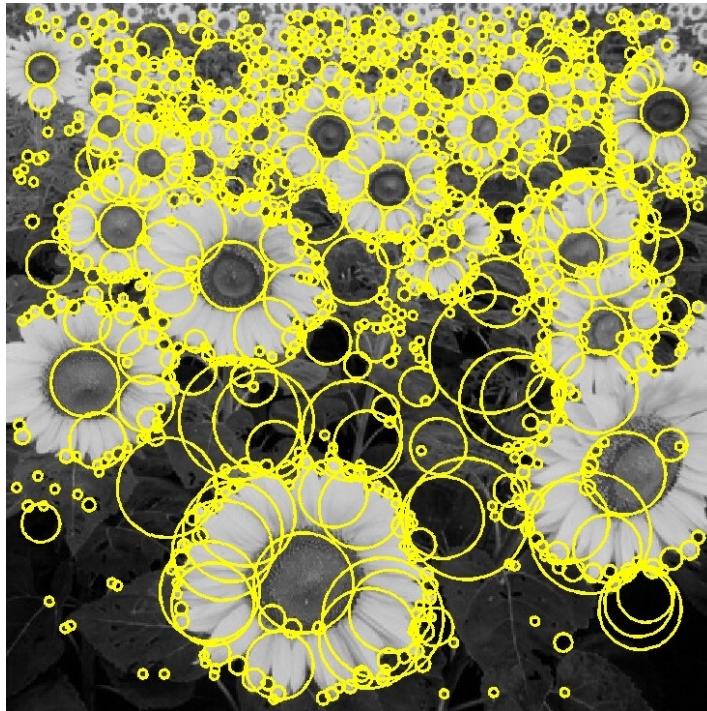
- We want to extract keypoints with *characteristic scales* that are *covariant* w.r.t. the image transformation



# Basic idea

---

- Convolve the image with a “blob filter” at multiple scales and look for extrema of filter response in the resulting *scale space*



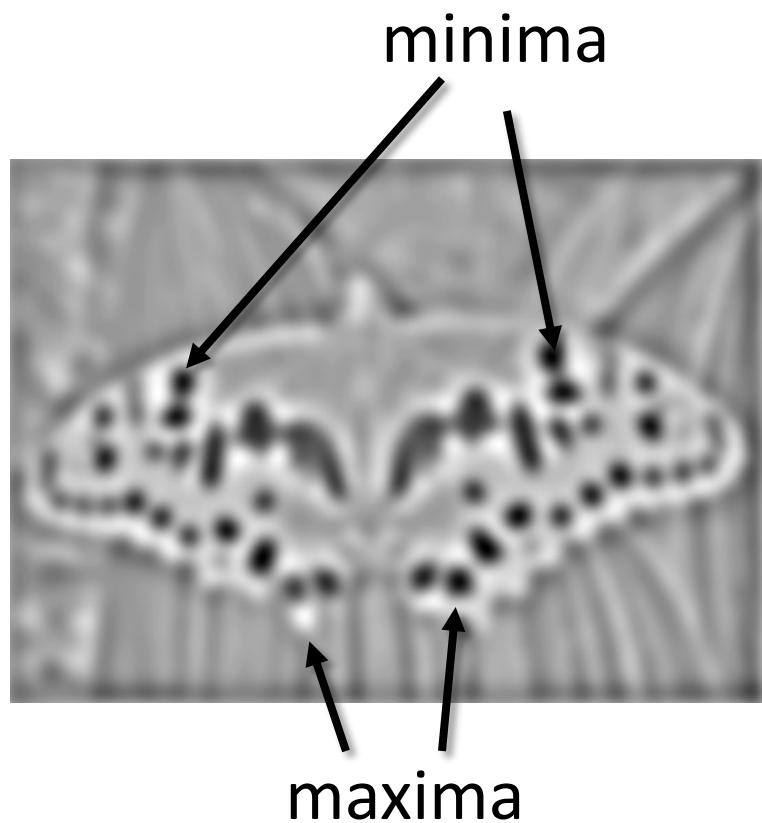
T. Lindeberg, [Feature detection with automatic scale selection](#),  
IJCV 30(2), pp 77-116, 1998

# Blob detection

---



$$* \quad \bullet =$$

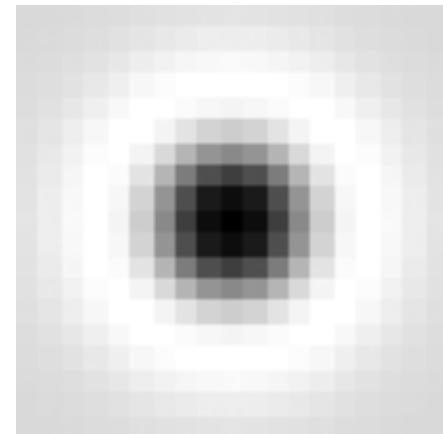
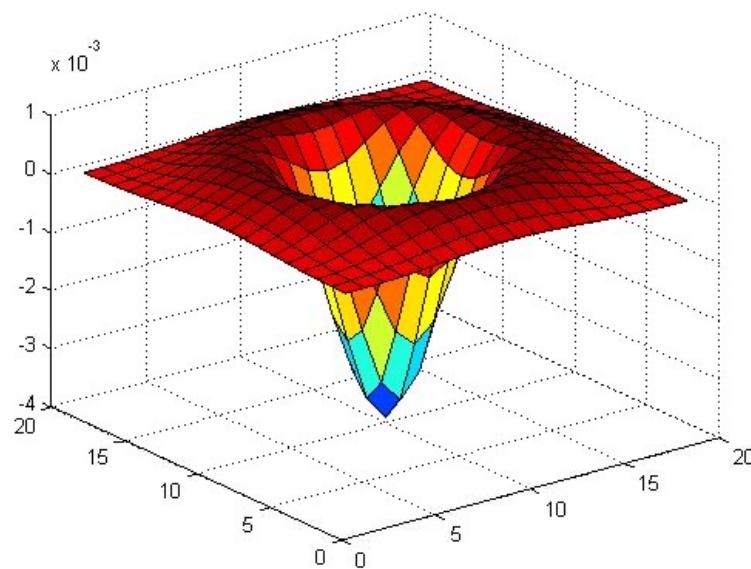


Find maxima *and minima* of blob filter response  
in space *and scale*

# Blob filter

---

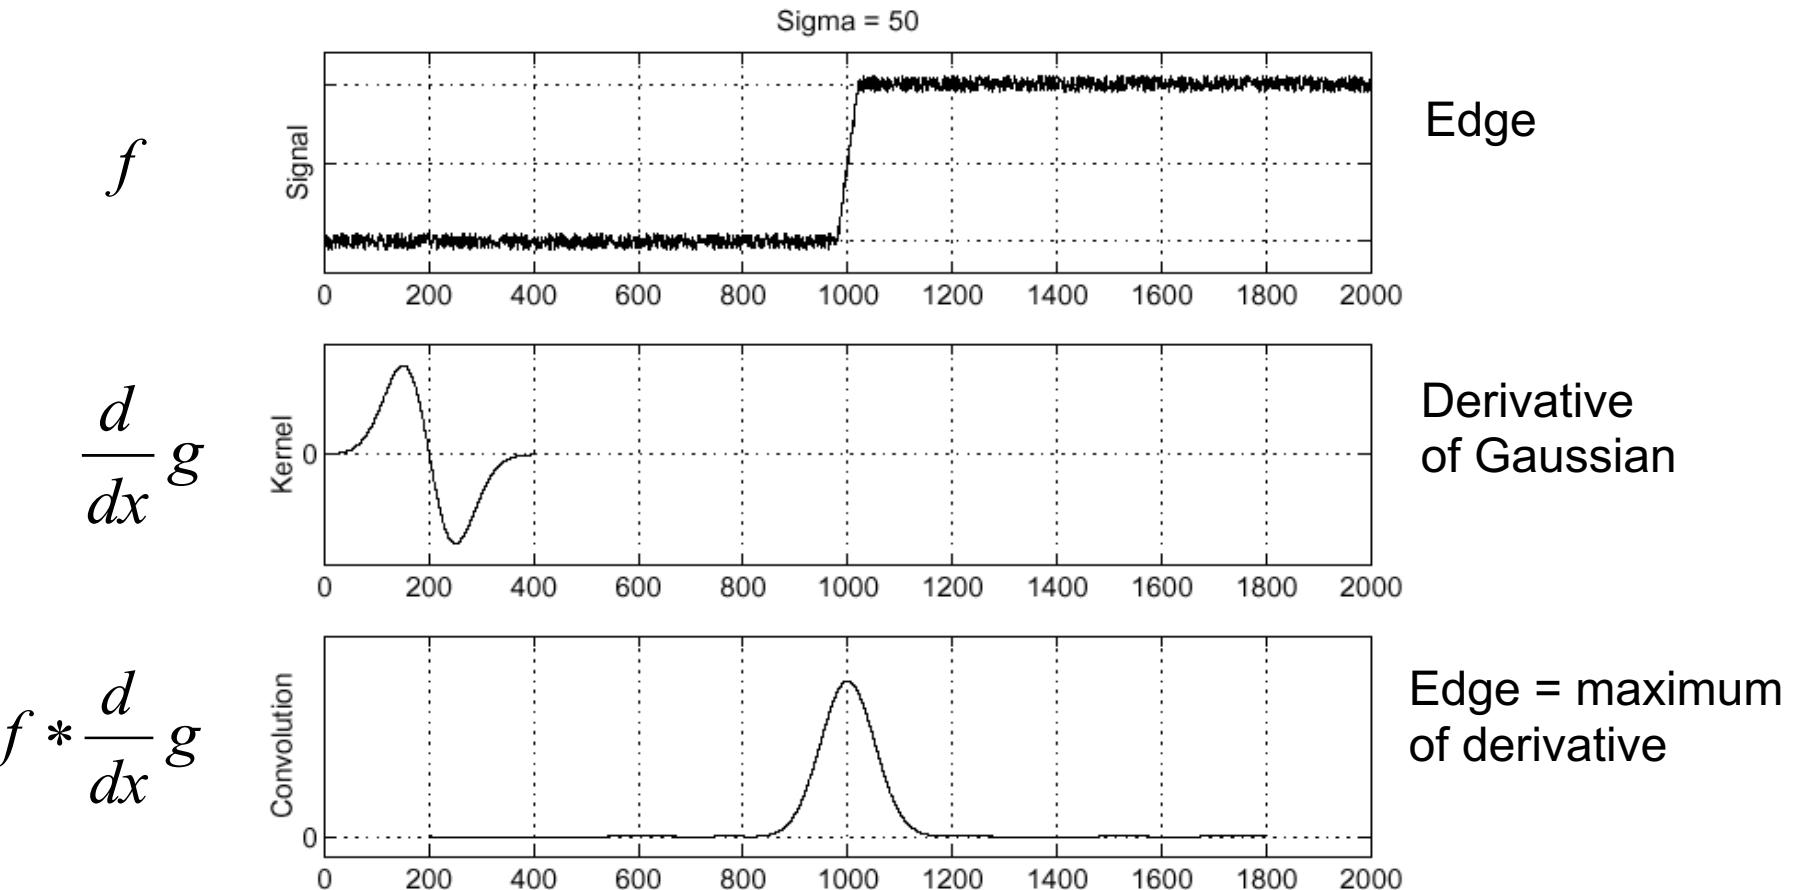
Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

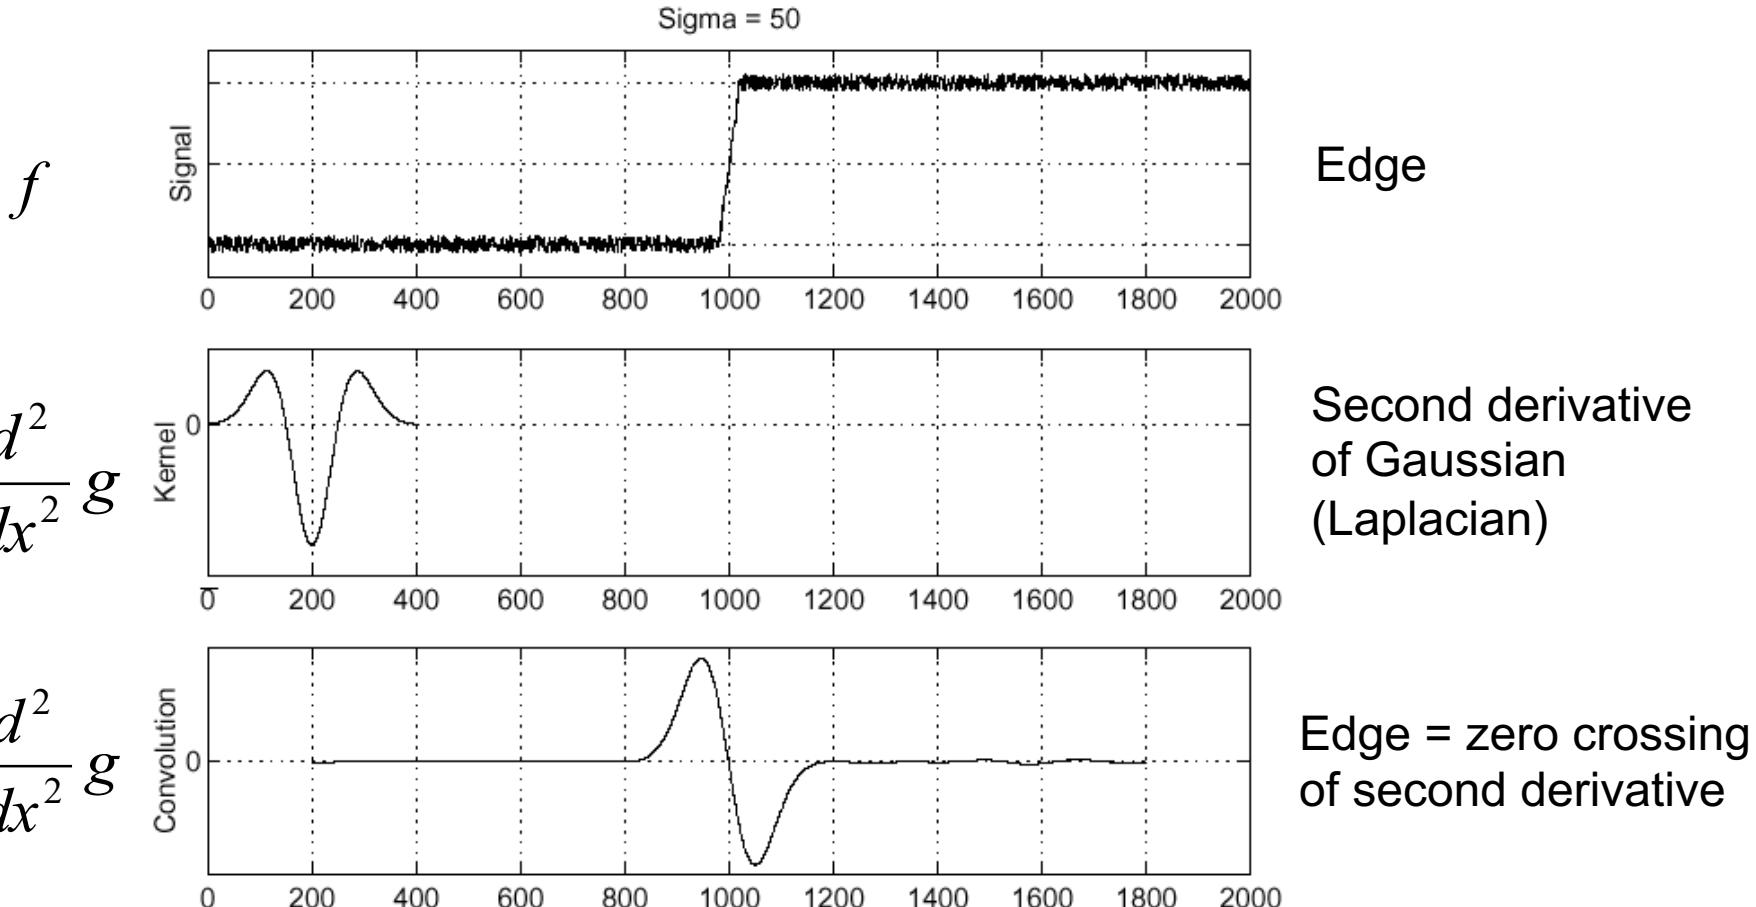
# Recall: Edge detection

---



# Edge detection, Take 2

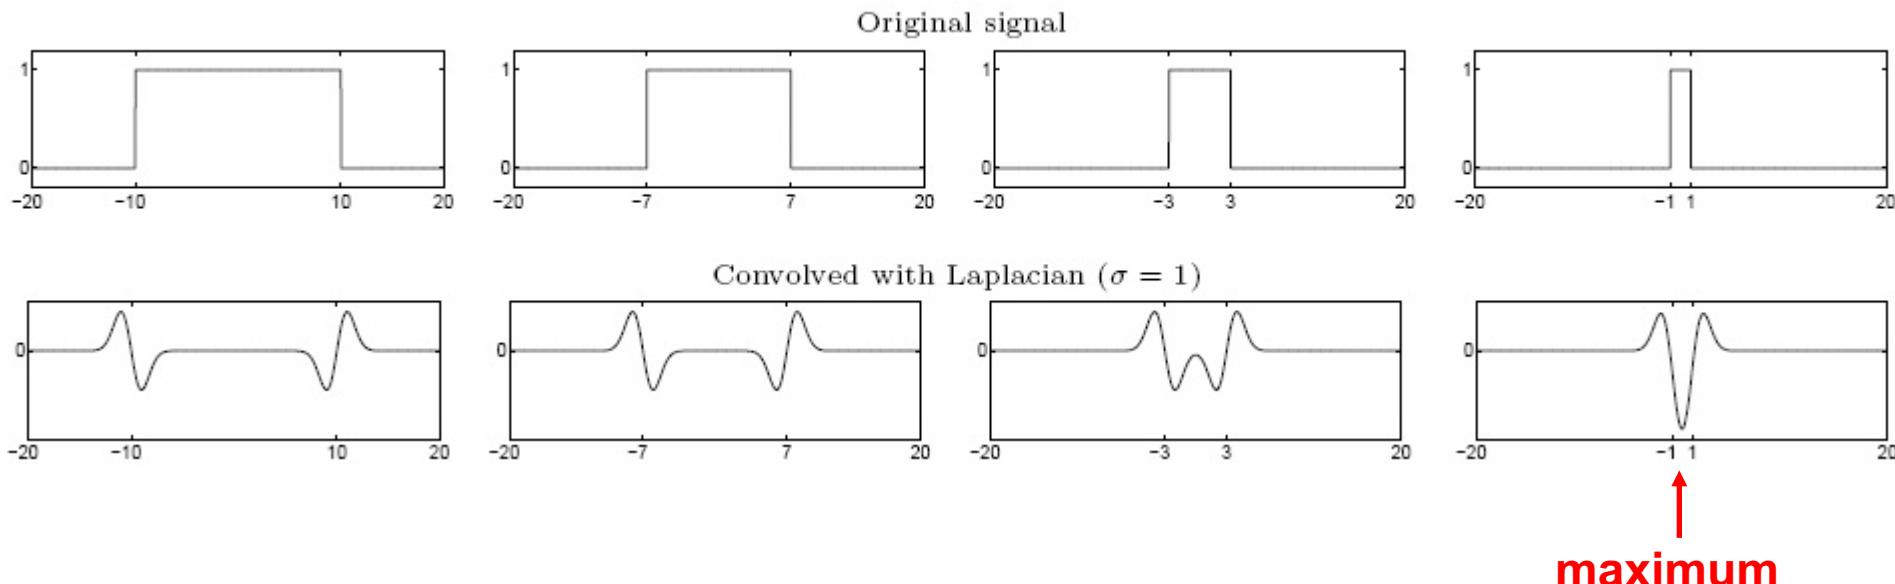
---



# From edges to blobs

---

- Edge = ripple
- Blob = superposition of two ripples

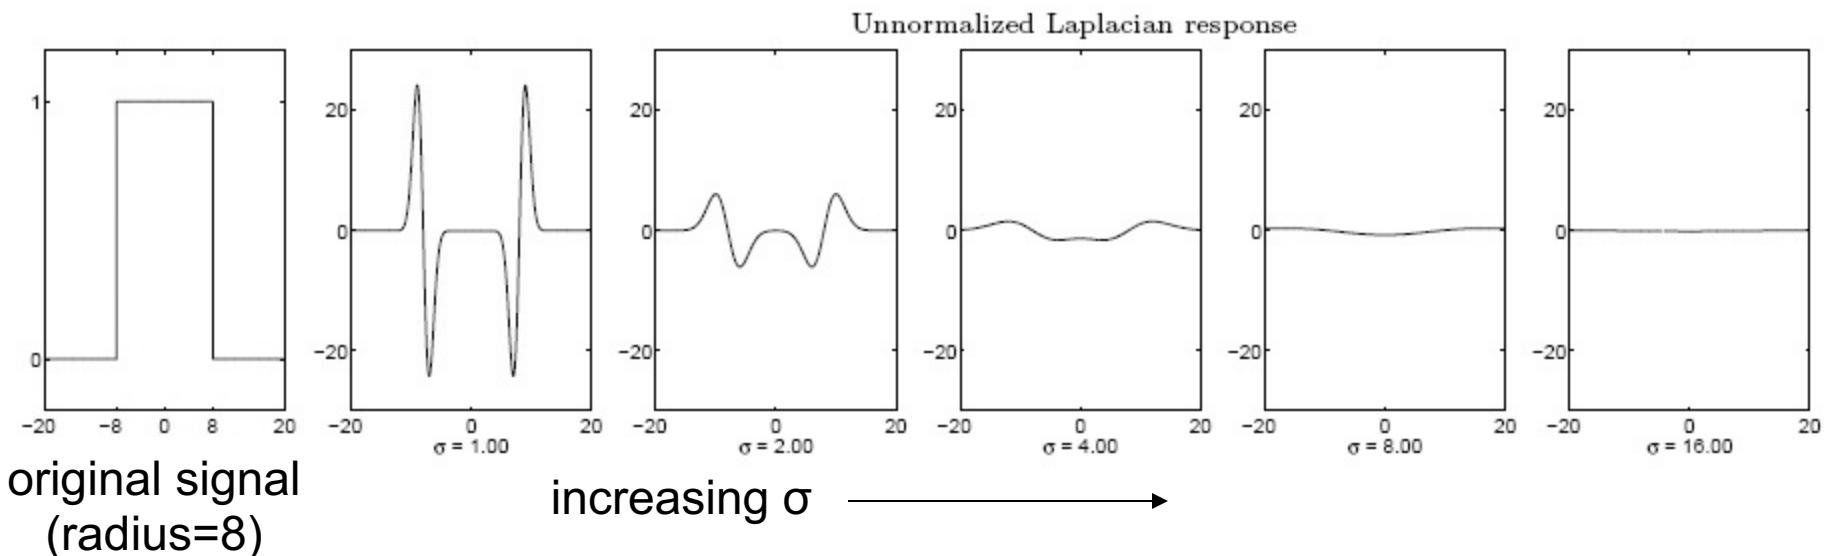


**Spatial selection:** the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob

# Scale selection

---

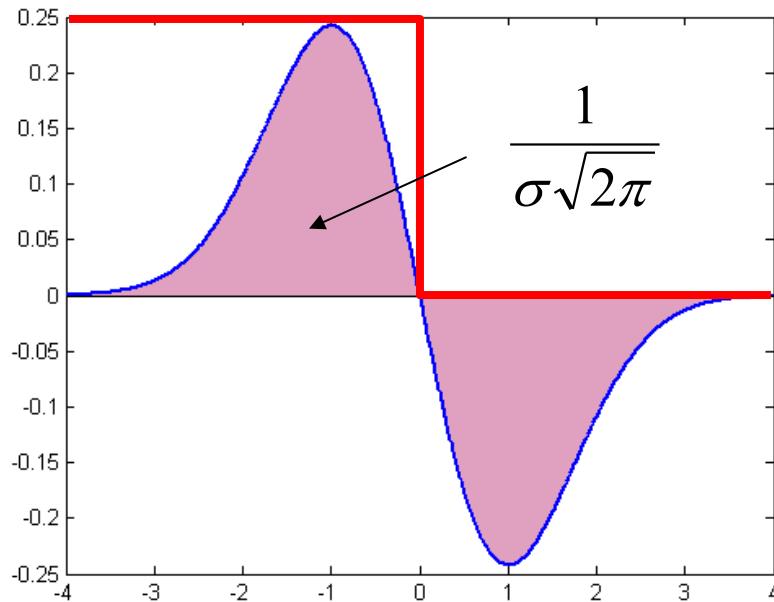
- We want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases:



# Scale normalization

---

- The response of a derivative of Gaussian filter to a perfect step edge decreases as  $\sigma$  increases:

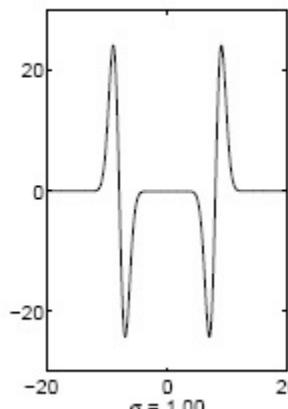
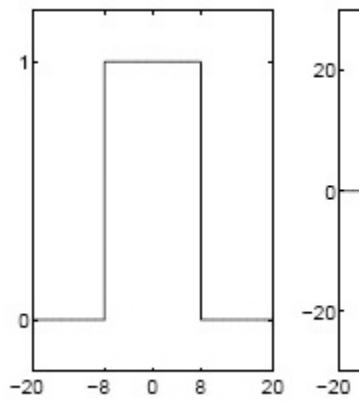


- To keep response the same (scale-invariant), must multiply Gaussian derivative by  $\sigma$
- Laplacian is the second Gaussian derivative, so it must be multiplied by  $\sigma^2$

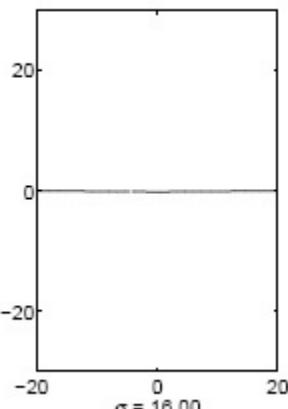
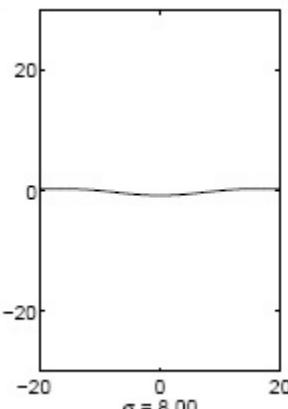
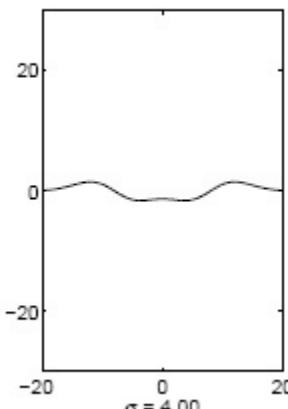
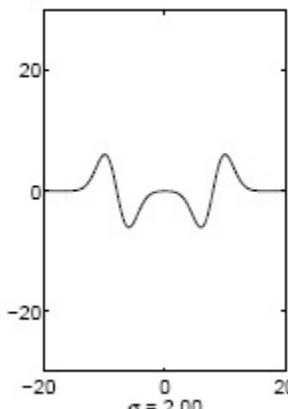
# Effect of scale normalization

---

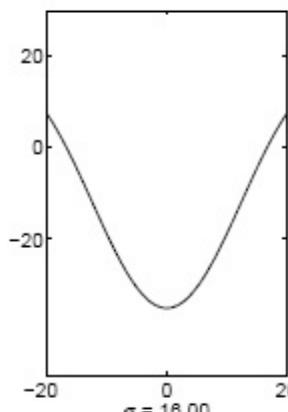
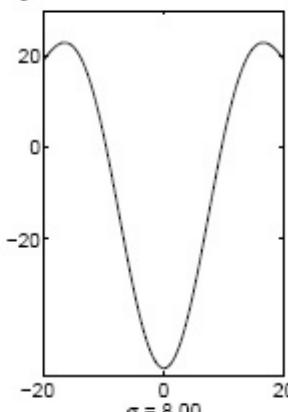
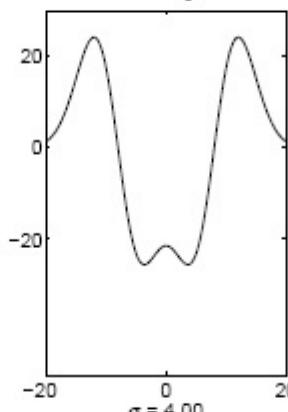
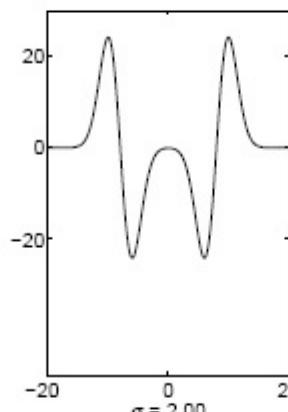
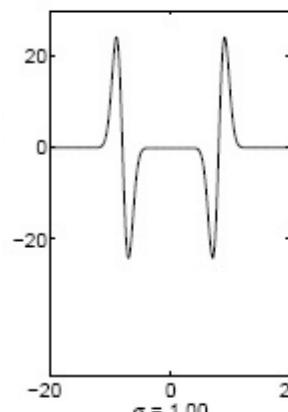
Original signal



Unnormalized Laplacian response



Scale-normalized Laplacian response



maximum

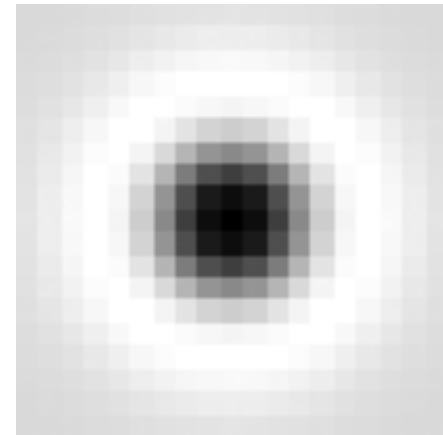
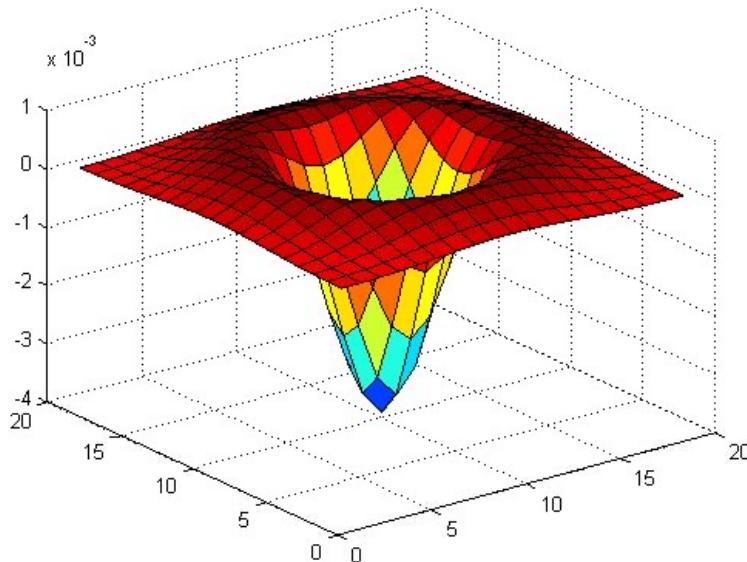


# Blob detection in 2D

---

- *Scale-normalized Laplacian of Gaussian:*

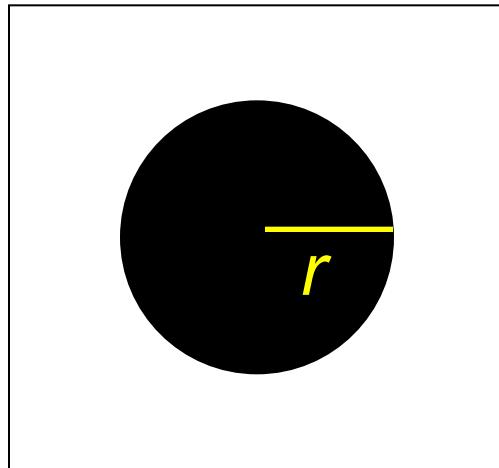
$$\nabla_{\text{norm}}^2 g = \sigma^2 \left( \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$



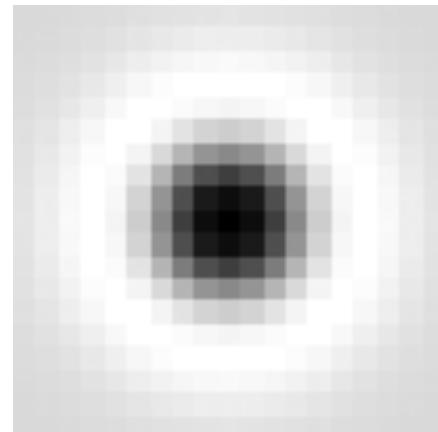
# Blob detection in 2D

---

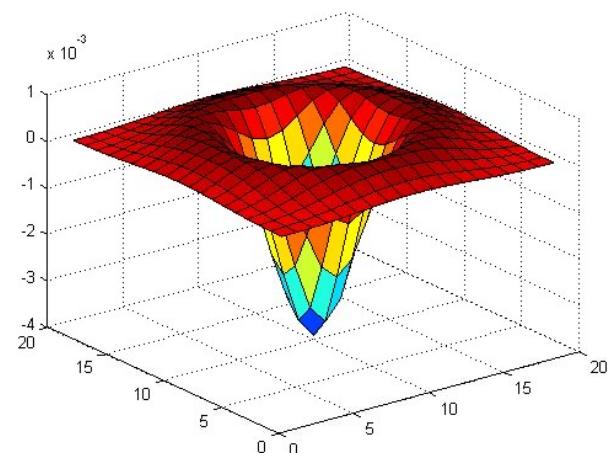
- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?



image



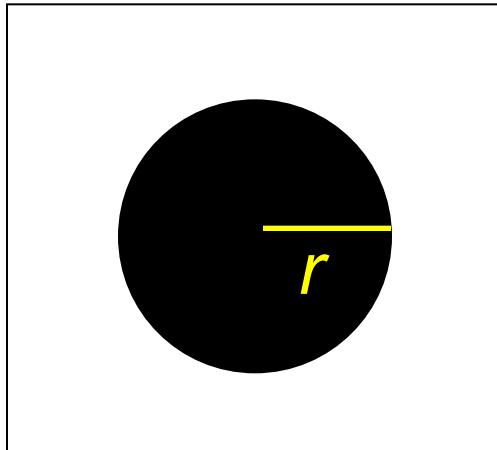
Laplacian



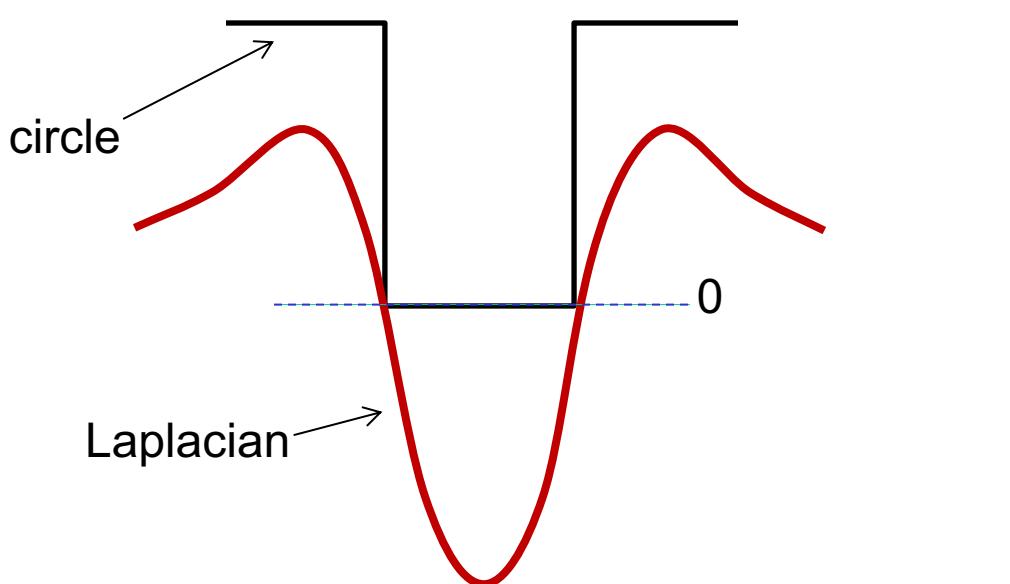
# Blob detection in 2D

---

- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle
- The Laplacian is given by (up to scale):
$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2+y^2)/2\sigma^2}$$
- Therefore, the maximum response occurs at  $\sigma = r / \sqrt{2}$ .



image



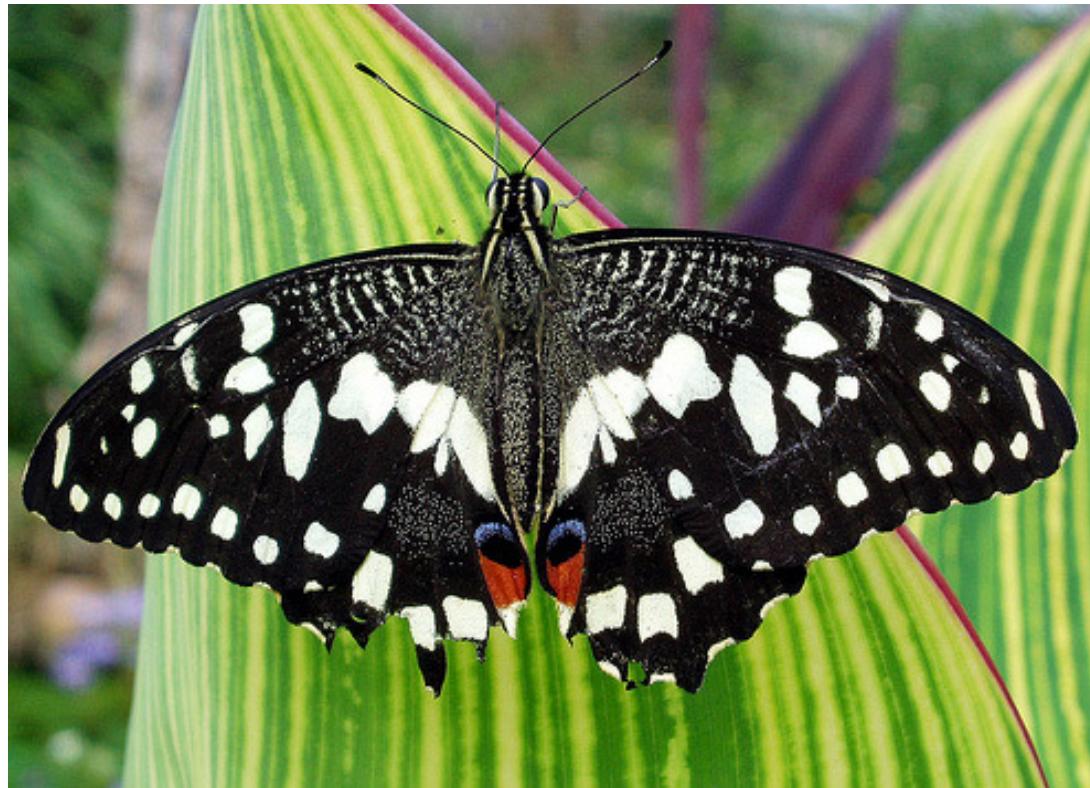
# Scale-space blob detector

---

1. Convolve image with scale-normalized Laplacian at several scales

# Scale-space blob detector: Example

---



# Scale-space blob detector: Example

---

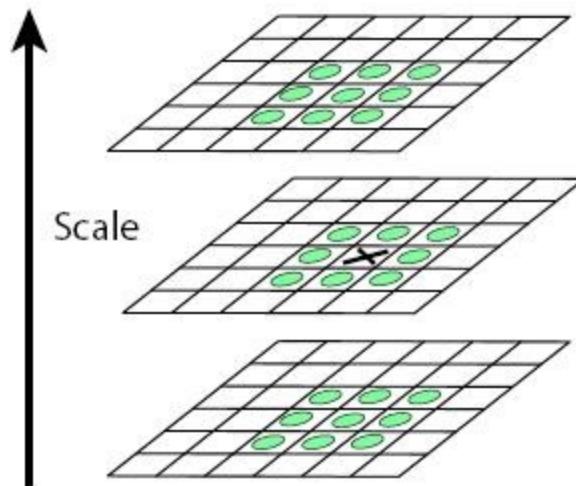


$\sigma = 11.9912$

# Scale-space blob detector

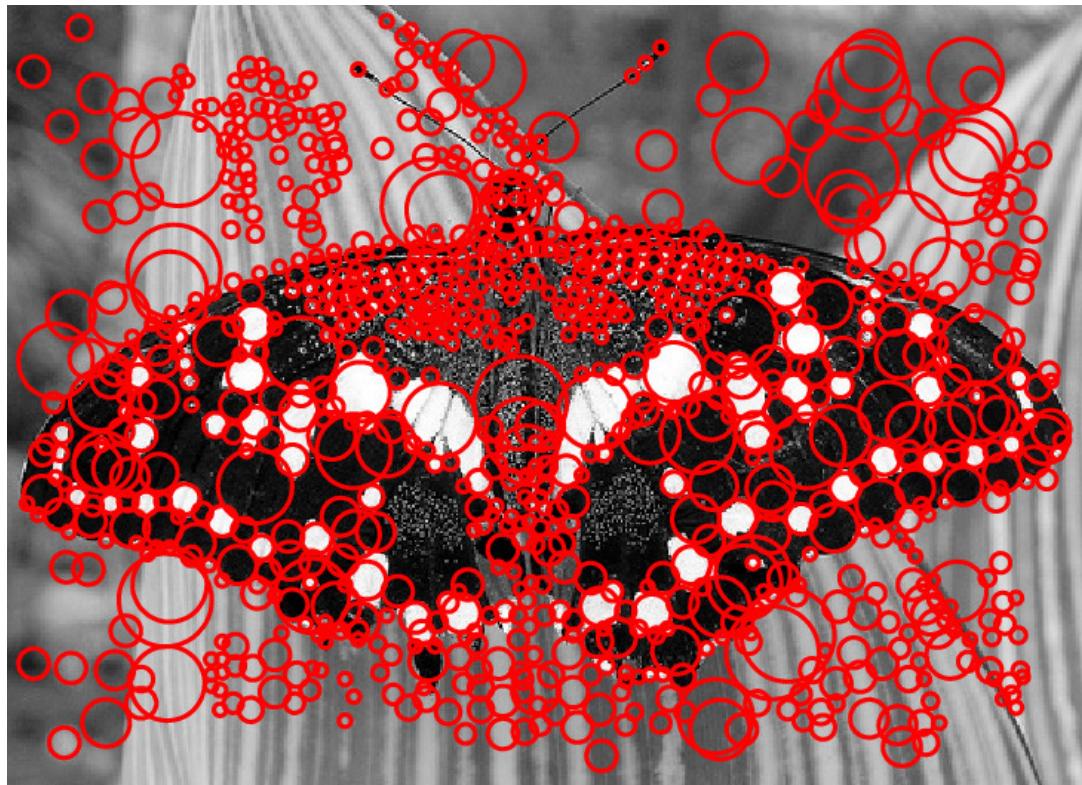
---

1. Convolve image with scale-normalized Laplacian at several scales
2. Find maxima of squared Laplacian response in scale-space



# Scale-space blob detector: Example

---



# Efficient implementation

---

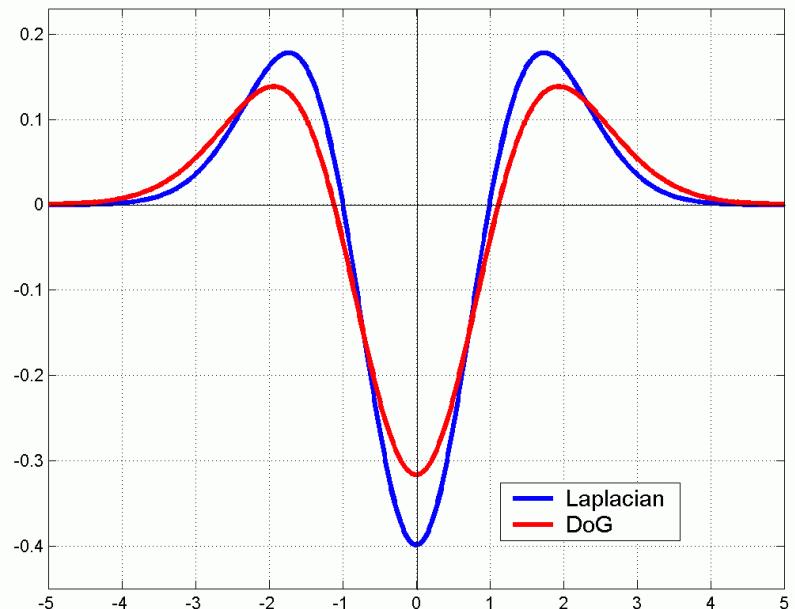
- Approximating the Laplacian with a difference of Gaussians:

$$L = \sigma^2 \left( G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

(Laplacian)

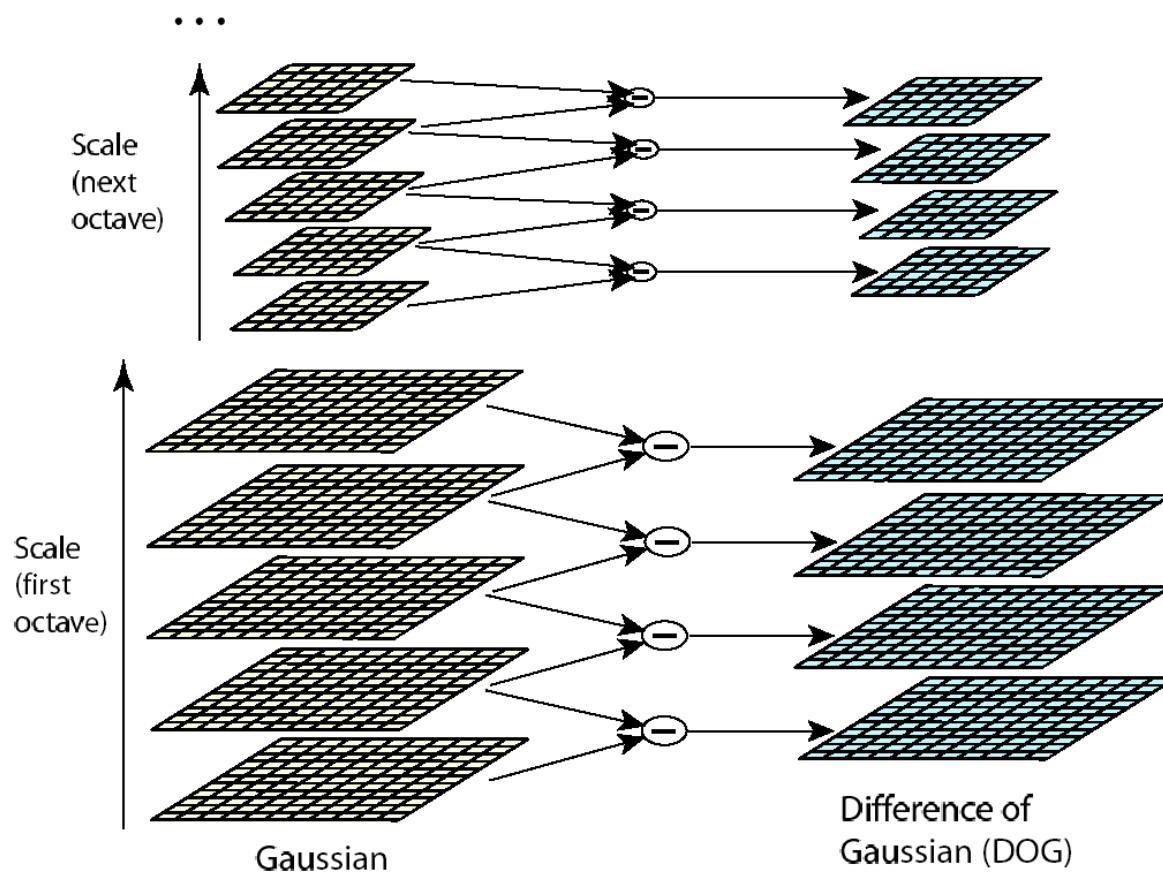
$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



# Efficient implementation

---

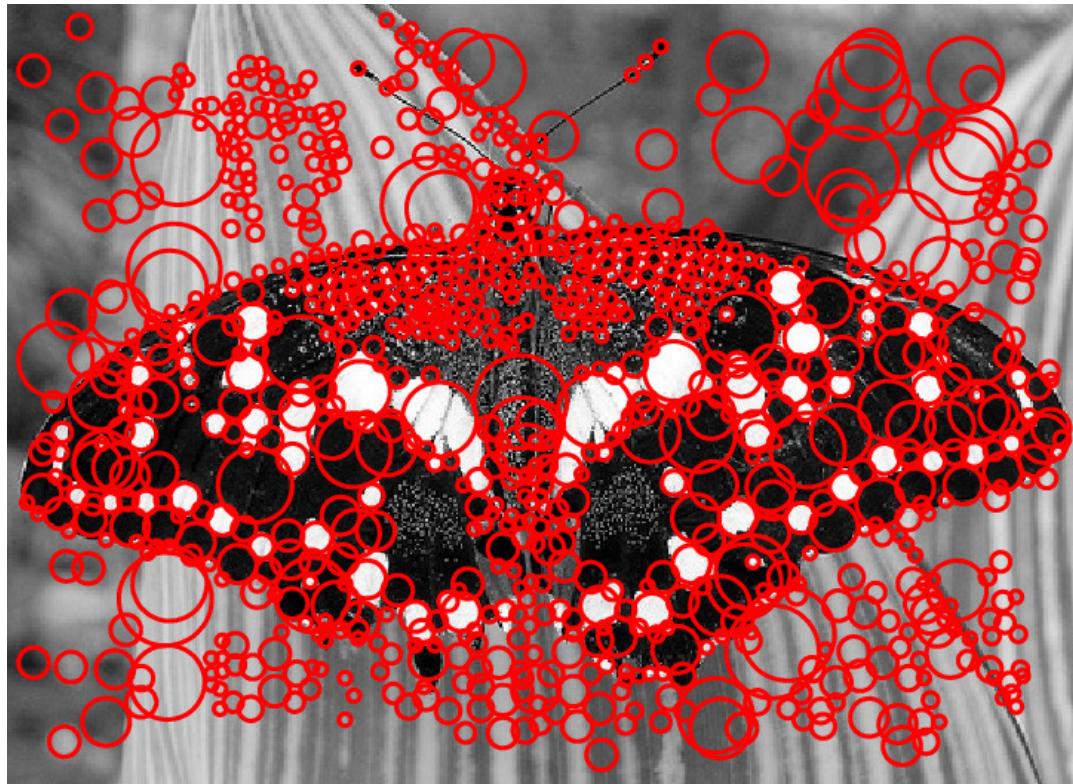


David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" IJCV 60 (2), pp. 91-110, 2004.

# Eliminating edge responses

---

- Laplacian has strong response along edges



# Eliminating edge responses

---

- Laplacian has strong response along edges

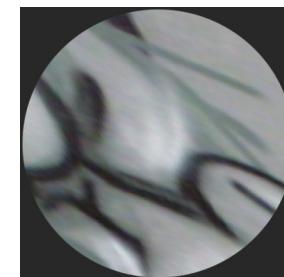
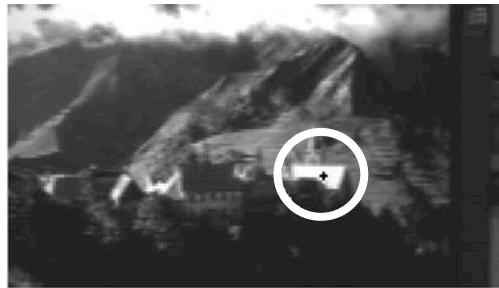
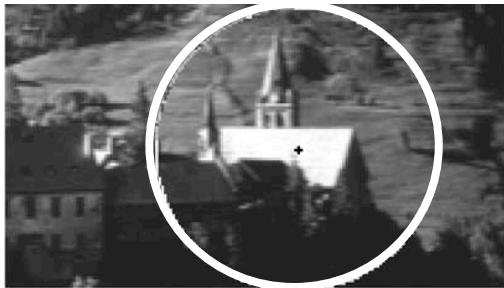


- Solution: filter based on Harris response function over neighborhoods containing the “blobs”

# From feature detection to feature description

---

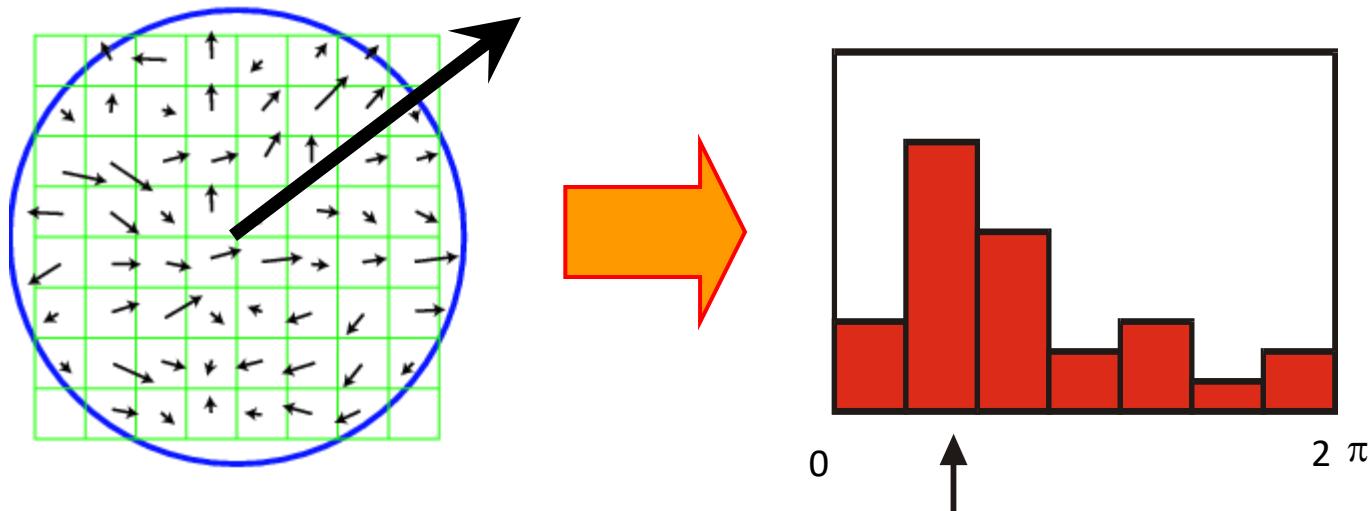
- To recognize the same pattern in multiple images, we need to match appearance “signatures” in the neighborhoods of extracted keypoints
  - But corresponding neighborhoods can be related by a scale change or rotation
  - We want to *normalize* neighborhoods to make signatures invariant to these transformations



# Finding a reference orientation

---

- Create histogram of local gradient directions in the patch
- Assign reference orientation at peak of smoothed histogram



# SIFT features

---

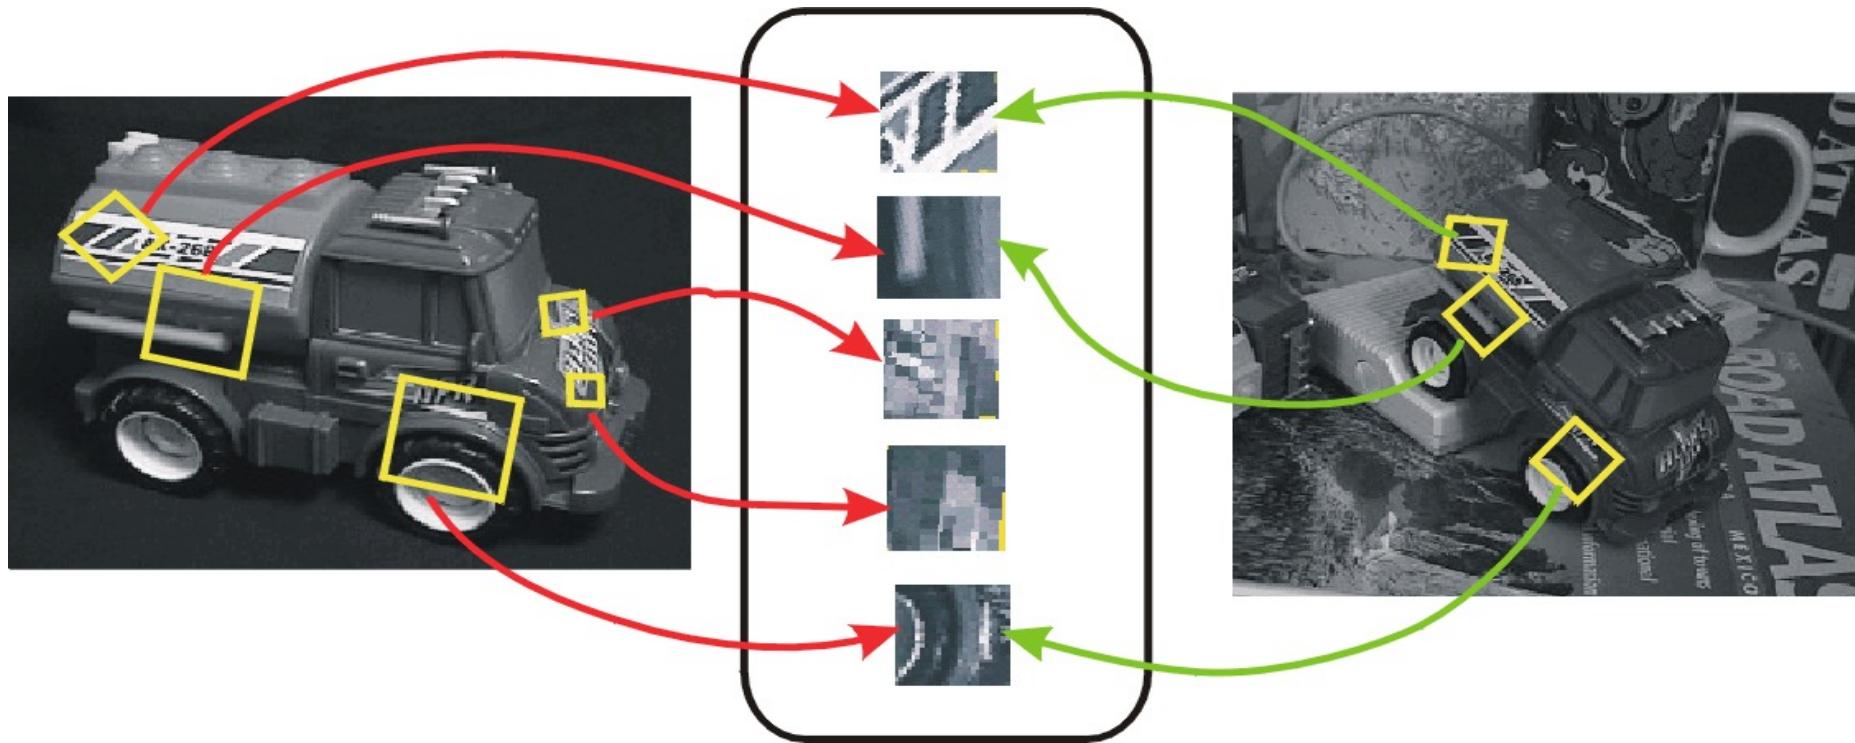
- Detected features with characteristic scales and orientations:



David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" *IJCV* 60 (2), pp. 91-110, 2004.

# From keypoint detection to feature description

---



Detection is *covariant*:

$$\text{features}(\text{transform}(\text{image})) = \text{transform}(\text{features}(\text{image}))$$

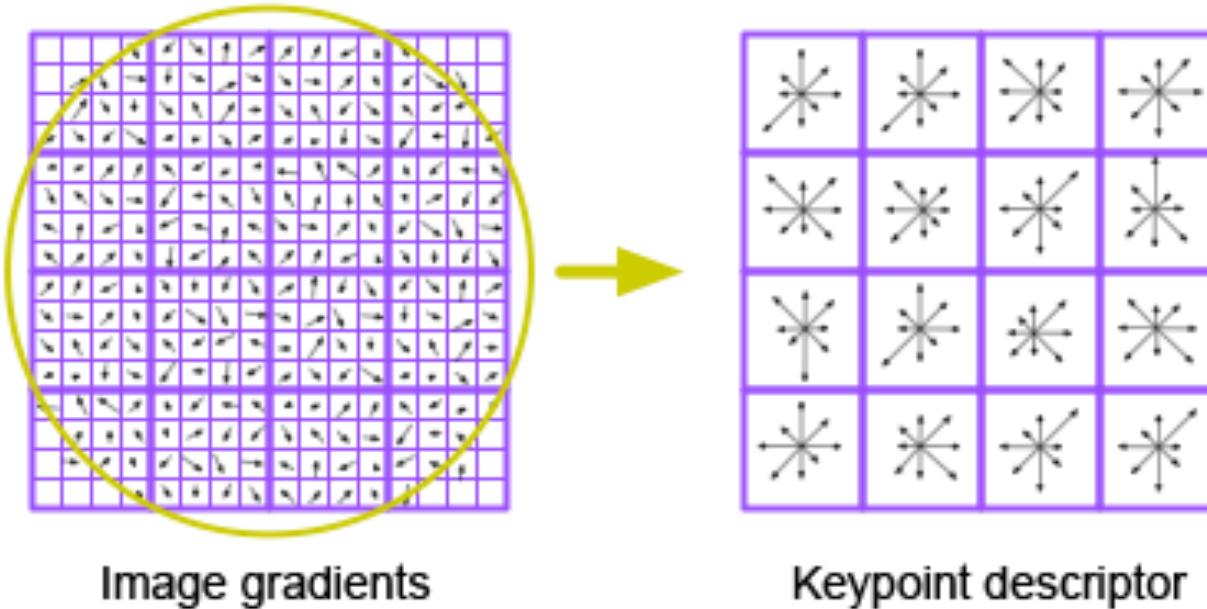
Description is *invariant*:

$$\text{features}(\text{transform}(\text{image})) = \text{features}(\text{image})$$

# SIFT descriptors

---

- Inspiration: complex neurons in the primary visual cortex



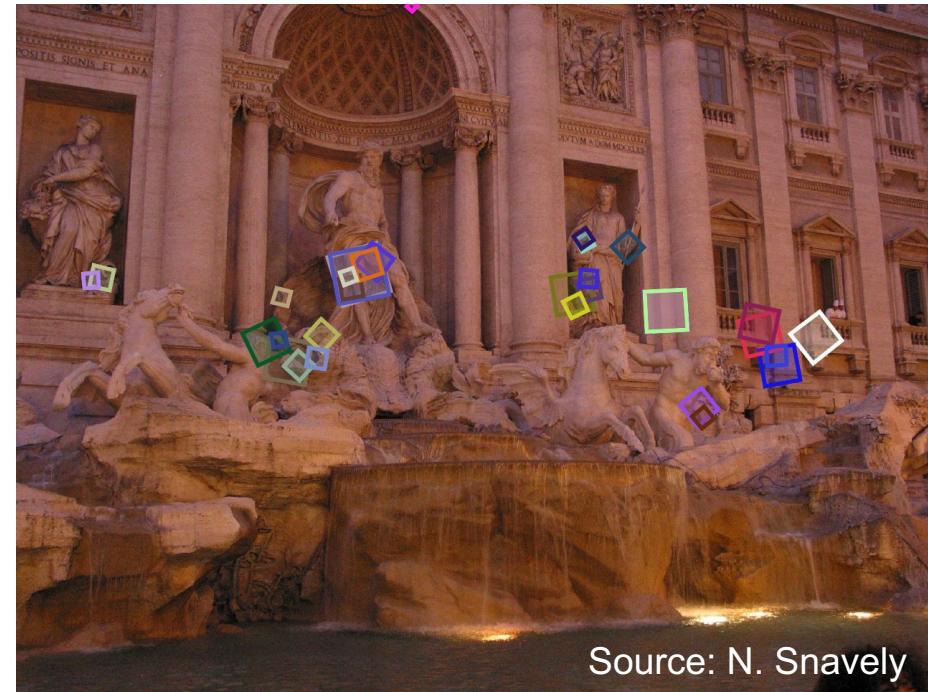
D. Lowe, [Distinctive image features from scale-invariant keypoints](#),  
IJCV 60 (2), pp. 91-110, 2004

# Properties of SIFT

---

Extraordinarily robust detection and description technique

- Can handle changes in viewpoint
  - Up to about 60 degree out-of-plane rotation
- Can handle significant changes in illumination
  - Sometimes even day vs. night
- Fast and efficient—can run in real time
- Lots of code available



Source: N. Snavely

# Fitting

Computer Vision  
CS 543 / ECE 549  
University of Illinois

# Fitting lines to point sets



In general, fit a function from a given function class to samples from the function

# Fitting: Overview

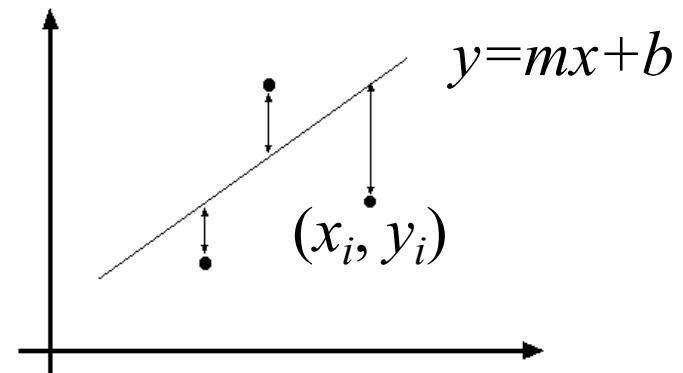
- If we know which points belong to the line, how do we find the “optimal” line parameters?
  - Least squares
- What if there are outliers?
  - Robust fitting, RANSAC
- What if there are many lines?
  - Voting methods: RANSAC, Hough transform
- What if we’re not even sure it’s a line?
  - Model selection (not covered)

# Simple example: Fitting a line

# Least squares line fitting

- Data:  $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation:  $y_i = mx_i + b$
- Find  $(m, b)$  to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$



$$E = \sum_{i=1}^n \left( \begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} - y_i \right)^2 = \left\| \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} - \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} \right\|^2 = \|\mathbf{Ap} - \mathbf{y}\|^2$$

$$= \mathbf{y}^T \mathbf{y} - 2(\mathbf{Ap})^T \mathbf{y} + (\mathbf{Ap})^T (\mathbf{Ap})$$

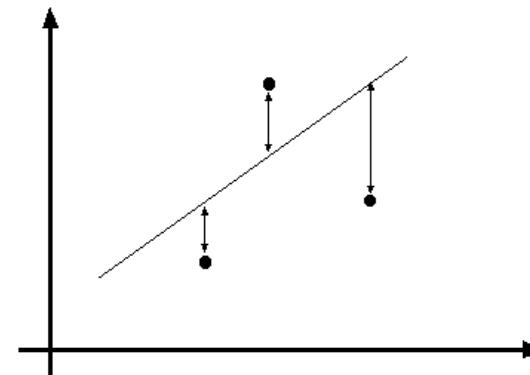
$$\frac{dE}{dp} = 2\mathbf{A}^T \mathbf{Ap} - 2\mathbf{A}^T \mathbf{y} = 0$$

Matlab:  $\mathbf{p} = \mathbf{A} \setminus \mathbf{y};$

$$\mathbf{A}^T \mathbf{Ap} = \mathbf{A}^T \mathbf{y} \Rightarrow \mathbf{p} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$$

# Problem with “vertical” least squares

- Not rotation-invariant
- Fails completely for vertical lines



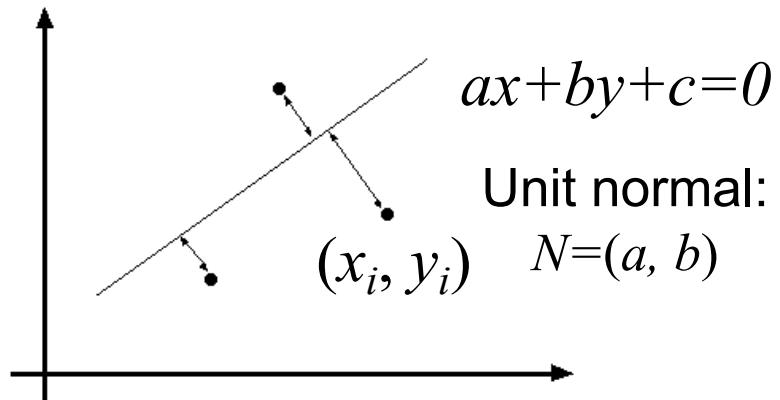
# Total least squares

If  $(a^2+b^2=1)$  then

Distance between point  $(x_i, y_i)$  and line  
 $ax+by+c=0$  is  $|ax_i + by_i + c|$

proof:

<http://mathworld.wolfram.com/Point-LineDistance2-Dimensional.html>

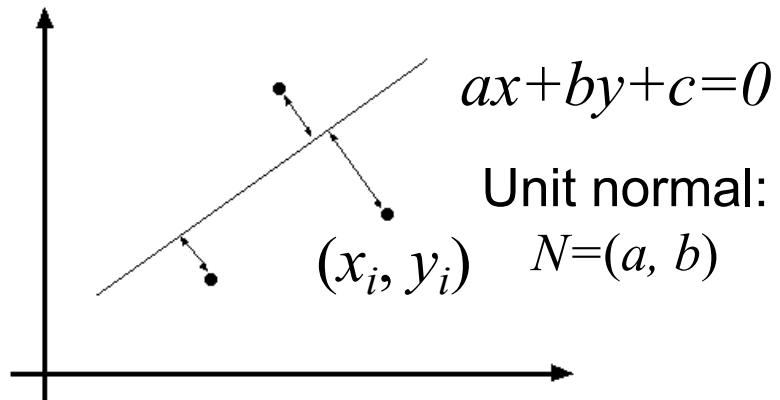


# Total least squares

If  $(a^2+b^2=1)$  then

Distance between point  $(x_i, y_i)$  and line  
 $ax+by+c=0$  is  $|ax_i + by_i + c|$

Find  $(a, b, c)$  to minimize the sum of squared perpendicular distances



$$E = \sum_{i=1}^n (ax_i + by_i + c)^2$$

# Total least squares

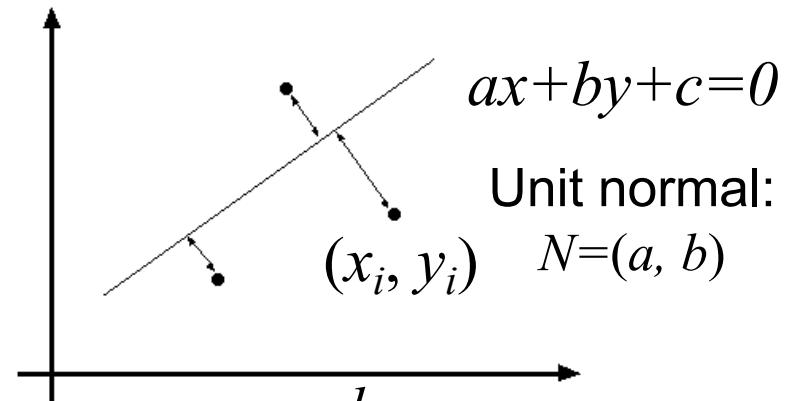
Find  $(a, b, c)$  to minimize the sum of squared perpendicular distances

$$E = \sum_{i=1}^n (ax_i + by_i + c)^2$$

$$\frac{\partial E}{\partial c} = \sum_{i=1}^n 2(ax_i + by_i + c) = 0$$

$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = \mathbf{p}^T \mathbf{A}^T \mathbf{A} \mathbf{p}$$

$$\text{minimize } \mathbf{p}^T \mathbf{A}^T \mathbf{A} \mathbf{p} \quad \text{s.t. } \mathbf{p}^T \mathbf{p} = 1 \quad \Rightarrow \quad \text{minimize } \frac{\mathbf{p}^T \mathbf{A}^T \mathbf{A} \mathbf{p}}{\mathbf{p}^T \mathbf{p}}$$



$$c = -\frac{a}{n} \sum_{i=1}^n x_i - \frac{b}{n} \sum_{i=1}^n y_i = -a\bar{x} - b\bar{y}$$

Solution is eigenvector corresponding to smallest eigenvalue of  $\mathbf{A}^T \mathbf{A}$

# Recap: Two Common Optimization Problems

## Problem statement

$$\text{minimize } \|\mathbf{Ax} - \mathbf{b}\|^2$$

least squares solution to  $\mathbf{Ax} = \mathbf{b}$

## Solution

$$\mathbf{x} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}$$

$\mathbf{x} = \mathbf{A} \setminus \mathbf{b}$  (matlab)

## Problem statement

$$\text{minimize } \mathbf{x}^T \mathbf{A}^T \mathbf{A} \mathbf{x} \text{ s.t. } \mathbf{x}^T \mathbf{x} = 1$$

$$\text{minimize } \frac{\mathbf{x}^T \mathbf{A}^T \mathbf{A} \mathbf{x}}{\mathbf{x}^T \mathbf{x}}$$

non-trivial lsq solution to  $\mathbf{Ax} = 0$

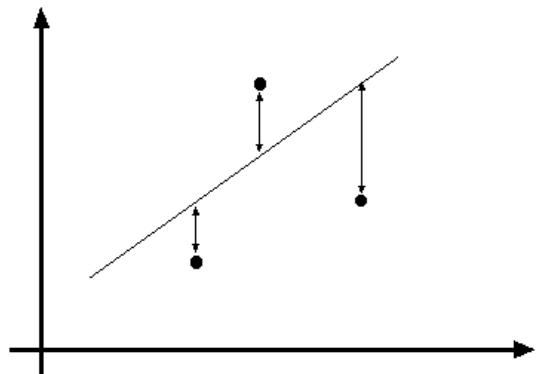
## Solution

$$[\mathbf{v}, \lambda] = \text{eig}(\mathbf{A}^T \mathbf{A})$$

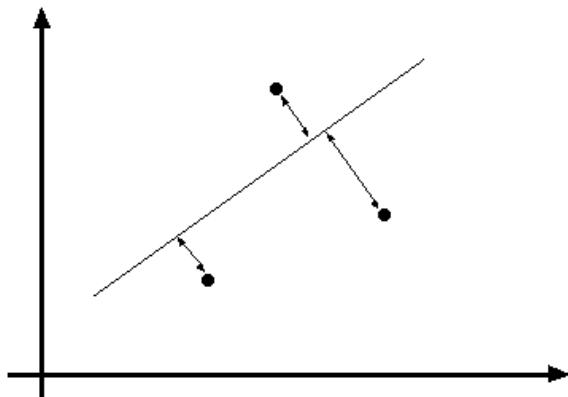
$$\lambda_1 < \lambda_{2..n} : \mathbf{x} = \mathbf{v}_1$$

# Recap: Fitting Lines

- a. fit  $y = mx + b$



- b. fit  $ax + by + c = 0$



Solution involves:

- 1. Eigen vector
  - 2. Pseudo-inverse
- 
- A.  $a \rightarrow 1, b \rightarrow 2$
  - B.  $a \rightarrow 2, b \rightarrow 1$
  - C.  $a \rightarrow 1, b \rightarrow 1$
  - D.  $a \rightarrow 2, b \rightarrow 2$

# Least squares (global) optimization

## Good

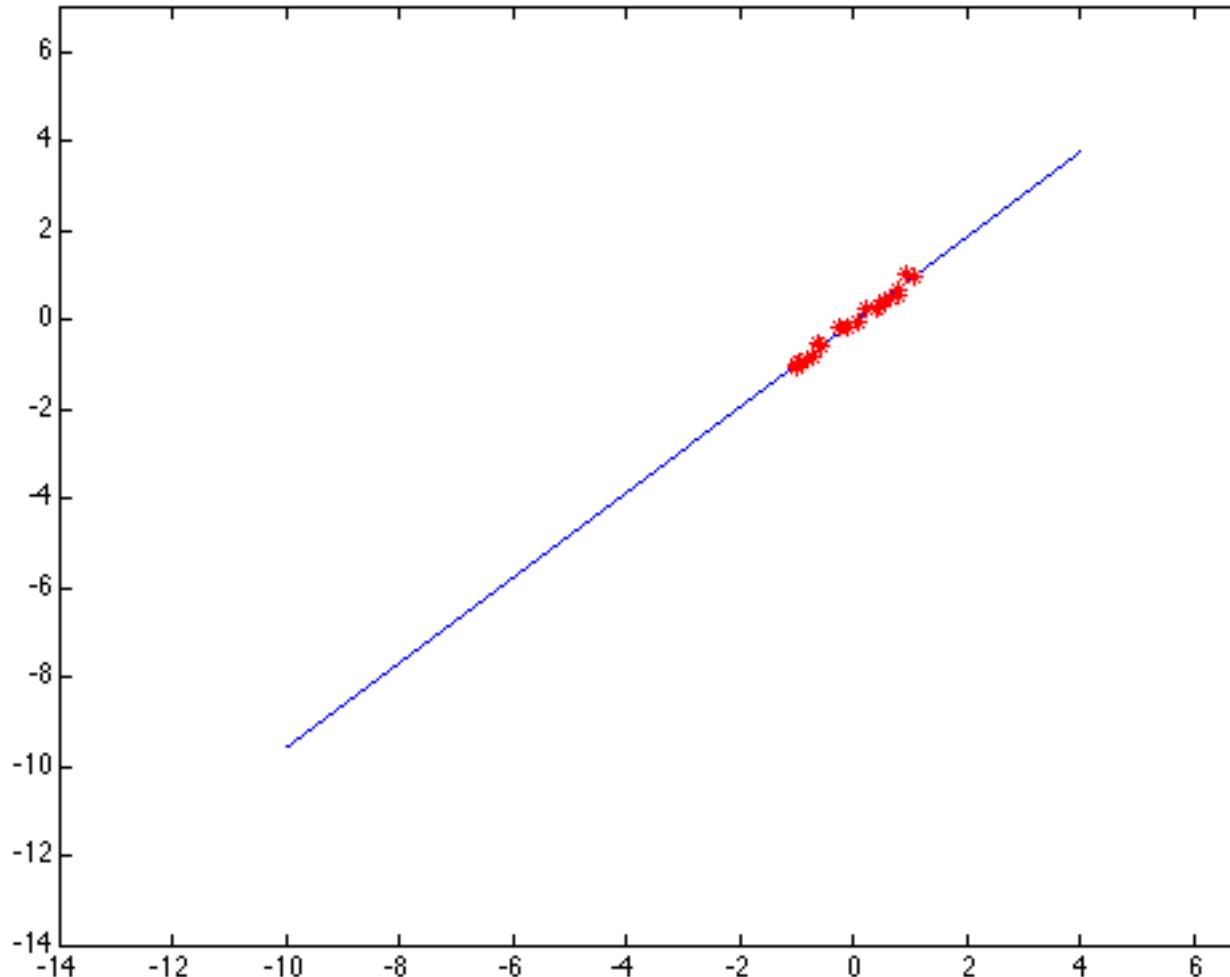
- Clearly specified objective
- Optimization is easy

## Bad

- May not be what you want to optimize
- Sensitive to outliers
  - Bad matches, extra points
- Doesn't allow you to get multiple good fits
  - Detecting multiple objects, lines, etc.

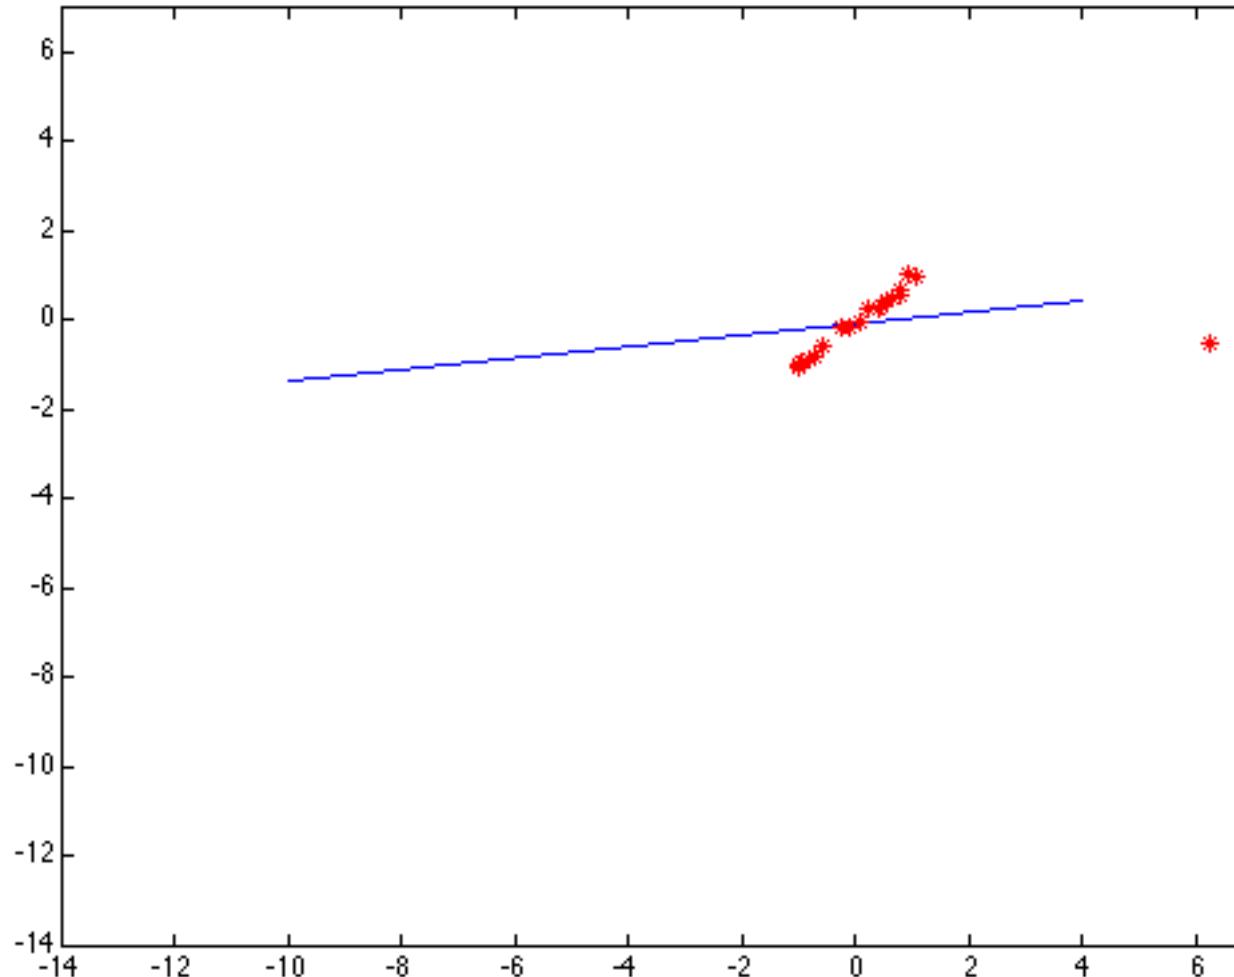
# Least squares: Robustness to noise

- Least squares fit to the red points:



# Least squares: Robustness to noise

- Least squares fit with an outlier:



Problem: squared error heavily penalizes outliers

# Robust least squares (to deal with outliers)

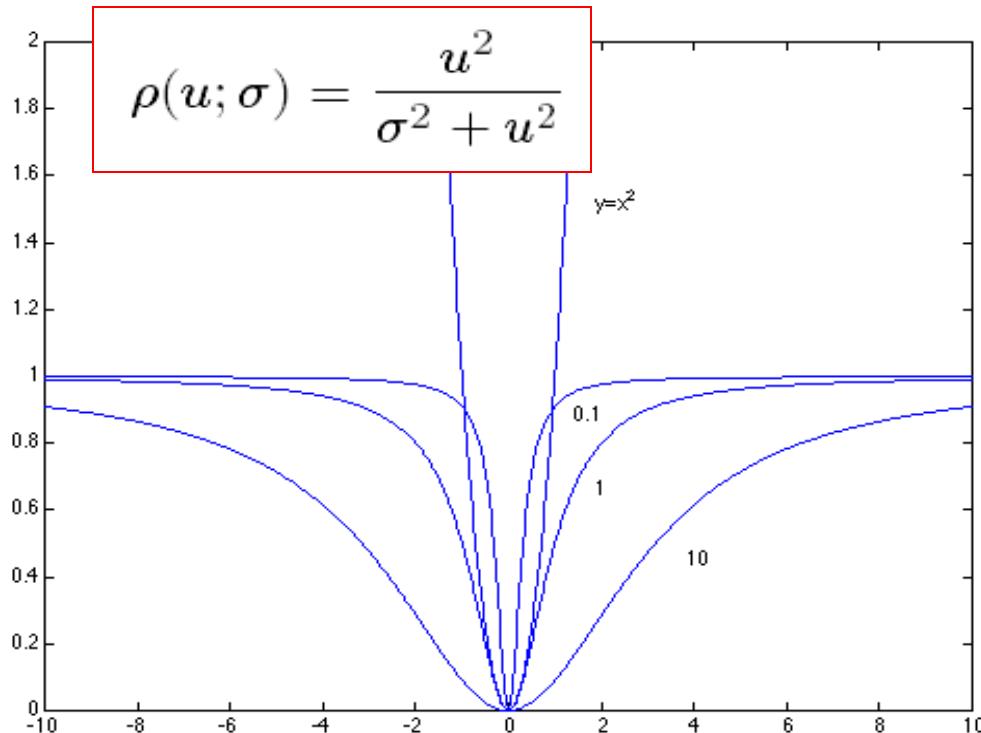
General approach:

minimize

$$\sum_i \rho(u_i(x_i, \theta); \sigma) \quad u^2 = \sum_{i=1}^n (y_i - mx_i - b)^2$$

$u_i(x_i, \theta)$  – residual of  $i^{\text{th}}$  point w.r.t. model parameters  $\vartheta$

$\rho$  – robust function with scale parameter  $\sigma$



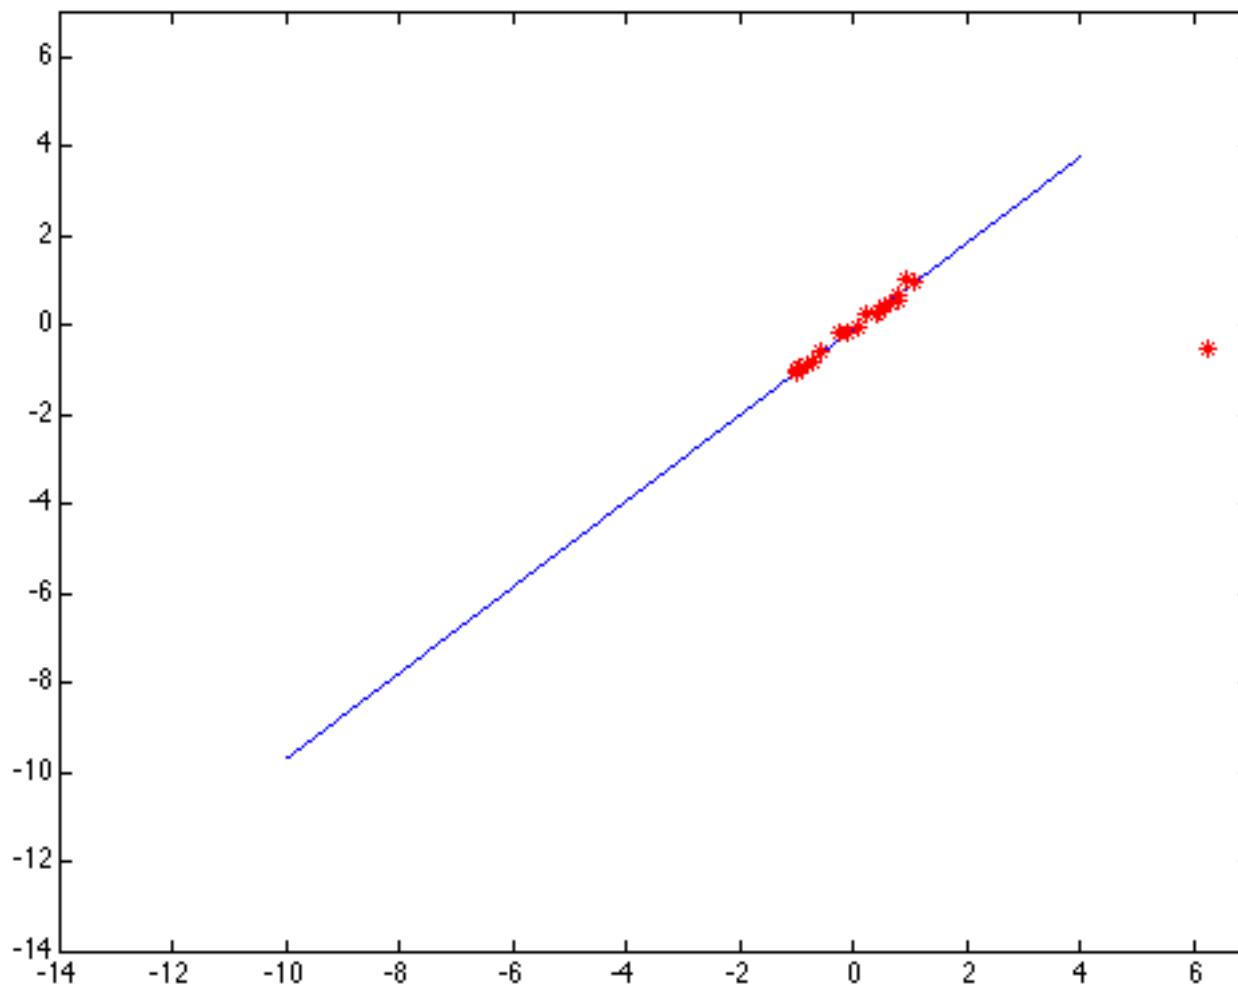
## The robust function $\rho$

- Favors a configuration with small residuals
- Constant penalty for large residuals

# Robust Estimator

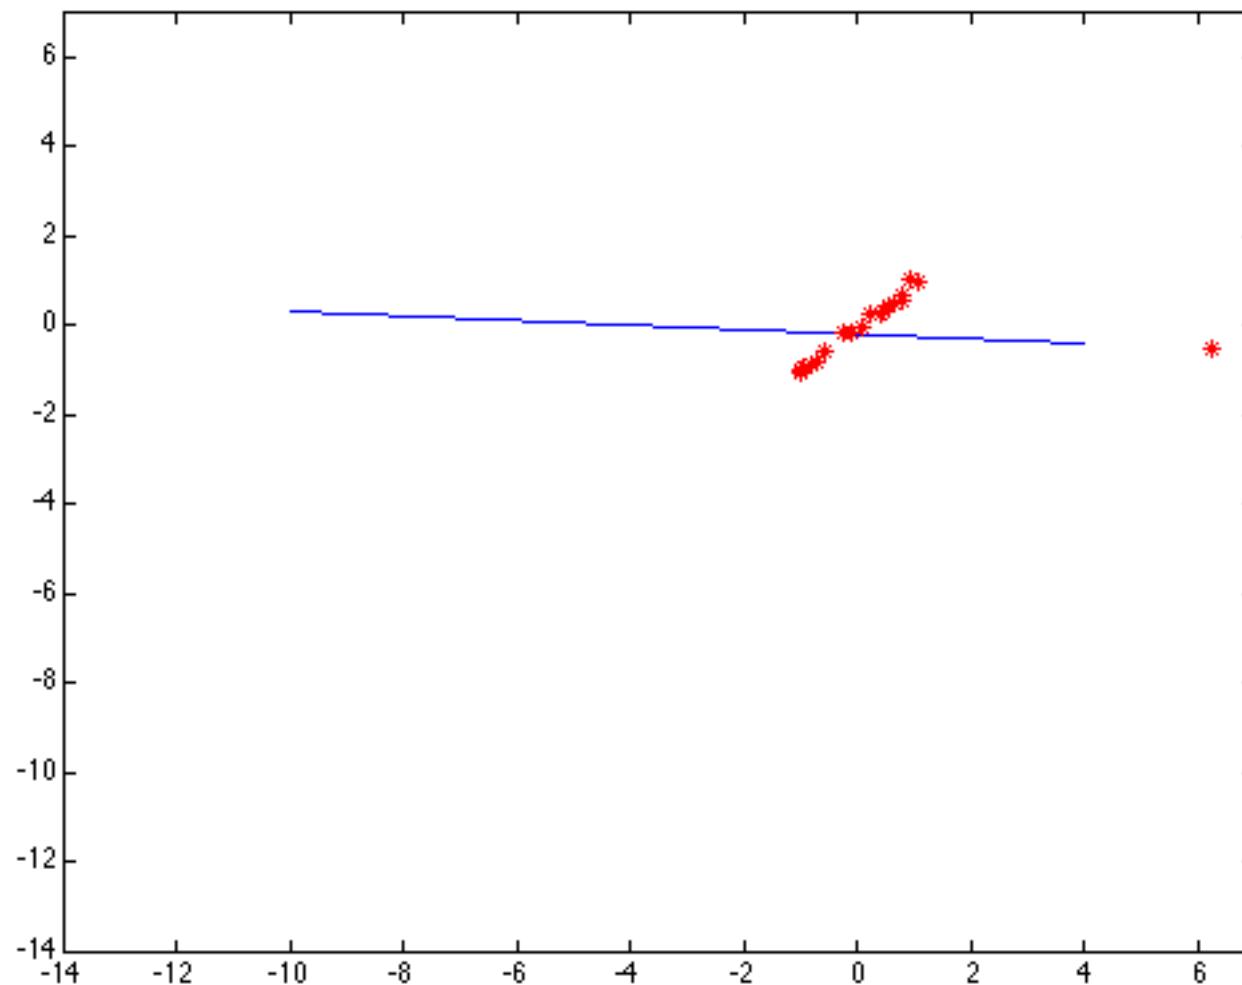
1. Initialize: e.g., choose  $\theta$  by least squares fit and  $\sigma = 1.5 \cdot \text{median}(\text{error})$
2. Choose params to minimize:  $\sum_i \frac{\text{error}(\theta, \text{data}_i)^2}{\sigma^2 + \text{error}(\theta, \text{data}_i)^2}$ 
  - E.g., numerical optimization
3. Compute new  $\sigma = 1.5 \cdot \text{median}(\text{error})$
4. Repeat (2) and (3) until convergence

# Choosing the scale: Just right



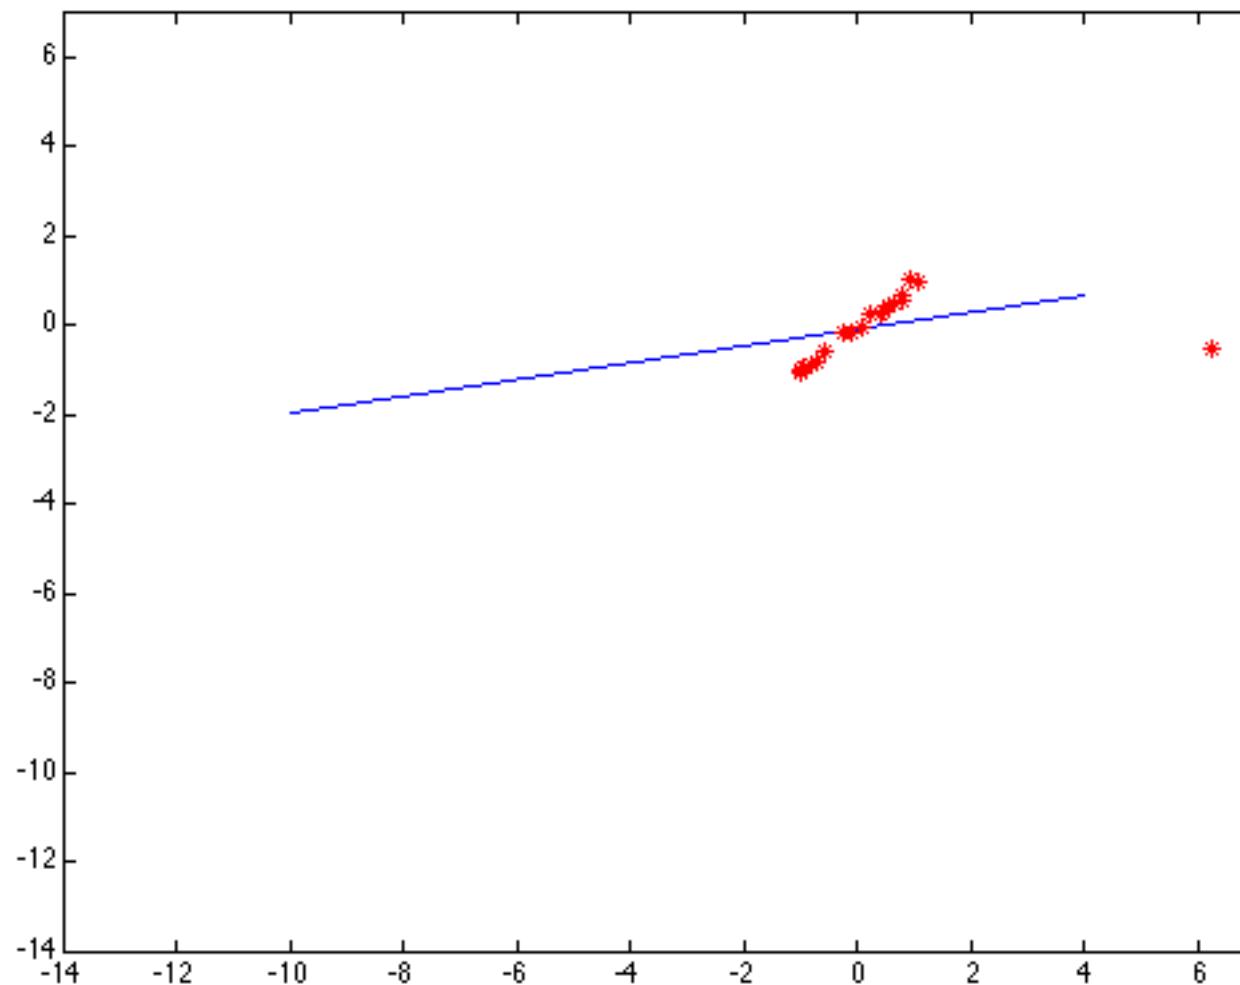
The effect of the outlier is minimized

# Choosing the scale: Too small



The error value is almost the same for every point and the fit is very poor

# Choosing the scale: Too large



Behaves much the same as least squares

# Fitting: Overview

- If we know which points belong to the line, how do we find the “optimal” line parameters?
  - Least squares
- What if there are outliers?
  - Robust fitting, RANSAC
- What if there are many lines?
  - Voting methods: RANSAC, Hough transform
- What if we’re not even sure it’s a line?
  - Model selection (not covered)

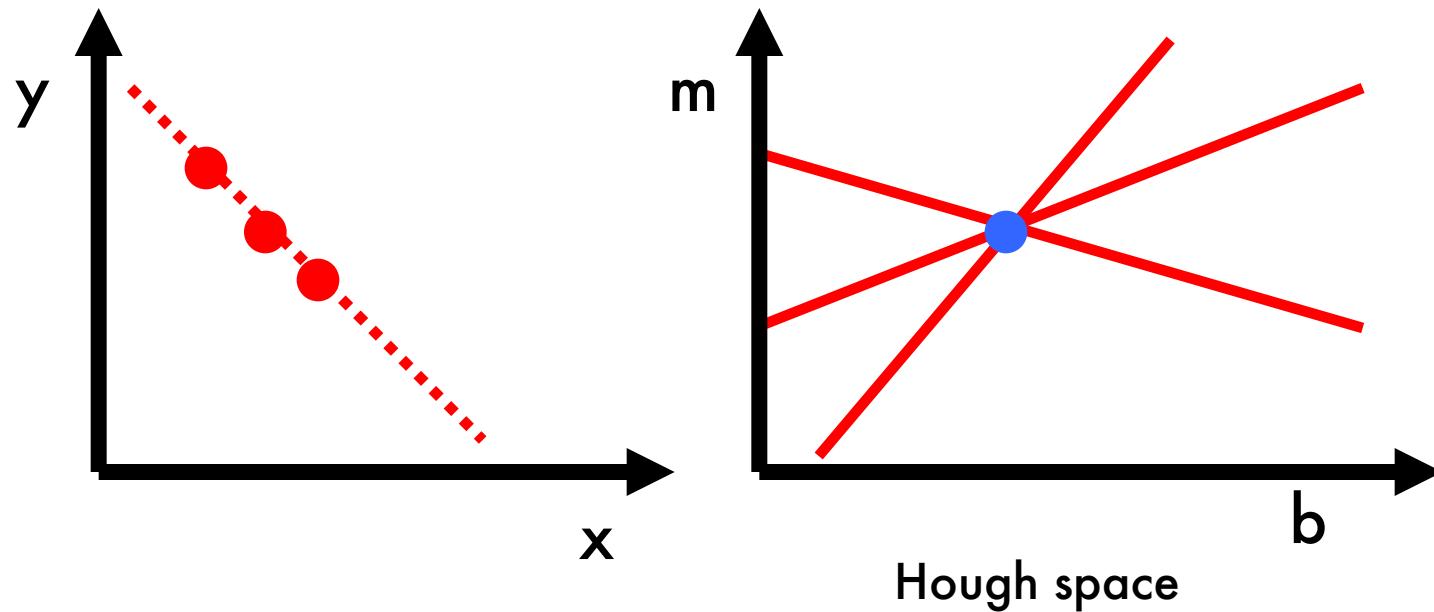
# Hough Transform: Outline

1. Create a grid of parameter values
2. Each point votes for a set of parameters, incrementing those values in grid
3. Find maximum or local maxima in grid

# Hough transform

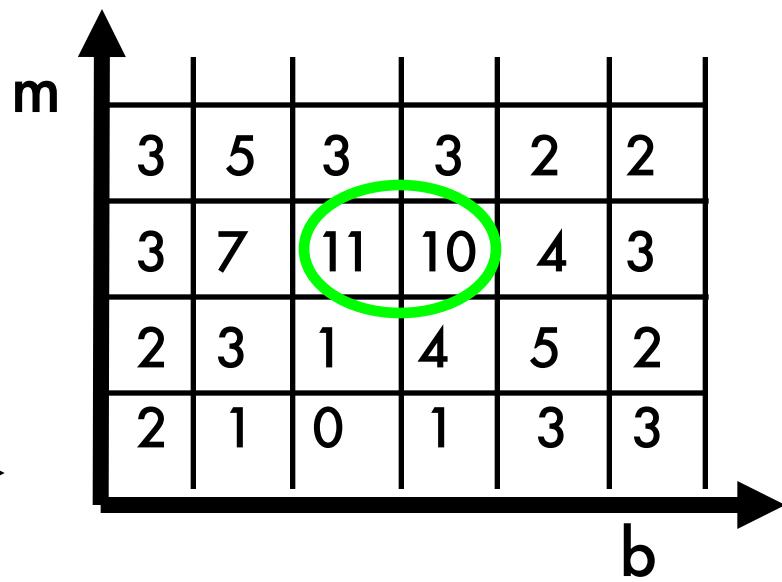
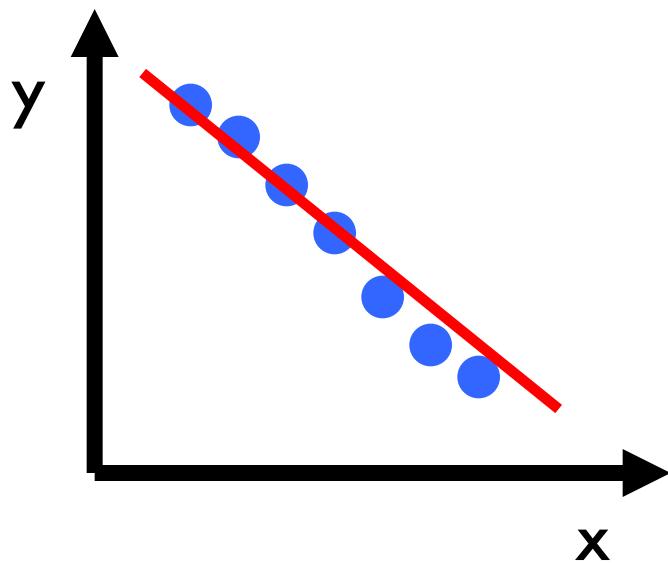
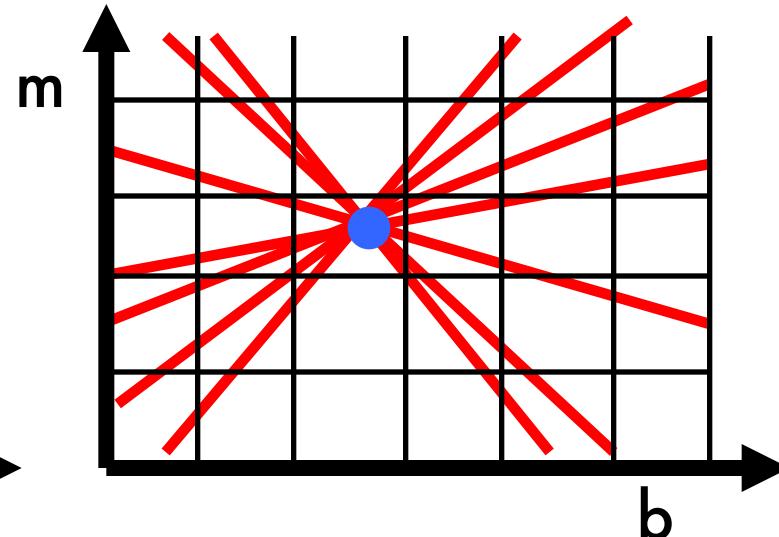
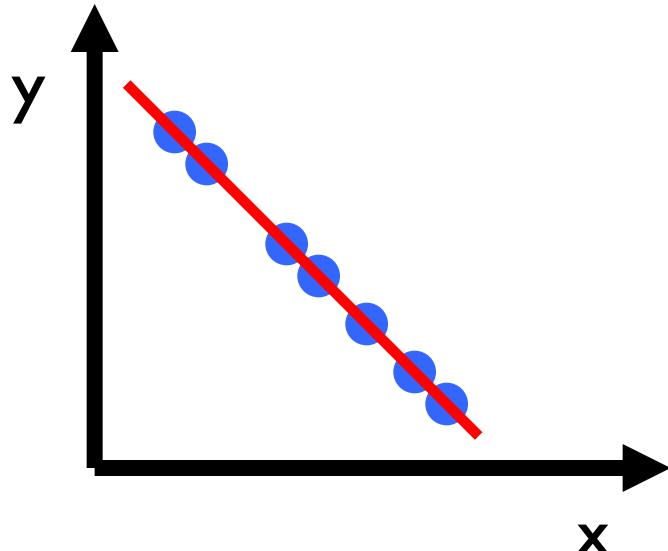
P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*, Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959

Given a set of points, find the curve or line that explains the data points best



$$y = m x + b$$

# Hough transform

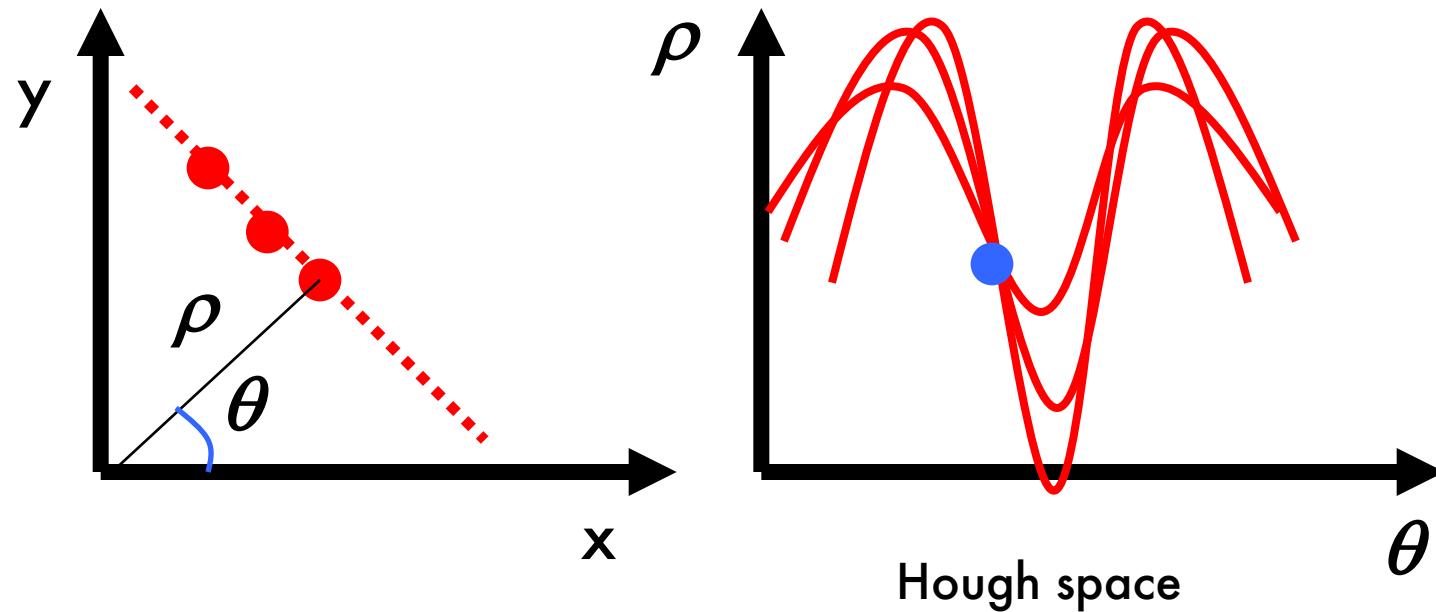


# Hough transform

P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*, Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959

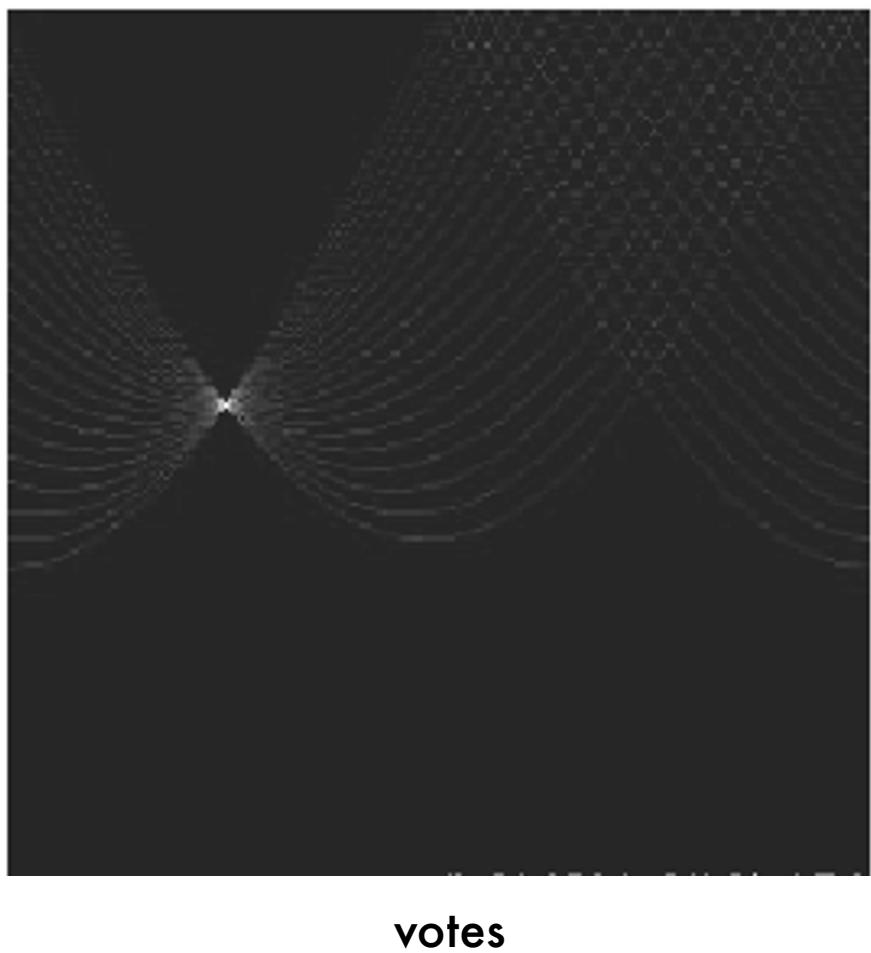
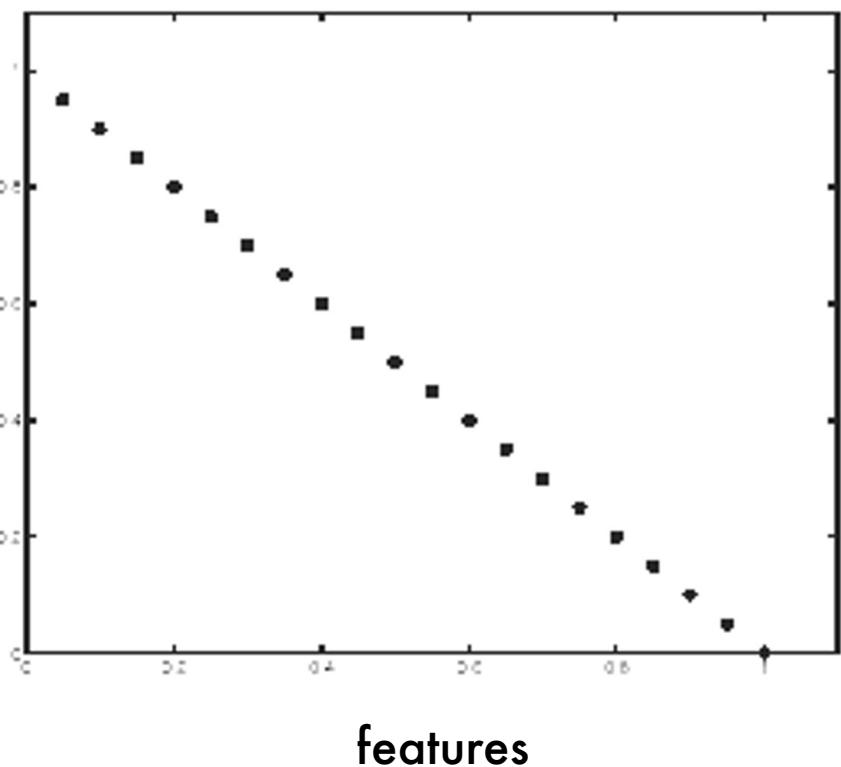
Issue : parameter space  $[m,b]$  is unbounded...

Use a polar representation for the parameter space



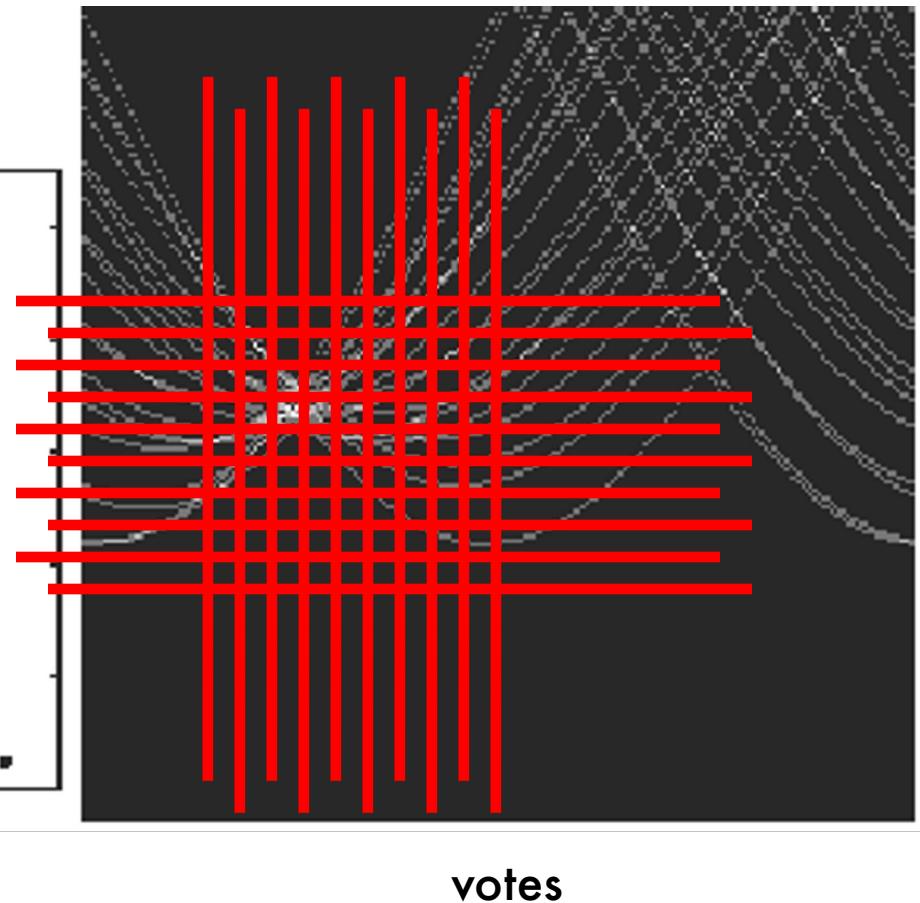
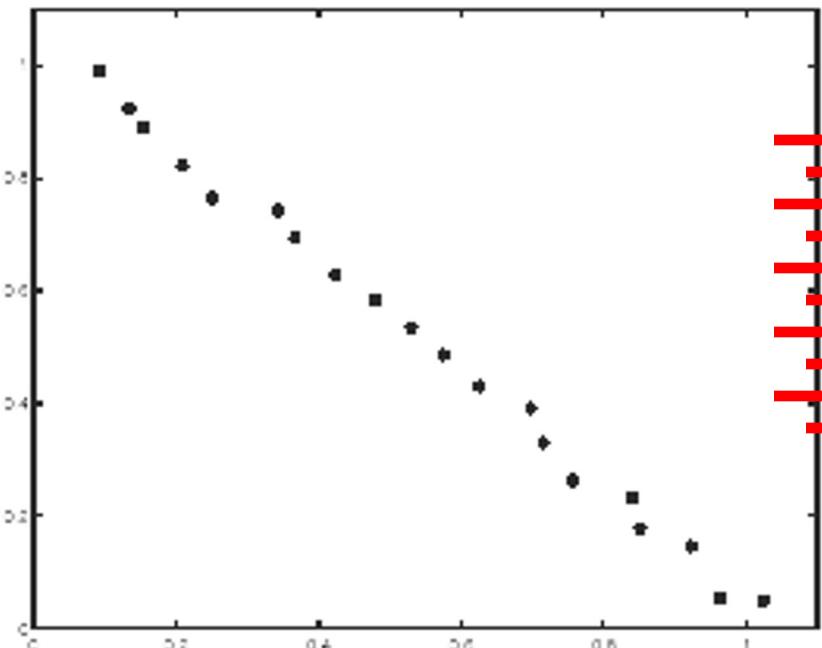
$$x \cos \theta + y \sin \theta = \rho$$

# Hough transform - experiments



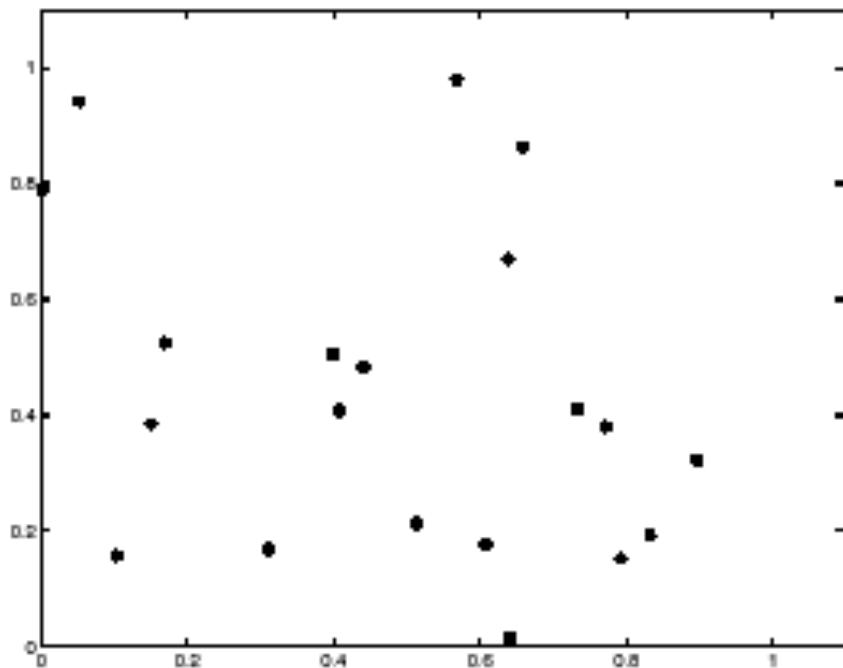
# Hough transform - experiments

Noisy data

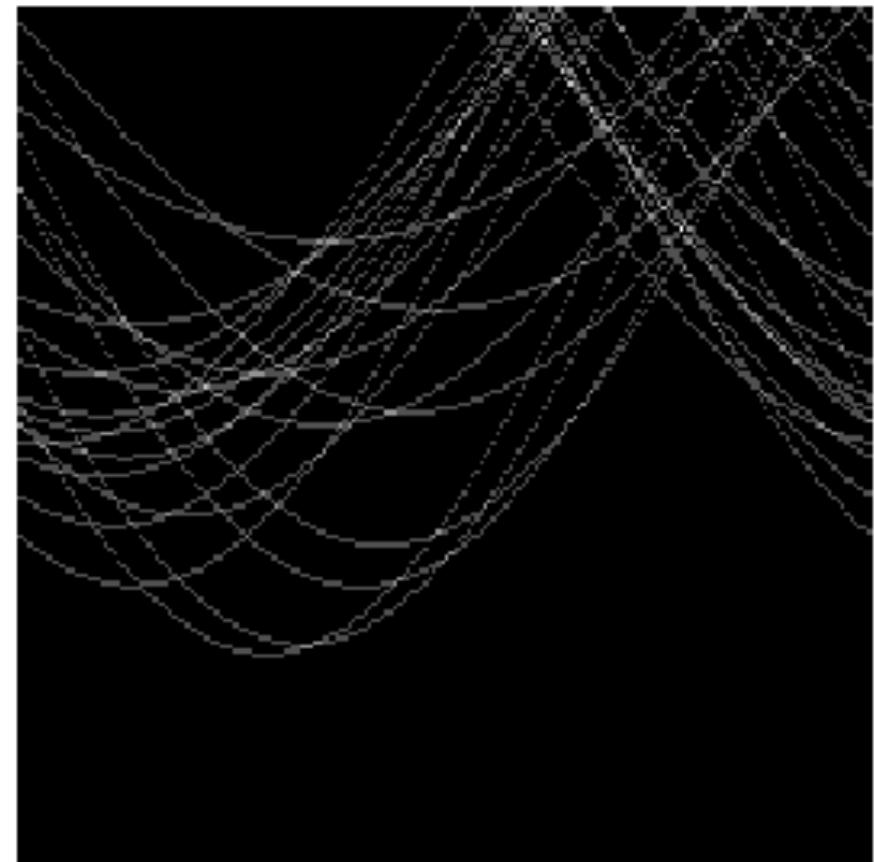


Need to adjust grid size or smooth

# Hough transform - experiments



features



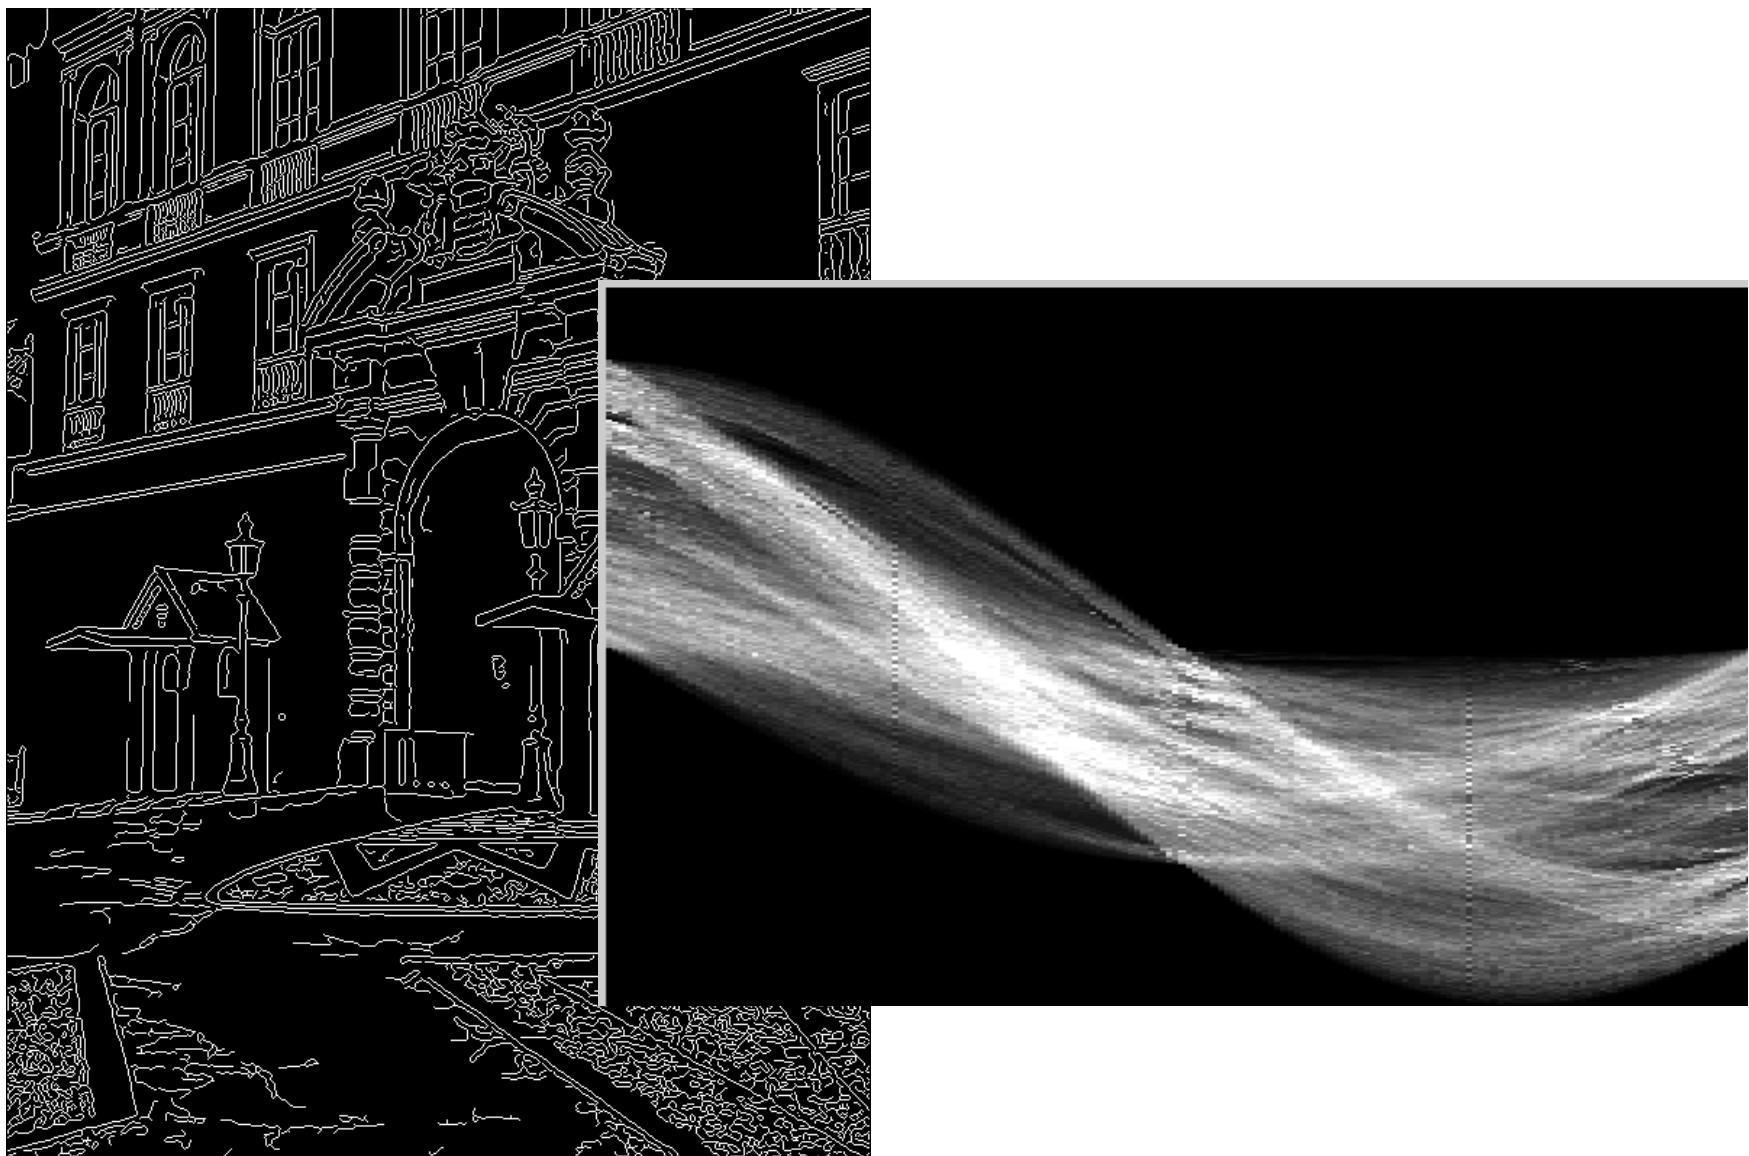
votes

**Issue: spurious peaks due to uniform noise**

# 1. Image → Canny

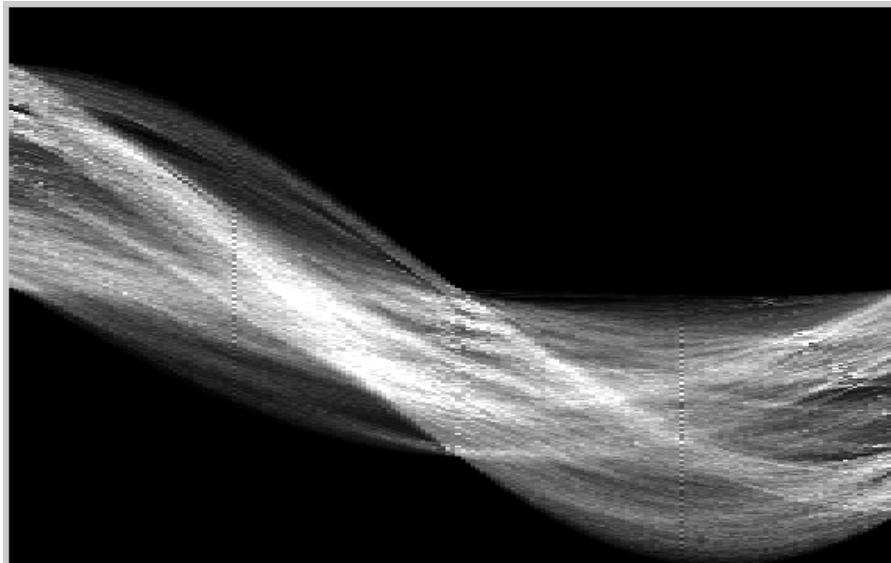


## 2. Canny → Hough votes

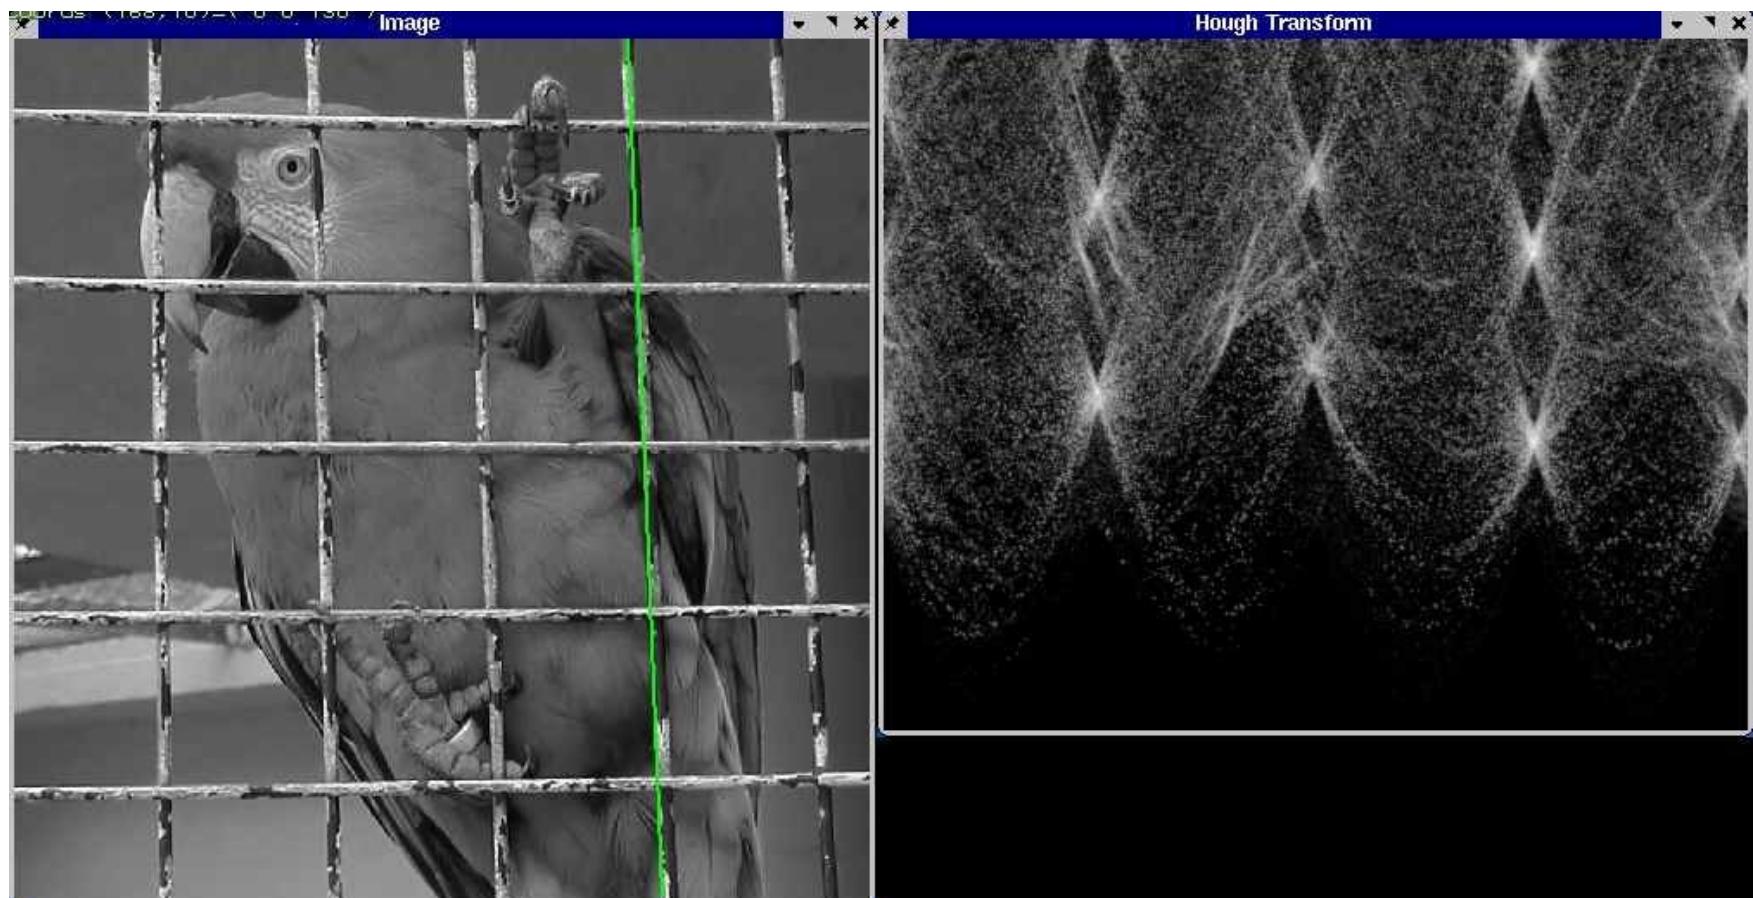


### 3. Hough votes → Edges

Find peaks and post-process



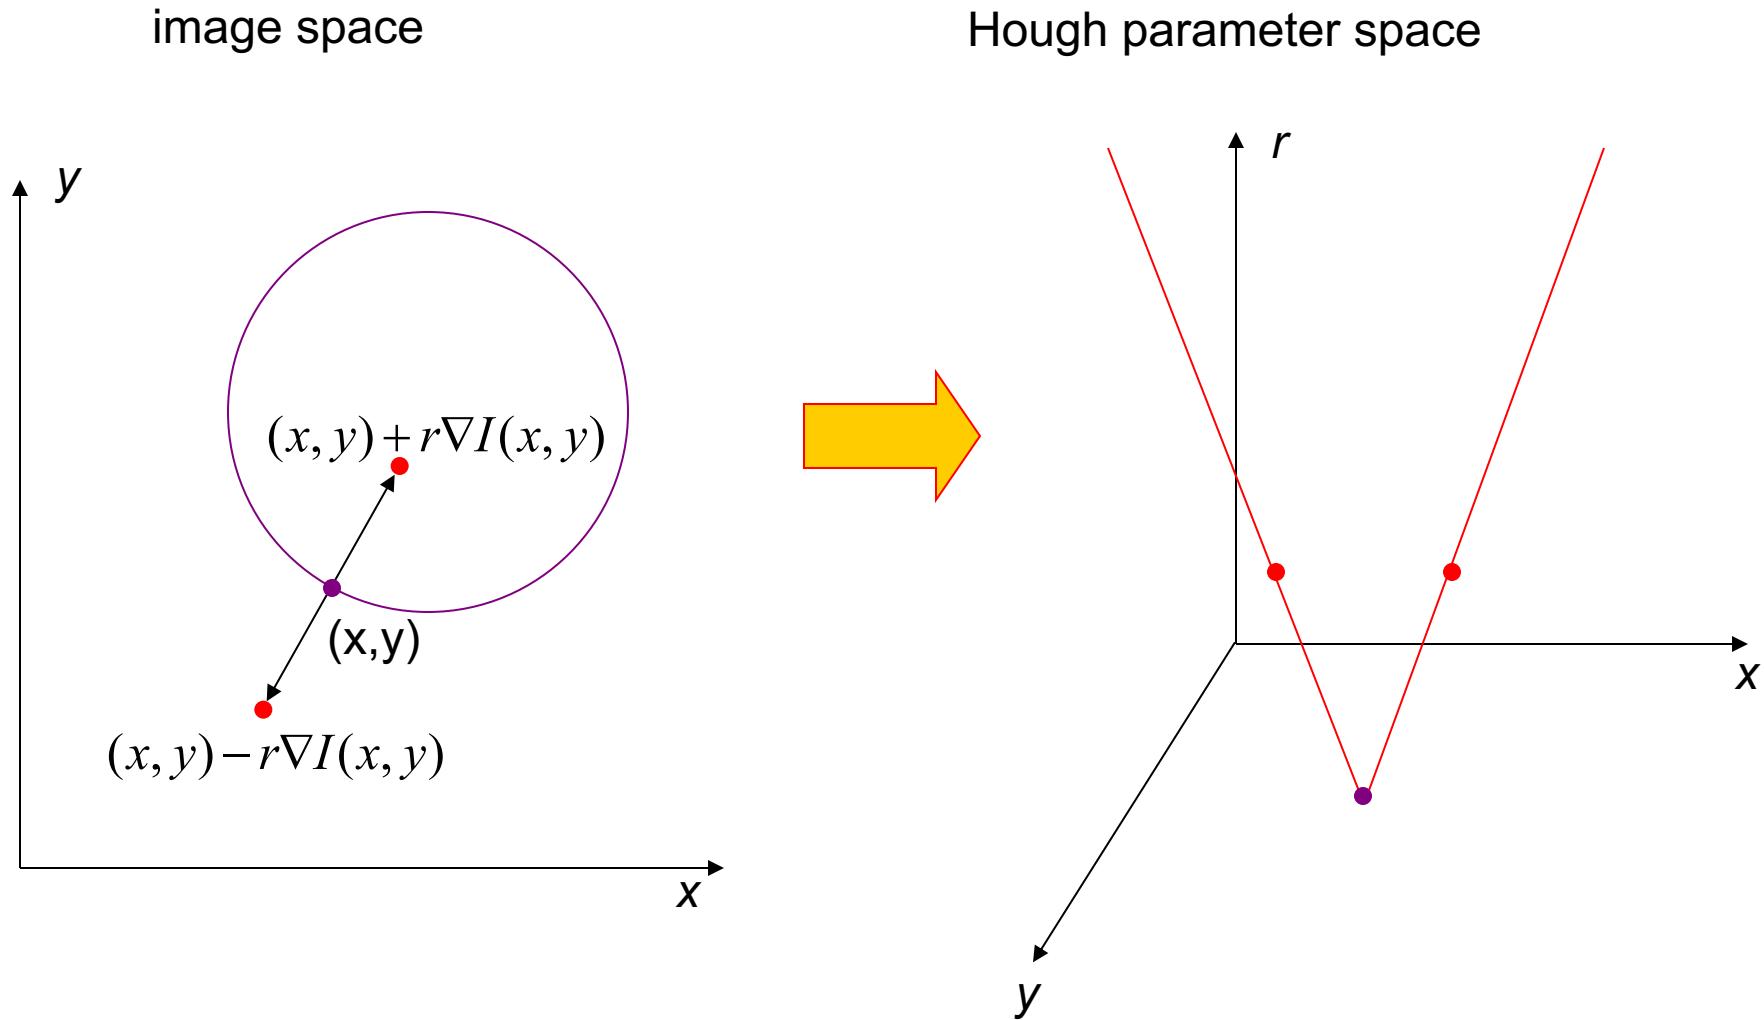
# Hough transform example



# Finding circles $(x_0, y_0, r)$ using Hough transform

- Fixed  $r$
- Variable  $r$

# Hough transform for circles

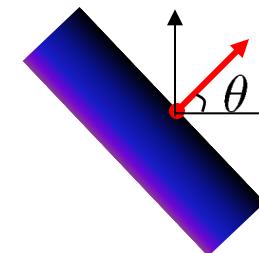


# Incorporating image gradients

- When we detect an edge point, we also know its gradient orientation
- How does this constrain possible lines passing through the point?
- Modified Hough transform:

- For each edge point  $(x, y)$   
 $\theta = \text{gradient orientation at } (x, y)$   
 $\rho = x \cos \theta + y \sin \theta$   
 $H(\theta, \rho) = H(\theta, \rho) + 1$

end



$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$
$$\theta = \tan^{-1} \left( \frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

# Hough transform conclusions

## Good

- Robust to outliers: each point votes separately
- Fairly efficient (much faster than trying all sets of parameters)
- Provides multiple good fits

## Bad

- Some sensitivity to noise
- Bin size trades off between noise tolerance, precision, and speed/memory
  - Can be hard to find sweet spot
- Not suitable for more than a few parameters
  - grid size grows exponentially

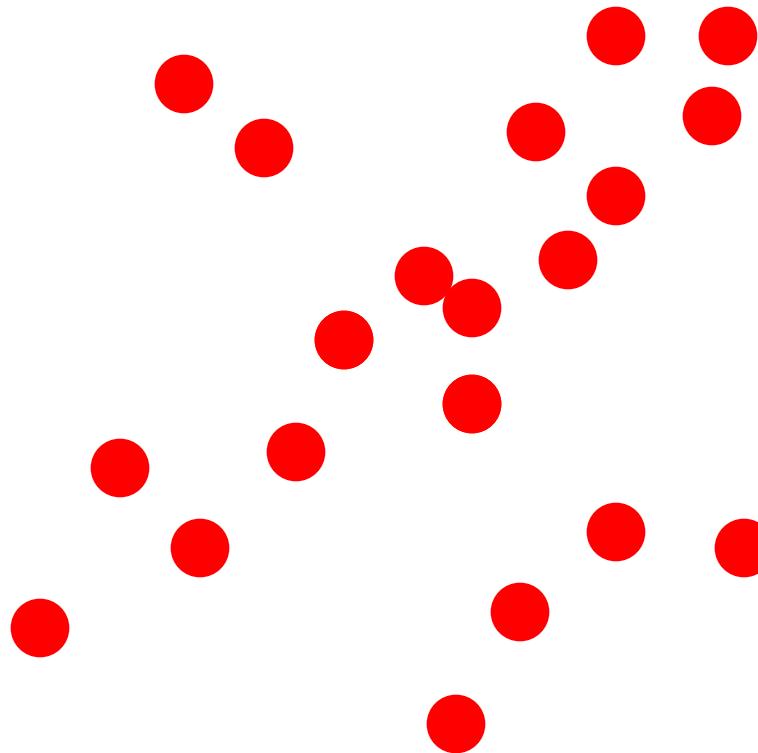
## Common applications

- Line fitting (also circles, ellipses, etc.)
- Object instance recognition (parameters are position/scale/orientation)
- Object category recognition (parameters are position/scale)

# RANSAC

(RANdom SAmples Consensus) :

Fischler & Bolles in '81.



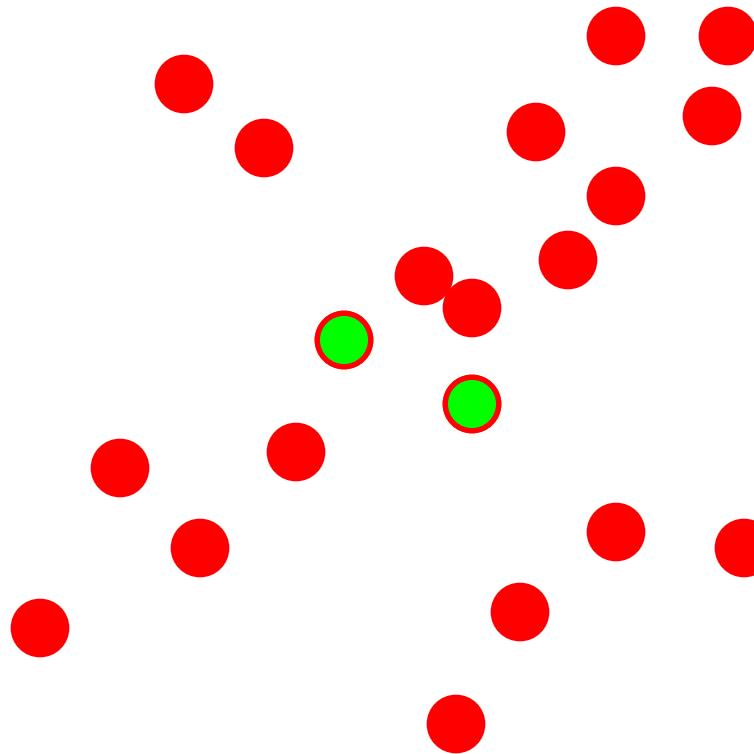
## Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# RANSAC

Line fitting example



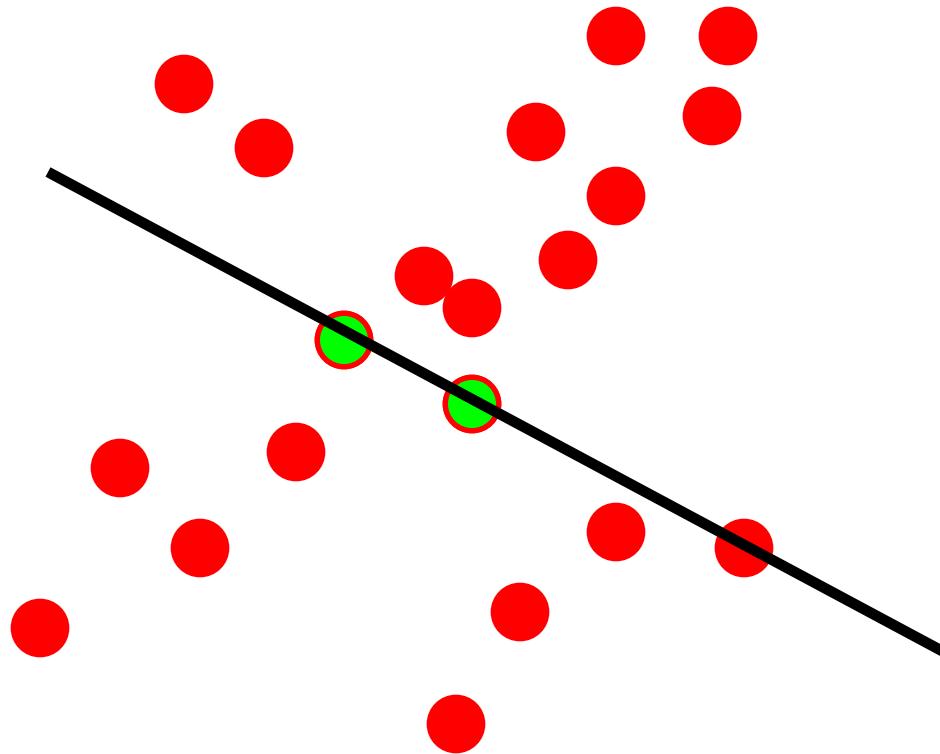
Algorithm:

1. **Sample** (randomly) the number of points required to fit the model (#=2)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# RANSAC

Line fitting example



Algorithm:

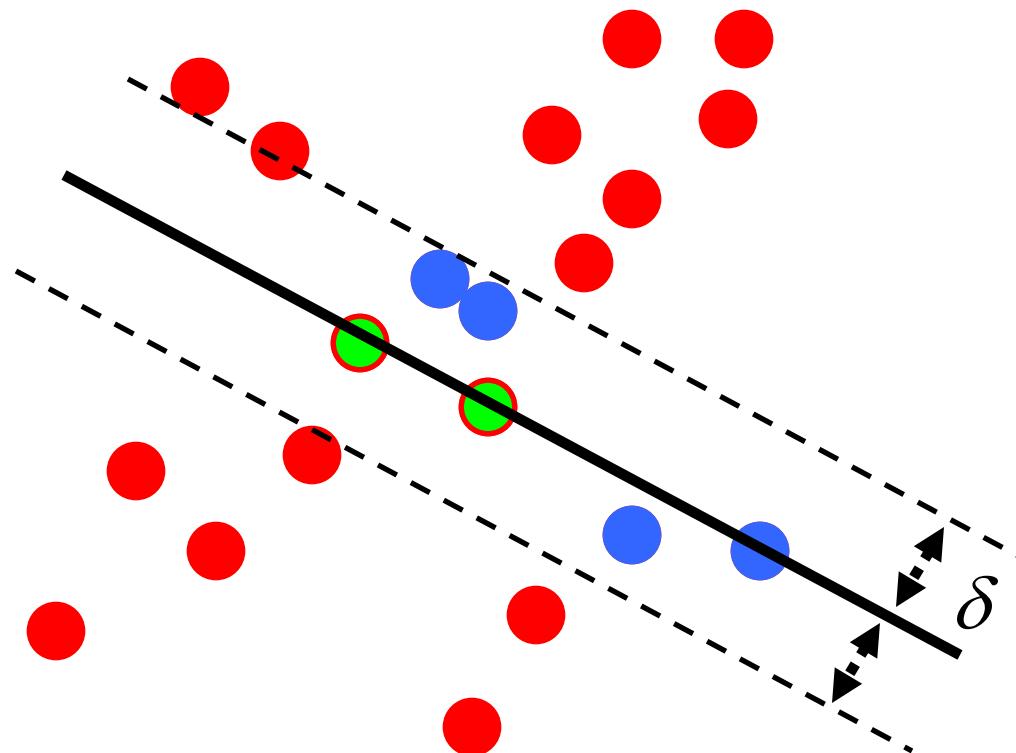
1. **Sample** (randomly) the number of points required to fit the model (#=2)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# RANSAC

Line fitting example

$$N_I = 6$$

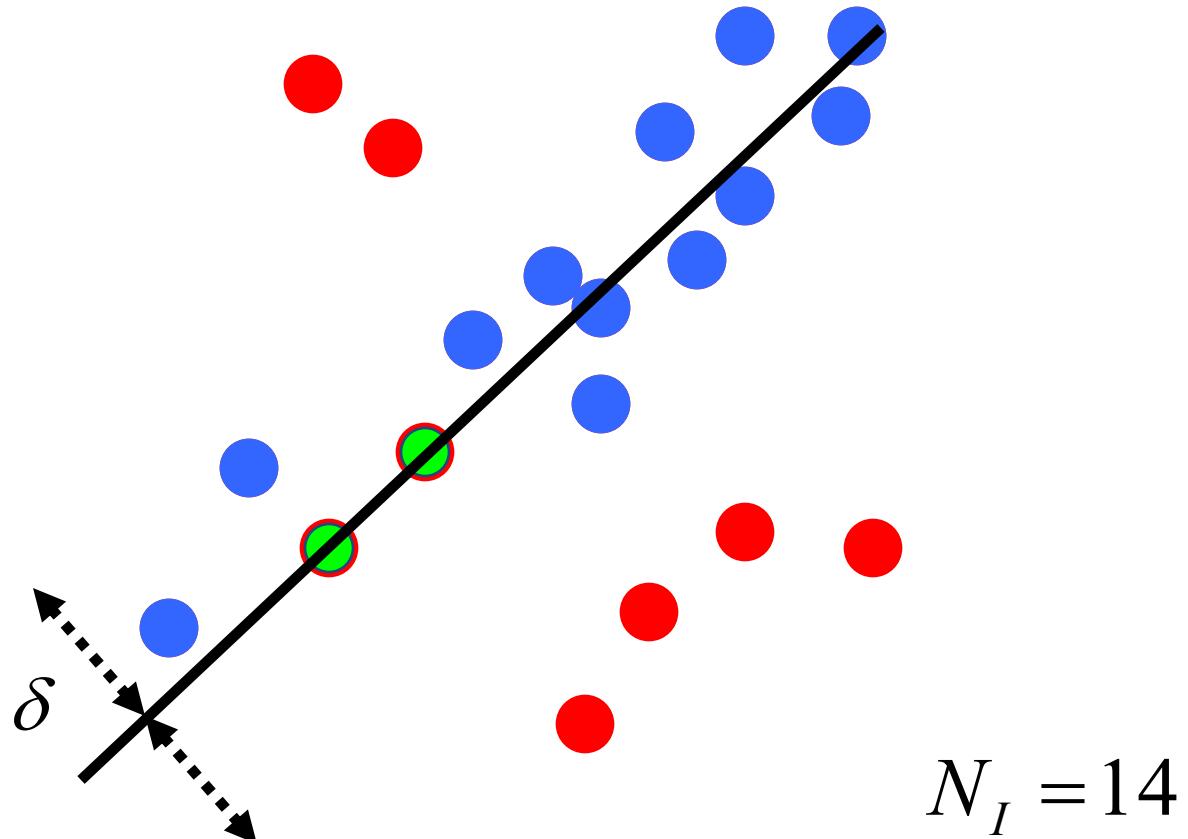


Algorithm:

1. **Sample** (randomly) the number of points required to fit the model (#=2)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# RANSAC



Algorithm:

1. **Sample** (randomly) the number of points required to fit the model (#=2)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

**Repeat** 1-3 until the best model is found with high confidence

# How to choose parameters?

- Number of samples  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p=0.99$ ) (outlier ratio:  $e$ )
- Number of sampled points  $s$ 
  - Minimum number needed to fit the model
- Distance threshold  $\delta$ 
  - Choose  $\delta$  so that a good point with noise is likely (e.g., prob=0.95) within threshold
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2=3.84\sigma^2$

$$N = \log(1-p)/\log\left(1-(1-e)^s\right)$$

s	proportion of outliers $e$							
	5%	10%	20%	25%	30%	40%	50%	
2	2	3	5	6	7	11	17	
3	3	4	7	9	11	19	35	
4	3	5	9	13	17	34	72	
5	4	6	12	17	26	57	146	
6	4	7	16	24	37	97	293	
7	4	8	20	33	54	163	588	
8	5	9	26	44	78	272	1177	

# RANSAC conclusions

## Good

- Robust to outliers
- Applicable for larger number of objective function parameters than Hough transform
- Optimization parameters are easier to choose than Hough transform

## Bad

- Computational time grows quickly with fraction of outliers and number of parameters
- Not as good for getting multiple fits (though one solution is to remove inliers after each fit and repeat)

## Common applications

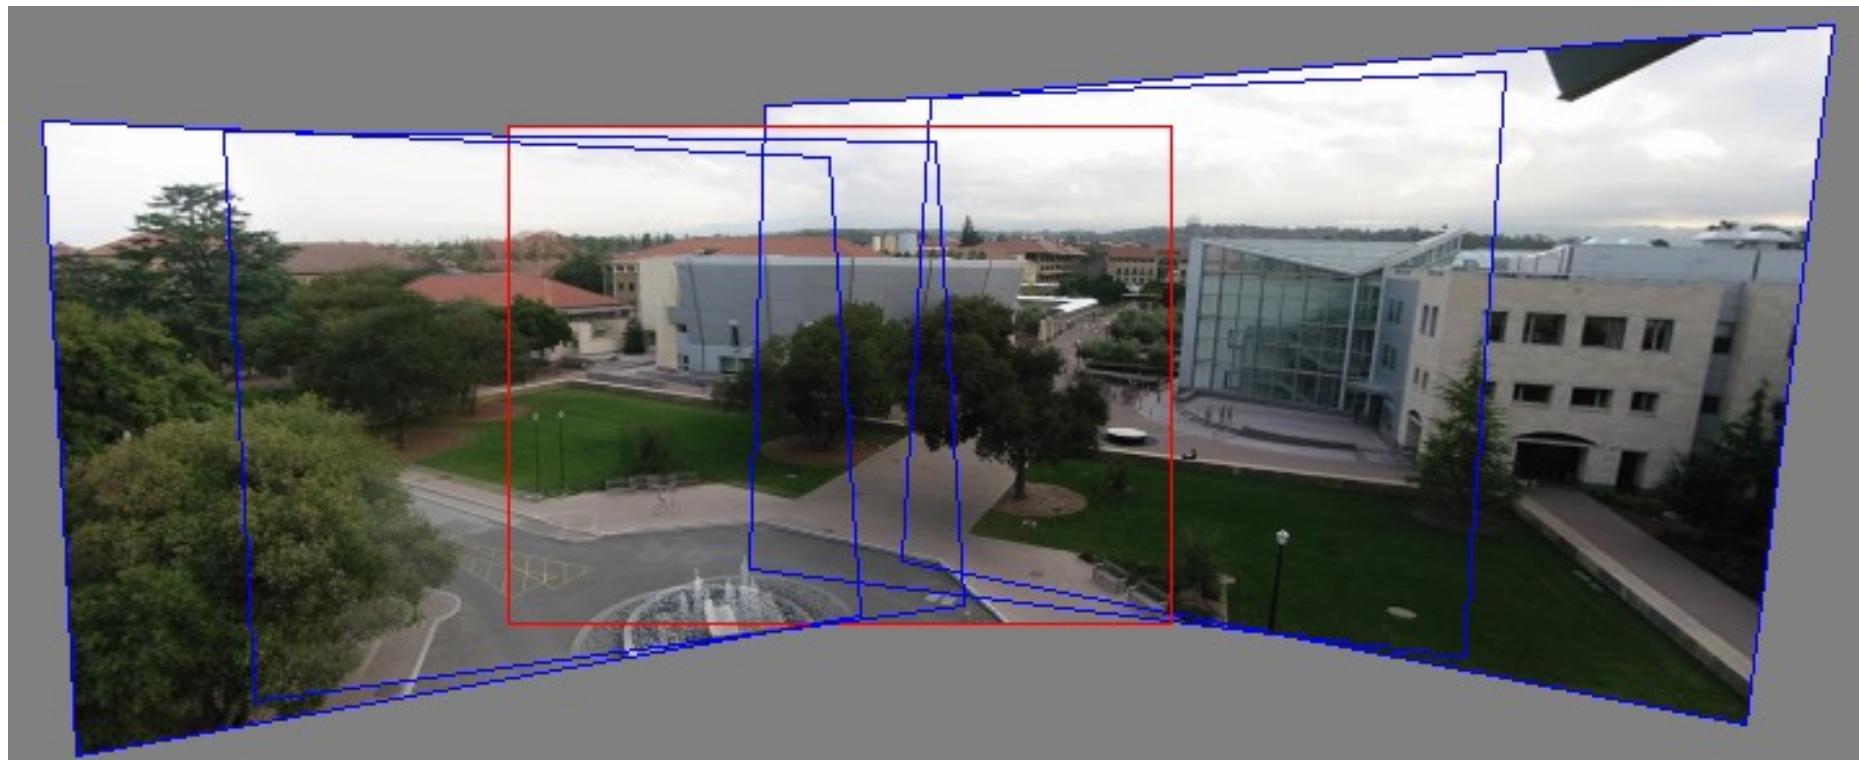
- Computing a homography (e.g., image stitching)
- Estimating fundamental matrix (relating two views)

# Fitting Summary

- Least Squares Fit
  - closed form solution
  - robust to noise
  - not robust to outliers
- Robust Least Squares
  - improves robustness to noise
  - requires iterative optimization
- Hough transform
  - robust to noise and outliers
  - can fit multiple models
  - only works for a few parameters (1-4 typically)
- RANSAC
  - robust to noise and outliers
  - works with a moderate number of parameters (e.g, 1-8)

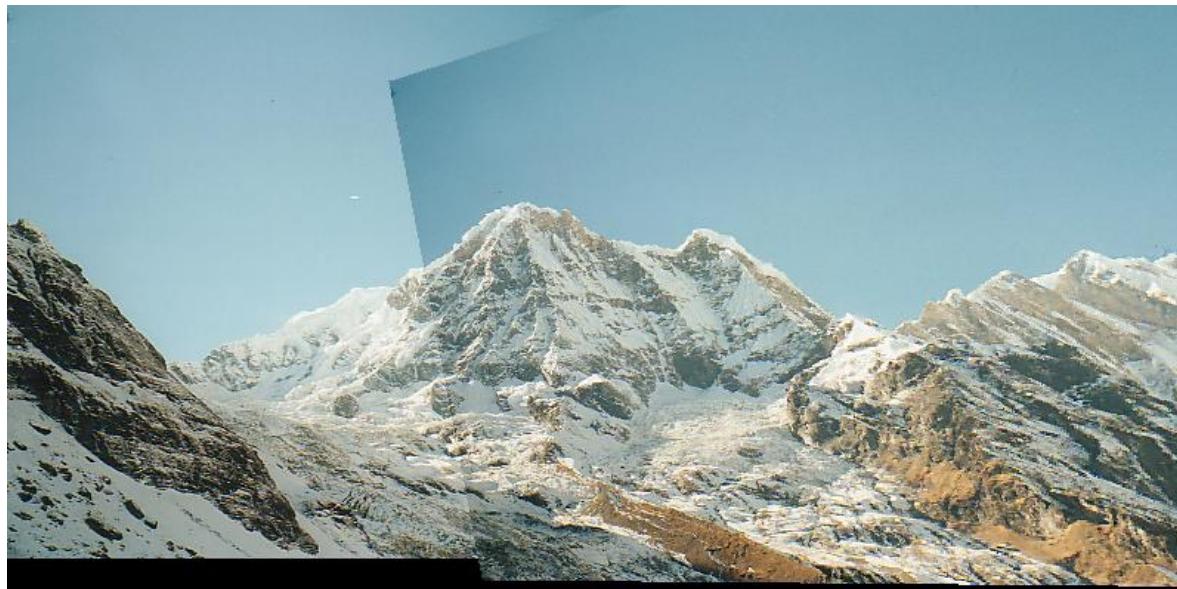
# Image alignment

---



# Image Alignment

---



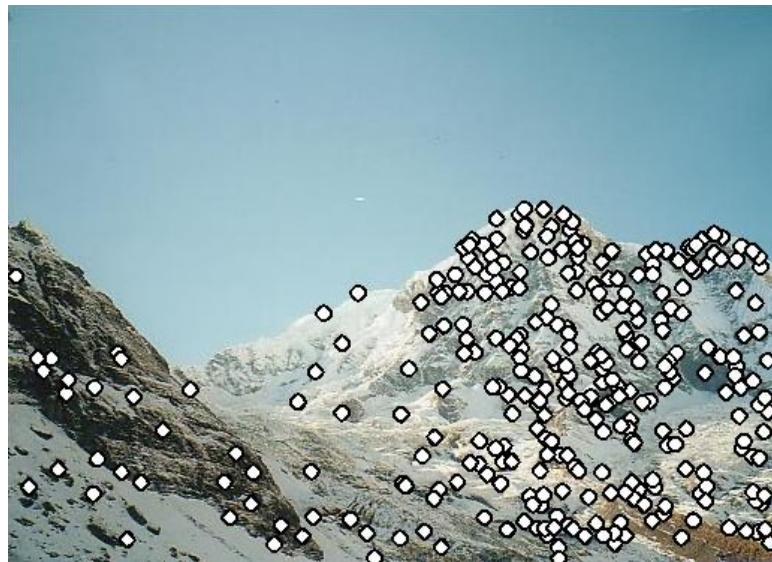
# Robust feature-based alignment

---



# Robust feature-based alignment

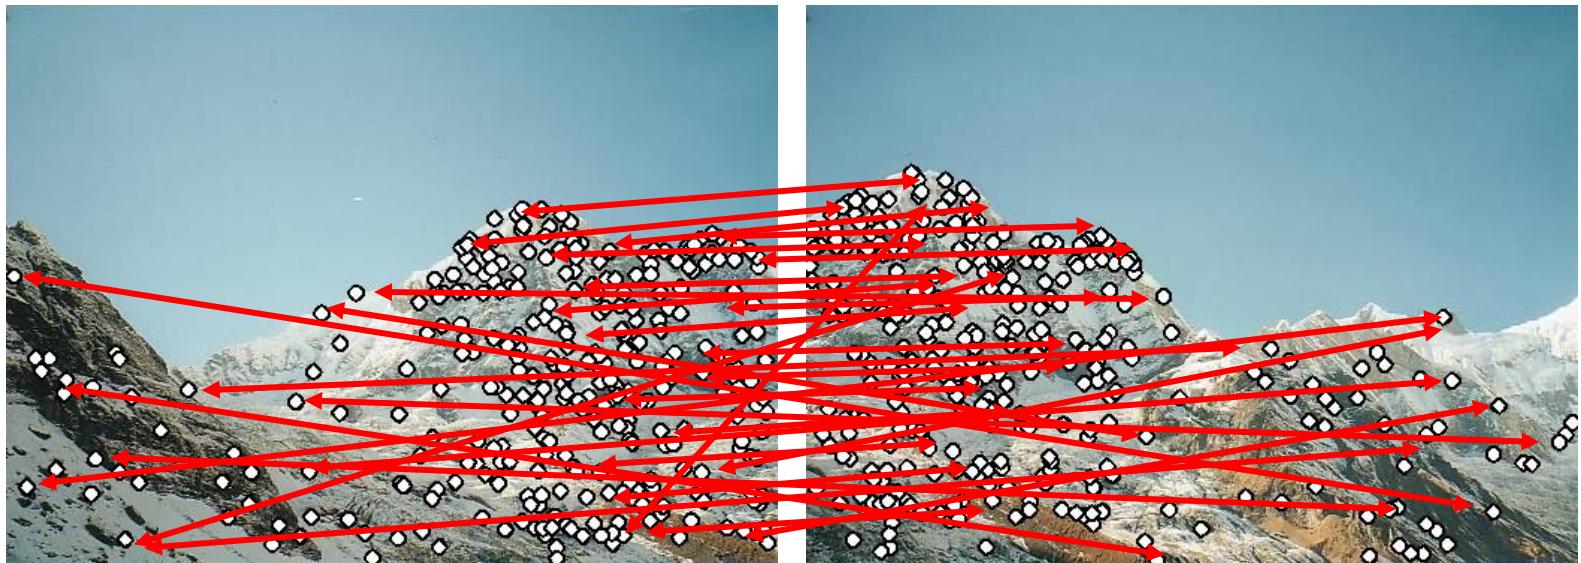
---



- Extract features

# Robust feature-based alignment

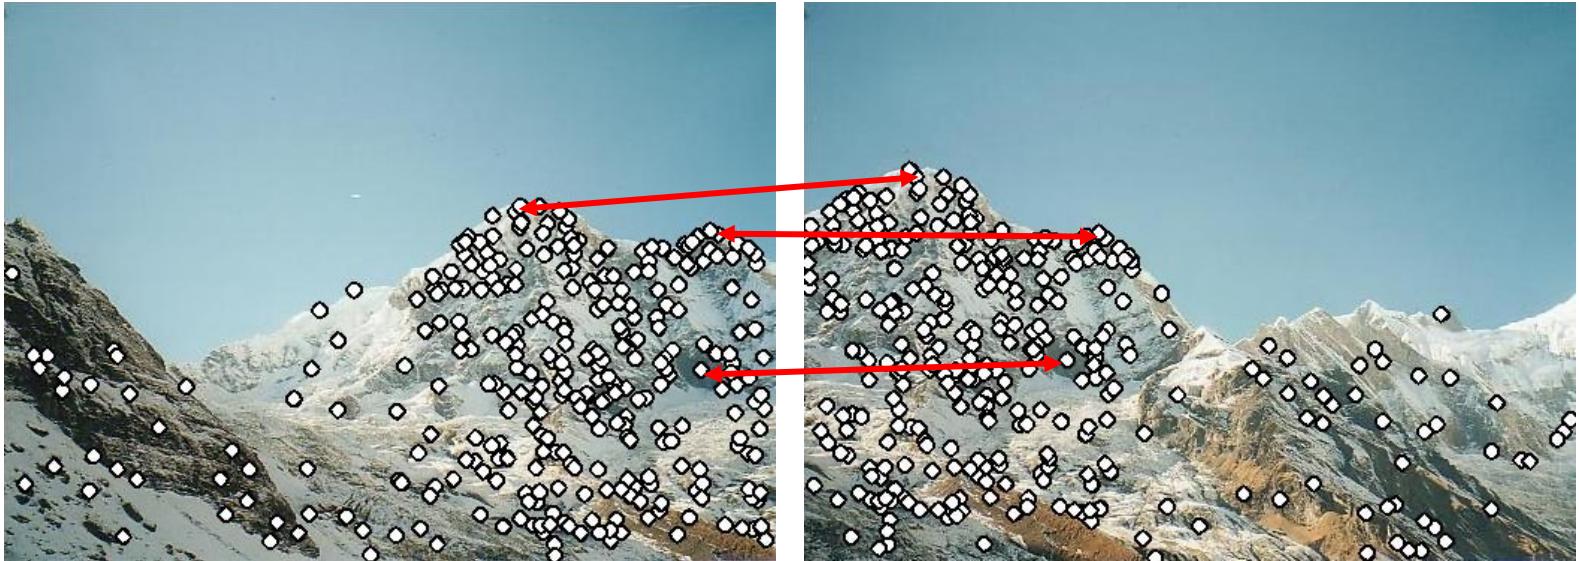
---



- Extract features
- Compute *putative matches*

# Robust feature-based alignment

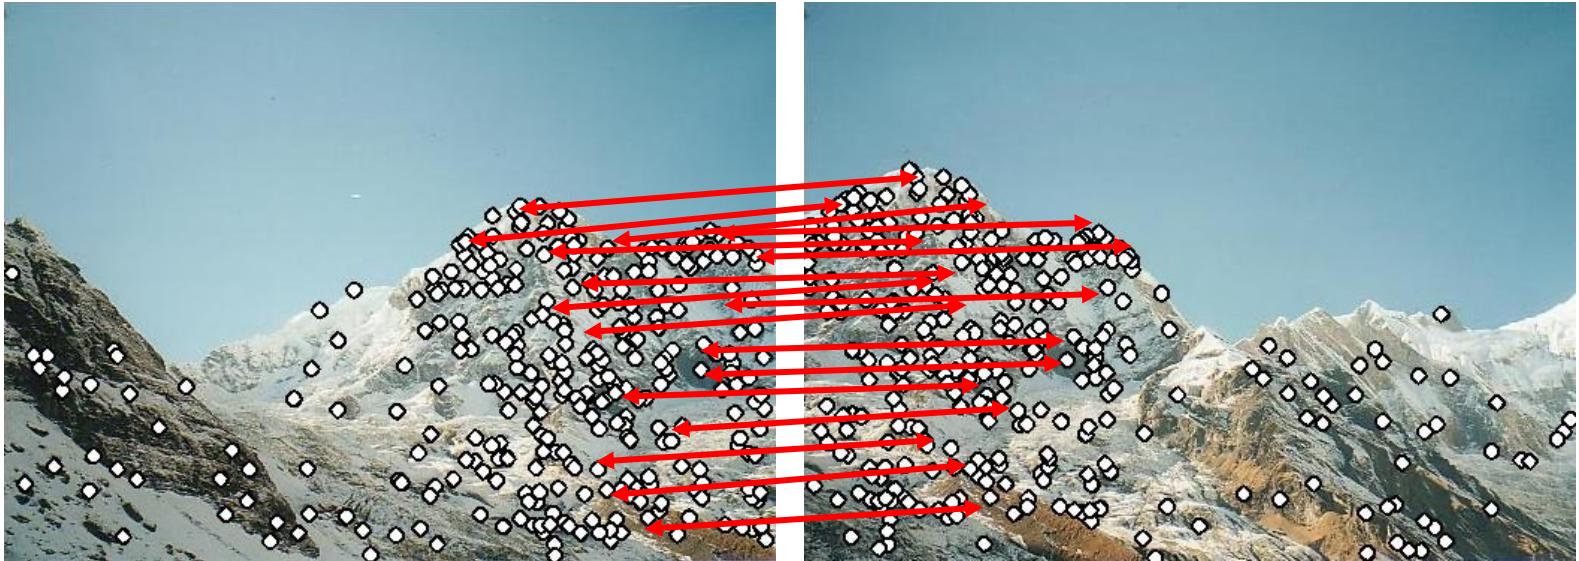
---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$

# Robust feature-based alignment

---



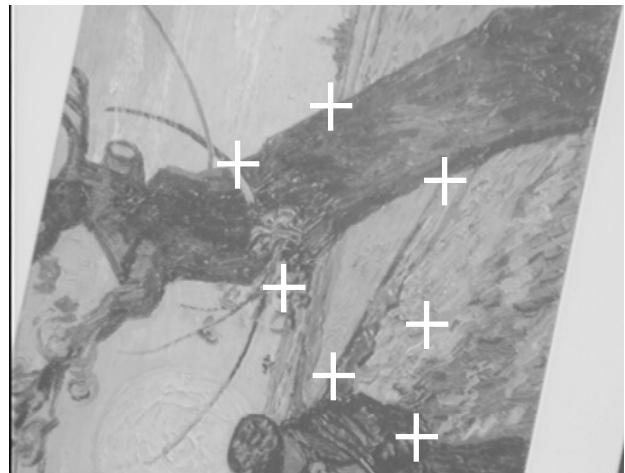
- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )

# Generating putative correspondences

---

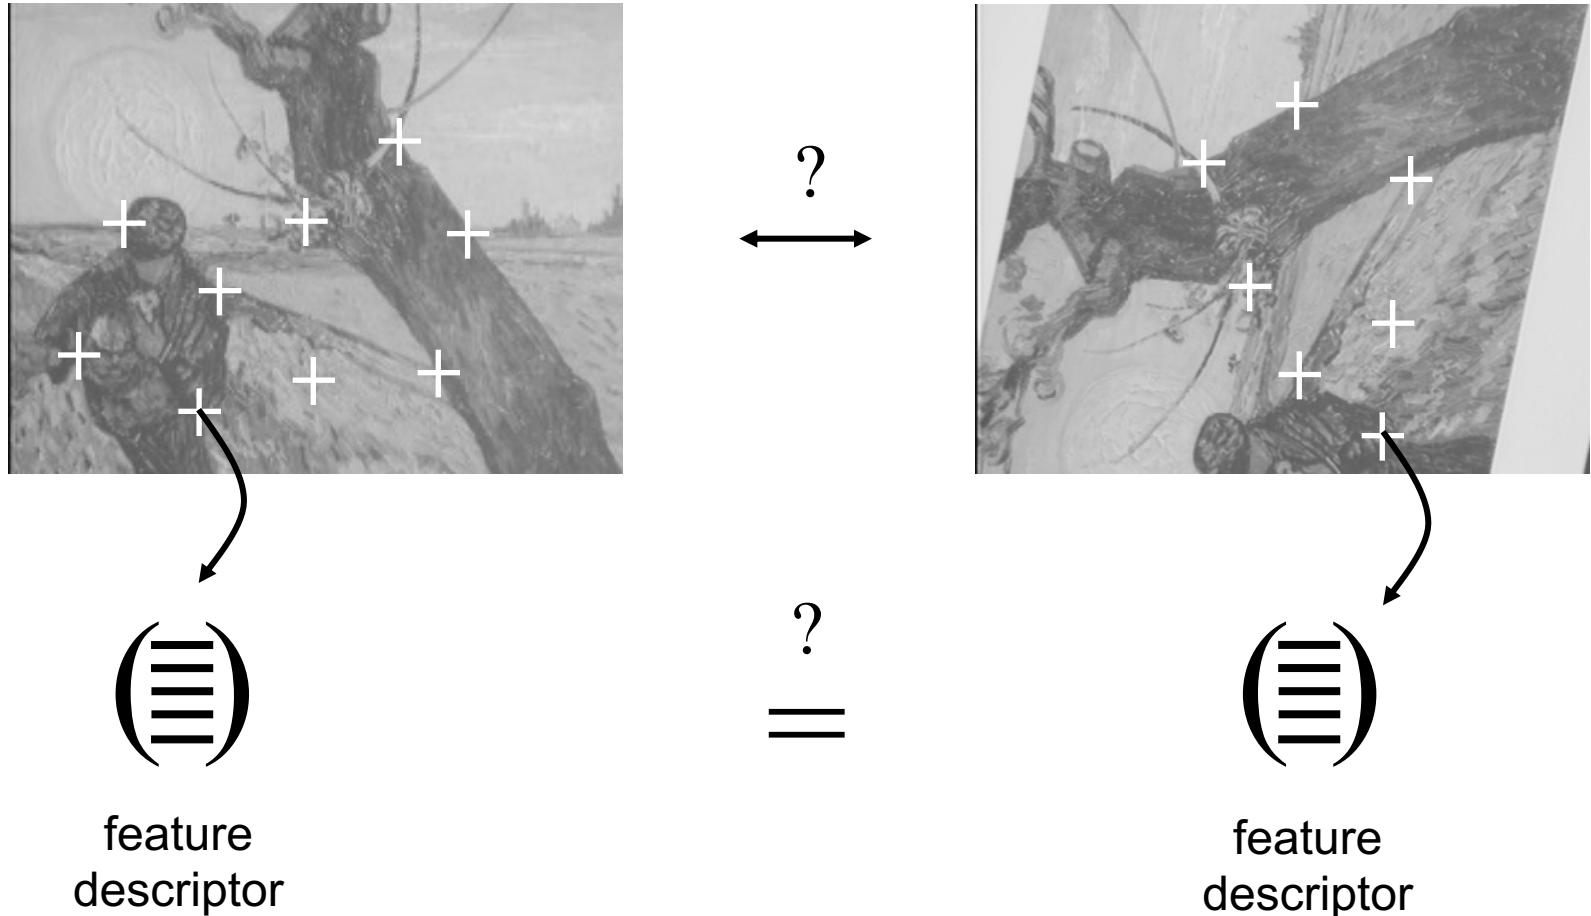


?  
↔



# Generating putative correspondences

---

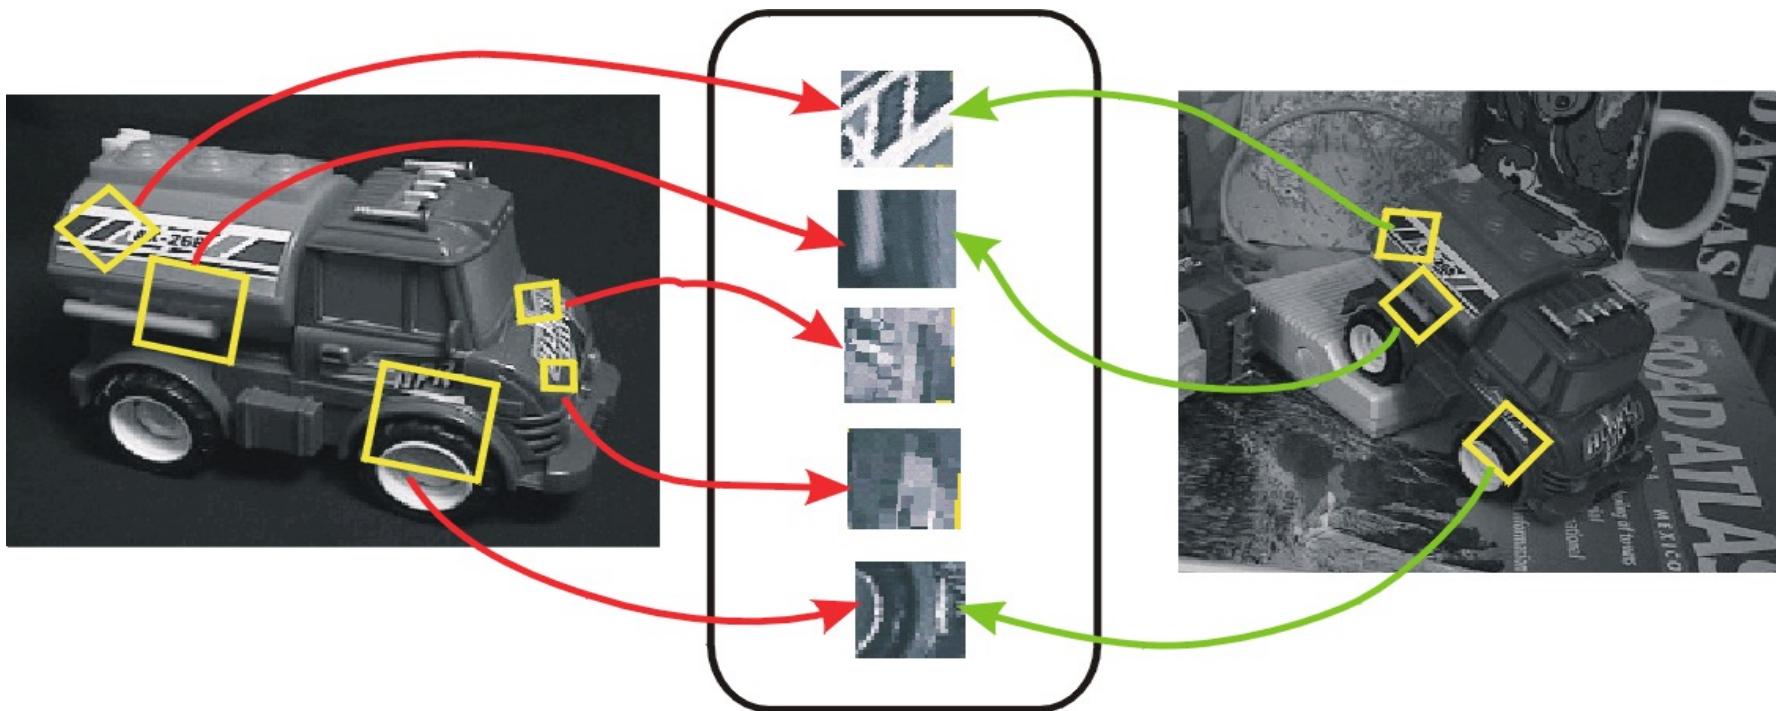


- Need to compare *feature descriptors* of local patches surrounding interest points

# Feature descriptors

---

- Recall: feature detection and description



# Feature descriptors

---

- Simplest descriptor: vector of raw intensity values
- How to compare two such vectors?
  - Sum of squared differences (SSD)

$$\text{SSD}(\mathbf{u}, \mathbf{v}) = \sum_i (u_i - v_i)^2$$

- Not invariant to intensity change
- Normalized correlation

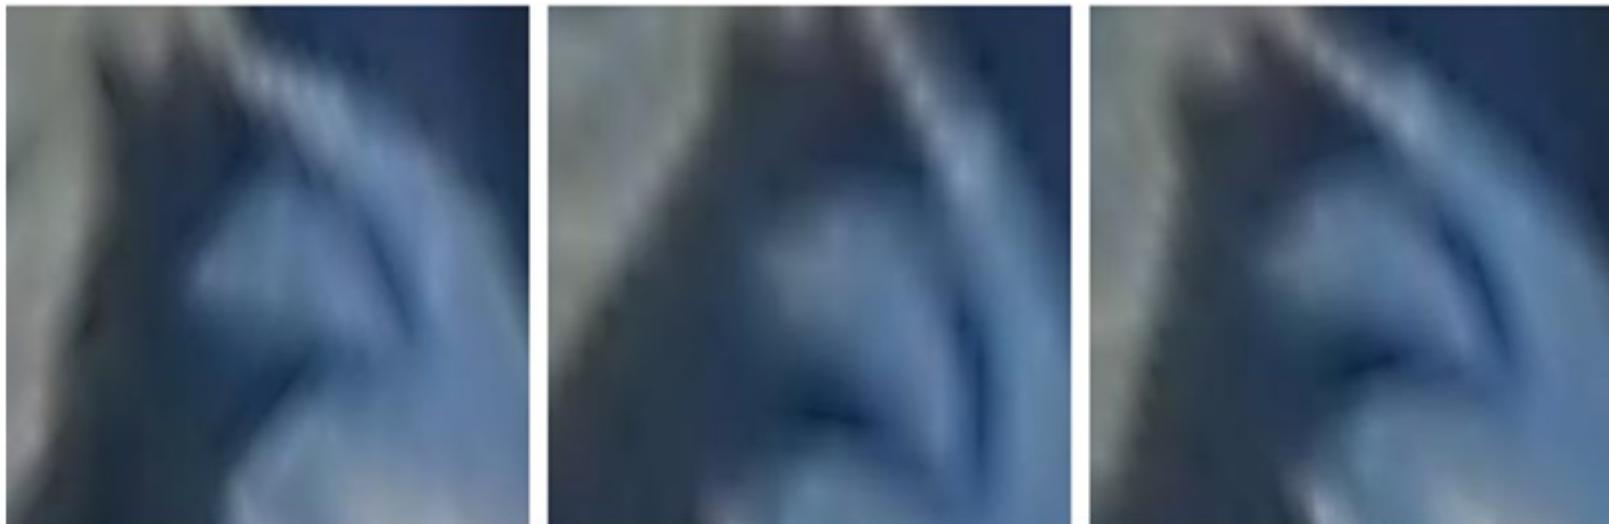
$$\rho(\mathbf{u}, \mathbf{v}) = \frac{(\mathbf{u} - \bar{\mathbf{u}}) \cdot (\mathbf{v} - \bar{\mathbf{v}})}{\|\mathbf{u} - \bar{\mathbf{u}}\| \|\mathbf{v} - \bar{\mathbf{v}}\|} = \frac{\sum_i (u_i - \bar{u})(v_i - \bar{v})}{\sqrt{\left( \sum_j (u_j - \bar{u})^2 \right) \left( \sum_j (v_j - \bar{v})^2 \right)}}$$

- Invariant to affine intensity change

# Disadvantage of intensity vectors as descriptors

---

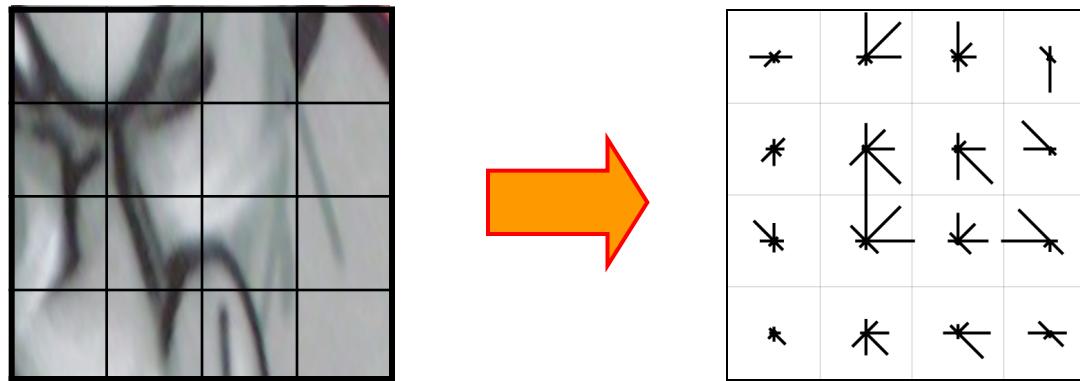
- Small deformations can affect the matching score a lot



# Feature descriptors: SIFT

---

- Descriptor computation:
  - Divide patch into  $4 \times 4$  sub-patches
  - Compute histogram of gradient orientations (8 reference angles) inside each sub-patch
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions



David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) IJCV 60 (2), pp. 91-110, 2004.

# Feature descriptors: SIFT

---

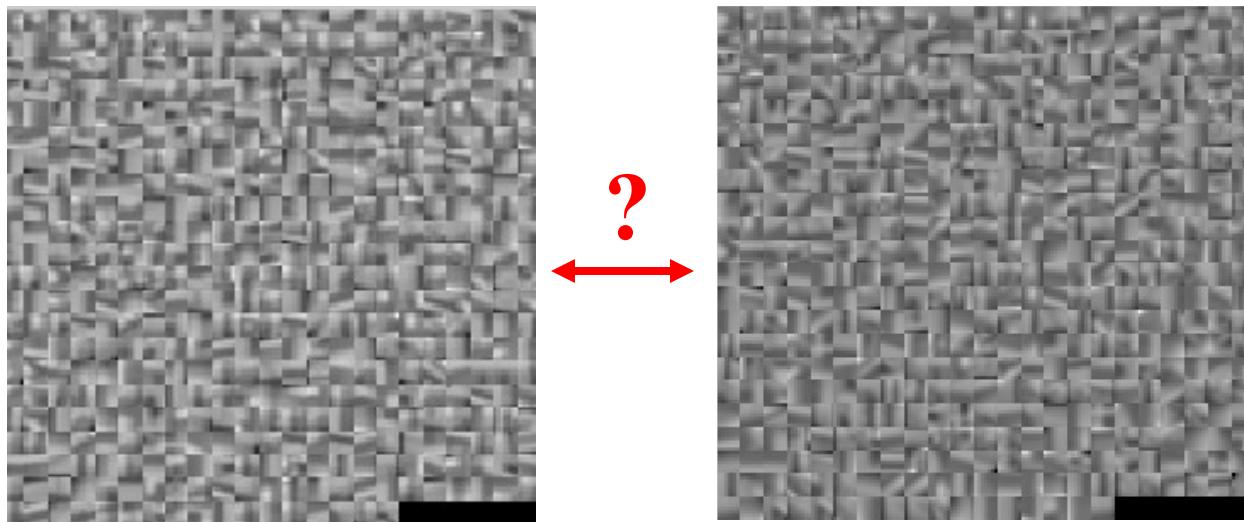
- Descriptor computation:
  - Divide patch into  $4 \times 4$  sub-patches
  - Compute histogram of gradient orientations (8 reference angles) inside each sub-patch
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions
- Advantage over raw vectors of pixel values
  - Gradients less sensitive to illumination change
  - Pooling of gradients over the sub-patches achieves robustness to small shifts, but still preserves some spatial information

David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" IJCV 60 (2), pp. 91-110, 2004.

# Feature matching

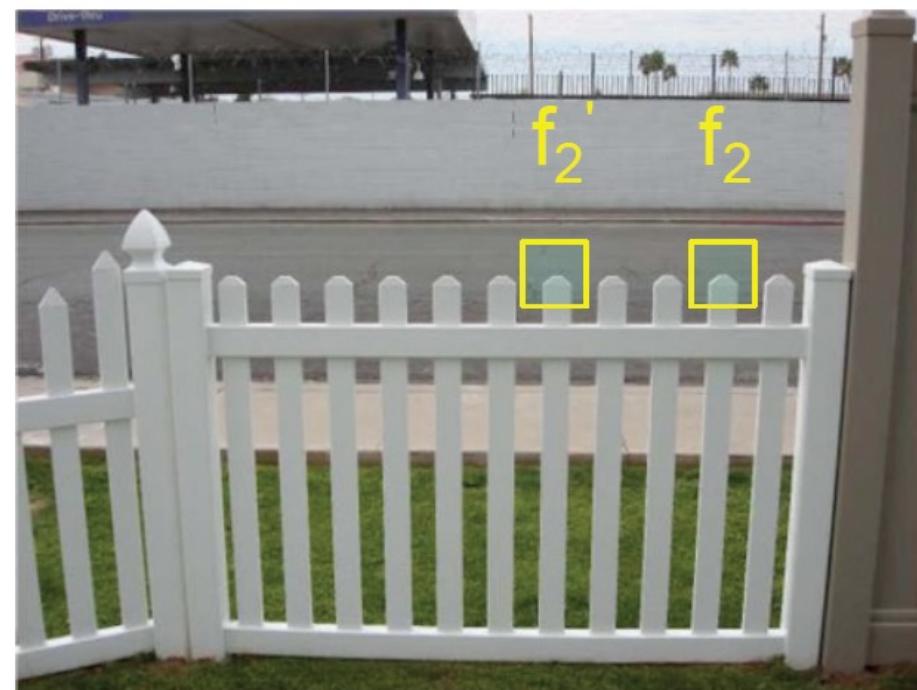
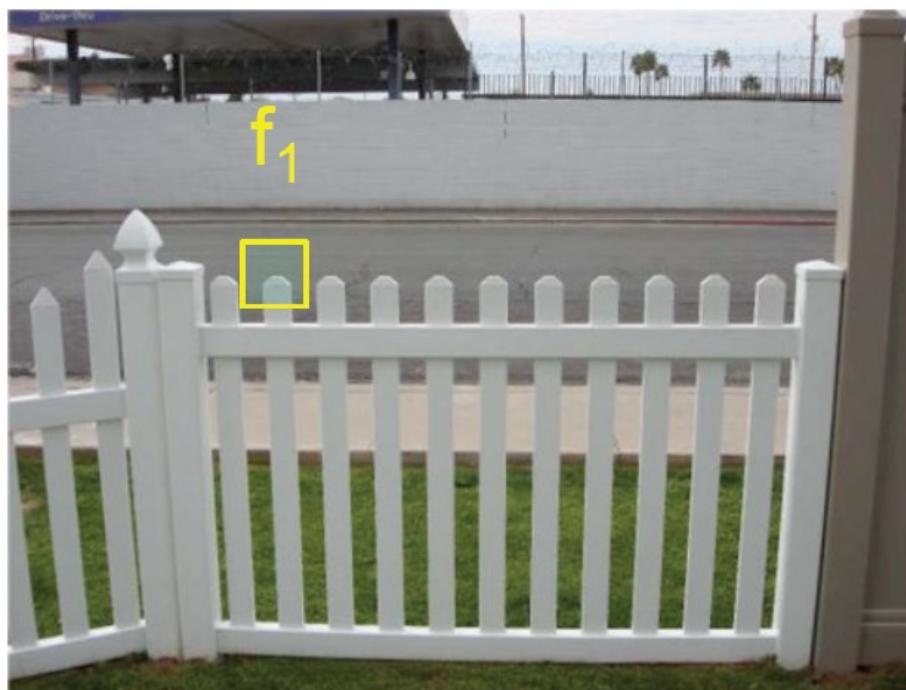
---

- Generating *putative matches*: for each patch in one image, find a short list of patches in the other image that could match it based solely on appearance



# Problem: Ambiguous putative matches

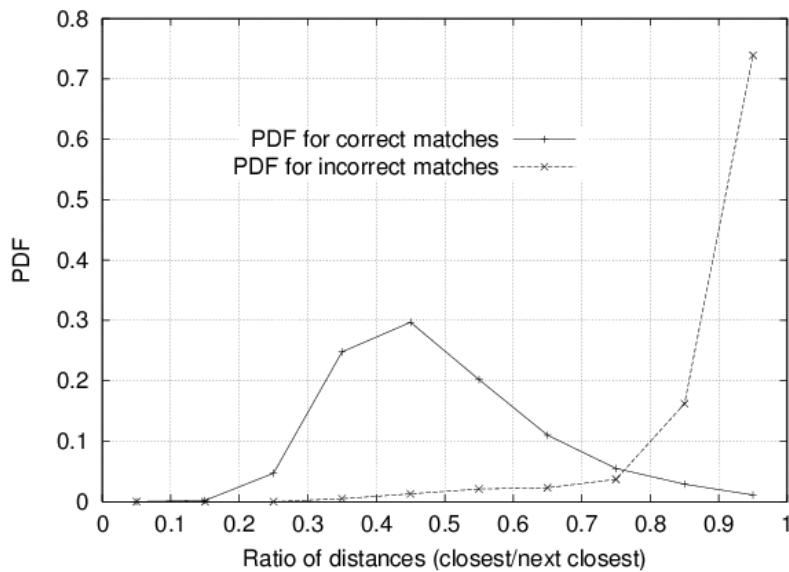
---



# Rejection of unreliable matches

---

- How can we tell which putative matches are more reliable?
- Heuristic: compare distance of **nearest** neighbor to that of **second** nearest neighbor
  - Ratio of closest distance to second-closest distance will be *high* for features that are *not* distinctive



**Threshold of 0.8 provides good separation**

# RANSAC

---

- The set of putative matches contains a very high percentage of outliers

## RANSAC loop:

1. Randomly select a *seed group* of matches
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers

Keep the transformation with the largest number of inliers

# Robust feature-based alignment

---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )

# Common transformations

---



original

Transformed



translation



rotation



aspect



affine



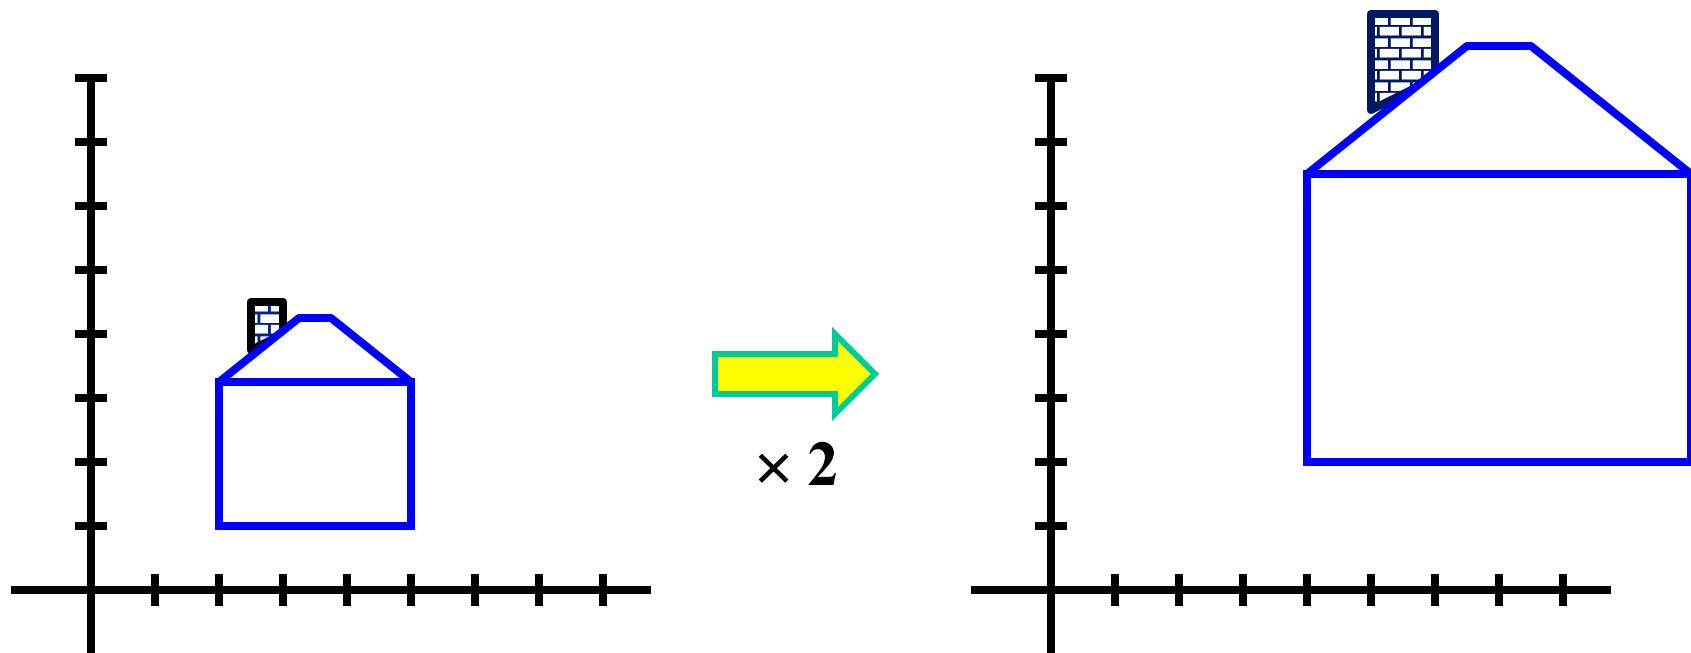
perspective

# Scaling

---

*Scaling* a coordinate means multiplying each of its components by a scalar

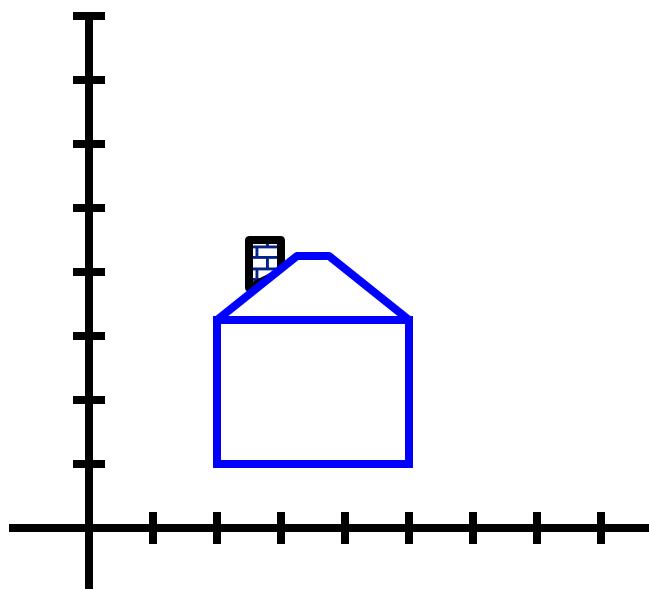
*Uniform scaling* means this scalar is the same for all components:



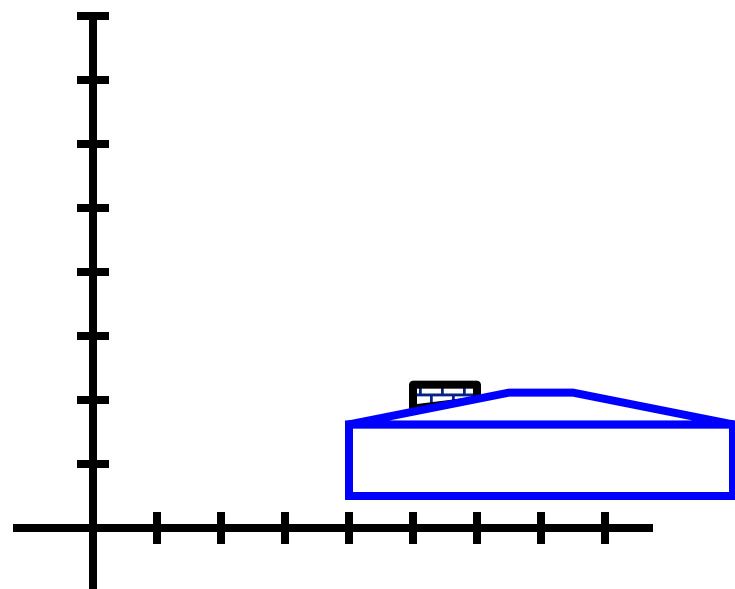
# Scaling

---

*Non-uniform scaling*: different scalars per component:



$\text{X} \times 2,$   
 $\text{Y} \times 0.5$



# Scaling

---

Scaling operation:

$$x' = ax$$

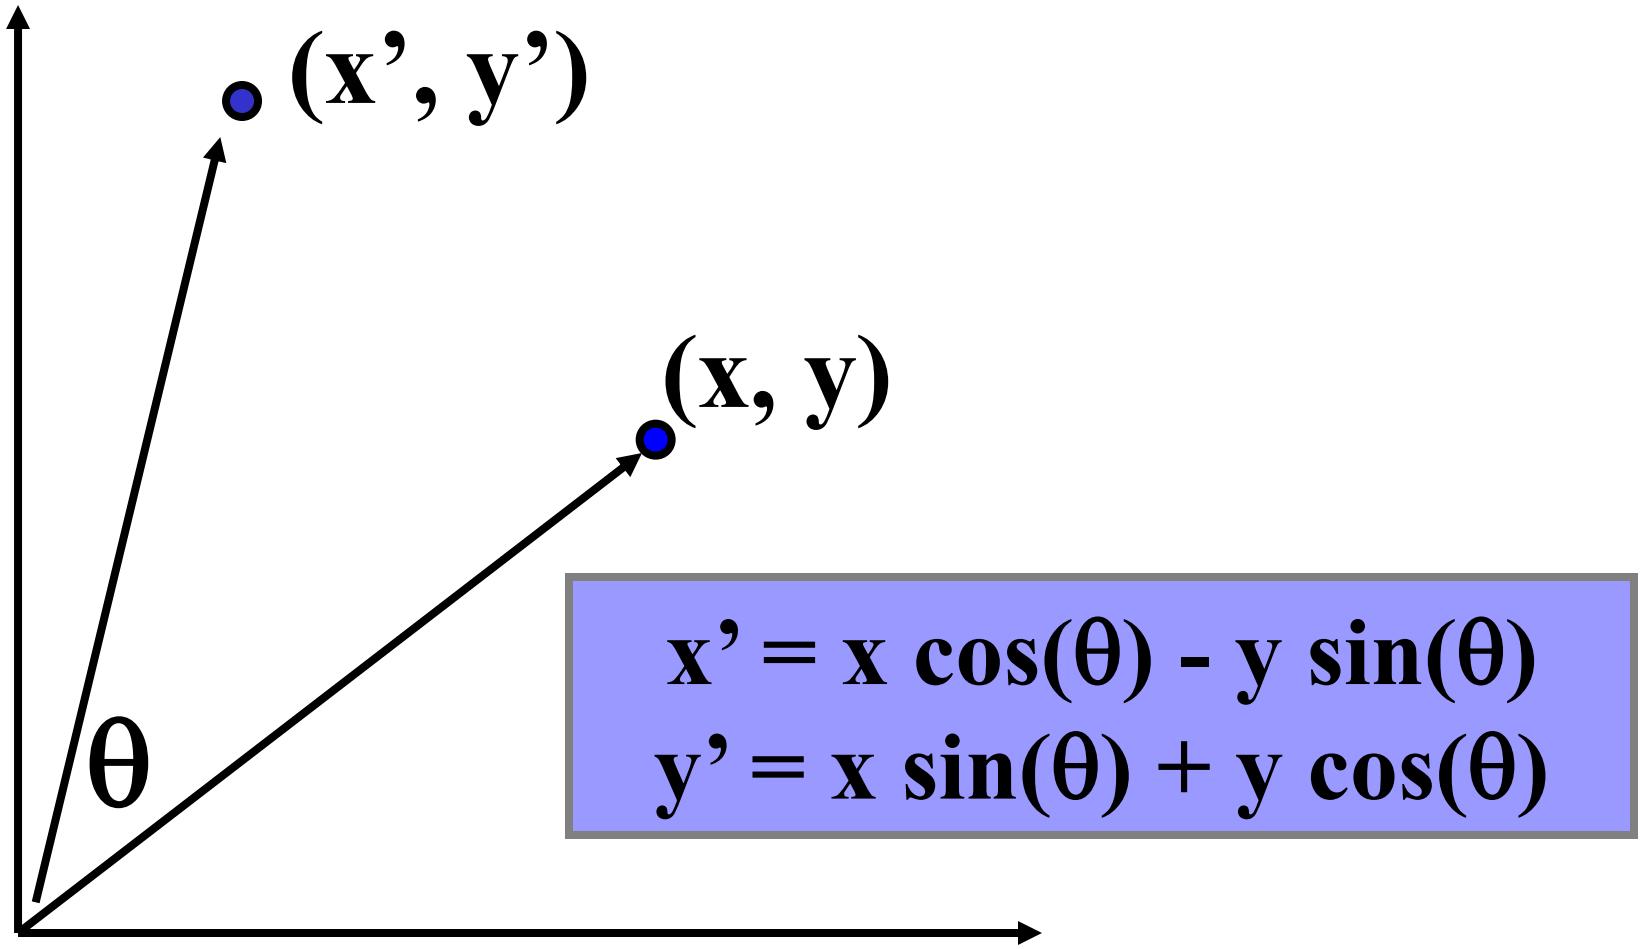
$$y' = by$$

Or, in matrix form:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \underbrace{\begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}}_{\text{scaling matrix } S} \begin{bmatrix} x \\ y \end{bmatrix}$$

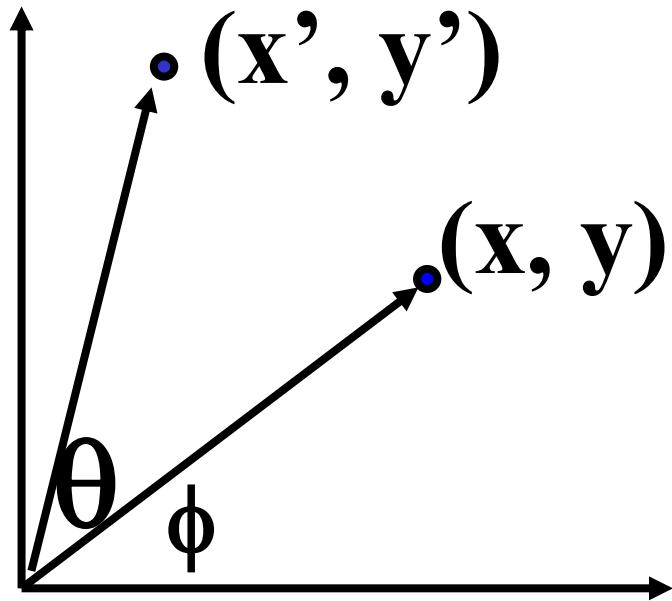
# 2-D Rotation

---



# 2-D Rotation

---



Polar coordinates...

$$x = r \cos(\phi)$$

$$y = r \sin(\phi)$$

$$x' = r \cos(\phi + \theta)$$

$$y' = r \sin(\phi + \theta)$$

Trig Identity...

$$x' = r \cos(\phi) \cos(\theta) - r \sin(\phi) \sin(\theta)$$

$$y' = r \sin(\phi) \cos(\theta) + r \cos(\phi) \sin(\theta)$$

Substitute...

$$x' = x \cos(\theta) - y \sin(\theta)$$

$$y' = x \sin(\theta) + y \cos(\theta)$$

# 2-D Rotation

---

This is easy to capture in matrix form:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \underbrace{\begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix}}_{\mathbf{R}} \begin{bmatrix} x \\ y \end{bmatrix}$$

Even though  $\sin(\theta)$  and  $\cos(\theta)$  are nonlinear functions of  $\theta$ ,

- $x'$  is a linear combination of  $x$  and  $y$
- $y'$  is a linear combination of  $x$  and  $y$

What is the inverse transformation?

- Rotation by  $-\theta$
- For rotation matrices  $\mathbf{R}^{-1} = \mathbf{R}^T$

# Basic 2D transformations

---

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Scale**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & \alpha_x \\ \alpha_y & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Shear**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos\Theta & -\sin\Theta \\ \sin\Theta & \cos\Theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Rotate**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

**Translate**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

**Affine**

**Affine is any combination of translation, scale, rotation, shear**

# Affine Transformations

---

Affine transformations are combinations of

- Linear transformations, and
- Translations

Properties of affine transformations:

- Lines map to lines
- Parallel lines remain parallel
- Ratios are preserved
- Closed under composition

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

or

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Projective Transformations

---

Projective transformations are combos of

- Affine transformations, and
- Projective warps

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

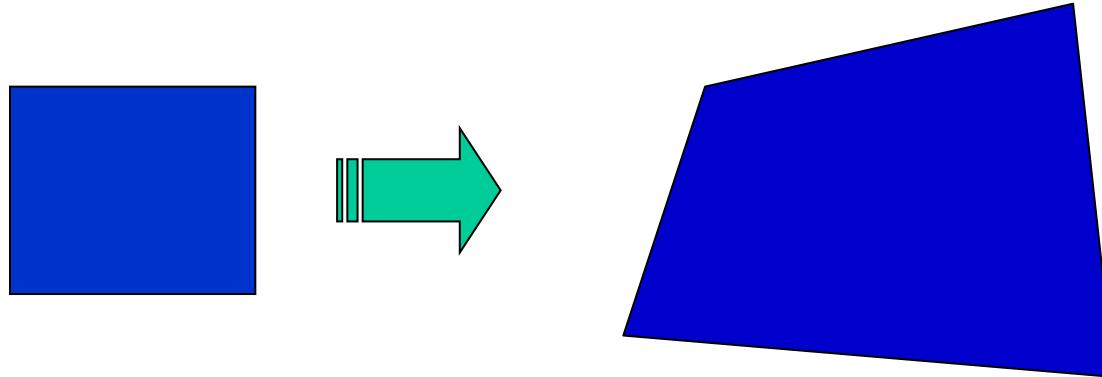
Properties of projective transformations:

- Lines map to lines
- Parallel lines do not necessarily remain parallel
- Ratios are not preserved
- Closed under composition
- **Projective matrix is defined up to a scale (8 DOF)**

# Fitting a plane projective transformation

---

- **Homography:** plane projective transformation  
(transformation taking a quad to another arbitrary quad)



# Homography

---

- The transformation between two views of a planar surface

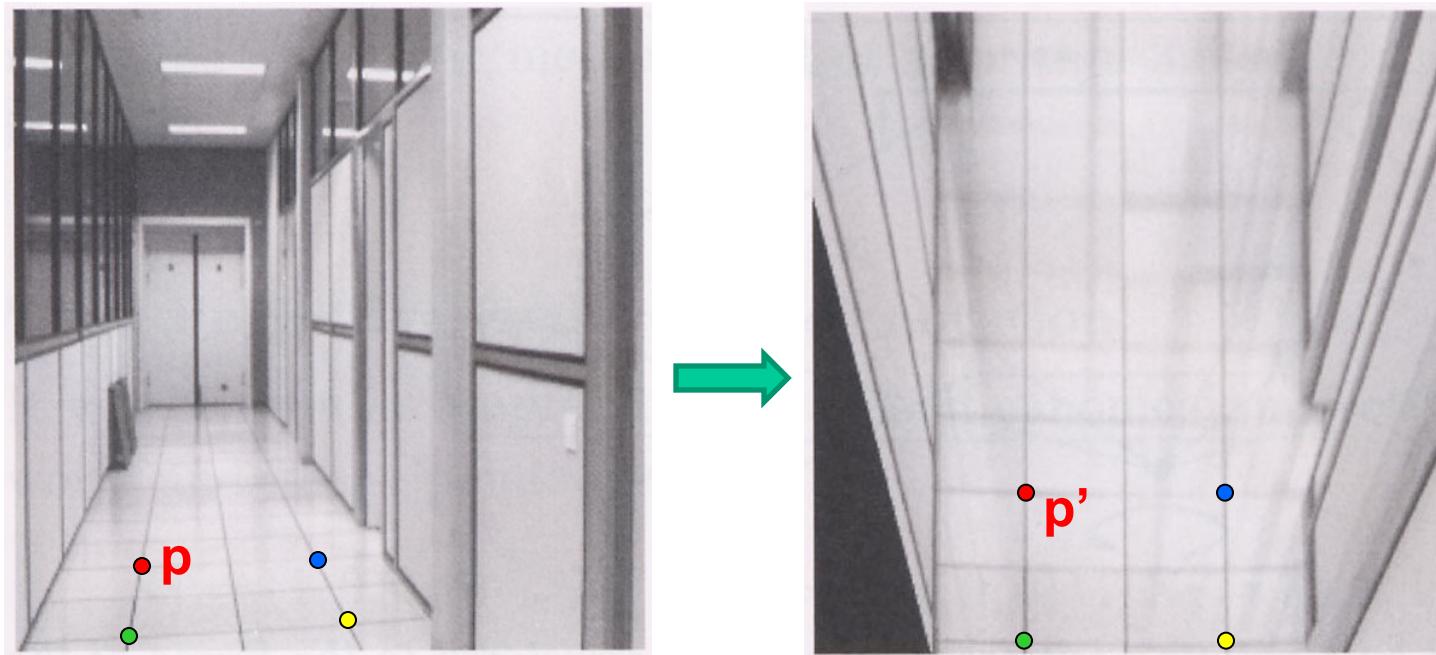


- The transformation between images from two cameras that share the same center



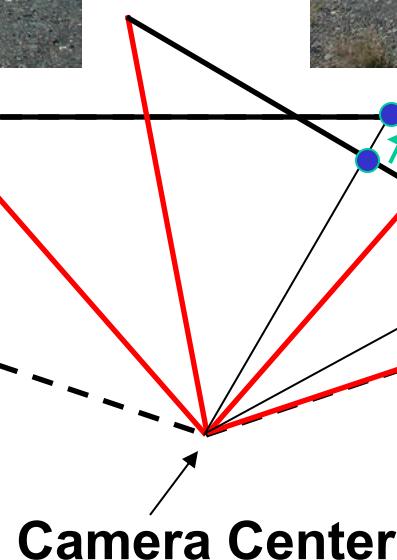
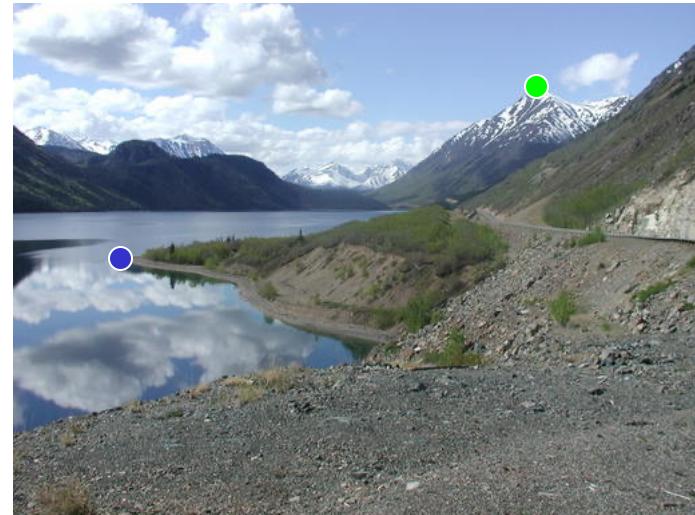
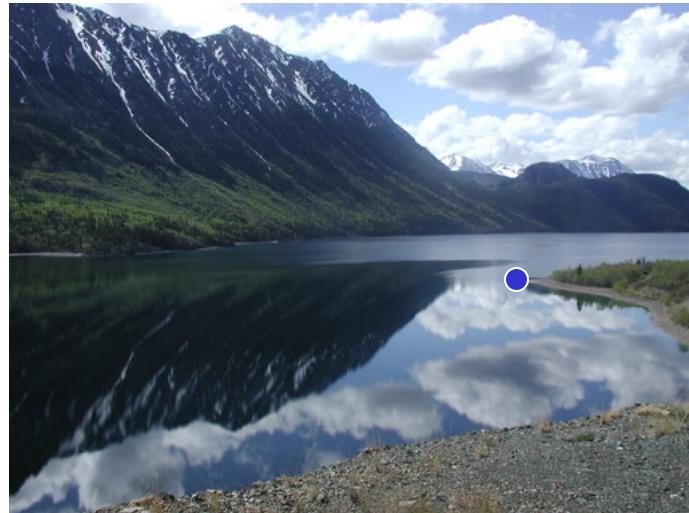
# Homography example: Image rectification

---



# Homography Example

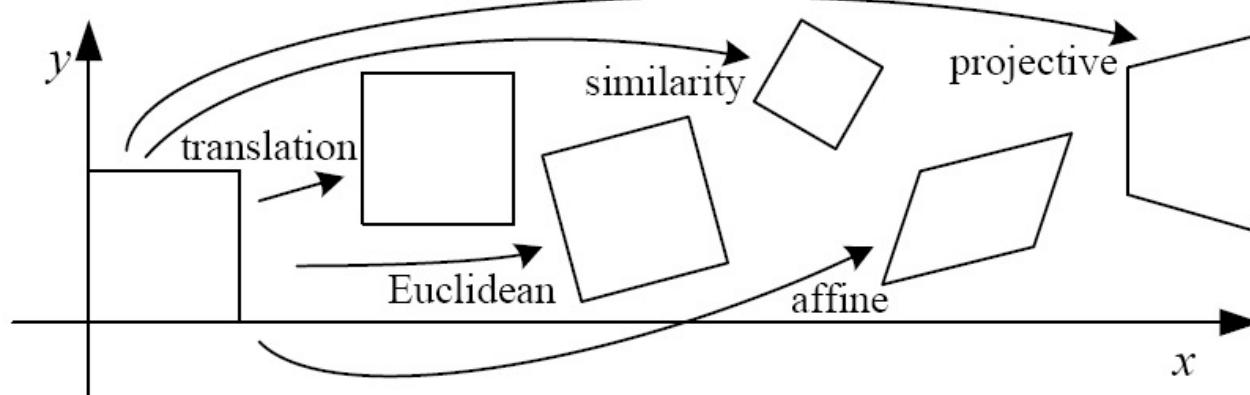
---



**Camera Center**

# 2D image transformations

---

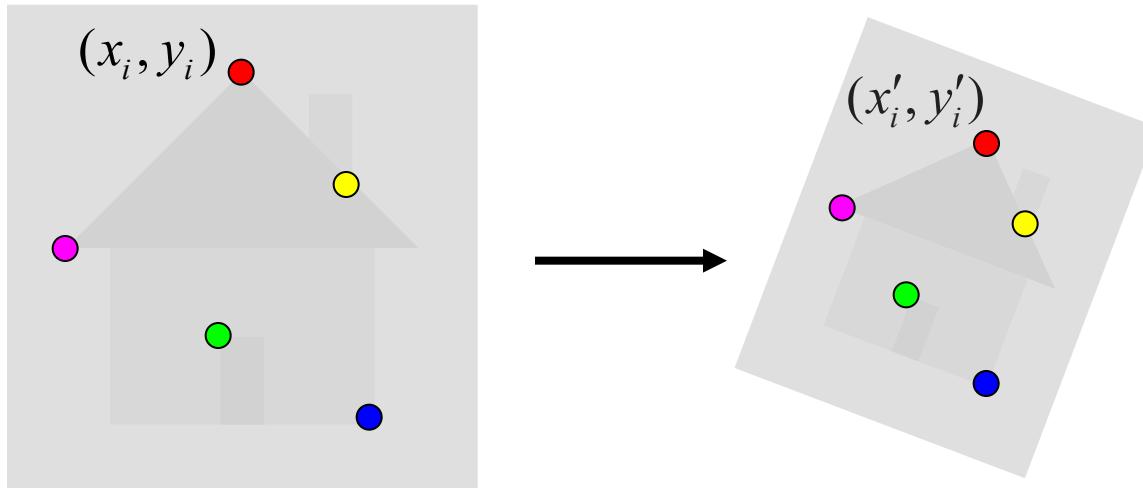


Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

# Fitting an affine transformation

---

- Assume we know the correspondences, how do we get the transformation?



$$\mathbf{x}'_i = \mathbf{M}\mathbf{x}_i + \mathbf{t}$$

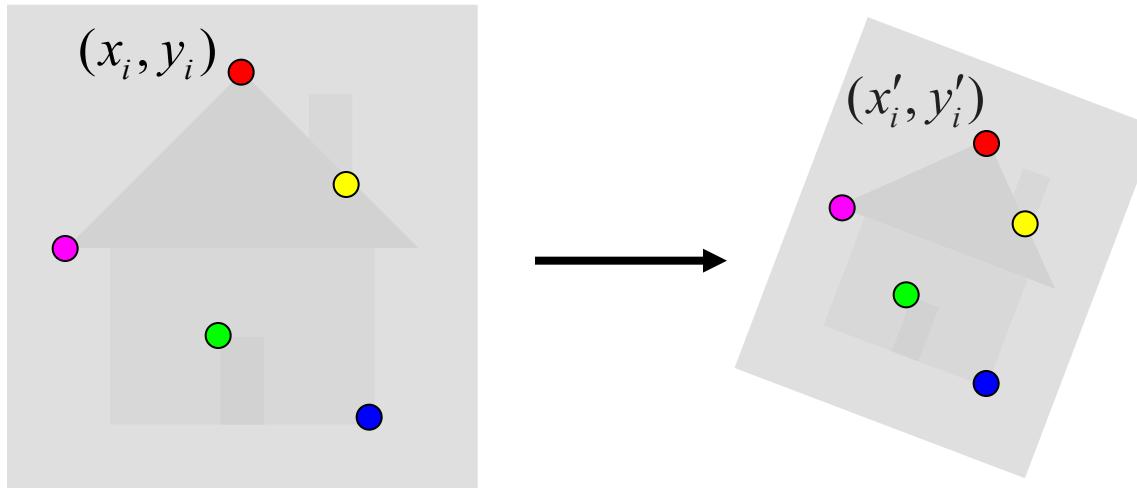
$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

Want to find  $\mathbf{M}$ ,  $\mathbf{t}$  to minimize

$$\sum_{i=1}^n \| \mathbf{x}'_i - \mathbf{M}\mathbf{x}_i - \mathbf{t} \|^2$$

# Fitting an affine transformation

- Assume we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\left[ \begin{array}{c} \vdots \\ m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{array} \right]$$

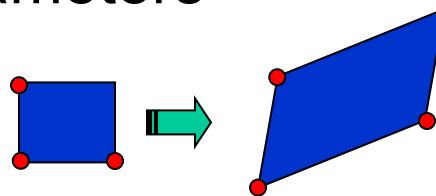
$$\left[ \begin{array}{c} \vdots \\ x'_i \\ y'_i \\ \dots \end{array} \right] = \left[ \begin{array}{c} \vdots \\ x'_i \\ y'_i \\ \dots \end{array} \right]$$

# Fitting an affine transformation

---

$$\begin{bmatrix} & & \cdots \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ & & \cdots \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \cdots \\ x'_i \\ y'_i \\ \cdots \end{bmatrix}$$

- Linear system with six unknowns
- Each match gives us two linearly independent equations: need at least three to solve for the transformation parameters



# Fitting a homography

---

- Recall: homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogeneous  
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogeneous  
image coordinates

# Fitting a homography

---

- Recall: homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogeneous  
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogeneous  
image coordinates

- Equation for homography:

$$\lambda \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Fitting a homography

---

- Equation for homography:

$$\lambda \begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \quad \lambda \mathbf{x}'_i = \mathbf{H} \mathbf{x}_i$$
$$\mathbf{x}'_i \times \mathbf{H} \mathbf{x}_i = 0$$

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} \times \begin{bmatrix} \mathbf{h}_1^T \mathbf{x}_i \\ \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_3^T \mathbf{x}_i \end{bmatrix} = \begin{bmatrix} y'_i \mathbf{h}_3^T \mathbf{x}_i - \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_1^T \mathbf{x}_i - x'_i \mathbf{h}_3^T \mathbf{x}_i \\ x'_i \mathbf{h}_2^T \mathbf{x}_i - y'_i \mathbf{h}_1^T \mathbf{x}_i \end{bmatrix}$$

$$\begin{bmatrix} 0^T & -\mathbf{x}_i^T & y'_i \mathbf{x}_i^T \\ \mathbf{x}_i^T & 0^T & -x'_i \mathbf{x}_i^T \\ -y'_i \mathbf{x}_i^T & x'_i \mathbf{x}_i^T & 0^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0$$

3 equations,  
only 2 linearly  
independent

# Fitting a homography

---

$$\begin{bmatrix} 0^T & \mathbf{x}_1^T & -y'_1 \mathbf{x}_1^T \\ \mathbf{x}_1^T & 0^T & -x'_1 \mathbf{x}_1^T \\ \dots & \dots & \dots \\ 0^T & \mathbf{x}_n^T & -y'_n \mathbf{x}_n^T \\ \mathbf{x}_n^T & 0^T & -x'_n \mathbf{x}_n^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0 \quad \mathbf{A h} = 0$$

- $\mathbf{H}$  has 8 degrees of freedom (9 parameters, but scale is arbitrary)
- One match gives us two linearly independent equations
- Homogeneous least squares: find  $\mathbf{h}$  minimizing  $\|\mathbf{Ah}\|^2$ 
  - Eigenvector of  $\mathbf{A}^T \mathbf{A}$  corresponding to smallest eigenvalue
  - Four matches needed for a minimal solution

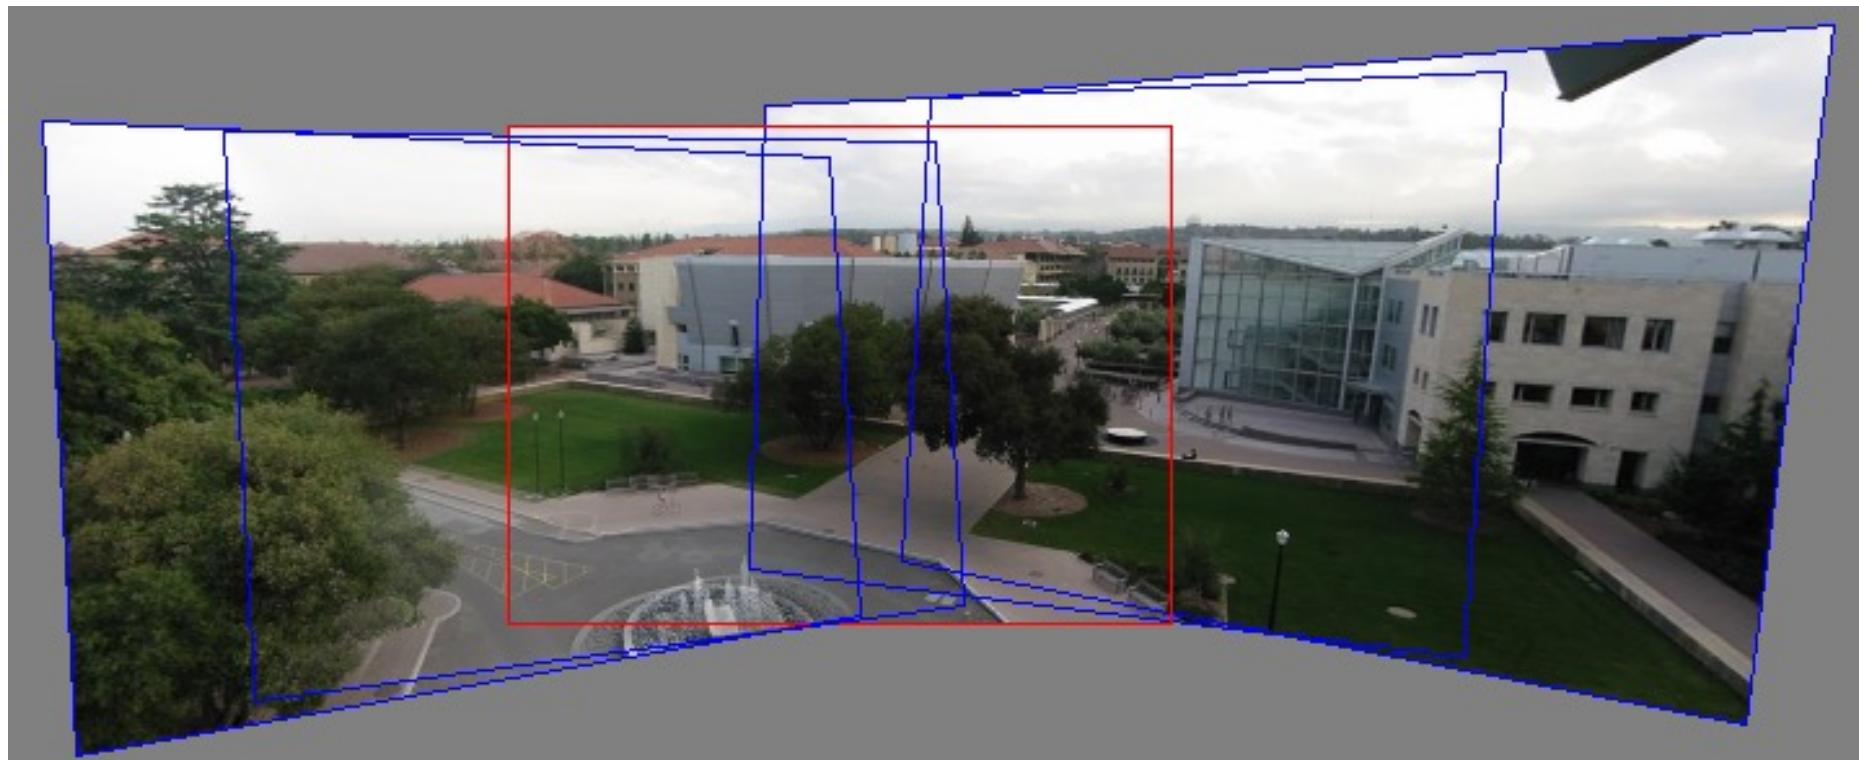
# Feature-based alignment: Overview

---

- Alignment as fitting
  - Affine transformations
  - Homographies
- Robust alignment
  - Descriptor-based feature matching
  - RANSAC
- Applications

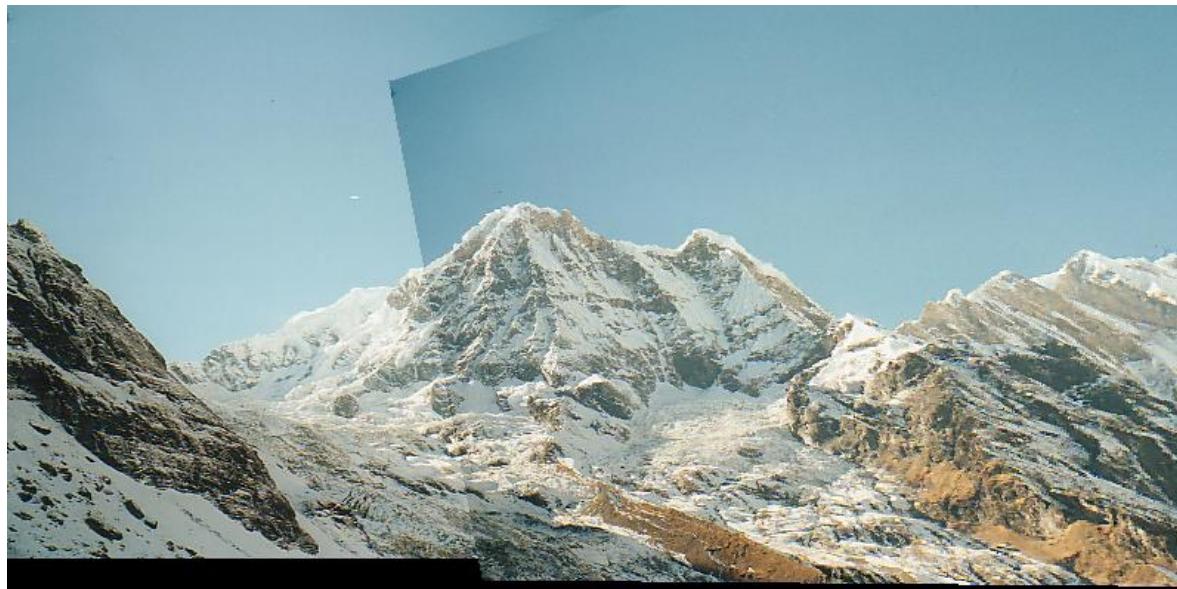
# Image alignment

---



# Image Alignment

---



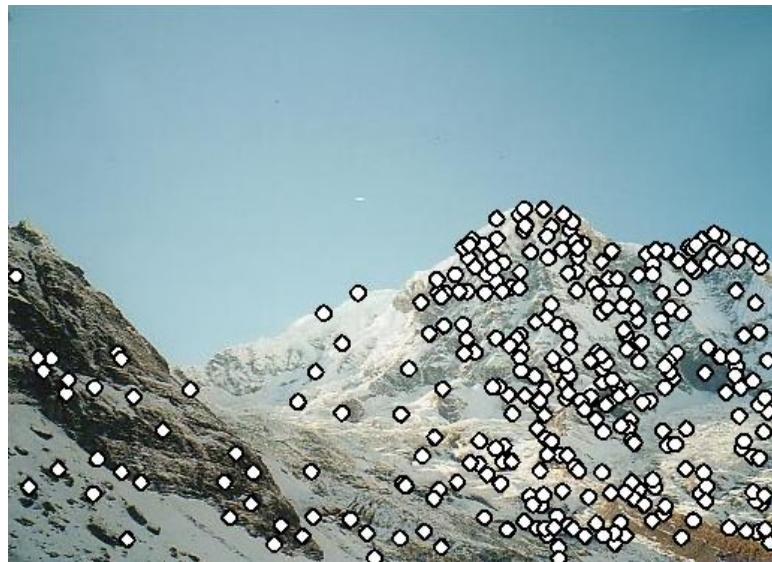
# Robust feature-based alignment

---



# Robust feature-based alignment

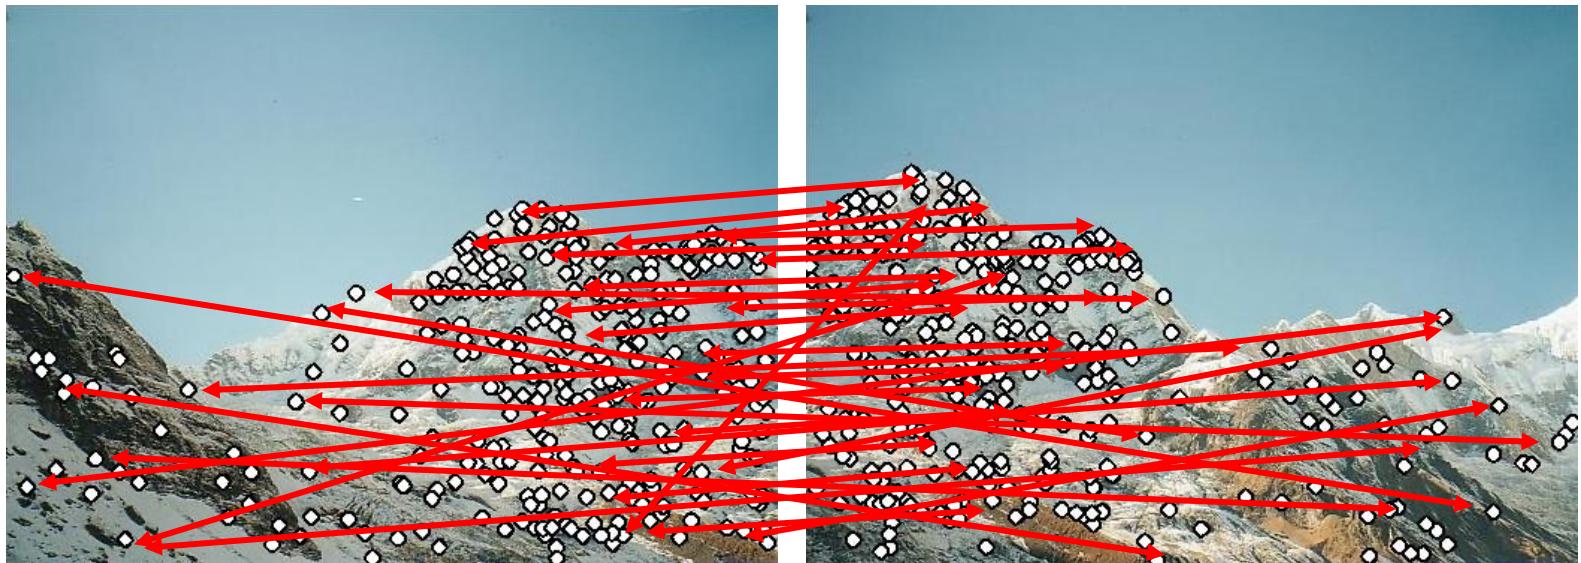
---



- Extract features

# Robust feature-based alignment

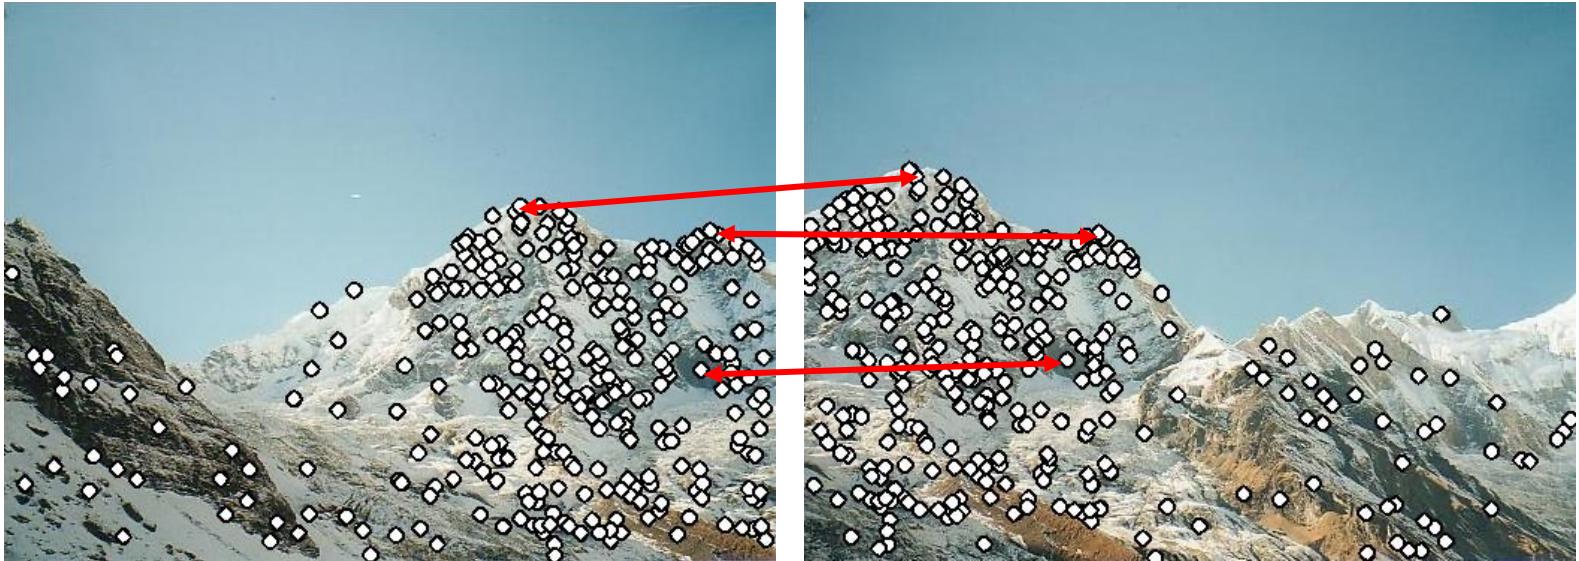
---



- Extract features
- Compute *putative matches*

# Robust feature-based alignment

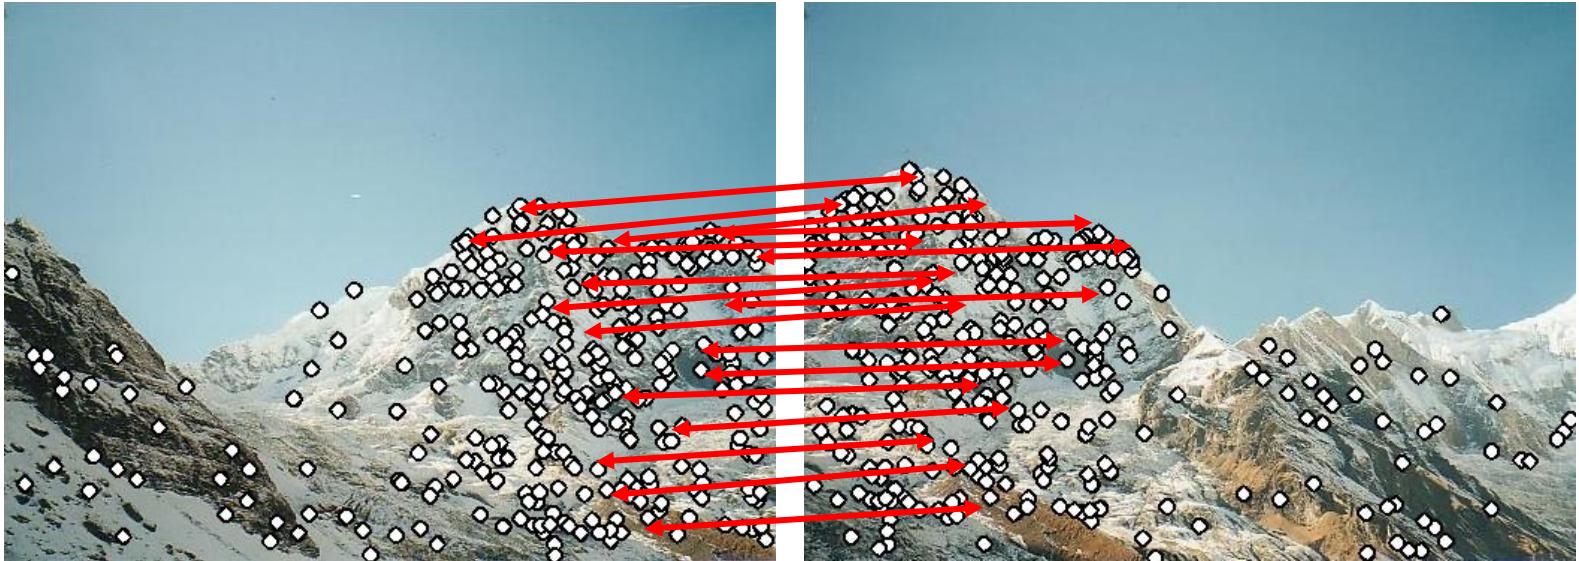
---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$

# Robust feature-based alignment

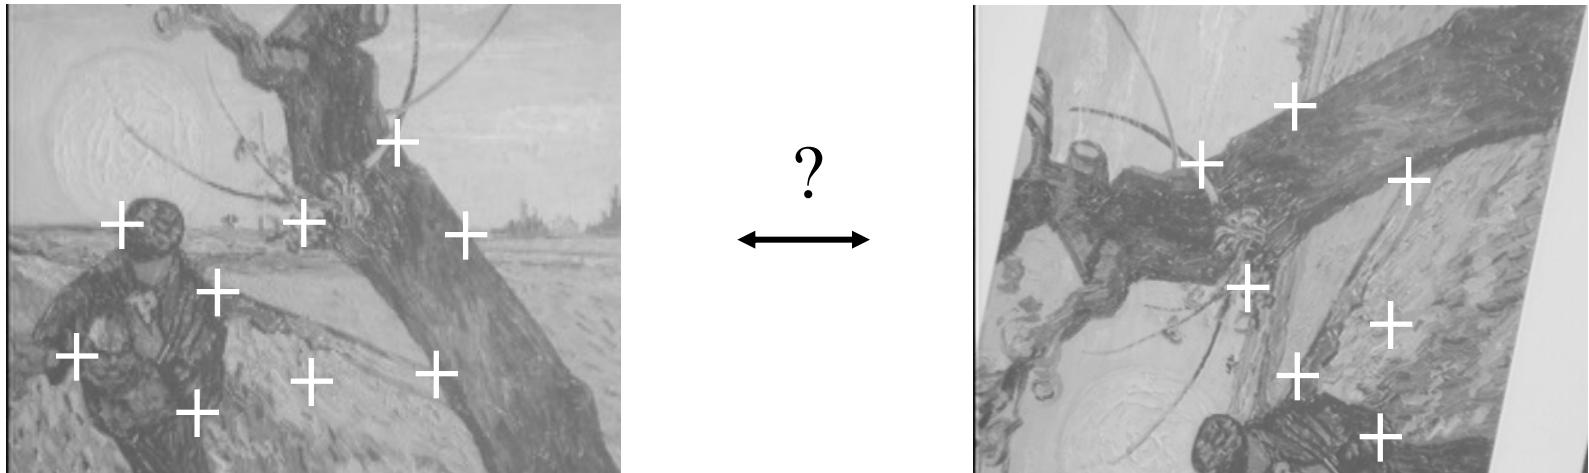
---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )

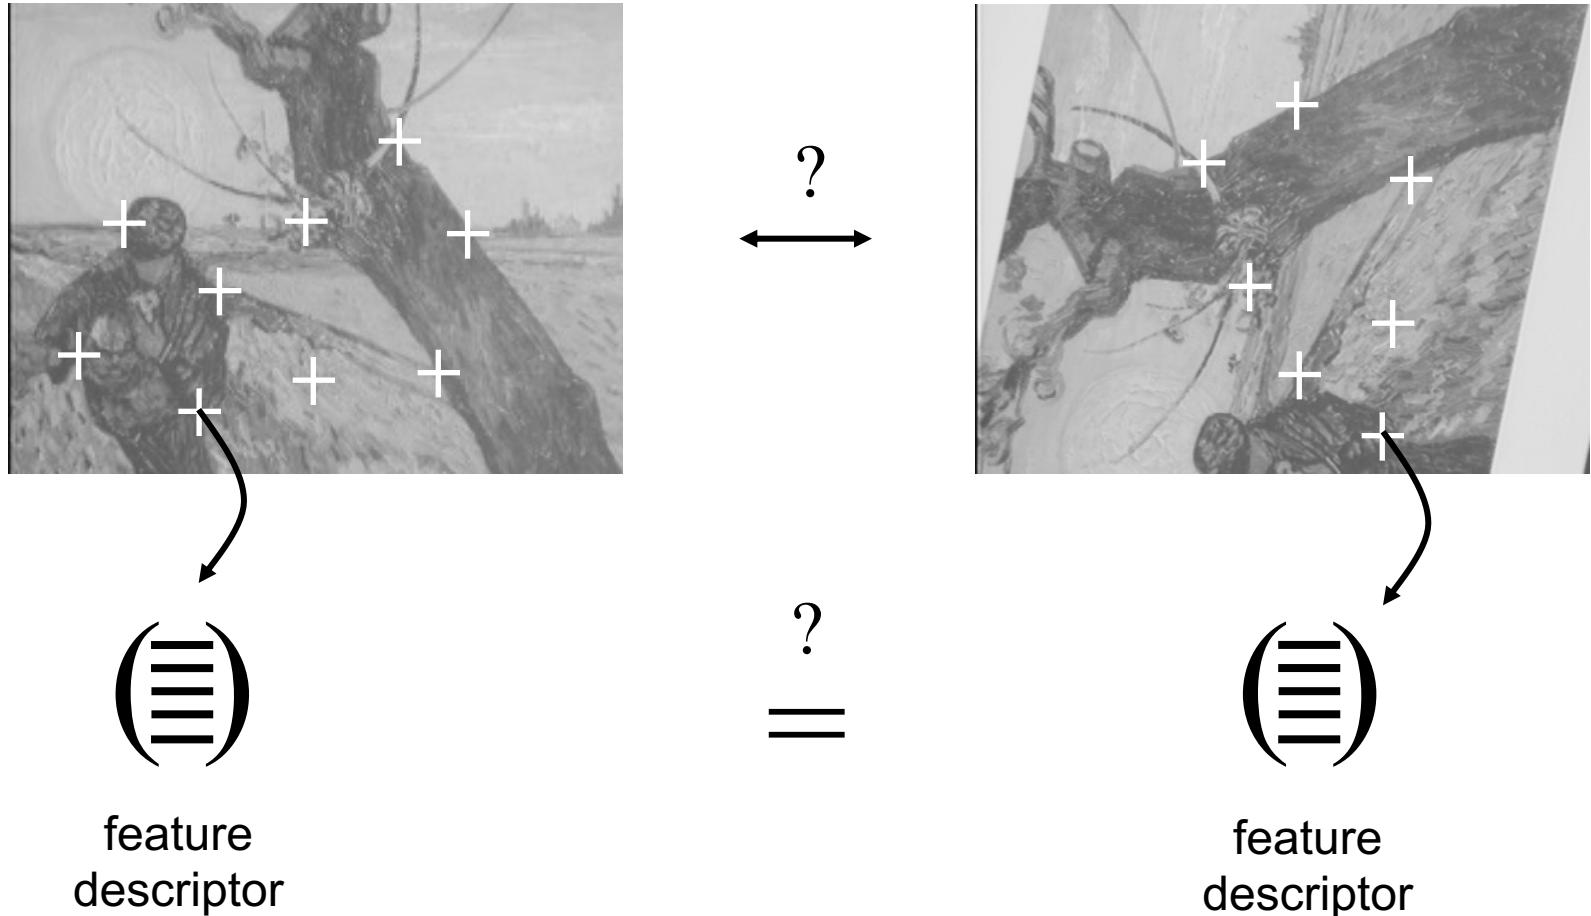
# Generating putative correspondences

---



# Generating putative correspondences

---

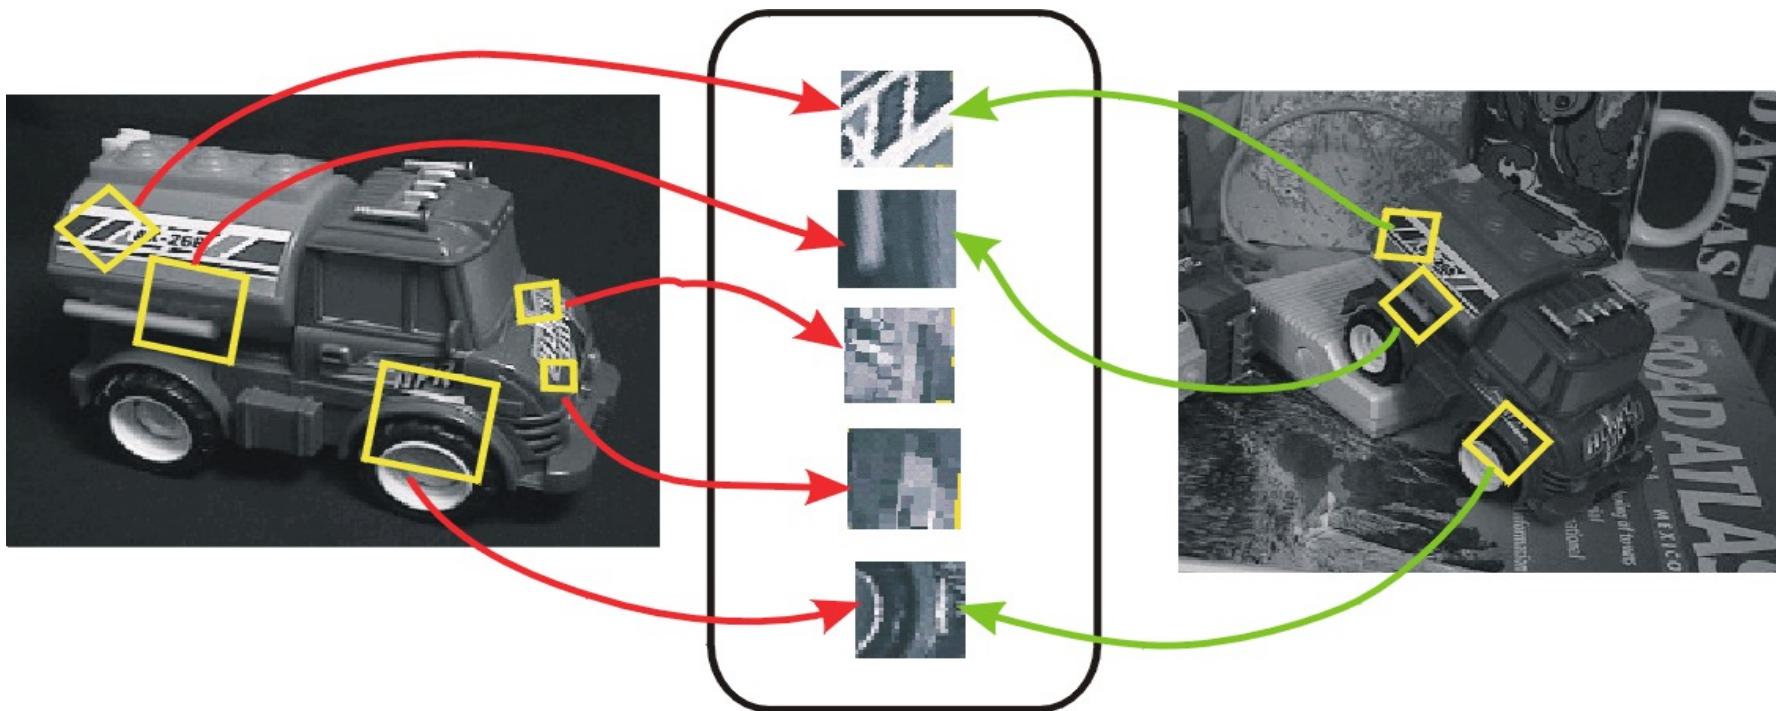


- Need to compare *feature descriptors* of local patches surrounding interest points

# Feature descriptors

---

- Recall: feature detection and description



# Feature descriptors

---

- Simplest descriptor: vector of raw intensity values
- How to compare two such vectors?
  - Sum of squared differences (SSD)

$$\text{SSD}(\mathbf{u}, \mathbf{v}) = \sum_i (u_i - v_i)^2$$

- Not invariant to intensity change
- Normalized correlation

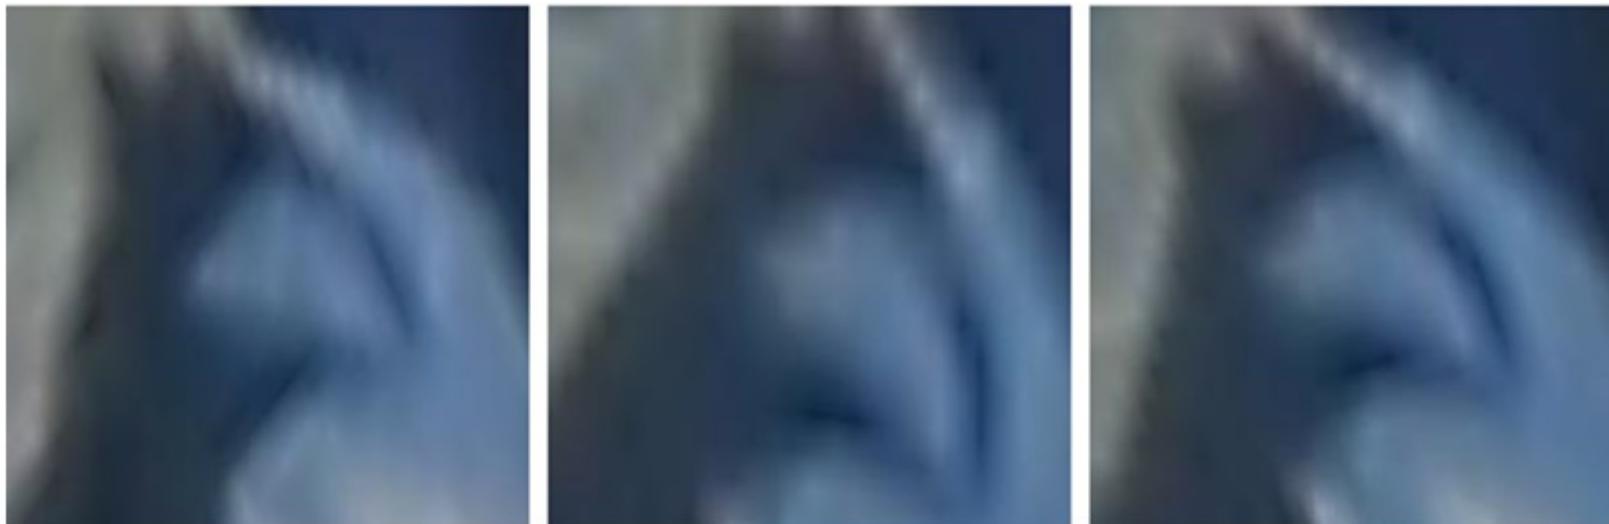
$$\rho(\mathbf{u}, \mathbf{v}) = \frac{(\mathbf{u} - \bar{\mathbf{u}}) \cdot (\mathbf{v} - \bar{\mathbf{v}})}{\|\mathbf{u} - \bar{\mathbf{u}}\| \|\mathbf{v} - \bar{\mathbf{v}}\|} = \frac{\sum_i (u_i - \bar{u})(v_i - \bar{v})}{\sqrt{\left( \sum_j (u_j - \bar{u})^2 \right) \left( \sum_j (v_j - \bar{v})^2 \right)}}$$

- Invariant to affine intensity change

# Disadvantage of intensity vectors as descriptors

---

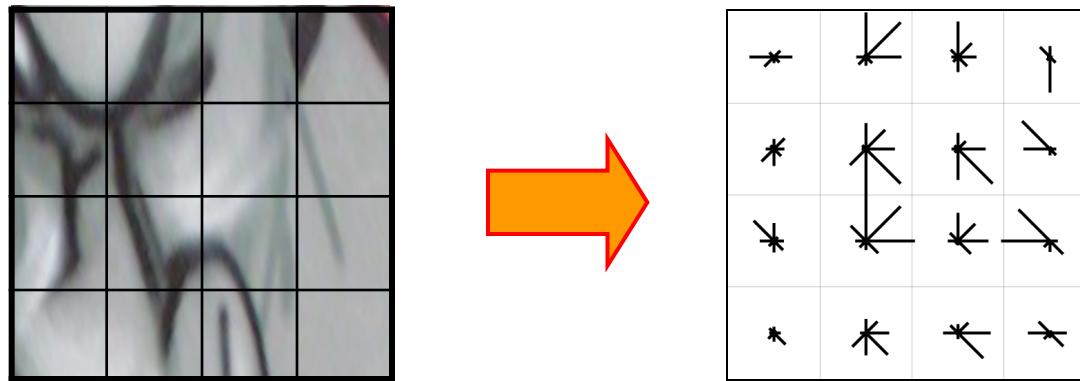
- Small deformations can affect the matching score a lot



# Feature descriptors: SIFT

---

- Descriptor computation:
  - Divide patch into  $4 \times 4$  sub-patches
  - Compute histogram of gradient orientations (8 reference angles) inside each sub-patch
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions



David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) IJCV 60 (2), pp. 91-110, 2004.

# Feature descriptors: SIFT

---

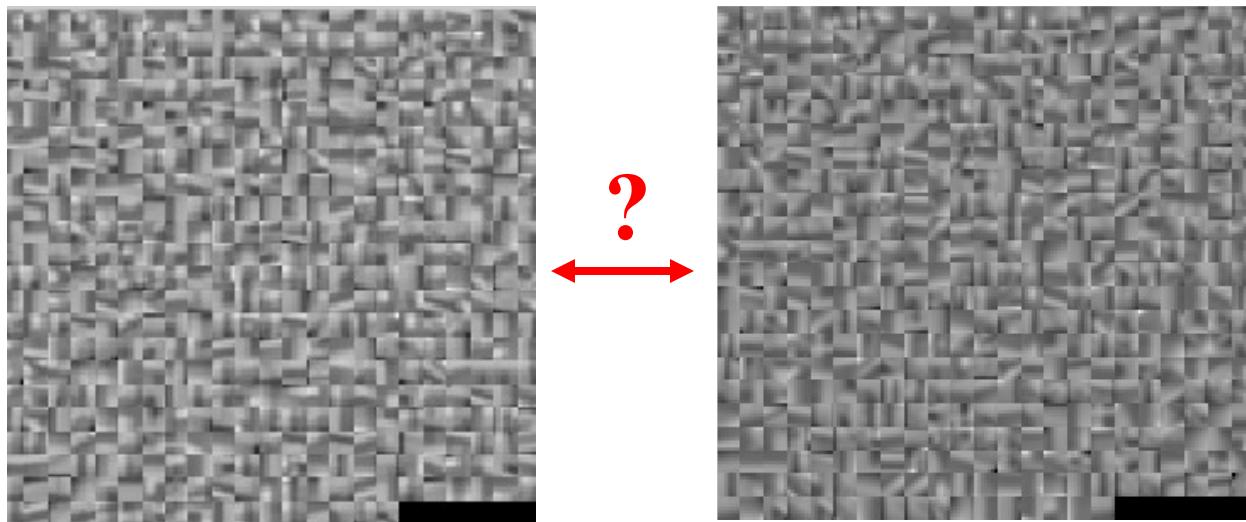
- Descriptor computation:
  - Divide patch into  $4 \times 4$  sub-patches
  - Compute histogram of gradient orientations (8 reference angles) inside each sub-patch
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions
- Advantage over raw vectors of pixel values
  - Gradients less sensitive to illumination change
  - Pooling of gradients over the sub-patches achieves robustness to small shifts, but still preserves some spatial information

David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" IJCV 60 (2), pp. 91-110, 2004.

# Feature matching

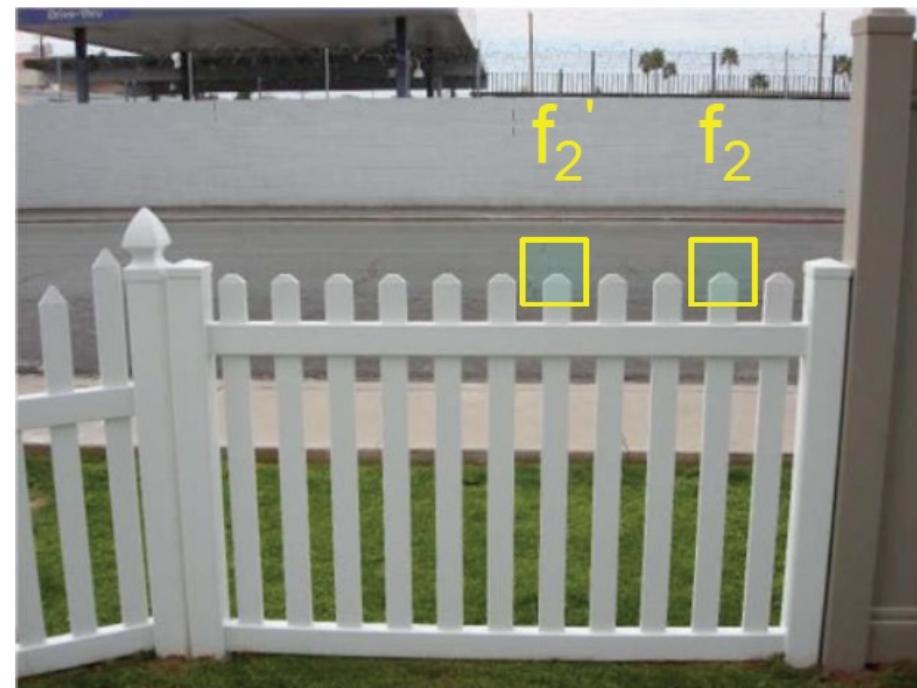
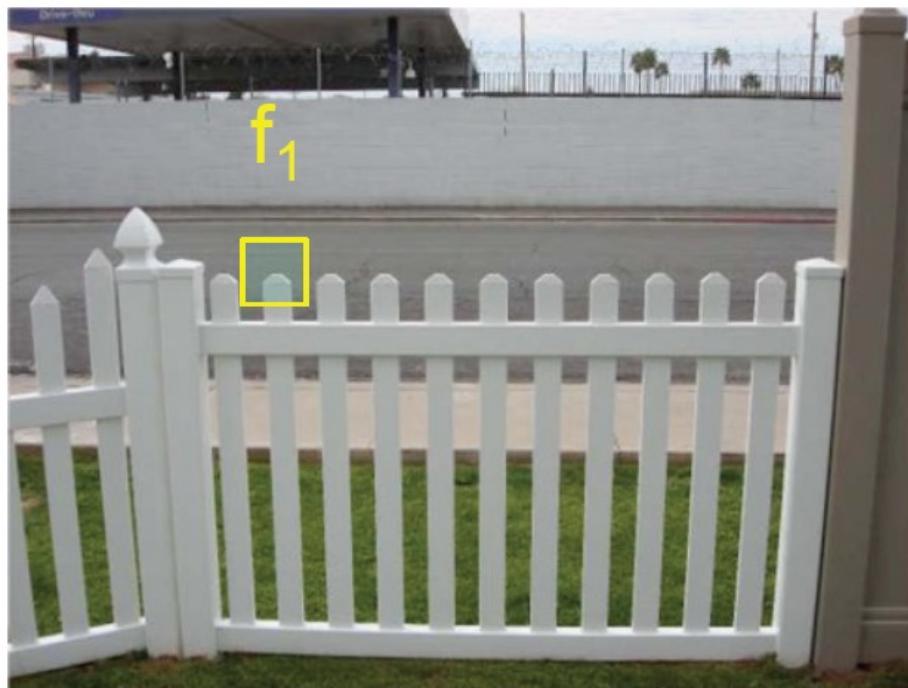
---

- Generating *putative matches*: for each patch in one image, find a short list of patches in the other image that could match it based solely on appearance



# Problem: Ambiguous putative matches

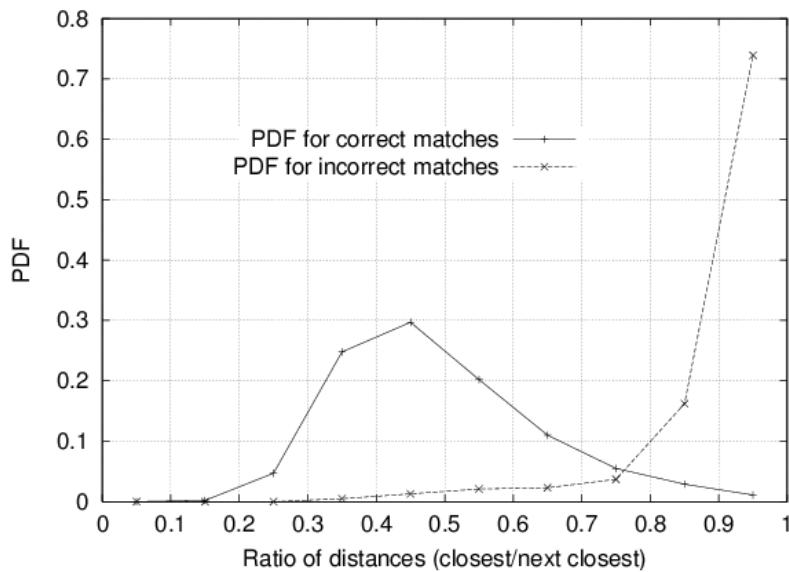
---



# Rejection of unreliable matches

---

- How can we tell which putative matches are more reliable?
- Heuristic: compare distance of **nearest** neighbor to that of **second** nearest neighbor
  - Ratio of closest distance to second-closest distance will be *high* for features that are *not* distinctive



**Threshold of 0.8 provides good separation**

# RANSAC

---

- The set of putative matches contains a very high percentage of outliers

## RANSAC loop:

1. Randomly select a *seed group* of matches
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers

Keep the transformation with the largest number of inliers

# Robust feature-based alignment

---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )

# Common transformations

---



original

Transformed



translation



rotation



aspect



affine



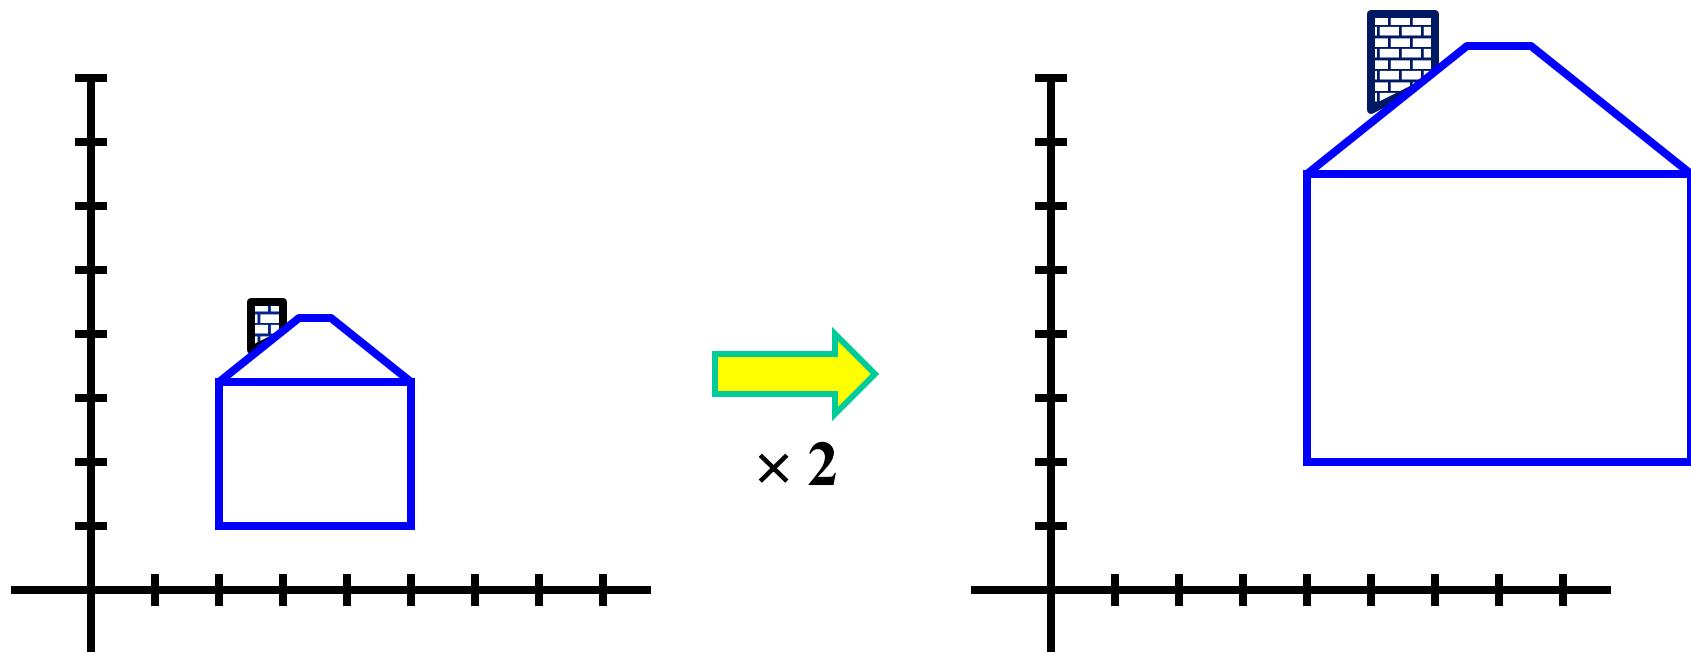
perspective

# Scaling

---

*Scaling* a coordinate means multiplying each of its components by a scalar

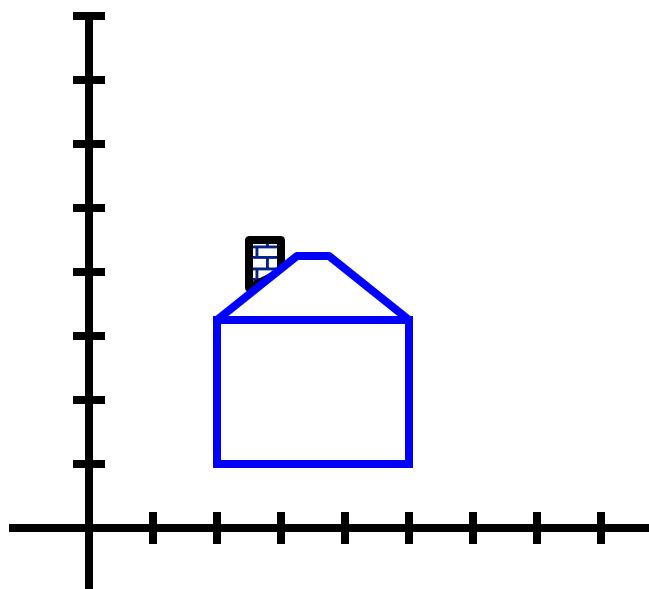
*Uniform scaling* means this scalar is the same for all components:



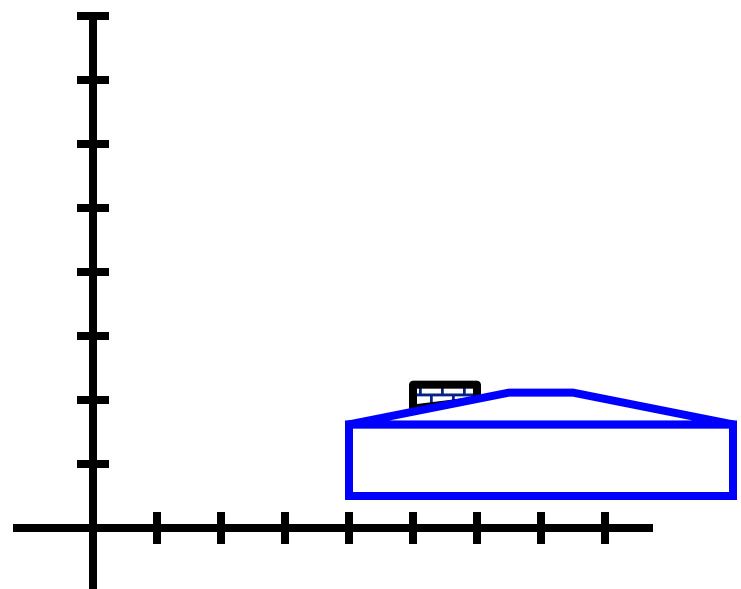
# Scaling

---

*Non-uniform scaling*: different scalars per component:



$\rightarrow$   
 $X \times 2,$   
 $Y \times 0.5$



# Scaling

---

Scaling operation:

$$x' = ax$$

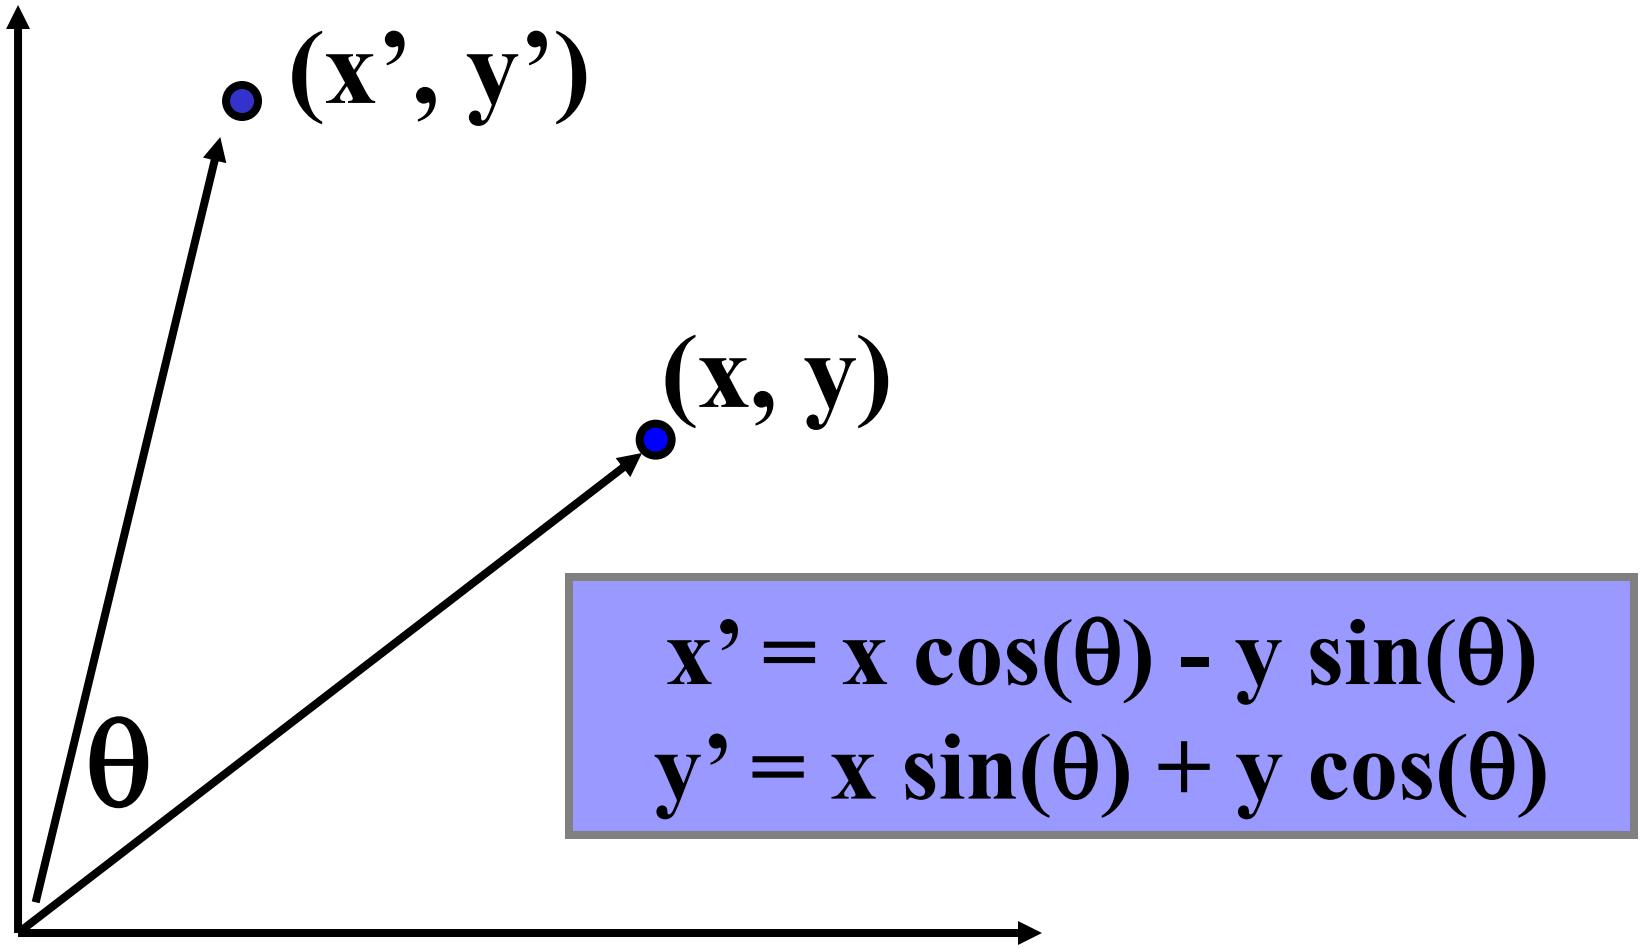
$$y' = by$$

Or, in matrix form:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \underbrace{\begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}}_{\text{scaling matrix } S} \begin{bmatrix} x \\ y \end{bmatrix}$$

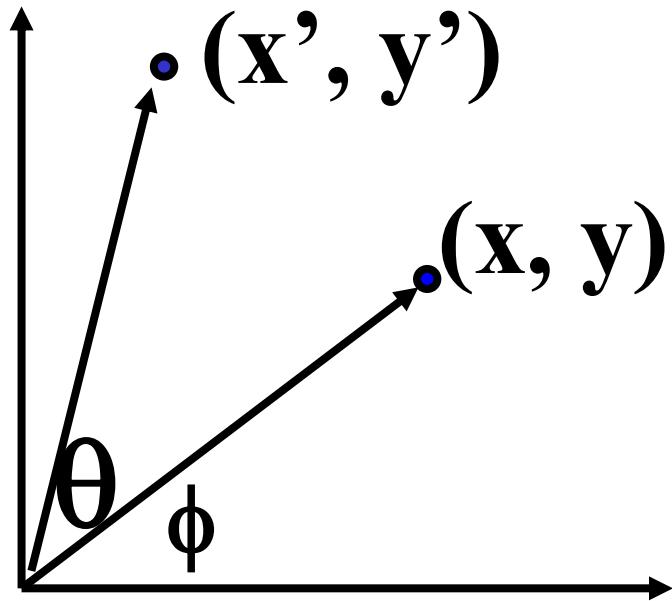
# 2-D Rotation

---



# 2-D Rotation

---



Polar coordinates...

$$x = r \cos(\phi)$$

$$y = r \sin(\phi)$$

$$x' = r \cos(\phi + \theta)$$

$$y' = r \sin(\phi + \theta)$$

Trig Identity...

$$x' = r \cos(\phi) \cos(\theta) - r \sin(\phi) \sin(\theta)$$

$$y' = r \sin(\phi) \cos(\theta) + r \cos(\phi) \sin(\theta)$$

Substitute...

$$x' = x \cos(\theta) - y \sin(\theta)$$

$$y' = x \sin(\theta) + y \cos(\theta)$$

# 2-D Rotation

---

This is easy to capture in matrix form:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \underbrace{\begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix}}_{\mathbf{R}} \begin{bmatrix} x \\ y \end{bmatrix}$$

Even though  $\sin(\theta)$  and  $\cos(\theta)$  are nonlinear functions of  $\theta$ ,

- $x'$  is a linear combination of  $x$  and  $y$
- $y'$  is a linear combination of  $x$  and  $y$

What is the inverse transformation?

- Rotation by  $-\theta$
- For rotation matrices  $\mathbf{R}^{-1} = \mathbf{R}^T$

# Basic 2D transformations

---

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Scale**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & \alpha_x \\ \alpha_y & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Shear**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos\Theta & -\sin\Theta \\ \sin\Theta & \cos\Theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Rotate**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

**Translate**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

**Affine**

**Affine is any combination of translation, scale, rotation, shear**

# Affine Transformations

---

Affine transformations are combinations of

- Linear transformations, and
- Translations

Properties of affine transformations:

- Lines map to lines
- Parallel lines remain parallel
- Ratios are preserved
- Closed under composition

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

or

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Projective Transformations

---

Projective transformations are combos of

- Affine transformations, and
- Projective warps

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

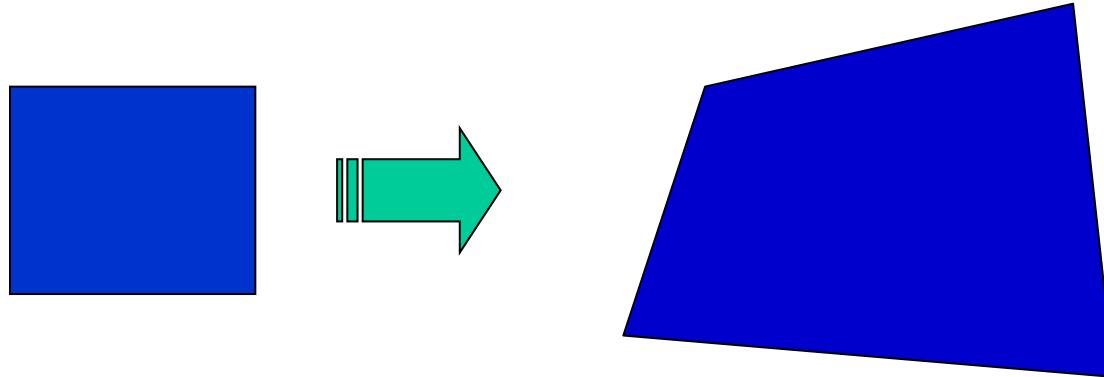
Properties of projective transformations:

- Lines map to lines
- Parallel lines do not necessarily remain parallel
- Ratios are not preserved
- Closed under composition
- **Projective matrix is defined up to a scale (8 DOF)**

# Fitting a plane projective transformation

---

- **Homography:** plane projective transformation  
(transformation taking a quad to another arbitrary quad)



# Homography

---

- The transformation between two views of a planar surface

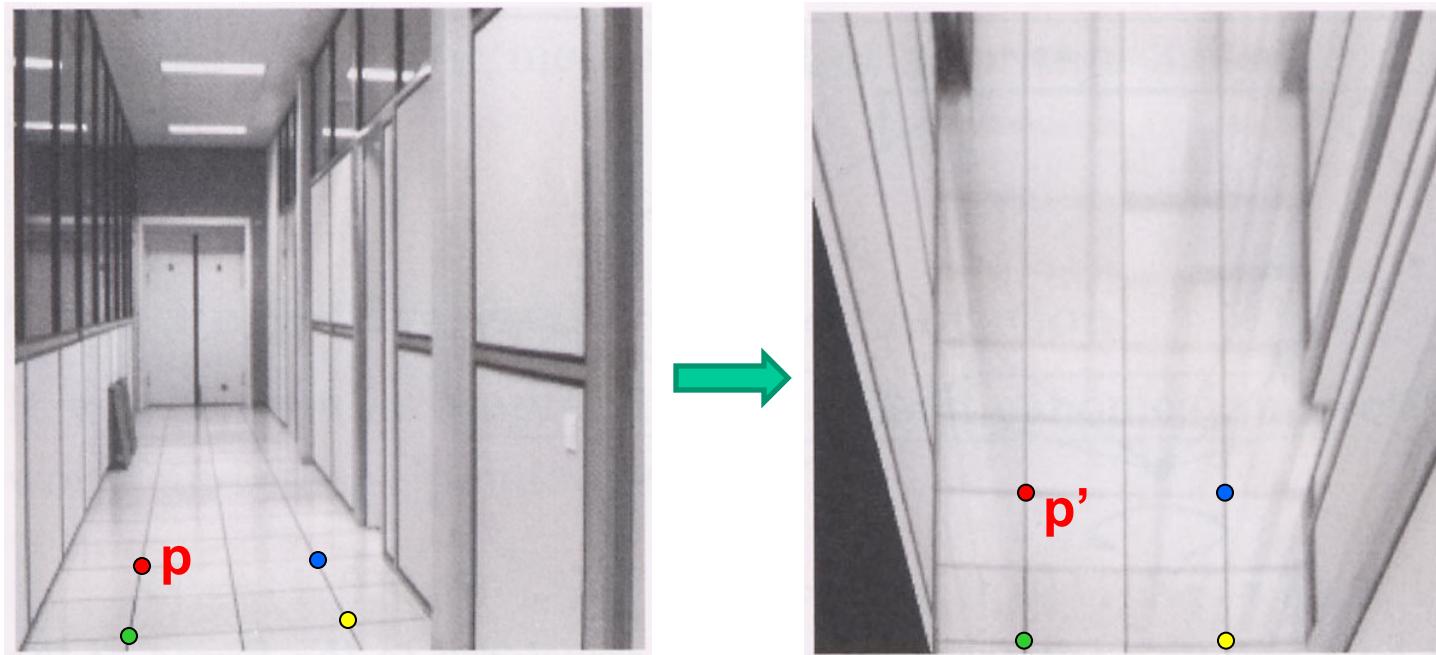


- The transformation between images from two cameras that share the same center



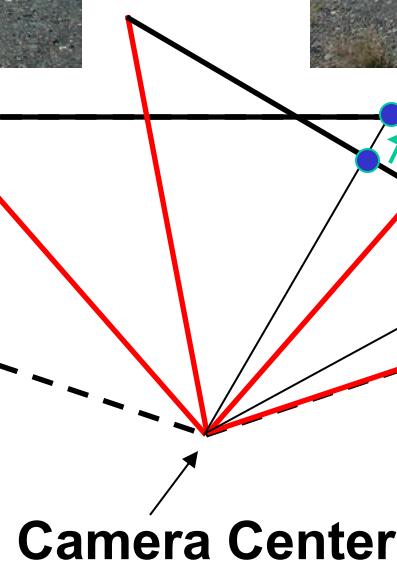
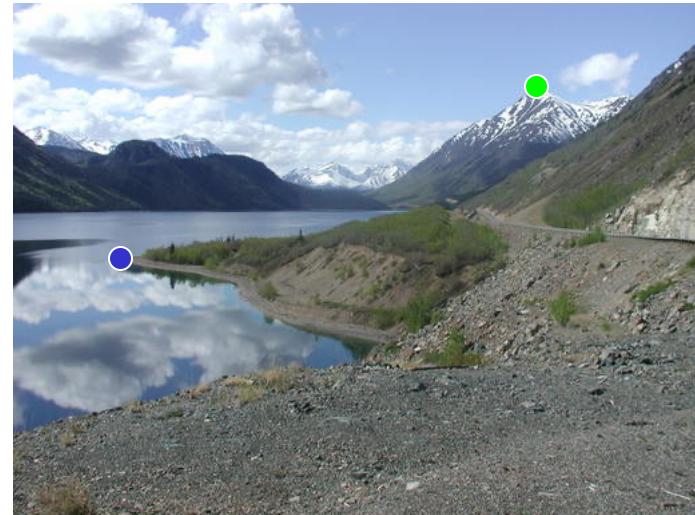
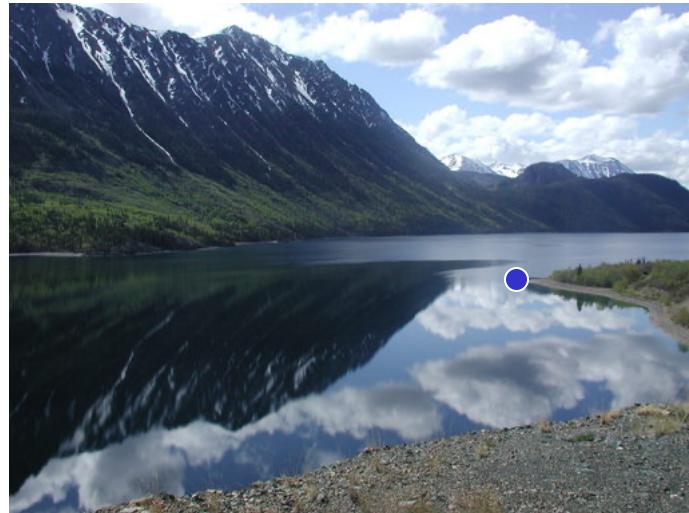
# Homography example: Image rectification

---



# Homography Example

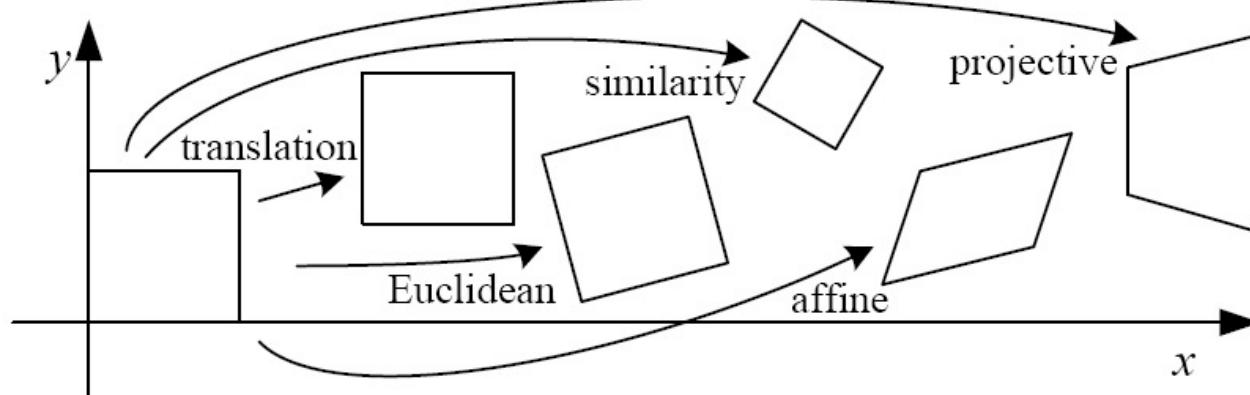
---



**Camera Center**

# 2D image transformations

---

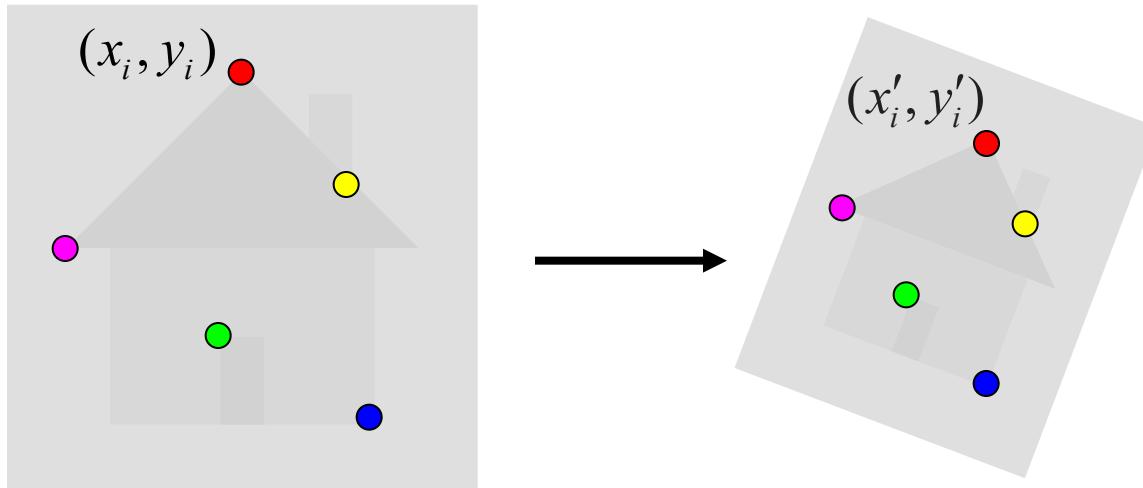


Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

# Fitting an affine transformation

---

- Assume we know the correspondences, how do we get the transformation?



$$\mathbf{x}'_i = \mathbf{M}\mathbf{x}_i + \mathbf{t}$$

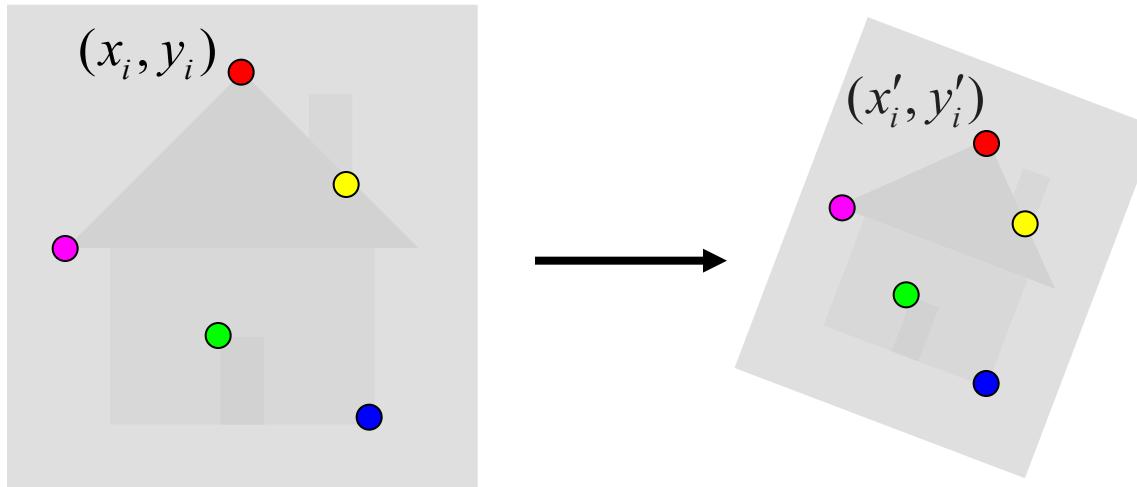
$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

Want to find  $\mathbf{M}$ ,  $\mathbf{t}$  to minimize

$$\sum_{i=1}^n \| \mathbf{x}'_i - \mathbf{M}\mathbf{x}_i - \mathbf{t} \|^2$$

# Fitting an affine transformation

- Assume we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\left[ \begin{array}{c} \vdots \\ m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{array} \right]$$

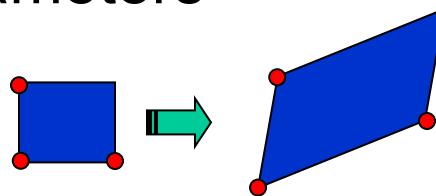
$$\left[ \begin{array}{c} \vdots \\ x'_i \\ y'_i \\ \dots \end{array} \right] = \left[ \begin{array}{c} \vdots \\ x'_i \\ y'_i \\ \dots \end{array} \right]$$

# Fitting an affine transformation

---

$$\begin{bmatrix} & & \cdots \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ & & \cdots \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \cdots \\ x'_i \\ y'_i \\ \cdots \end{bmatrix}$$

- Linear system with six unknowns
- Each match gives us two linearly independent equations: need at least three to solve for the transformation parameters



# Fitting a homography

---

- Recall: homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogeneous  
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogeneous  
image coordinates

# Fitting a homography

---

- Recall: homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogeneous  
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogeneous  
image coordinates

- Equation for homography:

$$\lambda \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Fitting a homography

---

- Equation for homography:

$$\lambda \begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \quad \lambda \mathbf{x}'_i = \mathbf{H} \mathbf{x}_i$$
$$\mathbf{x}'_i \times \mathbf{H} \mathbf{x}_i = 0$$

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} \times \begin{bmatrix} \mathbf{h}_1^T \mathbf{x}_i \\ \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_3^T \mathbf{x}_i \end{bmatrix} = \begin{bmatrix} y'_i \mathbf{h}_3^T \mathbf{x}_i - \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_1^T \mathbf{x}_i - x'_i \mathbf{h}_3^T \mathbf{x}_i \\ x'_i \mathbf{h}_2^T \mathbf{x}_i - y'_i \mathbf{h}_1^T \mathbf{x}_i \end{bmatrix}$$

$$\begin{bmatrix} 0^T & -\mathbf{x}_i^T & y'_i \mathbf{x}_i^T \\ \mathbf{x}_i^T & 0^T & -x'_i \mathbf{x}_i^T \\ -y'_i \mathbf{x}_i^T & x'_i \mathbf{x}_i^T & 0^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0$$

3 equations,  
only 2 linearly  
independent

# Fitting a homography

---

$$\begin{bmatrix} 0^T & \mathbf{x}_1^T & -y'_1 \mathbf{x}_1^T \\ \mathbf{x}_1^T & 0^T & -x'_1 \mathbf{x}_1^T \\ \dots & \dots & \dots \\ 0^T & \mathbf{x}_n^T & -y'_n \mathbf{x}_n^T \\ \mathbf{x}_n^T & 0^T & -x'_n \mathbf{x}_n^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0 \quad \mathbf{A h} = 0$$

- $\mathbf{H}$  has 8 degrees of freedom (9 parameters, but scale is arbitrary)
- One match gives us two linearly independent equations
- Homogeneous least squares: find  $\mathbf{h}$  minimizing  $\|\mathbf{Ah}\|^2$ 
  - Eigenvector of  $\mathbf{A}^T \mathbf{A}$  corresponding to smallest eigenvalue
  - Four matches needed for a minimal solution

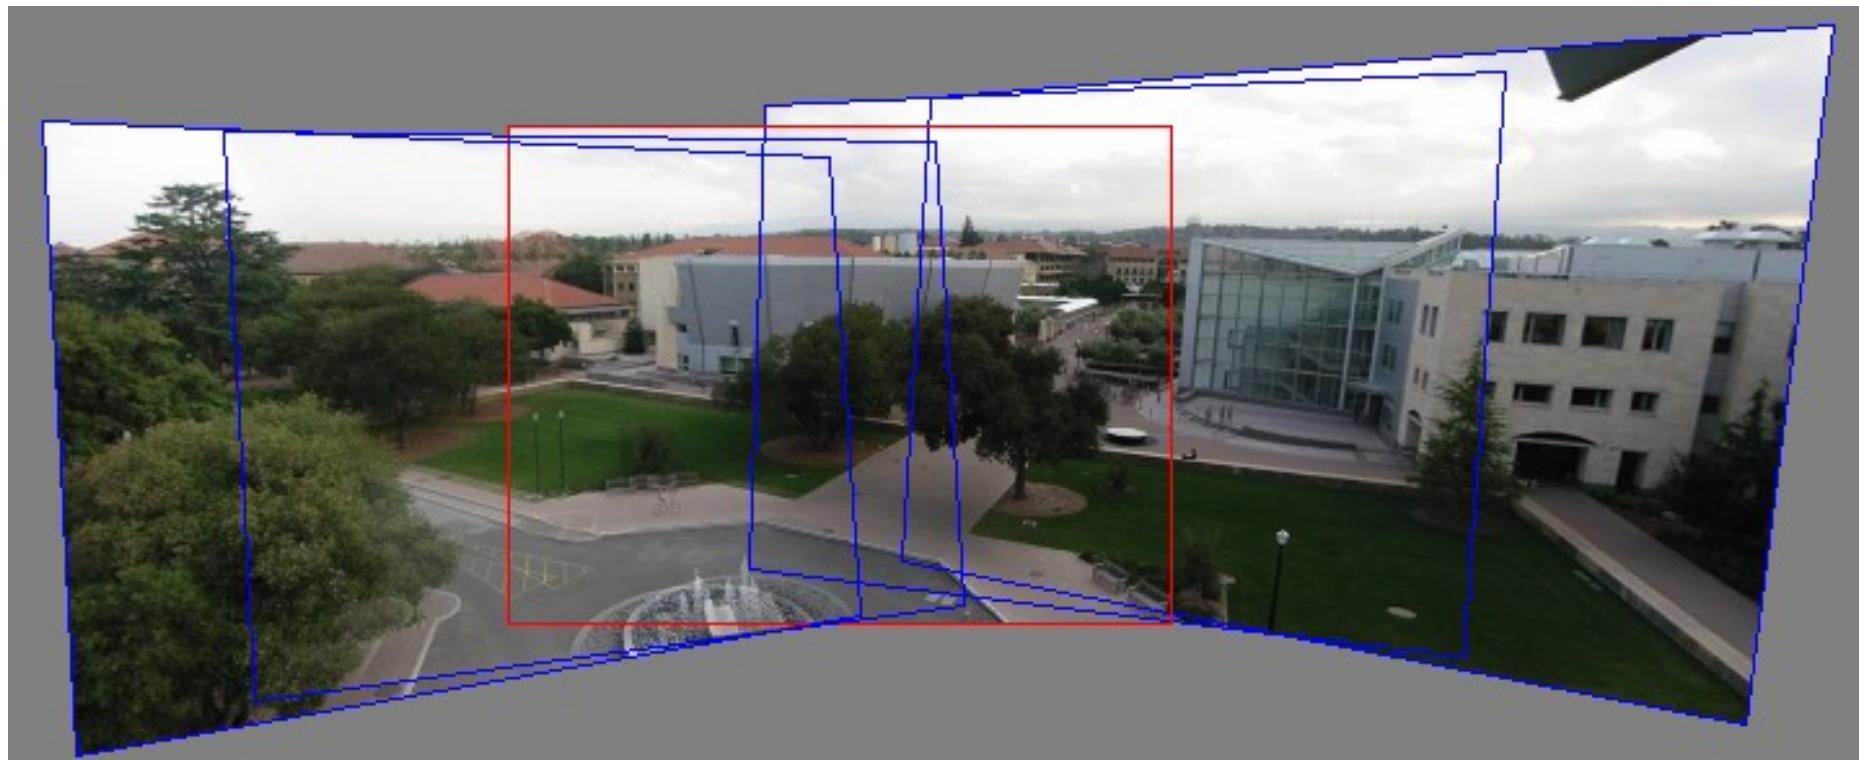
# Feature-based alignment: Overview

---

- Alignment as fitting
  - Affine transformations
  - Homographies
- Robust alignment
  - Descriptor-based feature matching
  - RANSAC
- Applications

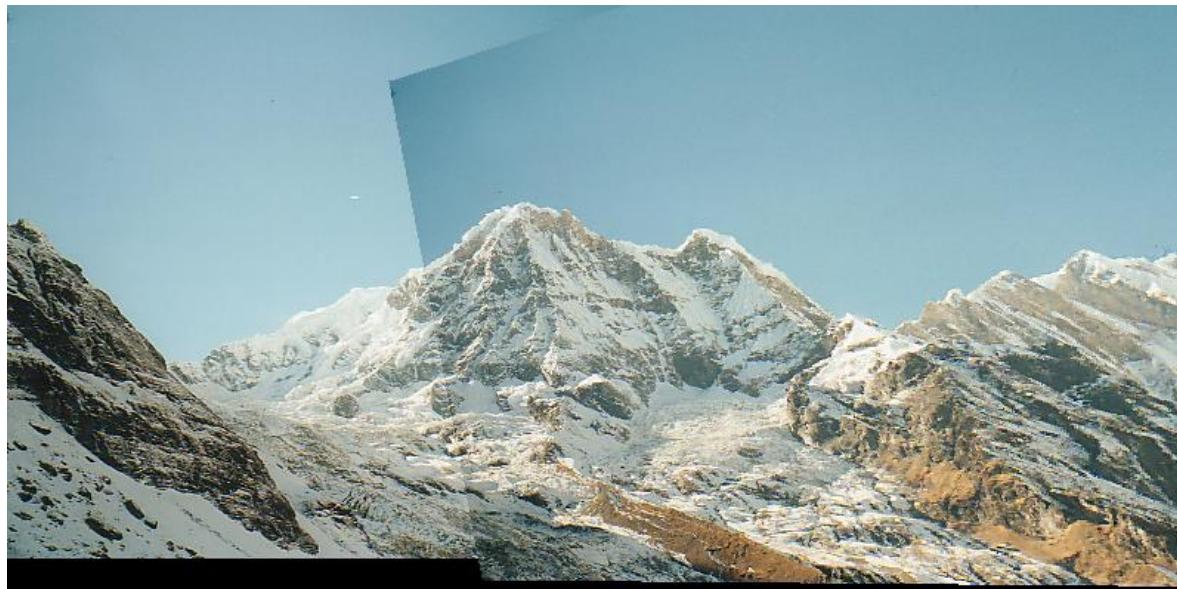
# Image alignment

---



# Image Alignment

---



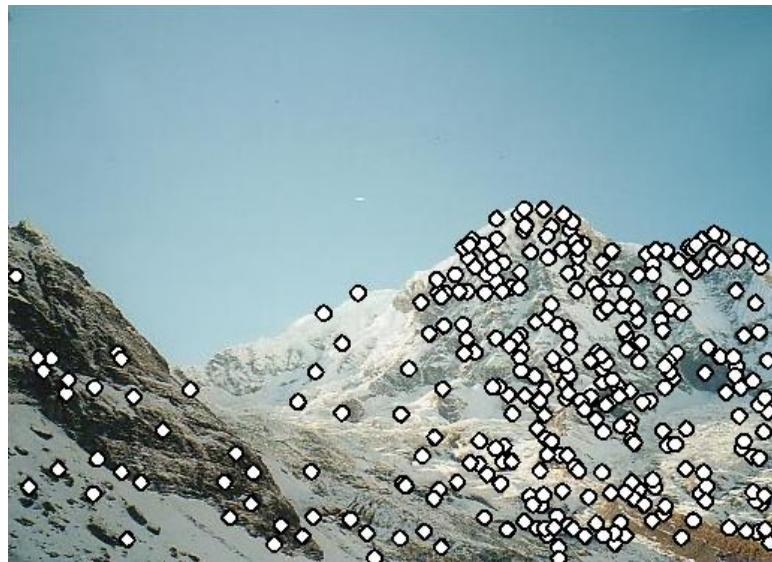
# Robust feature-based alignment

---



# Robust feature-based alignment

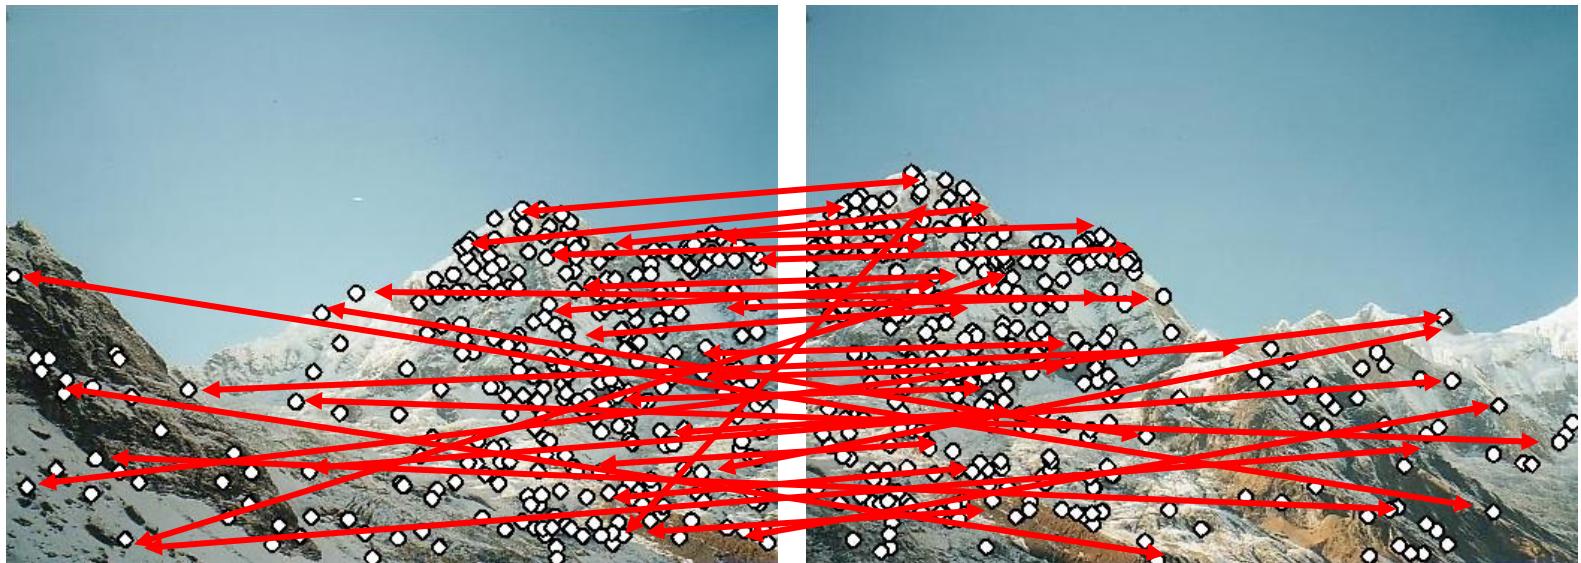
---



- Extract features

# Robust feature-based alignment

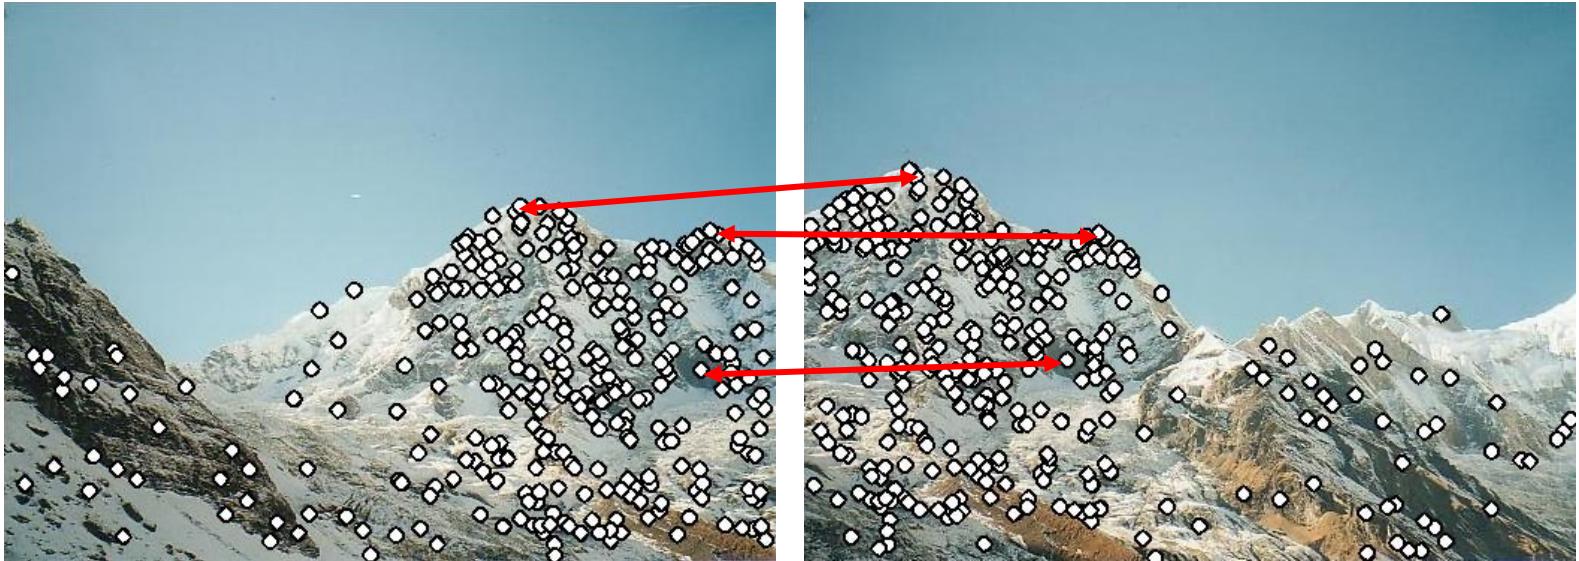
---



- Extract features
- Compute *putative matches*

# Robust feature-based alignment

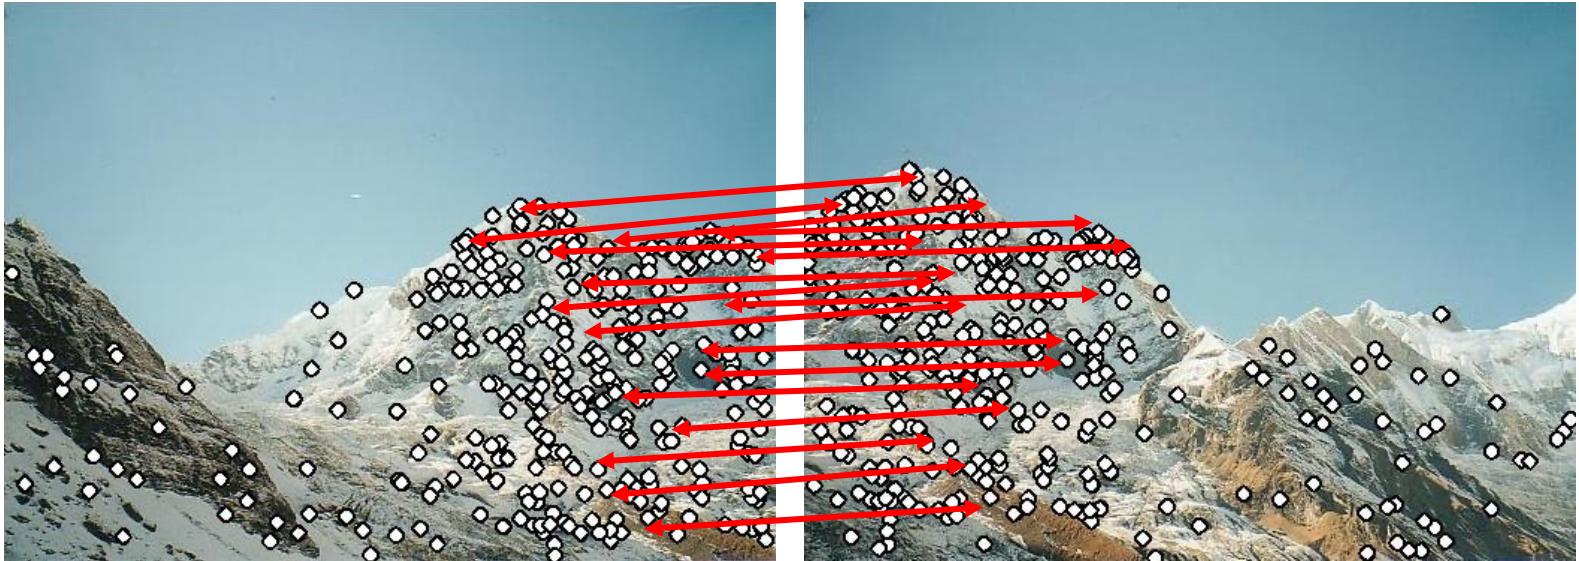
---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$

# Robust feature-based alignment

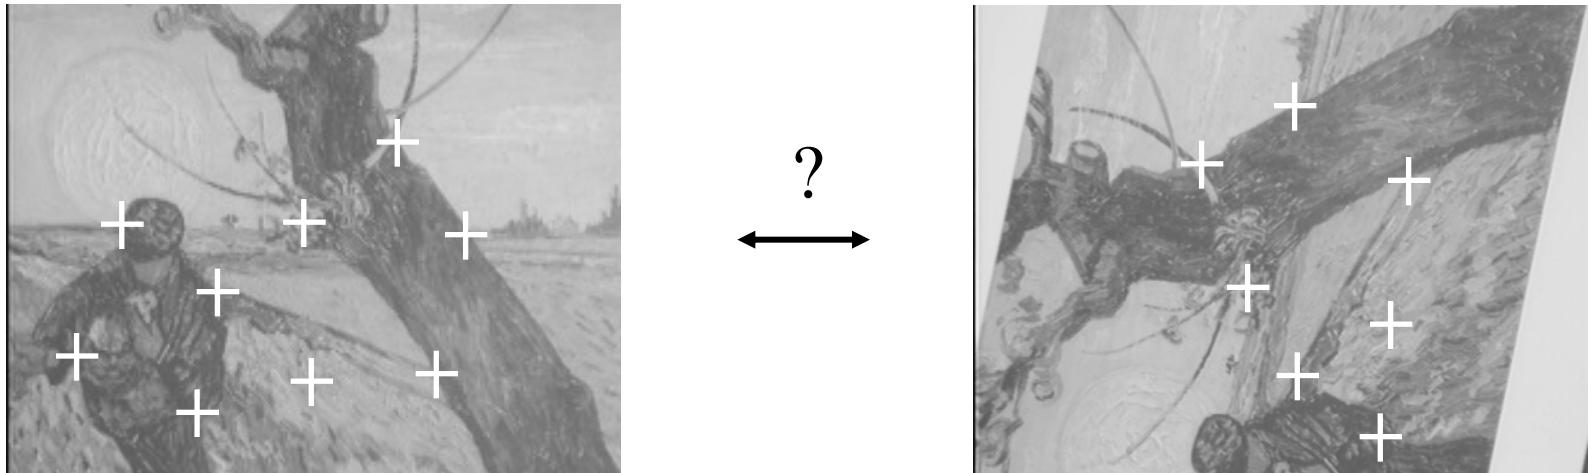
---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )

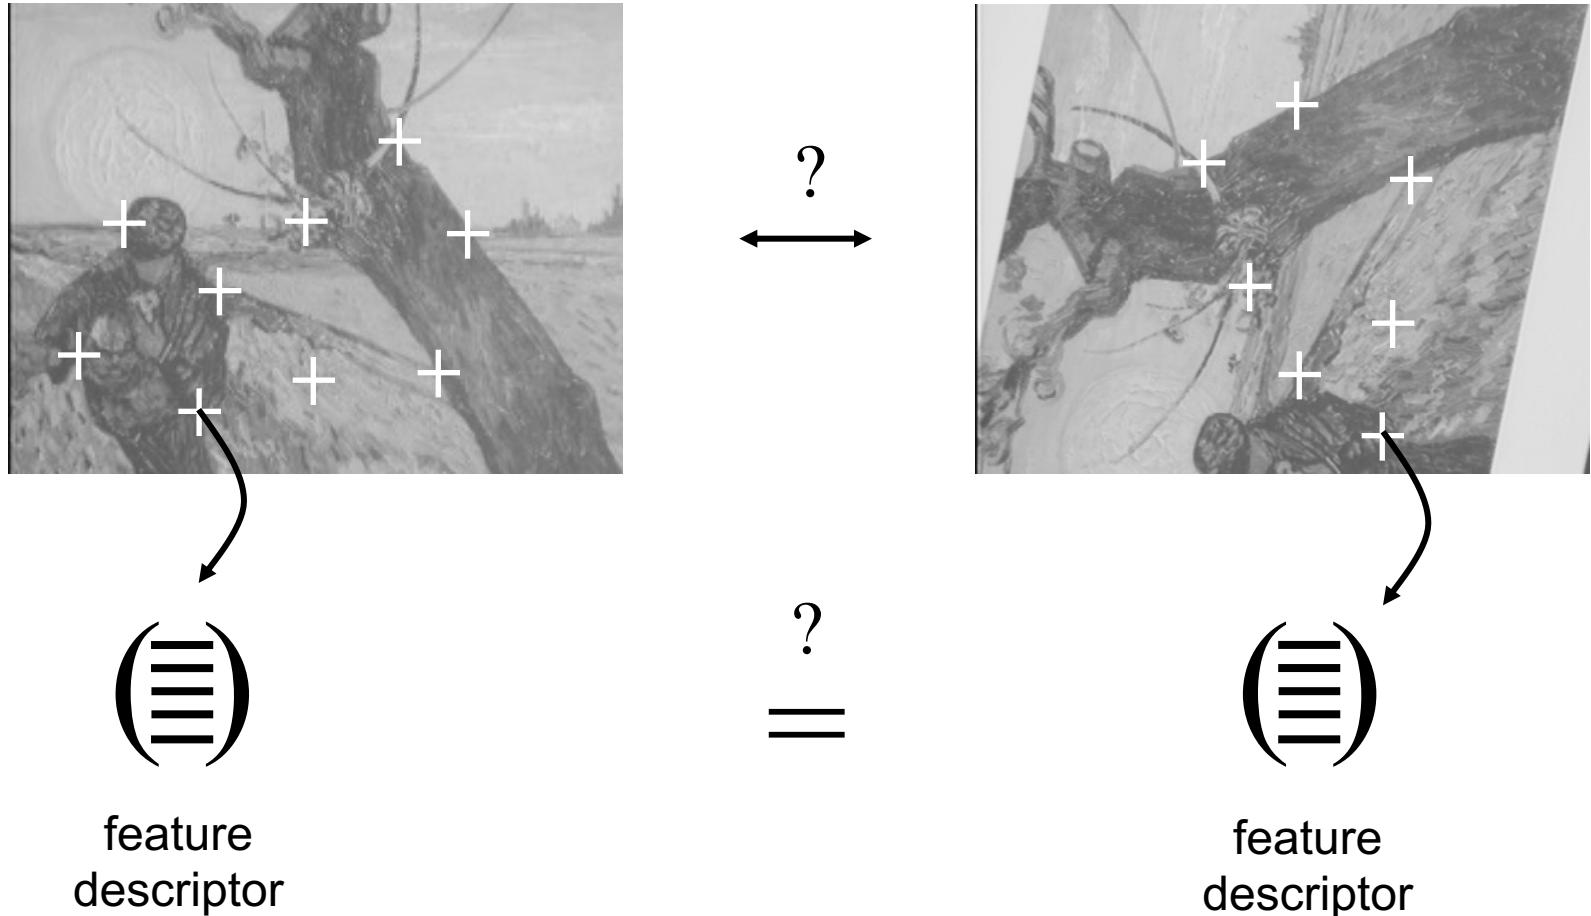
# Generating putative correspondences

---



# Generating putative correspondences

---

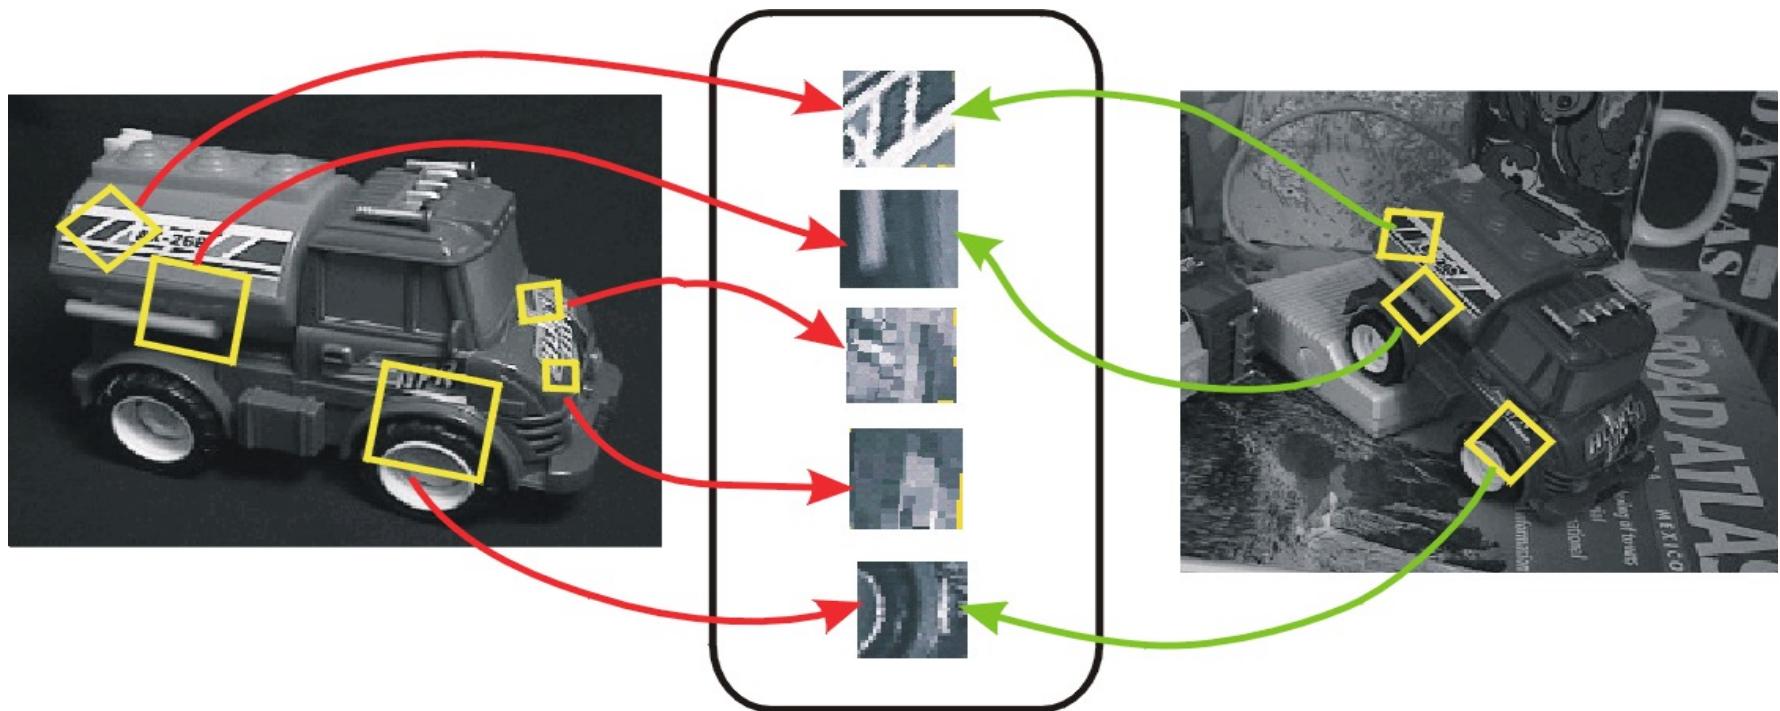


- Need to compare *feature descriptors* of local patches surrounding interest points

# Feature descriptors

---

- Recall: feature detection and description



# Feature descriptors

---

- Simplest descriptor: vector of raw intensity values
- How to compare two such vectors?
  - Sum of squared differences (SSD)

$$\text{SSD}(\mathbf{u}, \mathbf{v}) = \sum_i (u_i - v_i)^2$$

- Not invariant to intensity change
- Normalized correlation

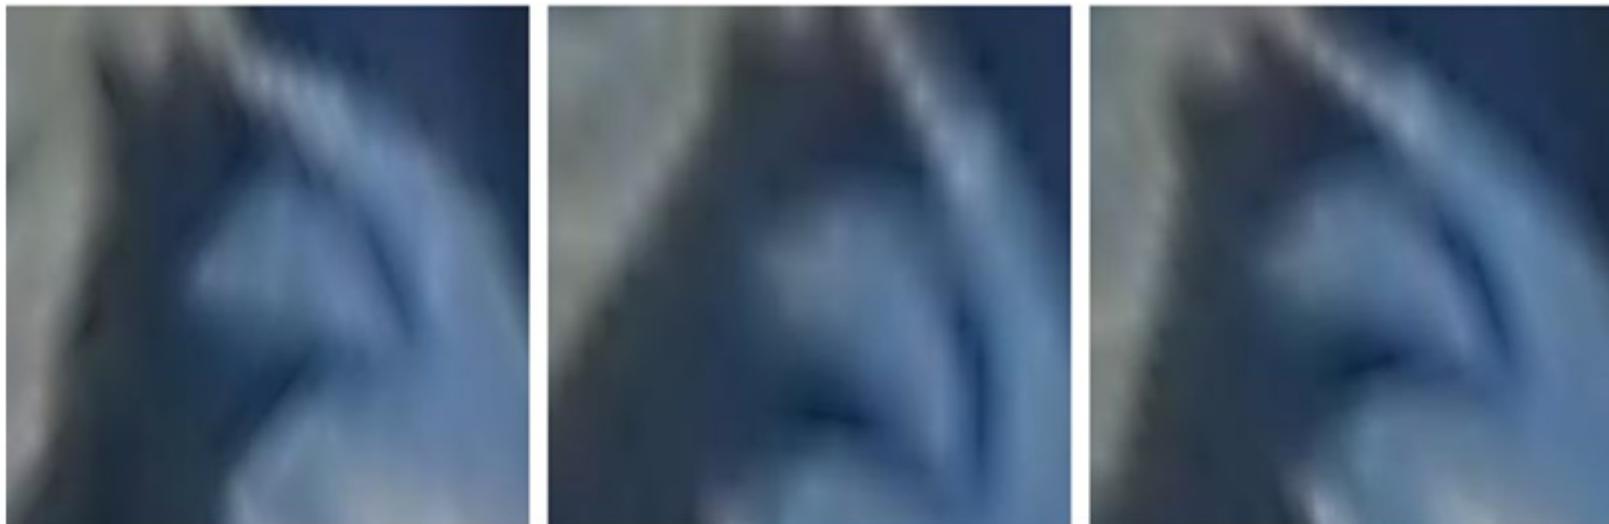
$$\rho(\mathbf{u}, \mathbf{v}) = \frac{(\mathbf{u} - \bar{\mathbf{u}}) \cdot (\mathbf{v} - \bar{\mathbf{v}})}{\|\mathbf{u} - \bar{\mathbf{u}}\| \|\mathbf{v} - \bar{\mathbf{v}}\|} = \frac{\sum_i (u_i - \bar{u})(v_i - \bar{v})}{\sqrt{\left( \sum_j (u_j - \bar{u})^2 \right) \left( \sum_j (v_j - \bar{v})^2 \right)}}$$

- Invariant to affine intensity change

# Disadvantage of intensity vectors as descriptors

---

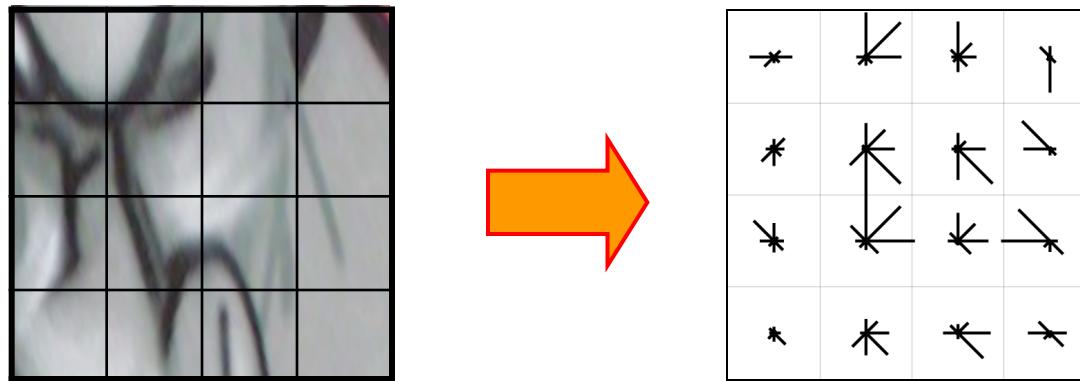
- Small deformations can affect the matching score a lot



# Feature descriptors: SIFT

---

- Descriptor computation:
  - Divide patch into  $4 \times 4$  sub-patches
  - Compute histogram of gradient orientations (8 reference angles) inside each sub-patch
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions



David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) IJCV 60 (2), pp. 91-110, 2004.

# Feature descriptors: SIFT

---

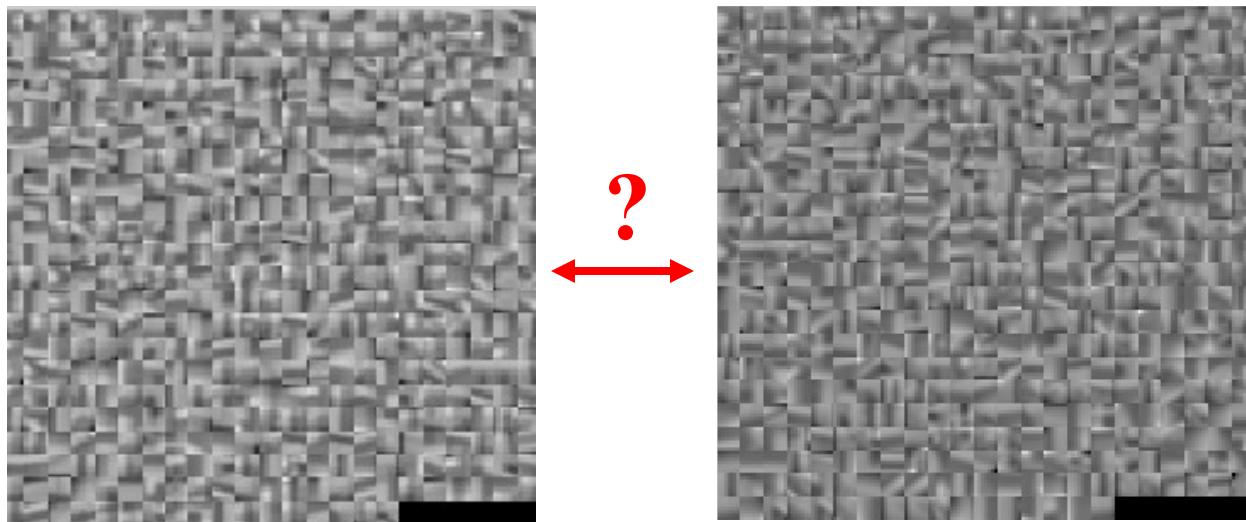
- Descriptor computation:
  - Divide patch into  $4 \times 4$  sub-patches
  - Compute histogram of gradient orientations (8 reference angles) inside each sub-patch
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions
- Advantage over raw vectors of pixel values
  - Gradients less sensitive to illumination change
  - Pooling of gradients over the sub-patches achieves robustness to small shifts, but still preserves some spatial information

David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" IJCV 60 (2), pp. 91-110, 2004.

# Feature matching

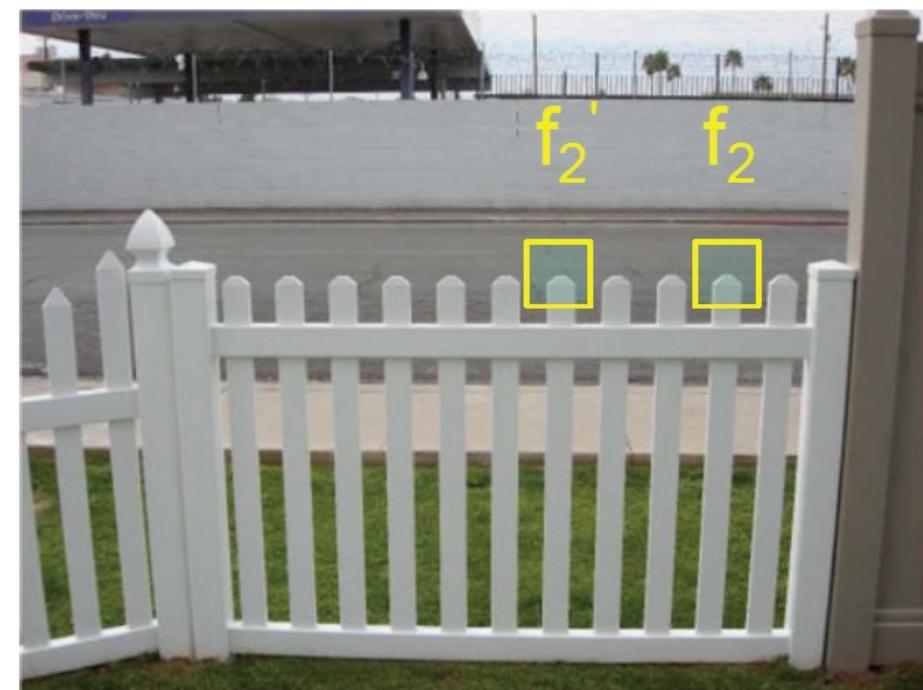
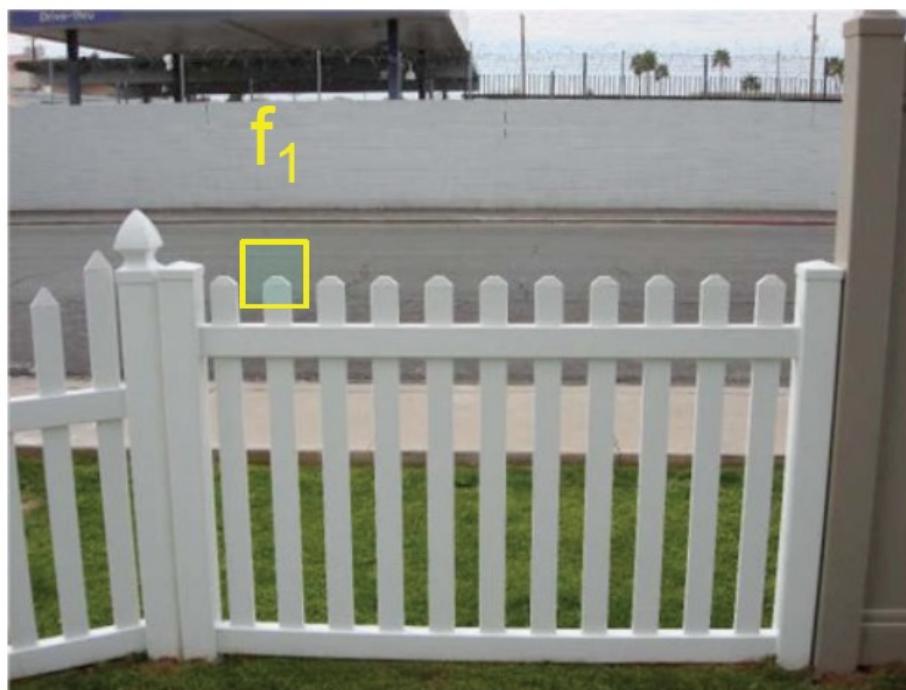
---

- Generating *putative matches*: for each patch in one image, find a short list of patches in the other image that could match it based solely on appearance



# Problem: Ambiguous putative matches

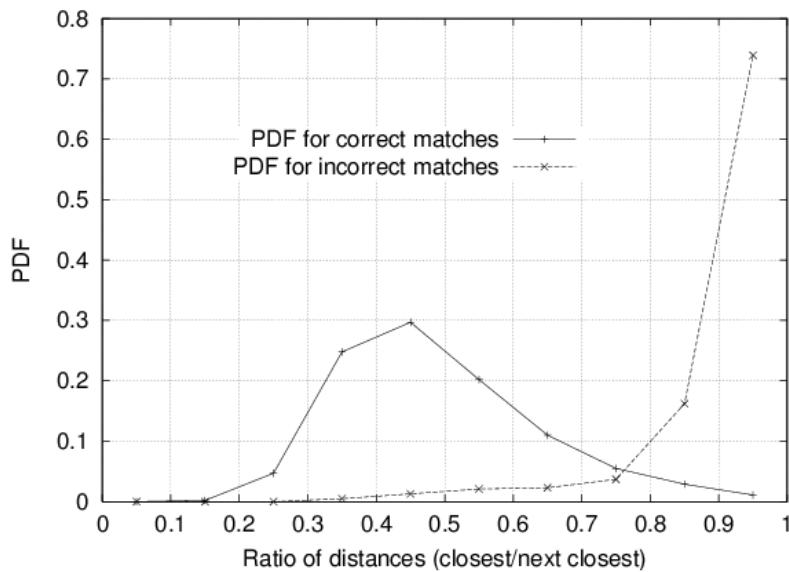
---



# Rejection of unreliable matches

---

- How can we tell which putative matches are more reliable?
- Heuristic: compare distance of **nearest** neighbor to that of **second** nearest neighbor
  - Ratio of closest distance to second-closest distance will be *high* for features that are *not* distinctive



**Threshold of 0.8 provides good separation**

# RANSAC

---

- The set of putative matches contains a very high percentage of outliers

## RANSAC loop:

1. Randomly select a *seed group* of matches
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers

Keep the transformation with the largest number of inliers

# Robust feature-based alignment

---



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )

# Common transformations

---



original

Transformed



translation



rotation



aspect



affine



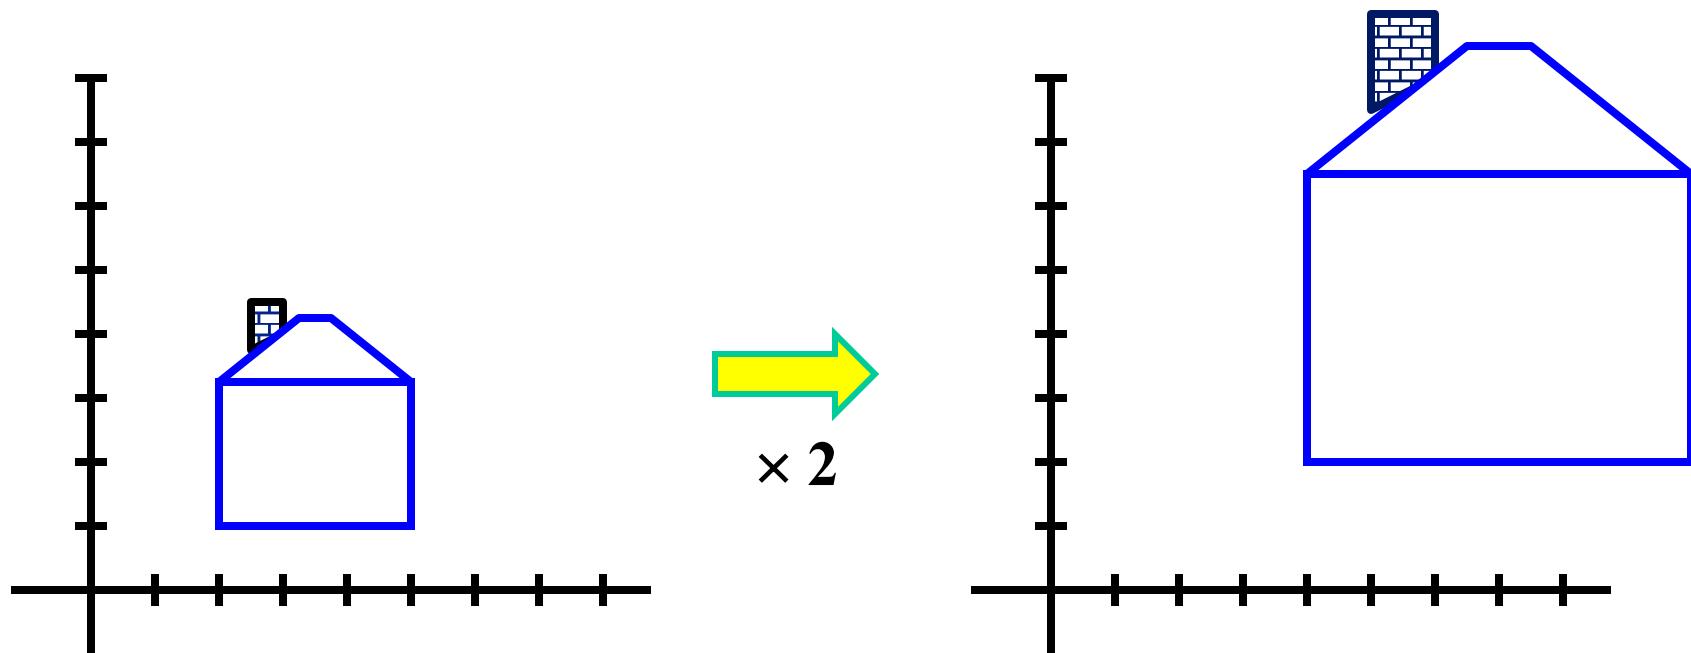
perspective

# Scaling

---

*Scaling* a coordinate means multiplying each of its components by a scalar

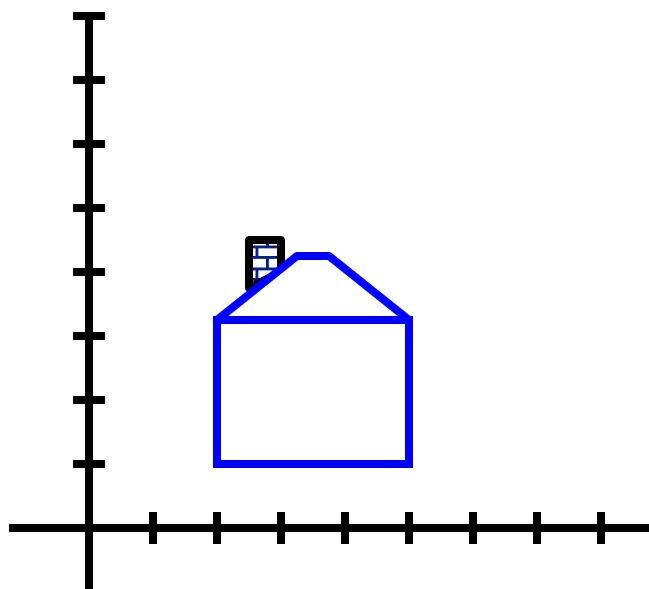
*Uniform scaling* means this scalar is the same for all components:



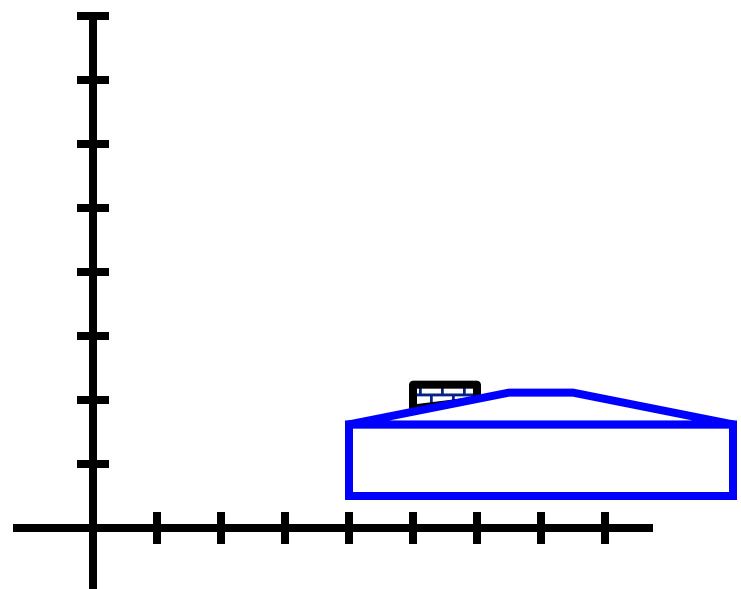
# Scaling

---

*Non-uniform scaling*: different scalars per component:



$\rightarrow$   
 $X \times 2,$   
 $Y \times 0.5$



# Scaling

---

Scaling operation:

$$x' = ax$$

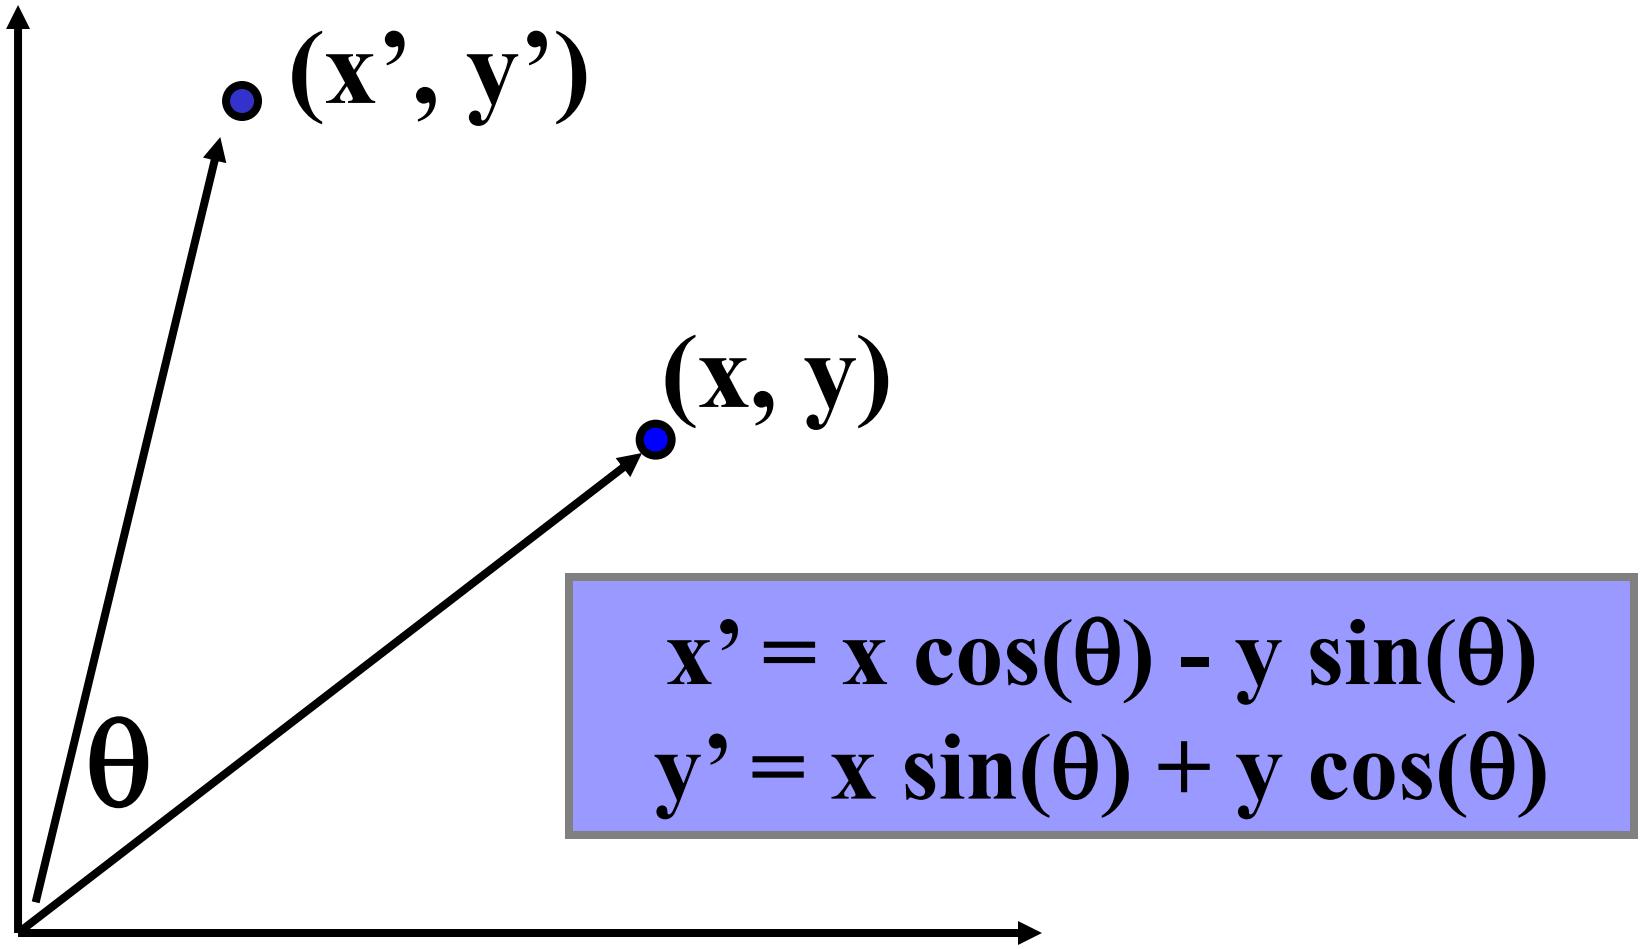
$$y' = by$$

Or, in matrix form:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \underbrace{\begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}}_{\text{scaling matrix } S} \begin{bmatrix} x \\ y \end{bmatrix}$$

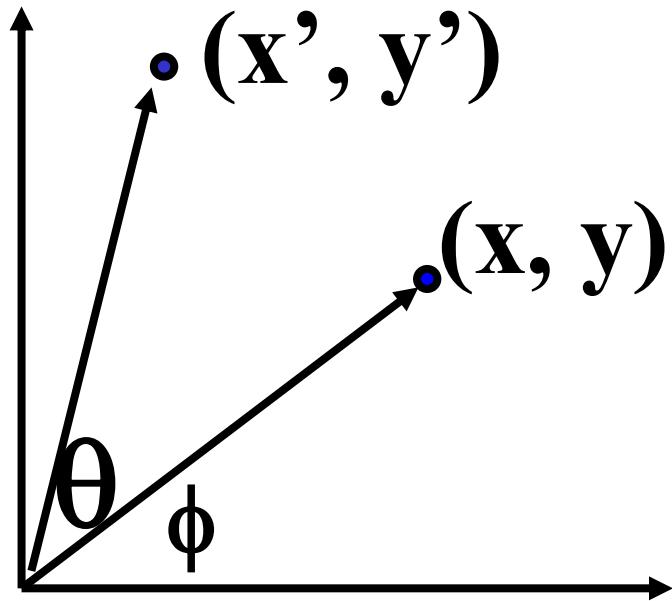
# 2-D Rotation

---



# 2-D Rotation

---



Polar coordinates...

$$x = r \cos(\phi)$$

$$y = r \sin(\phi)$$

$$x' = r \cos(\phi + \theta)$$

$$y' = r \sin(\phi + \theta)$$

Trig Identity...

$$x' = r \cos(\phi) \cos(\theta) - r \sin(\phi) \sin(\theta)$$

$$y' = r \sin(\phi) \cos(\theta) + r \cos(\phi) \sin(\theta)$$

Substitute...

$$x' = x \cos(\theta) - y \sin(\theta)$$

$$y' = x \sin(\theta) + y \cos(\theta)$$

# 2-D Rotation

---

This is easy to capture in matrix form:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \underbrace{\begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix}}_{\mathbf{R}} \begin{bmatrix} x \\ y \end{bmatrix}$$

Even though  $\sin(\theta)$  and  $\cos(\theta)$  are nonlinear functions of  $\theta$ ,

- $x'$  is a linear combination of  $x$  and  $y$
- $y'$  is a linear combination of  $x$  and  $y$

What is the inverse transformation?

- Rotation by  $-\theta$
- For rotation matrices  $\mathbf{R}^{-1} = \mathbf{R}^T$

# Basic 2D transformations

---

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Scale**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & \alpha_x \\ \alpha_y & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Shear**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos\Theta & -\sin\Theta \\ \sin\Theta & \cos\Theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

**Rotate**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

**Translate**

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

**Affine**

**Affine is any combination of translation, scale, rotation, shear**

# Affine Transformations

---

Affine transformations are combinations of

- Linear transformations, and
- Translations

Properties of affine transformations:

- Lines map to lines
- Parallel lines remain parallel
- Ratios are preserved
- Closed under composition

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

or

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Projective Transformations

---

Projective transformations are combos of

- Affine transformations, and
- Projective warps

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

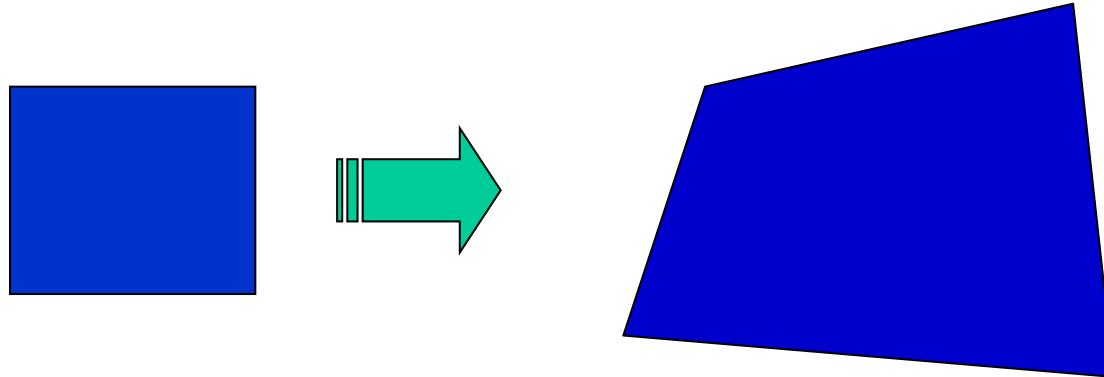
Properties of projective transformations:

- Lines map to lines
- Parallel lines do not necessarily remain parallel
- Ratios are not preserved
- Closed under composition
- **Projective matrix is defined up to a scale (8 DOF)**

# Fitting a plane projective transformation

---

- **Homography:** plane projective transformation  
(transformation taking a quad to another arbitrary quad)



# Homography

---

- The transformation between two views of a planar surface

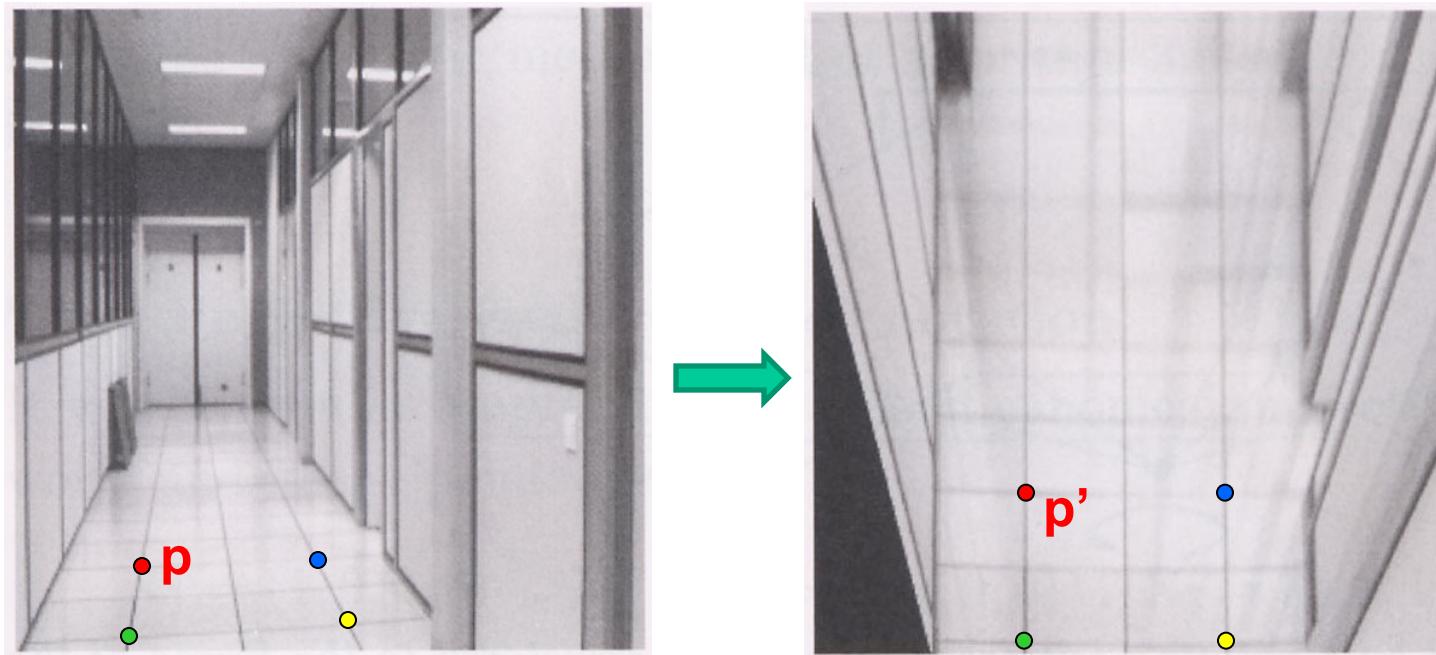


- The transformation between images from two cameras that share the same center



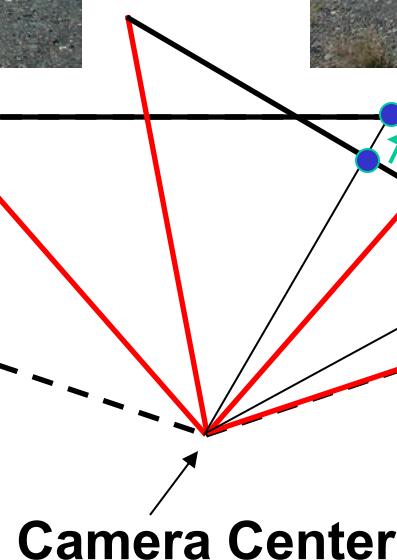
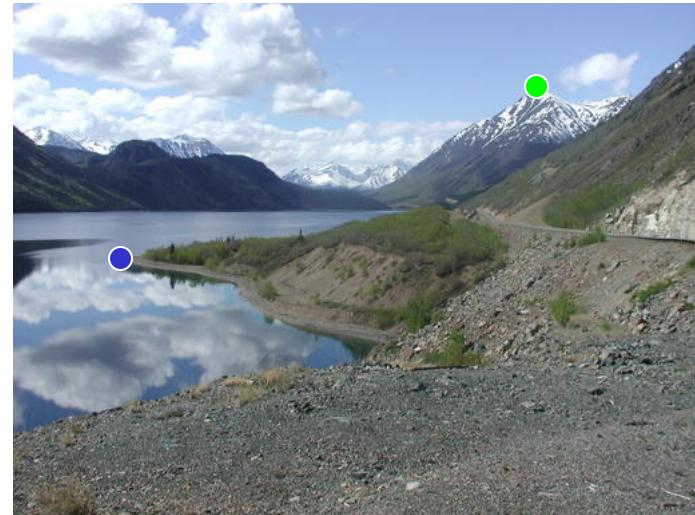
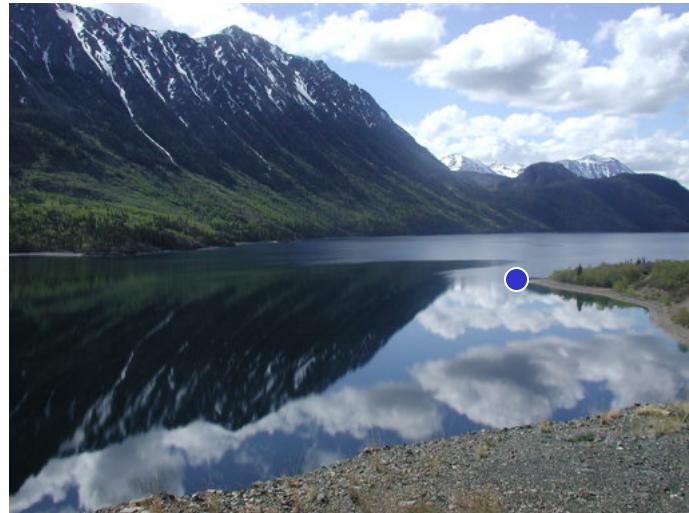
# Homography example: Image rectification

---



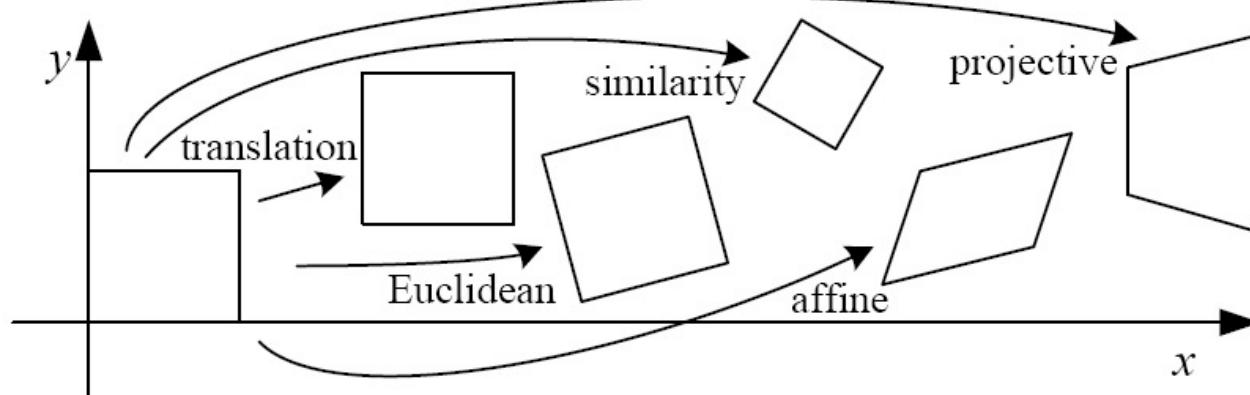
# Homography Example

---



# 2D image transformations

---

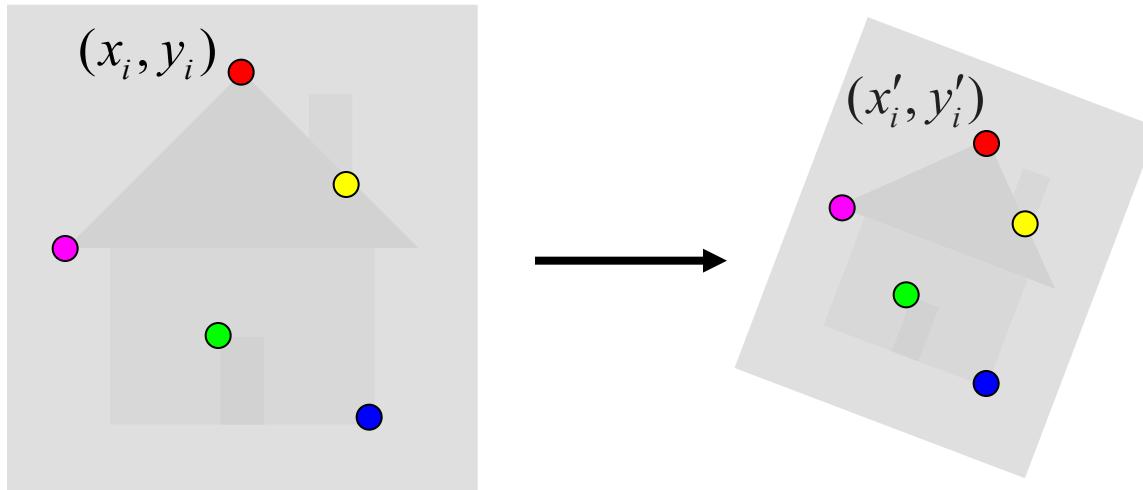


Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

# Fitting an affine transformation

---

- Assume we know the correspondences, how do we get the transformation?



$$\mathbf{x}'_i = \mathbf{M}\mathbf{x}_i + \mathbf{t}$$

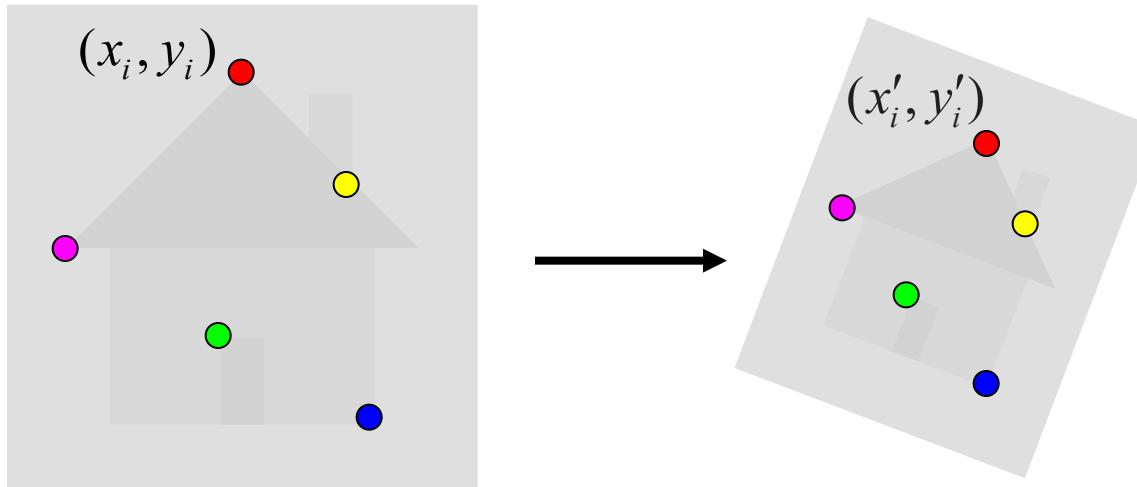
$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

Want to find  $\mathbf{M}$ ,  $\mathbf{t}$  to minimize

$$\sum_{i=1}^n \| \mathbf{x}'_i - \mathbf{M}\mathbf{x}_i - \mathbf{t} \|^2$$

# Fitting an affine transformation

- Assume we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\left[ \begin{array}{c} \vdots \\ m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{array} \right]$$

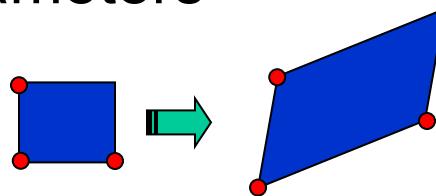
$$\left[ \begin{array}{c} \vdots \\ x'_i \\ y'_i \\ \dots \end{array} \right] = \left[ \begin{array}{c} \vdots \\ x'_i \\ y'_i \\ \dots \end{array} \right]$$

# Fitting an affine transformation

---

$$\begin{bmatrix} & & \cdots \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ & & \cdots \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \cdots \\ x'_i \\ y'_i \\ \cdots \end{bmatrix}$$

- Linear system with six unknowns
- Each match gives us two linearly independent equations: need at least three to solve for the transformation parameters



# Fitting a homography

---

- Recall: homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogeneous  
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogeneous  
image coordinates

# Fitting a homography

---

- Recall: homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogeneous  
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogeneous  
image coordinates

- Equation for homography:

$$\lambda \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Fitting a homography

---

- Equation for homography:

$$\lambda \begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \quad \lambda \mathbf{x}'_i = \mathbf{H} \mathbf{x}_i$$
$$\mathbf{x}'_i \times \mathbf{H} \mathbf{x}_i = 0$$

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} \times \begin{bmatrix} \mathbf{h}_1^T \mathbf{x}_i \\ \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_3^T \mathbf{x}_i \end{bmatrix} = \begin{bmatrix} y'_i \mathbf{h}_3^T \mathbf{x}_i - \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_1^T \mathbf{x}_i - x'_i \mathbf{h}_3^T \mathbf{x}_i \\ x'_i \mathbf{h}_2^T \mathbf{x}_i - y'_i \mathbf{h}_1^T \mathbf{x}_i \end{bmatrix}$$

$$\begin{bmatrix} 0^T & -\mathbf{x}_i^T & y'_i \mathbf{x}_i^T \\ \mathbf{x}_i^T & 0^T & -x'_i \mathbf{x}_i^T \\ -y'_i \mathbf{x}_i^T & x'_i \mathbf{x}_i^T & 0^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0$$

3 equations,  
only 2 linearly  
independent

# Fitting a homography

---

$$\begin{bmatrix} 0^T & \mathbf{x}_1^T & -y'_1 \mathbf{x}_1^T \\ \mathbf{x}_1^T & 0^T & -x'_1 \mathbf{x}_1^T \\ \dots & \dots & \dots \\ 0^T & \mathbf{x}_n^T & -y'_n \mathbf{x}_n^T \\ \mathbf{x}_n^T & 0^T & -x'_n \mathbf{x}_n^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0 \quad \mathbf{A h} = 0$$

- $\mathbf{H}$  has 8 degrees of freedom (9 parameters, but scale is arbitrary)
- One match gives us two linearly independent equations
- Homogeneous least squares: find  $\mathbf{h}$  minimizing  $\|\mathbf{Ah}\|^2$ 
  - Eigenvector of  $\mathbf{A}^T \mathbf{A}$  corresponding to smallest eigenvalue
  - Four matches needed for a minimal solution

# Feature-based alignment: Overview

---

- Alignment as fitting
  - Affine transformations
  - Homographies
- Robust alignment
  - Descriptor-based feature matching
  - RANSAC
- Applications