
Light, Camera and Shading

CS 543 / ECE 549 – Saurabh Gupta

Overview

- Cameras with lenses
 - Depth of field
 - Field of view
 - Lens aberrations
- Brightness of a pixel
 - Small taste of radiometry
 - In-camera transformation of light
 - Reflectance properties of surfaces
 - Lambertian reflection model
 - Shape from shading

Building a Real Camera

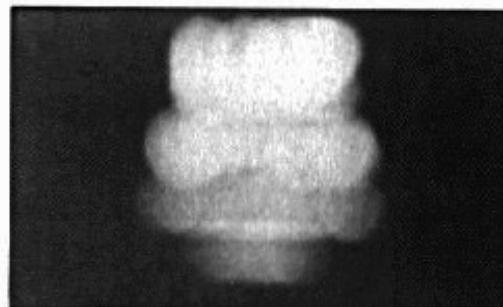


Home-made pinhole camera



<http://www.debevec.org/Pinhole/>

Shrinking the aperture



2 mm



1 mm



0.6mm

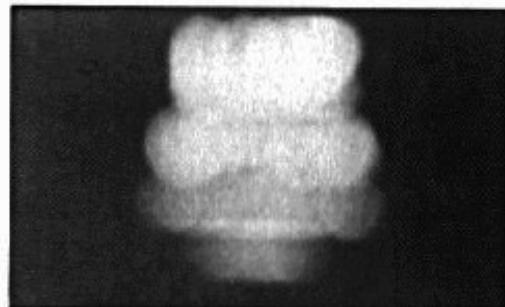


0.35 mm

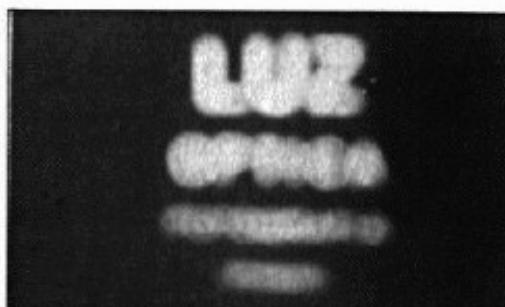
Why not make the aperture as small as possible?

- Less light gets through
- Diffraction effects...

Shrinking the aperture



2 mm



1 mm



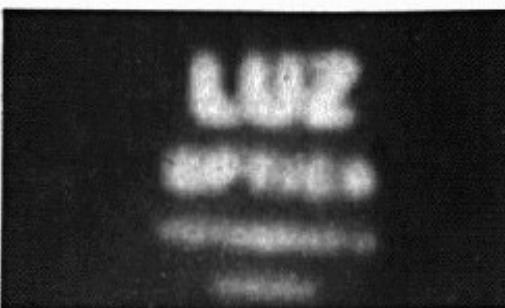
0.6mm



0.35 mm



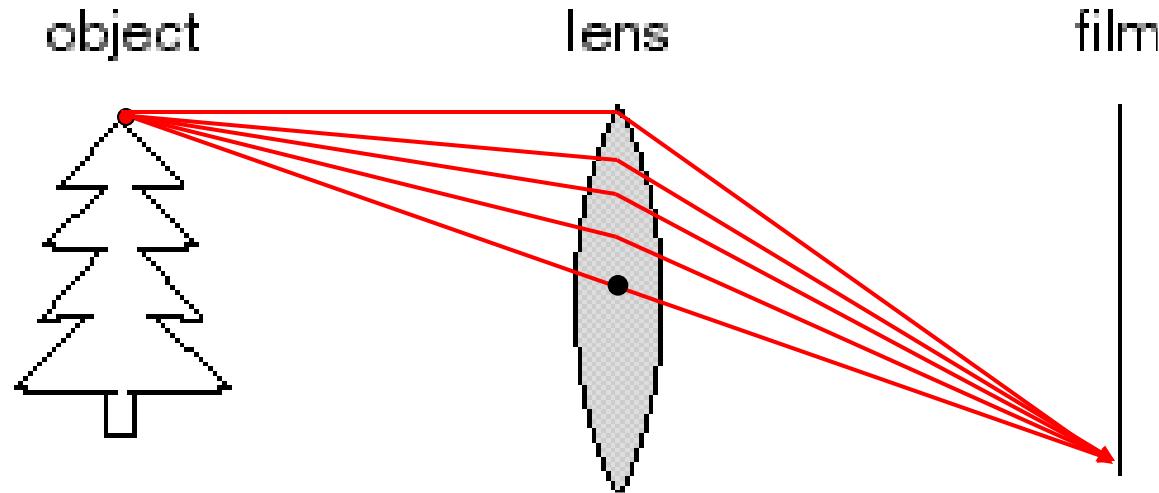
0.15 mm



0.07 mm

Adding a lens

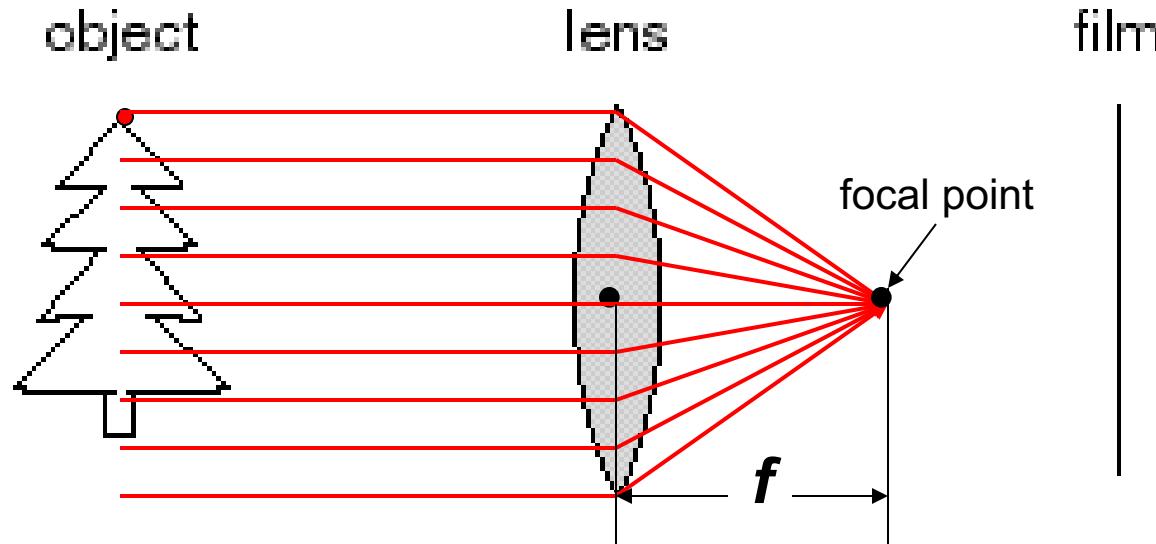
Adding a lens



A lens focuses light onto the film

- Thin lens model:
 - Rays passing through the center are not deviated (pinhole projection model still holds)

Adding a lens

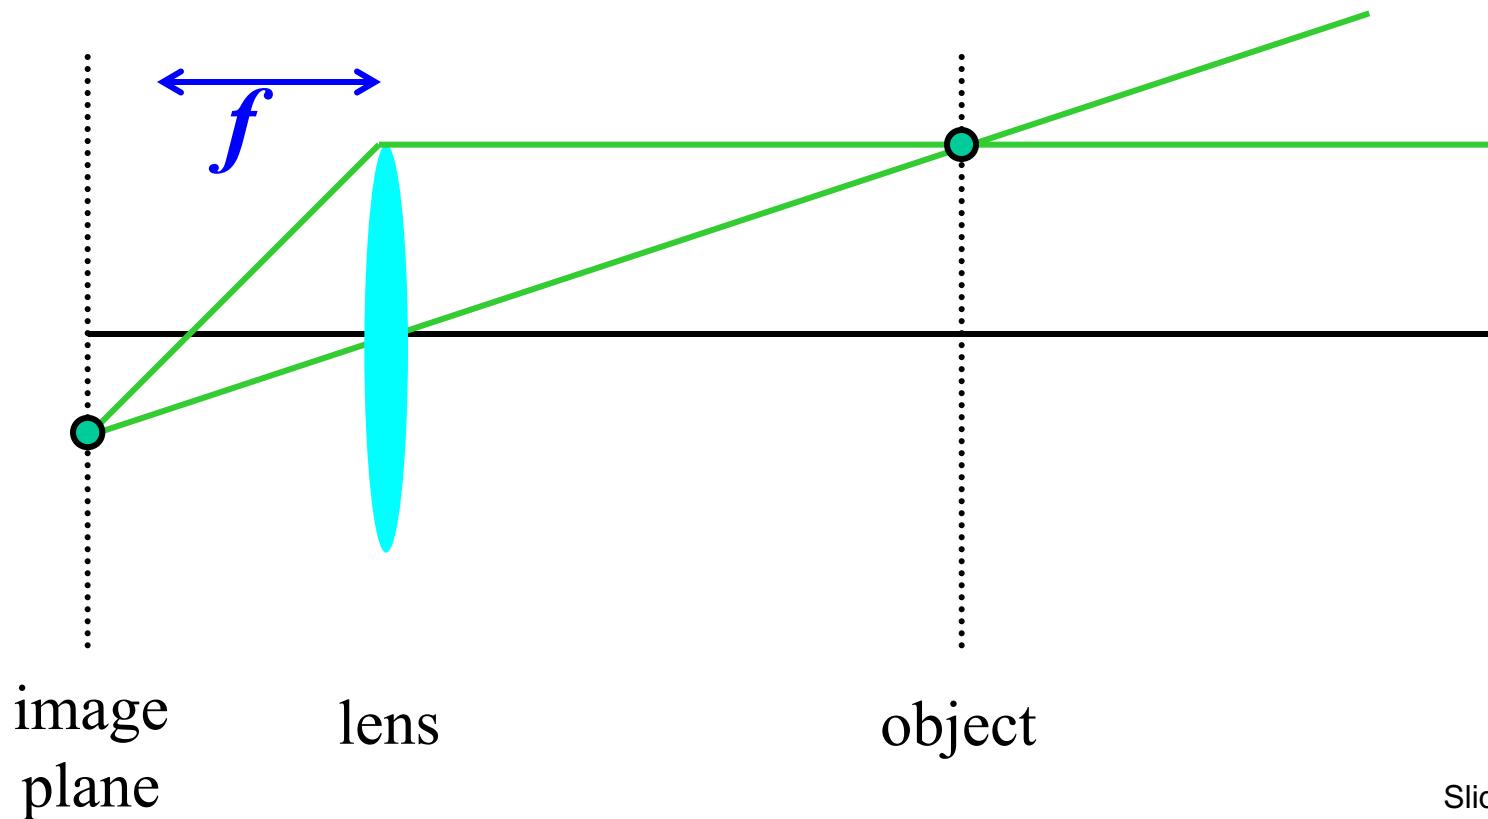


A lens focuses light onto the film

- Thin lens model:
 - Rays passing through the center are not deviated (pinhole projection model still holds)
 - All rays parallel to the optical axis pass through the *focal point*
 - All parallel rays converge to points on the *focal plane*

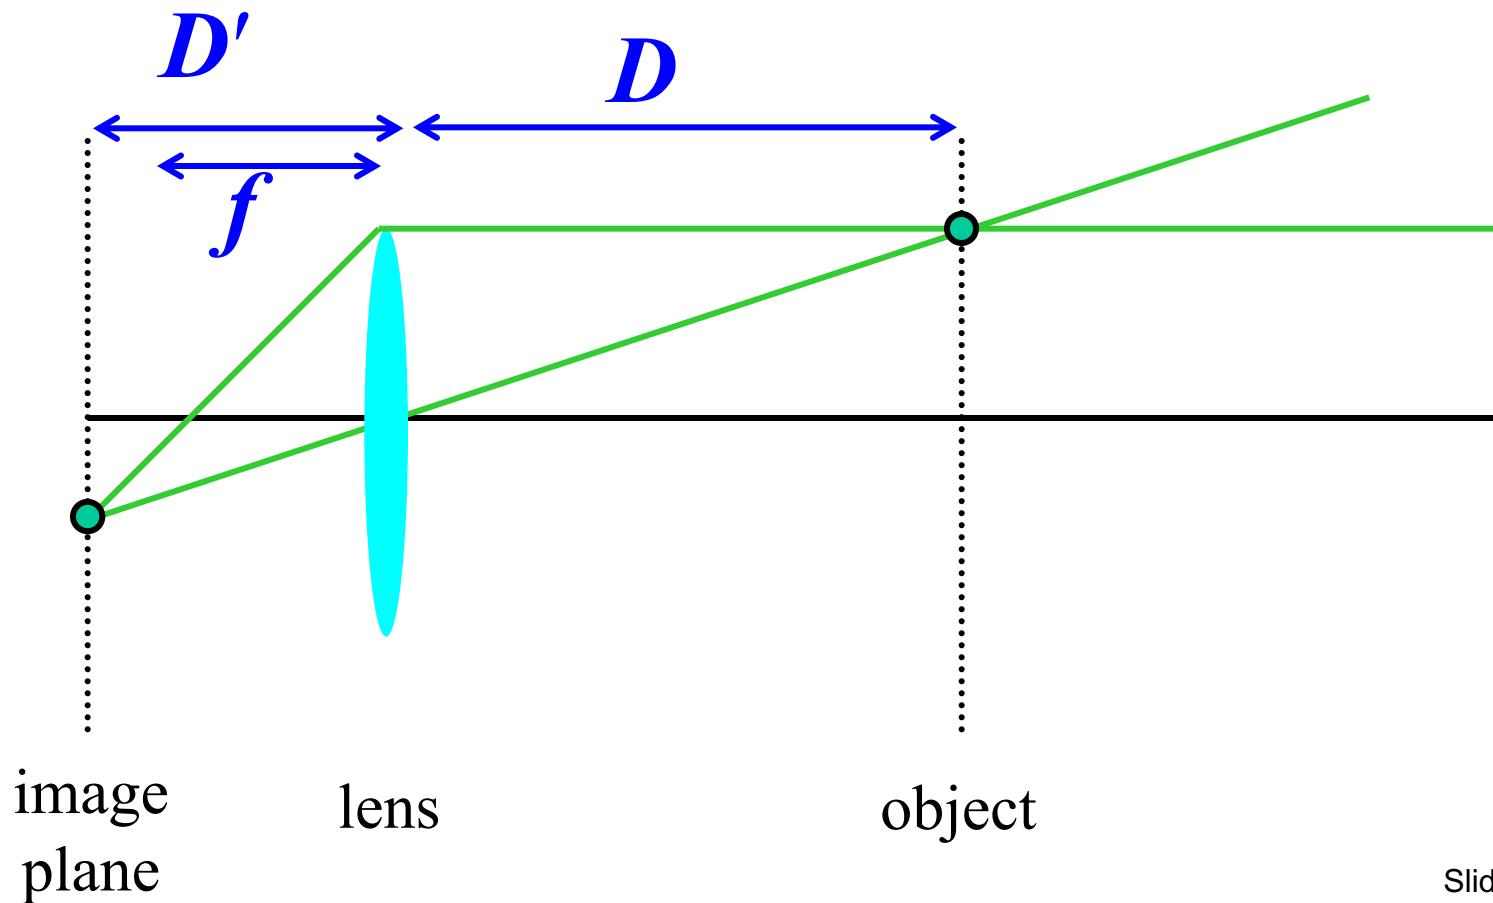
Thin lens formula

- Where does the lens focus the rays coming from a given point in the scene?



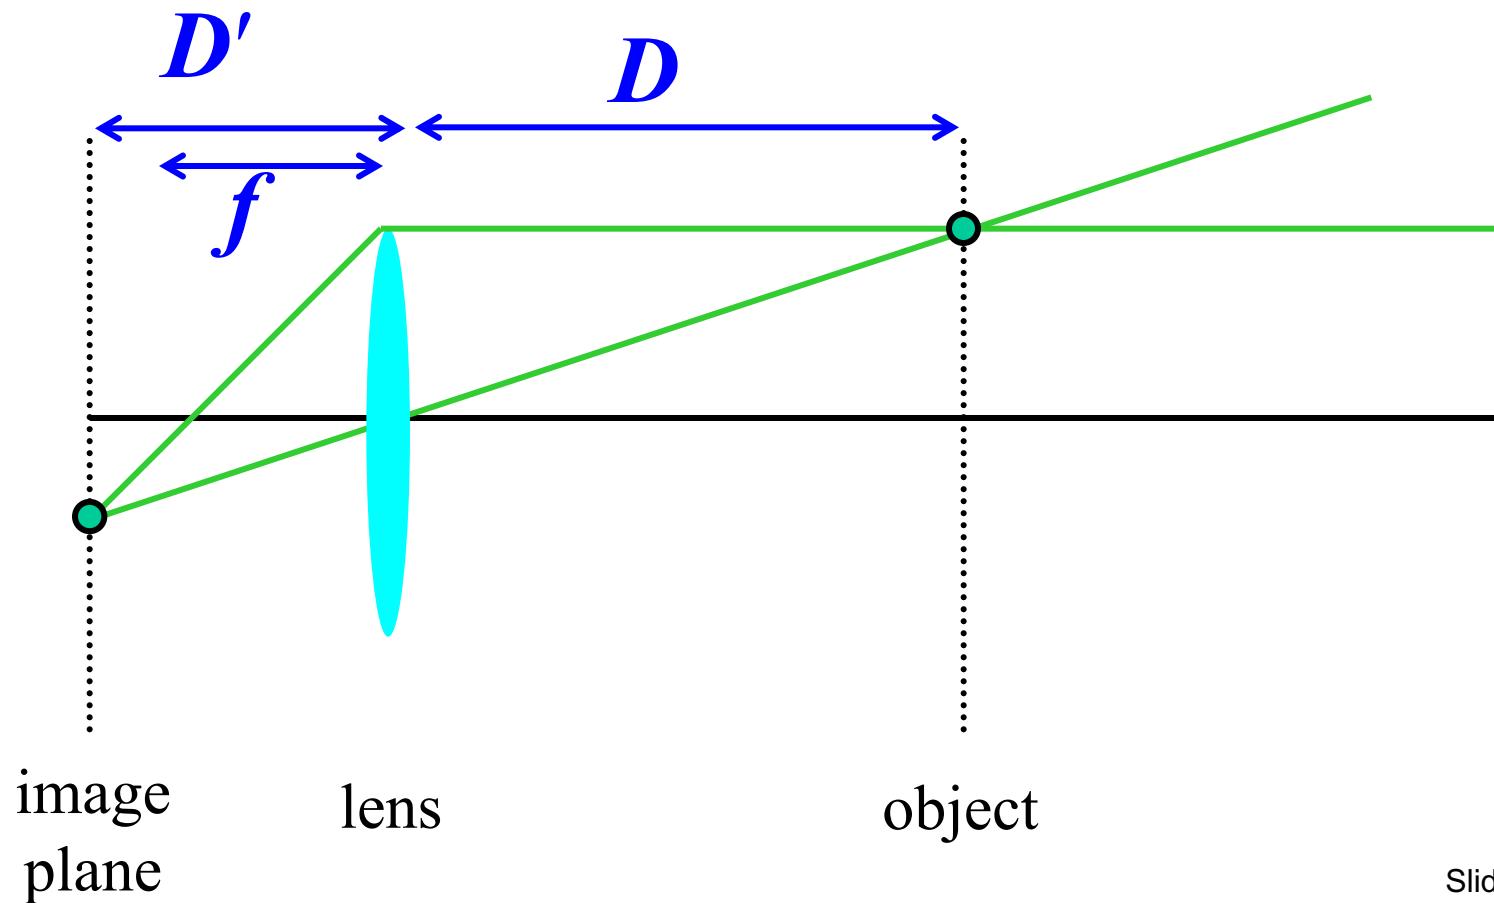
Thin lens formula

- What is the relation between the focal length (f), the distance of the object from the optical center (D), and the distance at which the object will be in focus (D')?



Thin lens formula

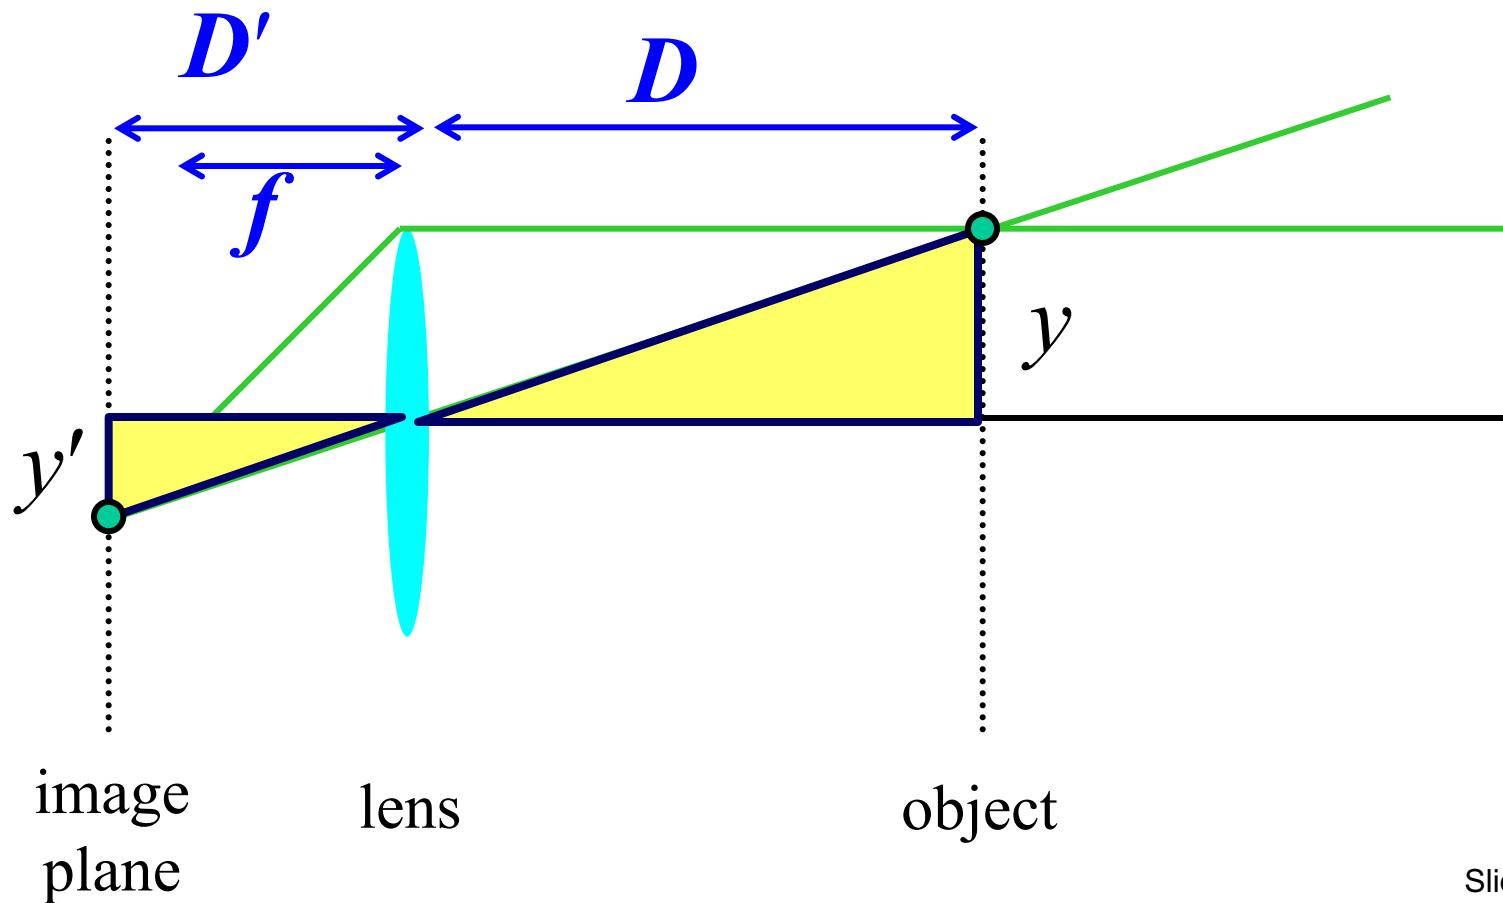
Similar triangles everywhere!



Thin lens formula

Similar triangles everywhere!

$$y'/y = D'/D$$

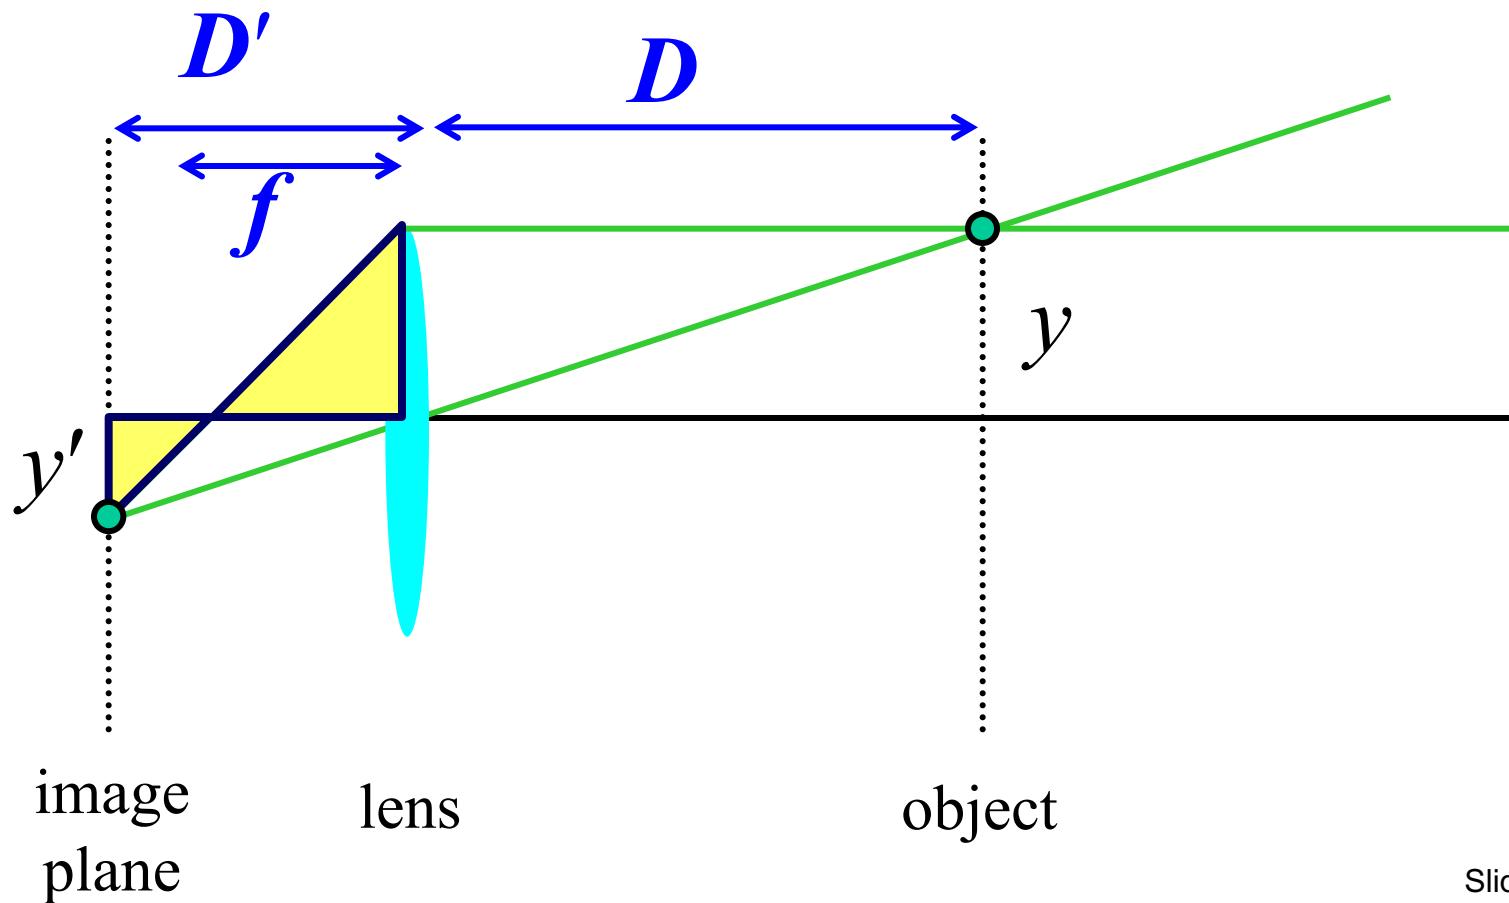


Thin lens formula

Similar triangles everywhere!

$$y'/y = D'/D$$

$$y'/y = (D' - f)/f$$

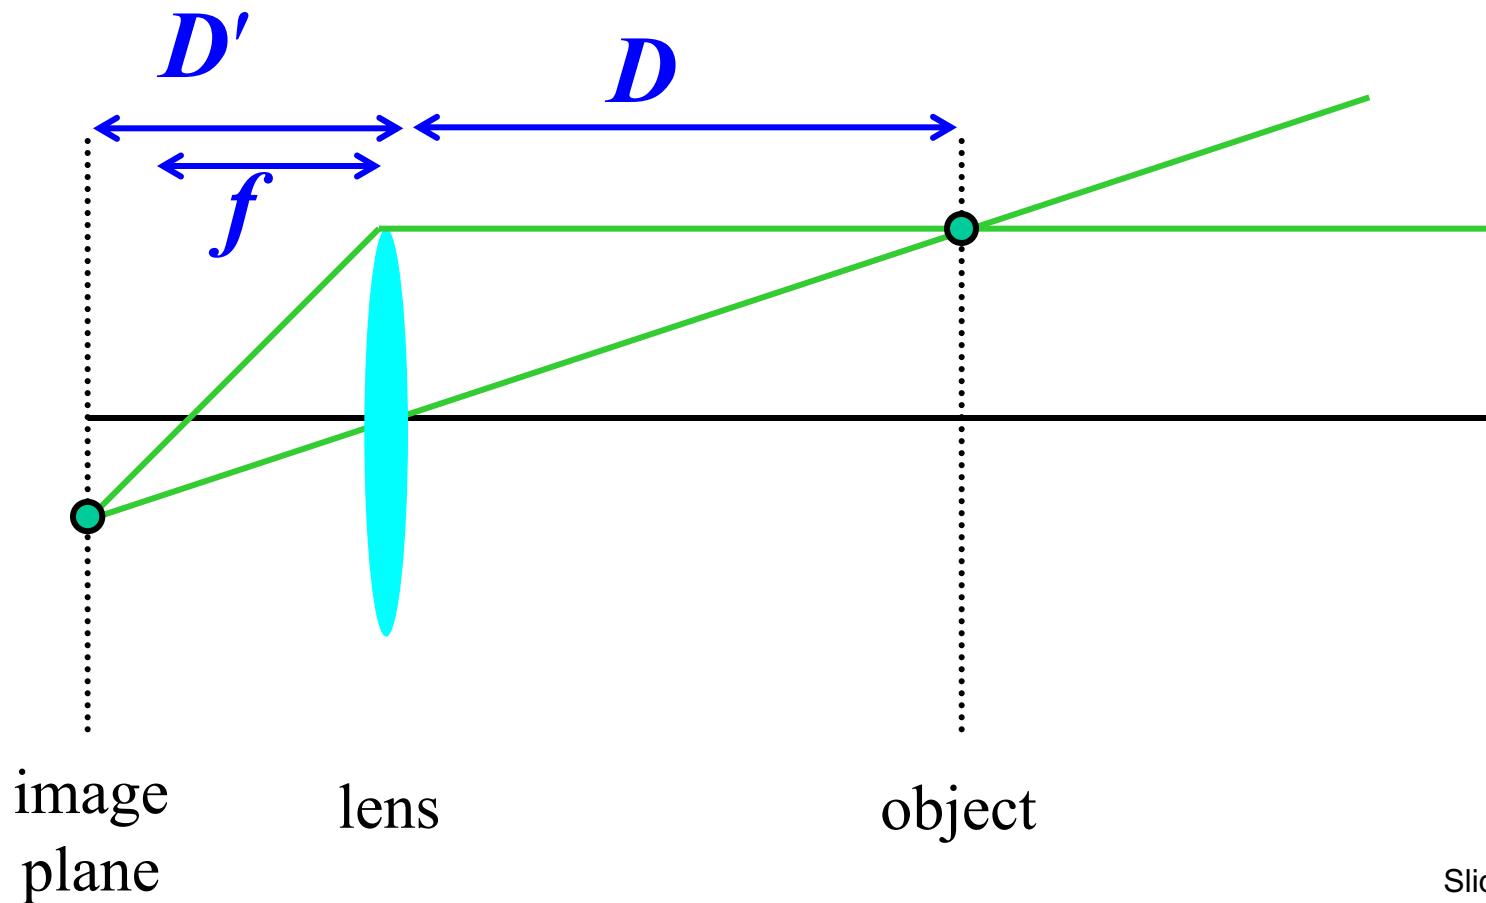


Thin lens formula

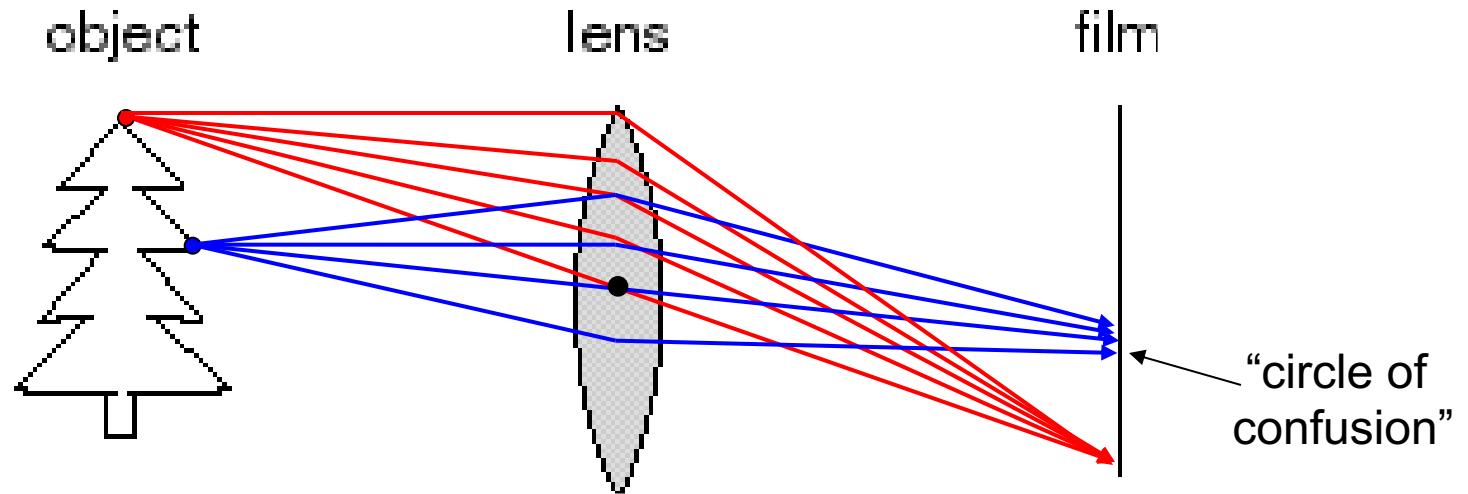
$$\frac{1}{D'} + \frac{1}{D} = \frac{1}{f}$$

Any point satisfying the thin lens equation is in focus.

What happens when D is very large?



Depth of Field



For a fixed focal length, there is a specific distance at which objects are “in focus”

- Other points project to a “circle of confusion” in the image

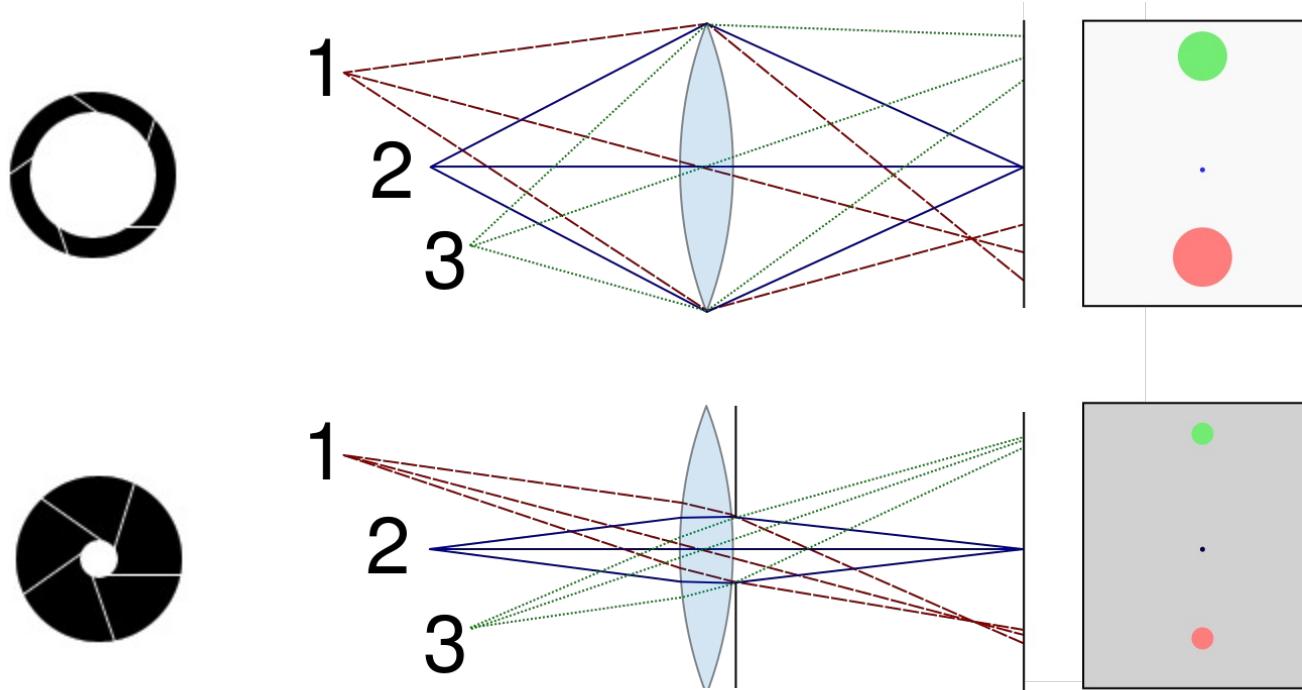
Depth of Field



DEPTH OF FIELD
DEPTH OF FIELD

<http://www.cambridgeincolour.com/tutorials/depth-of-field.htm>

Controlling depth of field



Changing the aperture size affects depth of field

- A smaller *aperture* increases the range in which the object is approximately in focus
- But small aperture reduces amount of light – need to increase *exposure*

Varying the aperture

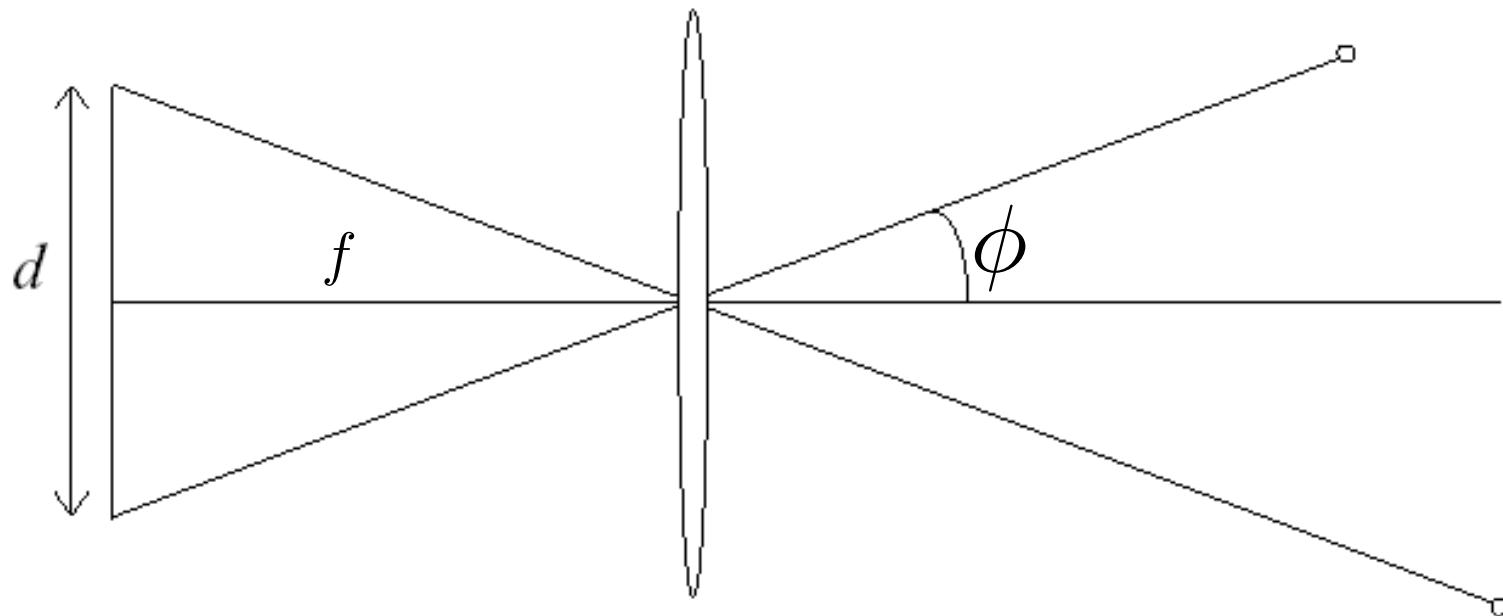


Large aperture = small DOF



Small aperture = large DOF

Field of View

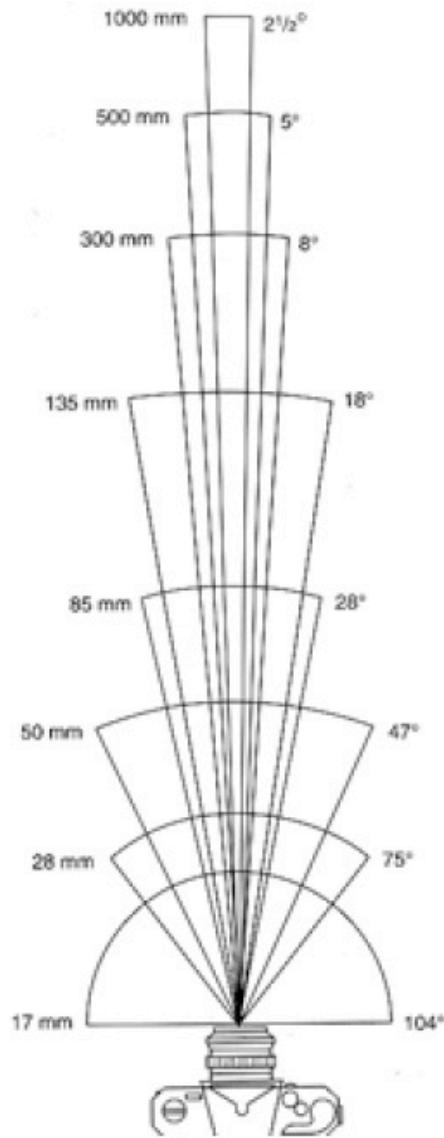


FOV depends on focal length and size of the camera retina

$$\phi = \tan^{-1}\left(\frac{d/2}{f}\right)$$

Larger focal length = smaller FOV

Field of View



17mm



28mm

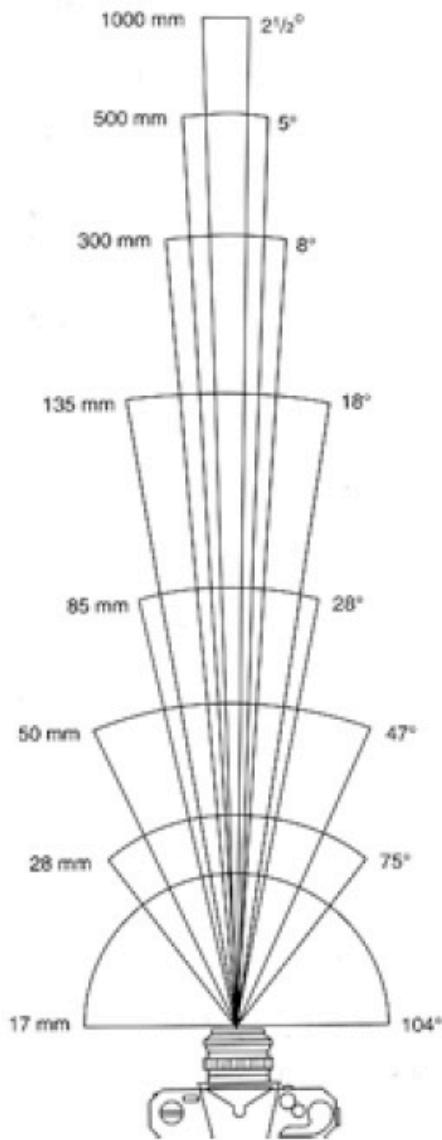


50mm

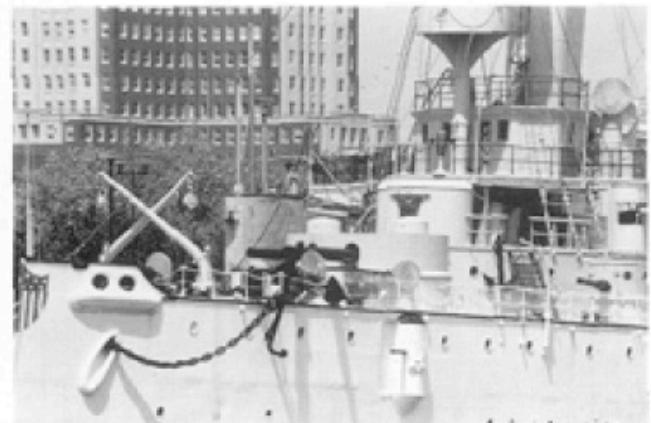


85mm

Field of View



135mm



300mm

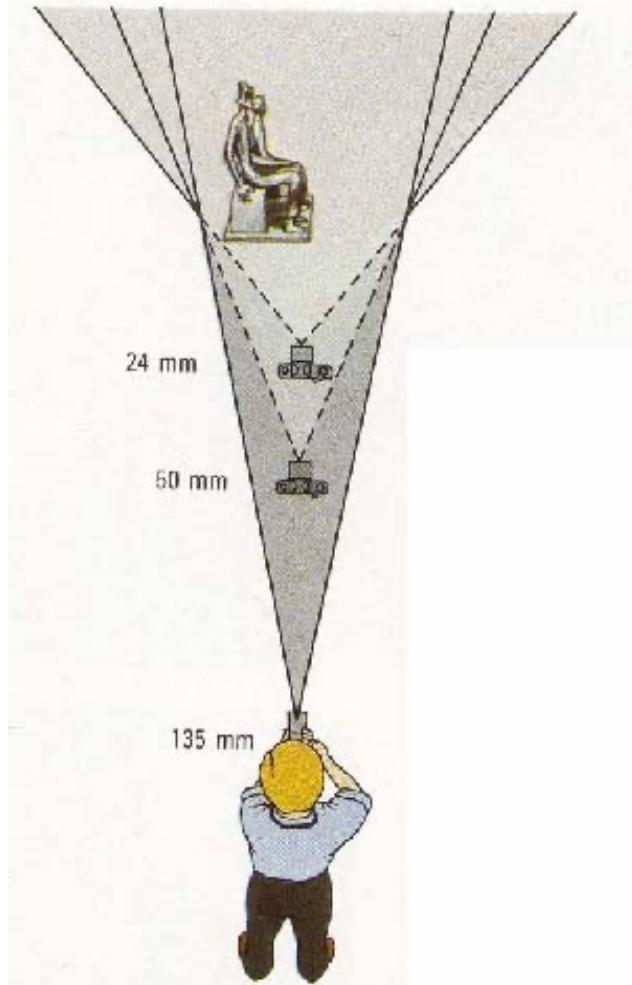


135mm



300mm

Field of View / Focal Length



Large FOV, small f
Camera close to car



Small FOV, large f
Camera far from the car

Same effect for faces



wide-angle

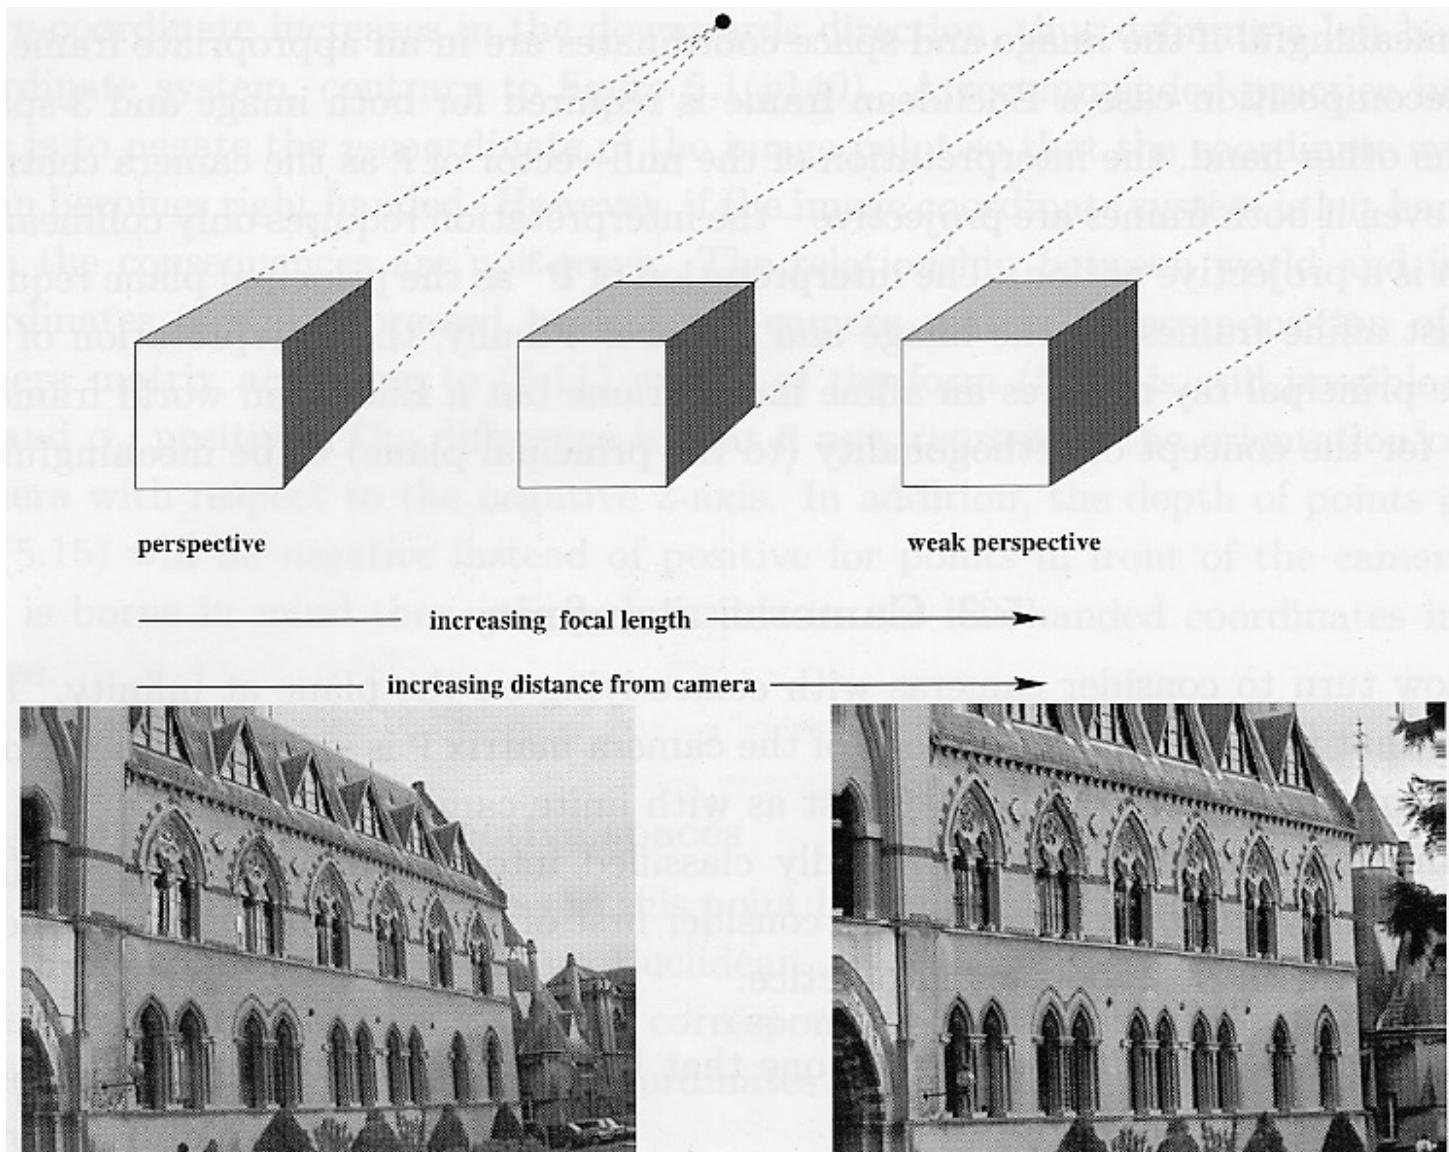


standard



telephoto

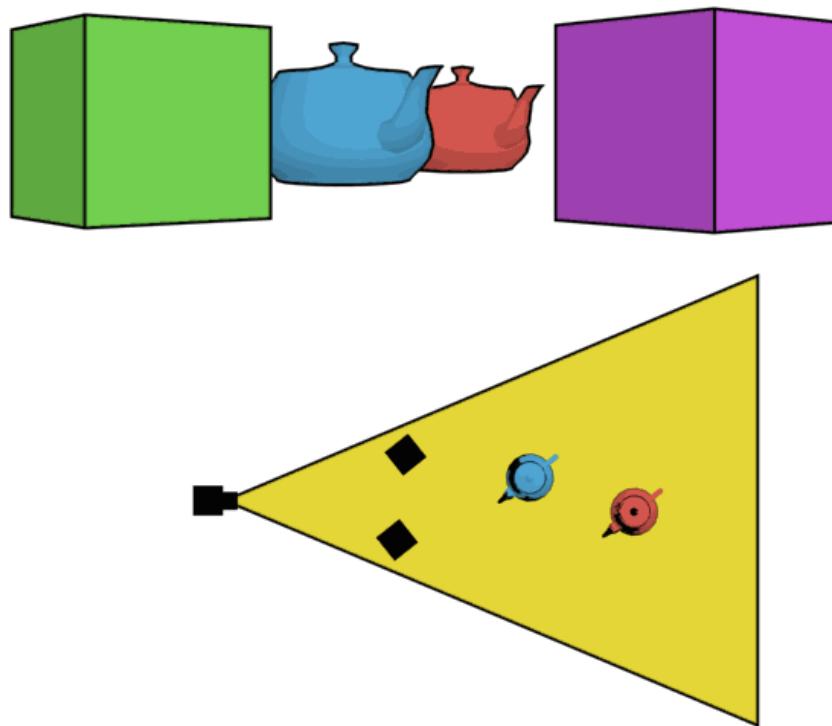
Approximating an orthographic camera



Source: Hartley & Zisserman

The dolly zoom

- Continuously adjusting the focal length while the camera moves away from (or towards) the subject



http://en.wikipedia.org/wiki/Dolly_zoom

The dolly zoom

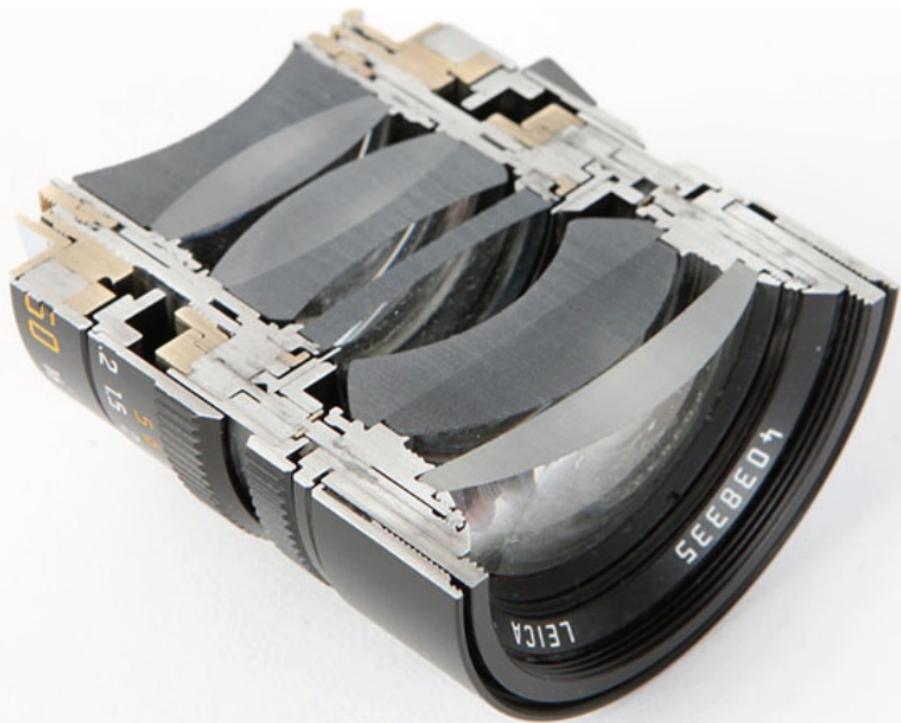
- Continuously adjusting the focal length while the camera moves away from (or towards) the subject
- “The Vertigo shot”



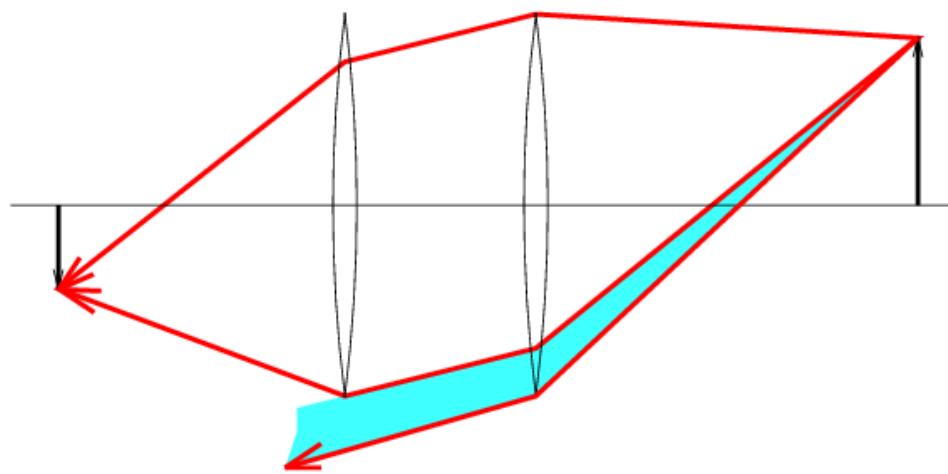
[Example of dolly zoom from *Goodfellas* \(YouTube\)](#)

[Example of dolly zoom from *La Haine* \(YouTube\)](#)

Real lenses

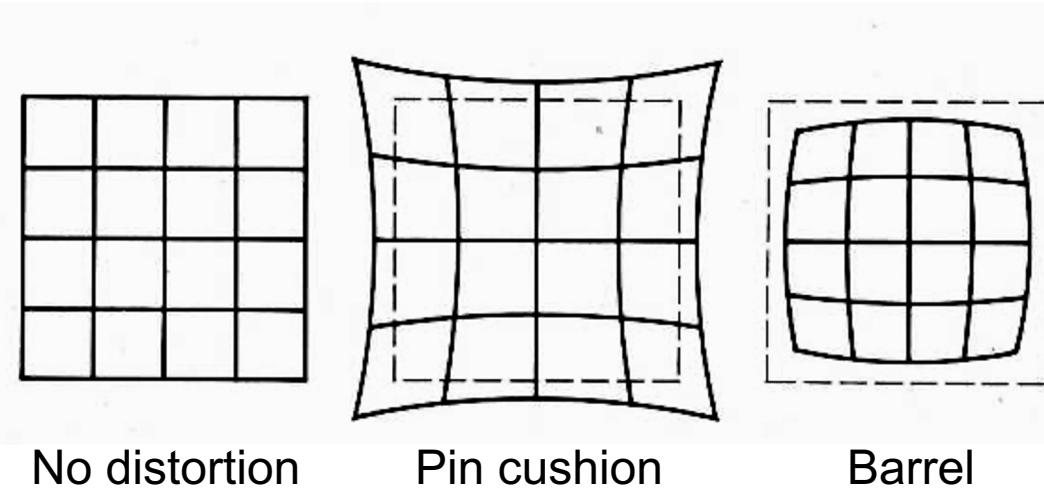


Lens flaws: Vignetting



Radial Distortion

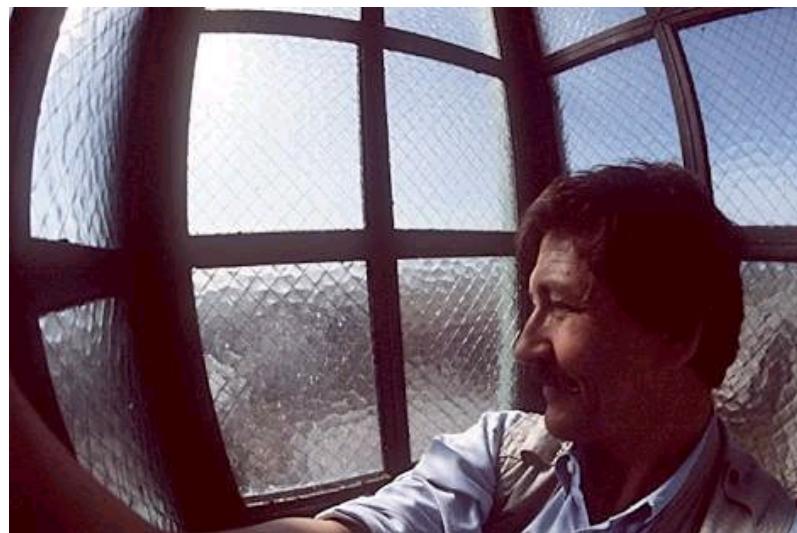
- Caused by imperfect lenses
- Deviations are most noticeable near the edge of the lens



No distortion

Pin cushion

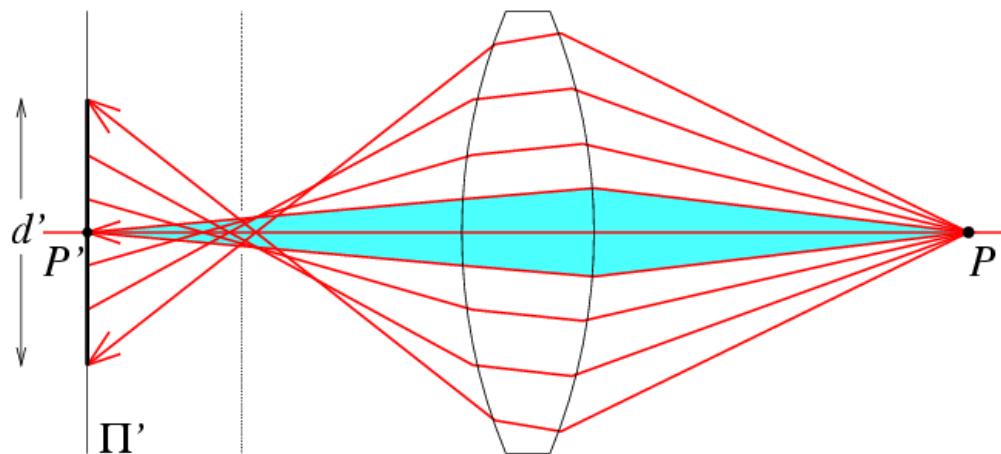
Barrel



Lens flaws: Spherical aberration

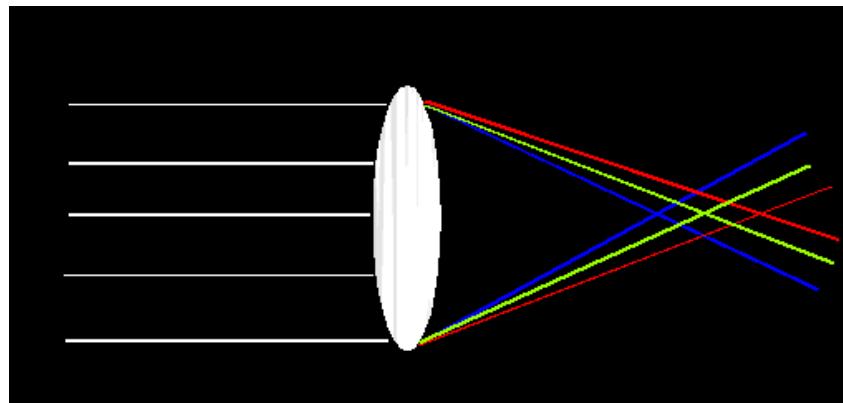
Spherical lenses don't focus light perfectly

Rays farther from the optical axis focus closer



Lens Flaws: Chromatic Aberration

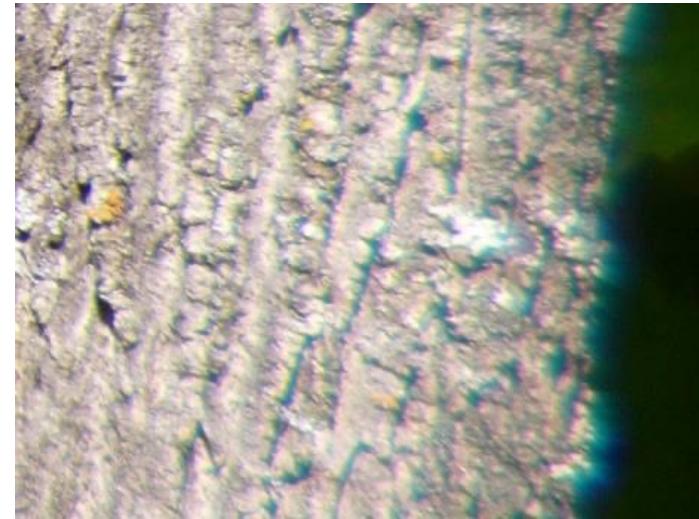
Lens has different refractive indices for different wavelengths: causes color fringing



Near Lens Center

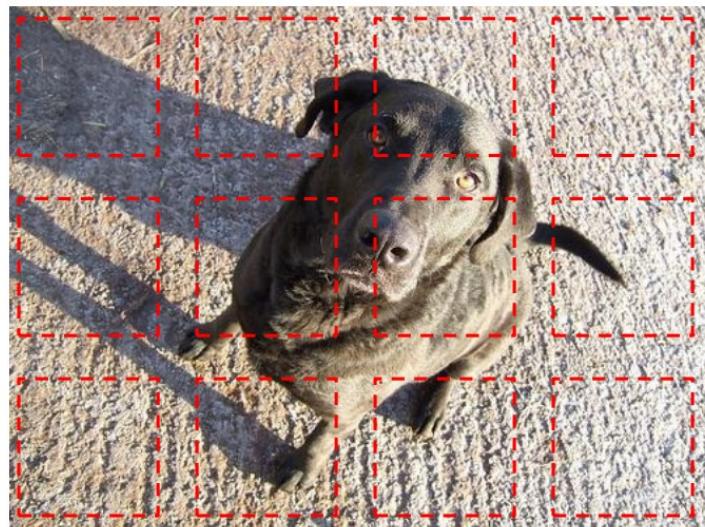


Near Lens Outer Edge



Lens Flaws: Chromatic Aberration

Researchers tried teaching a network about objects by forcing it to assemble jigsaws.



Initial layout, with sampled patches in red

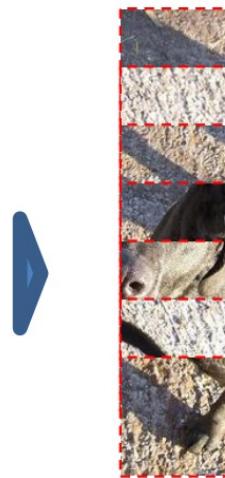


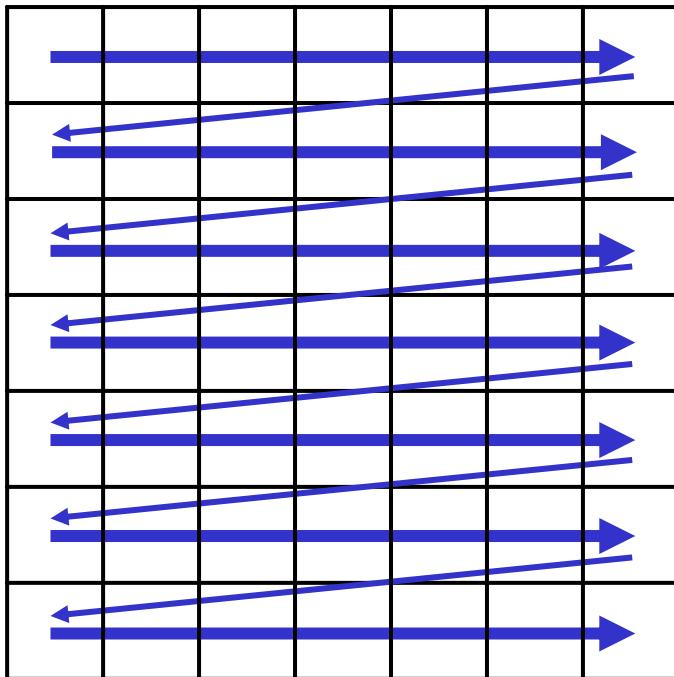
Image layout
is discarded



We can recover image layout automatically

Rolling Shutter

Rolling Shutter: pixels read in sequence
Can get global reading, but \$\$\$

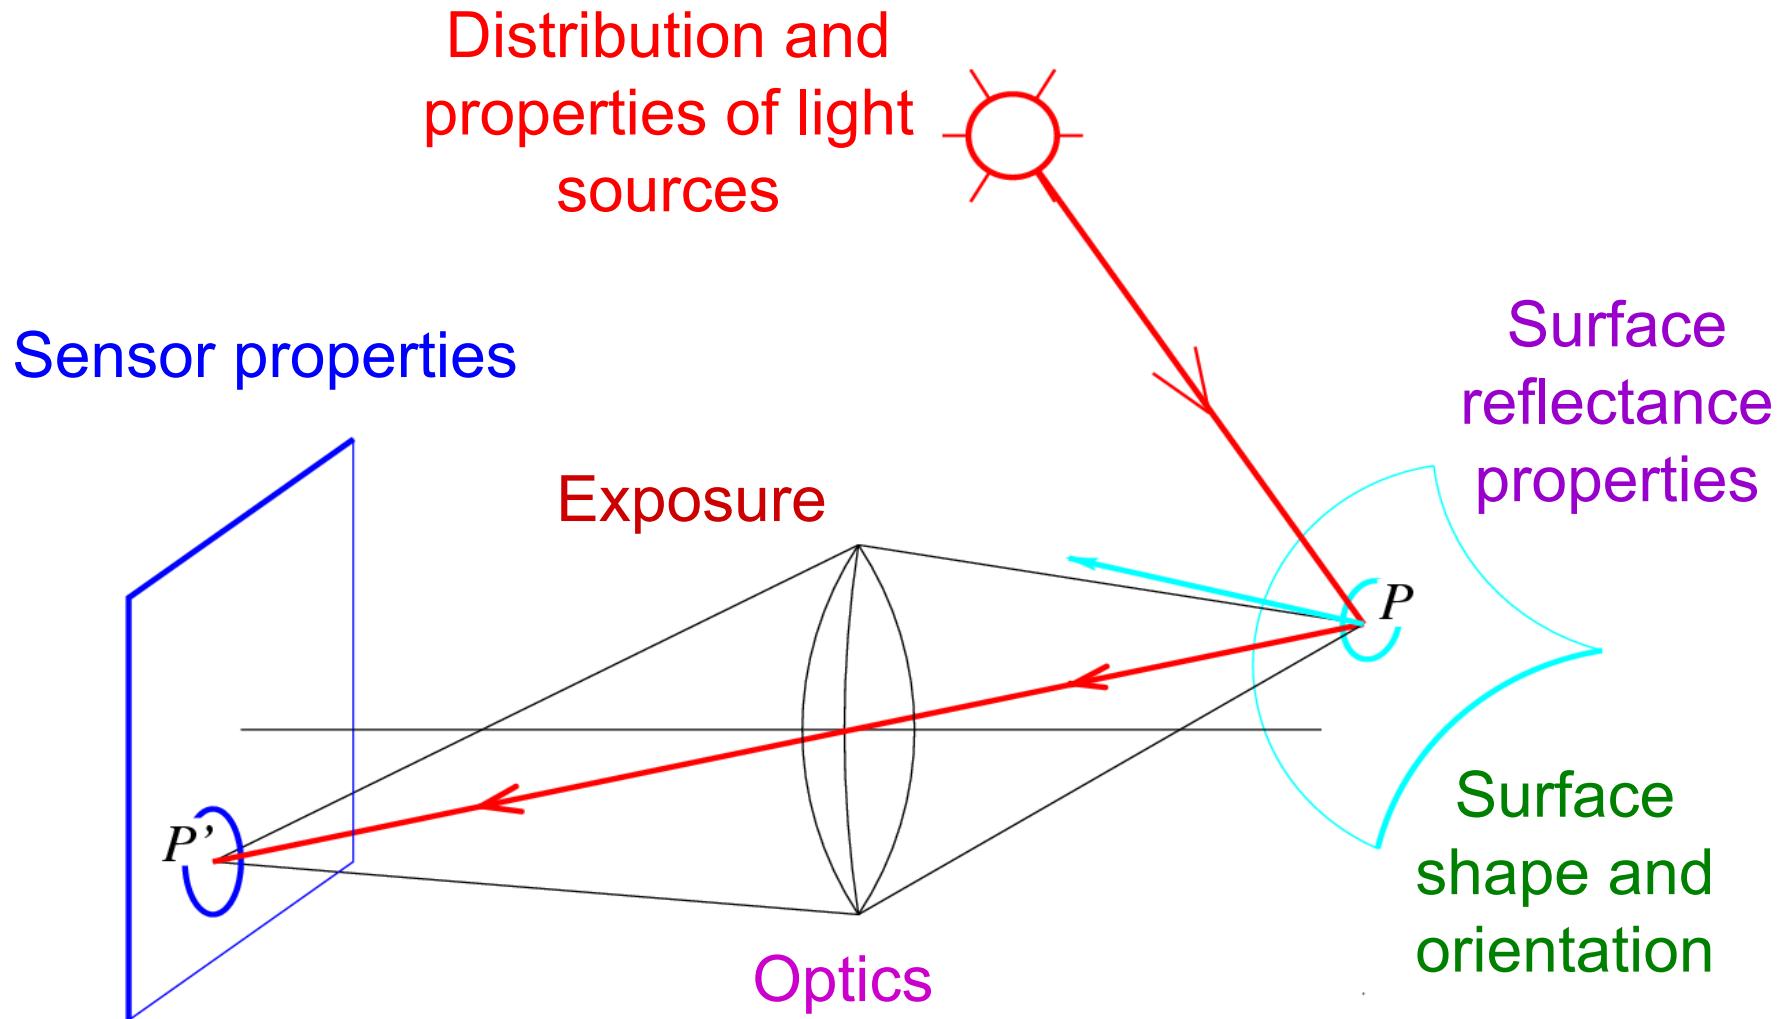


Overview

- Cameras with lenses
 - Depth of field
 - Field of view
 - Lens aberrations
- Brightness of a pixel
 - Small taste of radiometry
 - In-camera transformation of light
 - Reflectance properties of surfaces
 - Lambertian reflection model
 - Shape from shading

Image formation

What determines the brightness of an image pixel?



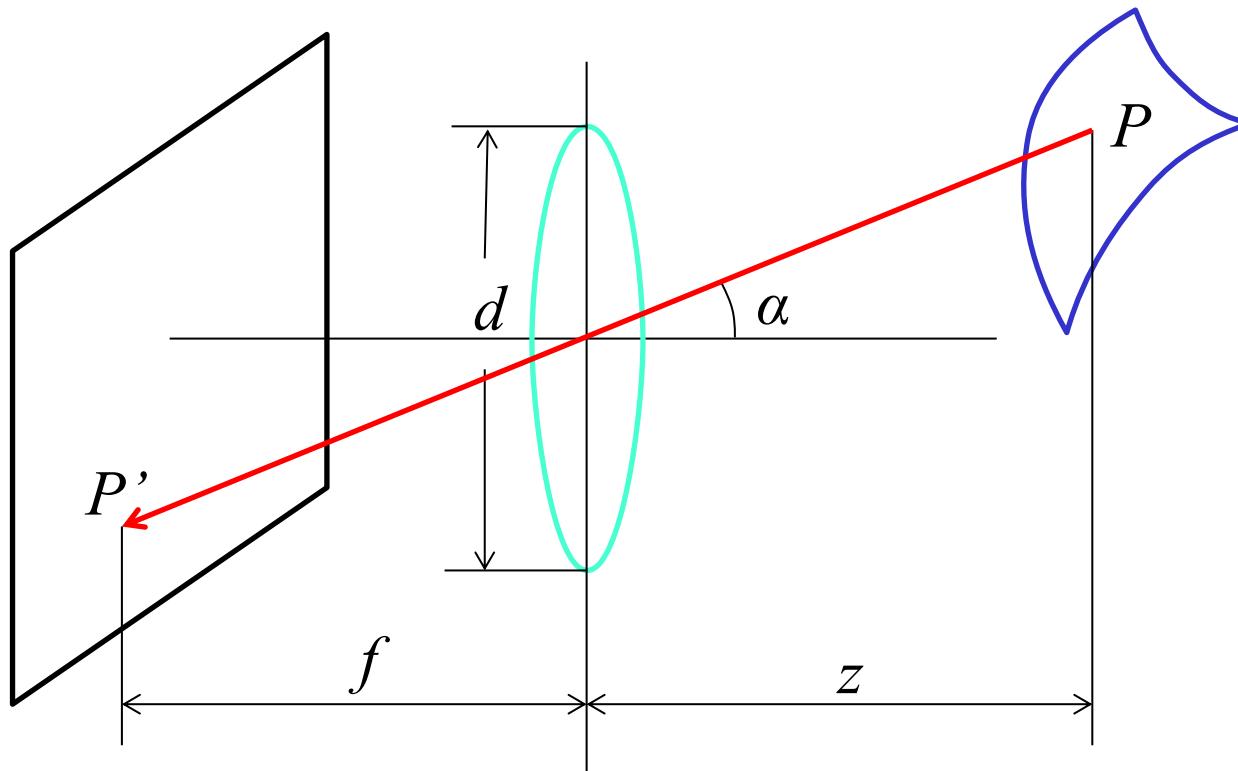
Fundamental radiometric relation

L: **Radiance** emitted from P toward P'

- Energy carried by a ray (Watts per sq. meter per steradian)

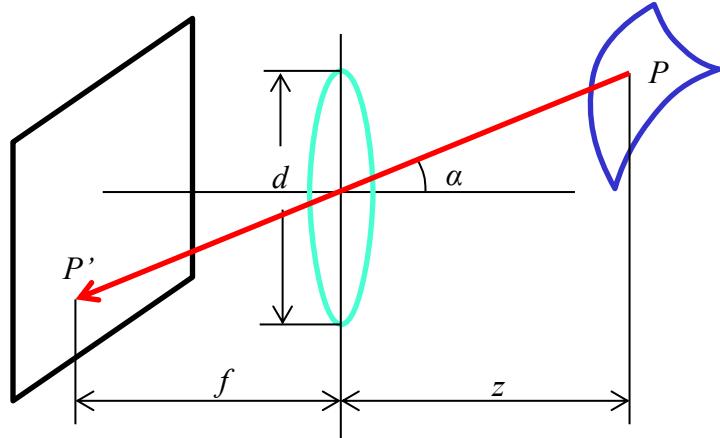
E: **Irradiance** falling on P' from the lens

- Energy arriving at a surface (Watts per sq. meter)



What is the relationship between E and L ?

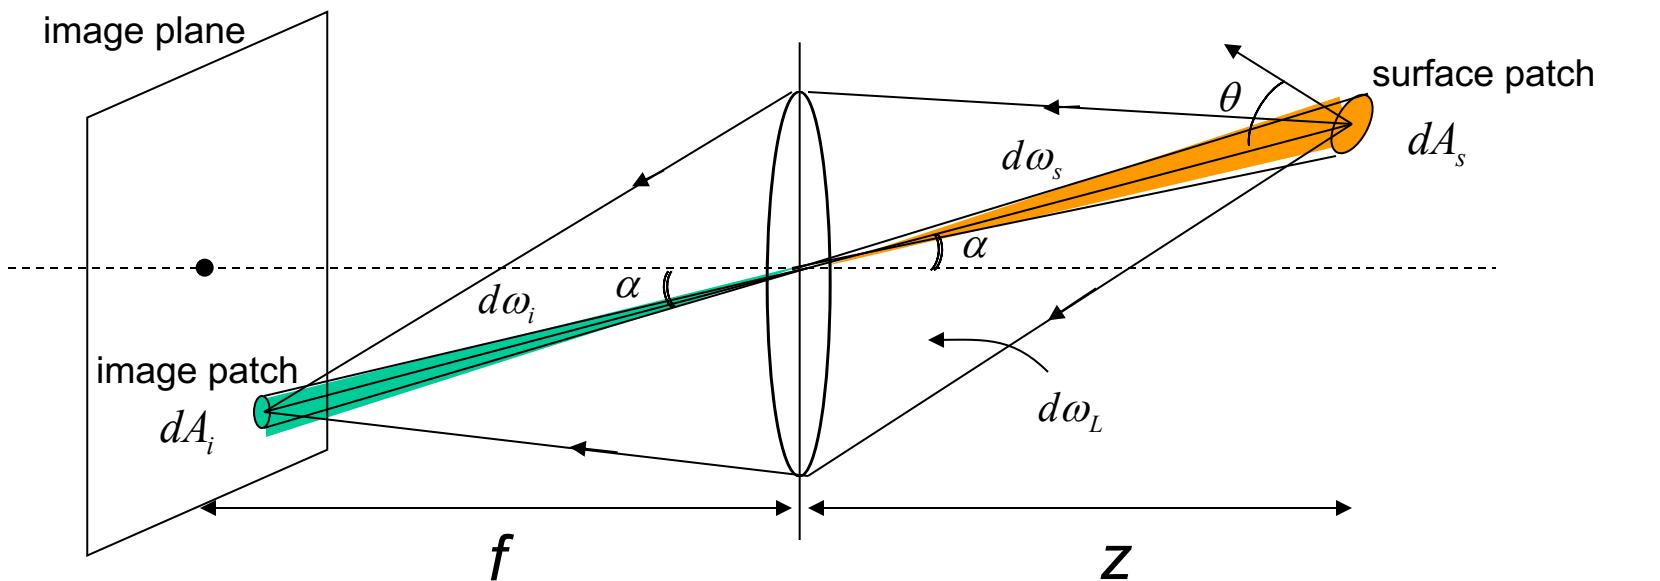
Fundamental radiometric relation



$$E = \left[\frac{\pi}{4} \left(\frac{d}{f} \right)^2 \cos^4 \alpha \right] L$$

- Image irradiance is linearly related to scene radiance
- Irradiance is proportional to the area of the lens and inversely proportional to the squared distance between the lens and the image plane
- The irradiance falls off as the angle between the viewing ray and the optical axis increases

Relation between Image Irradiance E and Scene Radiance L



- Solid angles of the double cone (orange and green):

$$d\omega_i = d\omega_s \quad \frac{dA_i \cos \alpha}{(f / \cos \alpha)^2} = \frac{dA_s \cos \theta}{(z / \cos \alpha)^2}$$

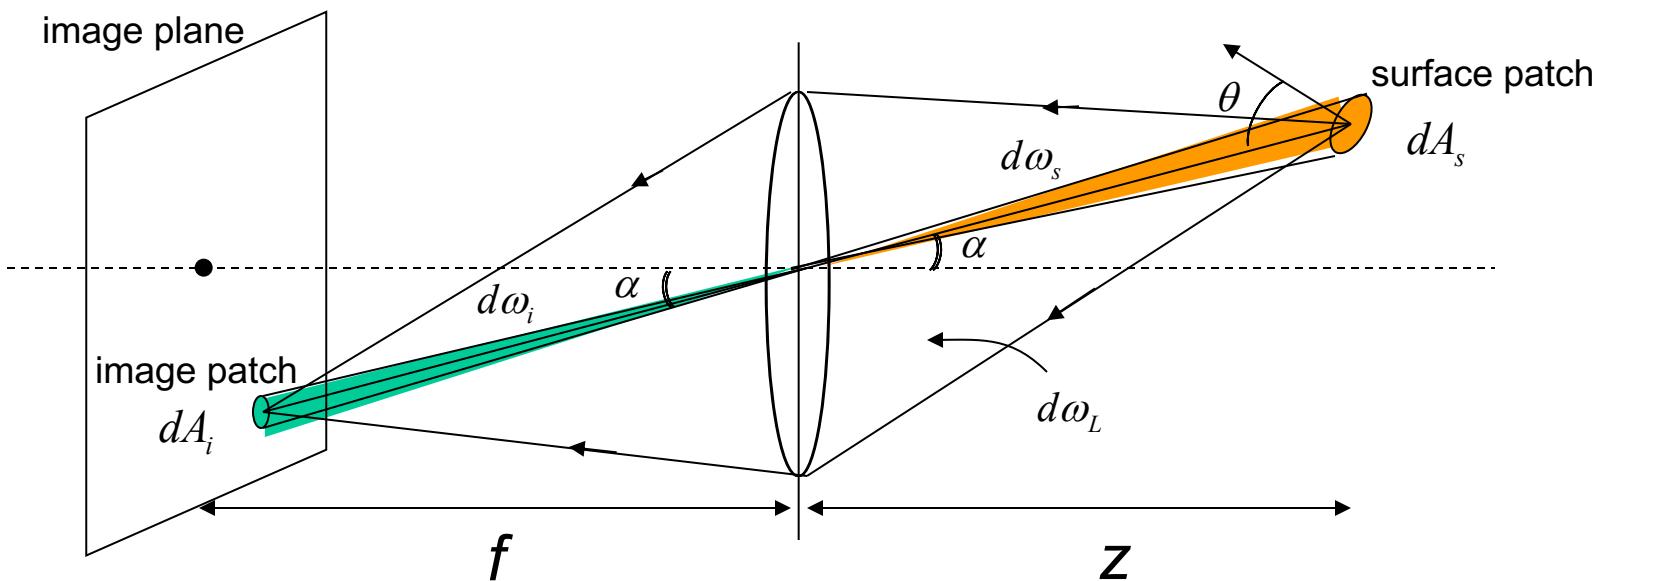
$$\frac{dA_s}{dA_i} = \frac{\cos \alpha}{\cos \theta} \left(\frac{z}{f} \right)^2$$

- Solid angle subtended by lens:

$$d\omega_L = \frac{\pi d^2}{4} \frac{\cos \alpha}{(z / \cos \alpha)^2} \rightarrow (2)$$

(1)

Relation between Image Irradiance E and Scene Radiance L



- Flux received by lens from dA_s = Flux projected onto image dA_i

$$L (dA_s \cos \theta) d\omega_L = E dA_i \rightarrow (3)$$

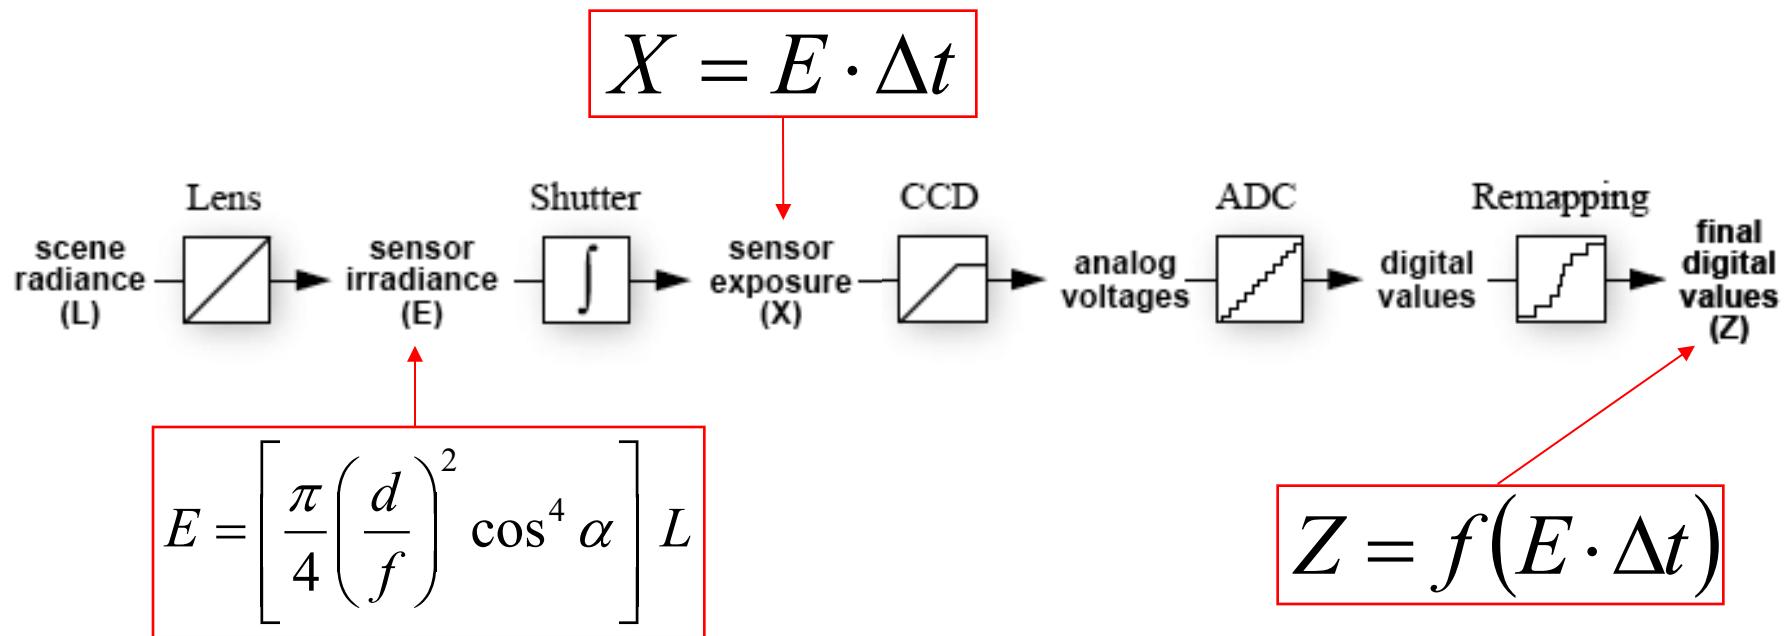
- From (1), (2), and (3):

$$E = L \frac{\pi}{4} \left(\frac{d}{f}\right)^2 \cos \alpha^4$$

- Image irradiance is proportional to Scene Radiance!

- Small field of view → Effects of 4th power of cosine are small.

From light rays to pixel values

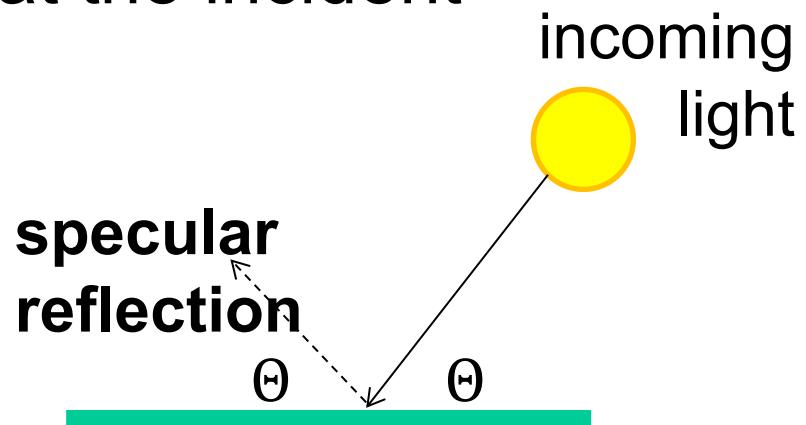


- Camera response function: the mapping f from irradiance to pixel values
 - Useful if we want to estimate material properties
 - Enables us to create *high dynamic range (HDR) images*
 - Classic reference: P. E. Debevec and J. Malik, [Recovering High Dynamic Range Radiance Maps from Photographs](#), SIGGRAPH 97

Basic models of reflection

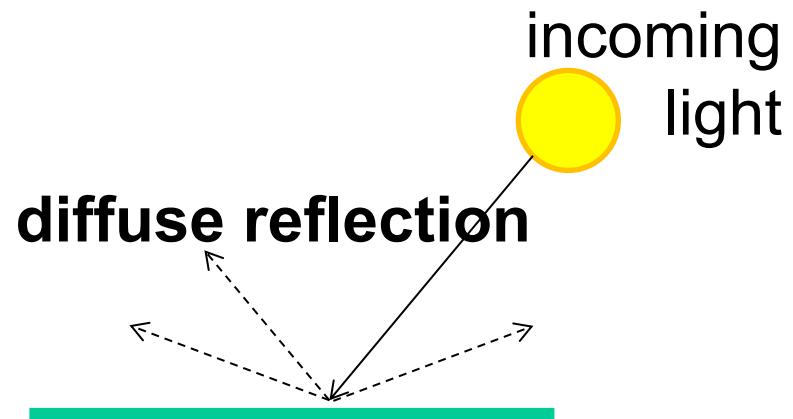
Specular: light bounces off at the incident angle

- E.g., mirror



Diffuse: light scatters in all directions

- E.g., brick, cloth, rough wood

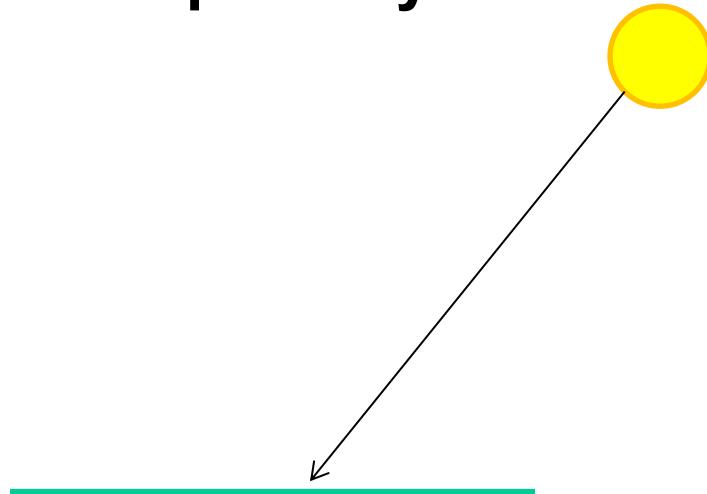


Other possible effects



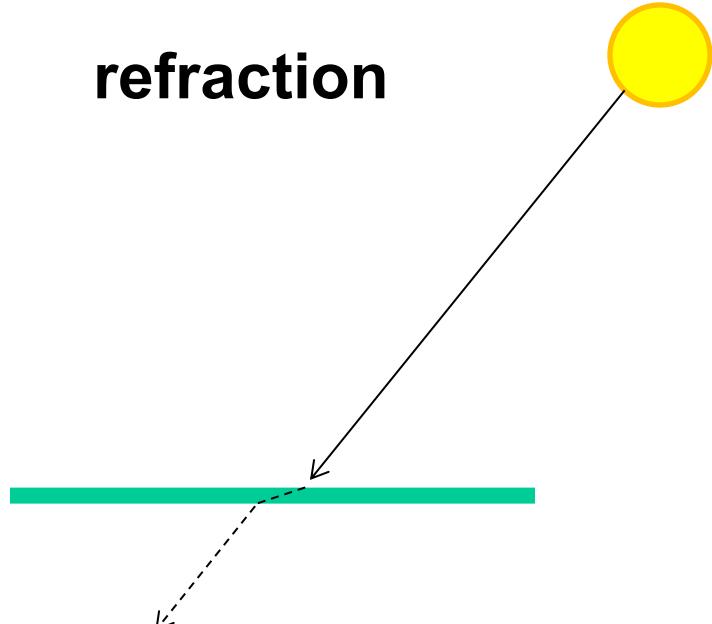
transparency

light source



refraction

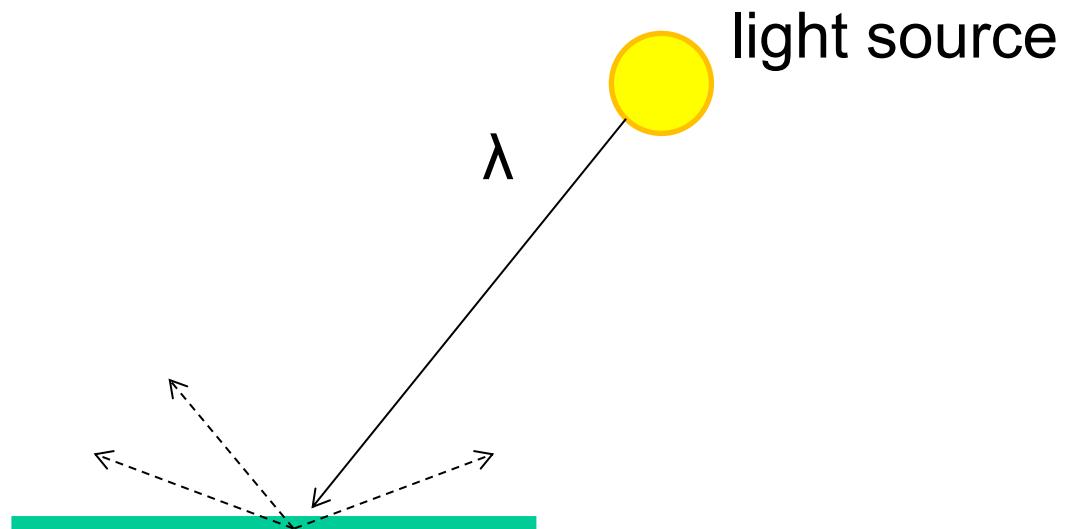
light source



Other possible effects



**subsurface
scattering**

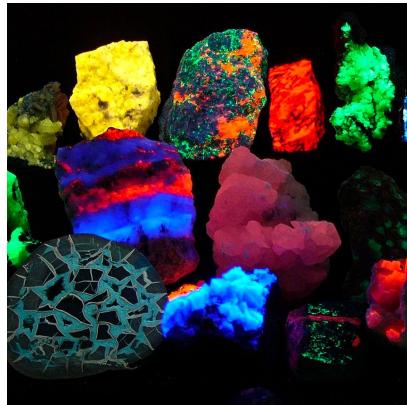


Slide from D. Hoiem

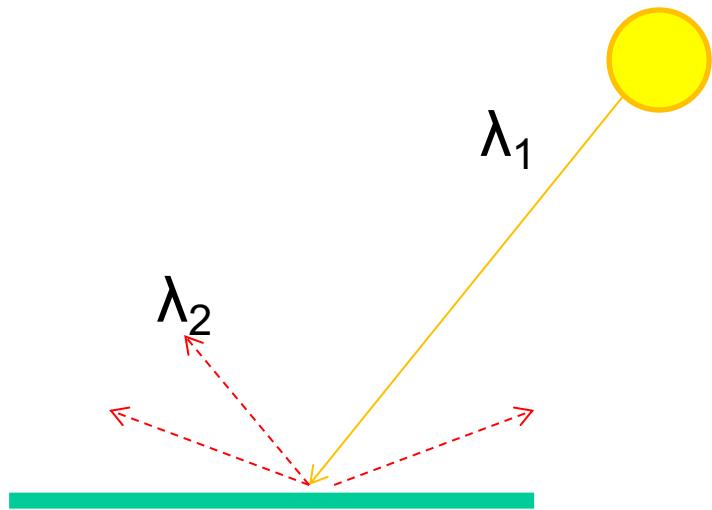
https://en.wikipedia.org/wiki/Subsurface_scattering#/media/File:Skin_Subsurface_Scattering.jpg

Other possible effects

fluorescence



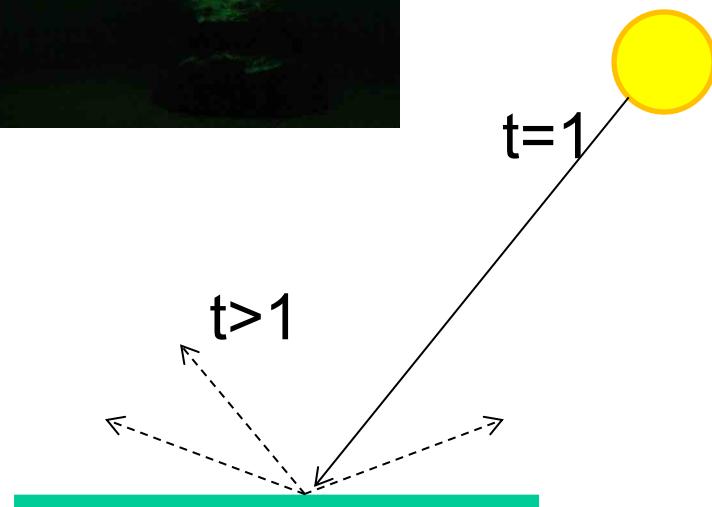
light source



phosphorescence



light source



Overview

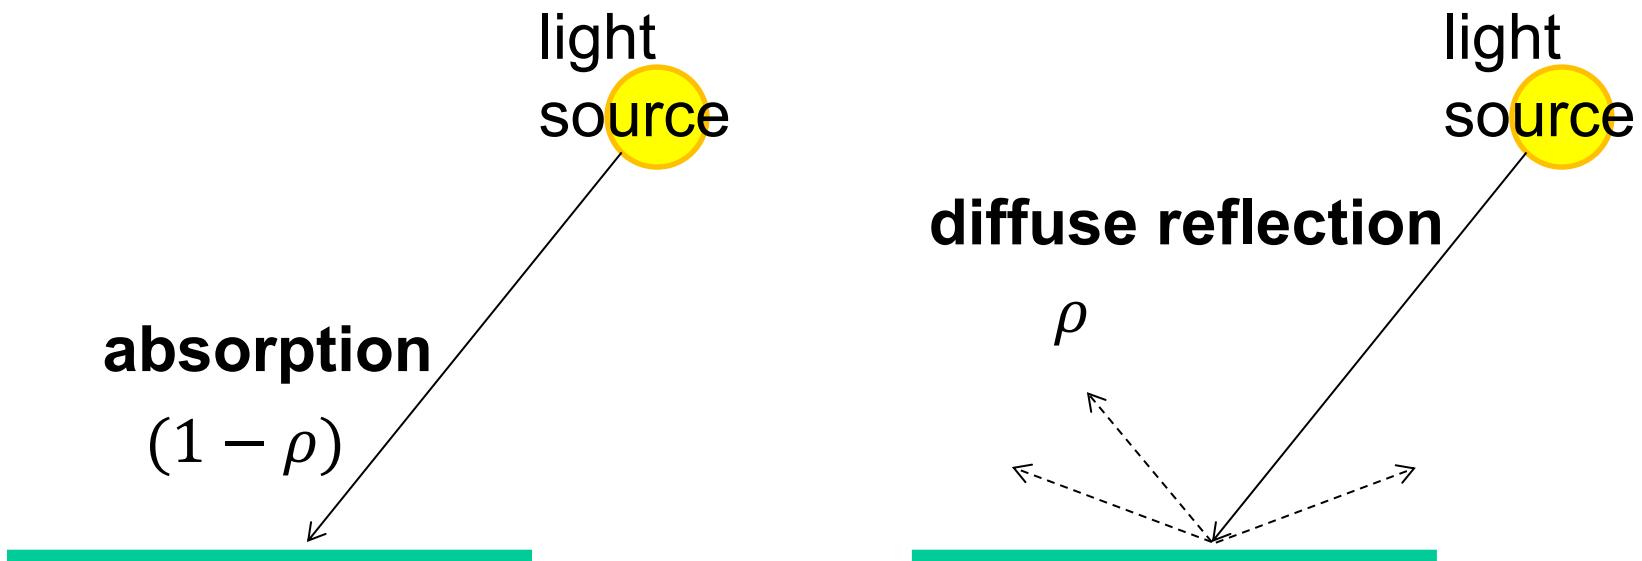
- Cameras with lenses
 - Depth of field
 - Field of view
 - Lens aberrations
- Brightness of a pixel
 - Small taste of radiometry
 - In-camera transformation of light
 - Reflectance properties of surfaces
 - Lambertian reflection model
 - Shape from shading

Lambertian reflectance model

Some light is absorbed (function of albedo ρ)

Remaining light is scattered (diffuse reflection)

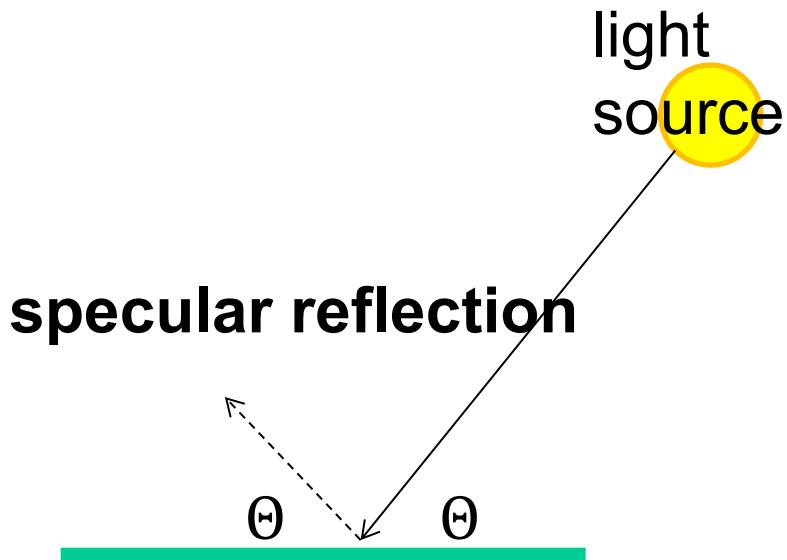
Examples: soft cloth, concrete, matte paints



Specular Reflection

Reflected direction depends on light orientation and surface normal

- E.g., mirrors are fully specular



Flickr, by suzysputnik



Flickr, by piratejohnny

Most surfaces have both specular and diffuse components

Specularity = spot where specular reflection dominates (typically reflects light source)



Photo: northcountryhardwoodfloors.com

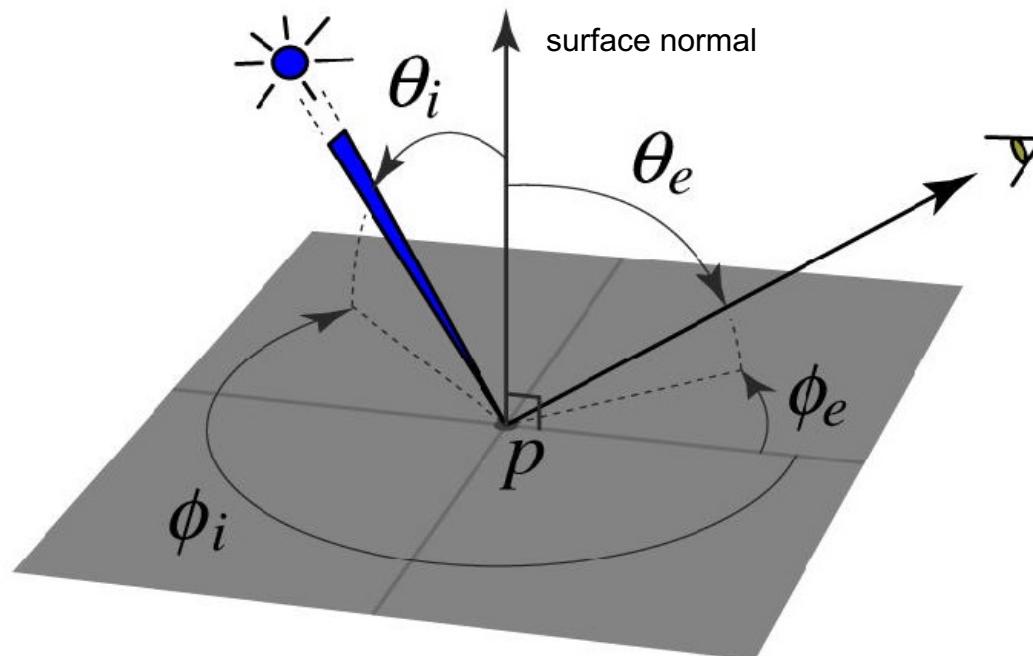


Typically, specular component is small

Slide from D. Hoiem

BRDF: Bidirectional Reflectance Distribution Function

- Model of local reflection that tells how bright a surface appears when viewed from one direction when light falls on it from another
- Definition: ratio of the radiance in the emitted direction to irradiance in the incident direction

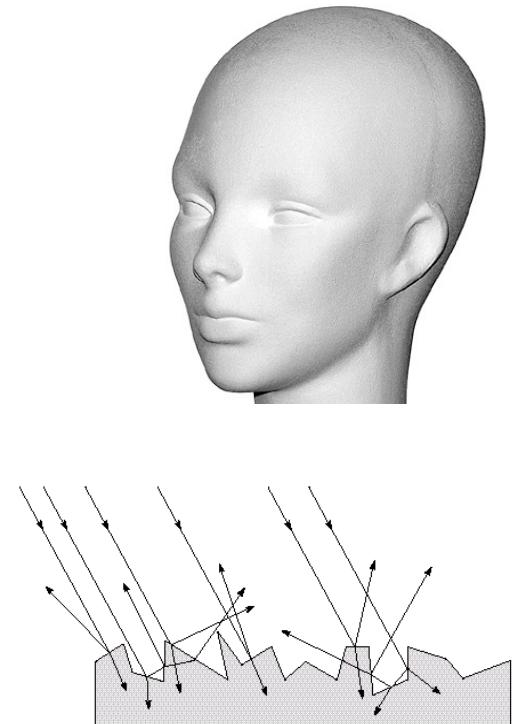
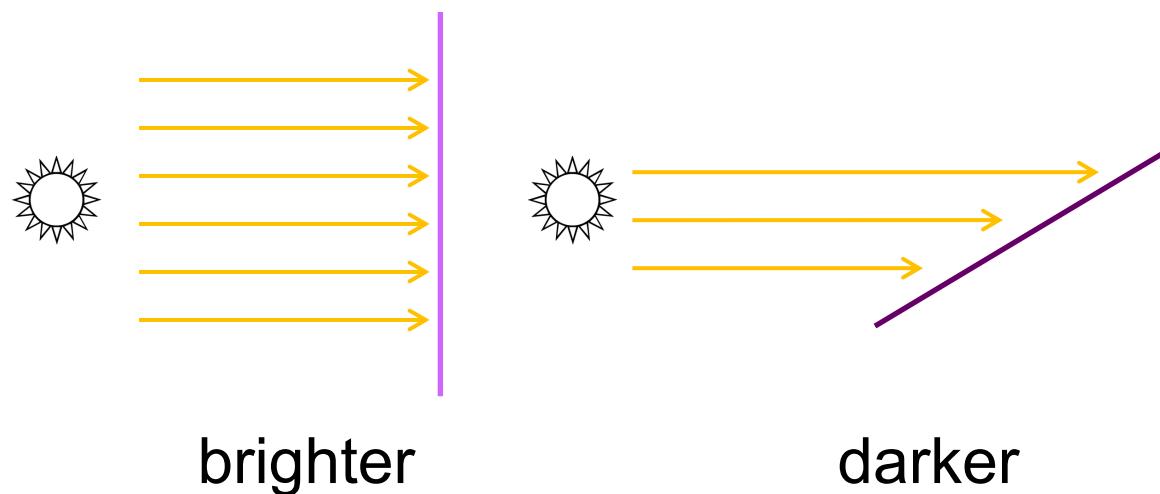


BRDFs can be incredibly complicated...

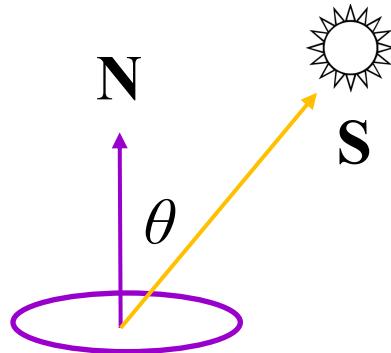


Diffuse reflection

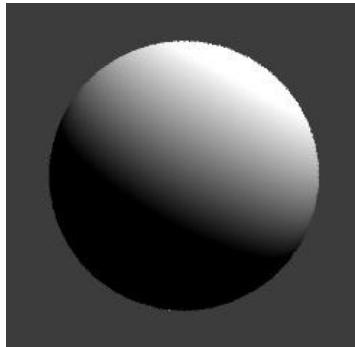
- Light is reflected equally in all directions
 - Dull, matte surfaces like chalk or latex paint
 - Microfacets scatter incoming light randomly
- Brightness of the surface depends on the incidence of illumination



Diffuse reflection: Lambert's law



$$\begin{aligned}B &= \rho \mathbf{N} \cdot \mathbf{S} \\&= \rho \|\mathbf{S}\| \cos \theta\end{aligned}$$

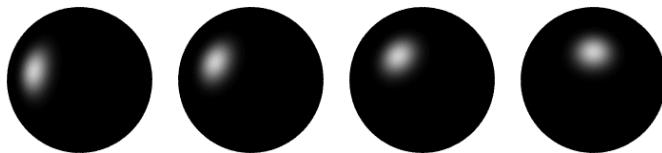


- B : radiosity (total power leaving the surface per unit area)
- ρ : albedo (fraction of incident irradiance reflected by the surface)
- N : unit normal
- S : source vector (magnitude proportional to intensity of the source)

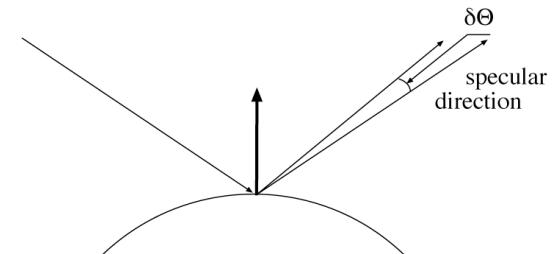
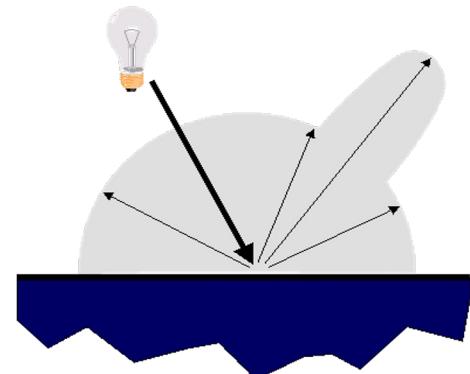
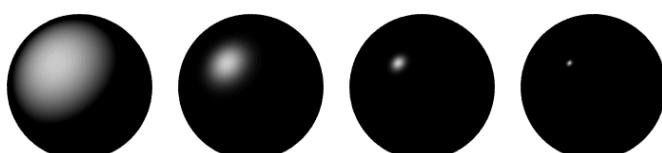
Specular reflection

- Radiation arriving along a source direction leaves along the **specular direction** (source direction reflected about normal)
- On real surfaces, energy usually goes into a lobe of directions
- **Phong model:** reflected energy falls off with $\cos^n(\delta\theta)$

Moving the light source



Changing the exponent



Specular reflection



[Picture source](#)

Lambertian + Specular Model

- $I(x) = \text{Ambient Term} + \text{Diffuse Term} + \text{Specular Term}$

Photometric stereo (shape from shading)

- Can we reconstruct the shape of an object based on shading cues?



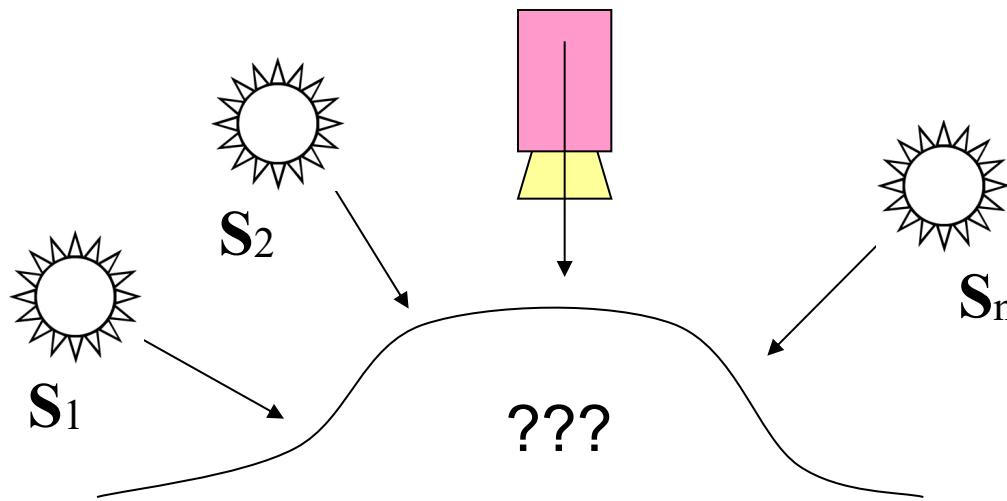
Luca della Robbia,
Cantoria, 1438

Photometric stereo

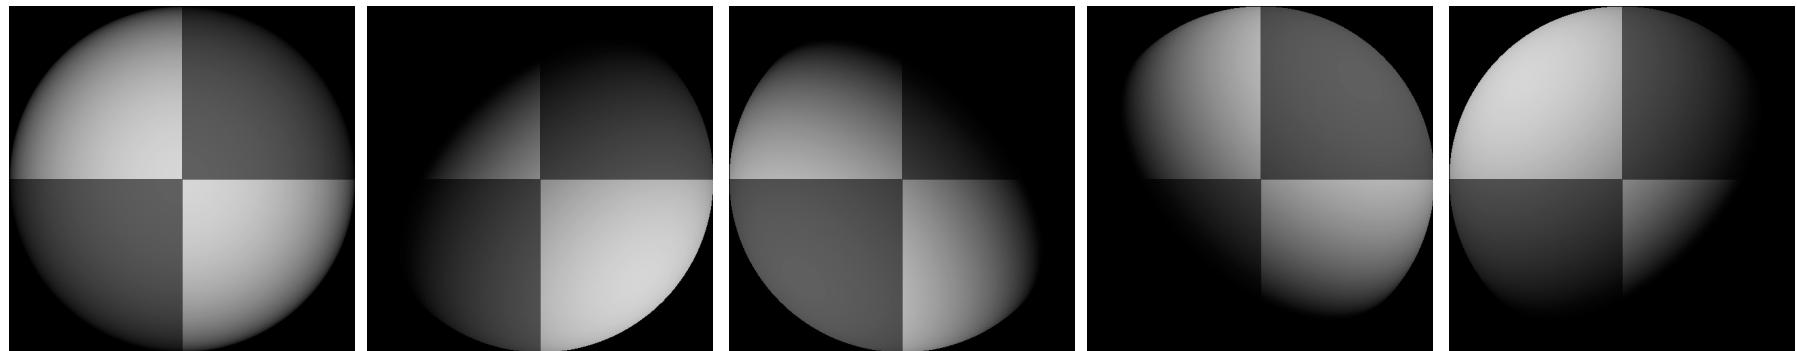
Assume:

- A Lambertian object
- A *local shading model* (each point on a surface receives light only from sources visible at that point)
- A set of *known* light source directions
- A set of pictures of an object, obtained in exactly the same camera/object configuration but using different sources
- Orthographic projection

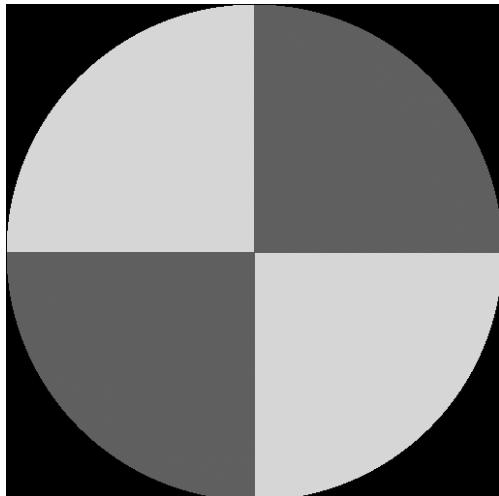
Goal: reconstruct object shape and albedo



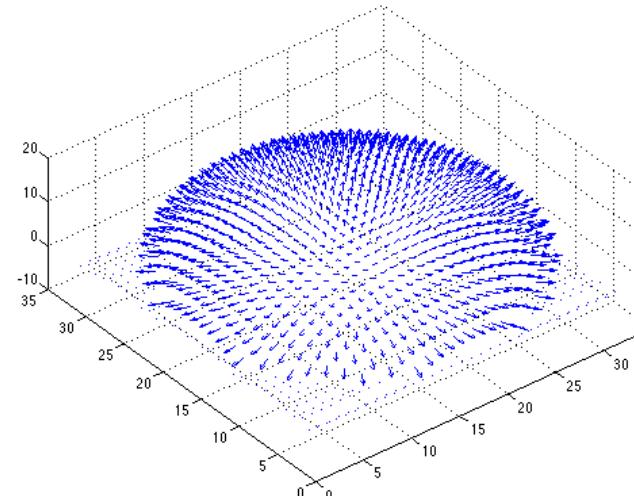
Example 1



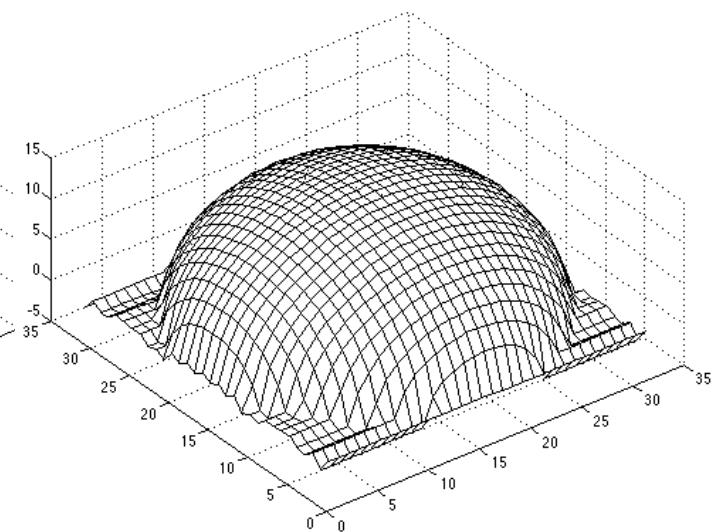
Recovered
albedo



Recovered normal
field



Recovered surface
model



Example 2

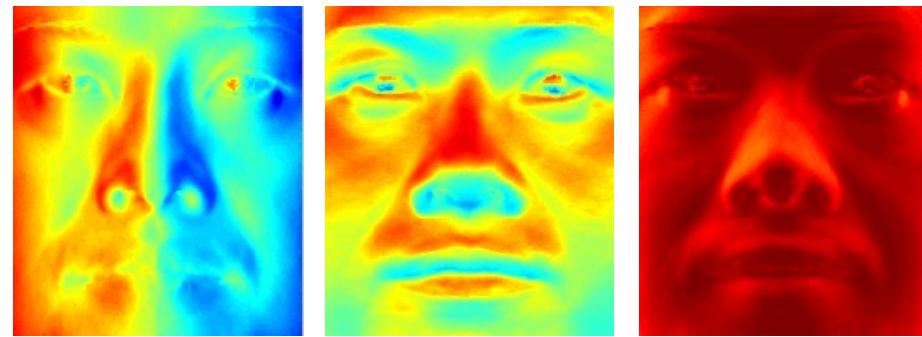
Input



Recovered
albedo



Recovered normal field



Recovered
surface model



Image model

- **Known:** source vectors S_j and pixel values $I_j(x,y)$
- **Unknown:** surface normal $\mathbf{N}(x,y)$ and albedo $\rho(x,y)$

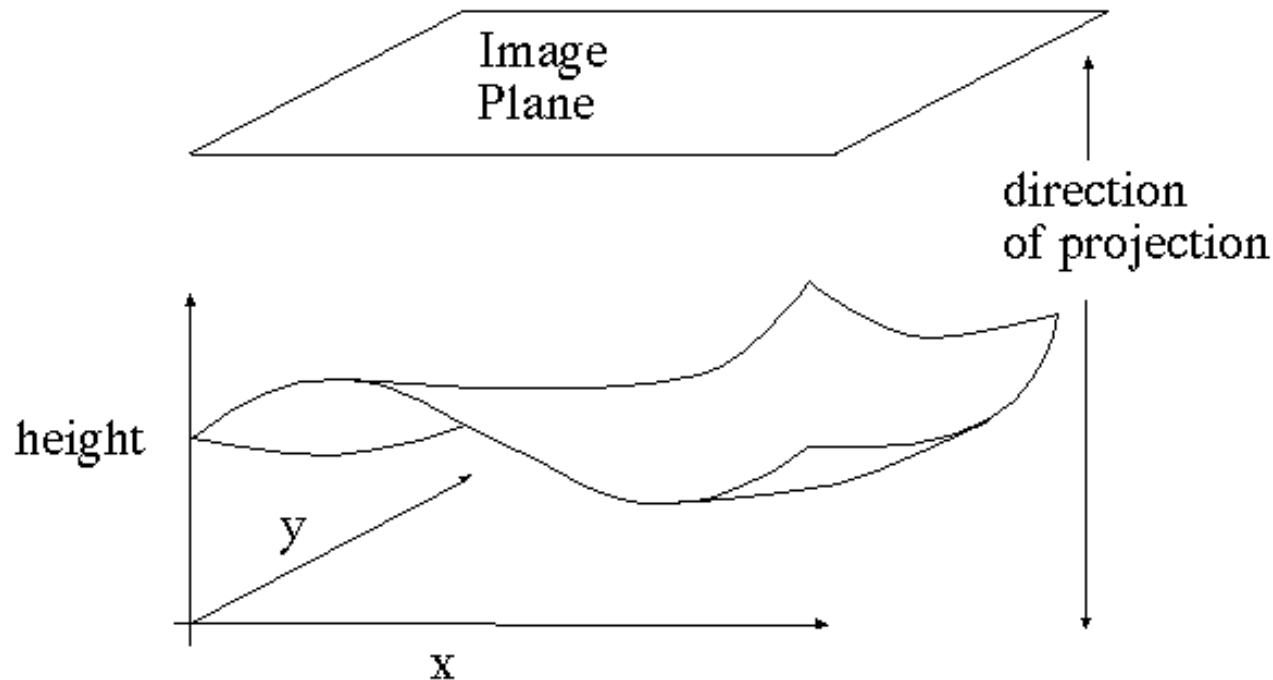


Image model

- **Known:** source vectors \mathbf{S}_j and pixel values $I_j(x,y)$
- **Unknown:** surface normal $\mathbf{N}(x,y)$ and albedo $\rho(x,y)$
- Assume that the response function of the camera is a linear scaling by a factor of k
- Lambert's law:

$$\begin{aligned} I_j(x,y) &= k \rho(x,y) (\mathbf{N}(x,y) \cdot \mathbf{S}_j) \\ &= (\rho(x,y) \mathbf{N}(x,y)) \cdot (k \mathbf{S}_j) \\ &= \mathbf{g}(x,y) \cdot \mathbf{V}_j \end{aligned}$$

Least squares problem

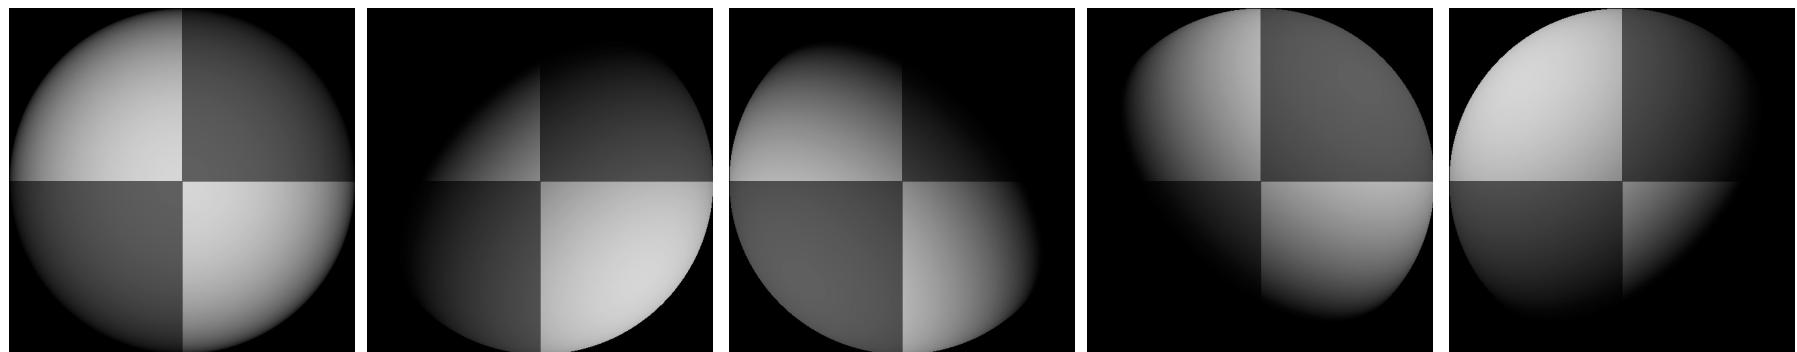
- For each pixel, set up a linear system:

$$\begin{bmatrix} I_1(x, y) \\ I_2(x, y) \\ \vdots \\ I_n(x, y) \end{bmatrix} = \begin{bmatrix} \mathbf{V}_1^T \\ \mathbf{V}_2^T \\ \vdots \\ \mathbf{V}_n^T \end{bmatrix} \mathbf{g}(x, y)$$

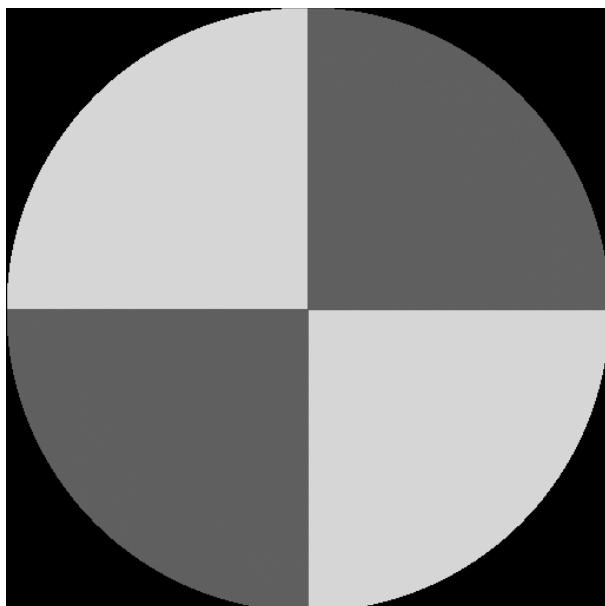
$(n \times 1)$ $(n \times 3)$ (3×1)
known known unknown

- Obtain least-squares solution for $\mathbf{g}(x, y)$ (which we defined as $\mathbf{N}(x, y) \rho(x, y)$)
- Since $\mathbf{N}(x, y)$ is the unit normal, $\rho(x, y)$ is given by the magnitude of $\mathbf{g}(x, y)$
- Finally, $\mathbf{N}(x, y) = \mathbf{g}(x, y) / \rho(x, y)$

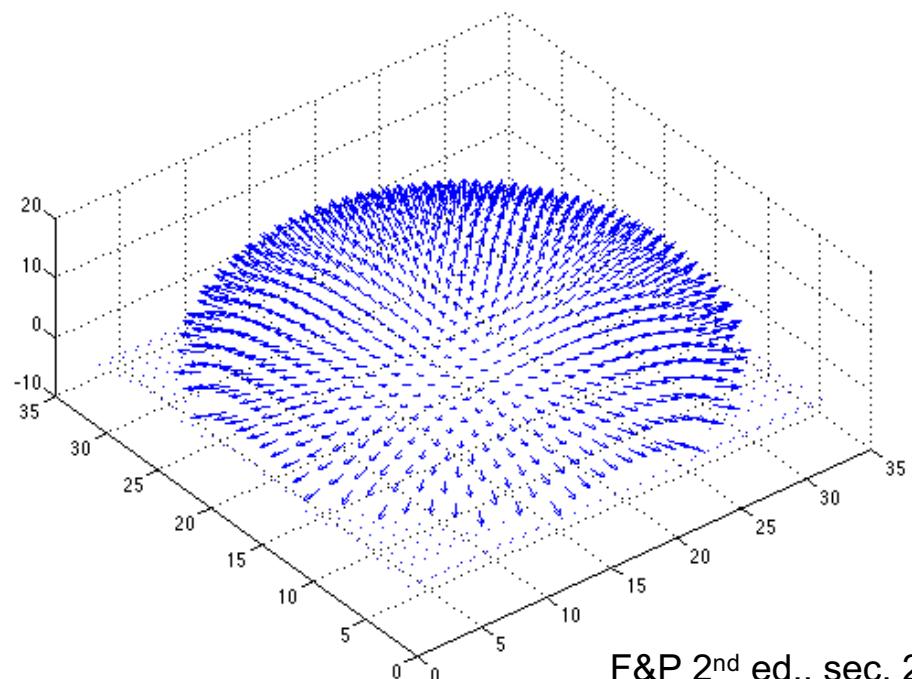
Synthetic example



Recovered albedo



Recovered normal field



Recovering a surface from normals

Recall the surface is written as

$$(x, y, f(x, y))$$

This means the normal has the form:

$$\mathbf{N}(x, y) = \frac{1}{\sqrt{f_x^2 + f_y^2 + 1}} \begin{pmatrix} f_x \\ f_y \\ 1 \end{pmatrix}$$

If we write the estimated vector g as

$$\mathbf{g}(x, y) = \begin{pmatrix} g_1(x, y) \\ g_2(x, y) \\ g_3(x, y) \end{pmatrix}$$

Then we obtain values for the partial derivatives of the surface:

$$f_x(x, y) = g_1(x, y) / g_3(x, y)$$

$$f_y(x, y) = g_2(x, y) / g_3(x, y)$$

Recovering a surface from normals

We can now recover the surface height at any point by integration along some path, e.g.

$$f(x, y) = \int_0^x f_x(s, 0) ds + \int_0^y f_y(x, t) dt + C$$

(for robustness, should take integrals over many different paths and average the results)

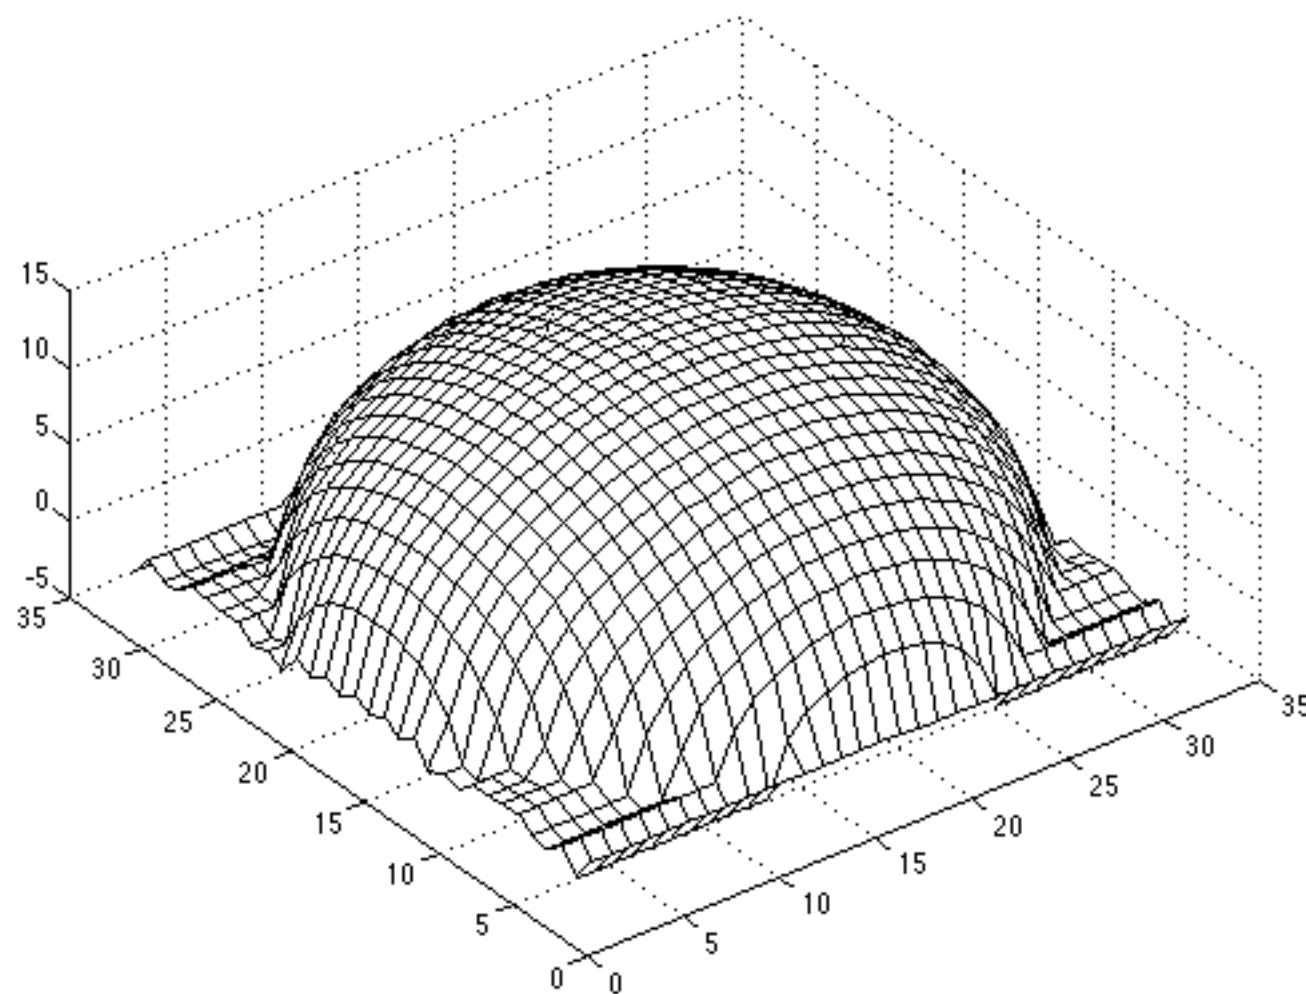
Integrability: for the surface f to exist, the mixed second partial derivatives must be equal:

$$\frac{\partial}{\partial y} (g_1(x, y) / g_3(x, y)) =$$

$$\frac{\partial}{\partial x} (g_2(x, y) / g_3(x, y))$$

(in practice, they should at least be similar)

Surface recovered by integration



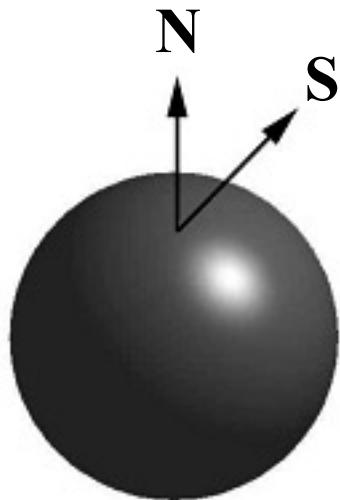
Limitations

- Orthographic camera model
- Simplistic reflectance and lighting model
- No shadows
- No interreflections
- No missing data
- Integration is tricky

Finding the direction of the light source

$$I(x,y) = \mathbf{N}(x,y) \cdot \mathbf{S}(x,y)$$

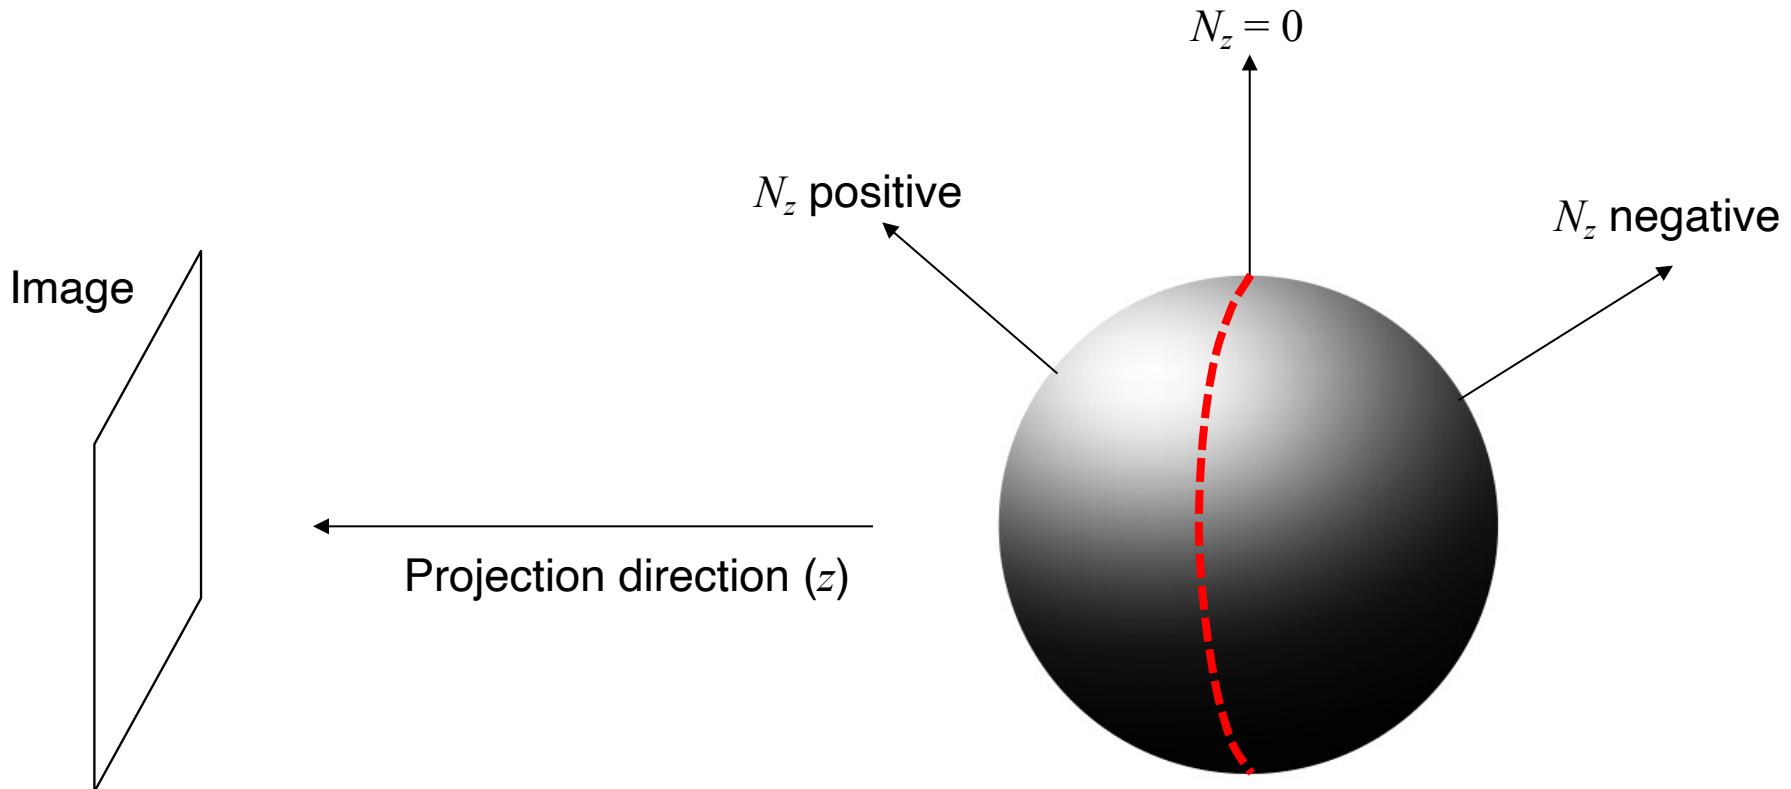
Full 3D case:



$$\begin{pmatrix} N_x(x_1, y_1) & N_y(x_1, y_1) & N_z(x_1, y_1) \\ N_x(x_2, y_2) & N_y(x_2, y_2) & N_z(x_2, y_2) \\ \vdots & \vdots & \vdots \\ N_x(x_n, y_n) & N_y(x_n, y_n) & N_z(x_n, y_n) \end{pmatrix} \begin{pmatrix} S_x \\ S_y \\ S_z \end{pmatrix} = \begin{pmatrix} I(x_1, y_1) \\ I(x_2, y_2) \\ \vdots \\ I(x_n, y_n) \end{pmatrix}$$

Finding the direction of the light source

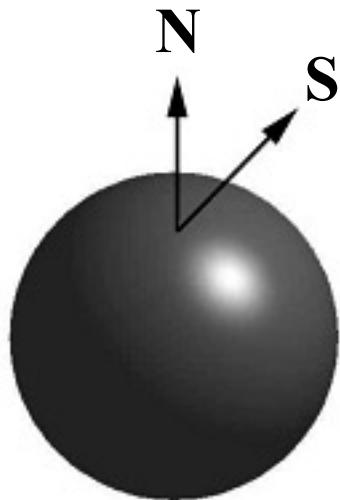
Consider points on the *occluding contour*:



Finding the direction of the light source

$$I(x,y) = \mathbf{N}(x,y) \cdot \mathbf{S}(x,y)$$

Full 3D case:



$$\begin{pmatrix} N_x(x_1, y_1) & N_y(x_1, y_1) & N_z(x_1, y_1) \\ N_x(x_2, y_2) & N_y(x_2, y_2) & N_z(x_2, y_2) \\ \vdots & \vdots & \vdots \\ N_x(x_n, y_n) & N_y(x_n, y_n) & N_z(x_n, y_n) \end{pmatrix} \begin{pmatrix} S_x \\ S_y \\ S_z \end{pmatrix} = \begin{pmatrix} I(x_1, y_1) \\ I(x_2, y_2) \\ \vdots \\ I(x_n, y_n) \end{pmatrix}$$

For points on the *occluding contour*, $N_z = 0$:

$$\begin{pmatrix} N_x(x_1, y_1) & N_y(x_1, y_1) \\ N_x(x_2, y_2) & N_y(x_2, y_2) \\ \vdots & \vdots \\ N_x(x_n, y_n) & N_y(x_n, y_n) \end{pmatrix} \begin{pmatrix} S_x \\ S_y \end{pmatrix} = \begin{pmatrix} I(x_1, y_1) \\ I(x_2, y_2) \\ \vdots \\ I(x_n, y_n) \end{pmatrix}$$

Finding the direction of the light source



P. Nillius and J.-O. Eklundh, “Automatic estimation of the projected light source direction,” CVPR 2001

Application: Detecting composite photos

Fake photo



Real photo



Color

CS 543 / ECE 549 – Saurabh Gupta

While we wait, what color is the dress?

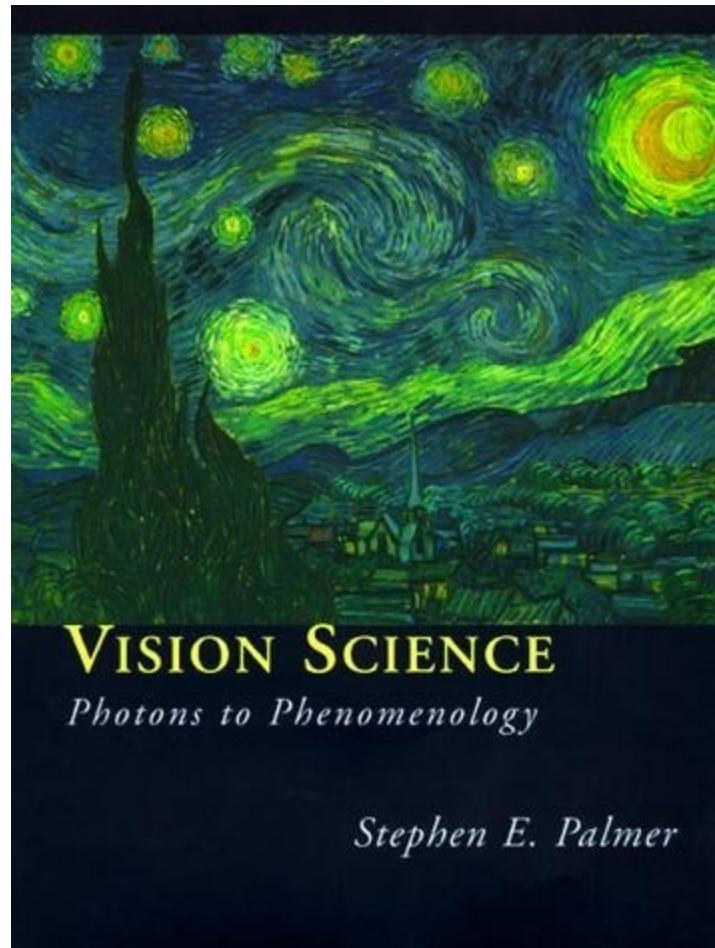


- A. White and gold
- B. Blue and Black

<https://www.wired.com/2015/02/science-one-agrees-color-dress/>

What is color?

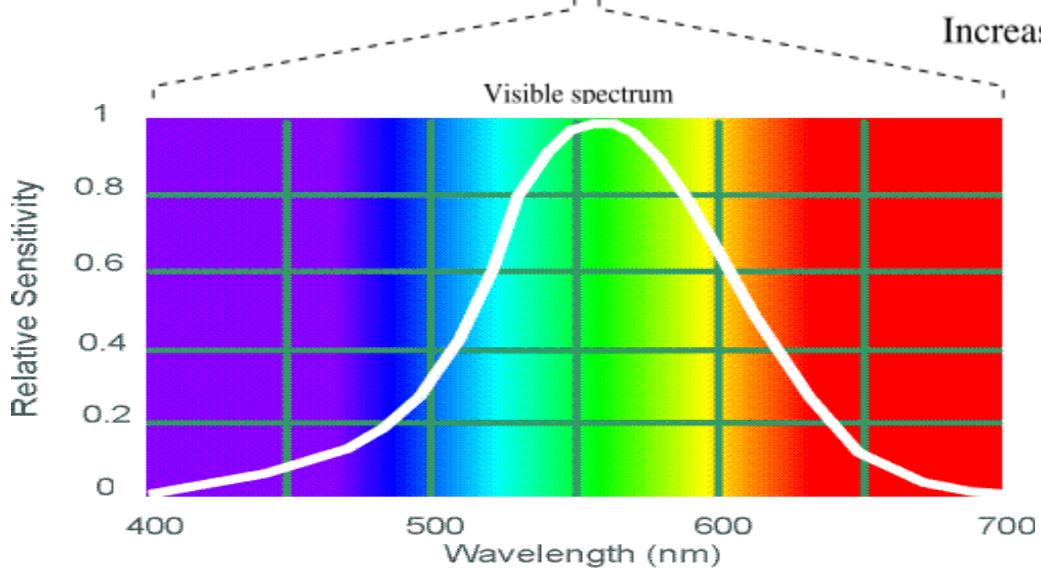
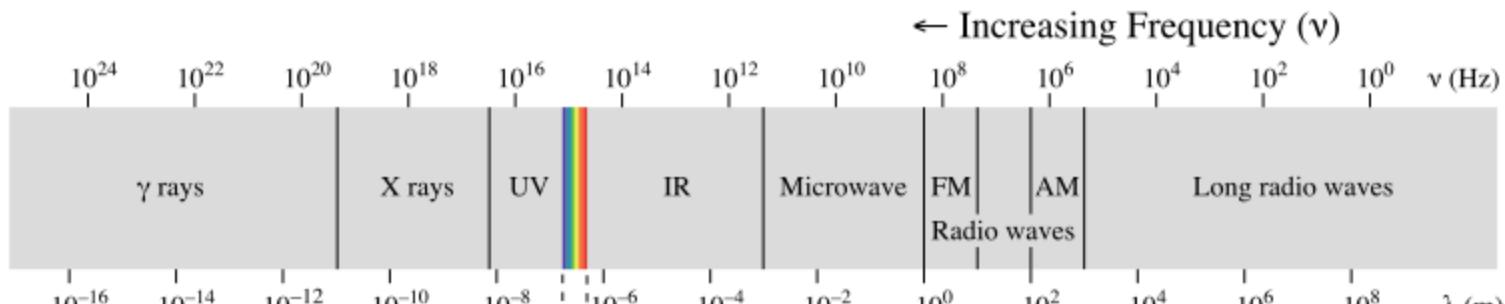
- Color is the result of interaction between physical light in the environment and our visual system
- Color is a psychological property of our visual experiences when we look at objects and lights, *not* a physical property of those objects or lights
(S. Palmer, *Vision Science: Photons to Phenomenology*)



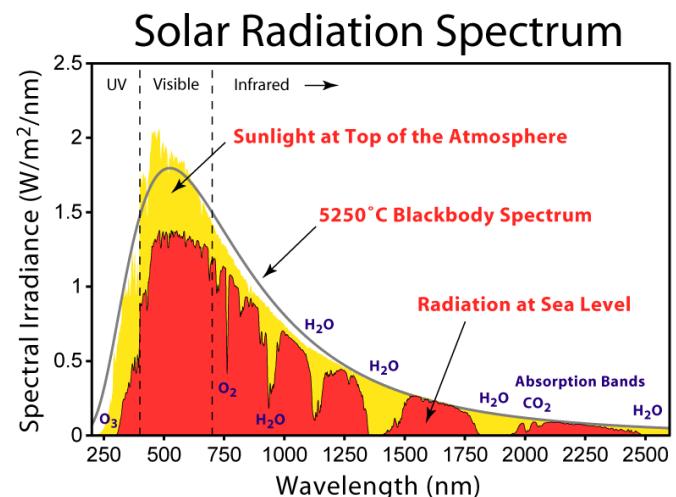
Outline

- Physical origin of color
- Spectra of sources and surfaces
- Physiology of color vision
- Trichromatic color theory
- Color spaces
- Color constancy, white balance

Electromagnetic spectrum

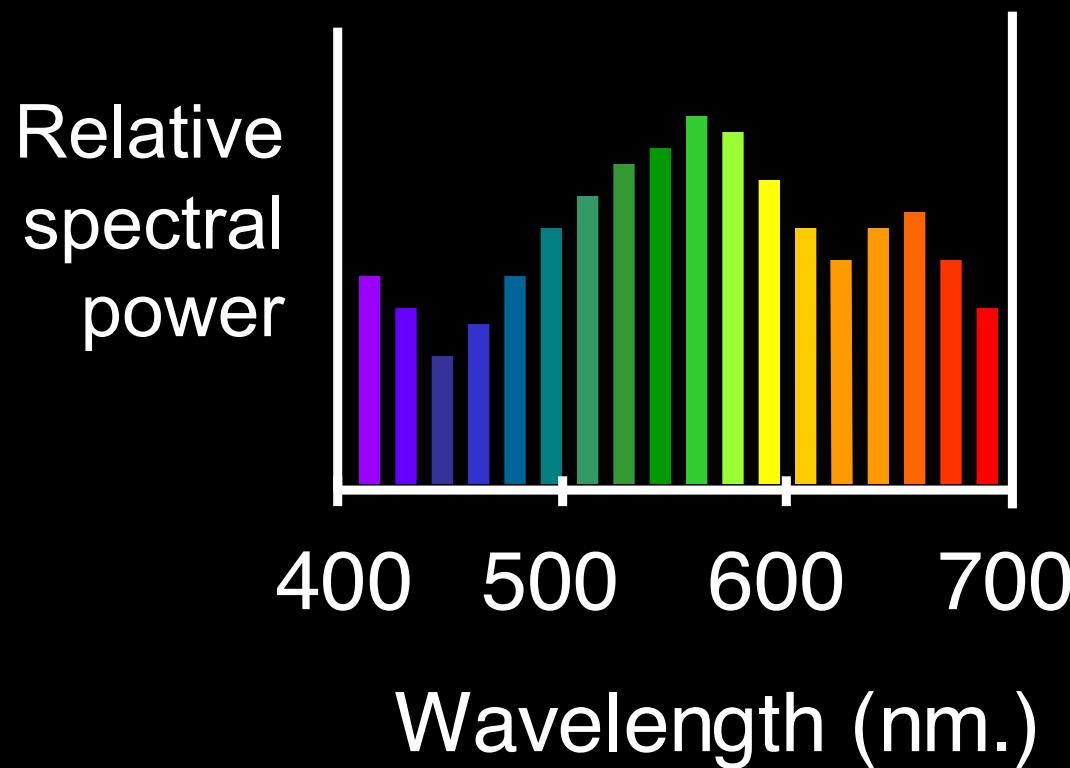


Human Luminance Sensitivity Function



The Physics of Light

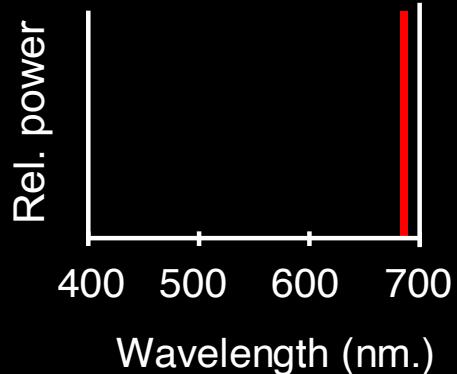
Any source of light can be completely described physically by its spectrum: the amount of energy emitted (per time unit) at each wavelength 400 - 700 nm.



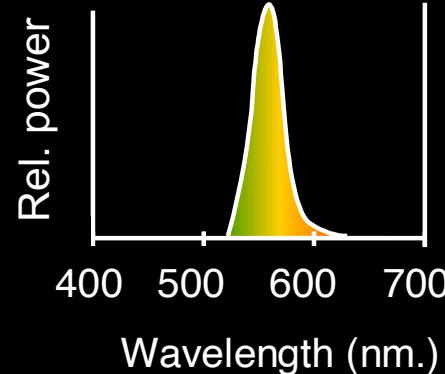
Spectra of Light Sources

Some examples of the spectra of light sources

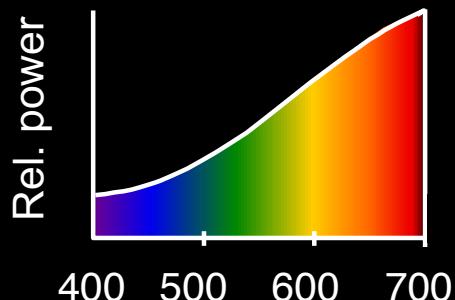
A. Ruby Laser



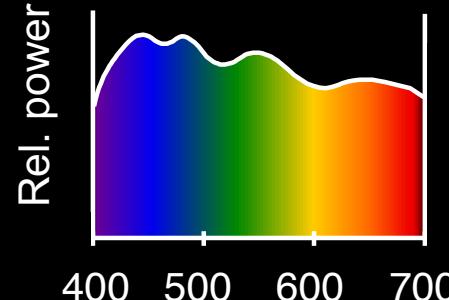
B. Gallium Phosphide Crystal



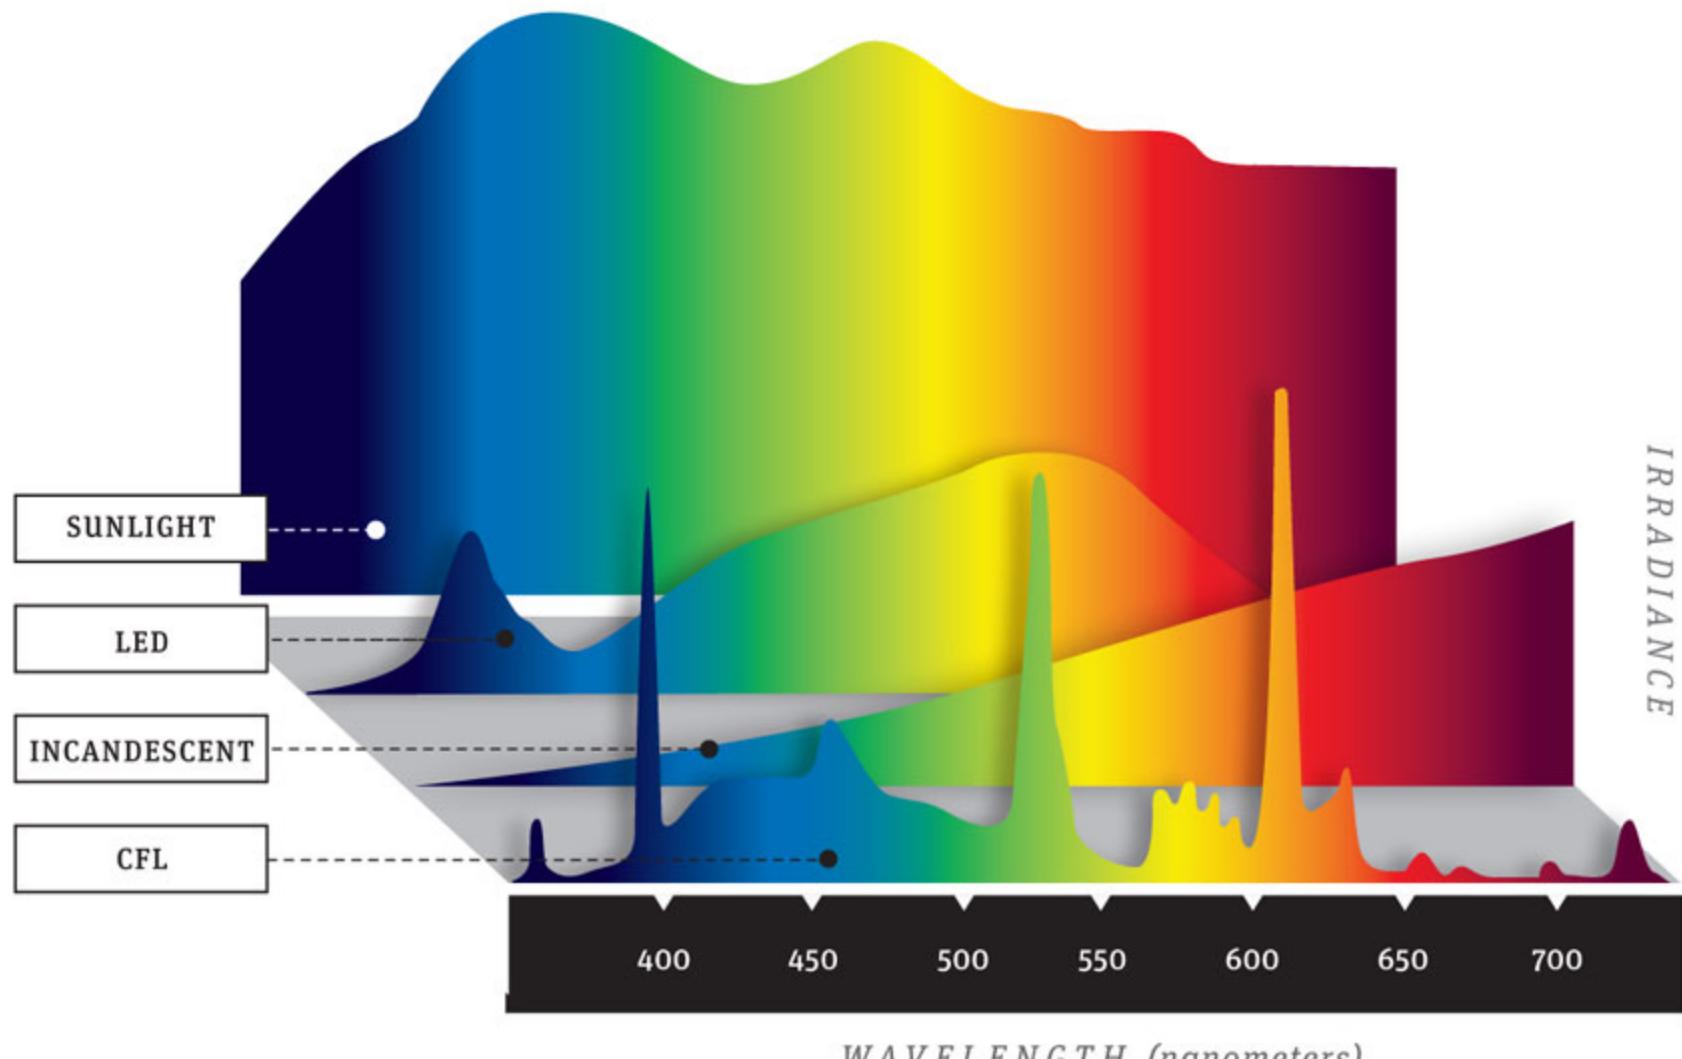
C. Tungsten Lightbulb



D. Normal Daylight



Spectra of light sources

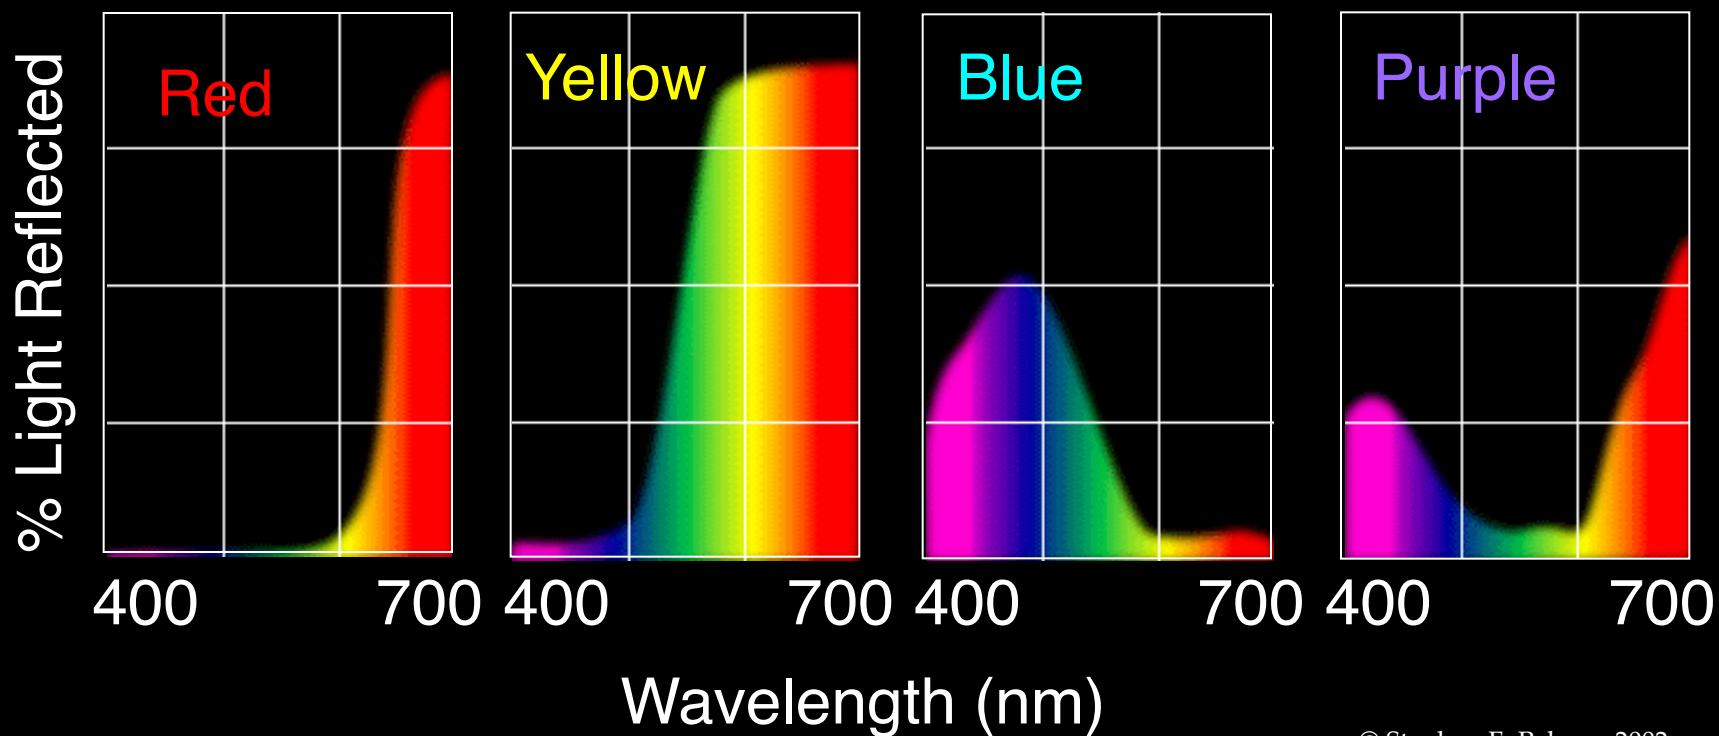


[Source: Popular Mechanics](#)

Slide by L. Lazebnik

Reflectance Spectra of Surfaces

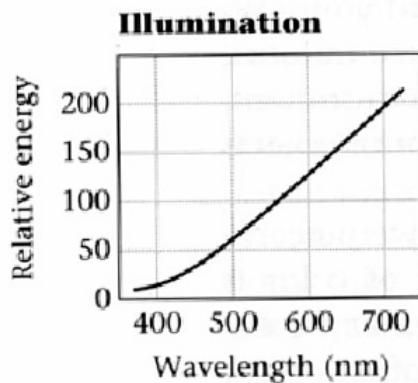
Some examples of the reflectance spectra of surfaces



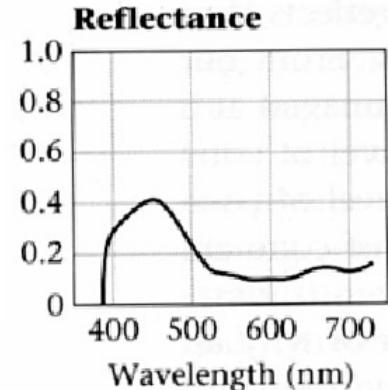
Interaction of light and surfaces



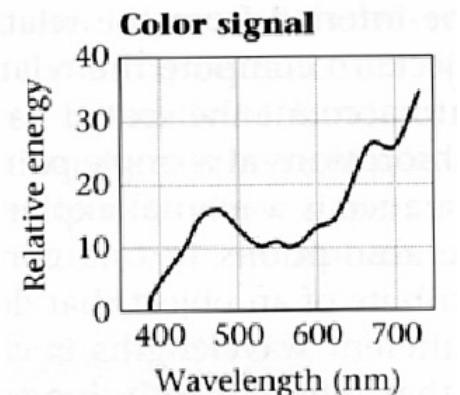
- Reflected color is the result of interaction of light source spectrum with surface reflectance



• *



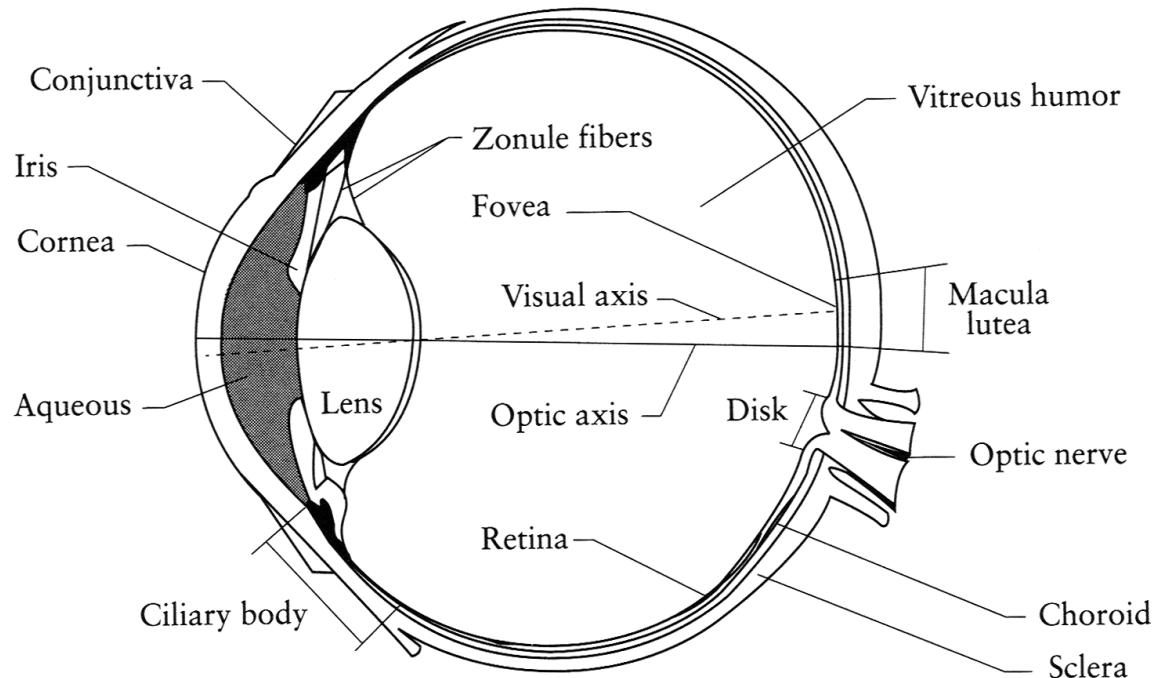
=



So, what is the dimensionality of color?

- A. 3
- B. $(255)^3$
- C. 255
- D. Infinite

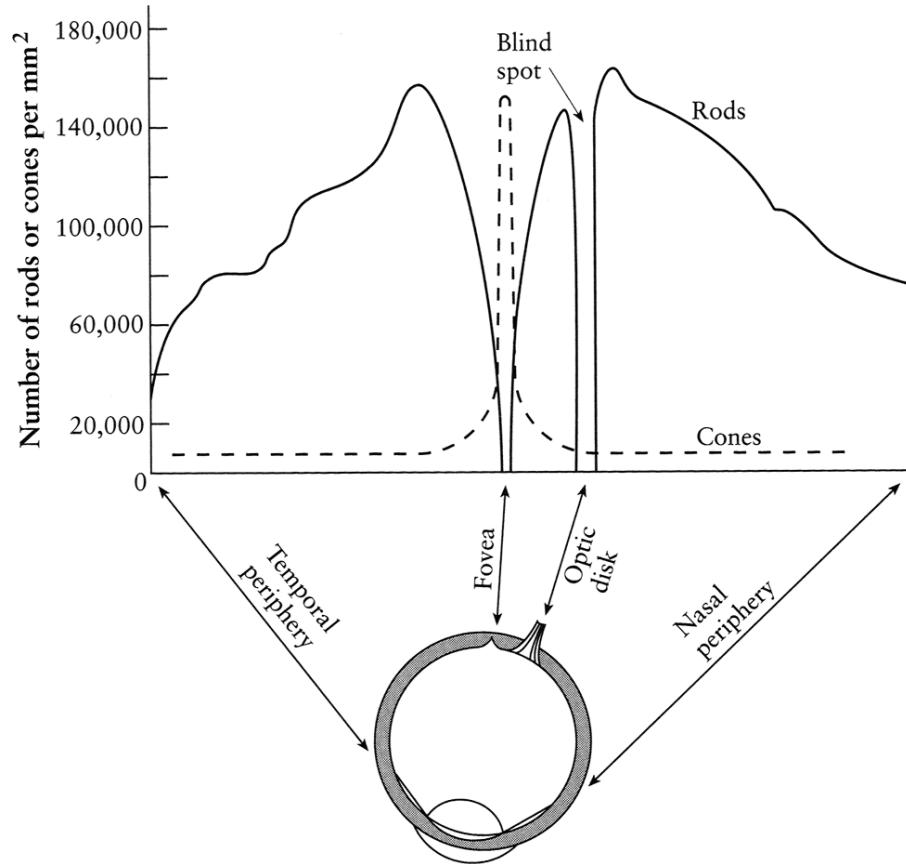
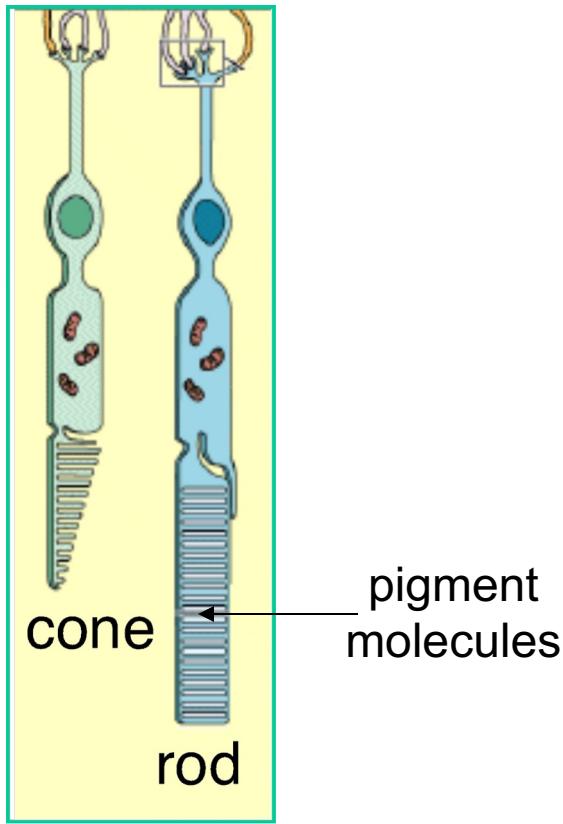
The Eye



The human eye is a camera!

- **Lens** - changes shape by using ciliary muscles (to focus on objects at different distances)
- **Pupil** - the hole (aperture) whose size is controlled by the iris
- **Iris** - colored annulus with radial muscles
- **Retina** - photoreceptor cells

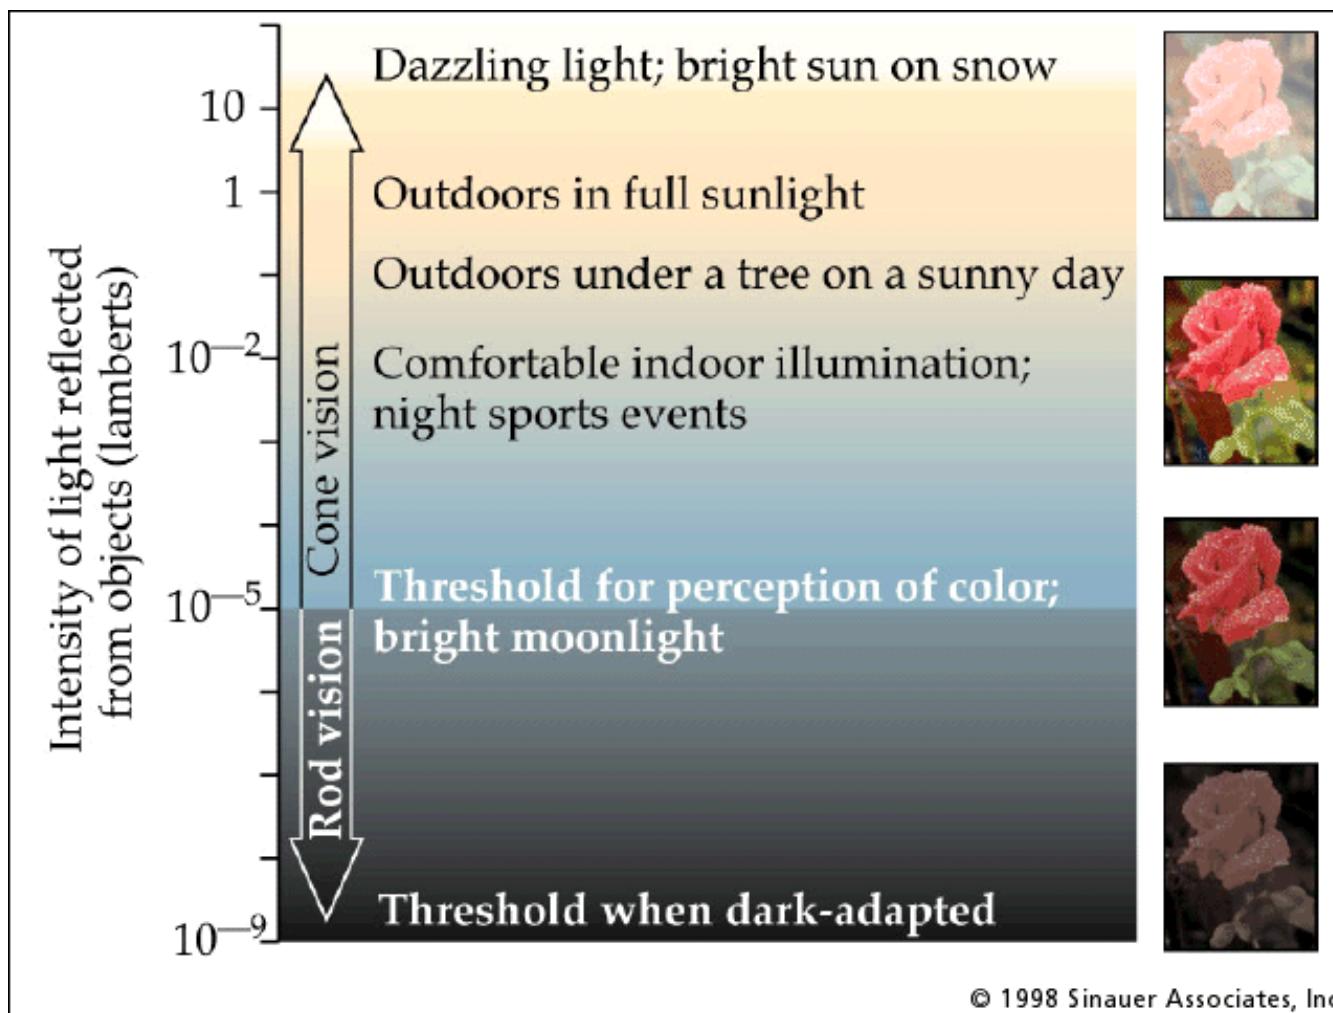
Rods and cones, fovea



Rods are responsible for intensity, cones for color perception
Rods and cones are *non-uniformly* distributed on the retina

- **Fovea** - Small region (1 or 2°) at the center of the visual field containing the highest density of cones – *and no rods*

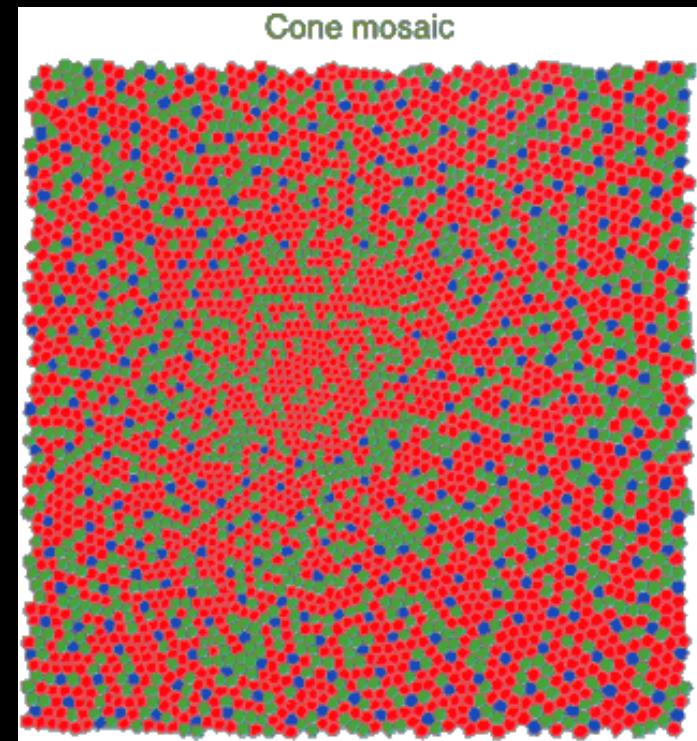
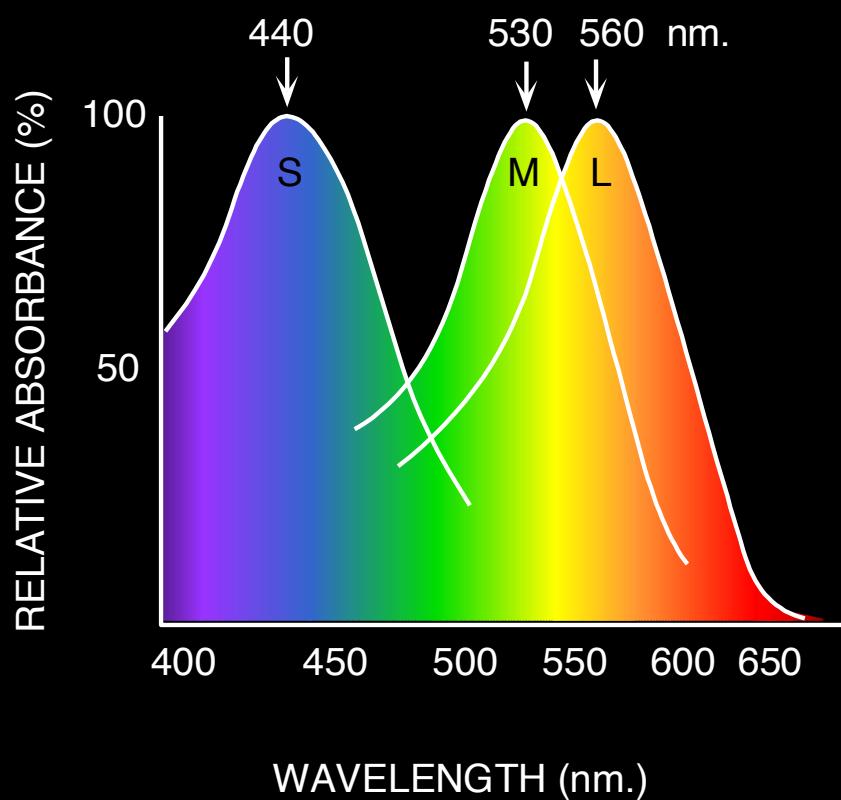
Rod / Cone sensitivity



Why can't we read in the dark?

Physiology of Color Vision

Three kinds of cones:



- Ratio of L to M to S cones: approx. 10:5:1
- Almost no S cones in the center of the fovea

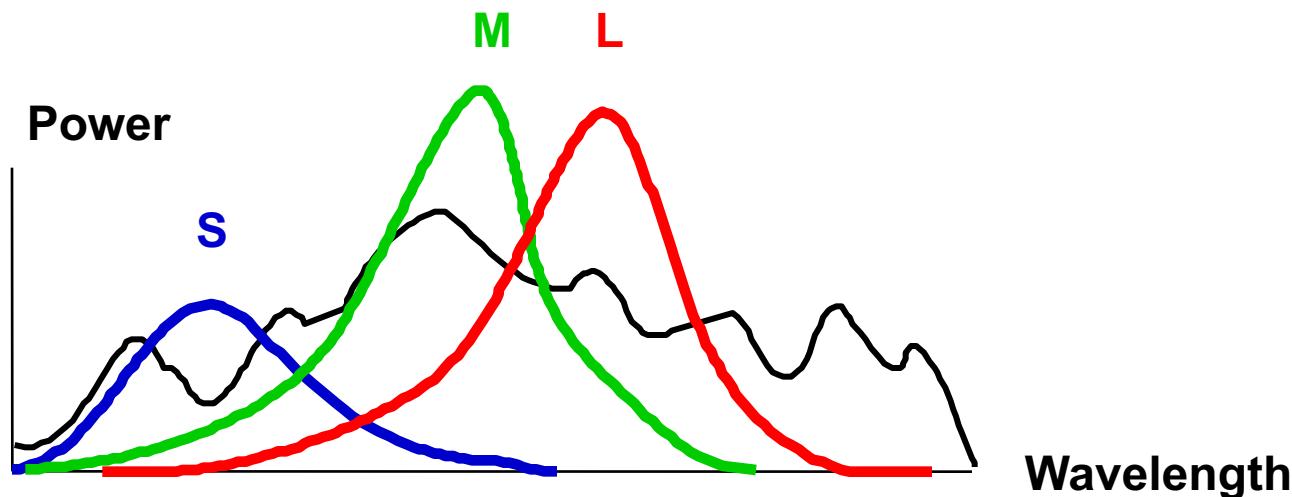
Physiology of Color Vision: Fun facts

- “M” and “L” pigments are encoded on the X-chromosome
 - That’s why men are more likely to be color blind
 - “L” gene has high variation, so some women may be *tetrachromatic*
- Some animals have one (night animals), two (e.g., dogs), four (fish, birds), five (pigeons, some reptiles/amphibians), or even 12 (mantis shrimp) types of cones

<http://ngm.nationalgeographic.com/2016/02/evolution-of-eyes-text>

http://en.wikipedia.org/wiki/Color_vision

Color perception

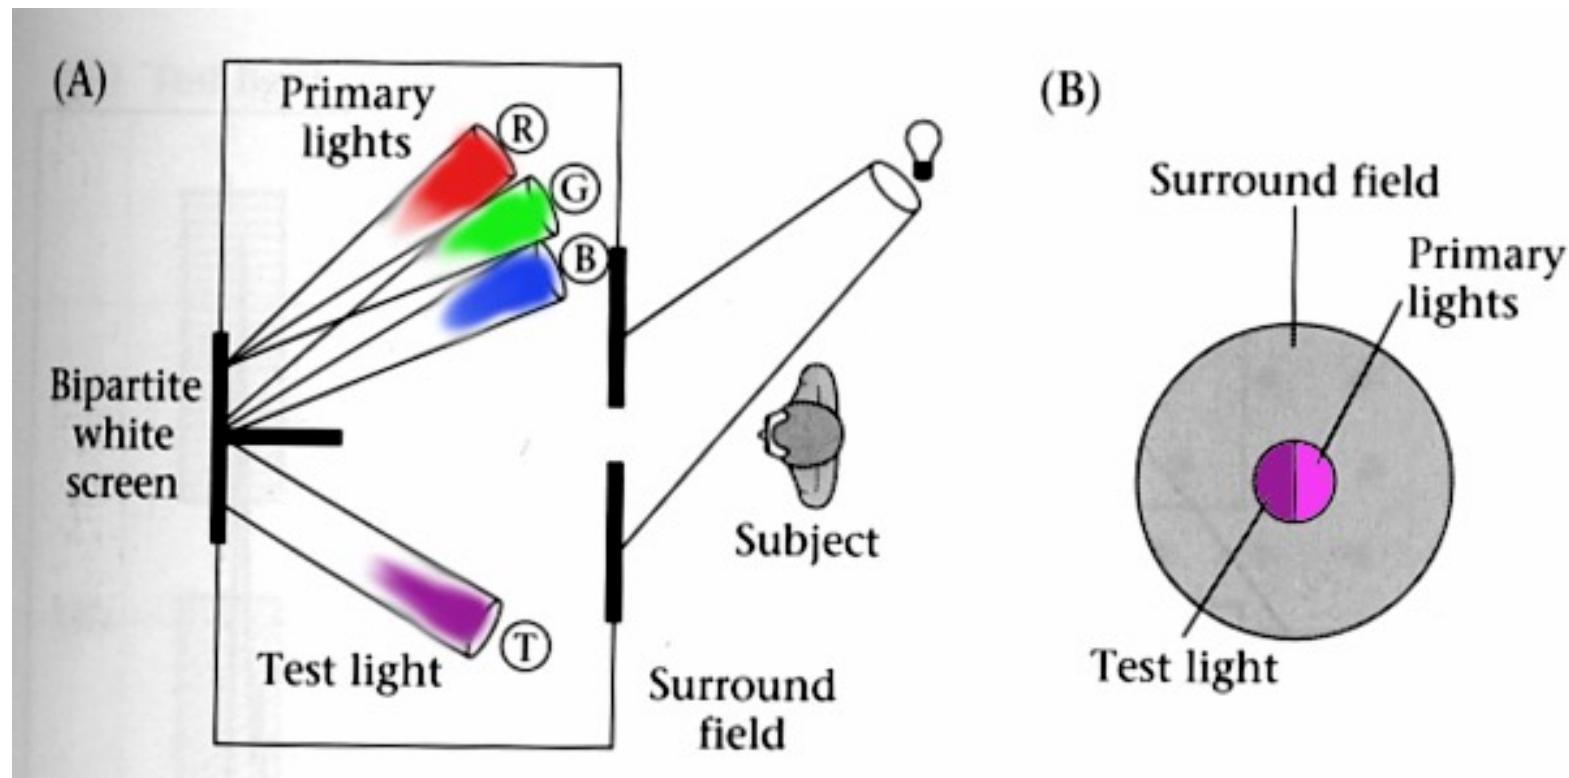


Rods and cones act as *filters* on the spectrum

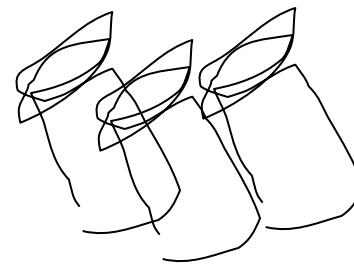
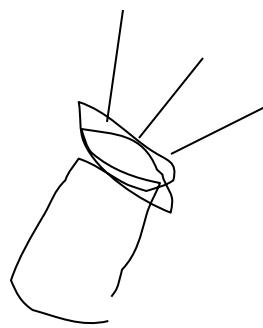
- To get the output of a filter, multiply its response curve by the spectrum, integrate over all wavelengths
 - Each cone yields one number
- How can we represent an entire spectrum with three numbers?
- We can't! Most of the information is lost
 - As a result, two different spectra may appear indistinguishable
 - » such spectra are known as **metamers**

Color matching experiments

- We would like to understand which spectra produce the same color sensation in people under similar viewing conditions

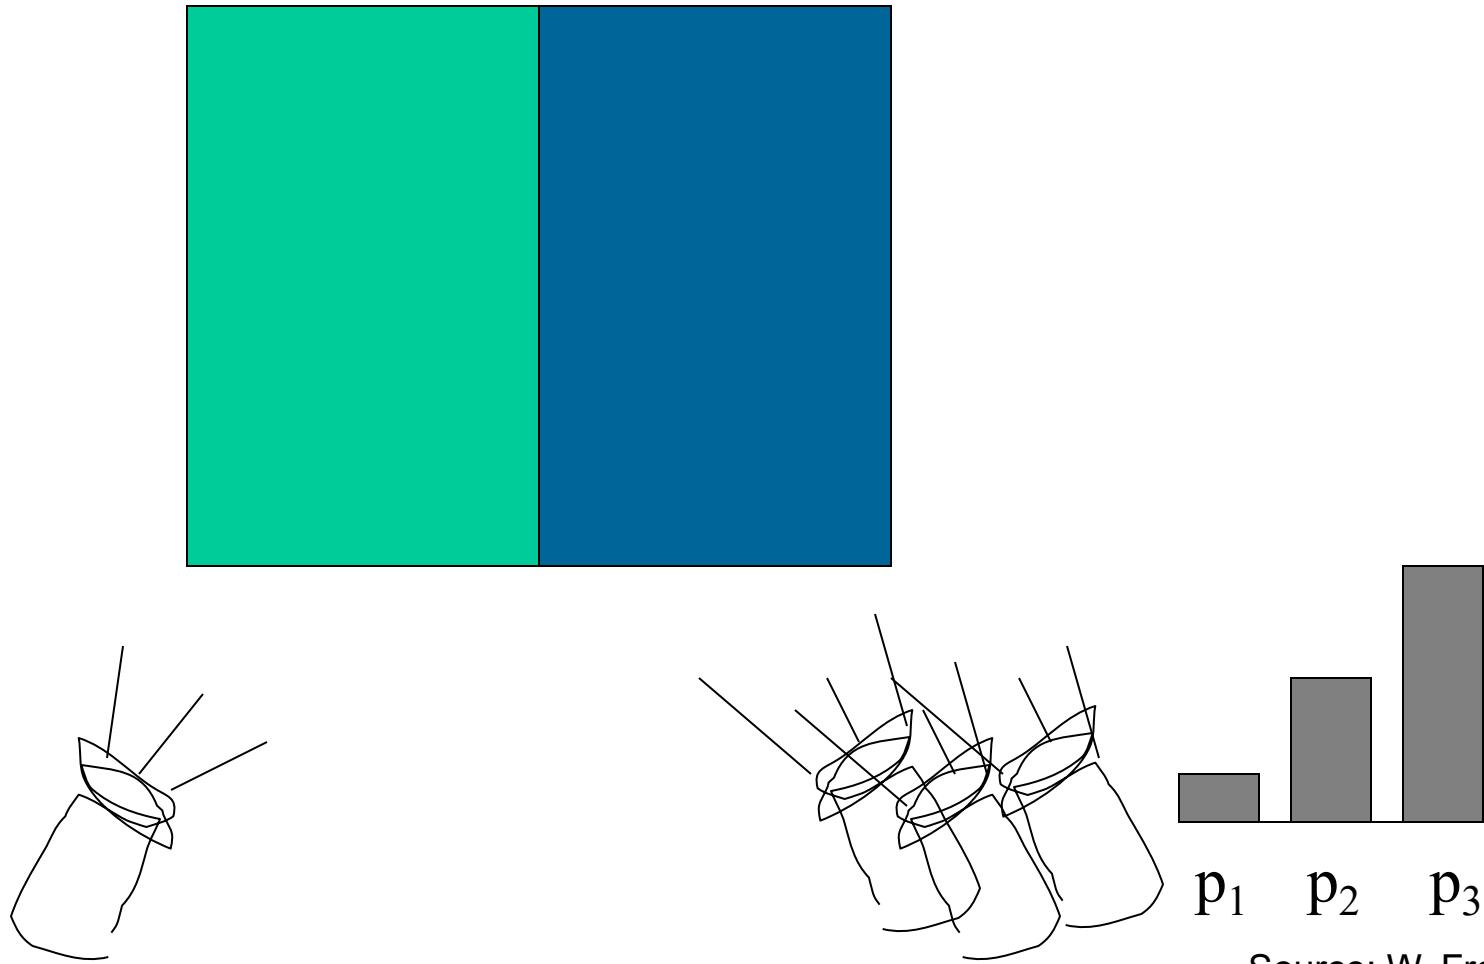


Color matching experiment 1



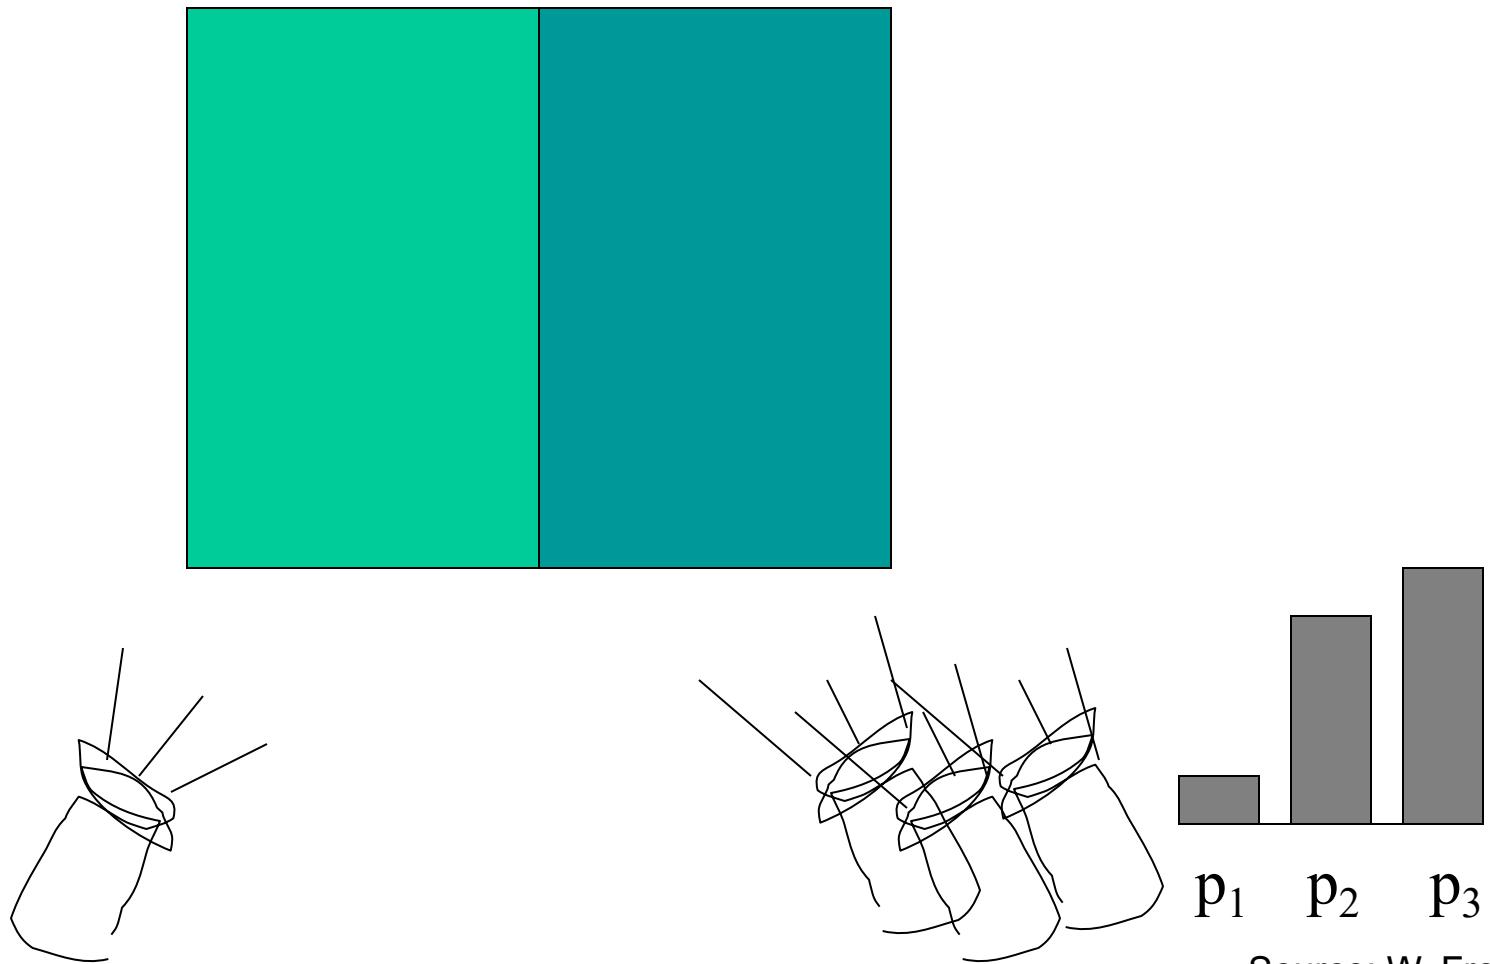
Source: W. Freeman

Color matching experiment 1



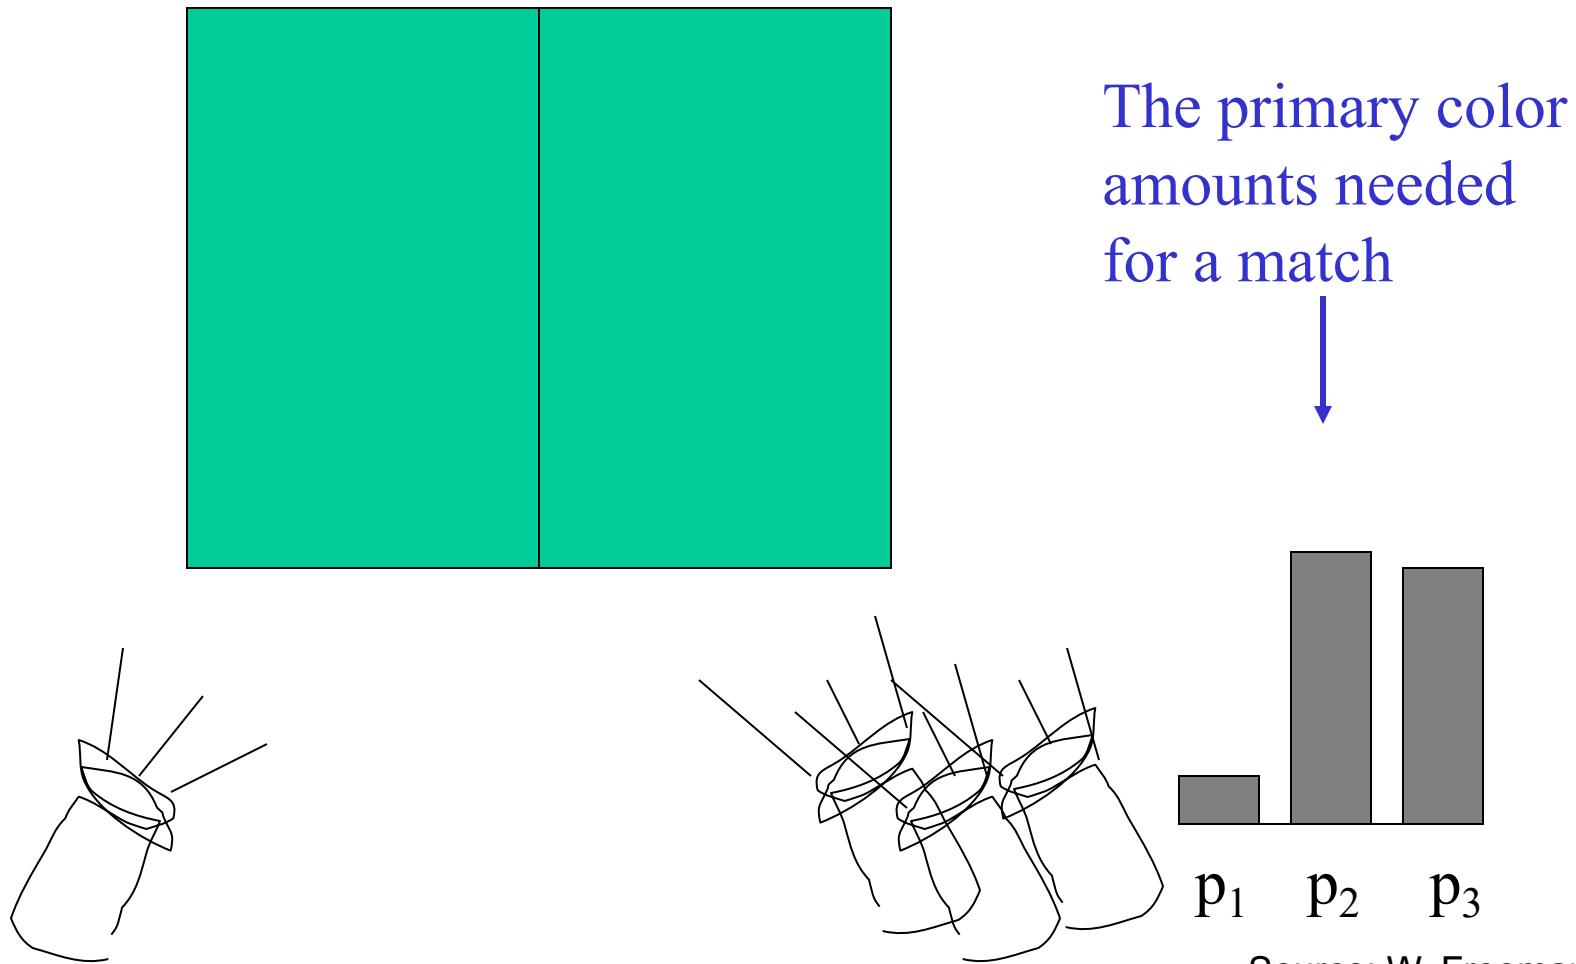
Source: W. Freeman

Color matching experiment 1

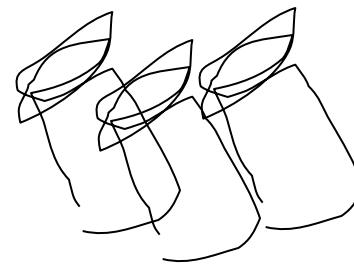
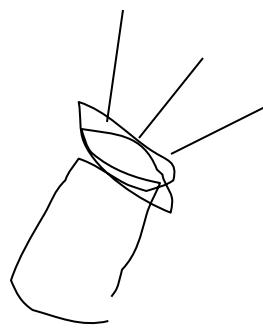


Source: W. Freeman

Color matching experiment 1

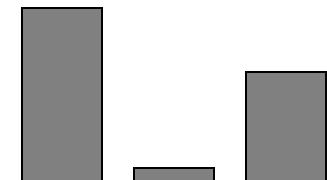
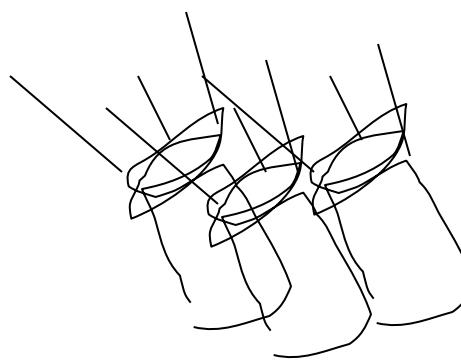
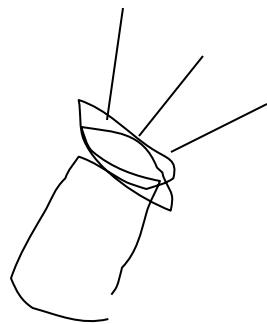
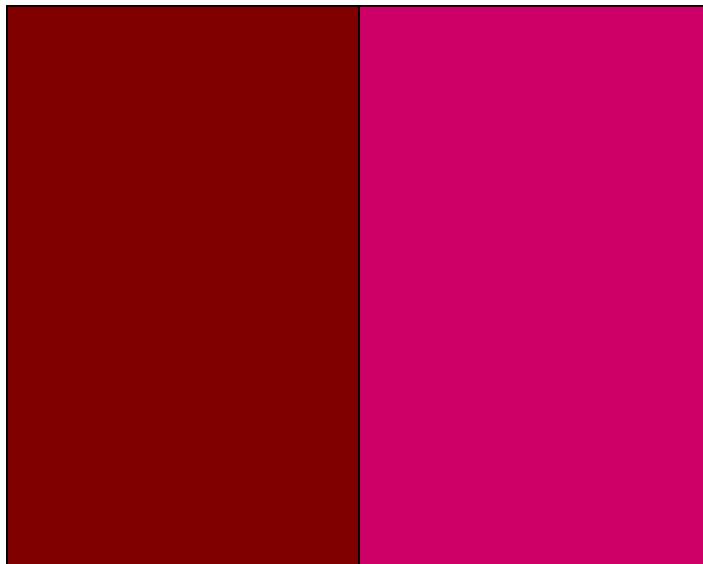


Color matching experiment 2



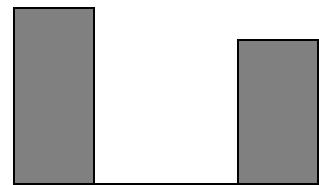
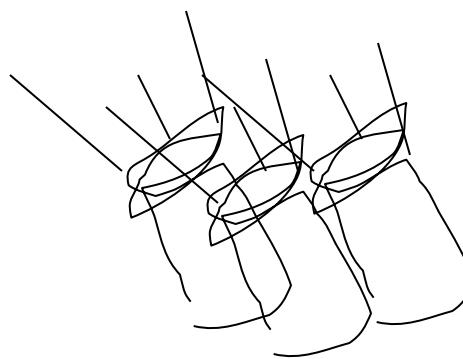
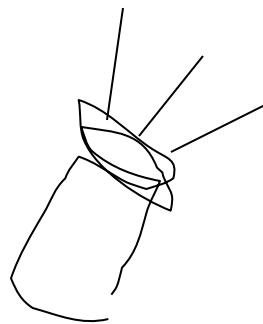
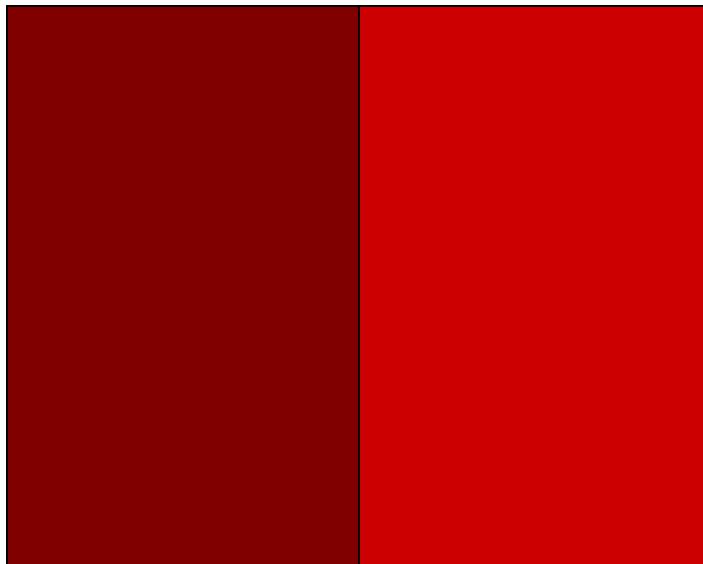
Source: W. Freeman

Color matching experiment 2



Source: W. Freeman

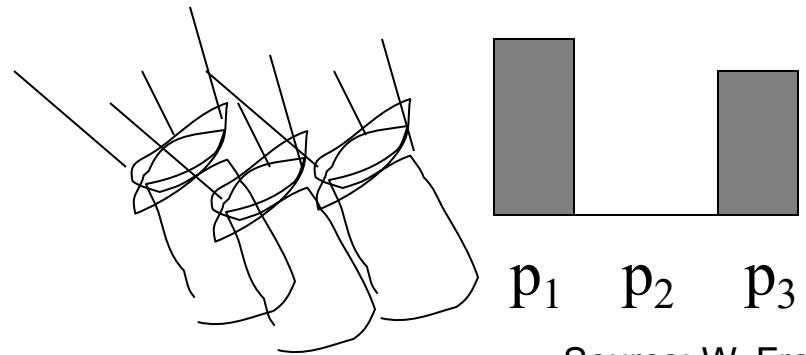
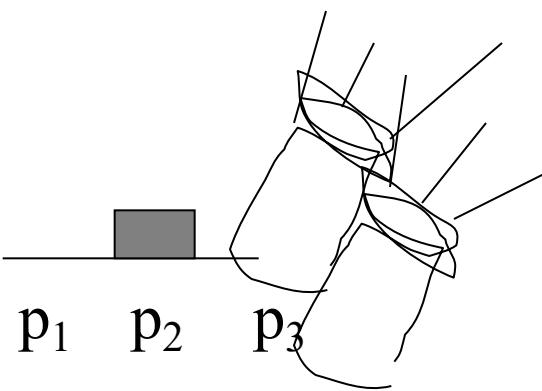
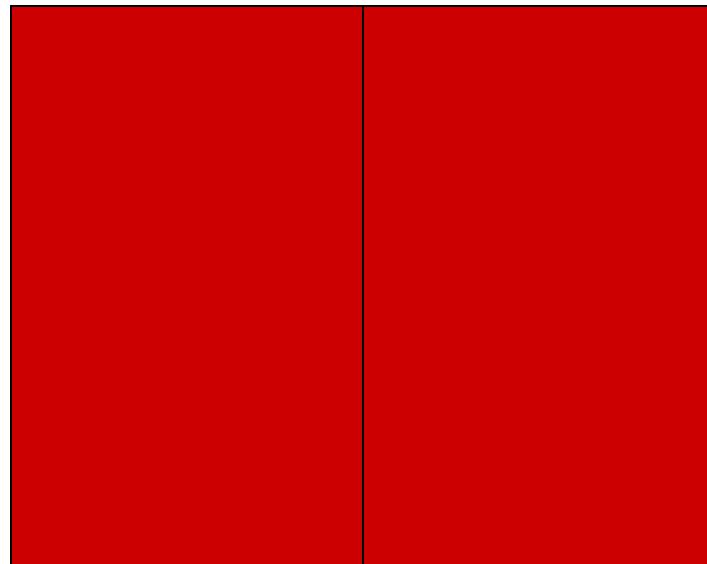
Color matching experiment 2



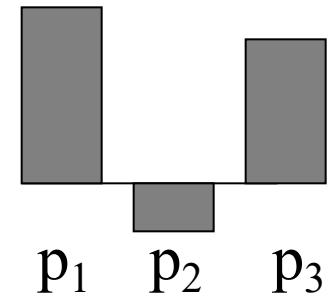
Source: W. Freeman

Color matching experiment 2

We say a “negative” amount of p_2 was needed to make the match, because we added it to the test color’s side.



The primary color amounts needed for a match:

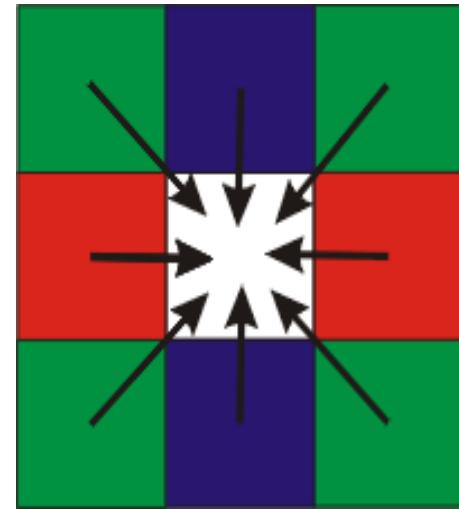
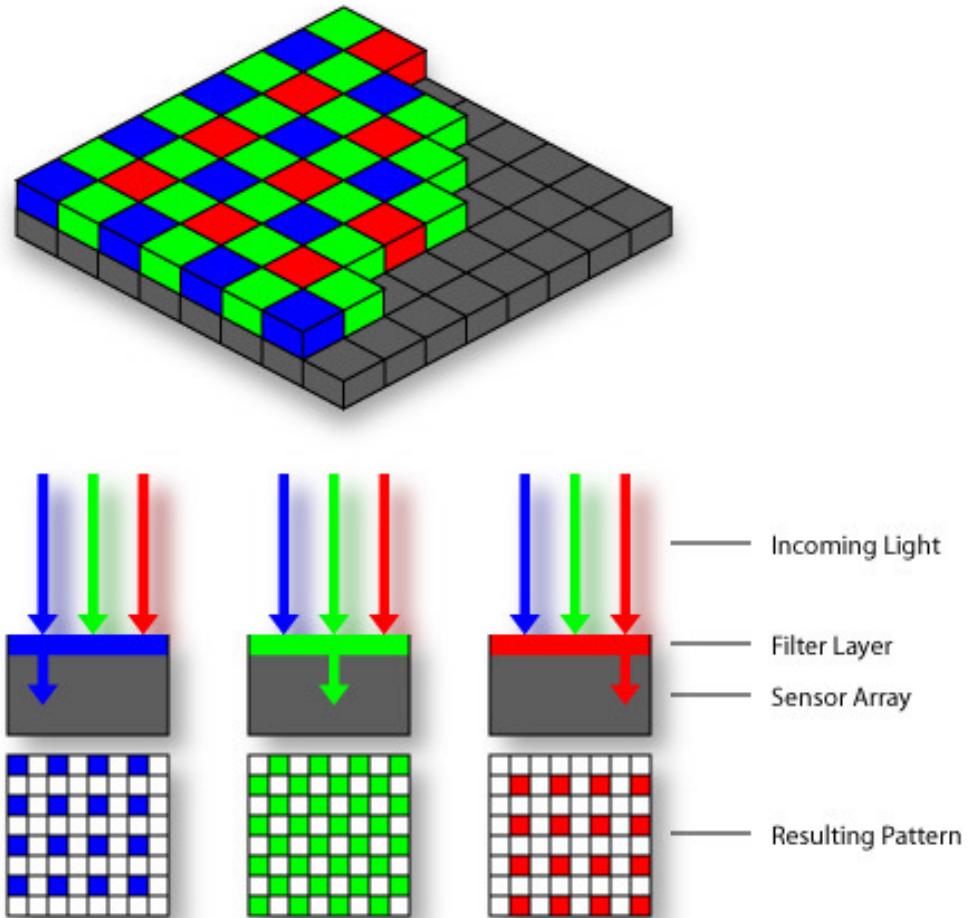


Trichromacy

- In color matching experiments, most people can match any given light with three primaries
 - Primaries must be *independent*
- For the same light and same primaries, most people select the same weights
 - Exception: color blindness
- Trichromatic color theory
 - Three numbers seem to be sufficient for encoding color
 - Dates back to 18th century (Thomas Young)

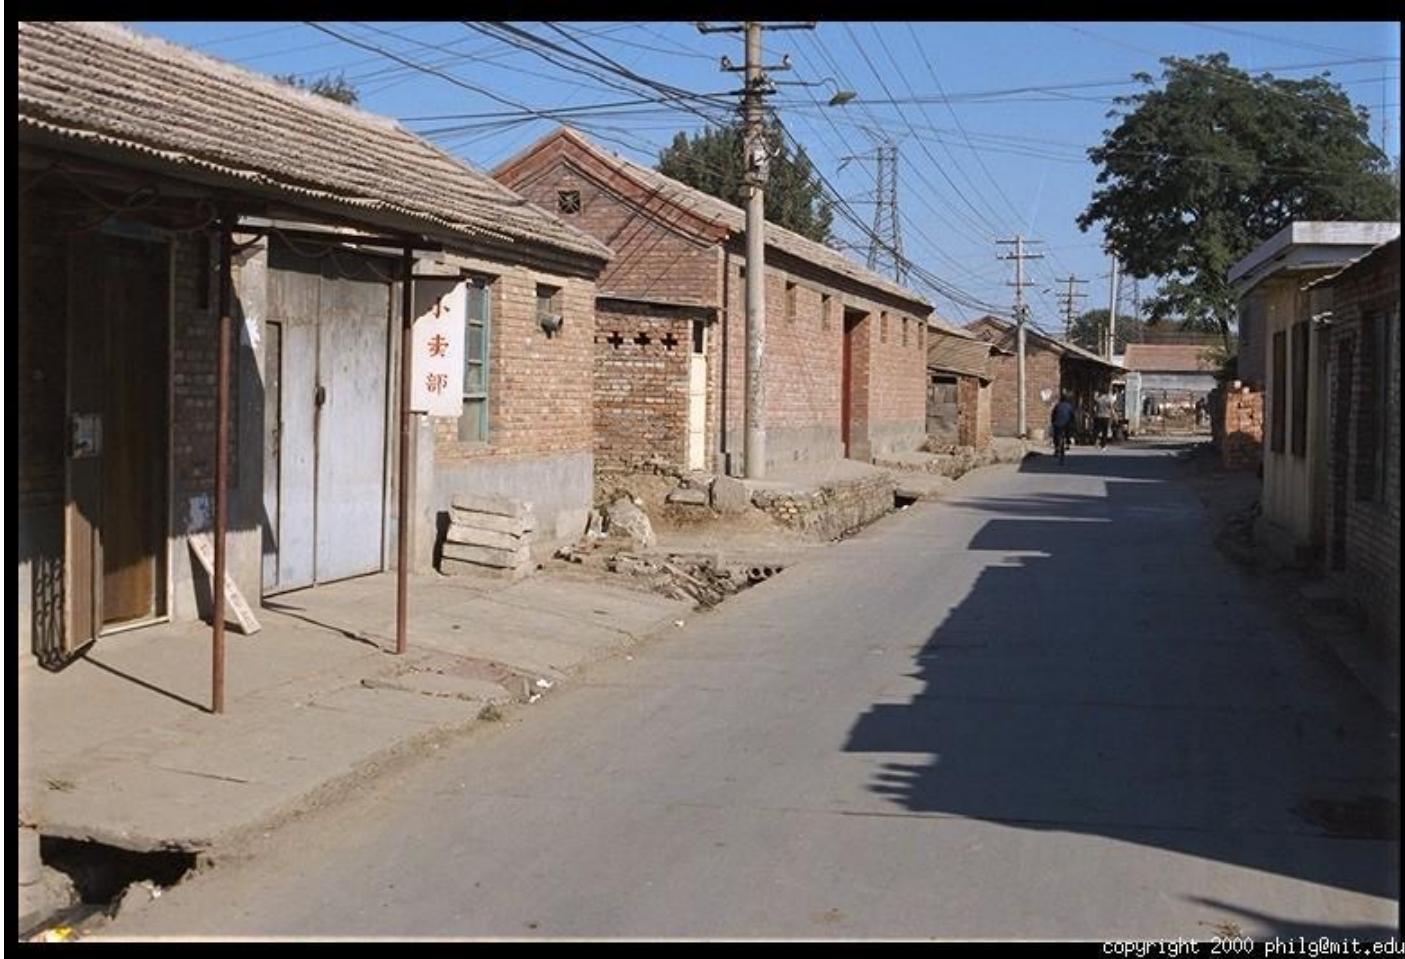
https://en.wikipedia.org/wiki/Young_Helmholtz_theory

Artificial Cones



Estimate RGB at 'G' cells from neighboring values.

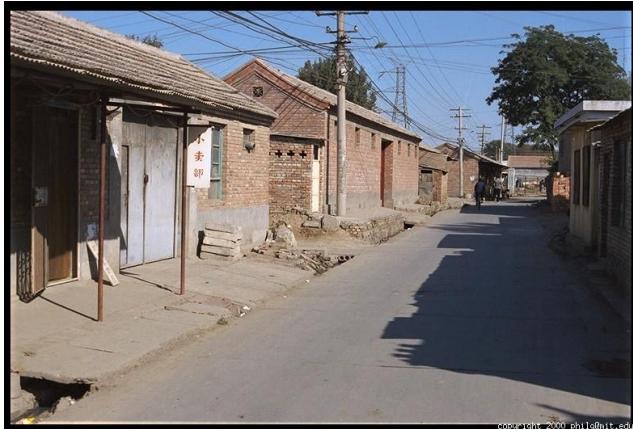
Color Image



copyright 2000 philg@mit.edu

Color Image

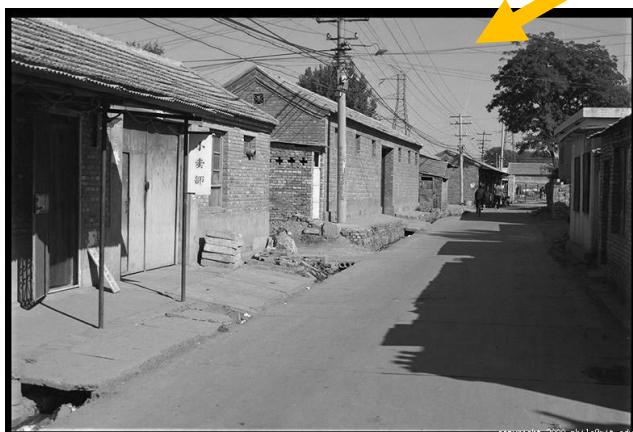
Combined



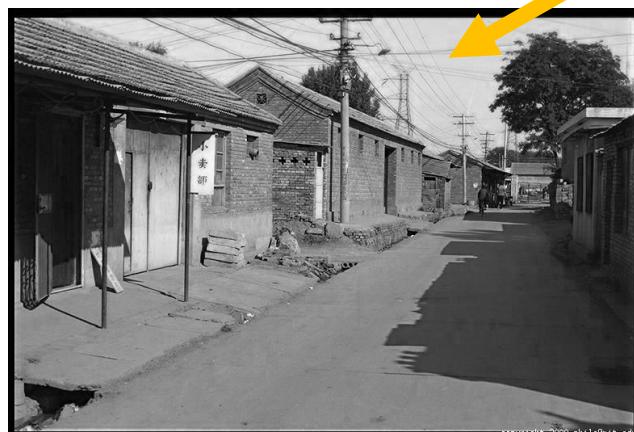
Red



Green



Blue



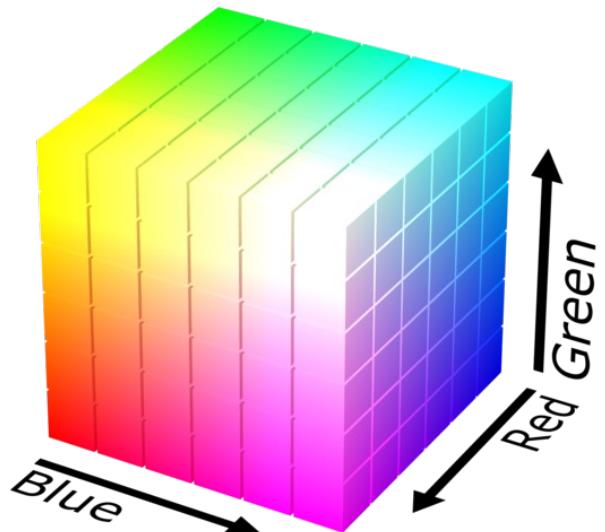
One Option: RGB

Pros

1. Simple
2. Common

Cons

1. Distances don't make sense
2. Correlated



R



G



B

RGB



Photo credit: J. Hays

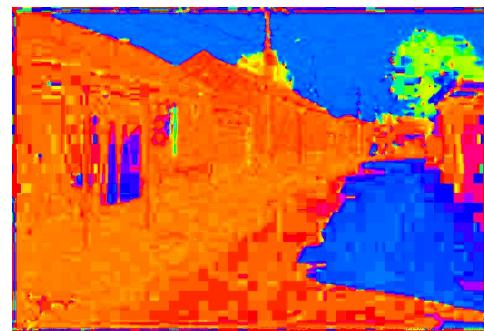
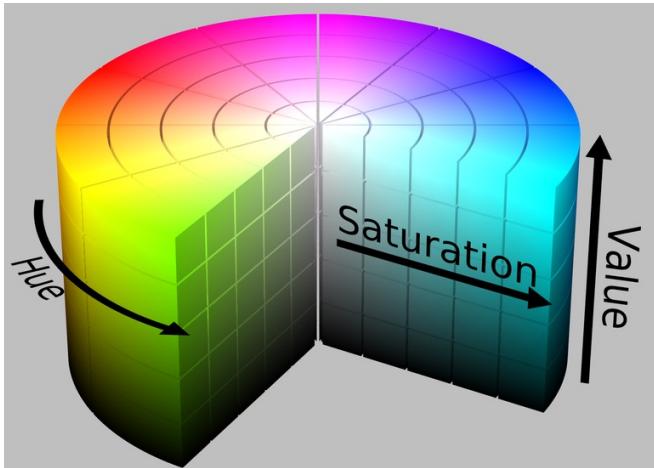
Another Option: HSV

Pros

1. Intuitive for picking colors
2. Sort of common
3. Fast to convert

Cons

1. Not as perceptual



HSV

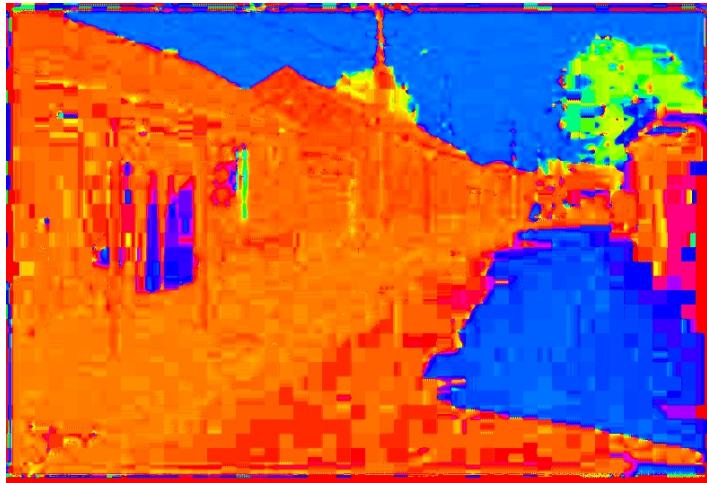


Photo credit: J. Hays

copyright 2000 phigmalit.edu

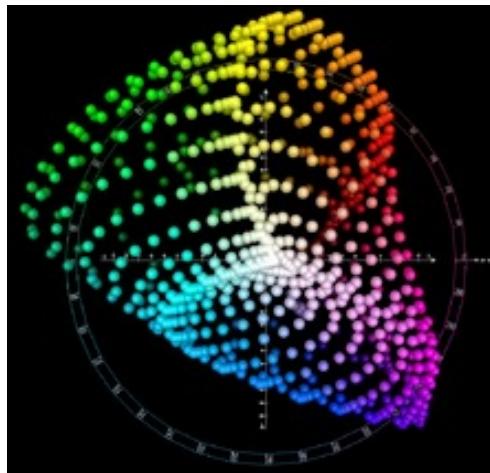
Another Option: Lab

Pros

1. Distances correspond with human judgment

Cons

1. Complex to calculate



L
(a=0,b=0)



a
(L=65,b=0)



b
(L=65,a=0)

Lab



Photo credit: J. Hays

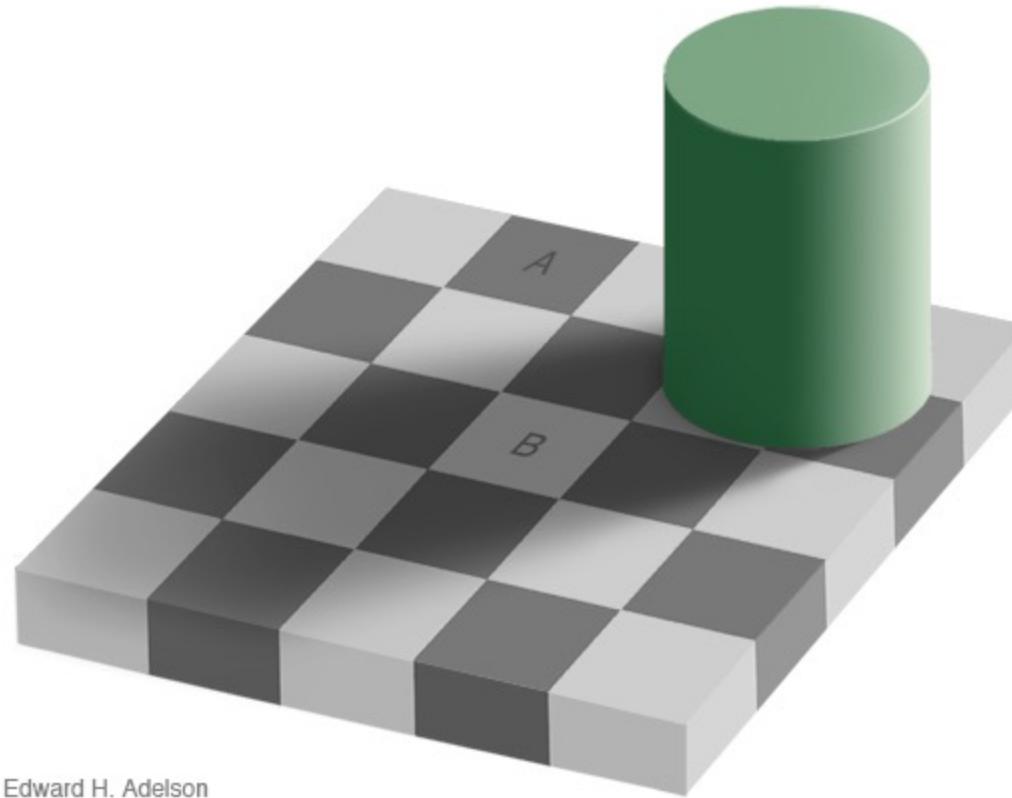
Why Are There So Many?

- Each serves different functions
 - RGB: sort of intuitive, standard, everywhere
 - HSV: good for picking, fast to compute
 - YCbCr/YUV: fast to compute, compresses well
 - Lab: the right(?) thing to do, but “slow” to compute
- Pick based on what you need and don’t sweat it: color really isn’t crucial

Color perception

- Color/lightness constancy
 - The ability of the human visual system to perceive the intrinsic reflectance properties of the surfaces despite changes in illumination conditions

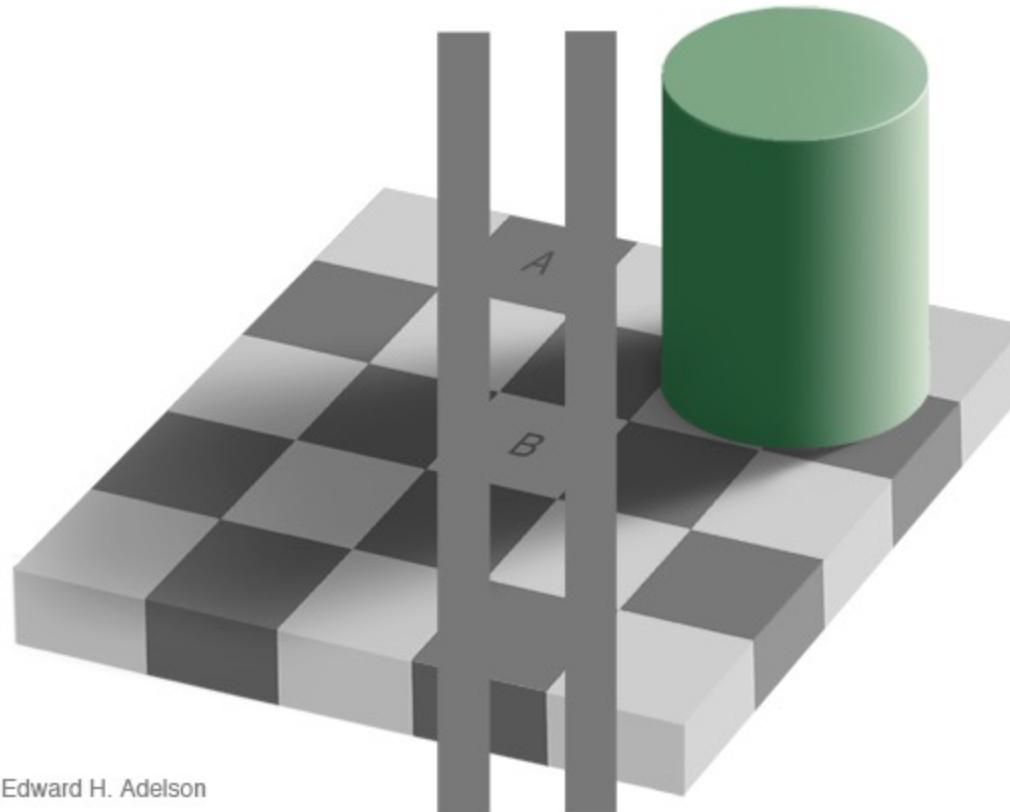
Checker shadow illusion



Edward H. Adelson

https://en.wikipedia.org/wiki/Checker_shadow_illusion

Checker shadow illusion



- Possible explanations
 - Simultaneous contrast
 - Reflectance edges vs. illumination edges

https://en.wikipedia.org/wiki/Checker_shadow_illusion

What color is the dress?



<https://www.wired.com/2015/02/science-one-agrees-color-dress/>

This strawberry cake has no red pixels!



<https://www.digitaltrends.com/photography/non-red-strawberries/>

White balance

- Analogous to color constancy mechanisms in human vision, cameras have mechanisms to adapt to the illumination in the environment so that neutral (white or gray) objects look neutral

Incorrect white balance

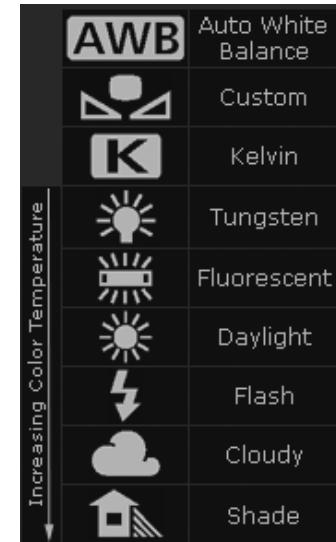


Correct white balance



White balance

- Film cameras:
 - Different types of film or different filters for different illumination conditions
- Digital cameras:
 - Automatic white balance
 - White balance settings corresponding to several common illuminants
 - Custom white balance using a reference object



White balance

- Von Kries adaptation: Multiply each channel by a *gain factor*
- Best way: gray card
 - Take a picture of a neutral object (white or gray)
 - If the object is recorded as r_w , g_w , b_w use weights $1/r_w$, $1/g_w$, $1/b_w$



White balance

- Without gray cards: we need to “guess” which pixels correspond to white objects
- Gray world assumption
 - The image average r_{ave} , g_{ave} , b_{ave} is gray
 - Use weights $1/r_{ave}$, $1/g_{ave}$, $1/b_{ave}$
- Brightest pixel assumption
 - Highlights usually have the color of the light source
 - Use weights inversely proportional to the values of the brightest pixels
- Gamut mapping
 - Gamut: convex hull of all pixel colors in an image
 - Find the transformation that matches the gamut of the image to the gamut of a “typical” image under white light
- Use image statistics, learning techniques

Mixed illumination

- When there are several types of illuminants in the scene, different reference points will yield different results



Reference: moon



Reference: stone

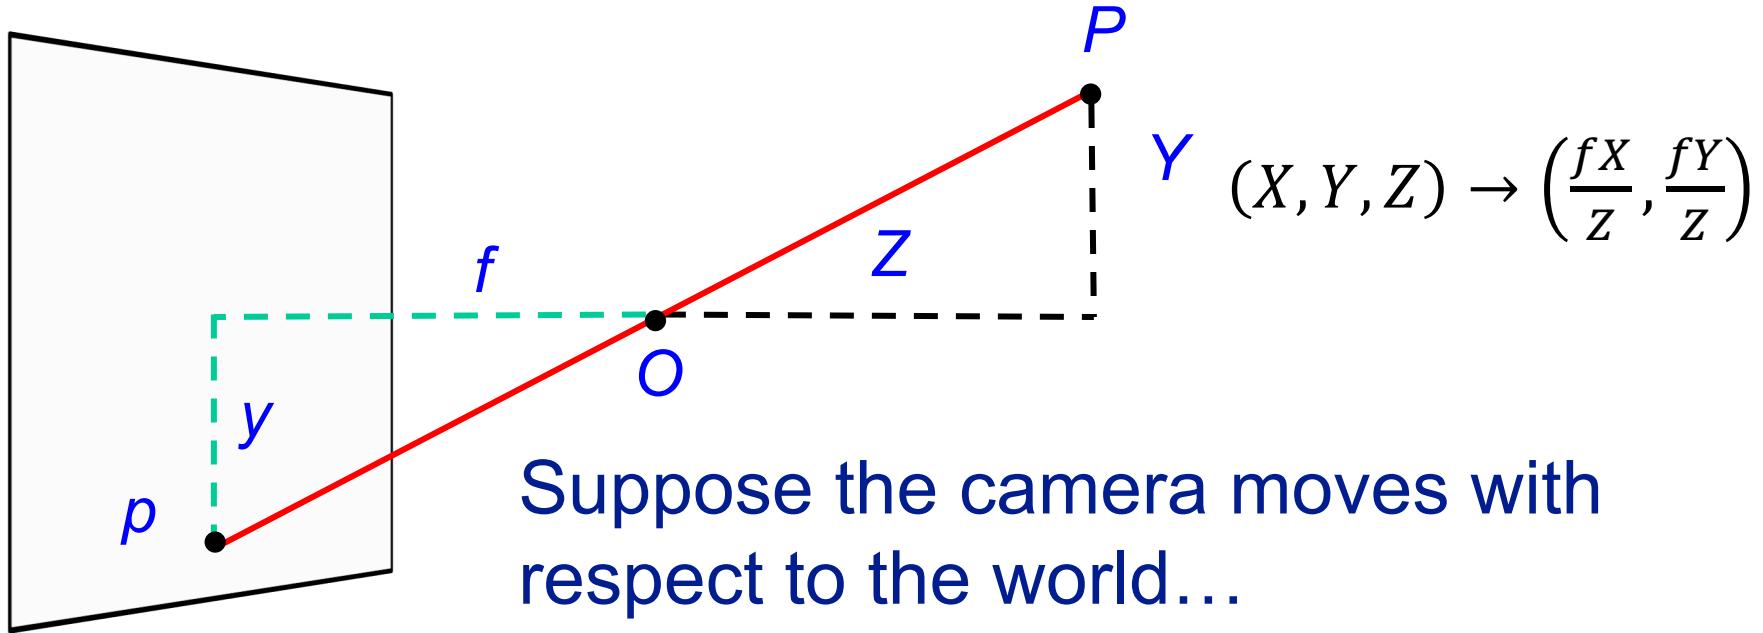
Outline

- Physical origin of color
- Spectra of sources and surfaces
- Physiology of color vision
- Trichromatic color theory
- Color spaces
- Color constancy, white balance

Dynamic Perspective

CS 543 / ECE 549 – Saurabh Gupta

Perspective Projection



- Point $P (X, Y, Z)$ in the world moves relative to the camera, its projection in the image (x, y) moves as well.
- This movement in the image plane is called **optical flow**.
- Suppose the image of the point (x, y) moves to $(x + \Delta x, y + \Delta y)$ in time Δt , then $\left(\frac{\Delta x}{\Delta t}, \frac{\Delta y}{\Delta t}\right)$ are the two components of the optical flow.

Outline

- Relate optical flow to camera motion
- Special cases

How does a point X in the scene move?

- Assume that the camera moves with a translational velocity $t = (t_x, t_y, t_z)$ and angular velocity $\omega = (\omega_x, \omega_y, \omega_z)$.
- Linear velocity of point $P = (X, Y, Z)$ is given by $\dot{P} = -t - \omega \times P$.

$$\begin{bmatrix} \dot{X} \\ \dot{Y} \\ \dot{Z} \end{bmatrix} = - \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} - \begin{bmatrix} \omega_y Z - \omega_z Y \\ \omega_z X - \omega_x Z \\ \omega_x Y - \omega_y X \end{bmatrix}$$

Now, lets consider the effect of projection

$$\begin{bmatrix} \dot{X} \\ \dot{Y} \\ \dot{Z} \end{bmatrix} = - \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} - \begin{bmatrix} \omega_y Z - \omega_z Y \\ \omega_z X - \omega_x Z \\ \omega_x Y - \omega_y X \end{bmatrix}$$

- Assume, $f = 1$, $x = \frac{X}{Z}$, $y = \frac{Y}{Z}$.
- $\dot{x} = \frac{\dot{X}Z - \dot{Z}X}{Z^2}$, $\dot{y} = \frac{\dot{Y}Z - \dot{Z}Y}{Z^2}$
- Substitute $\dot{X}, \dot{Y}, \dot{Z}$, from equation above:

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \dot{y} \end{bmatrix} = \frac{1}{Z} \begin{bmatrix} -1 & 0 & x \\ 0 & -1 & y \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} xy & -(1+x^2) & y \\ (1+y^2) & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

Dynamic Perspective Equations

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \dot{y} \end{bmatrix} = \frac{1}{Z} \begin{bmatrix} -1 & 0 & x \\ 0 & -1 & y \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} xy & -(1+x^2) & y \\ (1+y^2) & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

Translation Component

Rotation Component

Optical flow for pure rotation

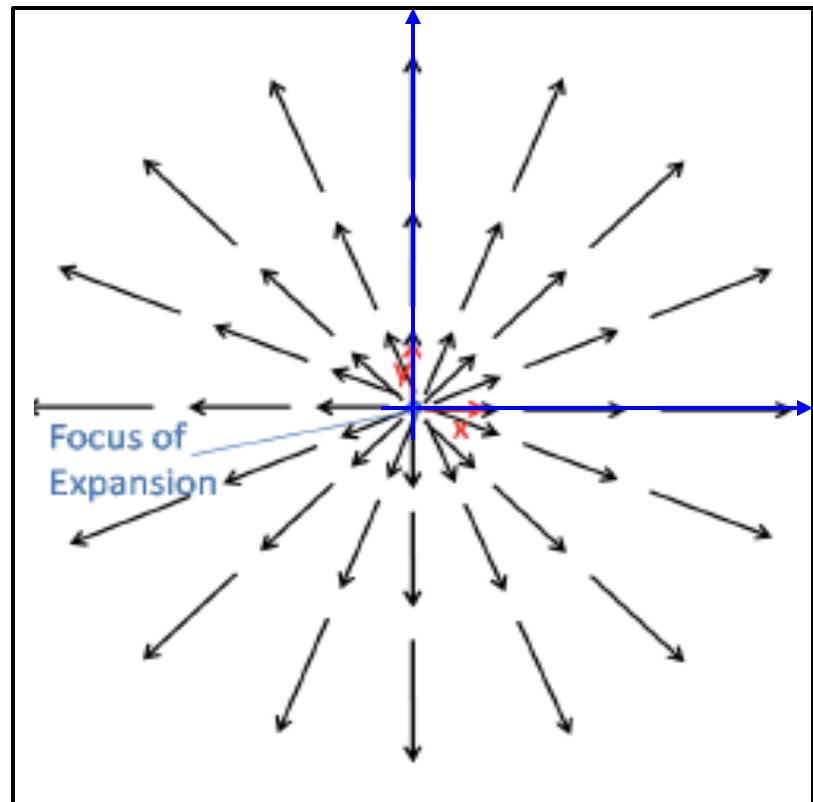
$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \dot{y} \end{bmatrix} = \frac{1}{Z} \begin{bmatrix} -1 & 0 & x \\ 0 & -1 & y \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} xy & -(1+x^2) & y \\ (1+y^2) & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

- $\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} xy & -(1+x^2) & y \\ (1+y^2) & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$
- We can determine ω from the flow field.
- Flow field is independent of $Z(x, y)$.

Optical flow for pure translation along Z-axis

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \dot{y} \end{bmatrix} = \frac{1}{Z} \begin{bmatrix} -1 & 0 & x \\ 0 & -1 & y \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} xy & -(1+x^2) & y \\ (1+y^2) & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

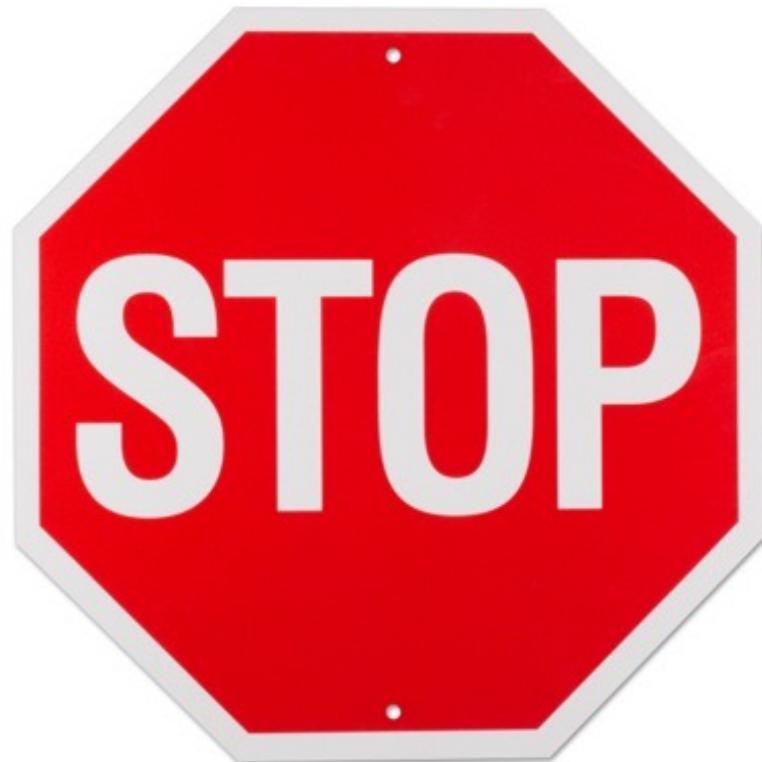
- $\begin{bmatrix} u \\ v \end{bmatrix} = \frac{t_z}{z(x,y)} \begin{bmatrix} x \\ y \end{bmatrix}$
- Optical flow vector is a scalar multiple of position vector.
- Scale factor ambiguity, if $t_z \rightarrow kt_z$, and $Z \rightarrow kZ$, optical flow remains unchanged.
- But, you can get time to collision, Z/t_z .







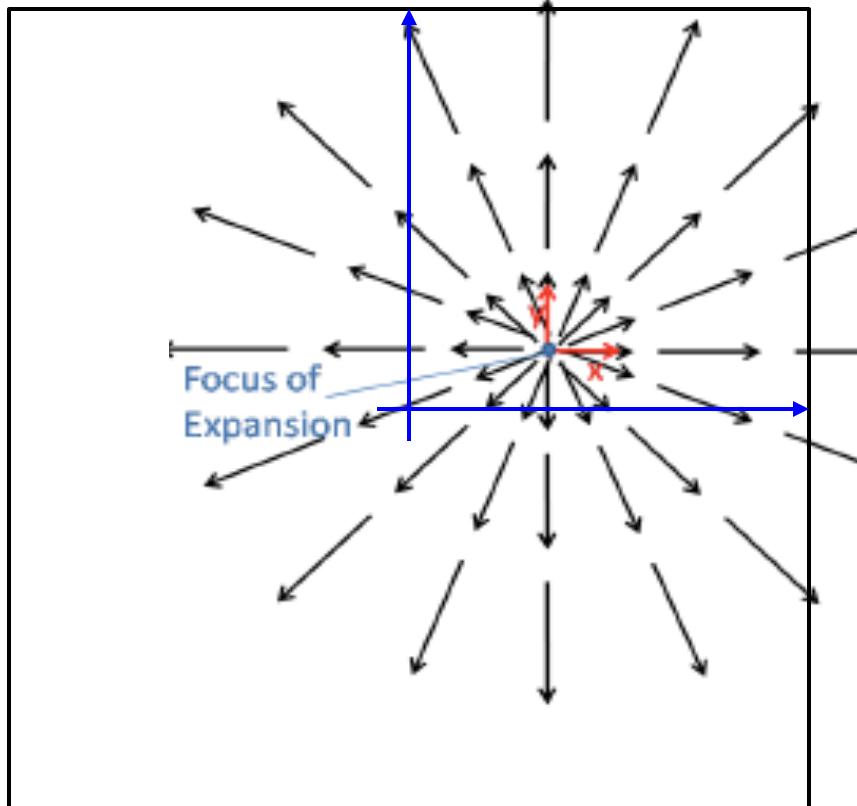




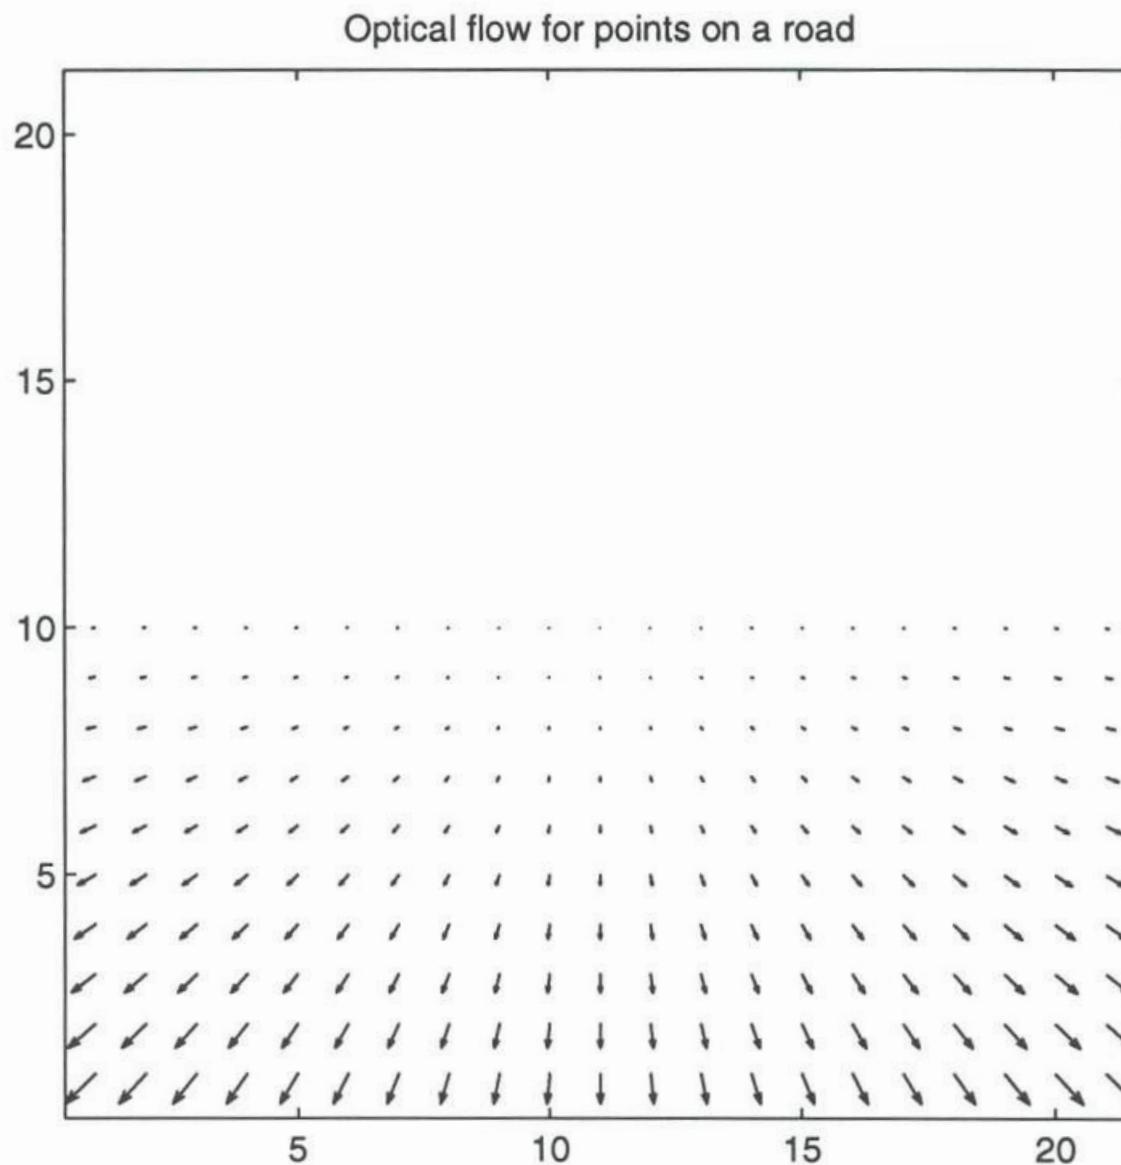
Optical flow for impure translation along Z-axis

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \dot{y} \end{bmatrix} = \frac{1}{Z} \begin{bmatrix} -1 & 0 & x \\ 0 & -1 & y \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} xy & -(1+x^2) & y \\ (1+y^2) & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

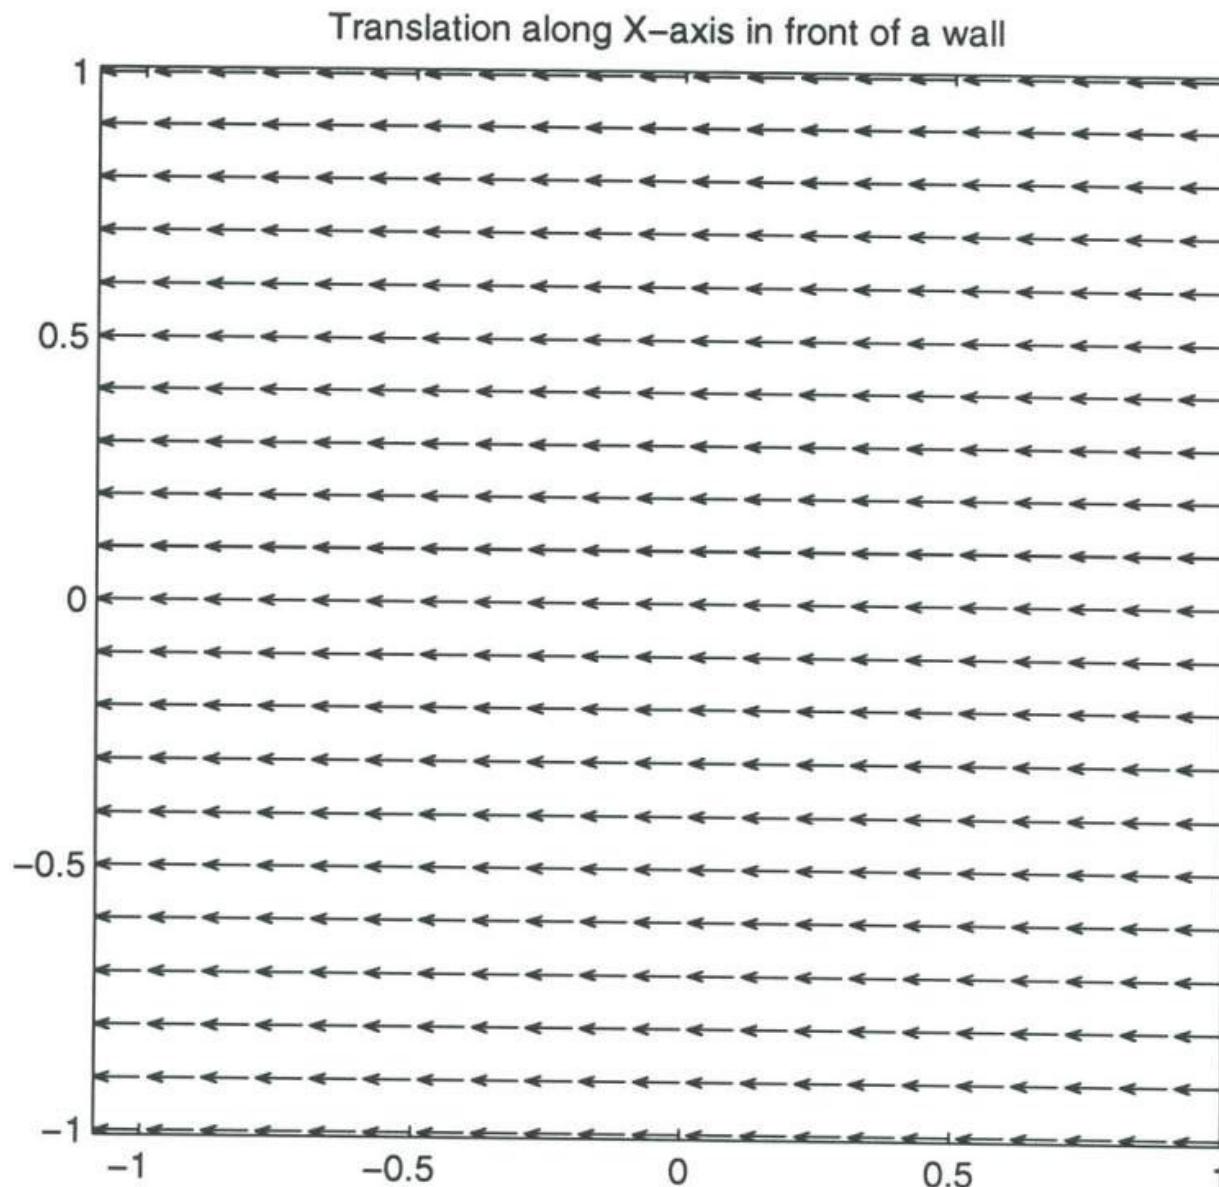
- $u = \frac{-t_x + xt_z}{Z(x,y)}$, $v = \frac{-t_y + yt_z}{Z(x,y)}$



Optical flow for points on a road



Translating along X-axis in front of a wall



Estimating Optical Flow from Images

Aperture Problem

Recap

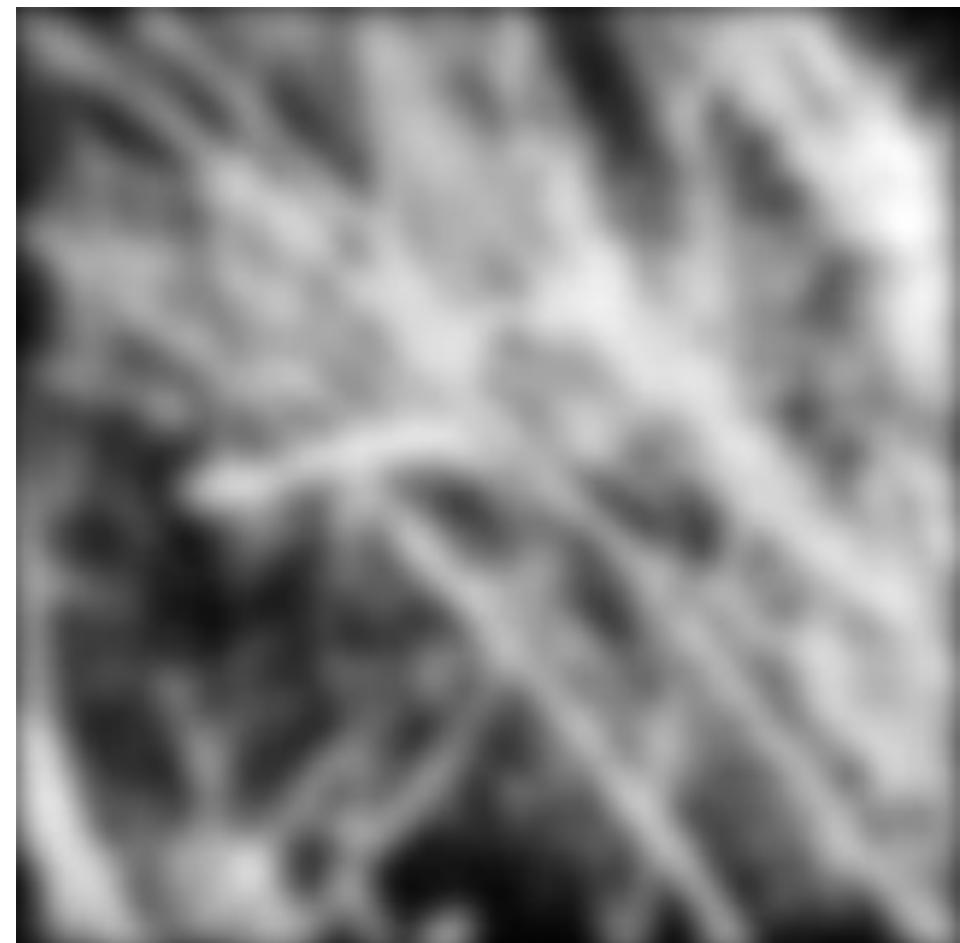
- Relate optical flow to camera motion

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \dot{x} \\ \dot{y} \end{bmatrix} = \frac{1}{z} \begin{bmatrix} -1 & 0 & x \\ 0 & -1 & y \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} xy & -(1+x^2) & y \\ (1+y^2) & -xy & -x \end{bmatrix} \begin{bmatrix} \omega_x \\ \omega_y \\ \omega_z \end{bmatrix}$$

- Special cases

- Pure rotation / pure translation / time to collision

Linear filtering



Motivation: Image denoising

- How can we reduce noise in a photograph?



Moving average

- Let's replace each pixel with a *weighted average* of its neighborhood
- The weights are called the *filter kernel*
- What are the weights for the average of a 3x3 neighborhood?

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

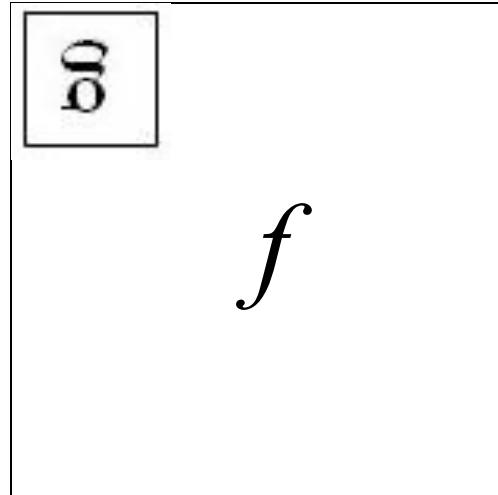
“box filter”

Defining convolution

- Let f be the image and g be the kernel. The output of convolving f with g is denoted $f * g$.

$$(f * g)[m, n] = \sum_{k,l} f[m-k, n-l]g[k, l]$$

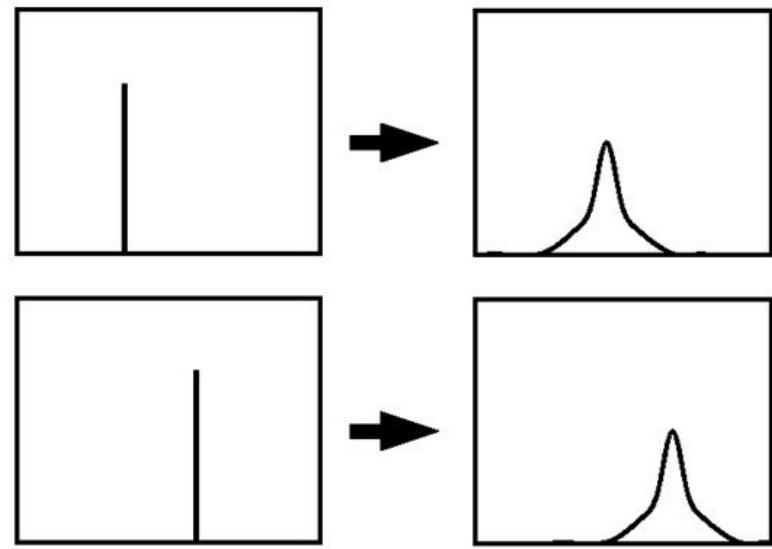
Convention:
kernel is “flipped”



Key properties

- **Shift invariance:** same behavior regardless of pixel location:

$$\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$$



- **Linearity:**

$$\begin{aligned}\text{filter}(f_1 + f_2) &= \\ \text{filter}(f_1) + \text{filter}(f_2)\end{aligned}$$

- Theoretical result: any linear shift-invariant operator can be represented as a convolution

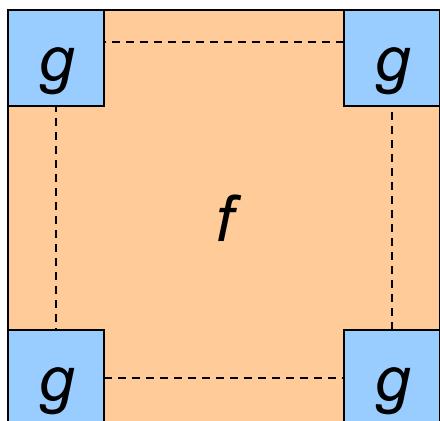
Properties in more detail

- Commutative: $a * b = b * a$
 - Conceptually no difference between filter and signal
- Associative: $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out: $ka * b = a * kb = k(a * b)$
- Identity: unit impulse $e = [..., 0, 0, 1, 0, 0, ...]$,
 $a * e = a$

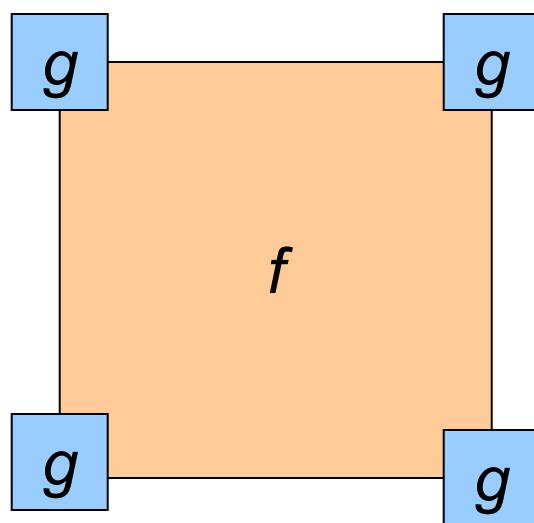
Dealing with edges

- If we convolve image f with filter g , what is the size of the output?

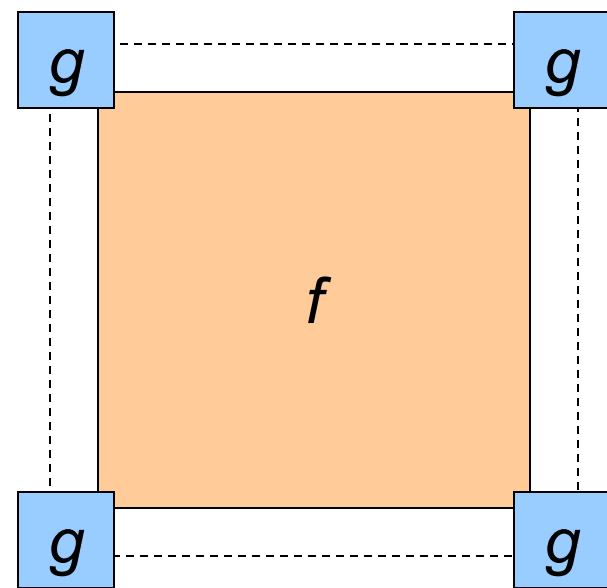
Output is smaller than input



Output is same size as input

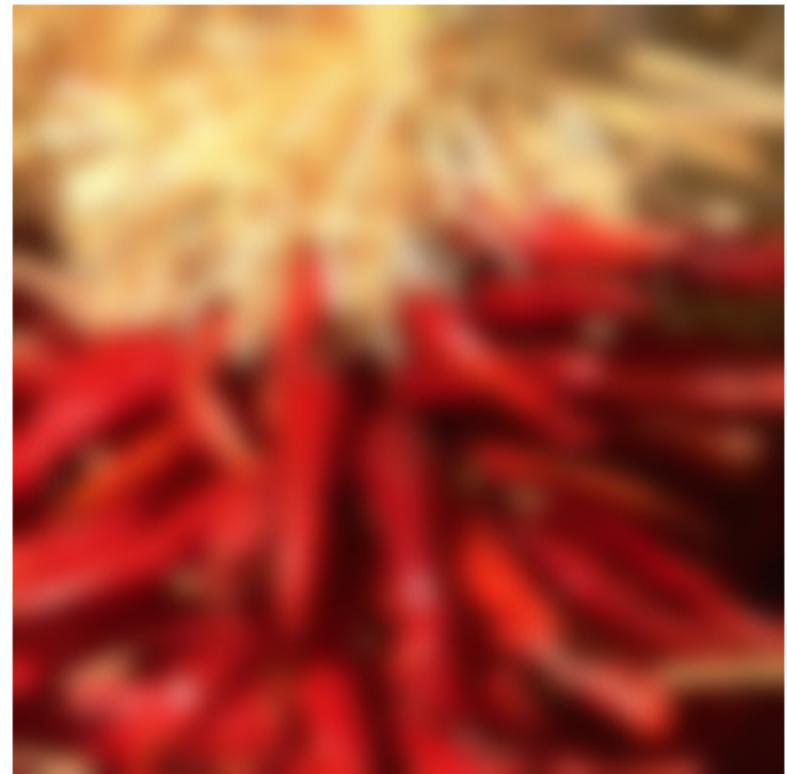


Output is larger than input

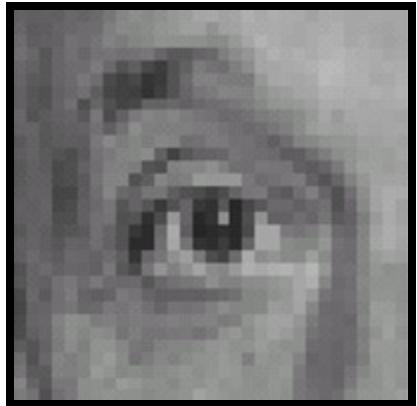


Dealing with edges

- If the filter window falls off the edge of the image, we need to pad the image
 - Zero pad (or clip filter)
 - Wrap around
 - Copy edge
 - Reflect across edge



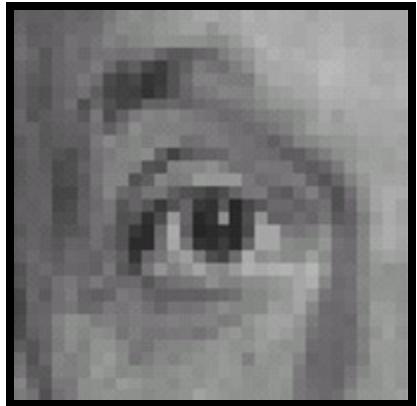
Practice with linear filters



0	0	0
0	1	0
0	0	0

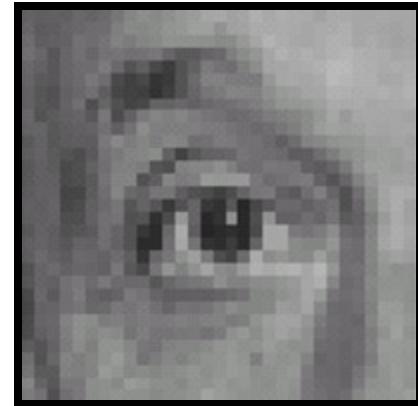
?

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters



0	0	0
0	0	1
0	0	0

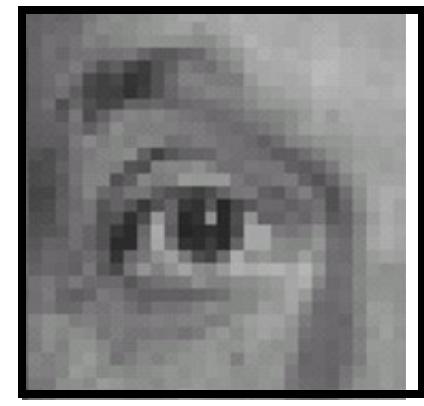
?

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted *left*
By 1 pixel

Practice with linear filters

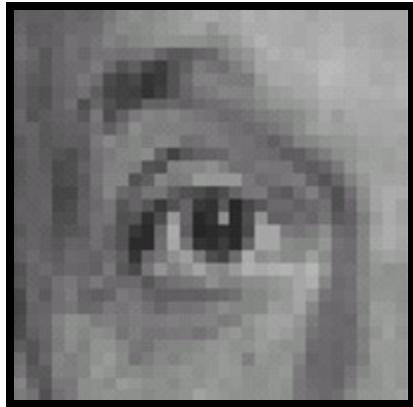


$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

?

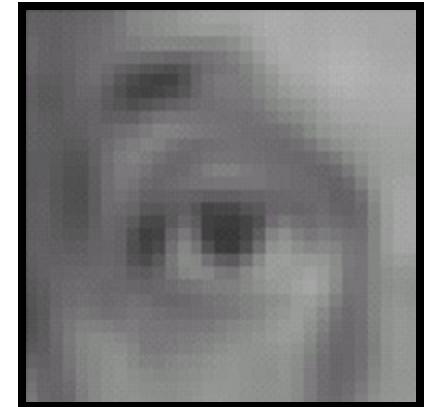
Original

Practice with linear filters



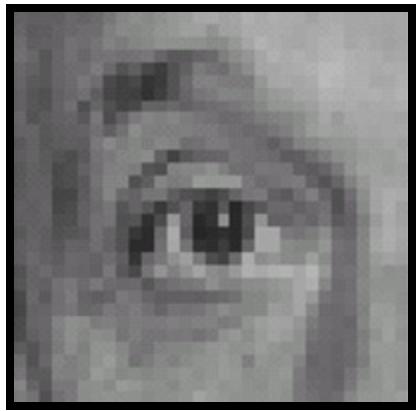
Original

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



Blur (with a
box filter)

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)

Practice with linear filters



Original

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

-

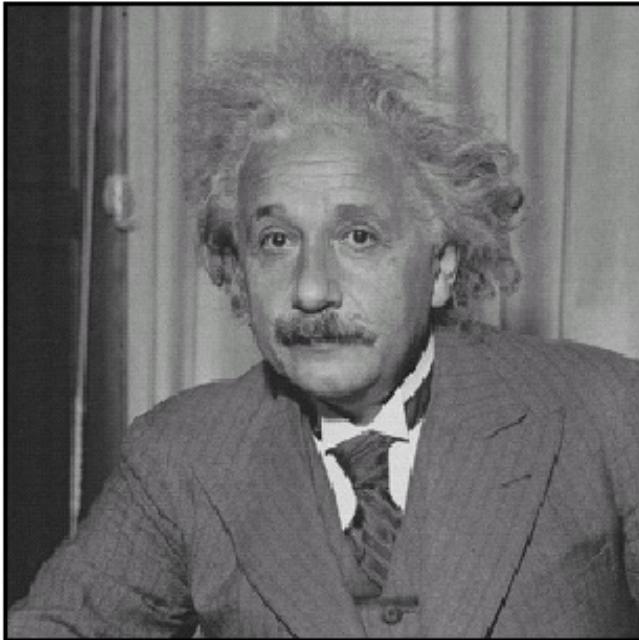
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



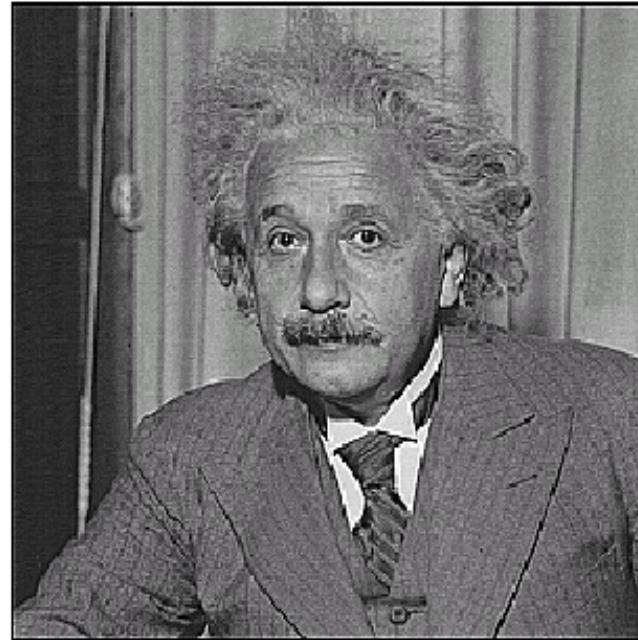
Sharpening filter

- Accentuates differences with local average

Sharpening



before



after

Sharpening

What does blurring take away?



Original



Smoothed



Detail

=

Let's add it back in.



Original



$+ \alpha$

Detail

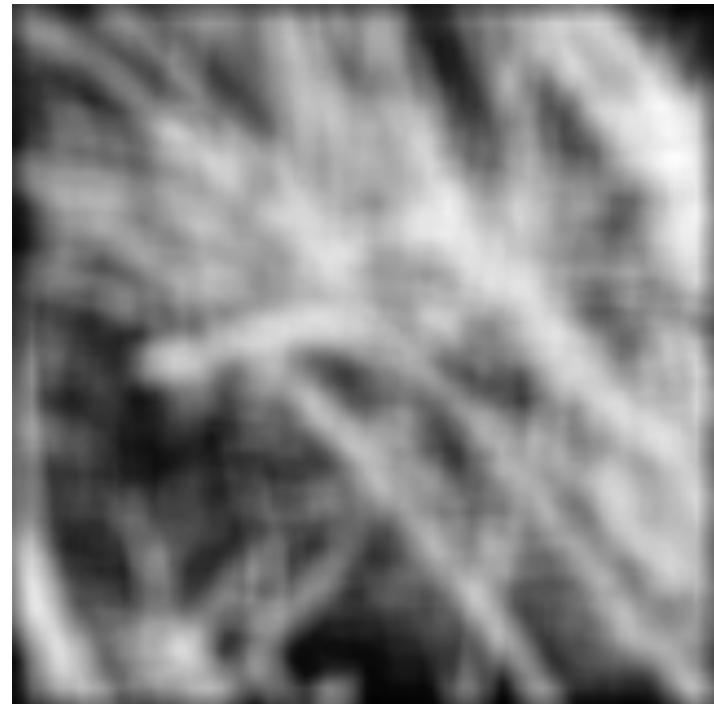
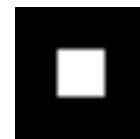


Sharpened

=

Smoothing with box filter revisited

- What's wrong with this picture?
- What's the solution?



Smoothing with box filter revisited

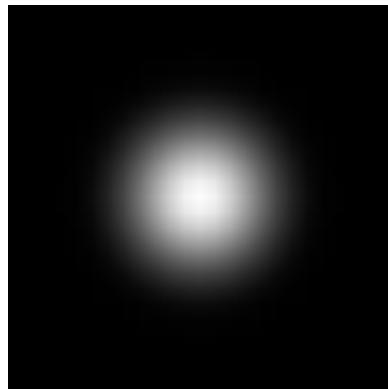
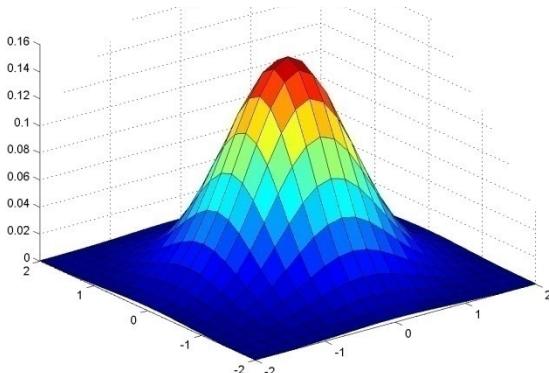
- What's wrong with this picture?
- What's the solution?
 - To eliminate edge effects, weight contribution of neighborhood pixels according to their closeness to the center



“fuzzy blob”

Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\left(\frac{x^2+y^2}{2\sigma^2}\right)}$$



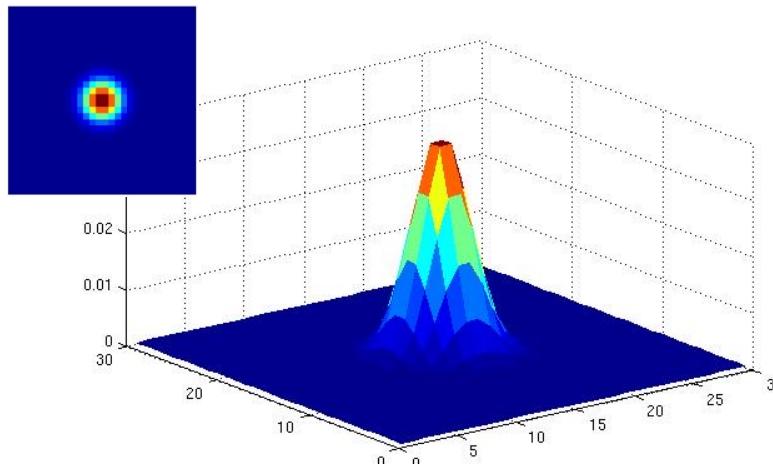
0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

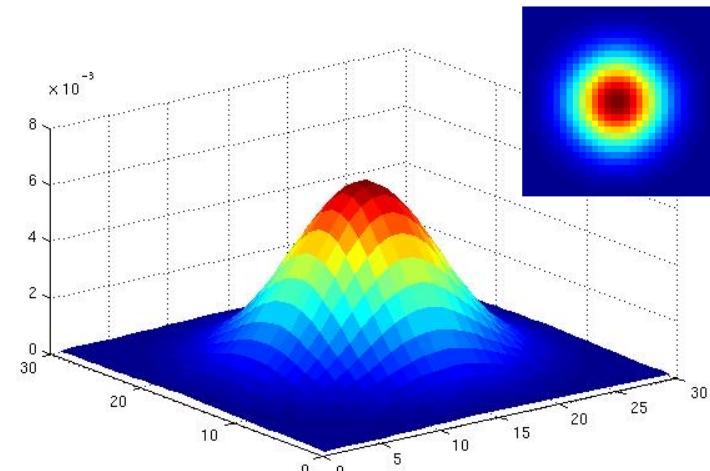
- Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case)

Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\left(\frac{x^2+y^2}{2\sigma^2}\right)}$$



$\sigma = 2$ with 30×30 kernel

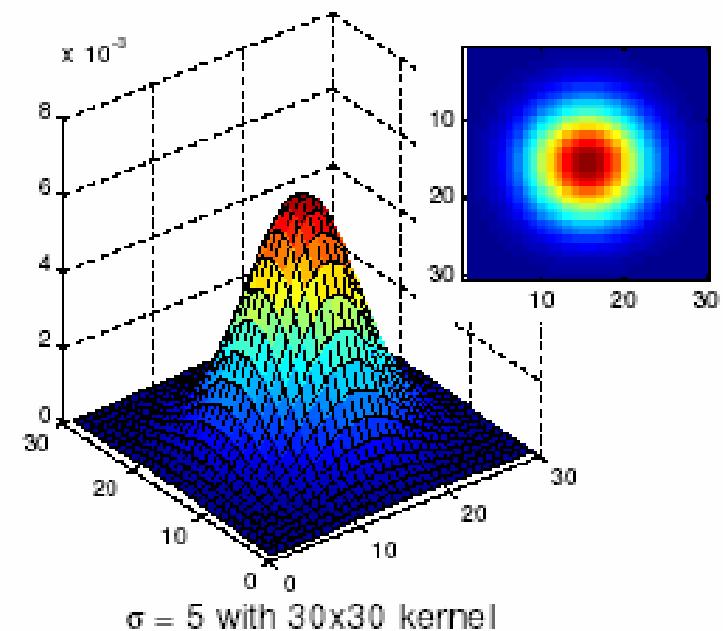
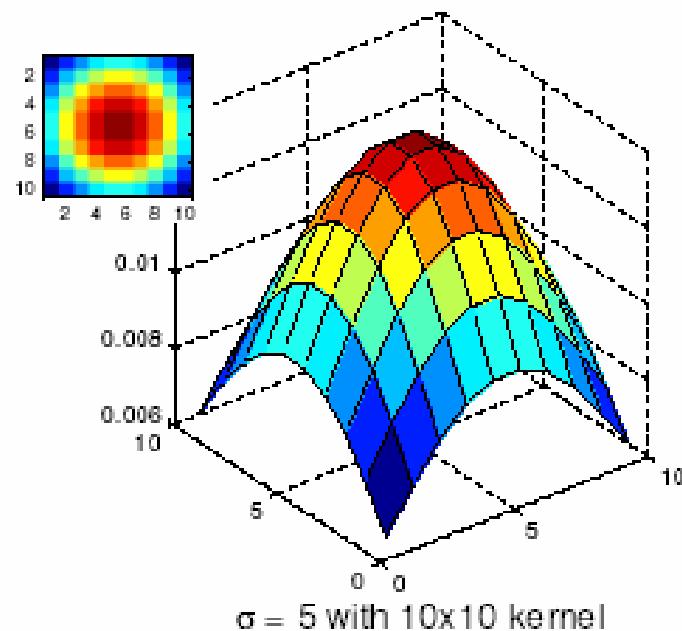


$\sigma = 5$ with 30×30 kernel

- Standard deviation σ : determines extent of smoothing

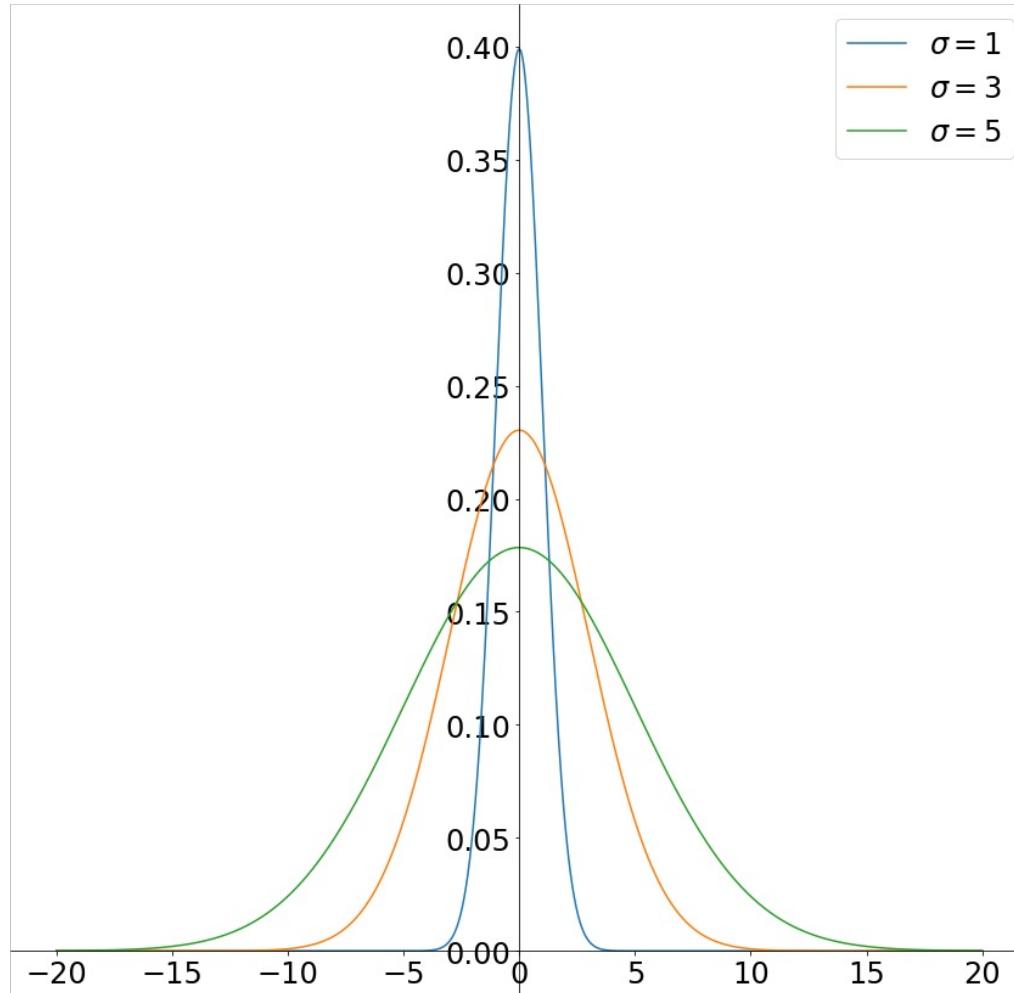
Choosing kernel width

- The Gaussian function has infinite support, but discrete filters use finite kernels

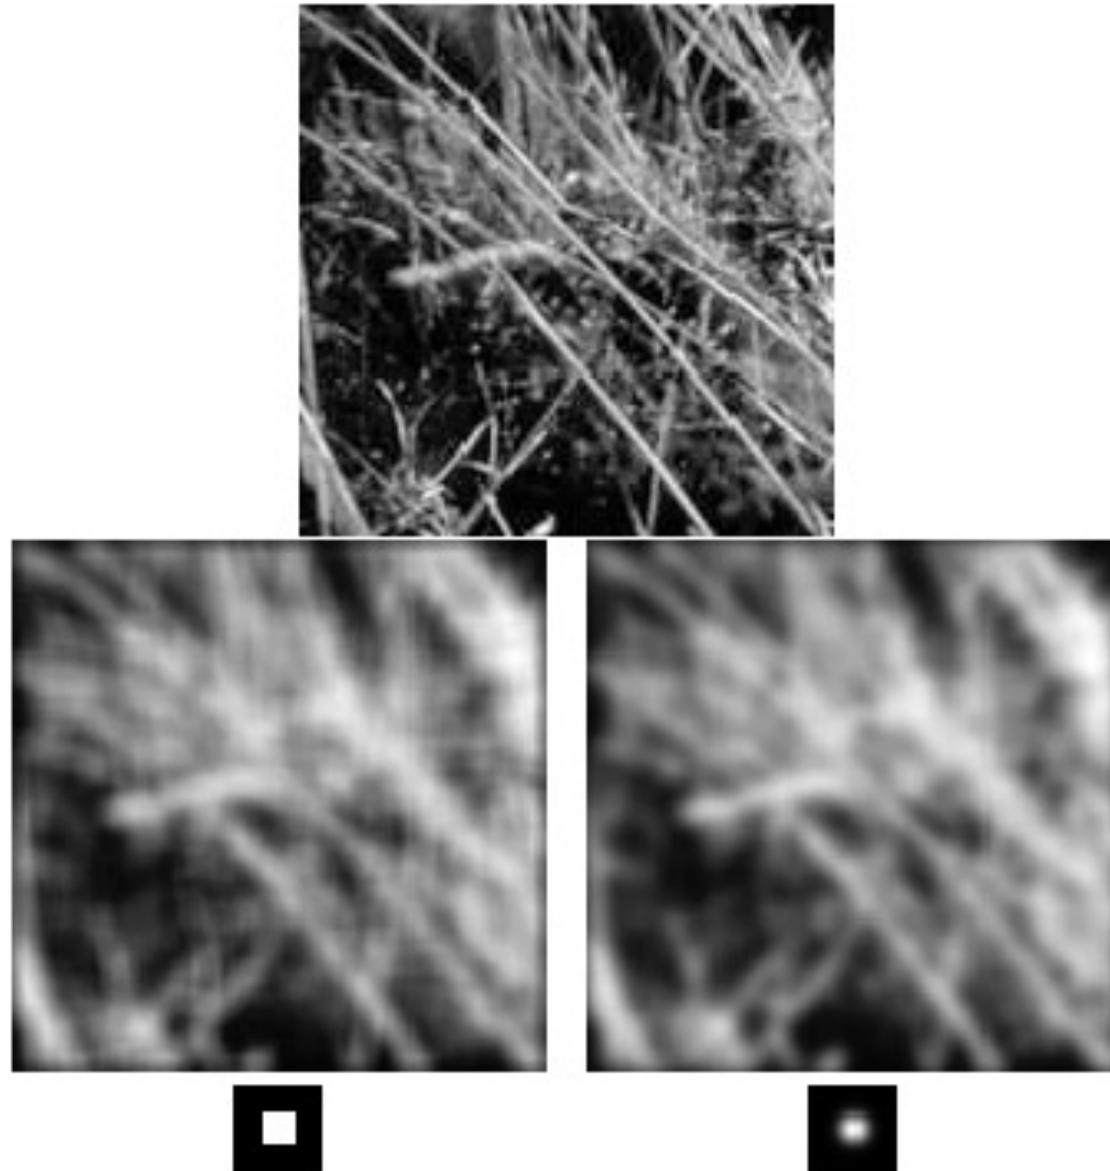


Choosing kernel width

- Rule of thumb: set filter half-width to about 3σ



Gaussian vs. box filtering



Gaussian filters

- Remove high-frequency components from the image (*low-pass filter*)
- Convolution with self is another Gaussian
 - So can smooth with small- σ kernel, repeat, and get same result as larger- σ kernel would have
 - Convolving two times with Gaussian kernel with std. dev. σ is same as convolving once with kernel with std. dev. $\sigma\sqrt{2}$
- **Separable kernel**
 - Factors into product of two 1D Gaussians
 - Discrete example:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} e^{-\left(\frac{x^2+y^2}{2\sigma^2}\right)} \\ &= \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

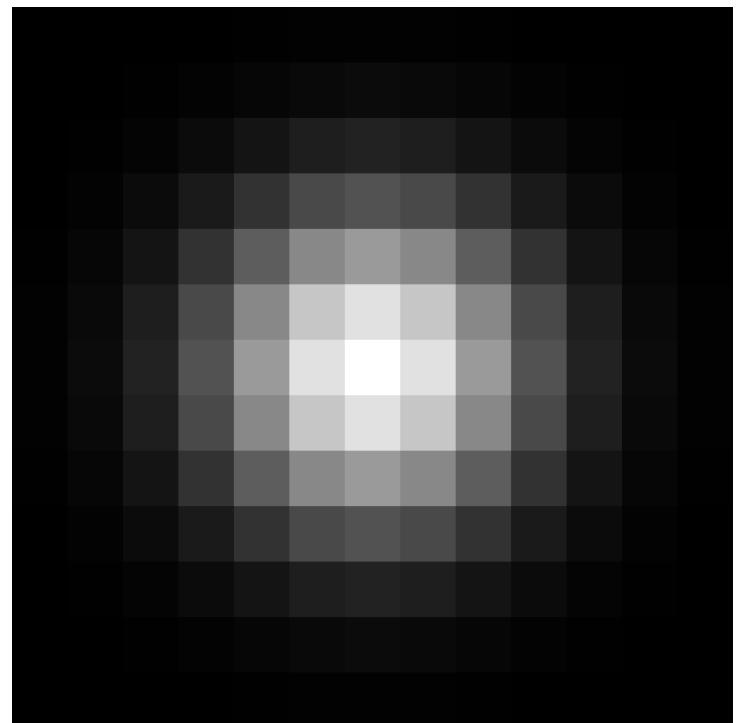
The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y.

In this case the two functions are the (identical) 1D Gaussian.

Separability

1D Gaussian * 1D Gaussian = 2D Gaussian

Image * 2D Gauss = Image * (1D Gauss * 1D Gauss)
= (Image * 1D Gauss) * 1D Gauss



Why is separability useful?

- Separability means that a 2D convolution can be reduced to two 1D convolutions (one along rows and one along columns)
- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
 - $O(n^2 m^2)$
- What if the kernel is separable?
 - $O(n^2 m)$

Noise



Original



Salt and pepper noise



Impulse noise

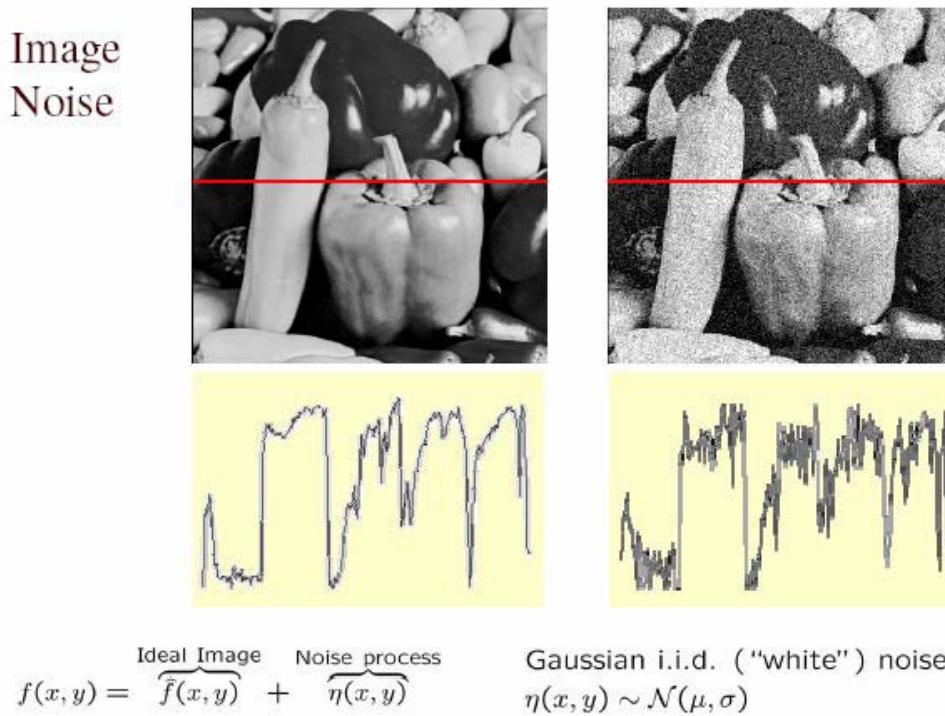


Gaussian noise

- **Salt and pepper noise:** contains random occurrences of black and white pixels
- **Impulse noise:** contains random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution

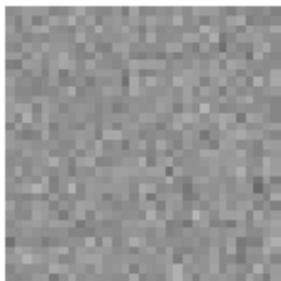
Gaussian noise

- Mathematical model: sum of many independent factors
- Good for small standard deviations
- Assumption: independent, zero-mean noise

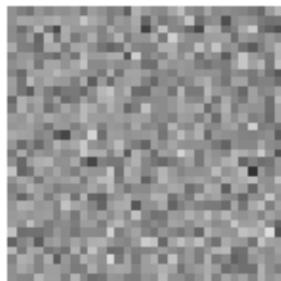


Reducing Gaussian noise

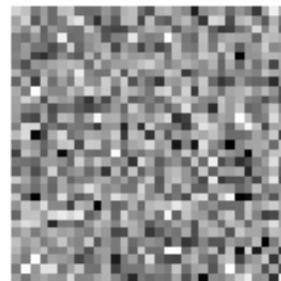
$\sigma=0.05$



$\sigma=0.1$



$\sigma=0.2$



no
smoothing



$\sigma=1$ pixel



$\sigma=2$ pixels

Smoothing with larger standard deviations suppresses noise,
but also blurs the image

Reducing salt-and-pepper noise

3x3



5x5



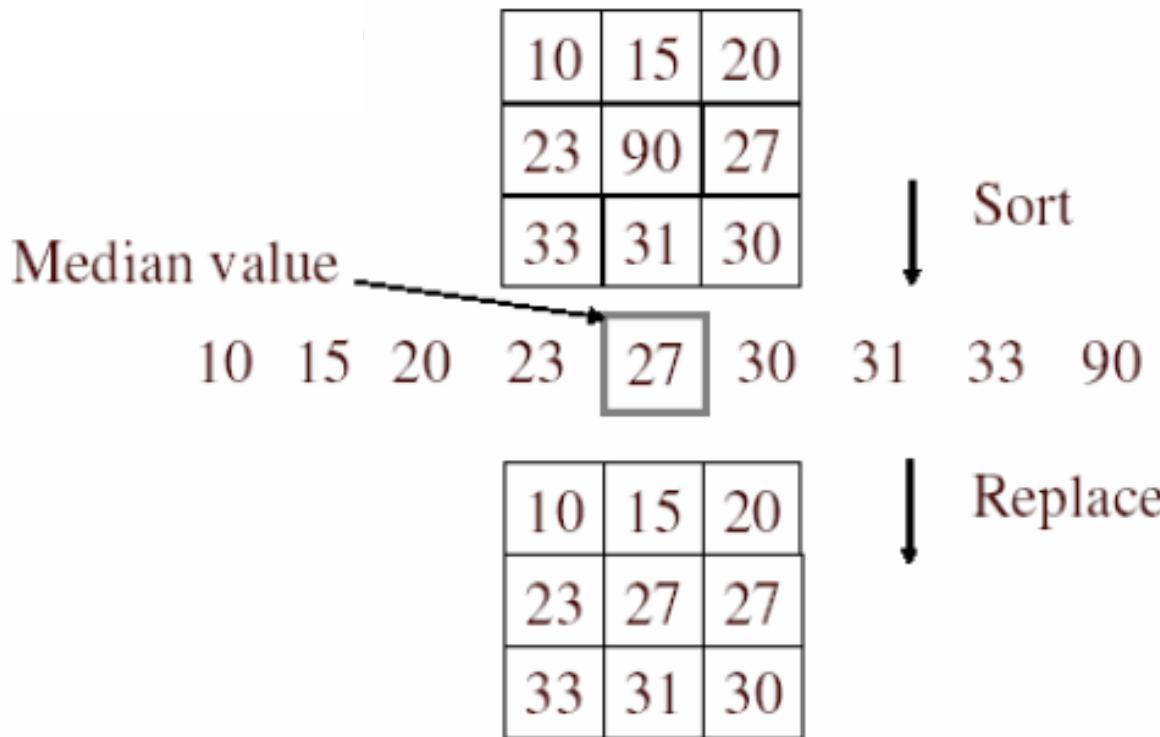
7x7



What's wrong with the results?

Alternative idea: Median filtering

- A **median filter** operates over a window by selecting the median intensity in the window



- Is median filtering linear?

Median filter

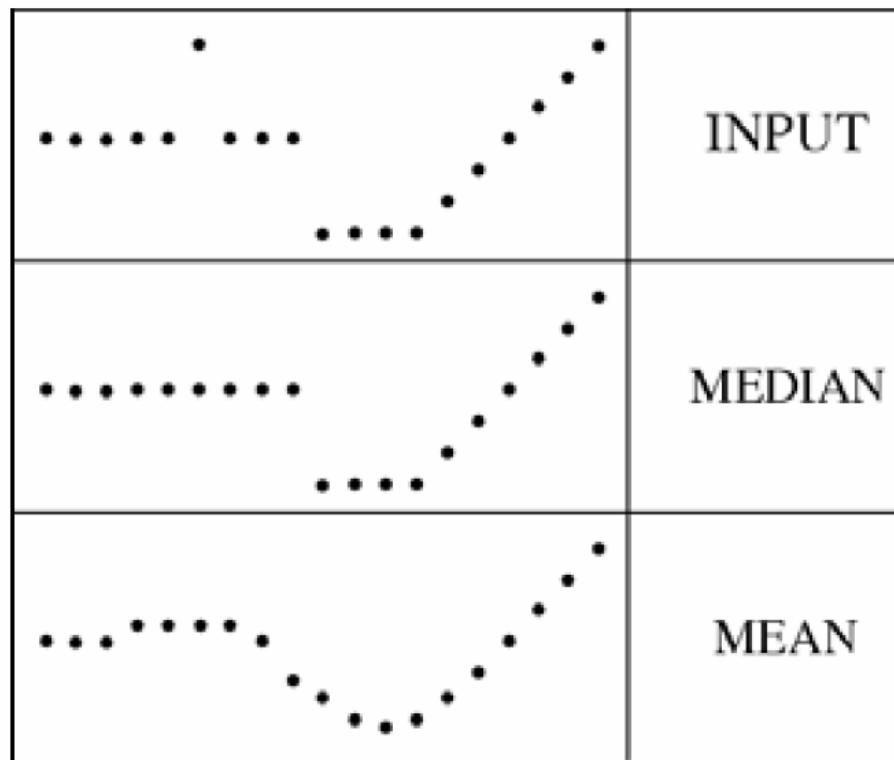
- Is median filtering linear?
- Let's try filtering

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 2 \\ 2 & 2 & 2 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Median filter

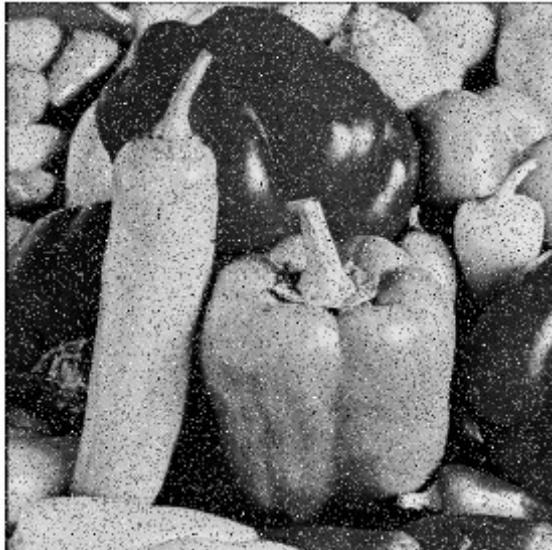
- What advantage does median filtering have over Gaussian filtering?
 - Robustness to outliers

filters have width 5 :

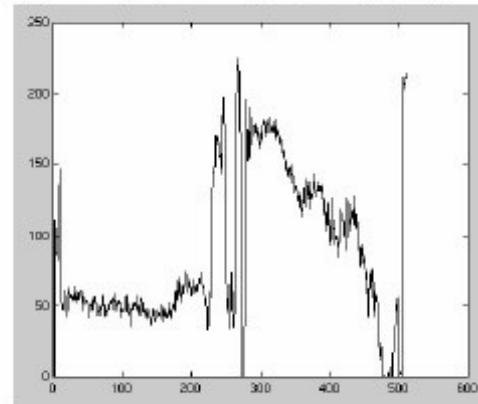
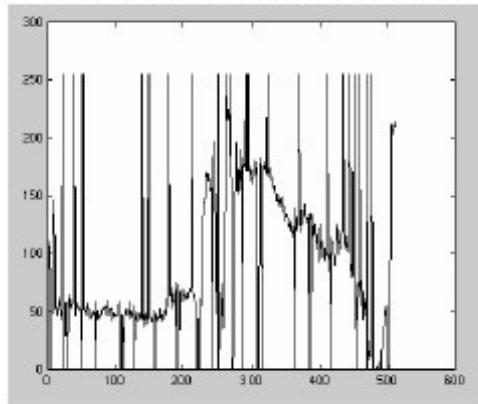


Median filter

Salt-and-pepper noise



Median filtered



Gaussian vs. median filtering

3x3



5x5



7x7

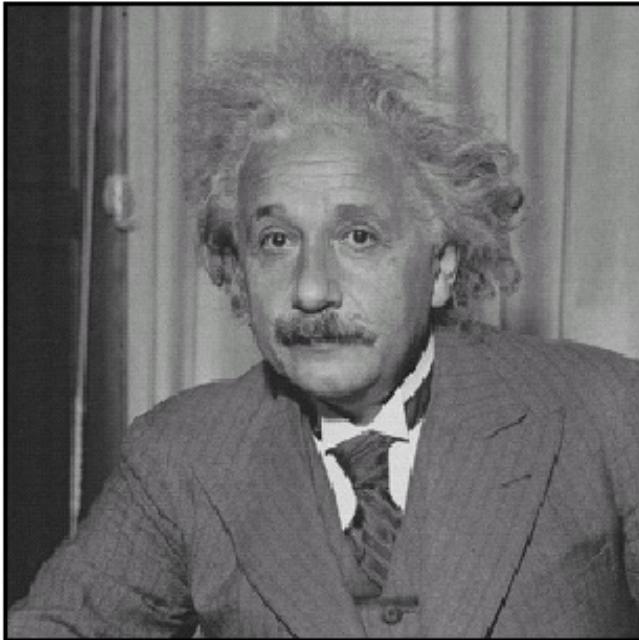


Gaussian

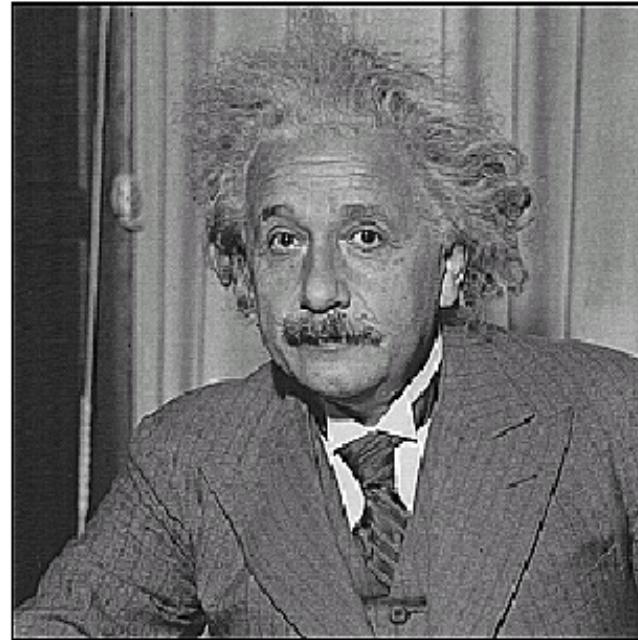
Median



Sharpening revisited



before



after

Sharpening

What does blurring take away?



Original



Smoothed



Detail

Let's add it back in.



Original



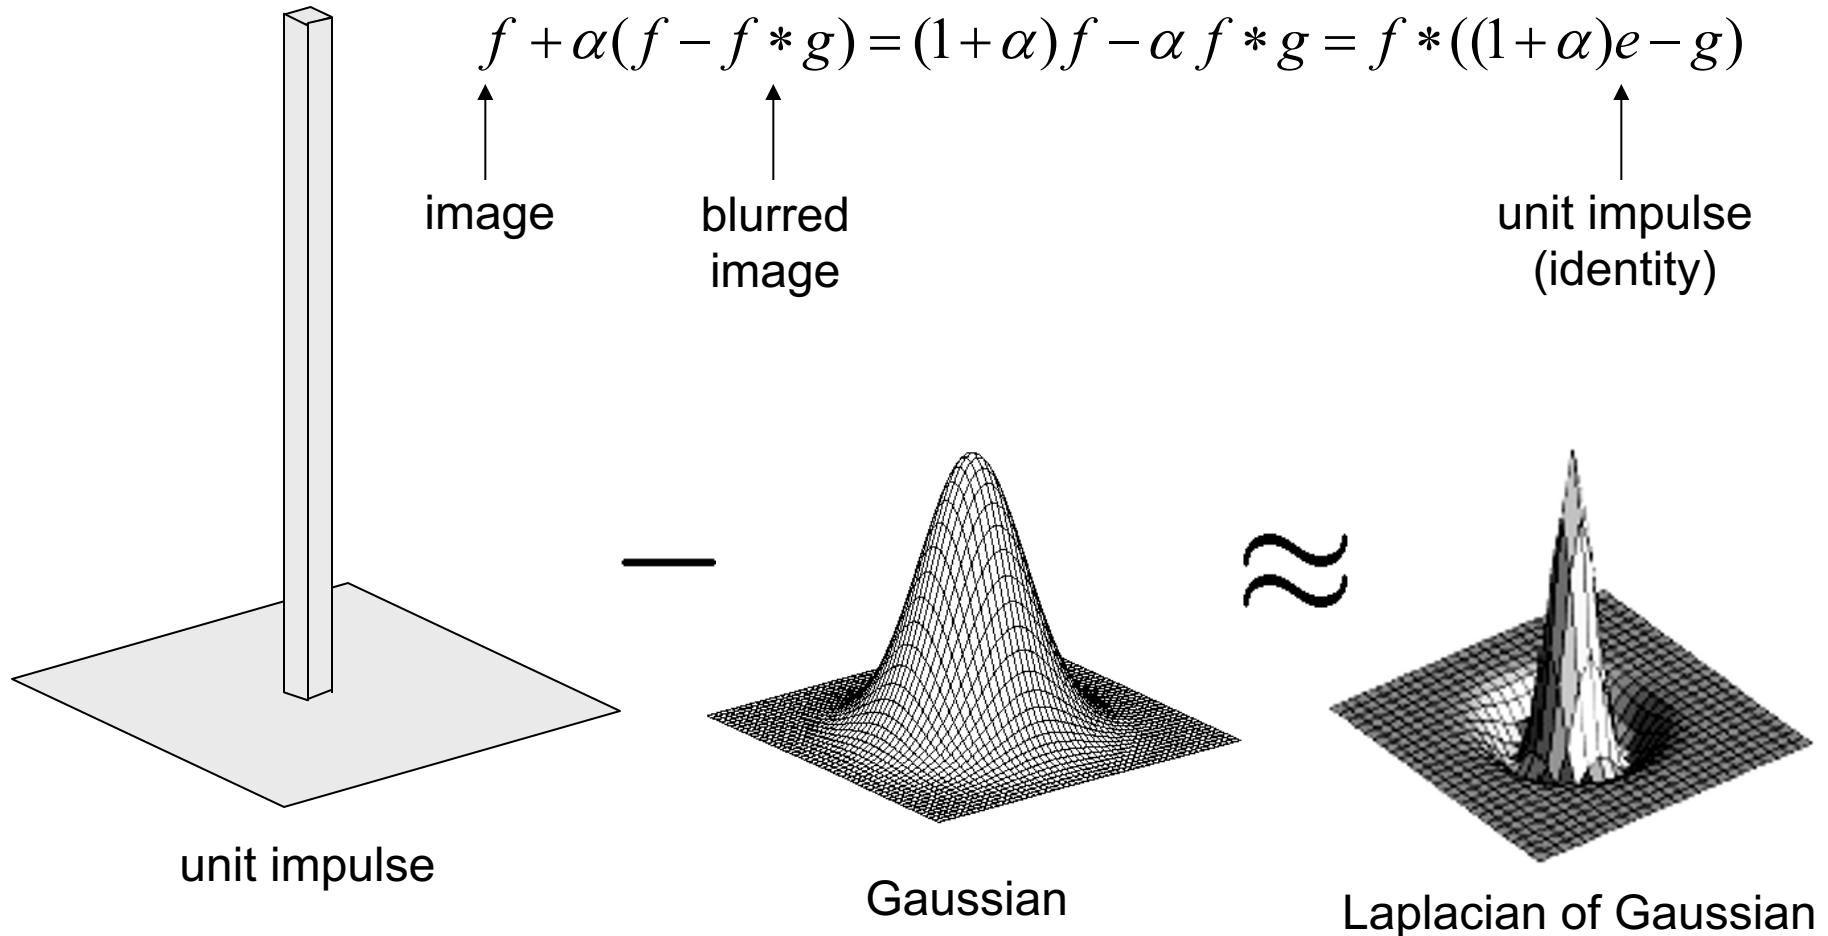
Detail



Sharpened

$$+\alpha$$

Unsharp mask filter



Next Class: Frequency view of filtering

Guess the filter



Guess the filter



Review: Image filtering

- Convolution
- Box vs. Gaussian filter
- Separability
- Median filter

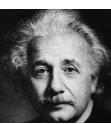
Filtering (in Frequency Domain)

while we wait,

Who do you see in
this picture?



A. Albert Einstein



B. Marilyn Monroe



C. None of these

Today's Class

- Fourier transforms
- Filtering in frequency domain
- Sampling
- Image Pyramids

Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

Gaussian



Box filter



Thinking in terms of frequency

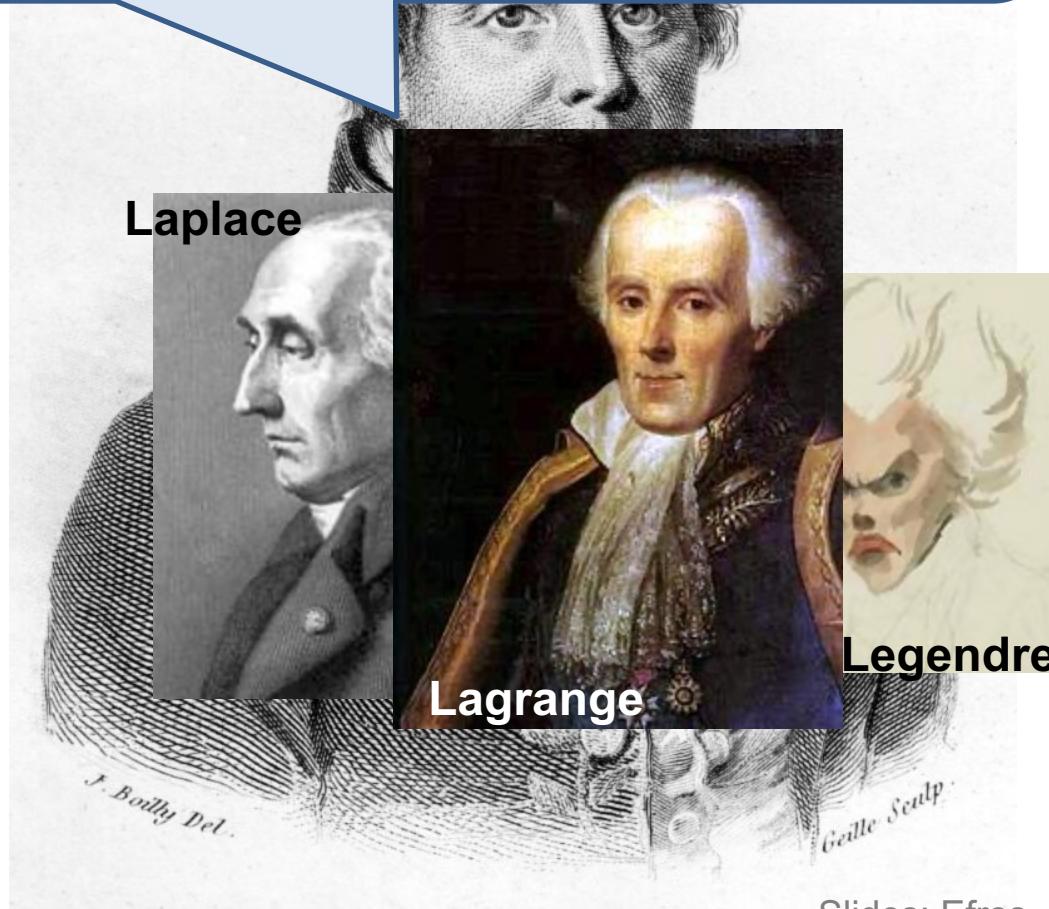
Jean Baptiste Joseph Fourier (1768-1830)

had crazy idea (1807):

Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.

...the manner in which the author arrives at these equations is not exempt of difficulties and...his analysis to integrate them still leaves something to be desired on the score of generality and even rigour.

- Don't believe it?
 - Neither did Lagrange, Laplace, Poisson and other big wigs
 - Not translated into English until 1878!
- But it's (mostly) true!
 - called Fourier Series
 - there are some subtle restrictions

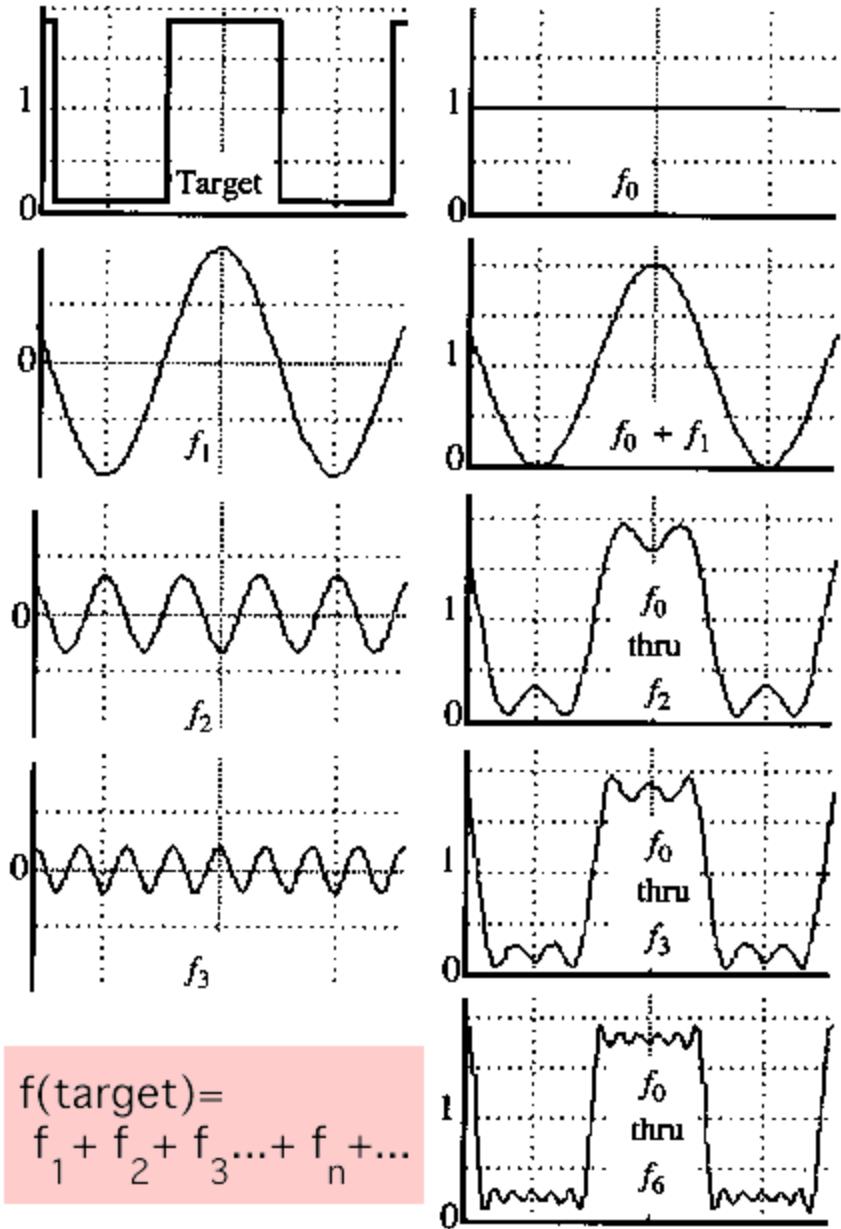


A sum of sines

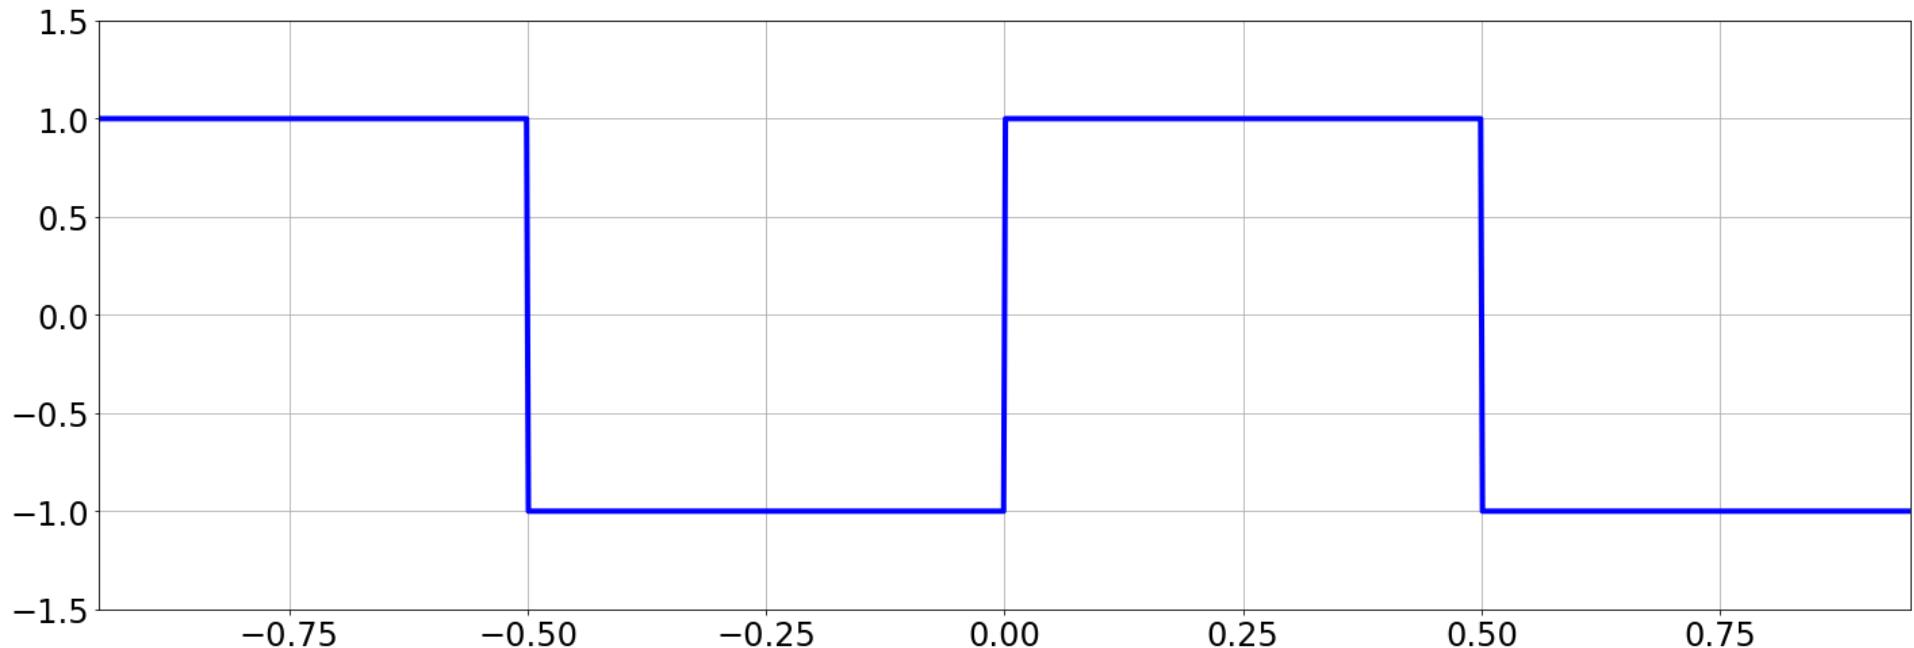
Our building block:

$$A \sin(\omega x + \phi)$$

Add enough of them to get any signal $f(x)$ you want!



Frequency Spectra

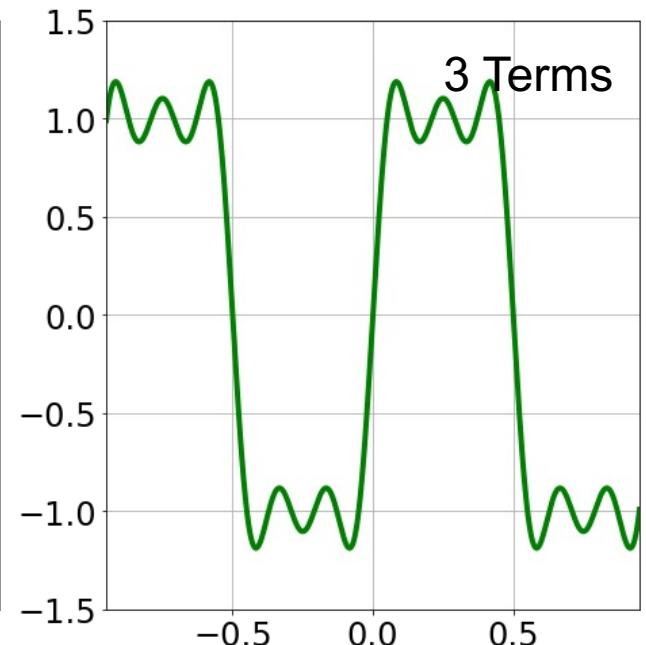
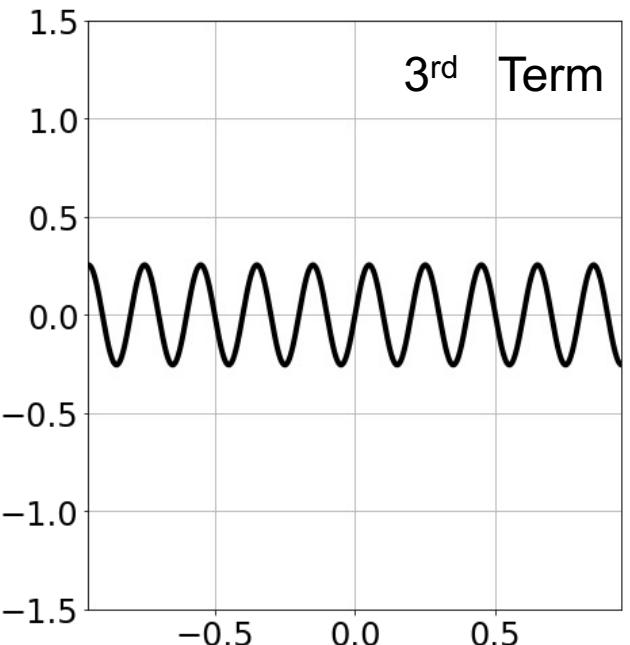
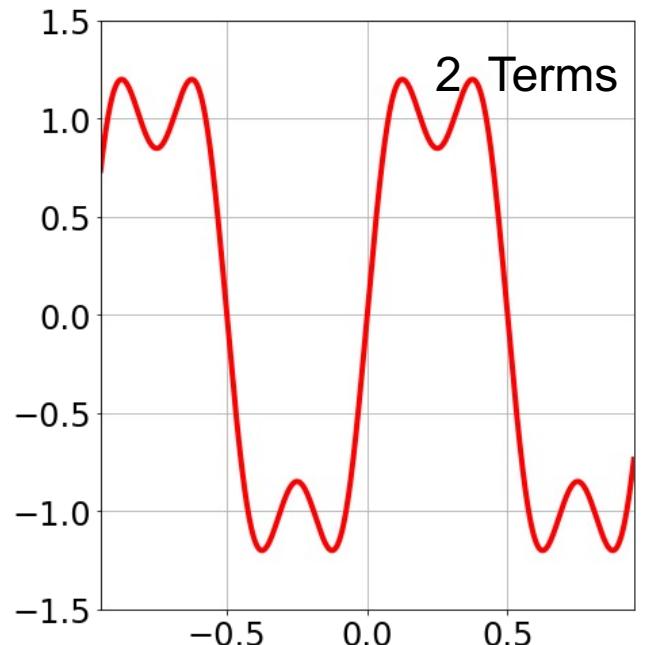
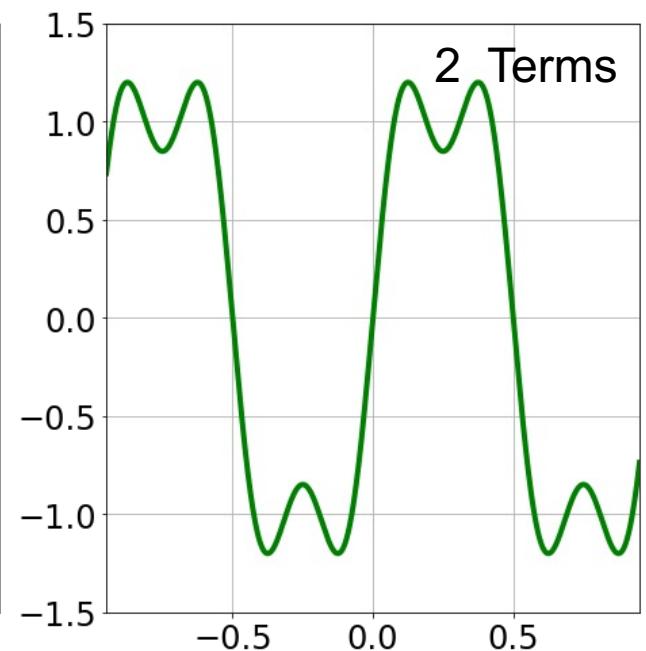
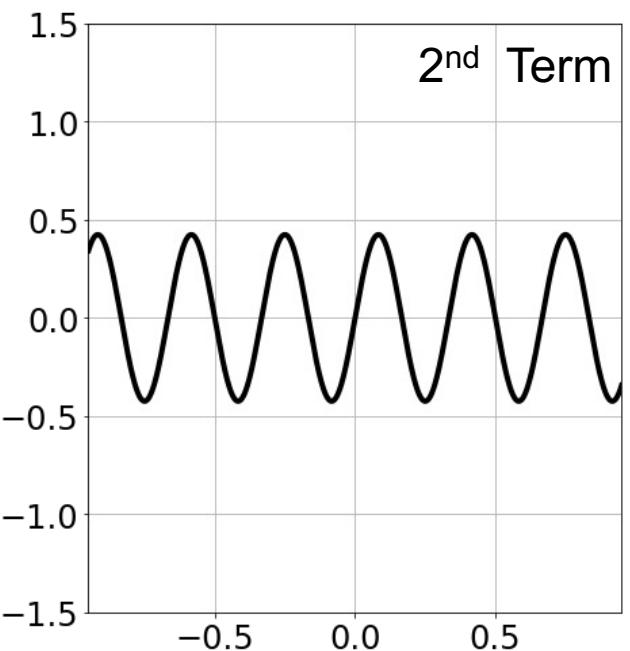
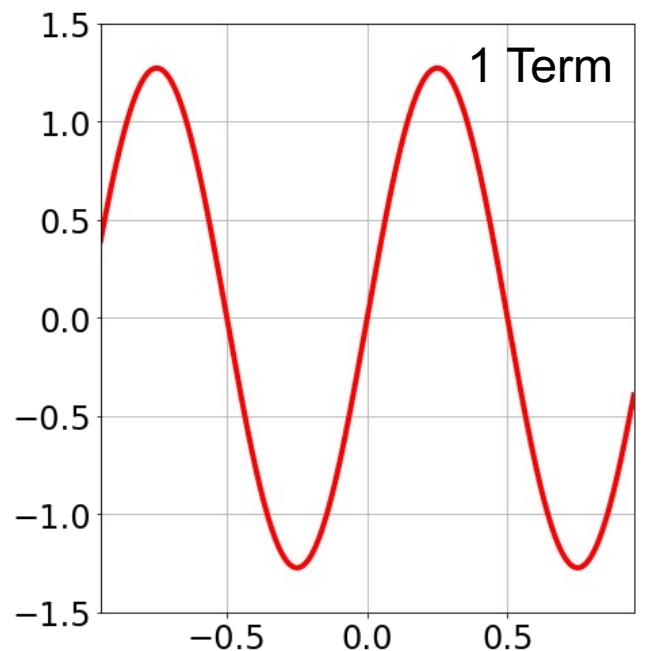


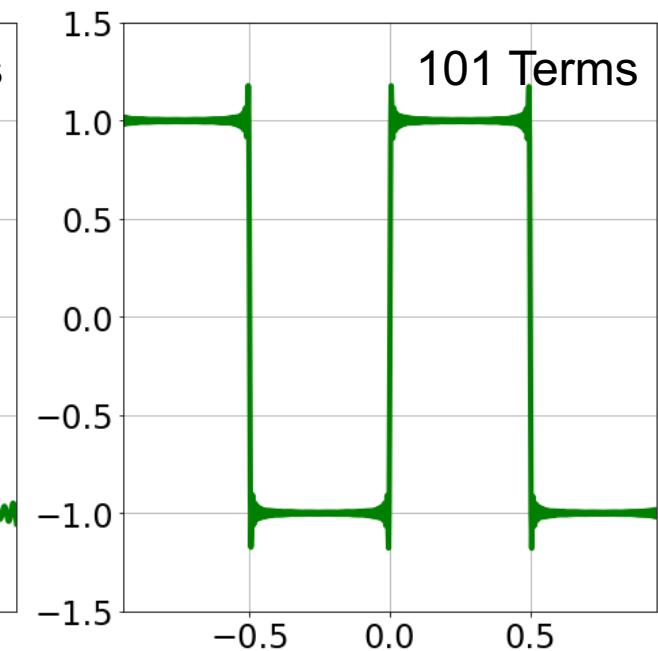
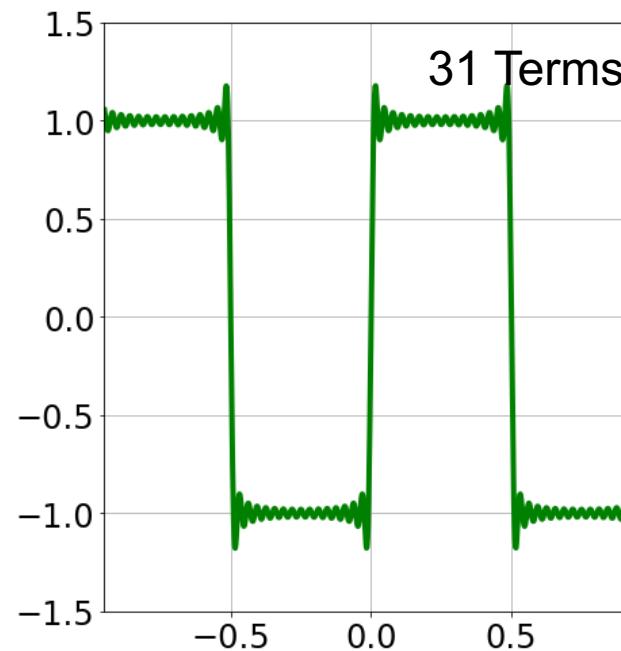
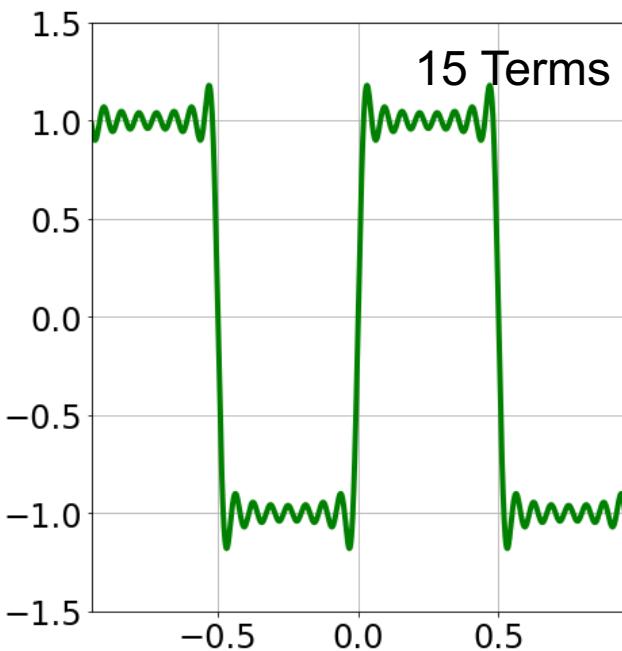
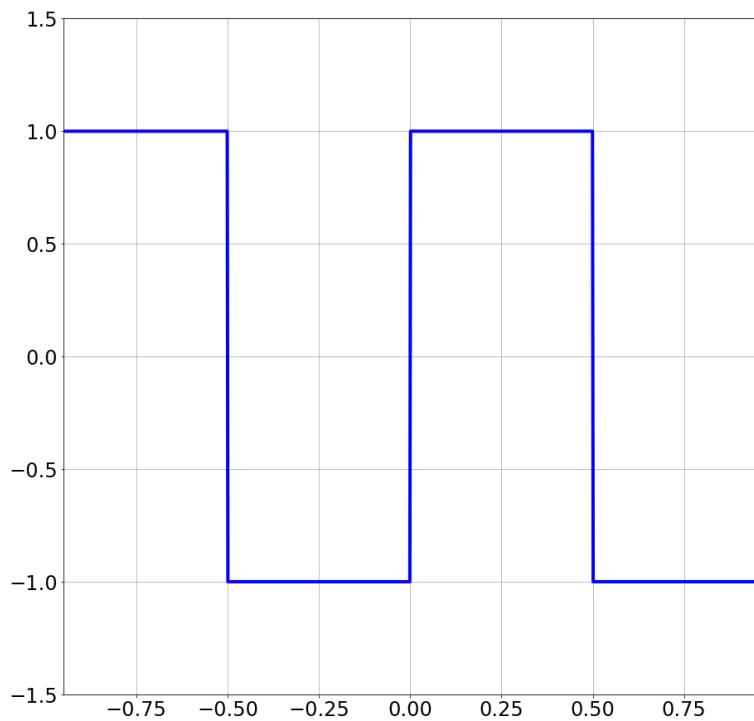
Square Wave:

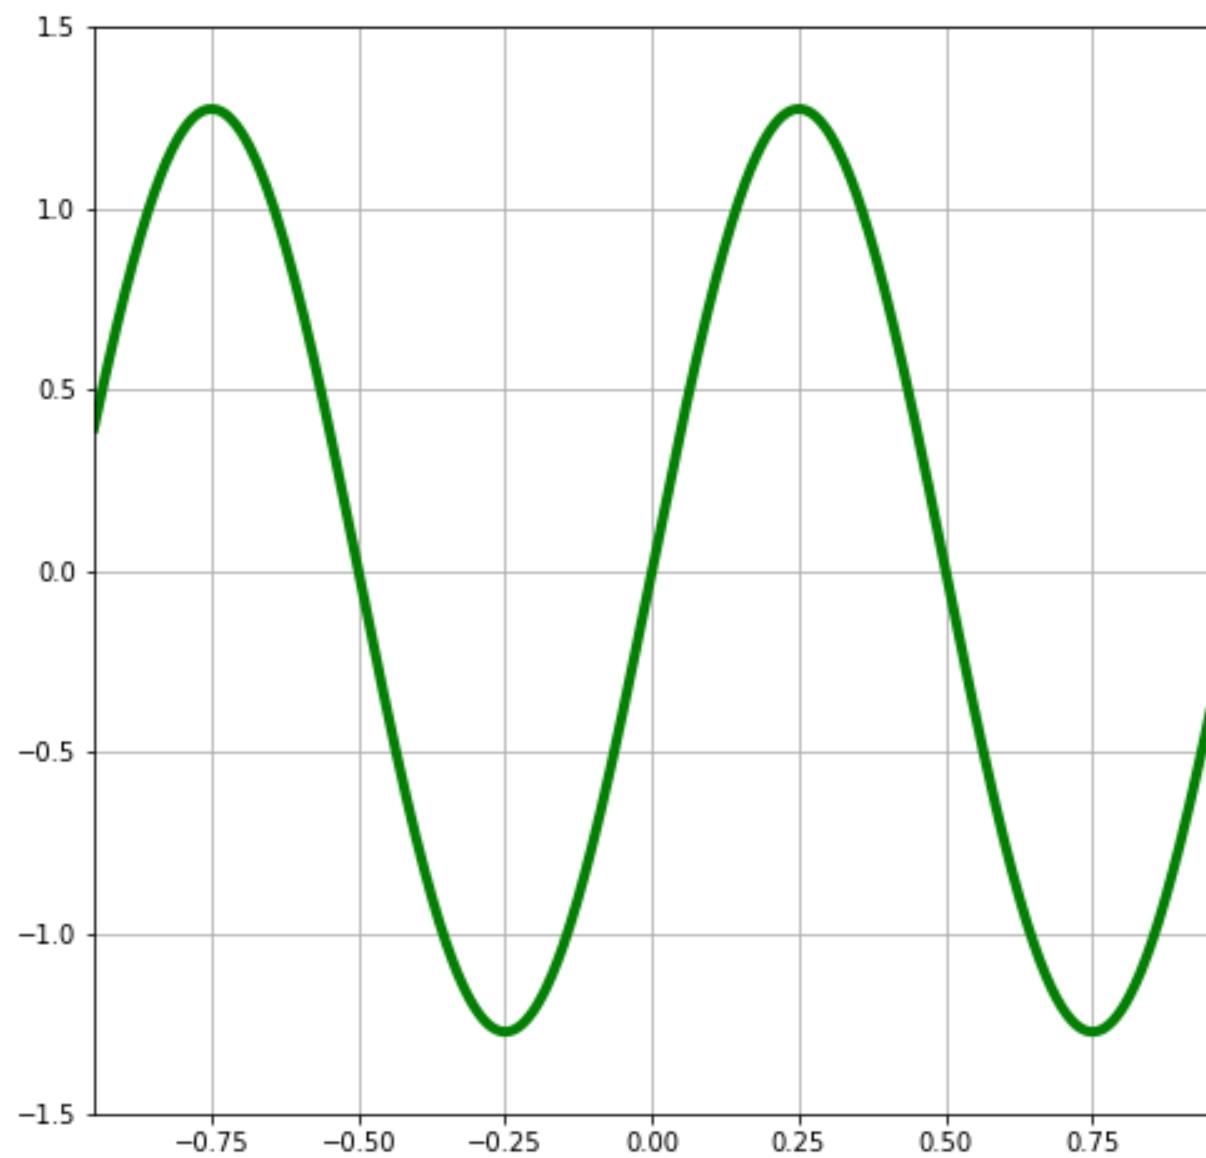
$$f(x) = \begin{cases} 1, & \text{if } \text{frac}(x) < 0.5 \\ -1, & \text{otherwise} \end{cases}$$

Fourier Transform:

$$\frac{4}{\pi} \left(\frac{\sin(2\pi \cdot 1 \cdot x)}{1} + \frac{\sin(2\pi \cdot 3 \cdot x)}{3} + \frac{\sin(2\pi \cdot 5 \cdot x)}{5} + \dots \right)$$







Fourier Transform

- Fourier transform stores the magnitude and phase at each frequency
 - Magnitude encodes how much signal there is at a particular frequency
 - Phase encodes spatial information (indirectly)
 - For mathematical convenience, this is often notated in terms of complex numbers
- Amplitude: $A = \sqrt{R(\omega)^2 + I(\omega)^2}$
- Phase: $\phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$

Computing the Fourier Transform

- $H(\omega) = \mathcal{F}[h(x)]$
- Continuous:
 - $H(\omega) = \int_{-\infty}^{\infty} h(x)e^{-j\omega x}dx$
- Discrete:
 - $H(k) = \frac{1}{N} \sum_{x=0}^{N-1} h(x)e^{-j2\pi kx/N}$
- Euler's Formula:
 - $e^{jnx} = \cos(nx) + j \sin(nx)$

Properties of Fourier Transforms

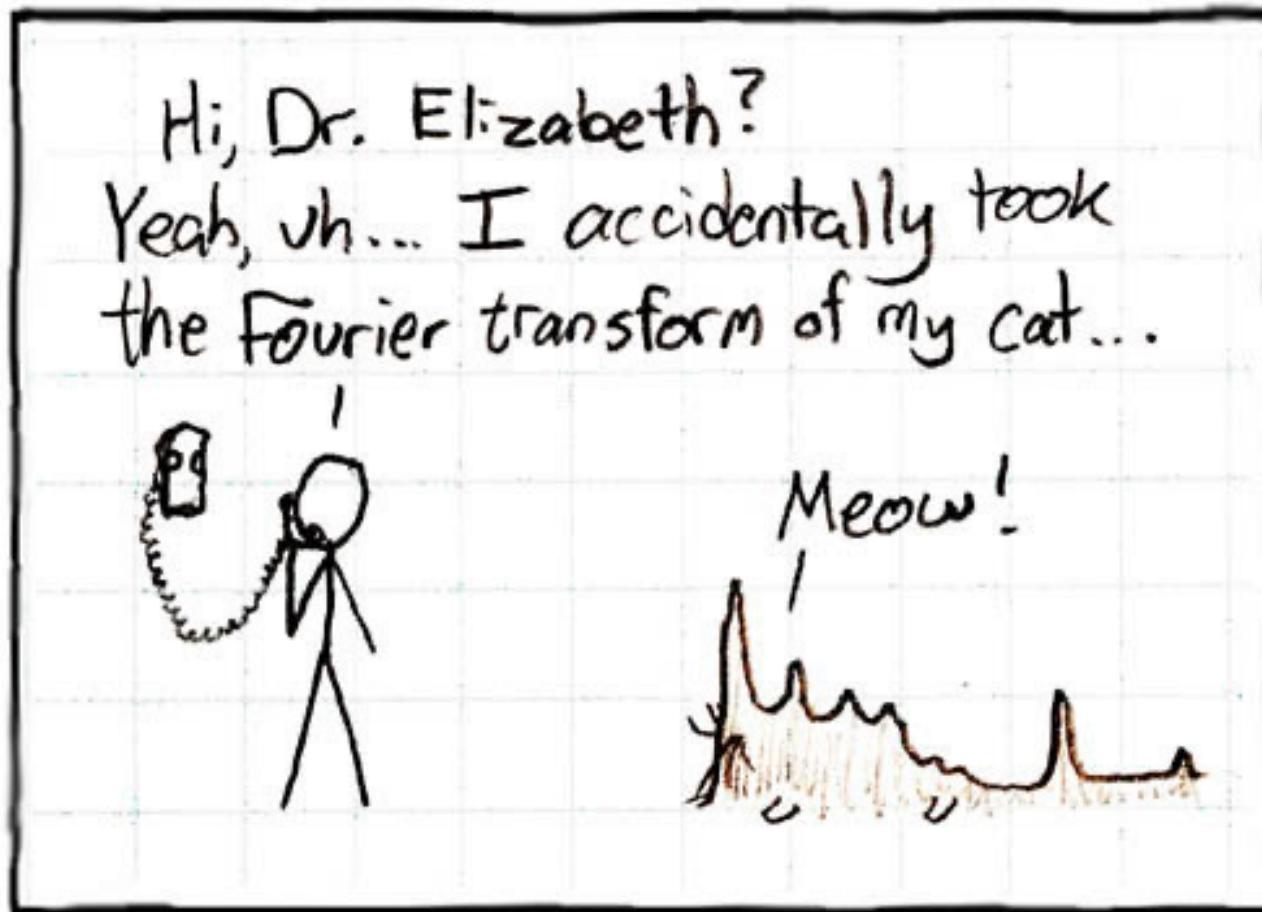
- Linearity:
 - $\mathcal{F}[ax(t) + by(t)] = a\mathcal{F}[x(t)] + b\mathcal{F}[y(t)]$
- Fourier transform of a real signal is symmetric about the origin
- The energy of the signal is the same as the energy of its Fourier transform

The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms
$$-\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$
- The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms
$$-\mathcal{F}^{-1}[gh] = \mathcal{F}^{-1}[g] * \mathcal{F}^{-1}[h]$$
- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

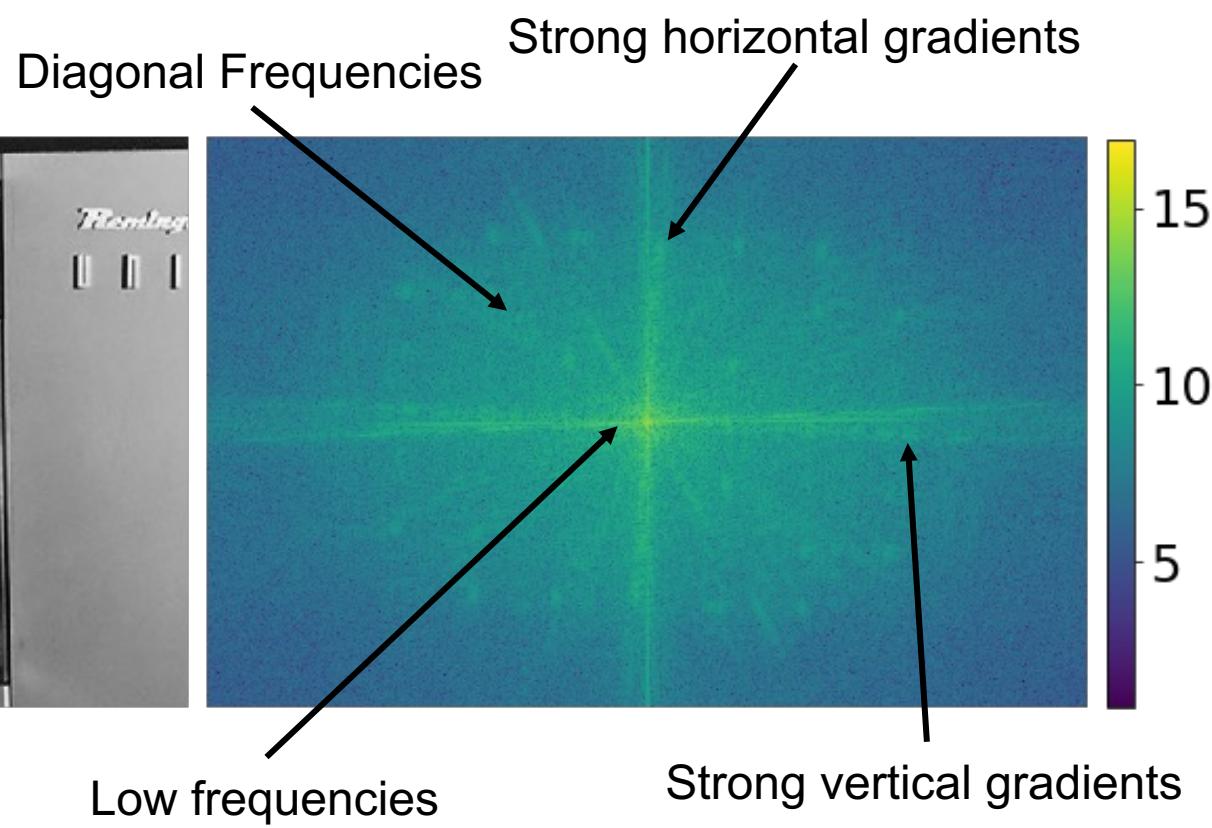
Other signals

- We can also think of all kinds of other signals the same way



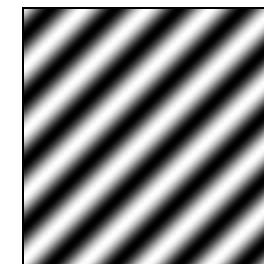
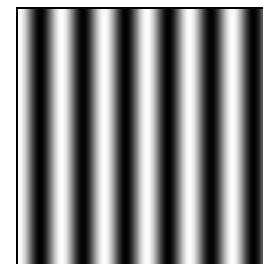
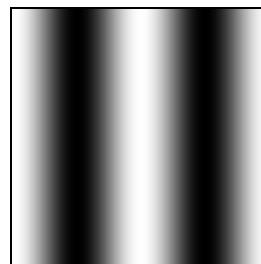
Images

$$H(k, l) = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} h(m, n) e^{-j2\pi(\frac{k}{M}m + \frac{l}{N}n)}$$

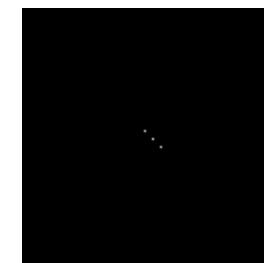
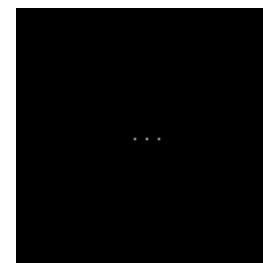
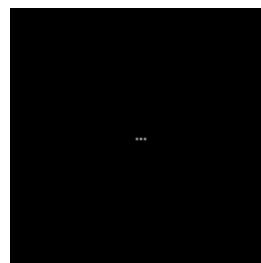


Fourier analysis in images

Intensity Image



Fourier Image

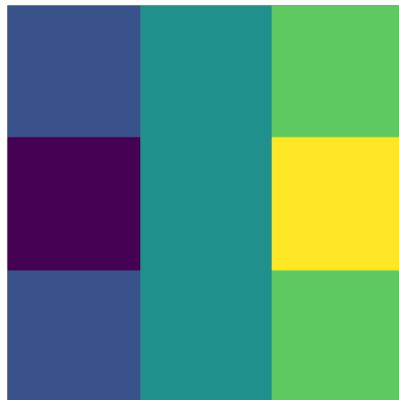


Filtering in spatial domain

Original



*



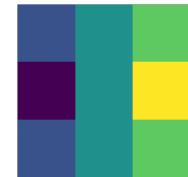
=

Filtered via Convolution

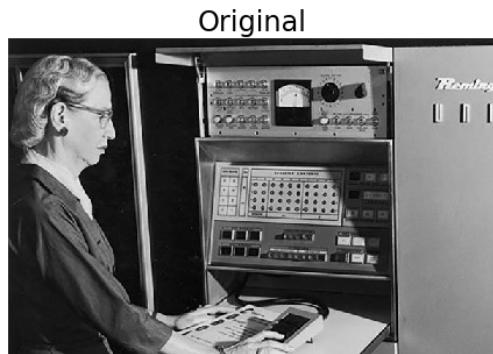


-1	0	1
-2	0	2
-1	0	1

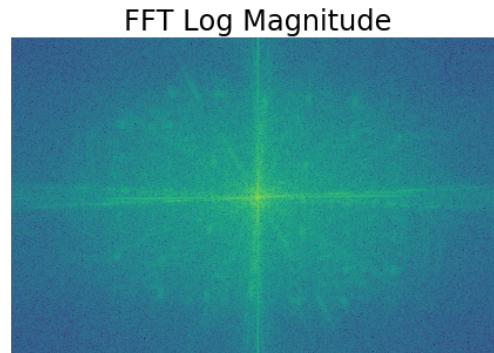
Filtering in frequency domain



FFT

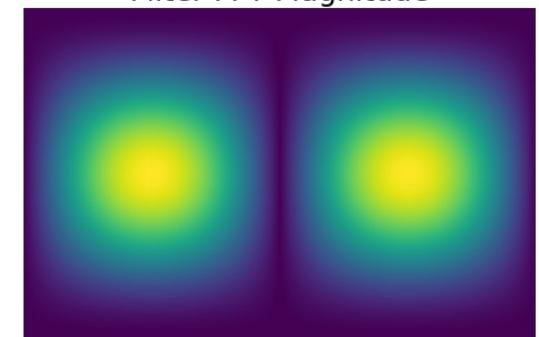


FFT



Filter FFT Magnitude

X



||



Inverse FFT

←

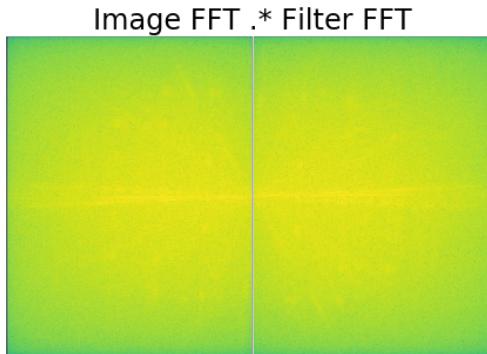


Image FFT .* Filter FFT

Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

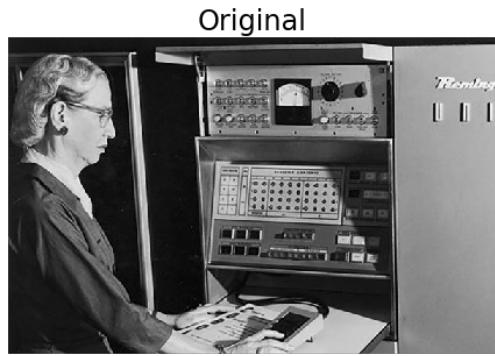
Gaussian



Box filter

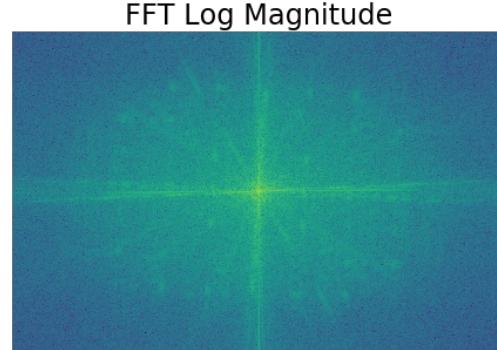


Filtering in frequency domain (Box)



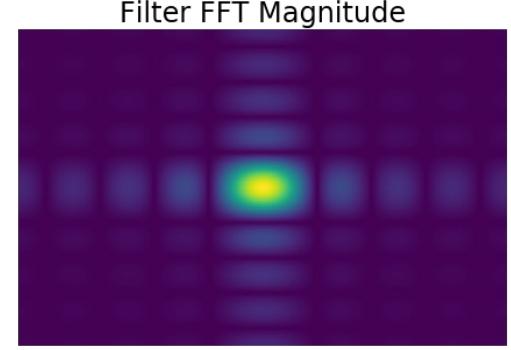
Original

FFT



FFT Log Magnitude

X



Filter FFT Magnitude

||



Filtered via FFT

Inverse FFT

←

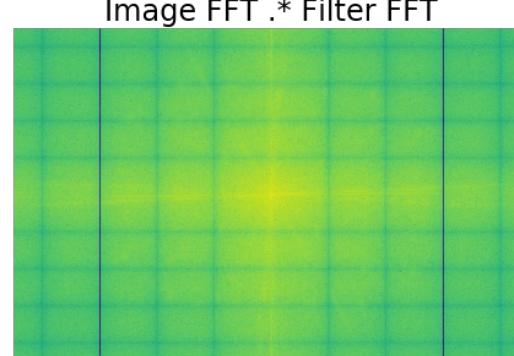
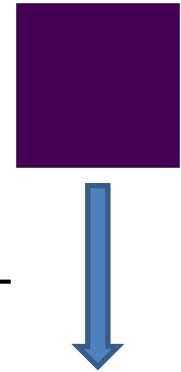
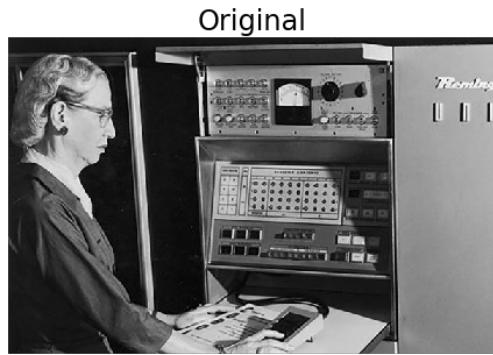


Image FFT .* Filter FFT

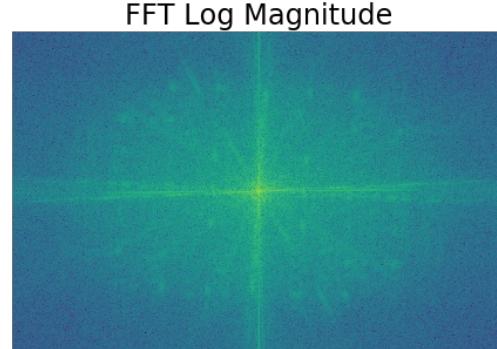


Filtering in frequency domain (Gaussian)

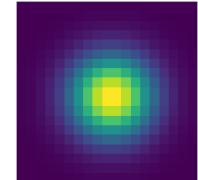


Original

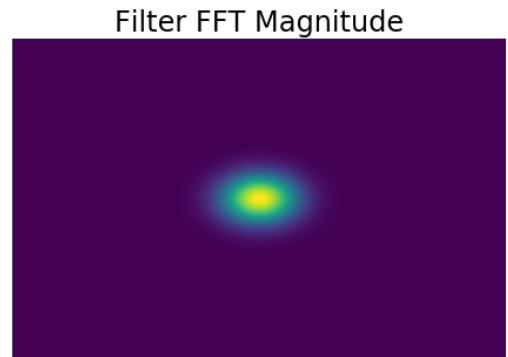
FFT



FFT Log Magnitude



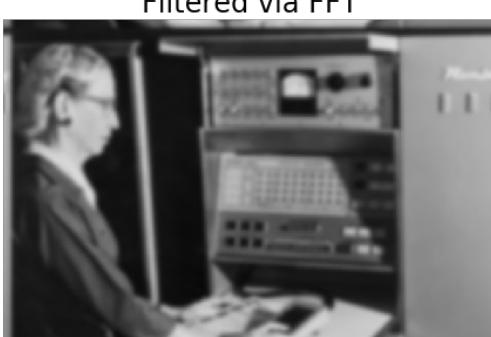
FFT



Filter FFT Magnitude

X

||



Filtered via FFT

Inverse FFT

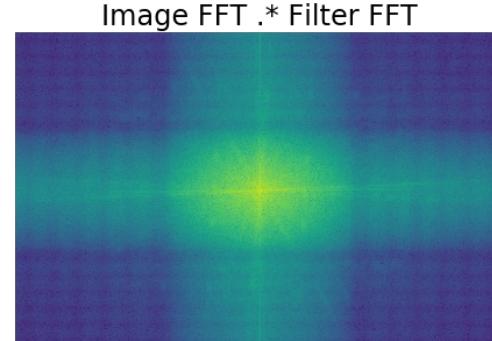
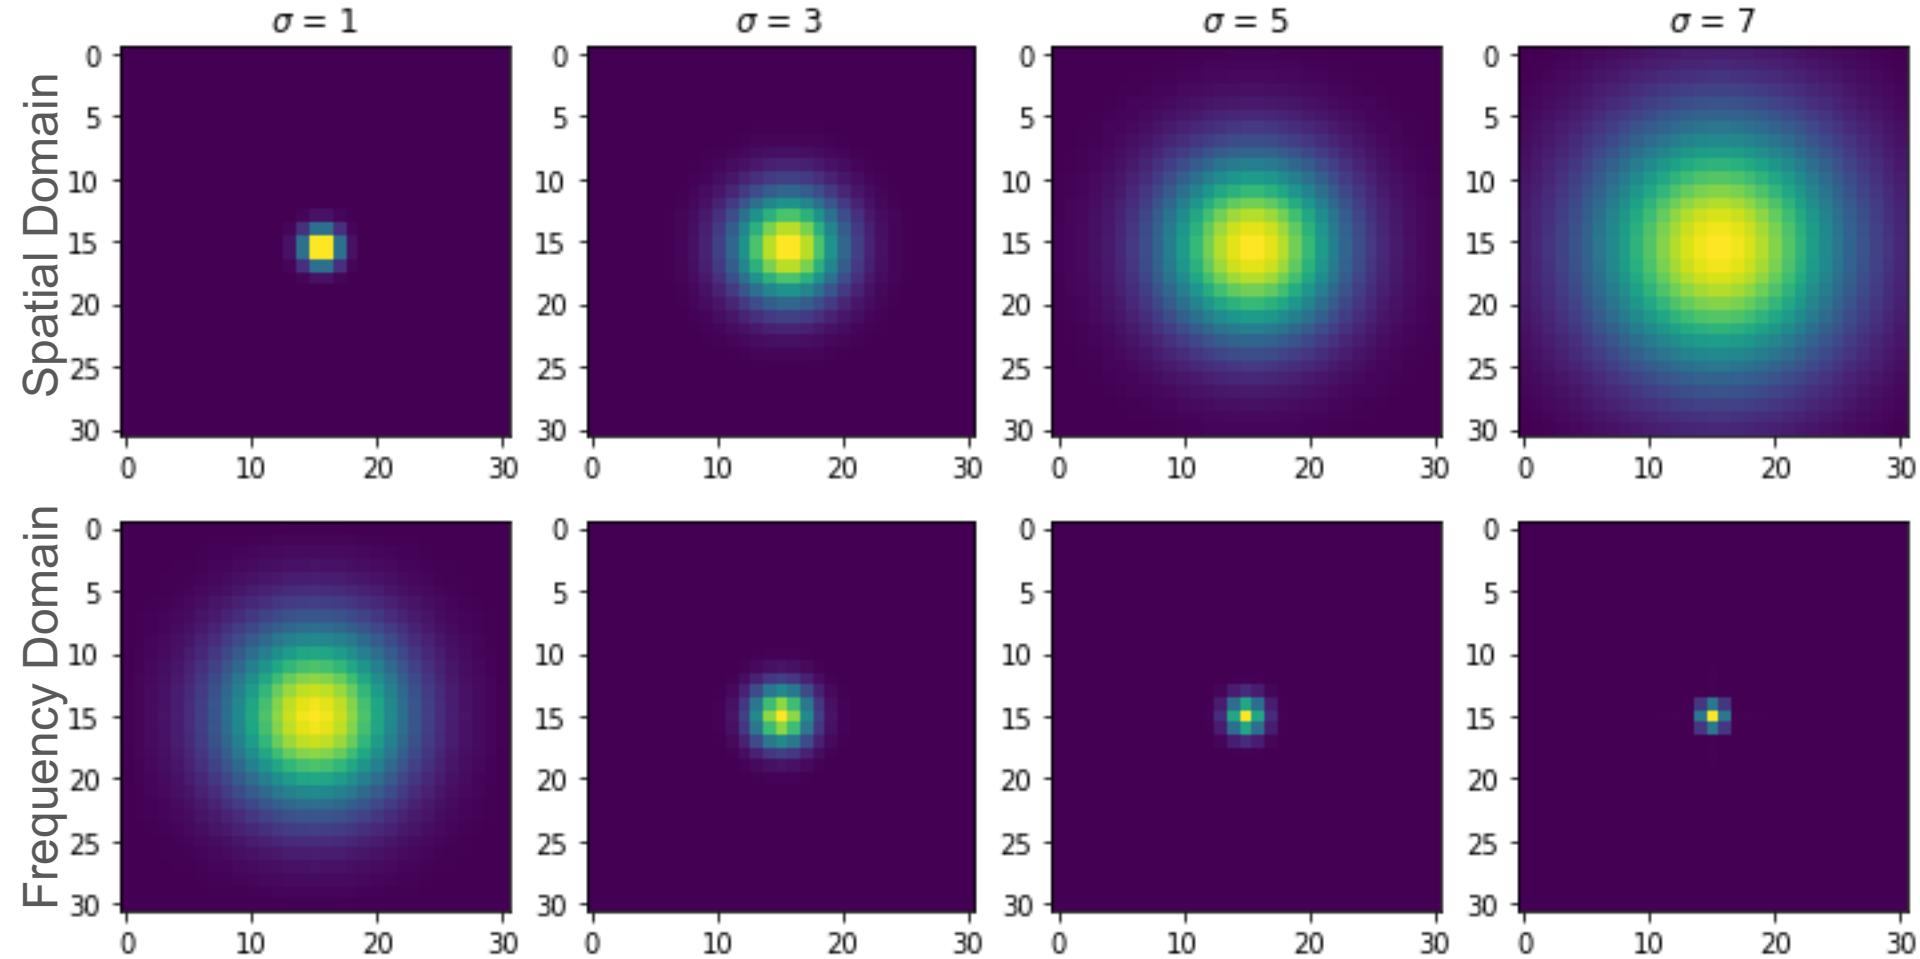
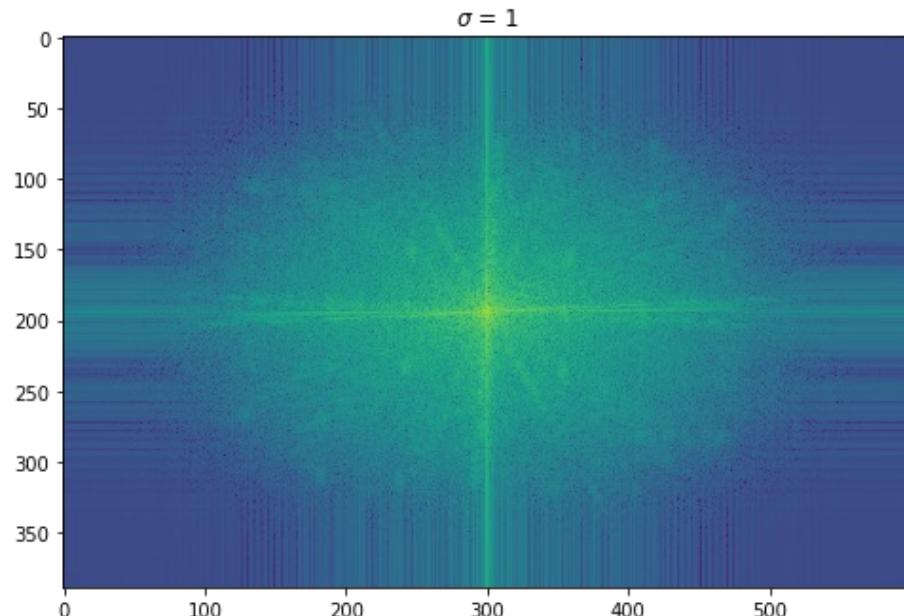


Image FFT .* Filter FFT

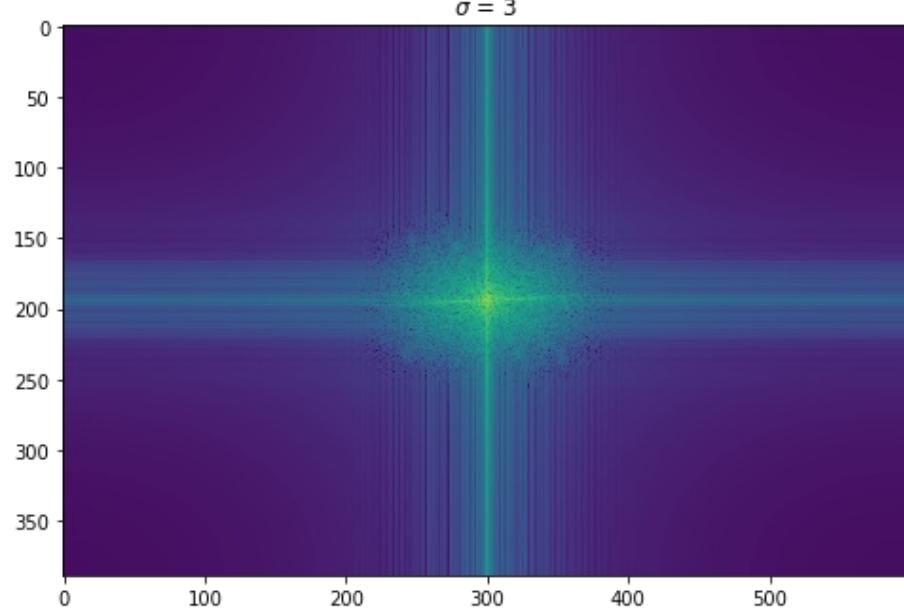
Filtering in frequency domain (Gaussian)



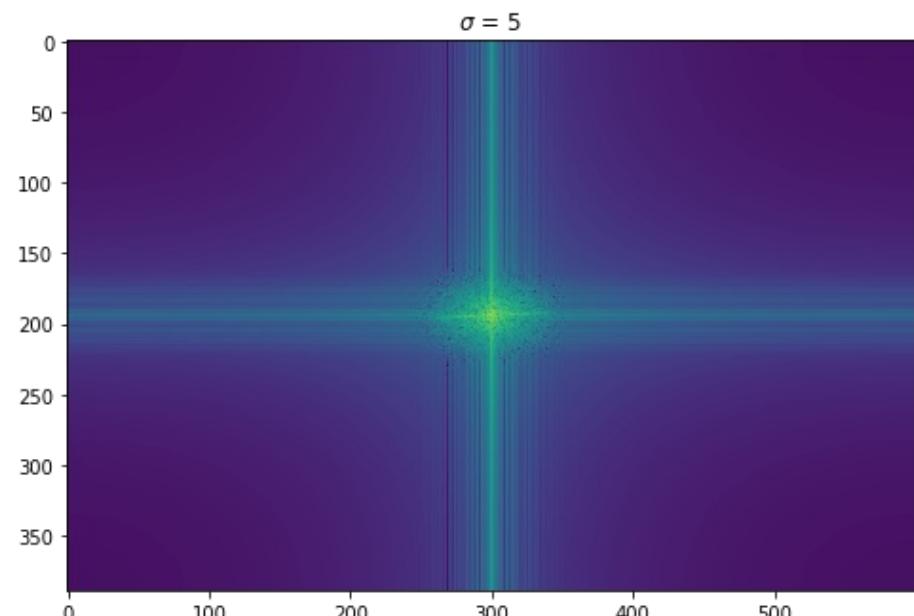
Filtering in frequency domain (Gaussian)



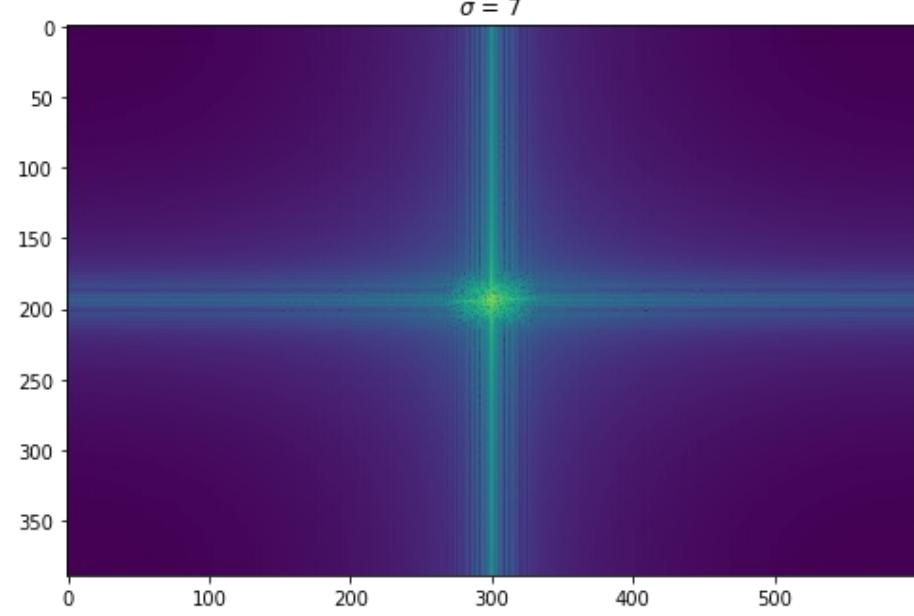
$\sigma = 1$



$\sigma = 3$

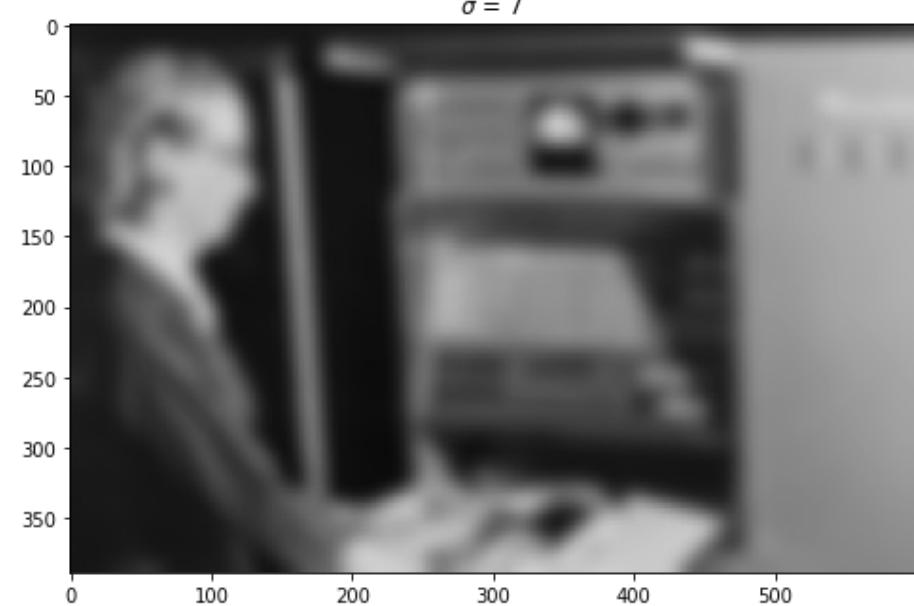
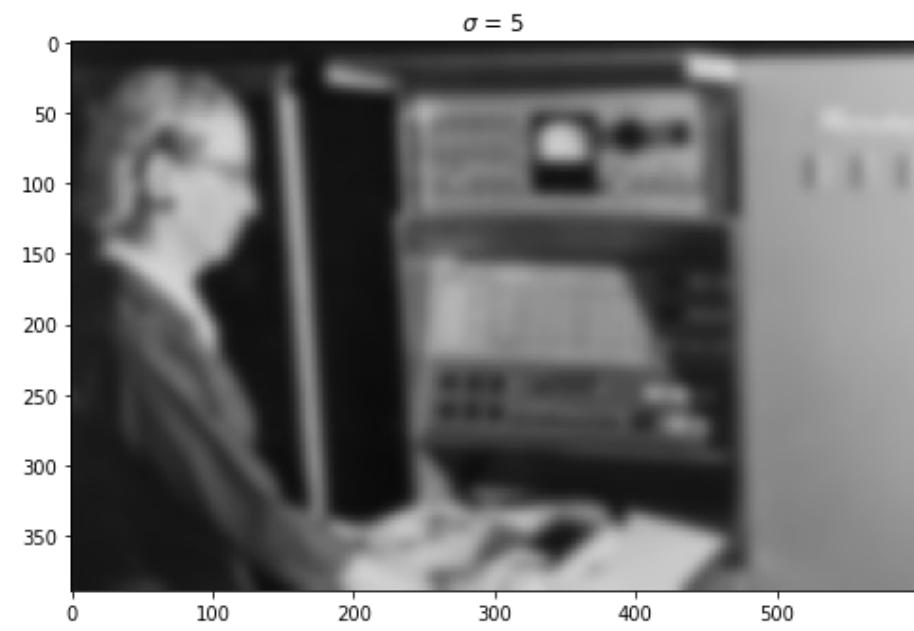
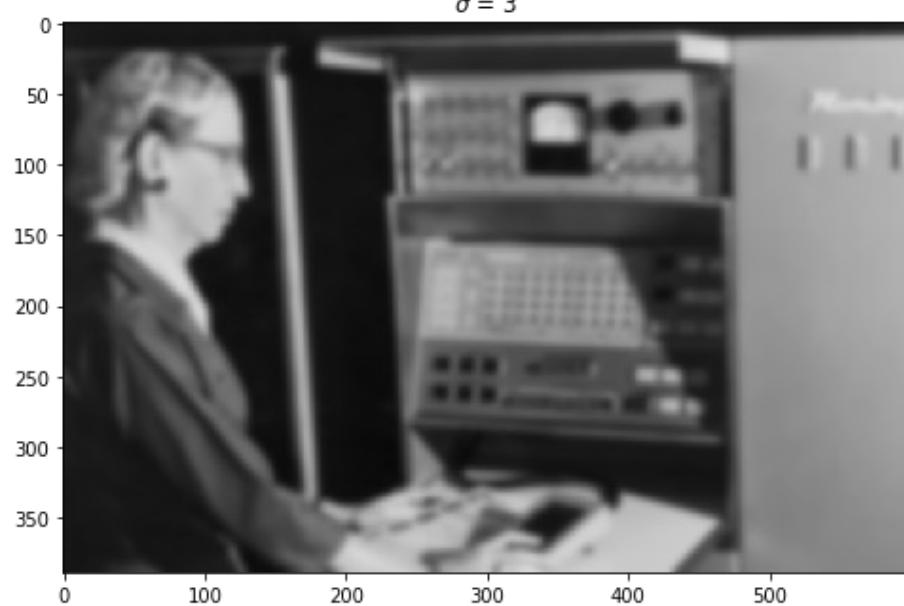
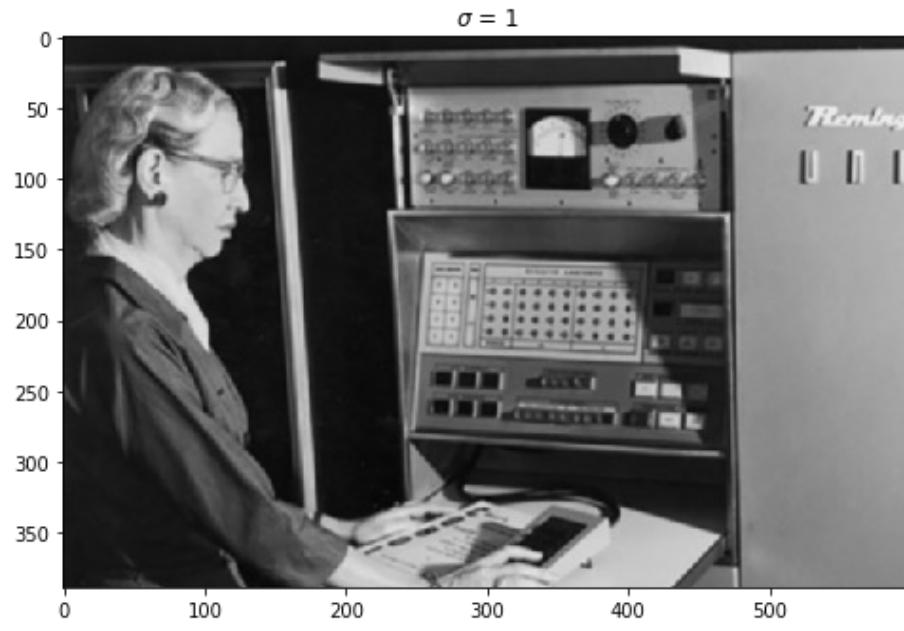


$\sigma = 5$



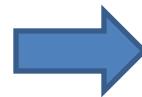
$\sigma = 7$

Filtering in frequency domain (Gaussian)

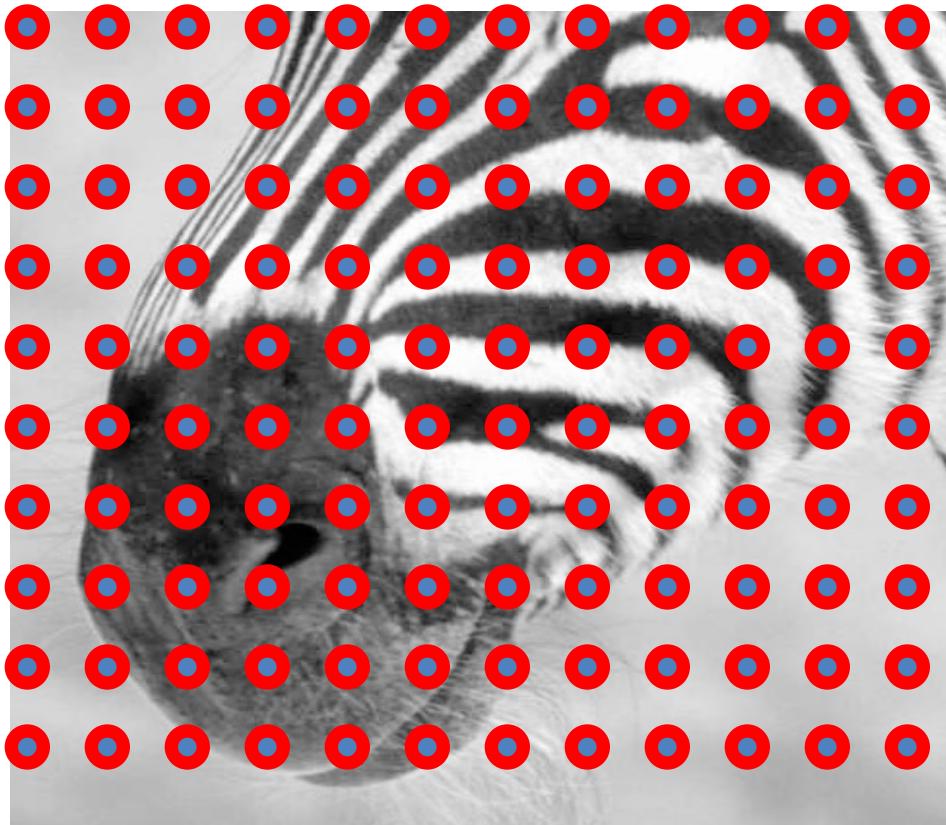


Sampling

Why does a lower resolution image still make sense to us? What do we lose?



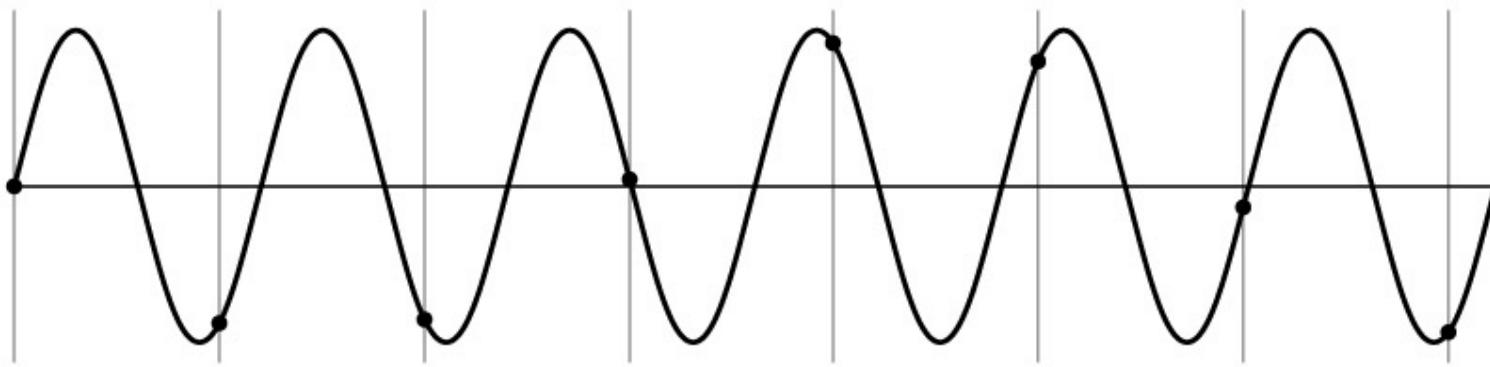
Subsampling by a factor of 2



Throw away every other row and column to create a 1/2 size image

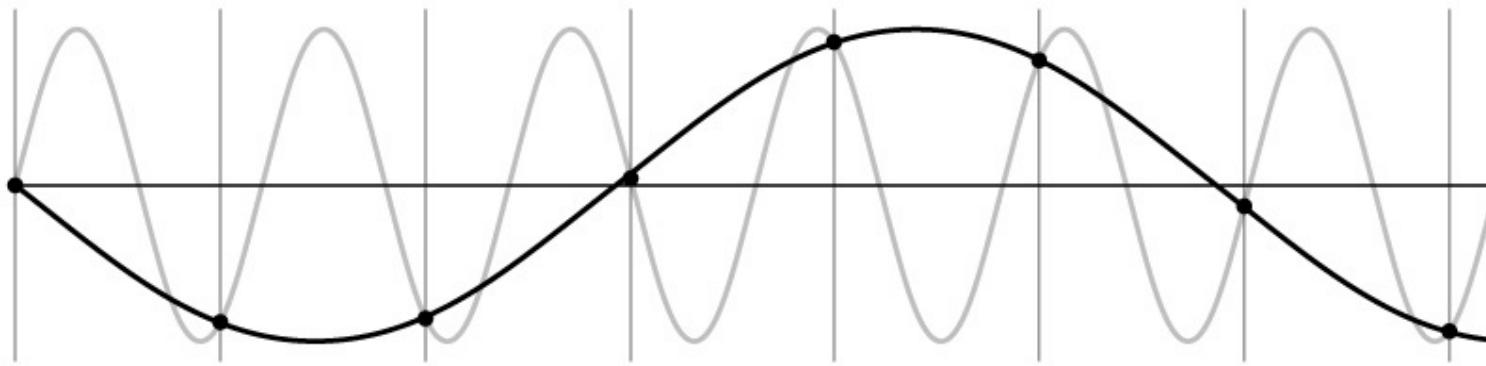
Aliasing problem

- 1D example (sinewave):



Aliasing problem

- 1D example (sinewave):



Aliasing problem

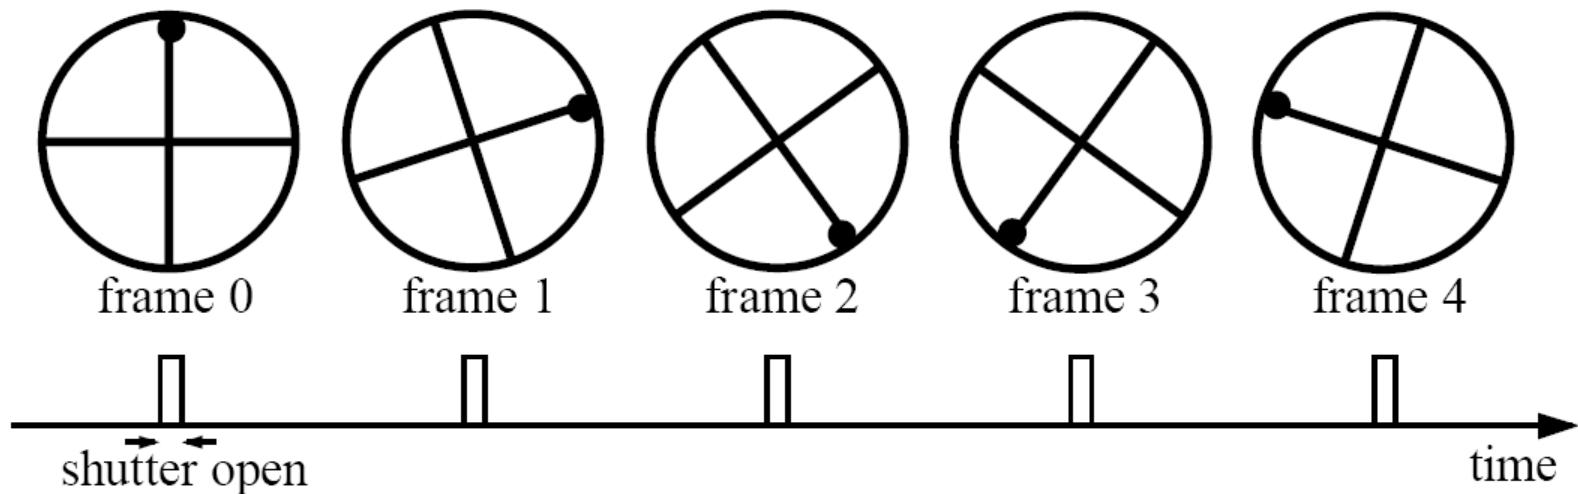
- Sub-sampling may be dangerous....
- Characteristic errors may appear:
 - “Wagon wheels rolling the wrong way in movies” [See](#)
 - “Checkerboards disintegrate in ray tracing”
 - “Striped shirts look funny on color television”

Aliasing in video

Imagine a spoked wheel moving to the right (rotating clockwise).

Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = $1/30$ sec. for video, $1/24$ sec. for film):



Without dot, wheel appears to be rotating slowly backwards!
(counterclockwise)

Aliasing in graphics



Sampling and aliasing

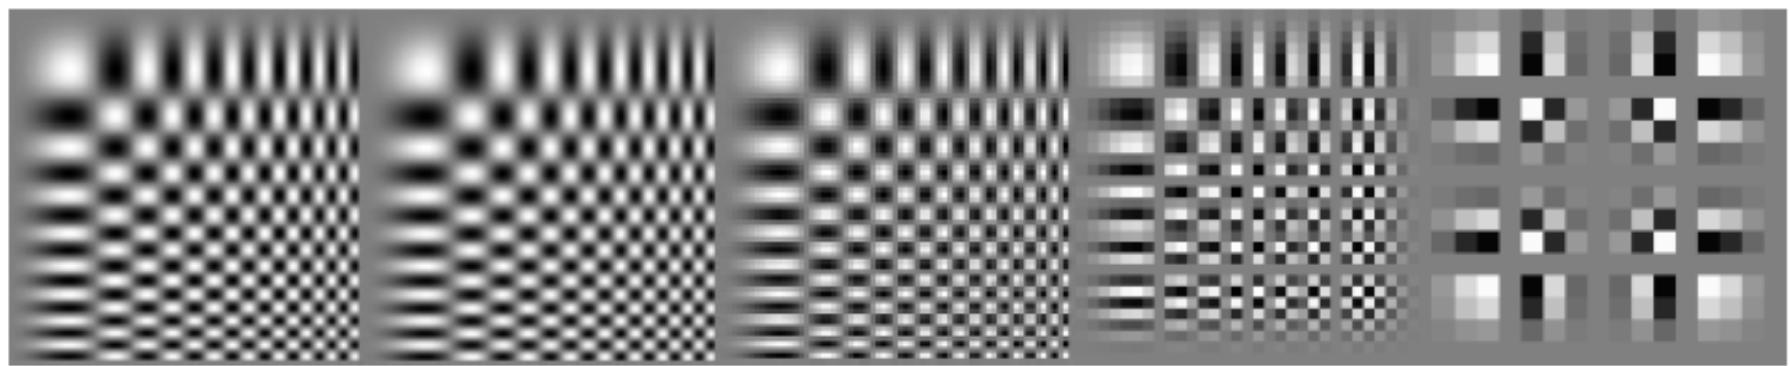
256x256

128x128

64x64

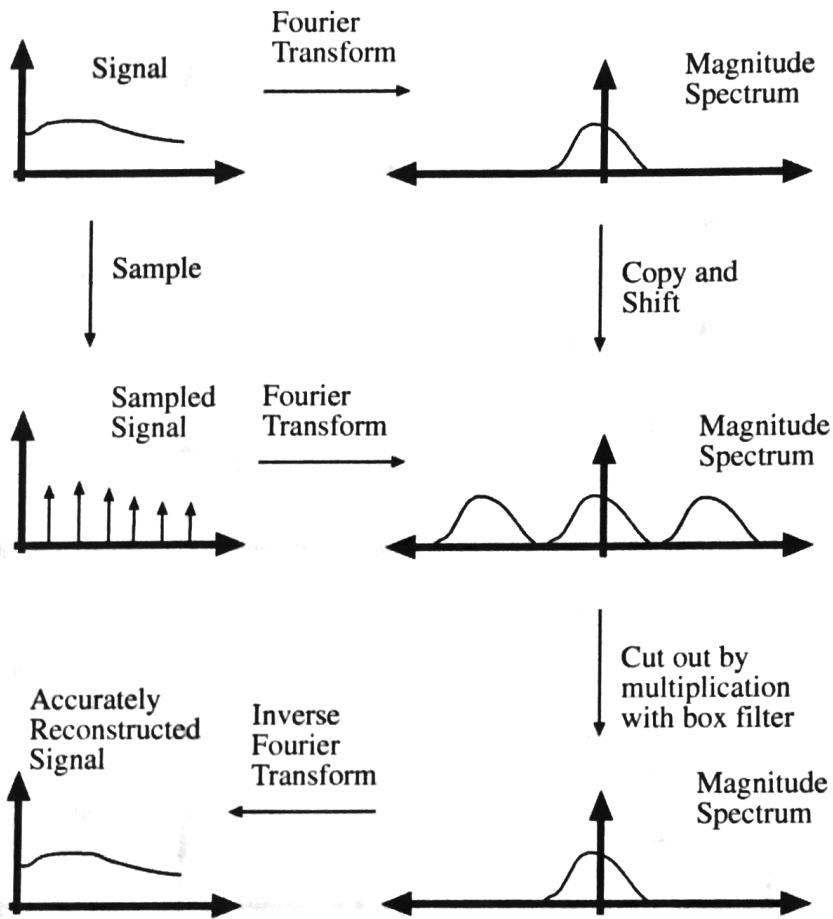
32x32

16x16

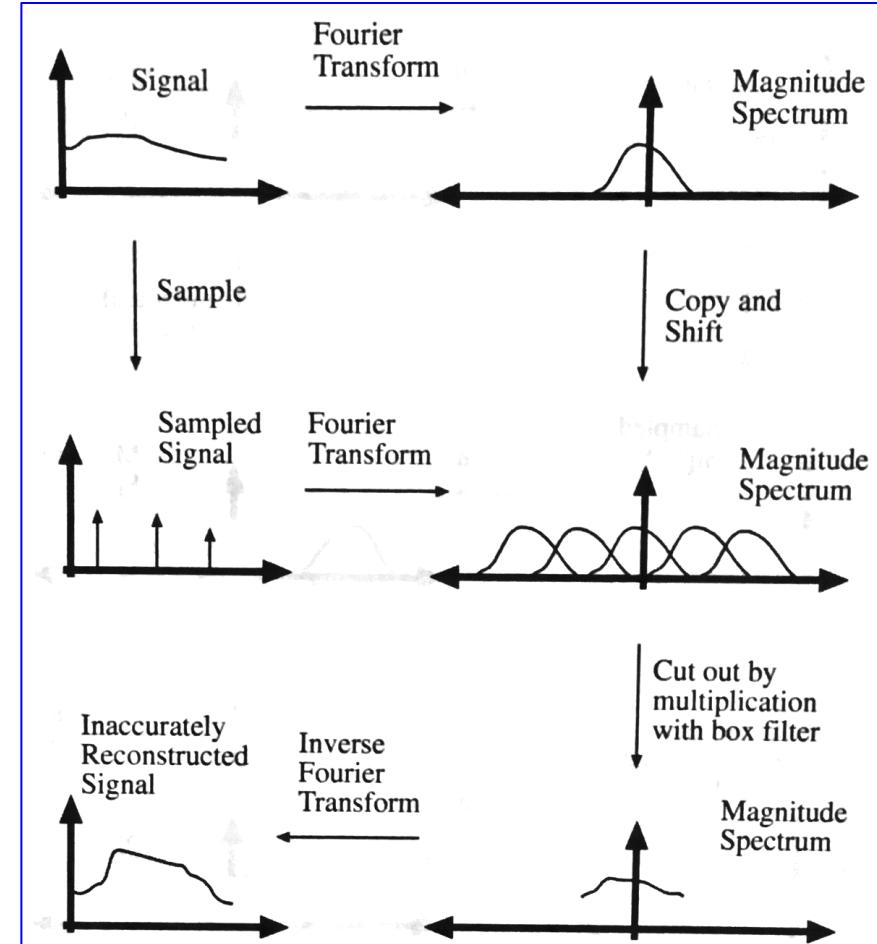


Aliasing in Frequency Domain

No Aliasing

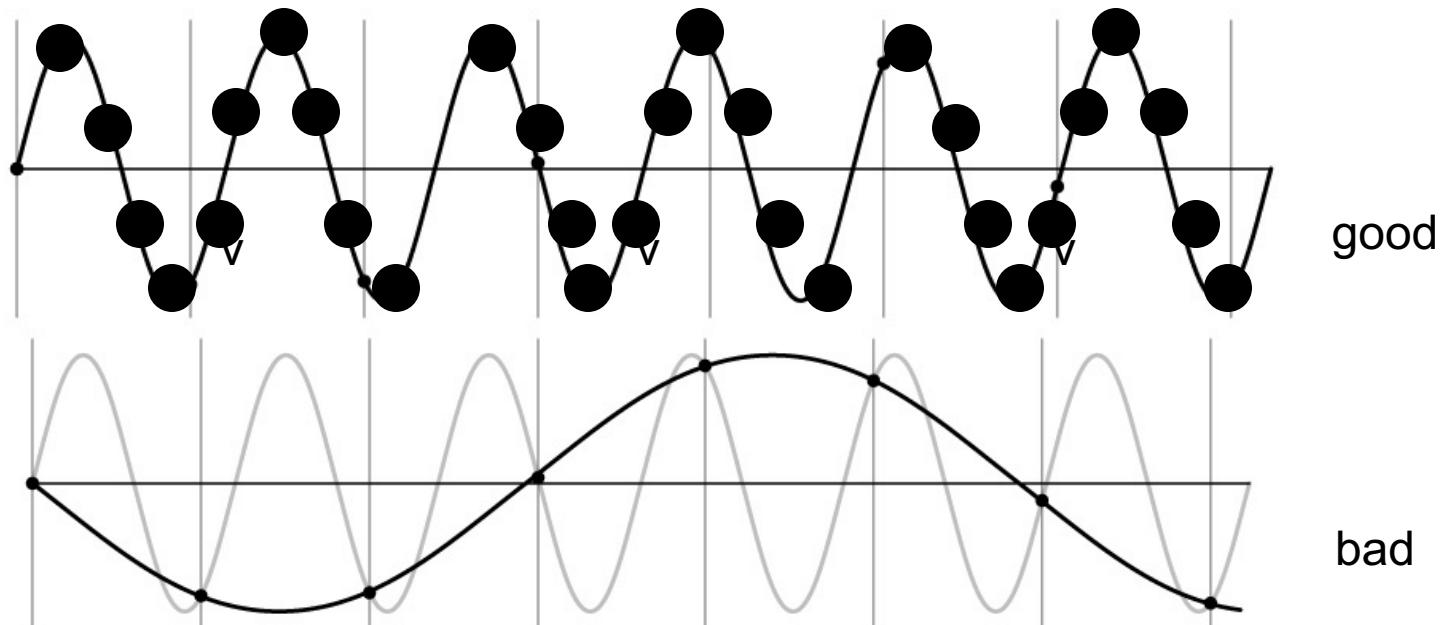


Aliasing



Nyquist-Shannon Sampling Theorem

- When sampling a signal at discrete intervals, the sampling frequency must be $\geq 2 \times f_{\max}$
- f_{\max} = max frequency of the input signal
- This will allow to reconstruct the original perfectly from the sampled version



Anti-aliasing

Solutions:

- Sample more often
- Get rid of all frequencies that are greater than half the new sampling frequency
 - Will lose information
 - But it's better than aliasing
 - Apply a smoothing filter

Algorithm for downsampling by factor of 2

1. Start with $\text{image}(h, w)$
2. Apply low-pass filter
3. Sample every other pixel

Anti-aliasing

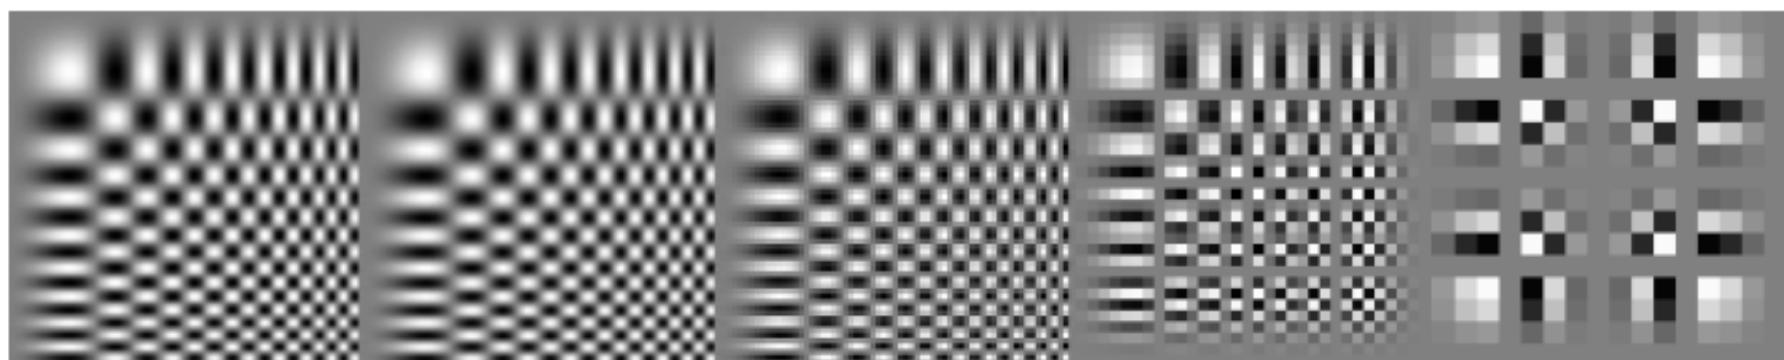
256x256

128x128

64x64

32x32

16x16



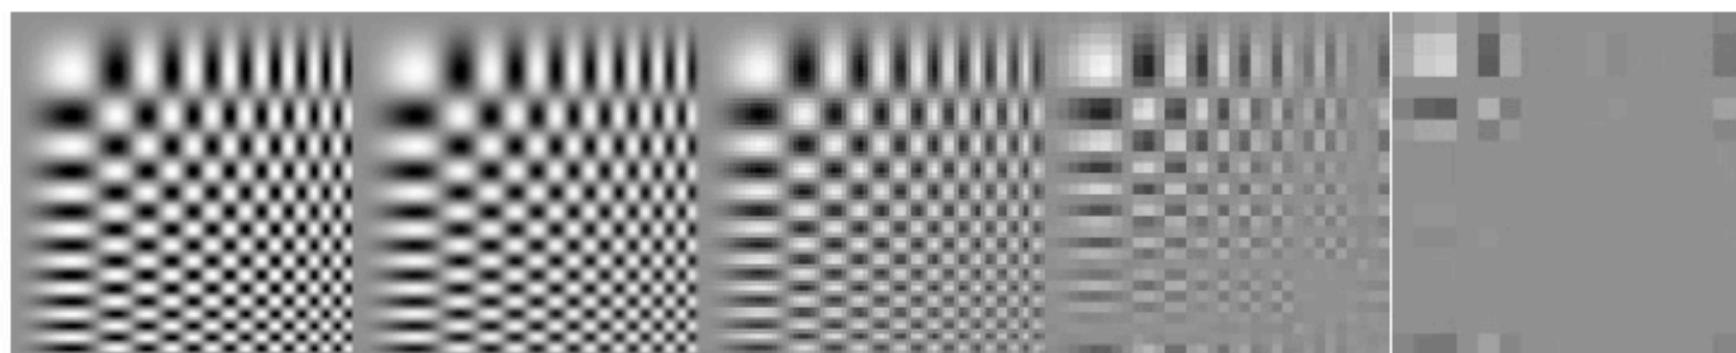
256x256

128x128

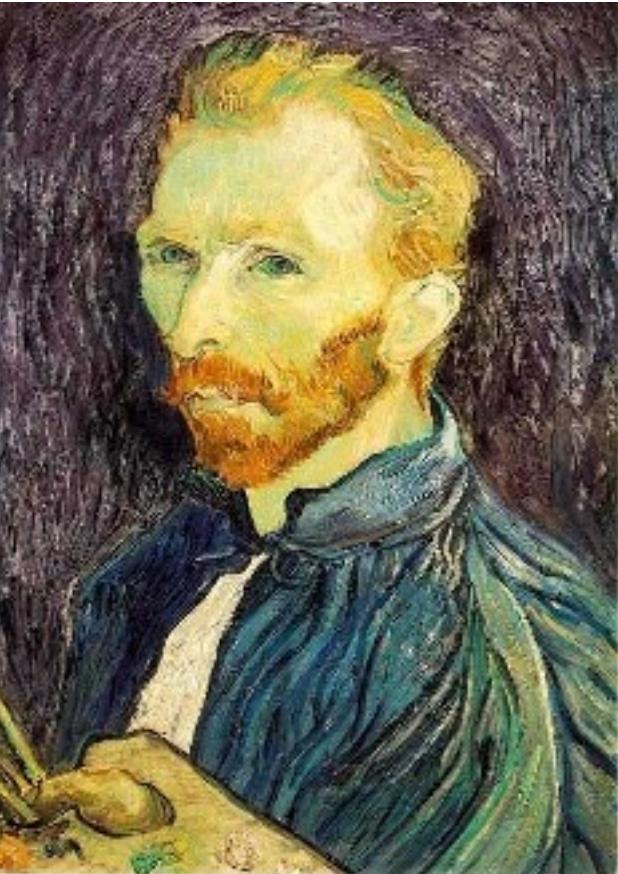
64x64

32x32

16x16



Subsampling without pre-filtering



1/2

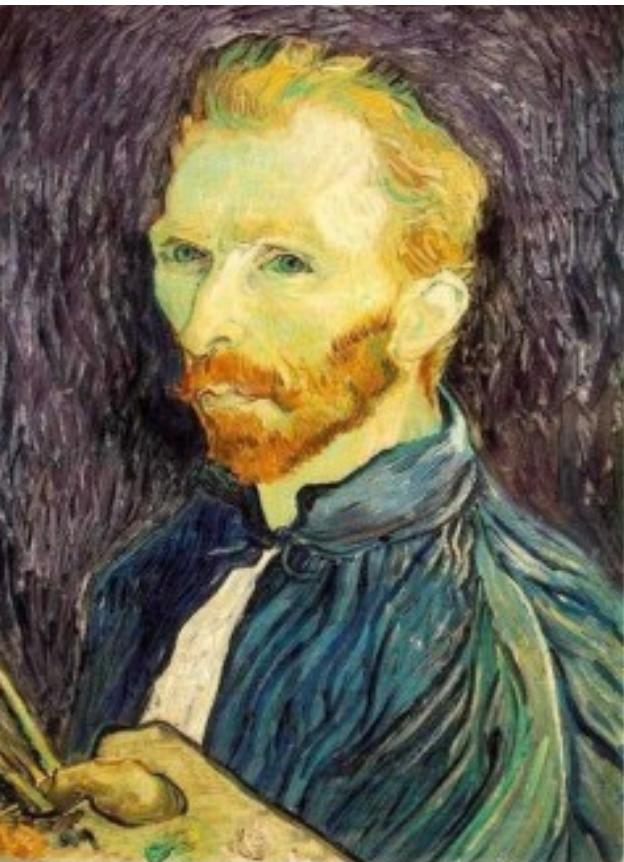


1/4 (2x zoom)

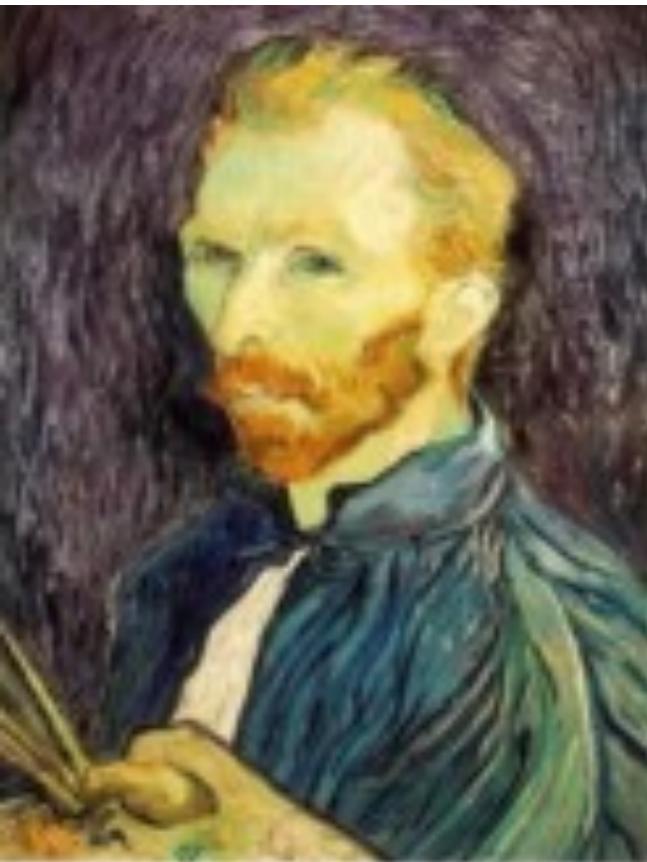


1/8 (4x zoom)

Subsampling with Gaussian pre-filtering



Gaussian 1/2

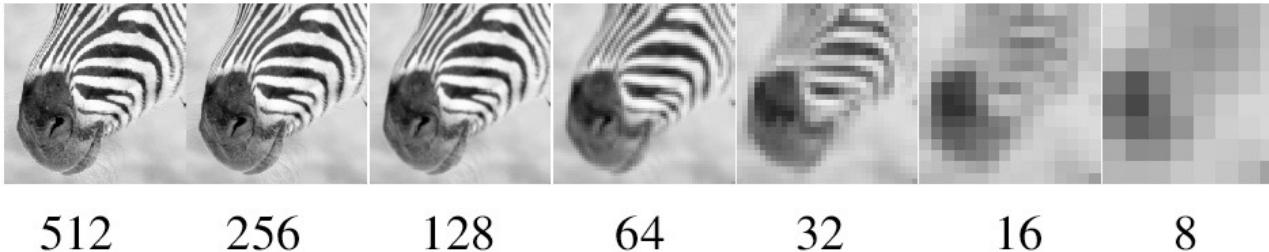


G 1/4



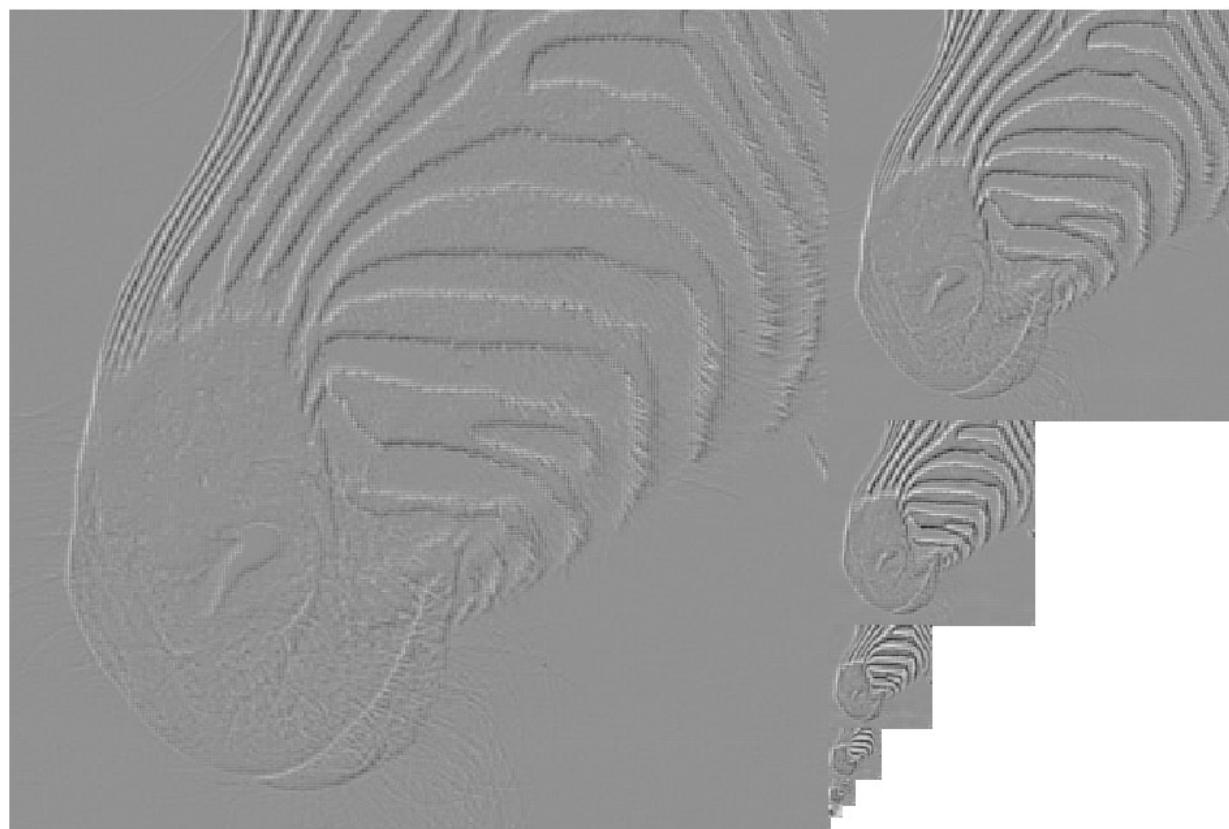
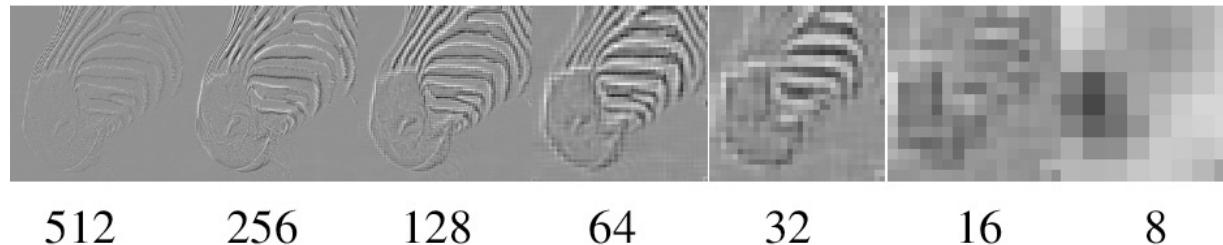
G 1/8

Gaussian pyramid (Repeated blurring and sampling)



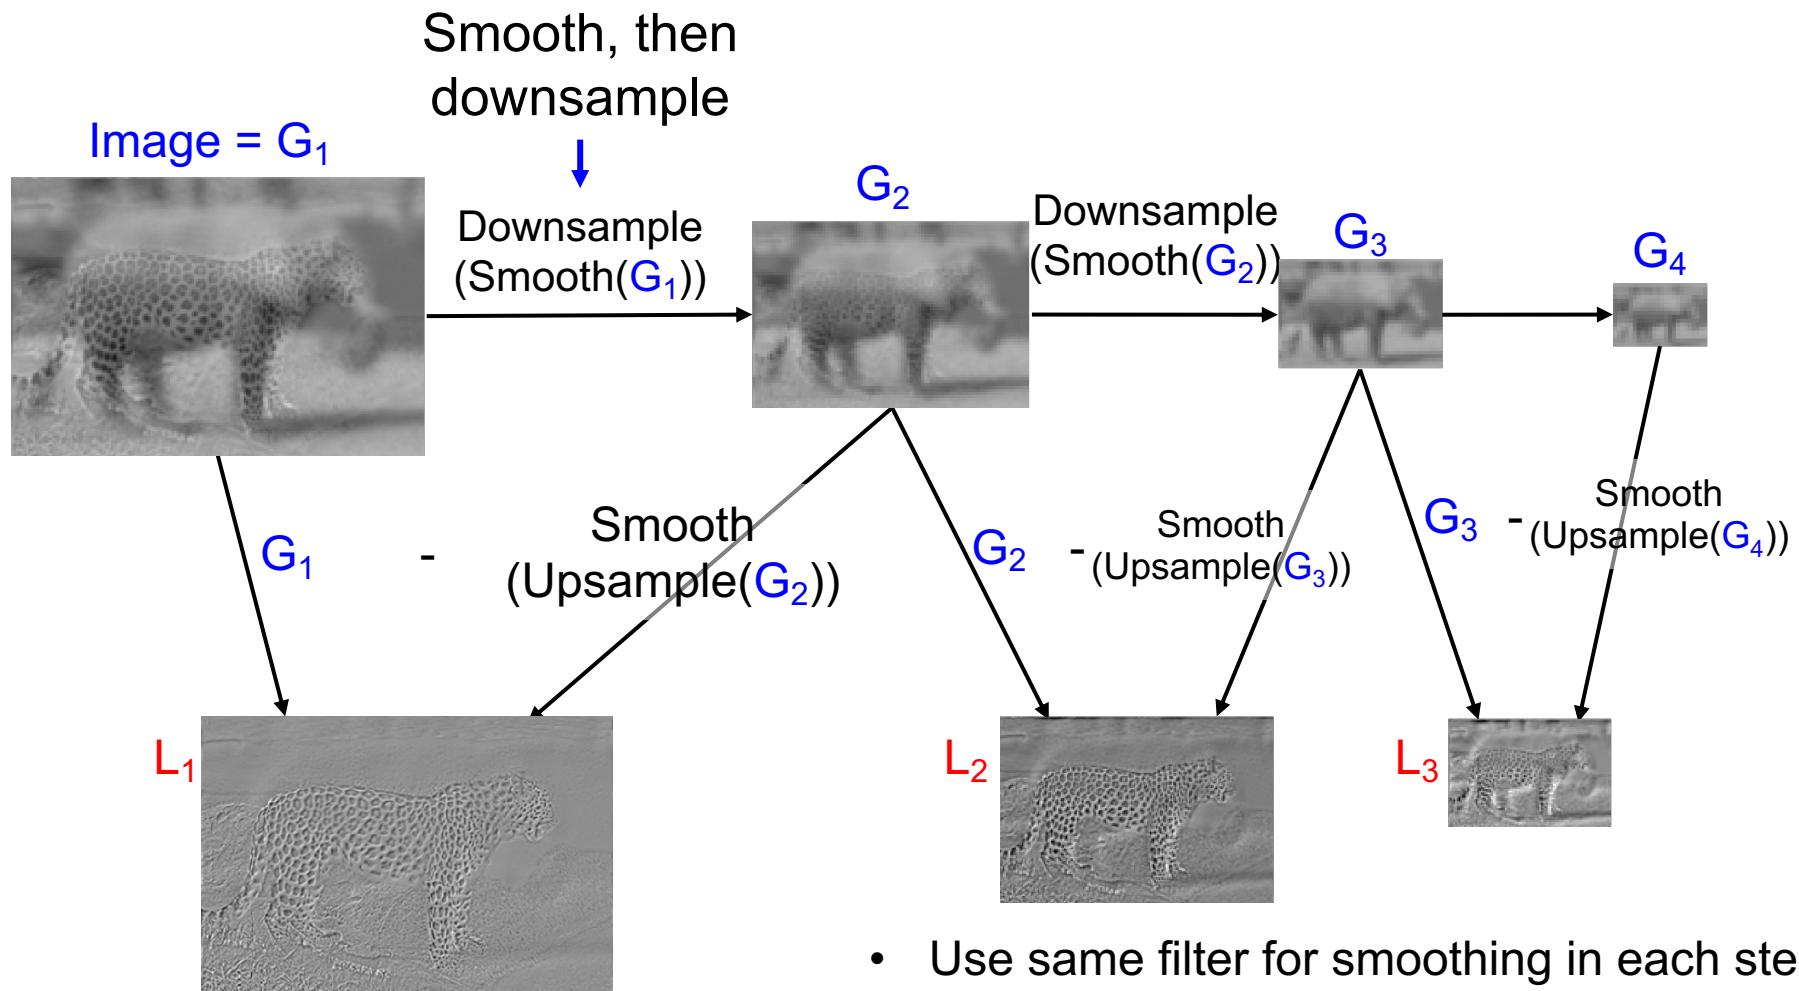
Source: Forsyth

Laplacian pyramid



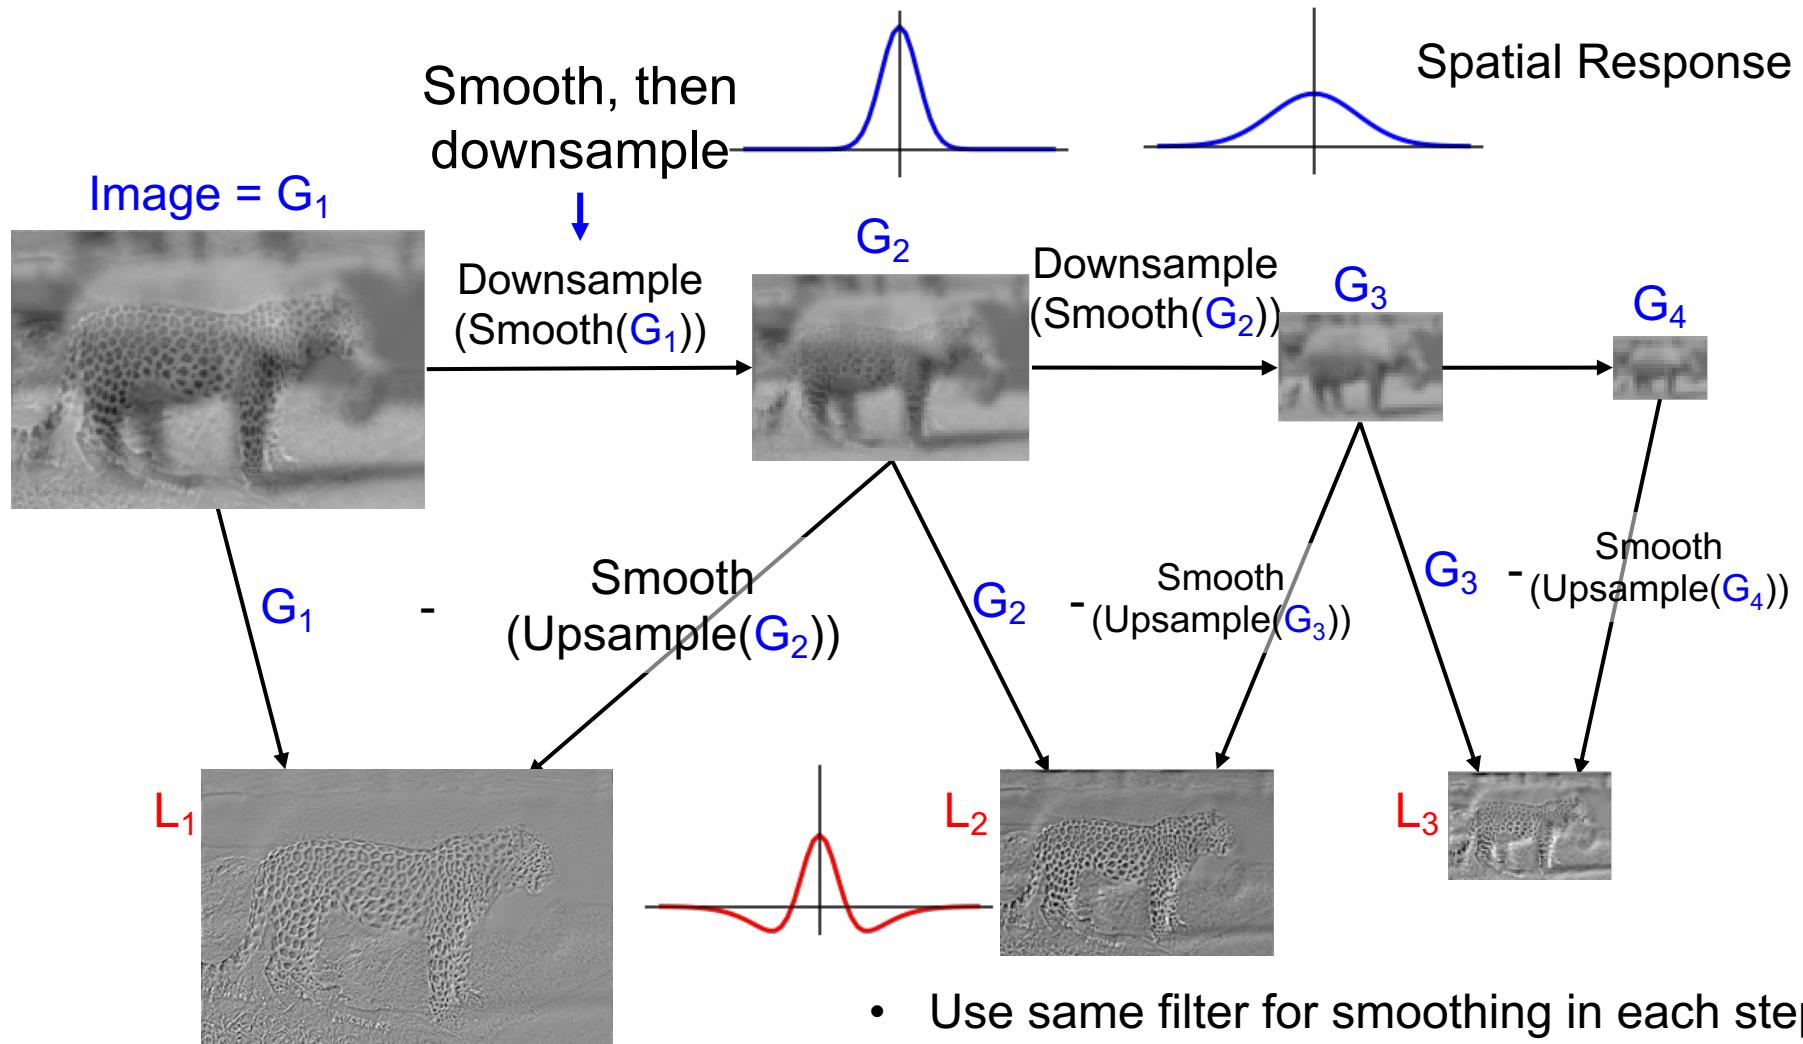
Source: Forsyth

Creating the Difference of Gaussian Pyramid



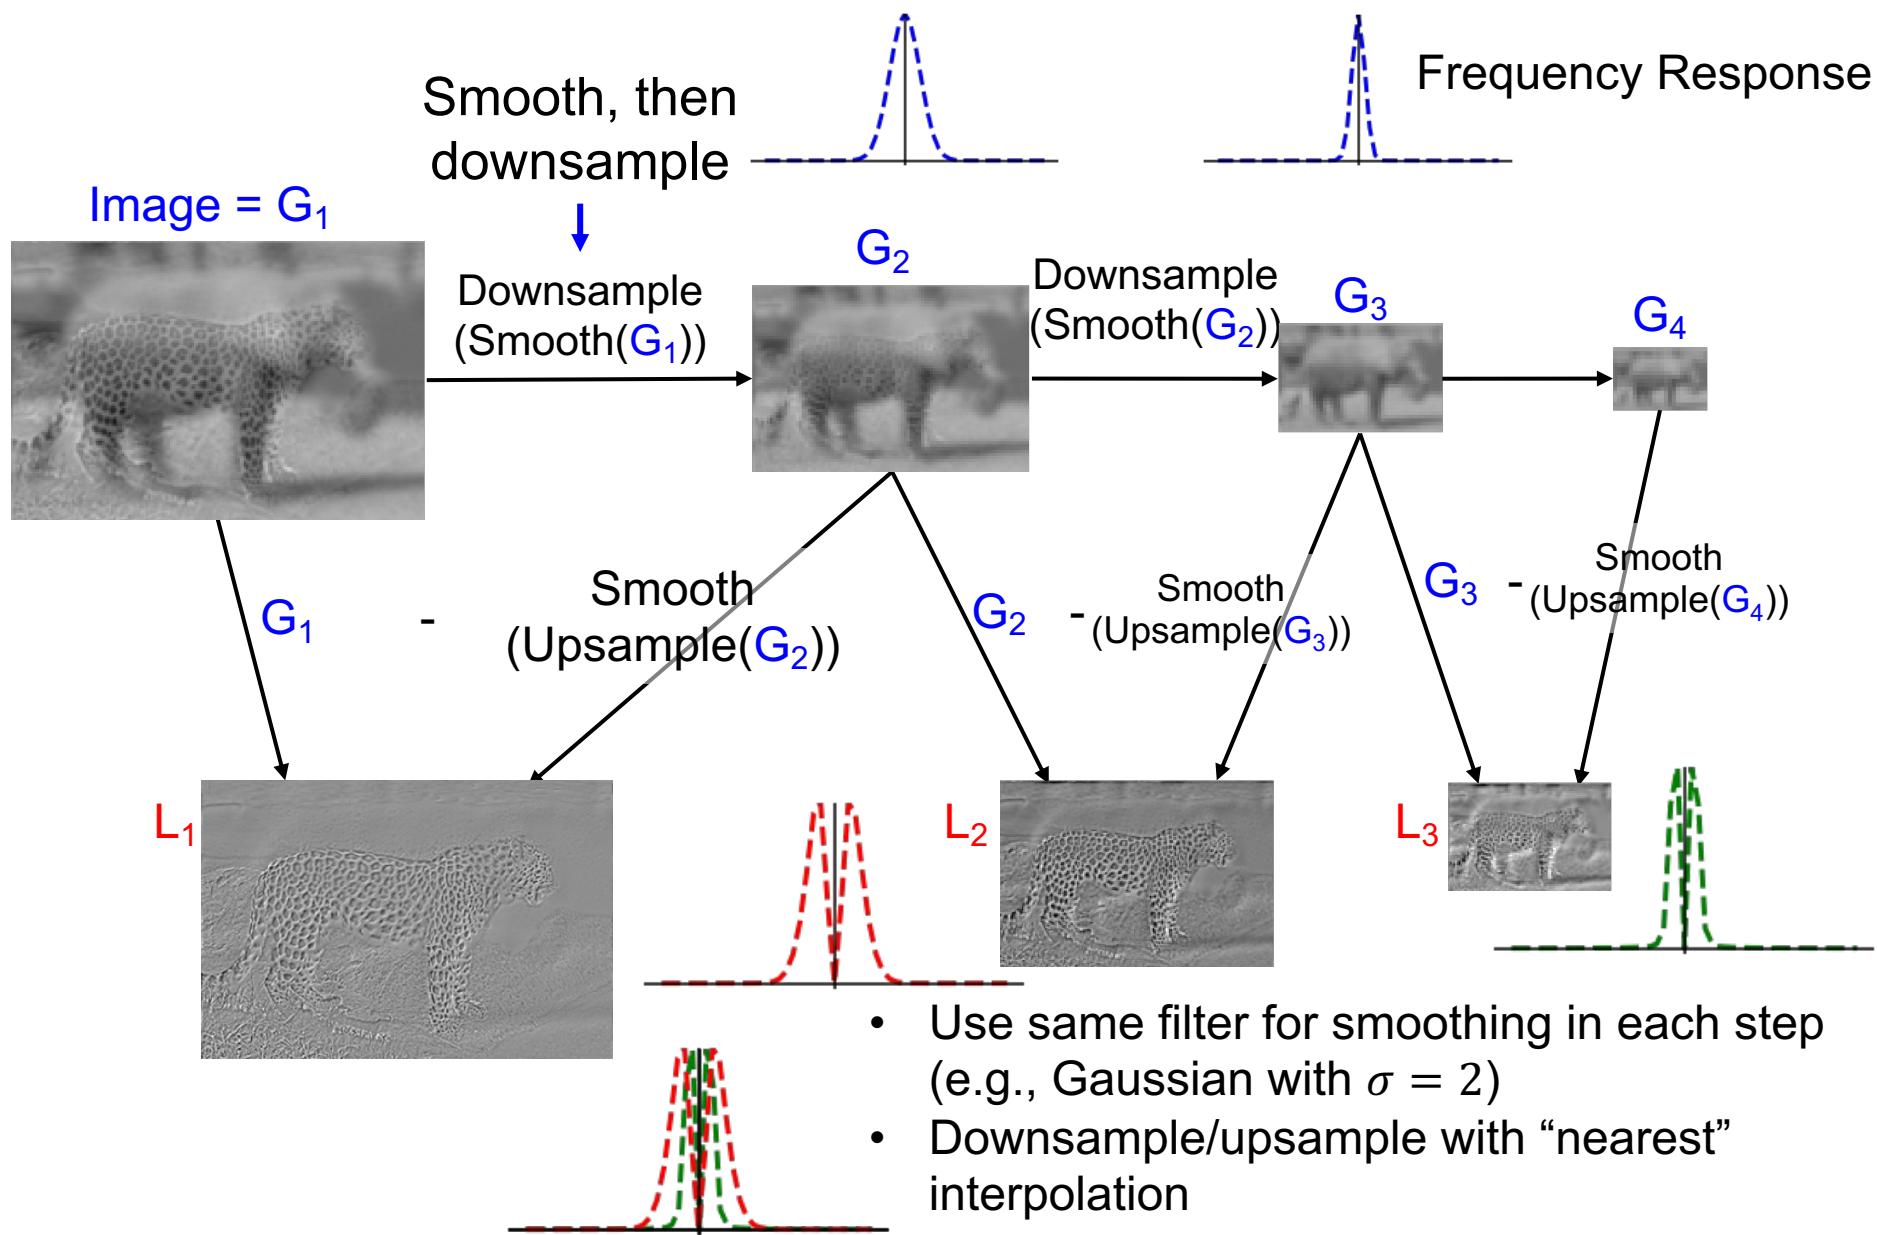
- Use same filter for smoothing in each step (e.g., Gaussian with $\sigma = 2$)
- Downsample/upsample with “nearest” interpolation

Creating the Difference of Gaussian Pyramid

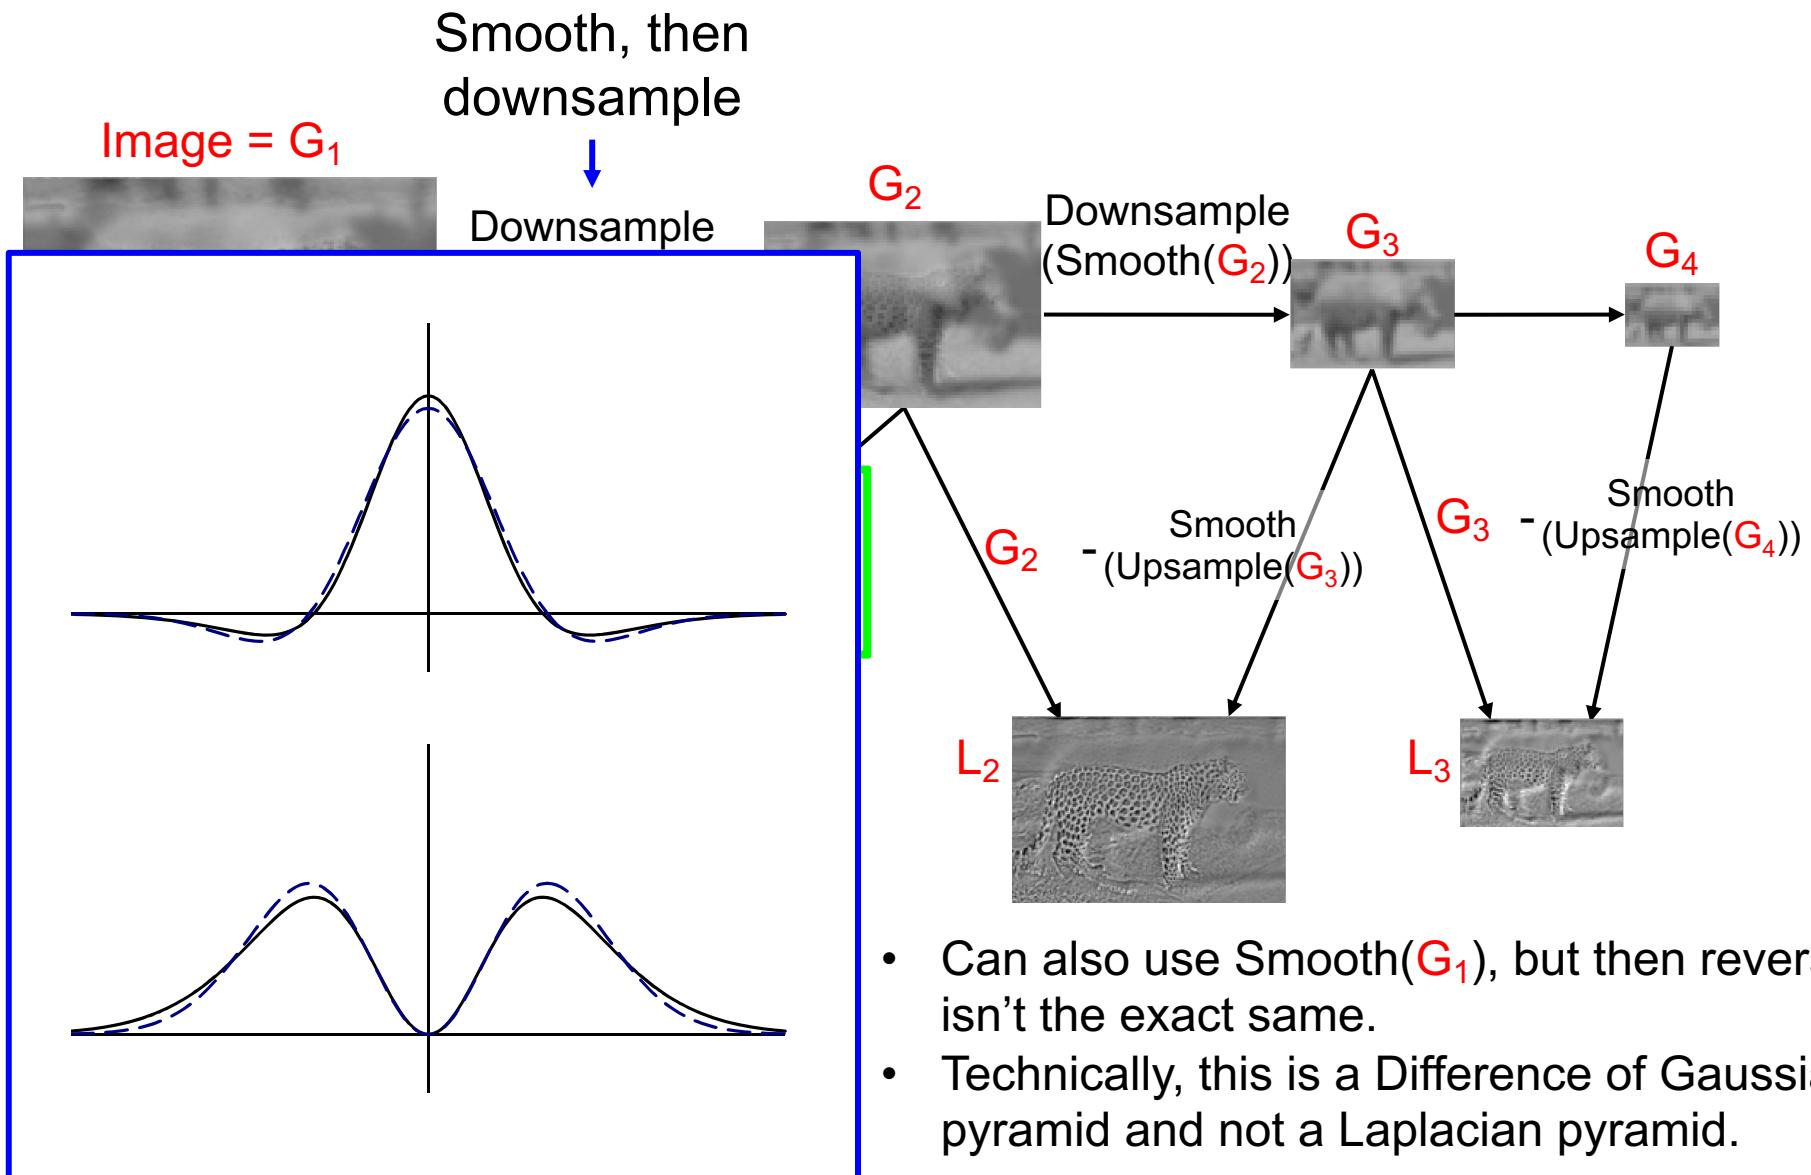


- Use same filter for smoothing in each step (e.g., Gaussian with $\sigma = 2$)
- Downsample/upsample with “nearest” interpolation

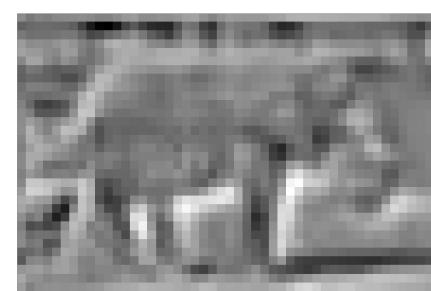
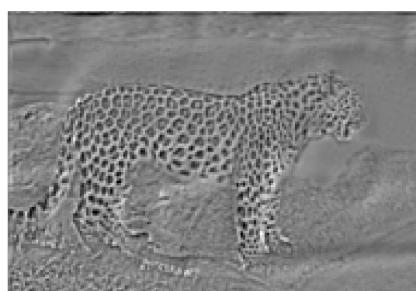
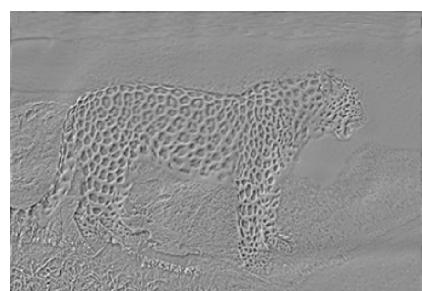
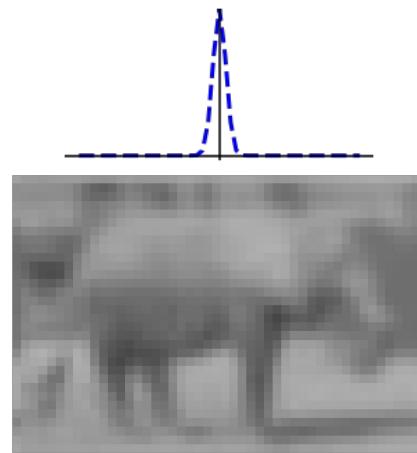
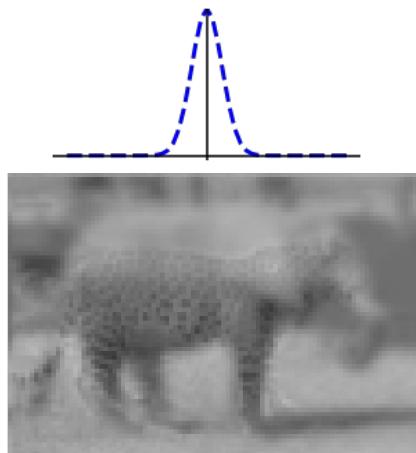
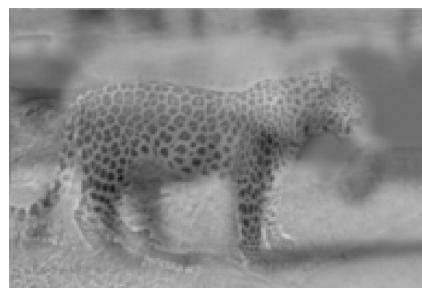
Creating the Difference of Gaussian Pyramid

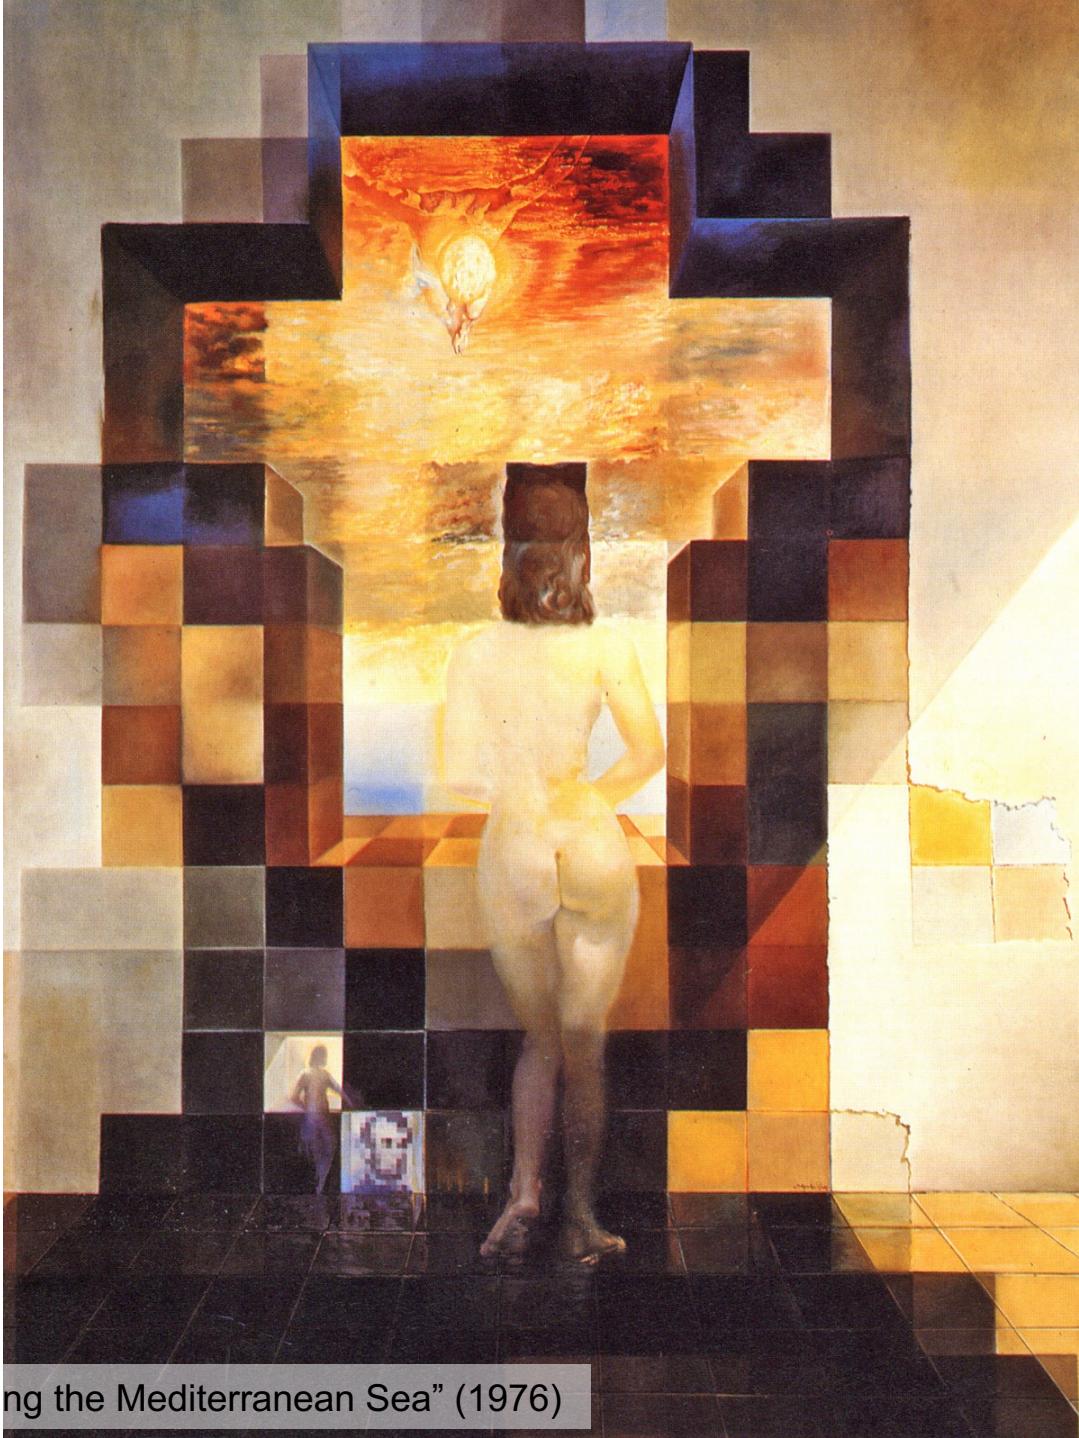


Creating the Difference of Gaussian Pyramid



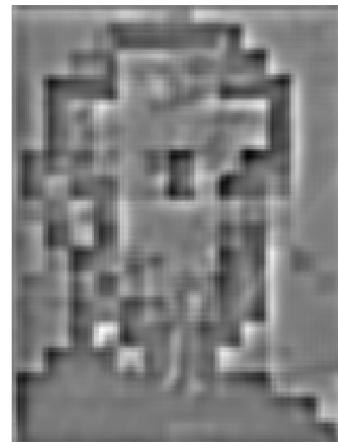
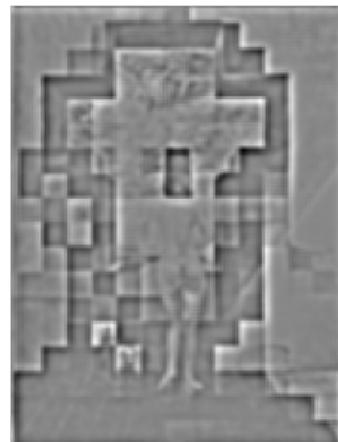
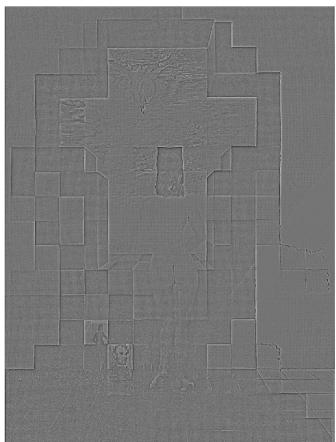
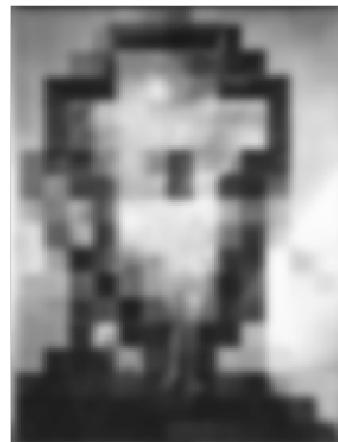
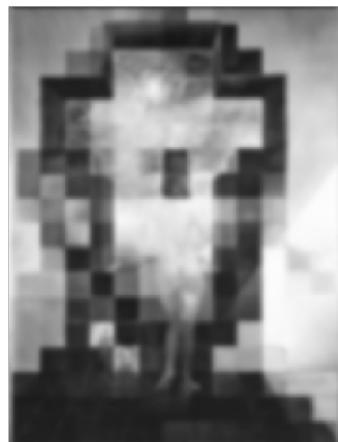
Images in a Difference of Gaussian Pyramid





Dali: "Gala Contemplating the Mediterranean Sea" (1976)

Images in a Difference of Gaussian Pyramid



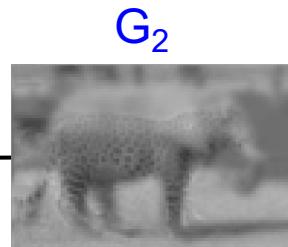
Dali: "Gala Contemplating the Mediterranean Sea" (1976)

Reconstructing from Diff of Gauss Pyramid

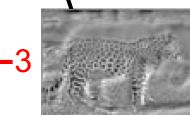
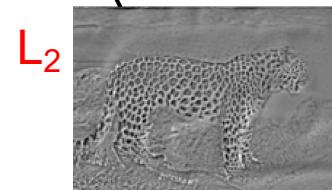
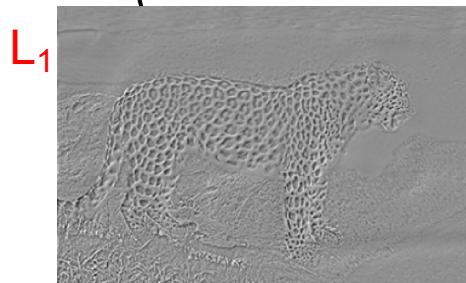
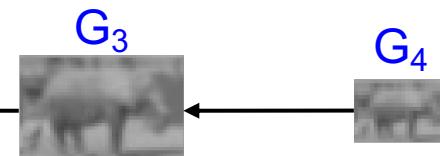
$$\begin{aligned} G_1 &= L_1 + \\ \text{Smooth}(&\text{Upsample}(G_2)) \\ \text{Image} &= G_1 \end{aligned}$$



$$\begin{aligned} G_2 &= L_2 + \\ \text{Smooth}(&\text{Upsample}(G_3)) \end{aligned}$$

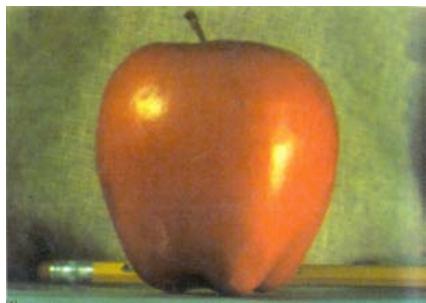


$$\begin{aligned} G_3 &= L_3 + \\ \text{Smooth}(&\text{Upsample}(G_4)) \end{aligned}$$



- Use same filter for smoothing as in deconstruction
- Upsample with “nearest” interpolation
- Reconstruction will be nearly lossless

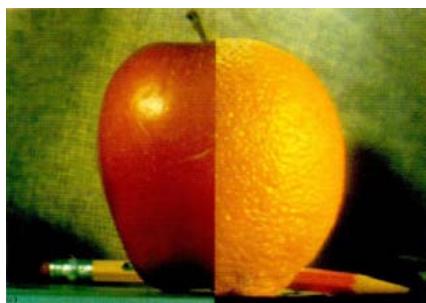
Application: Image Blending



(a)



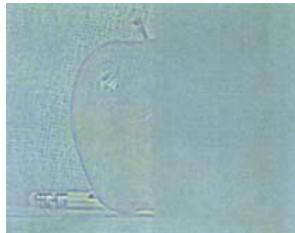
(b)



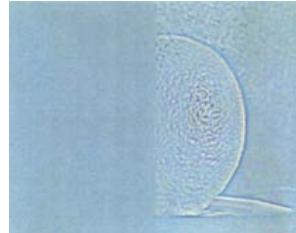
(c)



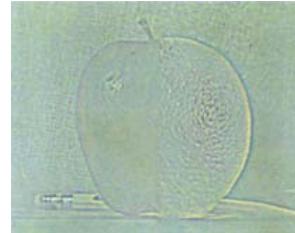
(d)



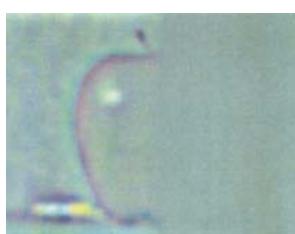
(a)



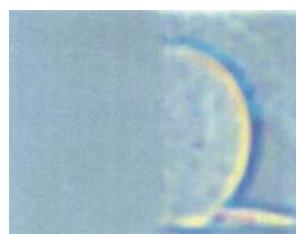
(b)



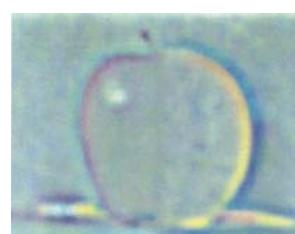
(c)



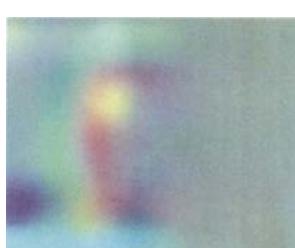
(d)



(e)



(f)



(g)



(h)

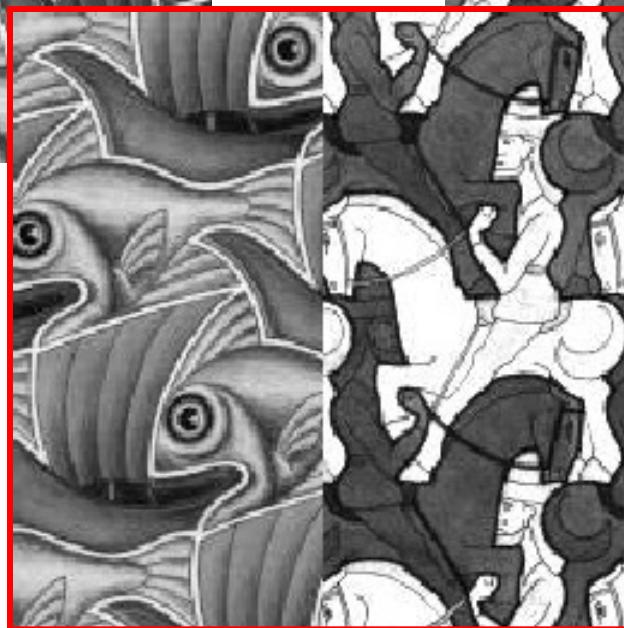
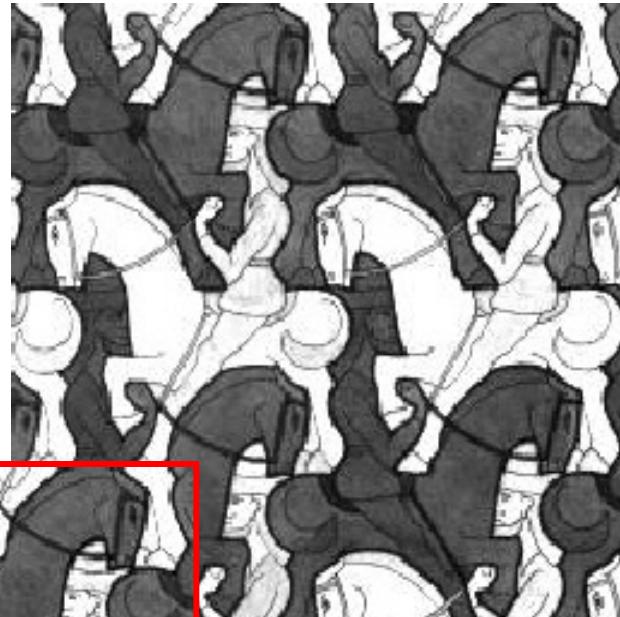
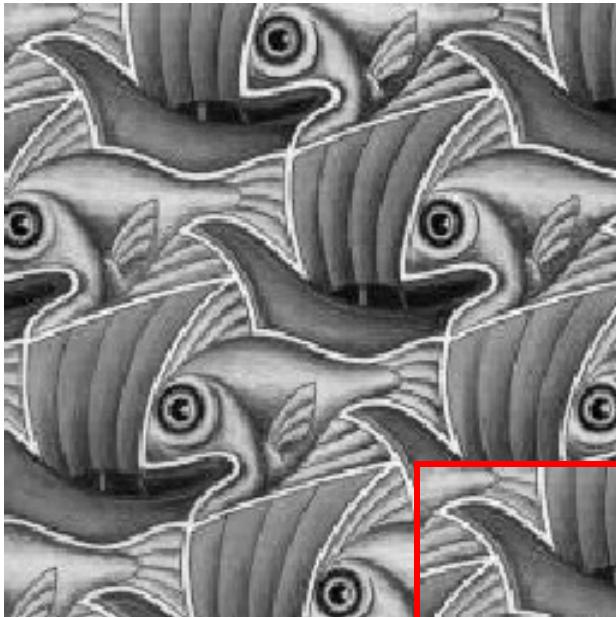


(i)



Laplacian pyramid blending (Burt and Adelson 1983b)

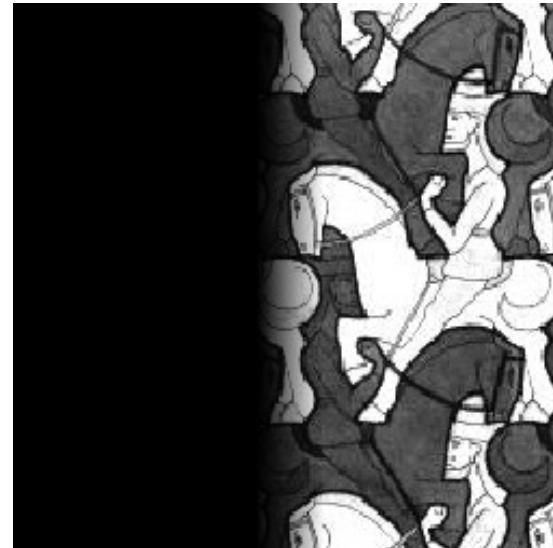
Blending



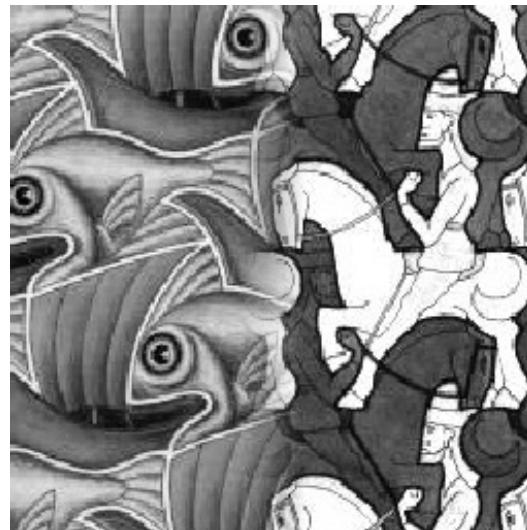
Alpha Blending / Feathering



+

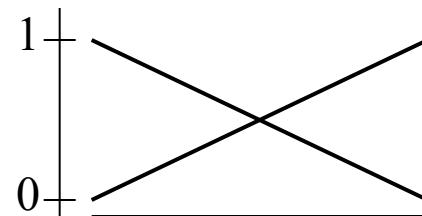
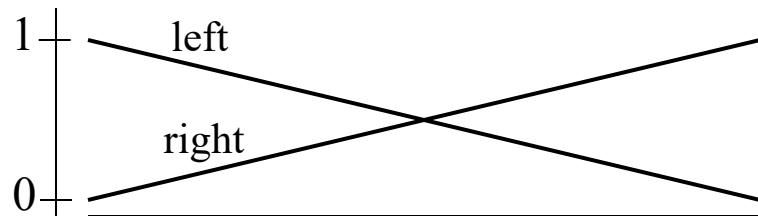
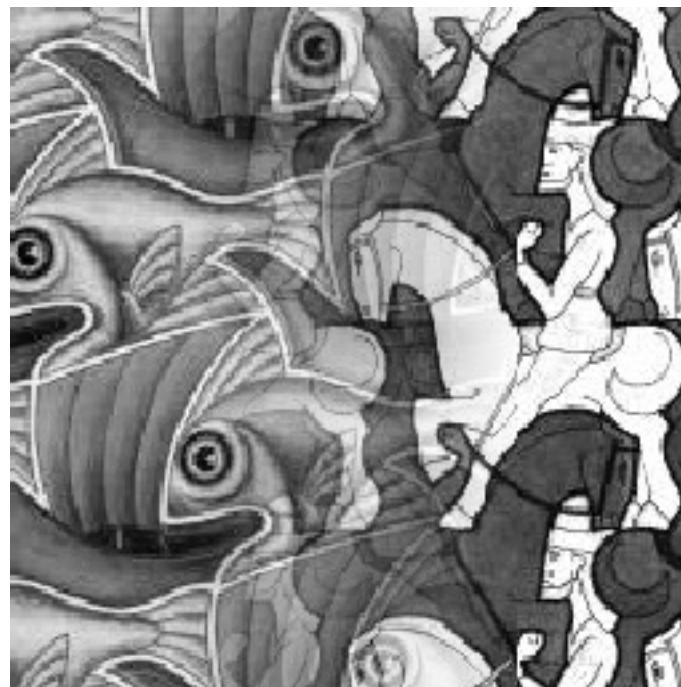
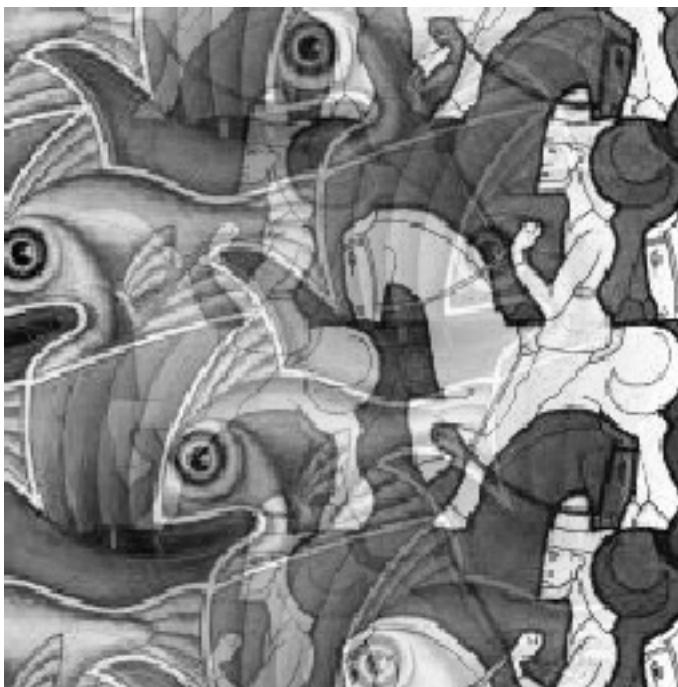


=

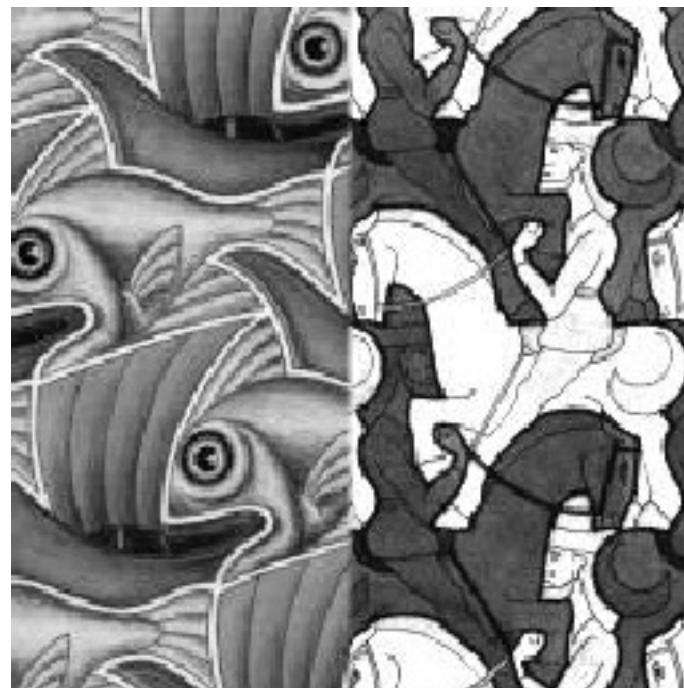
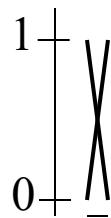
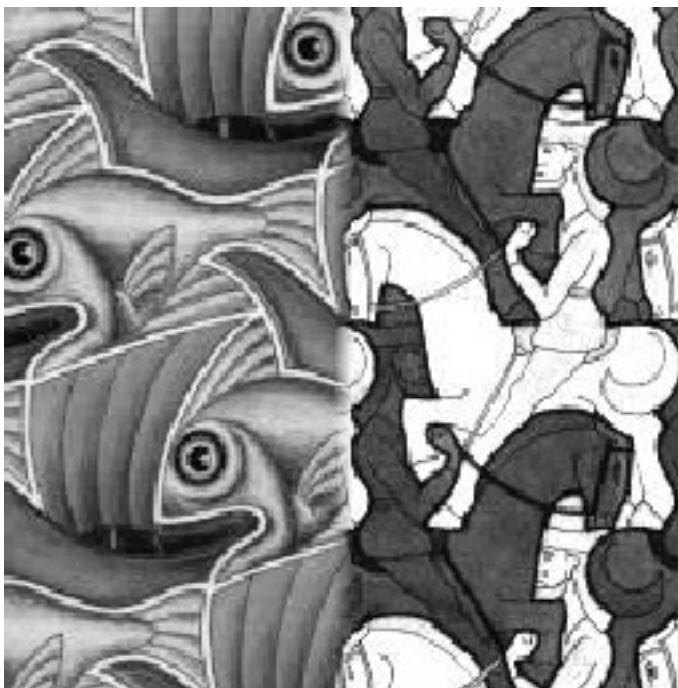


$$I_{\text{blend}} = \alpha I_{\text{left}} + (1-\alpha) I_{\text{right}}$$

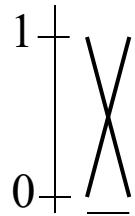
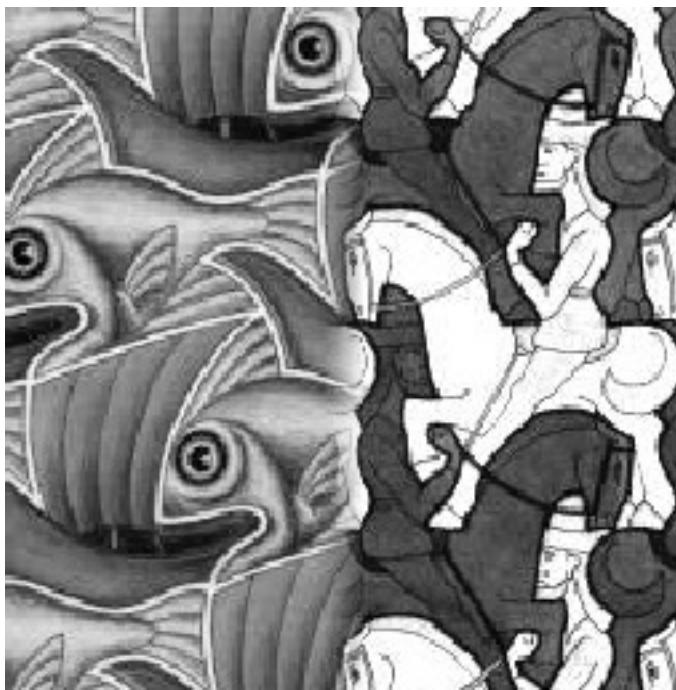
Affect of Window Size



Affect of Window Size



Good Window Size



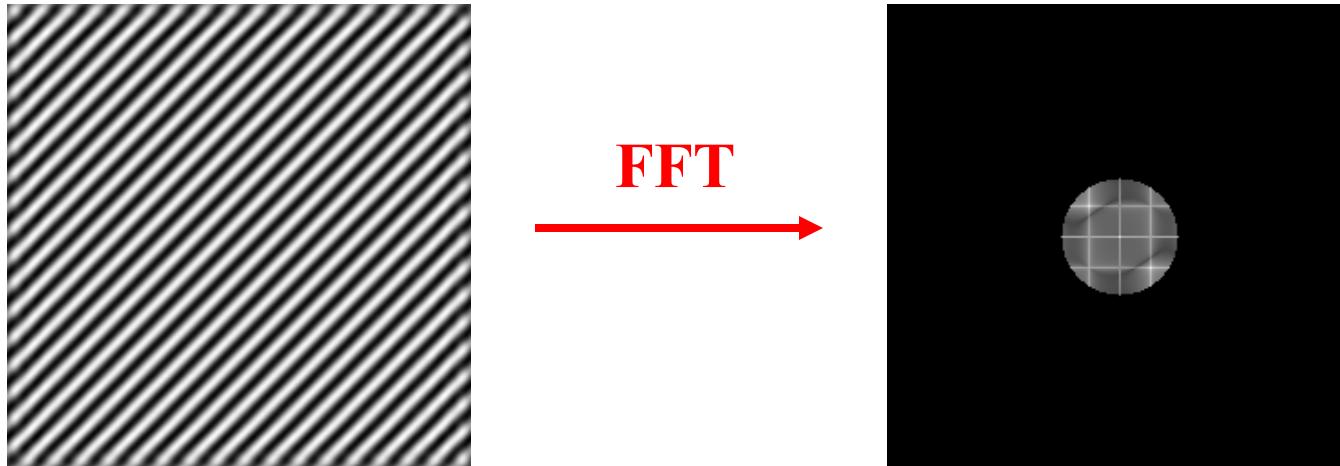
“Optimal” Window: smooth but not ghosted

What is the Optimal Window?

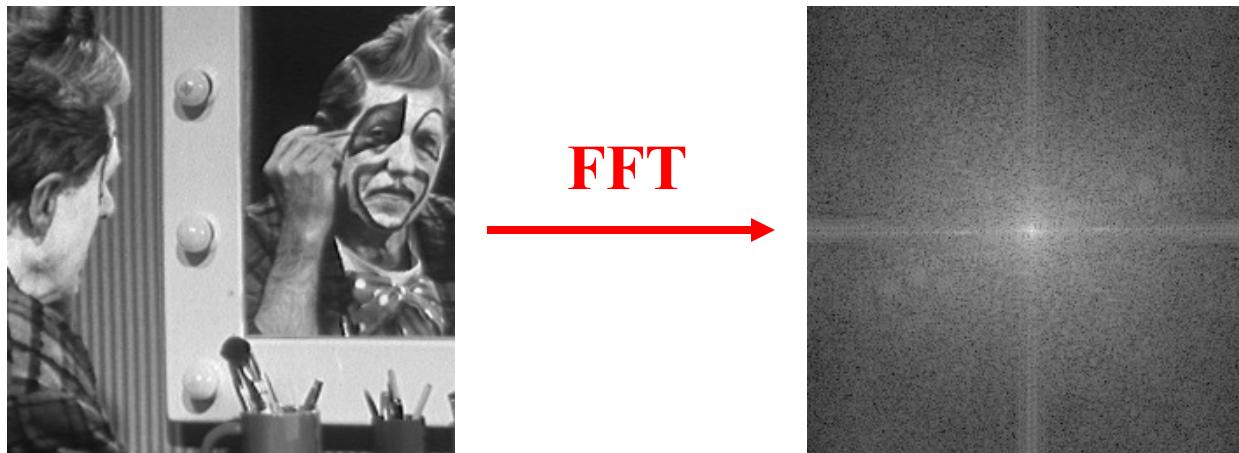
- To avoid seams
 - window = size of largest prominent feature
- To avoid ghosting
 - window $\leq 2 \times$ size of smallest prominent feature

Natural to cast this in the *Fourier domain*

- largest frequency $\leq 2 \times$ size of smallest frequency
- image frequency content should occupy one “octave” (power of two)



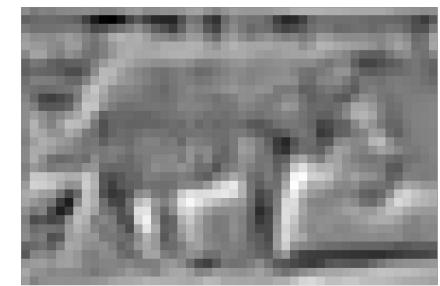
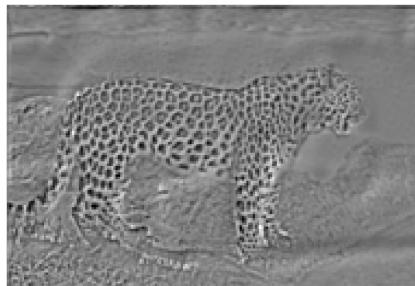
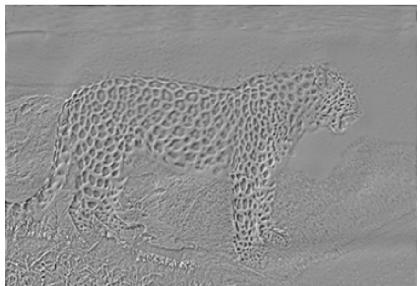
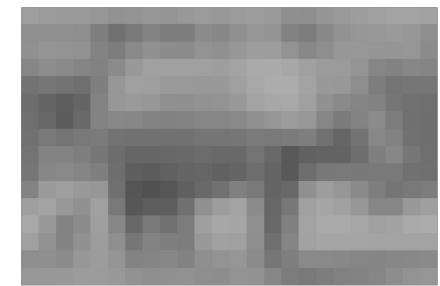
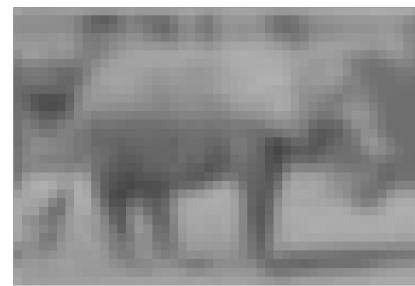
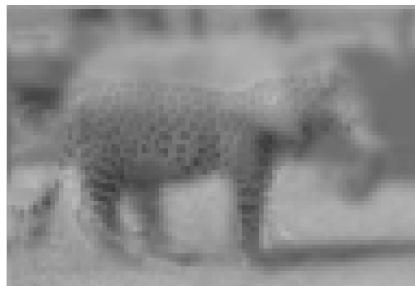
What if the Frequency Spread is Wide



- Idea (Burt and Adelson)
 - Compute $F_{\text{left}} = \text{FFT}(I_{\text{left}})$, $F_{\text{right}} = \text{FFT}(I_{\text{right}})$
 - Decompose Fourier image into octaves (bands)
 - $F_{\text{left}} = F_{\text{left}}^1 + F_{\text{left}}^2 + \dots$
 - Feather corresponding octaves F_{left}^i with F_{right}^i
 - Can compute inverse FFT and feather in spatial domain
 - Sum feathered octave images in frequency domain
 - Better implemented in *spatial domain*

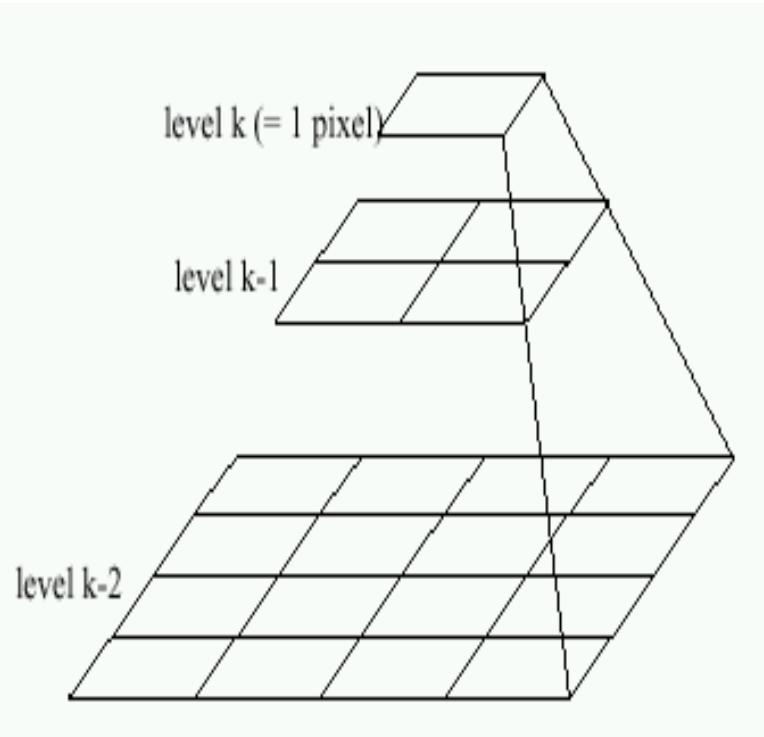
Octaves in the Spatial Domain

Lowpass Images

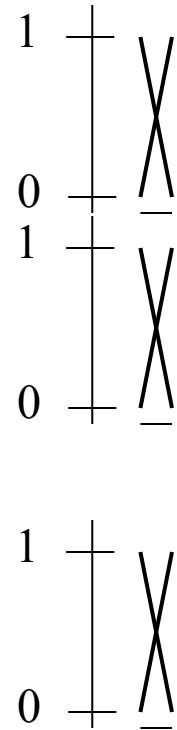


- Bandpass Images

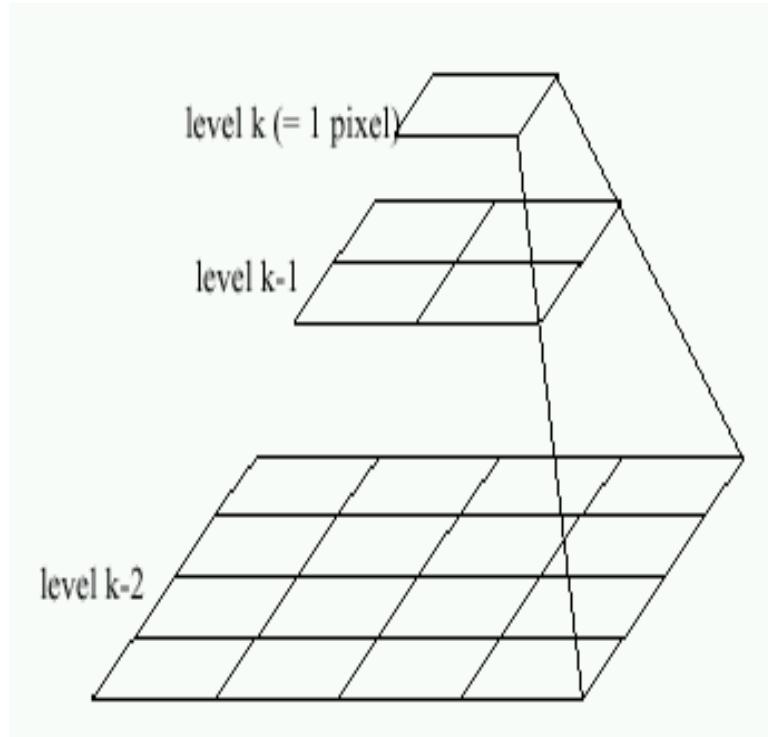
Pyramid Blending



Left pyramid

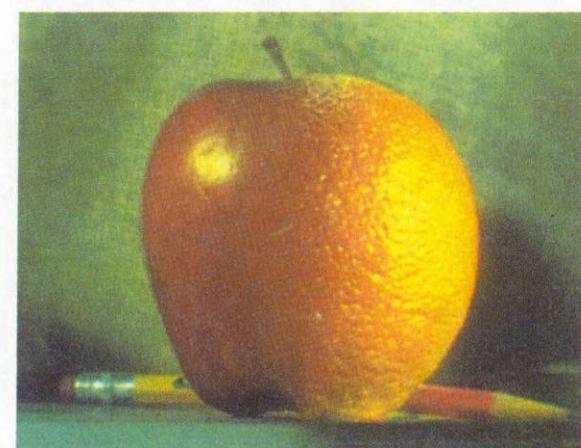
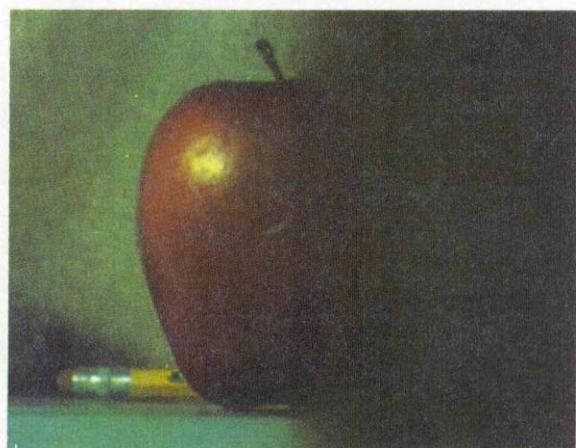
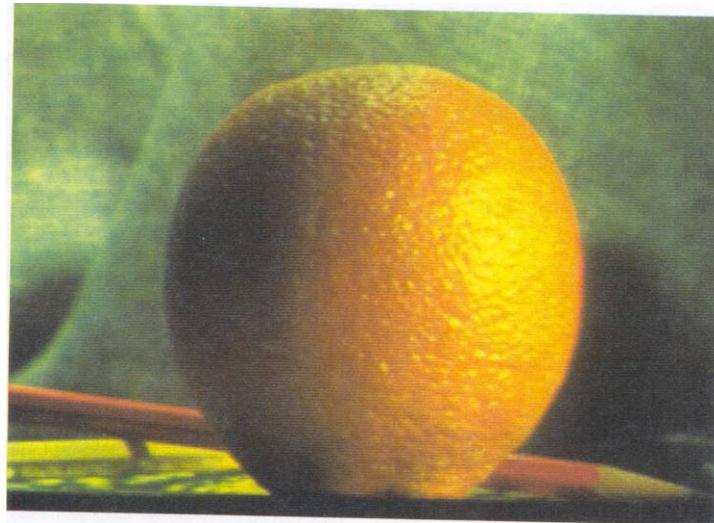
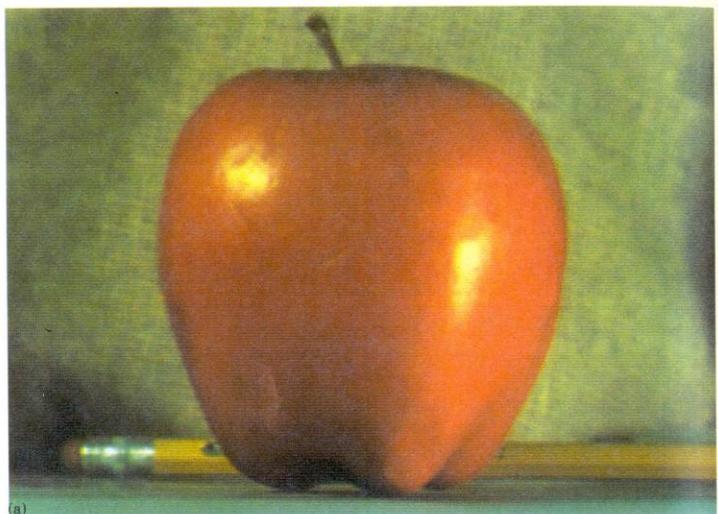


blend

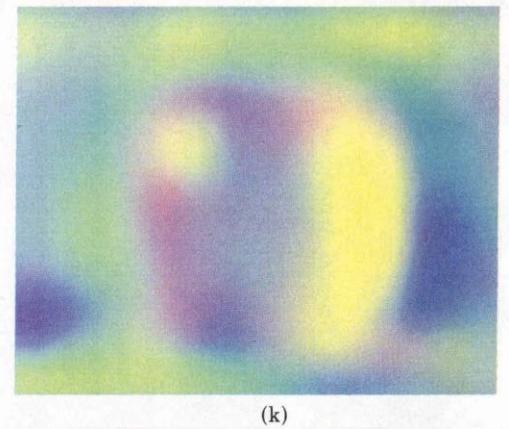
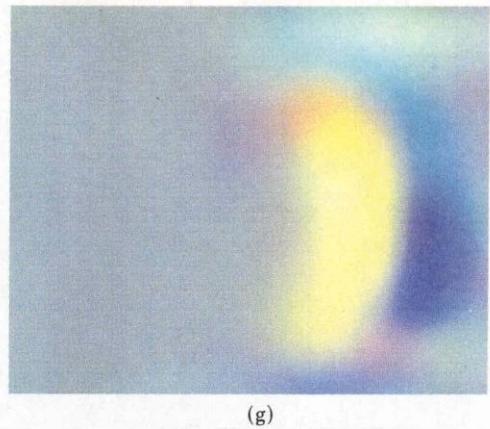
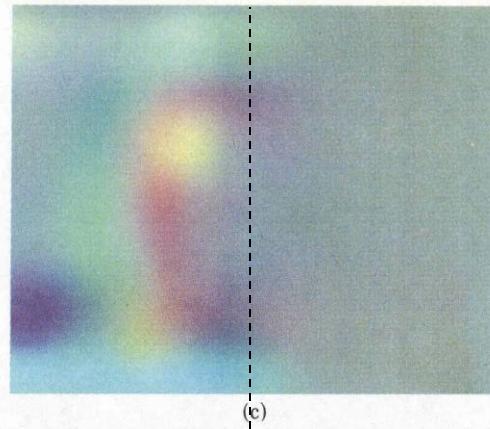


Right pyramid

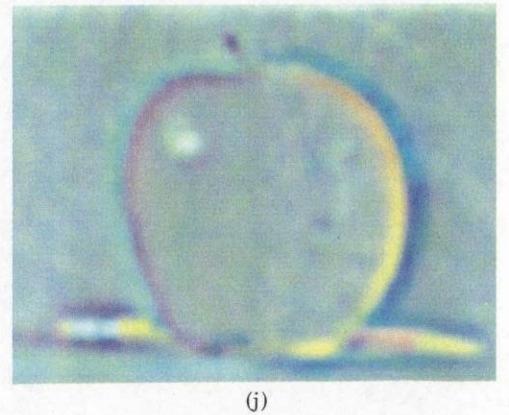
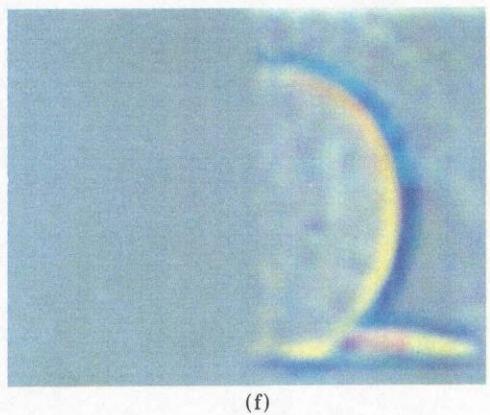
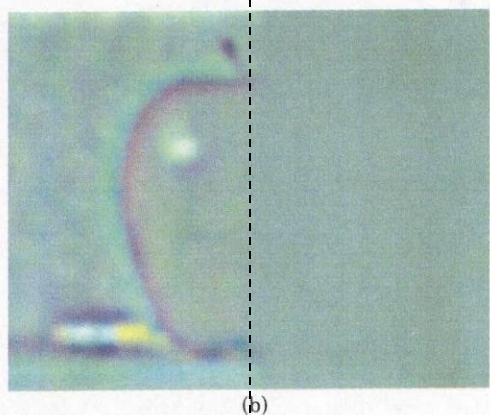
Pyramid Blending



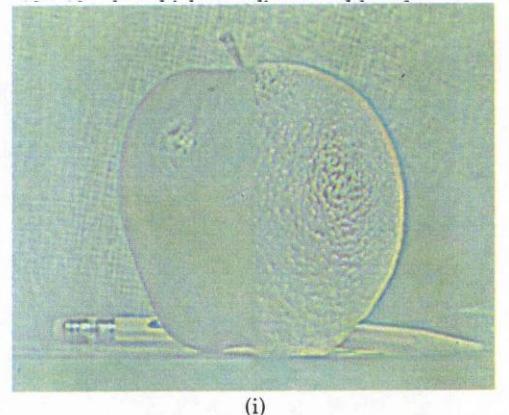
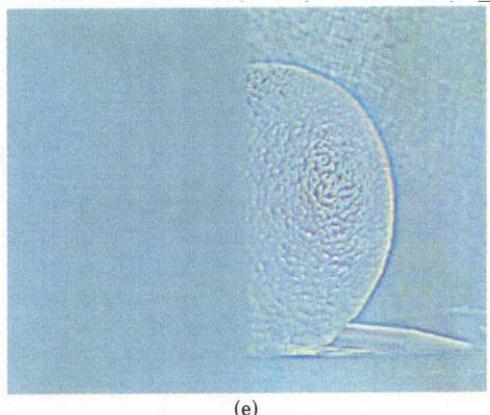
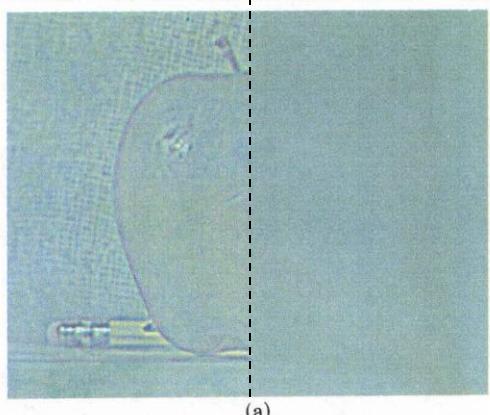
laplacian
level
4



laplacian
level
2



laplacian
level
0

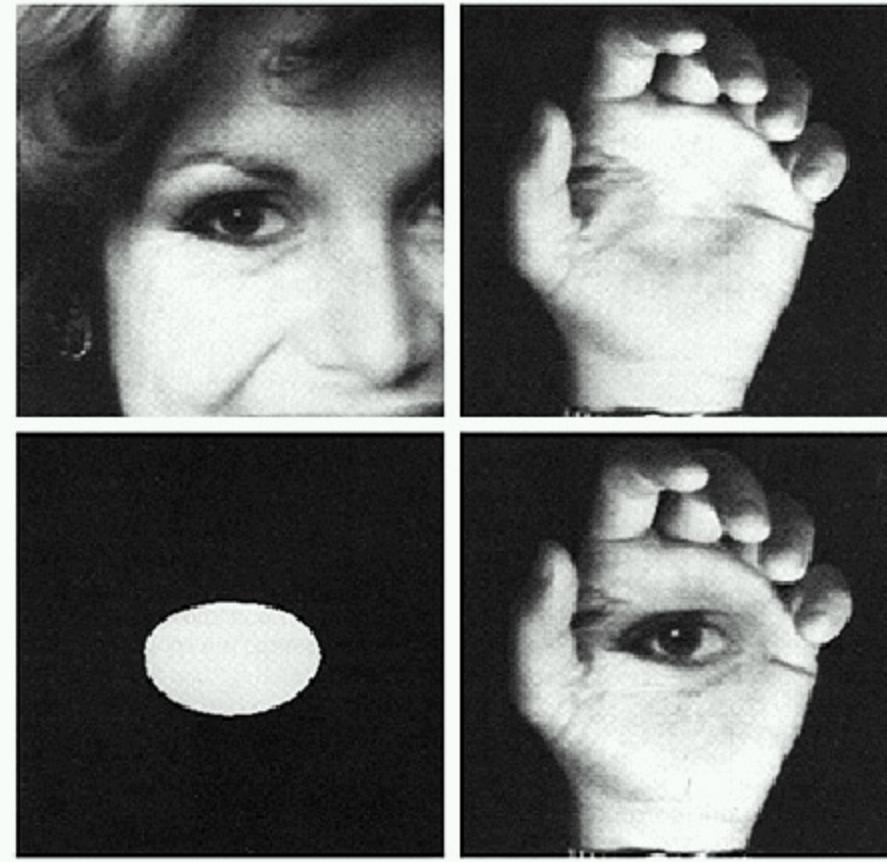


left pyramid

right pyramid

blended pyramid

Blending Regions



Laplacian Pyramid: Blending

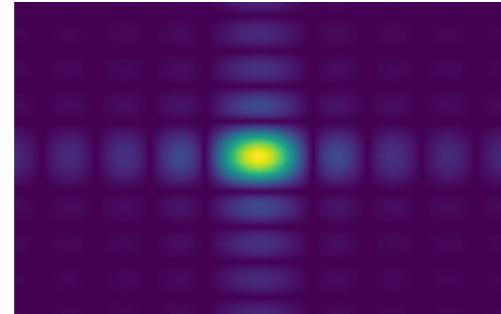
- General Approach:
 1. Build Laplacian pyramids LA and LB from images A and B
 2. Build a Gaussian pyramid GR from selected region R
 3. Form a combined pyramid LS from LA and LB using nodes of GR as weights:
 - $LS(i,j) = GR(i,j) * LA(i,j) + (1 - GR(i,j)) * LB(i,j)$
 4. Collapse the LS pyramid to get the final blended image

Major uses of image pyramids

- Compression
- Object detection
 - Scale search
 - Features
- Detecting stable interest points
- Registration
 - Course-to-fine

Recap

- Sometimes it makes sense to think of filtering in the frequency domain
 - Fourier analysis



- Sampling and Aliasing
- Image Pyramids

